

Preface: Preface

preface — Preface

Why does an institution designed to serve the public so often end up serving itself? Why does the reformer who replaces a corrupt system reproduce it — with different personnel and the same logic? Why can a society know, with considerable empirical precision, that a specific arrangement is producing mass suffering — and leave it intact for decades? Why do societies that developed independently arrive at strikingly similar solutions to the same coordination problems? Why does the political actor who thinks clearly as an individual make, in an institutional capacity, decisions they would never endorse as a person? Why do movements that succeed in winning their stated goals so often become what they originally opposed? These are not rhetorical questions. Each has a structural explanation — one that does not depend on the intentions or character of the people involved, but on the configuration of incentives, power, and material conditions within which they operate. Recognizing that these patterns are not random, and that they can be explained, requires deciding where the explanation should begin. Every framework that has managed to reorient how people think about the world began from a single displacement — a shift in where the analysis starts. Hegel displaced the analysis of history from the actions of kings and the will of God to the movement of consciousness working out its own contradictions. Marx took that displacement and grounded it further, insisting that what Hegel had described in the realm of ideas was actually happening in the realm of material production. Both moves were generative not because they were entirely right, but because they changed where the questions came from. This framework attempts another displacement of the same kind — one that follows the materialist logic further than either of them did, to the question neither asked with sufficient force: downstream of what, exactly? This framework does not claim membership in that tradition or continuity with either figure. It uses the displacement as an analytical device — a way of locating where the starting point of the analysis is and why it matters — not as a statement of intellectual genealogy. The materialist tradition identified something important: that social arrangements, political forms, and ideologies are shaped by conditions that their participants do not fully control and often do not recognize. What it underspecified is how far back those conditions go. Economic organization is one layer of material constraint. Biology is a deeper one. The analysis here begins at the deeper layer and moves upward, which means the economic conclusions it reaches are derived rather than foundational — and therefore more vulnerable to revision when conditions change, which is the appropriate relationship between doctrine and framework. What follows is a working document, not a finished system. The starting point it establishes is meant to be generative — capable of producing analysis th

at could not have been produced from any other starting point. The specific positions taken are subject to revision. The starting point is the thing worth arguing about. One clarification of language before the analysis begins. The word ‘material’ appears throughout this document, and it carries a specific meaning here that differs from its most common association in political thought. In the Marxist tradition, material conditions refers primarily to the economic relations of production — to who owns what, who works for whom, how the products of labor are distributed. That is a subset of what this framework means by material. Here, material refers to the full range of physical and biological conditions that constrain and shape what is possible for organisms and the systems they build: the structure of the nervous system, the biology of pain and pleasure, the ecology of resource scarcity, the epidemiology of disease, the architecture of incentive structures. Economic relations are one layer of this. They are not the foundation. The biological conditions that make economic relations possible and that constrain what forms they can take — that is the foundation this framework begins from. A distinction that runs through everything that follows: the framework is the methodology, and the doctrine is its current application to existing conditions. The framework aspires to the durability of a genuine reorientation. The doctrine is conditional, temporal, and expected to change as material conditions change. What is communicated publicly is primarily the doctrine. The framework

operates behind it, available to those who engage seriously enough to encounter it. Part XI illustrates what this distinction looks like in practice in the domain where it is most frequently confused: economics. The framework provides diagnostic questions that apply across arrangements and historical periods. The doctrine selects instruments based on specific material conditions. Neither replaces the other. What is offered here is not an endpoint. It is a position in a longer movement — more adequate than what preceded it, less adequate than what will follow. The measure of its success is not whether it survives intact but whether it generates something more capable than itself that renders it obsolete. A framework that defends its own permanence has already begun to betray its foundations. The same principle applies to every claim within it. Wrong analysis should be corrected; incomplete analysis extended; a more adequate framework, when it arrives, embraced rather than resisted. The framework that makes itself permanent has already failed its own standard. The one that cannot be surpassed is not the strongest — it is the one that has stopped moving. This follows as a logical consequence from the epistemological and biological foundations, not as a performance of humility. An organism that cannot adapt dies. A framework that cannot be surpassed has made itself into exactly the kind

of static arrangement that this analysis identifies as a primary source of structural suffering.

how-to-read-this-document — How to Read This Document

This framework covers a large amount of terrain. The sequence is deliberate, but it is not a simple linear argument in which each part follows from the previous and adds one new element. Several concepts recur across multiple parts — not because the structure is poorly organized, but because the same concepts must be shown operating at different levels of analysis. The map below describes the sequence and explains what each level contributes. Eleven levels of engagement are available below. They are cumulative rather than exclusive: a reader who enters at Level 3 will have missed the foundations of Levels 1 and 2, but the framework is designed so that each level is internally coherent enough to be legible on its own terms, with references back to what it presupposes. A note on the visual conventions used throughout the document follows the level descriptions. Level 1 — Parts I and II: How we know, and what exists. The epistemological foundation (Part I) and the ontological foundation (Part II) establish the conditions under which the rest of the analysis is possible. Part I asks: what makes claims about reality valid, and what makes them revisable? Part II asks: what kind of thing is reality, and does it have a fundamental level? These two foundations determine what kind of claims the framework can make and what kind it cannot. Every subsequent part inherits their constraints. If Part I is wrong — if knowledge cannot be revised by evidence — then the rest of the framework's claims to be empirically grounded are empty. If Part II is wrong — if reality is constituted by consciousness rather than primary — then the biological foundation of Part III is derivative rather than foundational. Level 2 — Part III: What we are. The biological foundation establishes the specific starting point for evaluation: organisms that can suffer, and the consequences of that fact for everything they build and organize. Section 3.0 develops the conditional normative framework — clarifying why PMN's use of biology does not commit the naturalistic fallacy, does not reduce all analysis to biology, and does not smuggle in teleological assumptions about development. Section 3.0 also establishes the two-level architecture that runs through everything that follows: life is the ontological foundation (what exists), suffering is the evaluative criterion (what counts as bad and what counts as wrong when it is unnecessary), and becoming is shorthand for the anti-foreclosure criterion: what the floor criterion requires when its causal implications are traced temporally, namely that institutional arrangements not foreclose the adaptive response capacity of future populations. The biological floor is a floor, not a ceiling: it establishes the minimum constraint, not the full content of what the framework values. Section 3.3 now explicitly bridges the is-ought gap for the anti-foreclosure criterion as well as the floor: existential suffering (the Last Man condition) meets the four cri

teria of unnecessary suffering and generates physiological and social harm. Section 4.5d (Becoming as Epistemic Capacity) specifies that sustained collective learning about novel threats is a structural component of becoming, not an optional addition to the minimal anchor. Level 3 — Parts IV and V: What matters and what it means. The ethics (Part IV) and the metaphysical dimension (Part V) develop the evaluative commitment established in Part III into a fuller account of value and meaning. Part IV specifies the minimal anchor — the reduction of unnecessary suffering — and the aspirational horizon that it points toward: genuine development, not merely the absence of pain. Part V addresses the dimension of meaning that the materialist tradition has consistently handled badly: religion, consciousness, and the interface between material reality and subjective experience. Both parts involve concepts that recur throughout the framework, because ethics and meaning are not confined to a single analytical level — they appear wherever the analysis asks what the purpose of a given arrangement is. Level 4 — Parts VI and VII: How power operates, and what institutions do. Parts VI and VII apply the framework to political reality. Power (Part VI) cannot be left implicit: every claim about what arrangements are better or worse involves claims about who can do what to whom. Part VII develops the institutional analysis, including the custodian problem (7.3), capture mechanisms (7.3b–c), and cosmological capture (7.3c-ii): the failure mode in which actors captured by belief frameworks cannot experience themselves as captured. The anti-capture design section (7.3d) specifies the institutional mechanisms for resisting the capture sequence. The cost-asymmetry analysis provides the material complement to the cosmological capture diagnostic. Level 5 — Parts VIII and IX: Society, culture, and the world. Parts VIII and IX extend the analysis to the domains where structural conditions are most directly experienced: the organization of social life (Part VIII) and the geopolitical arrangements that shape what any given society can do (Part IX). The concepts from earlier levels — material analysis, evaluative criterion, power dynamics, institutional accountability — are applied here to questions of culture, identity, gender, colonial inheritance, and the material pressures that drive convergence and divergence between societies. Section 9.3b grounds the multipolarity analysis in the specific mechanisms of the current transition: semiconductor access controls, de-dollarization dynamics, and the emergence of critical minerals as a new constraint layer — read as material variables, not ideological positions. Level 6 — Parts X and XI: How systems change, and what economics means. Part X analyzes the dynamics of change: how stable arrangements become unstable, how change propagates non-linearly, how counter-power accumulates, and how windows of political opportunity open and close

. New sections include: beliefs as intervening variable between conditions and behavior (3.8b extension), vertical versus horizontal mobilization diagnostic (10.5b), and the reform-versus-revolution decision procedure (10.6). Section 10.7 provides the systematic engagement with the Gramscian tradition. Sections 10.10c-i and 10.10c-ii extend the agent typology and develops the watcher problem. Part XI applies the framework to economic doctrine. The accountability and contestability diagnostic is now explicit as the eighth economic diagnostic (11.0), alongside 11.5 (power beyond ownership: expertise, information, and network effects). Level 7 — Parts XII and XIII: How the framework operates, and what it cannot resolve. Part XII addresses the analytical methodology in practice, including the misuse resistance architecture (12.5d), the technocratic drift safeguard (12.8), and the comparative demonstration (12.9). Part XIII addresses the permanent tensions that cannot be resolved but must be managed, including the moral-functional tension (13.4h): the conditions under which dismantling an institution that fails moral evaluation destroys latent coordination infrastructure, and the grounded realism requirement that applies equally to transformation and preservation programs. Level 8 — Part XV: The formula architecture (Structural Tendency Analysis — STA). Part XV makes explicit the variable architecture implied by the preceding analysis. Section 15.0b provides canonical term definitions for all key variables (S, R, B, V, Becoming, Minimal Anchor, Layer/Level). The formula set is decomposed through domain-specific applications and applied to historical cases (15.14). Section 15.15 provides the compressed core: the three-idea minimum viable version for non-expert contexts. Read 15.13 (what the formulas cannot do) before applying any formula. Level 9 — Part XVI: Technology as structural variable. Part XVI treats technology

not as a tool but as a constraint-modifier: it changes which institutional arrangements are stable, which actors can coordinate, and at what cost. It covers the capture problem at technological leverage points, AI as institutional architecture (including the analysis of algorithmic governance as advanced V — symbolic violence via the ideology of objectivity), biotechnology's implications for the biological floor, and the acceleration problem in meaning infrastructure. Section 16.4b includes the falsification trigger for effective accelerationism: the framework's evaluative criterion applies to present populations in present conditions, not to projected future endpoints — a constraint that temporal displacement arguments must satisfy before any generational sacrifice argument is valid. Readers who entered through the formula architecture (Part XV) will find that Part XVI provides the empirical content for several of the technology variables in the formula set. Level 10 — Part XVII: Cases and the individual. Part XVII applies the framework

at two scales simultaneously: historical cases that illustrate structural patterns (the Prohibition failure, labour movement capture, land reform without implementation infrastructure, revolutionary custodian problem), and the individual-level demands of holding the framework honestly — inside-outside decisions, self-drift detection, meaning sustainability, and communication across frameworks. This Part is where the analytical and the practical converge. It can be read after any other Part, but it is most useful after Levels 4–7 have established the structural concepts the cases instantiate. Level 11 — Part XIV: Framework, doctrine, and what remains open. Part XIV is the living summary — positioned last because it must be read last to be read honestly. It synthesizes the analytical commitments of Parts I–XIII with the operationalization of Parts XV–XVII, states the current doctrinal positions across all major domains (14.2), distinguishes what the framework has argued from what it assumes and what it does not know (14.3), specifies the conditions under which it would be surpassed rather than merely revised (14.4), and addresses the institutional requirements for transmitting the framework without capturing it (14.5–14.6). Part XIV is updated as the framework develops: it is not a conclusion but a current position. The three-tier axiom hierarchy in 14.3 (Foundational/Structural/Empirical-hypothetical) is the most condensed statement of the framework's epistemological architecture. Read 14.4b before applying the framework to any situation where you suspect the framework itself may be the problem. How to read the visual elements Blue passages. Passages with a blue left border and light blue background mark theoretical distinctions and structural claims — things that are analytically important but can be returned to. They elaborate the argument; they are not the argument itself. Red passages. Passages with a red left border and light red background are operational diagnostics: numbered questions for applying the framework to a specific case or situation. They are most useful when you are working with the framework on real material, not on a first read. Bold inline labels within paragraphs. Bold labels at the start of a body paragraph (e.g., “The productive response is not a reason to permit suffering:”) mark a named concept, phenomenon, or structural claim being defined. Non-bold labels (e.g., “The practical consequence:”) mark derivations, applications, or consequences that follow from the analysis. The distinction is intentional throughout. Why concepts recur: epistemology reappears in Part VIII because the framework's standards for knowledge apply to cultural analysis as much as to any other domain. Ethics reappears in Part X because the question of what change is for cannot be separated from the analysis of how change happens. Power reappears in Parts VI and X because it is a variable that cannot be analyzed once and then left as settled. The re

currences are not repetitions. They are applications of the same tools to different levels of analysis. A reader who finds a concept reappearing should ask: what does it look like at this level that it did not look like at the previous one? Six Entry Points by Reader Intent If you want to understand the foundation before anything else Read in order: Part I → Part III (3.1–3.4b) → Part IV (4.1–4.5b). This gives you the epistemological commitment, the biological foundation, and the evaluative criterion. Everything else in the document is an application or extension of what these three sequences establish. The Preface provides the intellectual genealogy; it is not required for the argument but clarifies the direction of departure from prior frameworks. If you want to be able to use the framework operationally After the foundation sequence above, read: Part VI (Power) → Part X (How Systems Change, especially 10.10c for the agent typology) → Part XII (How This

Framework Operates) → Part XVII (Cases and the Individual). Part XVII's individual-level sections (17.7b–17.7e) are the operational complement to the structural analysis: they address what to do with the framework inside the situations it describes. Part X gives you the mechanism of change. Part XII gives you the methodology for applying all of it honestly — including 12.5, which identifies the specific ways this framework can be used badly, and 12.1, which gives you the four stress-testing procedures. Do not skip 12.5, and read 12.5d immediately after it: 12.5 depends on the analyst catching their own failure, and 12.5d does not. If you are analyzing a specific political or social situation Part VI → Part VII (7.0–7.3c, including cosmological capture at 7.3c) → Part X (especially 10.5, 10.7, 10.10c) → Part XV (read 15.13 first, then 15.0–15.3, then 15.14 for applied cases). The agent typology at 10.10c is directly operational: it maps the ecosystem of actors relevant to any specific situation analysis. If you want to understand the economic analysis Part III (3.4–3.5) → Part IV (4.2–4.3) → Part XI in full. Part XI is doctrine, not framework: it is the current best application of the framework to economic arrangements given present material conditions. It is explicitly provisional in ways the framework is not. Read 11.0 as the diagnostic checklist, 11.3 as the strategic horizon, and 11.4 as the applied illustration. If you disagree with Part XI conclusions, the correct response is to apply the eight diagnostics in 11.0 and show where the analysis is wrong — not to dispute the framework that generated it. If you want to evaluate the framework itself honestly Part XII (12.5–12.7) → Part XIII (13.4–13.5) →. Section 12.5 identifies the specific failure modes available to this framework. Section 12.7 is the genealogical interrogation of the framework itself — who produced it, under what conditions, and what structural blind spots follow. Part XIV (14.3–14.4) is positioned at the end of the document,

after Part XVII, for this reason: its self-assessment is most credible once the entire argument has been read. Section 14.4 specifies what surpassing this framework would actually look like. These are the document's most honest passages. They read differently — more convincingly, or more disturbingly — after everything else has been read first. If you are an individual actor trying to use this framework in practice Part XVII directly. Start with 17.7 (the individual-level demands), then read whichever historical case in 17.1–17.4 most closely resembles your current situation. Then return to whichever Part of the main framework addresses the structural level at which you are operating. Part XVII is designed to be the bridge between structural analysis and individual decision, not a replacement for either. Five sections that must not be skipped regardless of entry point: 3.0 (the two-level architecture — life/suffering/becoming — that structures everything that follows); 12.5 (failure modes), 12.5d (the misuse triggers that operate without depending on the analyst's own honesty), 13.1c (tension vs contradiction diagnostic), and 15.13 (what the formulas cannot do). These are the document's primary safeguards against its own misuse. A reader who understands the framework's conclusions without understanding its limits and failure modes has not understood the framework. They have acquired a more sophisticated vocabulary for prior commitments.

I: How We Know: The Epistemological Foundation

1.1 — The Prior Question

Before any claim about what exists, what is right, or how to act, there is a question that most frameworks answer implicitly rather than explicitly: how do we know? The implicit answer is usually that we know because the framework tells us — because the internal consistency of the system is treated as its own validation. Hegel was explicit about this in a way that most frameworks are not: the system validates itself by completing itself, by showing that every apparent contradiction is a moment in a larger movement that resolves it. The grandeur of this move is also its weakness. A system that validates itself by completing itself has no external check. Its errors become features of the completion rather than evidence against it. (Hegel 1807) Marx grounded this differently. For Marx, the test of a framework was not its internal coherence but its capacity to illuminate actual historical processes — to explain why things happened as they did and to point toward the forces that would shape what happens next. This was a more honest epistemological position, and it produced analysis that was more durable precisely because it was more falsifiable. When the predictions were wrong, there was at least the possibility of asking why — of returning to the material conditions and asking what had been missed. What follows extends that move. Reality has patterns that exist independently of our perception of them, and the capacity to read those patterns is always partial, always developing, and always capable of error — including errors that will only become visible from positions the analysis has not yet reached. Building a system that cannot be wrong has been tried, repeatedly, and the systems that resulted were wrong in ways that caused serious harm precisely because they had foreclosed the mechanisms for detecting it. The requirement, then, is a methodology for being wrong productively: for identifying errors, revising conclusions, and continuing without abandoning the project. What distinguishes this from mere fallibilism is that it is an active commitment, not a passive acknowledgment — it requires building the revision mechanisms into the analysis rather than leaving them as theoretical possibilities that are never practically invoked. (Peirce 1878) The prior question is prior in two senses. It is logically prior: before you can assess whether a claim is well-supported, you need an account of what counts as support. And it is practically prior: frameworks that do not ask it explicitly tend to inherit whatever epistemological assumptions were most convenient at the moment of their construction — assumptions that then operate invisibly, shaping what counts as evidence and what counts as refutation without ever being examined. The damage done by unexamined epistemological assumptions is not random. It is systematically biased toward confirming whatever the framework already holds. What does it look like when a framework answers the prior question badly? Three

failure modes are diagnostic. The first is self-sealing: the framework's internal logic is used to evaluate claims made against the framework itself. Challenges are processed through the framework's own categories, which means they can only be resolved in the framework's own favor. Marxist-Leninist parties that explained every policy failure as sabotage by class enemies, market fundamentalists who explained every market failure as insufficient deregulation — both deployed self-sealing systematically. The second failure mode is criterion-shifting: the standards of evidence change depending on whose claims are being evaluated. Claims that support the framework's conclusions are held to lower evidentiary standards; challenges are held to higher ones. The asymmetry is not always deliberate. It operates through what the framework trains its users to notice and what it trains them to explain away. The third failure mode is unfalsifiability by design: the framework's core claims are formulated in ways that preclude any evidence from counting against them. This is different from a framework that has not yet been falsified. An unfalsifiable claim is one where no possible evidence could, even in principle, count as disconfirming. The distinction matters because unfalsifiability is often presented as a strength — the framework is said to apply universally, to account for all cases, to have no exceptions. The framework presented here makes the opposite claim: where its claims cannot be specified precisely enough to be falsifiable, that is a weakness to be named, not a feature to be celebrated. Asking the

prior question explicitly produces two immediate results. First, it requires the framework to commit to criteria for its own evaluation before those criteria are applied to anything else. The criteria cannot be derived from the framework's conclusions — that is circular — but must be grounded in something prior to the framework itself. Here that grounding is in material consequences: claims are evaluated by whether they make accurate predictions about outcomes that can be independently checked, by whether they survive contact with evidence generated by conditions the framework did not design, and by whether they perform better than available alternatives on the same domain. Second, asking the prior question makes the framework's epistemological commitments visible to the people who use it, so those commitments can themselves be evaluated rather than assumed. The alternative to asking the prior question explicitly is not to have no answer to it — it is to have an implicit answer that cannot be examined. Every framework operates with some set of epistemological assumptions. The question is whether those assumptions are chosen deliberately and named honestly, or inherited unreflectively and protected by the difficulty of making them visible. This framework chooses the former. That choice imposes costs: it makes the framework's foundations available for criticism

and it requires that the framework be revised when its foundations do not survive that criticism. Those costs are the price of intellectual honesty.

1.2 — The Primary Diagnostic

The most important test for any framework, ideology, or institutional arrangement is what it does when its predictions are contradicted by evidence. One response is revision: the model updates, the analysis continues from the corrected position, and the framework becomes more accurate. The other is insulation — the evidence gets reinterpreted, the conditions under which the test occurred are found inadequate, the people who were supposed to implement the model are found to have failed it, and the model itself remains untouched. Every system has a tendency toward one or the other, and that tendency is more diagnostic than any of its substantive claims. Dialectical materialism was supposed to be immune to the second response because the dialectic itself guaranteed that contradictions would be productive — that apparent failures were moments in a larger process of development. In practice, this guarantee functioned as the most sophisticated form of insulation available. Every failure could be absorbed as a necessary moment of negation on the way to a higher synthesis. The framework could not be wrong because wrongness was built into its conception of how truth develops. The result was a tradition that became progressively less capable of learning from its own failures — which is the opposite of what its epistemological commitments implied. From dialectical materialism, the genuine insight is preserved: reality moves and analysis must move with it. What is rejected is the teleological structure that made the dialectic self-insulating. The movement of reality does not guarantee productive contradictions. It guarantees that models built at one moment will become inadequate as conditions change, and that the work of revision is permanent rather than a transitional phase before the system completes itself. This criterion applies to this framework with the same force it applies to everything else. Any portion of what follows that explains its own failures by reference to external interference rather than by revising its conclusions is exhibiting exactly the failure mode identified here. That is not a caveat added for appearance. It is the framework's most important self-imposed constraint. The distinction between revision and insulation is the single most useful tool for evaluating frameworks, ideologies, and institutional claims, because it can be applied without prior agreement on any substantive question. Two people who disagree entirely about economics, politics, or ethics can nonetheless agree on whether a given framework revises or insulates in response to contradicting evidence. That agreement is possible because the diagnostic is procedural, not substantive — it evaluates how a framework responds to evidence, not what the evidence shows. What insulation looks like in practice: Insulation takes characteristically different forms depending on the type of claim being protected. For empirical predictions, the most common insulating move is auxiliary a

djustment: the prediction failed not because the model was wrong but because the conditions under which it was tested were not the right conditions. The model still applies — it just did not apply here, in this case, under these specific circumstances. Applied consistently, this move renders the model untestable in practice while preserving its theoretical integrity. For normative claims, the most common insulating move is implementation failure: the ideal was not wrong, the implementation was. True socialism has never been tried. True free markets have never existed. The commitment to the ideal is preserved by locating every failure in the distance between ideal and implementation rather than in the structure of the ideal itself. What revision looks like in practice: Genuine revision has three features that distinguish it from performed revision — the kind that accommodates surface criticism while preserving core commitments unchanged. First, the revised position makes different predictions than the original. A revision that produces exactly the same predictions as the original has not actually changed anything analytically useful. Second, the revision is costly: it requires abandoning a position previously held, not merely adding qualifications. Third, the revision is not triggered by strategic pressure alone. Frameworks that revise only when continued insulation becomes politically untenable — when the cost of defending the original position exceeds the cost of appearing to update — are not genuinely responsive to evidence. They are responsive to social pressure, which is a different mechanism and produces different epistemic qualities. The hardest version of the diagnostic is its application to one's own framework. Every analyst who uses a framework has an investment in its validity — a commitment, often prior to any specific empirical question, that the framework's general orientation is correct. That investment produces a systematic bias toward finding confirming evidence and away from weighting disconfirming evidence appropriately. The bias is not usually conscious. It operates through what questions get asked, what cases get selected for analysis, and what anomalies get treated as requiring explanation versus what anomalies get treated as noise. The framework that genuinely applies the primary diagnostic to itself must have an explicit procedure for generating disconfirming evidence (Popper 1959) — not just a willingness to accept it if it arrives, but active production of conditions under which it would arrive. The postmodern and post-structuralist critique of grand narratives applies to this framework with genuine force and deserves more than deflection. The critique is right that frameworks which claim to explain everything typically explain nothing specifically, that analytical authority can function as a form of epistemic violence against positions it cannot incorporate, and that the history of totalizing systems is substantially a h

istory of analytical capture. What this framework claims in response is not immunity from that history but a specific structural difference: it is built around its own failure modes (12.5), specifies what evidence would require it to revise or abandon its core claims (14.3), and treats the genealogy of its own production as material evidence about its blind spots (12.7). The difference between a totalizing narrative and an empirically grounded framework with acknowledged limits is not that the second is free from the first's failure modes — it is that the second has built the mechanisms for detecting those failures into the analysis itself. Whether this framework has actually succeeded in doing that is, as 12.5 makes clear, a question the framework cannot answer about itself. The primary diagnostic does not require that a framework be frequently revised to be good. A framework can be stable over long periods if the evidence consistently supports it. What it requires is that revision is possible in principle and has happened in practice when evidence warranted it. A framework that has never revised a core position, regardless of how much evidence has accumulated across how many cases and how many domains, has failed the diagnostic — not because stability is bad, but because perfect stability in the face of accumulating evidence from complex social systems is not a plausible epistemic outcome. It is the signature of insulation.

1.3 — Probabilistic Determinism

Apparent contingency in historical and social events conceals structural forces that can be identified and worked with, even when they cannot be precisely predicted — which is not determinism, and not the fatalism

that sometimes masquerades as realism, but something more demanding than either: the recognition that structure shapes what is possible without eliminating the difference that agency makes within those constraints. Historical materialism was right about this in its domain: the specific forms of economic organization in a given era shape what political arrangements are possible, what ideologies are compelling, and what kinds of change can actually occur. The insight generalizes beyond the economic domain, and what follows applies it more broadly — but the basic move, from apparent contingency to structural explanation, is one of historical materialism's genuine contributions. What the framework adds is the recognition that the structural forces extend deeper than the economic. The material conditions that shape history are not only the relations of production. They include the biological conditions of the organisms who produce — their vulnerabilities, their cognitive architecture, their social dependencies, the specific forms of suffering and flourishing that their physiology makes possible. These conditions constrain what economic arrangements can be stable, what ideological commitments can be sustained, and what kinds of institutional development are available at any given moment. Probabilistic determinism is the position that structure shapes what is possible without eliminating the difference that agency makes. It is not the claim that outcomes are fixed in advance — that would be hard determinism, which the available evidence does not support and which would make political analysis pointless. It is not the claim that outcomes are open in ways indifferent to structural conditions — that would be voluntarism, which produces analysis that consistently overestimates the reach of organized will and underestimates the constraints within which it operates. The position is more demanding than either: structural forces constrain the distribution of probable outcomes, and that distribution can be analyzed and worked with, even though which outcome within it actually occurs depends on contingencies that cannot be predicted in advance. The practical implication is that analysis must operate simultaneously at two levels that are easily confused. At the structural level: what configurations of conditions make which outcomes probable, what thresholds once crossed shift the distribution of possibilities, what the direction of current movement is. At the conjunctural level: what specific contingencies — leadership, timing, coalition, error — determine which probable outcome actually materializes. Analysis that operates only at the structural level will be right about direction and wrong about timing. Analysis that operates only at the conjunctural l

evel will be right about specific events and unable to explain the pattern of which events were possible in the first place. The framework operates primarily at the structural level, and is explicit about what that means for the precision of its claims. The distinction from fatalism matters because it is frequently collapsed by both critics and practitioners. Fatalism is the use of structural analysis to justify inaction: if outcomes are structurally determined, then agency cannot change them, so there is no point in trying. Probabilistic determinism produces the opposite conclusion. If structure shapes the distribution of probable outcomes, then organized effort that shifts structural conditions shifts the distribution — making previously improbable outcomes more probable and previously probable ones less so. The appropriate response to structural analysis is not resignation but the identification of which interventions, at which points, move the distribution in the direction the evaluative criterion requires. This is why the framework insists on distinguishing probabilistic determinism from both hard determinism and from contingency as a primary analytical category. Hard determinism — the claim that structural forces produce specific outcomes with necessity — fails because it cannot account for the genuine variation between cases with similar structural configurations, and because it produces strategic paralysis: if the outcome is determined, the question of what to do dissolves. Contingency as a primary category — the claim that historical outcomes are fundamentally open and that structural analysis cannot generate reliable expectations — fails because it leaves the analyst unable to distinguish between actions that are likely to produce their intended effects and those that are likely to be absorbed, redirected, or counterproductive. Neither position is analytically adequate to what political actors actually need from structural analysis. The operational consequence: Probabilistic determinism specifies that structural forces establish the range of likely outcomes and the conditions under which different parts of that range become more or less probable — without specifying which outcome within the range will

actually obtain. This is the form of structural knowledge that is most useful for actors trying to navigate and influence systems: it tells them which interventions are likely to produce effects, which conditions make the system more or less available for the changes they want to produce, and which configurations of variables are most leverage-sensitive. It does not tell them that the outcome they want will obtain; it tells them what needs to be true for it to be in the probable range and what would be required to shift the probabilities. The analogy with weather forecasting is precise: a competent forecast does not specify exactly what will happen but accurately characterizes the probability distribution over what might happen and identif

ies the conditions that are most determinative of where in that distribution the actual outcome falls. The forecast is not falsified by an event that falls in the unlikely tail of the distribution; it is falsified by a systematic pattern of miscalibration — consistent overconfidence or underconfidence about specific types of outcome. Structural analysis in the framework operates the same way: it is not falsified by individual cases that deviate from structural expectations, but by systematic patterns of deviation that indicate the structural variables have been misspecified. Connection to 1.1 (the prior question) and 1.2 (the primary diagnostic): the prior question establishes why analytical frameworks are necessary for engaging with material reality. The primary diagnostic establishes the specific question (who benefits and who bears costs?) that orients the analysis. Probabilistic determinism establishes the epistemological register in which that analysis can be reliably conducted: not certain prediction, not mere description, but calibrated probability assessment over structural dynamics. The logical form this probabilistic determinism takes is defeasible inference rather than strict implication: structural forces make certain outcomes probable unless specific defeating conditions obtain. The full specification of this reasoning type is developed in 12.1.

1.4 — Convergence Under Material Pressure

‘De facto universalism’ is a descriptive observation, not a normative claim. It describes what happens when different societies face similar material pressures over time: they tend toward similar solutions, not because anyone persuaded them to, but because the alternatives were less viable. The convergence is real and is driven by material selection, not ideological agreement. What this means for normative universalism — whether some standards apply universally — is a separate question addressed in Parts IV and VIII. One of historical materialism's most important observations was that the spread of capitalism across the world was not the result of ideas traveling — not the persuasion of one culture by the superior values of another. It was the result of material pressure: economic and military advantages that made certain forms of organization dominant regardless of what any particular culture preferred. The convergence of different societies on similar institutional forms was the residue of that selection process. This observation is more general than Marx applied it. The convergence of different civilizations on similar solutions to coordination problems — law, taxation, standing armies, markets — preceded capitalism by millennia. It was produced by the same mechanism: societies that failed to develop certain coordination capacities were absorbed, destroyed, or transformed by those that had. The universalism is not ideological. It is the result of material forces selecting for particular kinds of organization across all contexts in which those forces operate. The implication is that the choice is not between universalism and pluralism. The material convergence is already happening. The question is whether it is directed consciously — whether the forces driving convergence are understood well enough to reduce the suffering involved in the transition — or whether it continues to be driven entirely by the selection mechanism, which has historically been very effective and extremely cruel. The distinction between descriptive and normative convergence matters more than it might initially appear, because the conflation of the two is one of the most common moves in both progressive and conservative political argument. The conservative version: existing arrangements are what they are because material pressures produced them; therefore they are what they must be; therefore challenging them is utopian. The progressive version: some standards are universal because rational agents in idealized conditions would converge on them; therefore any arrangement that deviates from those standards is irrational. Both moves convert a

descriptive observation about convergence under real conditions into a normative claim about what conditions must produce — which does not follow. What de facto universalism actually establishes: It establishes that across a wide range of cultural, political, and historical contexts, certain functional requirements reliably emerge. Socie

ties require some means of coordinating collective action, some system for managing violent conflict internally, some institutional form for transmitting knowledge across generations, some arrangement for managing the production and distribution of material necessities. The specific institutional forms that address these requirements vary enormously. What converges under material pressure is the functional requirement, not the institutional form. This has analytical implications: it suggests that the functional requirement is robust — it reflects something durable about the material conditions of human social life — while leaving the space of adequate institutional responses much larger than any single tradition's history would suggest. The selection mechanism: De facto convergence is produced by a material selection mechanism: arrangements that do not adequately address functional requirements generate pressures — internal conflict, productive inefficiency, reduced capacity to resist external competition — that tend to destabilize them over time. This is not a guarantee of progress. The selection mechanism operates on fitness for current conditions, which may or may not produce arrangements that are adequate to conditions that do not yet exist. And the mechanism is slow, operating across generations and through processes of conflict and transformation that are not predictable in their specific form. What it is not is teleological: there is no endpoint toward which the selection process tends, no final arrangement that material pressure is moving toward. The process is adaptive without being directional. The normative question — what arrangements are better rather than merely more stable — cannot be settled by the descriptive observation of convergence. Arrangements can be stable for long periods while producing structural suffering at scale, as the preceding analysis of long-term stability shows. The convergence that matters normatively is not convergence on what material selection produces but convergence on what adequate reflection from multiple starting points endorses. Those are different criteria and need not produce the same results. De facto universalism is useful analytically — it identifies functional requirements that any adequate institutional arrangement must address — but it is not sufficient normatively, and treating it as sufficient is one of the characteristic errors of frameworks that naturalize the status quo by treating material selection as moral validation. The practical use of de facto universalism in analysis: when assessing whether a proposed institutional change is viable, asking whether similar arrangements have emerged under analogous material pressures across different contexts is a useful evidentiary check. Arrangements that have never appeared anywhere under any conditions that resemble the target context face a higher burden of proof than arrangements that have appeared repeatedly. This is not an argument against i

nnovation — the conditions under which an innovation is proposed may genuinely differ from all previous conditions in ways that make the historical record uninformative. But the absence of historical analogs should sharpen scrutiny, not trigger rejection.

1.5 — Inaction and Its Costs

Marx's critique of Feuerbach — that philosophers had only interpreted the world and the point was to change it — contained an epistemological claim as well as a practical one. The stance of pure interpretation is not neutral. It reproduces the conditions it describes. This framework takes that claim seriously while extending it: inaction is a decision, it has material consequences, and those consequences systematically favor whoever benefits from current conditions. The use of uncertainty as justification for inaction must be evaluated by the same standards as any proposed action — not exempted from evaluation by the appearance of caution. The mechanism is straightforward. Existing arrangements produce existing distributions of resources, opportunity, and suffering. Those distributions are sustained not only by active force but by the normal functioning of institutions whose continuation requires no additional effort from anyone. When actors who

have the capacity to change those arrangements choose not to act — because the analysis is uncertain, because the moment is not right, because the costs are unclear — the arrangements continue to produce what they were producing. The inaction that looks like caution is indistinguishable, in its material consequences, from active support for the status quo. The asymmetry is structural: action against existing arrangements requires effort; inaction sustains them automatically. This does not imply that every moment is equally appropriate for action, or that uncertainty can never justify delay. It implies that the decision to delay must be subjected to the same rigorous scrutiny as the decision to act. The question is not 'is the analysis certain enough to act?' but 'what are the material consequences of acting versus not acting under current conditions, given current levels of certainty?' Framed that way, uncertainty becomes a reason to assess consequences more carefully, not a reason to exempt inaction from assessment. An actor who consistently finds that their uncertainty about acting happens to align with the costs they would bear from acting has not applied this standard honestly. The costs of inaction are distributed unevenly. The populations who experience the most structural suffering from existing arrangements bear the costs of inaction most directly. The populations who benefit from existing arrangements experience inaction as the preservation of conditions they prefer. The claim that caution is neutral is therefore not politically neutral — it systematically favors those who benefit from the conditions being cautiously preserved. The claim that inaction is a decision with material consequences is not primarily a moral claim, though it has moral implications. It is an analytical one: the default state of a social arrangement under no external intervention is not preservation of its current form but reproduction of its current dynamics. If those dynamics are producing structural suffering, the absence

of intervention allows the reproduction to continue. The person who does not act has not thereby avoided causing harm — they have chosen not to interrupt a causal process that continues to produce harm through their inaction. The harm is real regardless of whether it was intended. The asymmetry in how uncertainty is deployed: Uncertainty is a genuine feature of complex social systems, and this framework does not minimize it. But the deployment of uncertainty in practice is systematically asymmetric: uncertainty about the consequences of action is treated as a reason for inaction, while uncertainty about the consequences of inaction is not treated as a reason for action. The asymmetry is not epistemically justified. If uncertainty is a reason for caution, it should apply equally in both directions. If it applies only to action and not to inaction, it is functioning as a bias toward the status quo, not as a genuinely cautious epistemic stance. The framework that deploys uncertainty selectively, in ways that consistently produce the same policy conclusion regardless of the evidence, has failed the primary diagnostic in 1.2: it is insulating, not updating. Who benefits from the asymmetric deployment of uncertainty: This is not a neutral question. The asymmetric deployment of uncertainty — the systematic weighting of action-risks over inaction-risks — consistently produces conclusions that favor whoever benefits from current conditions. When analysts warn that we cannot be certain of the effects of raising the minimum wage while treating the effects of not raising it as the baseline that requires no justification, when environmental uncertainty is used to delay action while the costs of inaction accumulate asymmetrically on those least able to bear them, when institutional reform is held to standards of proof that the existing institution was never required to meet — in each case, uncertainty is functioning as a structural resource for maintaining arrangements that benefit specific actors. Identifying this does not require imputing bad faith. It requires tracking who benefits from the conclusion that uncertainty is a reason for inaction, which is the interest-alignment test that the framework applies to all analytical claims. The epistemic corollary: frameworks that treat inaction as the neutral default are not epistemically more cautious than frameworks that require justification for inaction as well as action. They are epistemically biased in a specific direction. Genuine epistemic caution requires acknowledging that the costs of inaction are uncertain too, that those costs fall on specific populations in ways that can be analyzed and measured, and that the comparison between the expected costs of action and the expected costs of inaction is the analytically relevant question — not the comparison between action and a baseline treated as costless. Diagnostic for the analyst: when uncertainty is invoked as a reason for a policy conclusion, apply the

test symmetrically, and state the result in the recommendation itself: a conclusion that invokes uncertainty must report its confidence level for the inaction baseline in the same document, not hold it as an internal check. What is the confidence level in the claim that inaction is safe? What evidence would change that assessment? If the uncertainty is genuine, it should produce roughly equal caution about both action and inaction. If it consistently produces caution only about action, the uncertainty is being deployed asymmetrically — and the interest-alignment question applies.

1.6 — Truth, Communication, and the Management Problem

This section might read as advice about communication craft — as if the framework were offering tips for persuasion. It is not. It is making a structural claim with material consequences: in conditions of epistemic asymmetry, the same true claim produces different outcomes depending on how it enters the information environment. This is not a licence for manipulation. It is an acknowledgement that truth does not automatically win, and that a framework built on epistemic discipline has an obligation to account for the actual conditions of information reception rather than the idealized conditions under which truth would speak for itself. Ignoring this is not principled — it is a form of epistemic negligence that allows correct analysis to be systematically defeated by inferior analysis packaged better. For truth to produce better outcomes, the way it is communicated must account for the actual conditions under which people receive and process information. Careful communication is not a retreat from honesty. A true claim communicated in a way that ensures it will be misunderstood or rejected is not, by any practical standard, an improvement over more careful communication. The distinction between strategic communication and manipulation is real: communication that preserves the recipient's capacity to evaluate what is being said is strategic but not manipulative; communication that works by bypassing evaluation is manipulation regardless of whether the content is true. There is a harder problem here that this framework inherits from the tradition of political thought it engages with. Both Hegel and Marx recognized that historical actors are not transparent to themselves — that people act from motivations they do not fully understand, in service of forces they cannot see, producing outcomes they did not intend. The implication is that political communication always operates in a context where the gap between what people believe they are doing and what they are materially doing is significant. Managing that gap requires a form of communication that is more than the simple transmission of information: helping people understand the material consequences of their actions and beliefs, working with how they actually process reality rather than against it. The line between this and manipulation is real and easily crossed. The management problem this generates is specific: how to communicate true claims in conditions where the reception of those claims is structured by forces that the communicator cannot fully control and the audience cannot fully see. The problem is not whether to tell the truth. It is that the same true claim, communicated differently, produces different practical effects — and those effects are part of what the framework must account for. A true claim that is communicated in a way that guarantees it will be received as threatening, will activate defensive rather than receptive processing, and will be stored as evidence of the communicator

's bad faith rather than as information to be evaluated has not, by any practical standard, been successfully communicated. The truth has been stated; the communication has failed. This does not license strategic deception. The framework's epistemological commitments rule that out directly: a communicator who withholds or distorts true claims to produce better short-term reception is undermining the trust infrastructure that makes reliable communication possible at scale — and that infrastructure, once undermined, is difficult to rebuild. The distinction is between communication that is honest and communication that is honest and effective — that accounts for the actual conditions of reception without sacrificing the content being communicated. Careful communication is not a compromise between truth and manipulation. It is what truth-telling looks like when it is serious about actually producing the understanding it aims at. The management problem has a structural dimension that individual communicators cannot resolve alone. When

information environments are controlled by actors whose interest is in preventing accurate reception of specific true claims — when the channels through which claims travel are designed to sort, distort, and discredit before reception occurs — careful individual communication is insufficient. The problem requires the analysis of information infrastructure as a material condition, not only the refinement of individual communicative technique. This is the connection to Part I's account of epistemic infrastructure: the conditions under which true claims can be accurately received are themselves produced and maintained by structural arrangements that can be analyzed, and that are currently distributed in ways that systematically advantage specific kinds of claims over others. The three specific management problems: The first management problem is translation: the same true claim requires different formulation to reach different epistemic communities without losing the substance that makes it true. A structural claim about the relationship between wage suppression and aggregate demand requires different entry points for an economic policy audience, an organized labor audience, and a general public audience — not because the claim is different for each but because the prior knowledge, the trusted sources, and the interpretive frameworks that determine whether the claim is received as credible differ. The second management problem is timing: true claims that enter information environments before the conditions for receiving them are present are absorbed by the existing interpretive framework and lose their disruptive potential. The third management problem is sequencing: complex structural arguments require the simpler claims on which they depend to be established before the complex claim is introduced — the order in which true claims are presented determines whether each claim builds on established foundations or requires its

interlocutor to accept multiple new claims simultaneously, which exceeds most audiences' epistemic tolerance. None of these management problems licenses misrepresentation. The framework is explicit that the discipline of accurate analysis — the commitment to following the evidence regardless of whether it supports preferred conclusions — is both epistemically and strategically valuable. The management problems are problems of communication strategy within the constraint of honest content, not arguments for abandoning that constraint. A framework that solves its management problems through selective truth-telling loses its epistemic credibility at the moment its selectivity is identified — and in an information environment with many actors motivated to identify it, that moment arrives sooner than the framework's proponents typically expect.

1.6b — Belief and Doubt: The Productive Equilibrium

A specific failure mode of analytical frameworks, including this one, is the progressive conversion of working hypotheses into orthodoxies through the suppression of genuine doubt. Why it belongs here is structural, not psychological. Frameworks that cannot genuinely doubt their own conclusions have lost the adaptive mechanism that makes them useful — they become ideological in the pejorative sense: self-confirming, closed to revision, and increasingly misaligned with the reality they were designed to analyze. The tension between belief and doubt is not a philosophical nicety. It is the mechanism by which PMN avoids the capture-by-certainty failure mode that has closed off most of the analytical traditions it engages with. Belief and doubt are not opposites to be balanced but tensions to be maintained simultaneously. Neither alone is sufficient. Pure confidence stops testing its conclusions; pure skepticism cannot distinguish better from worse. Together, held in productive tension, they constitute the only honest posture available to analysis that takes seriously both the existence of something more objective to move toward and the permanent incompleteness of any given position. Belief — genuine confidence that there is something more objective to reach, that analysis can get closer to reality through revision and inquiry, that human capacity for understanding is not exhausted by what is currently understood — is what makes the entire project possible. Without it, the analysis collapses into a relativism that cannot distinguish better from worse, and the commitment to continue loses its justification. Belief provides direction. It sustains the commitment to keep moving when current positions become uncomfortable or when the analysis points

toward conclusions that challenge what was previously held. But belief without doubt produces exactly the failure mode this framework is designed to avoid. The framework that is too confident in its own conclusions stops testing them. It begins to insulate its core commitments from evidence rather than revising them. It starts to read every challenge as confirmation and every contrary evidence as proof of the challenger's bad faith. Fanaticism — in its intellectual as well as its political forms — is belief that has consumed doubt and left only certainty behind. The result is always the same: a framework that has become more interested in defending its conclusions than in tracking reality. Doubt — genuine critical pressure applied to what is currently held — is what keeps the analysis honest. It forces conclusions against evidence, holds open the possibility that what seems true is incomplete, prevents the accumulation of certainty that mistakes a position for a destination. The difference between belief with doubt and belief without it is the difference between a direction and an identity — between something that can be wrong and learn, and something that has foreclosed learning by making itself unfalsifiable.

The productive equilibrium between the two is not a stable midpoint. It is a dynamic tension that has to be actively maintained. In conditions of intellectual or political pressure, the tendency is always toward one extreme or the other — either the false security of certainty or the paralysis of total skepticism. Maintaining the tension requires recognizing that what sometimes looks like retreat — revising a position, acknowledging an error, moving toward a conclusion that contradicts what was previously held — is often the form that genuine progress takes. Moving toward greater objectivity sometimes requires moving away from what was previously believed, and that movement can feel like loss before it feels like advance. (Peirce 1878) The anti-dogmatic principle follows from the same epistemological foundation: commit fully to the direction while holding specific positions lightly enough that evidence can revise them. The commitment is to the movement, not to any particular point along it. The productive character of doubt: Doubt is productive in a specific sense: it is the mechanism through which frameworks update in response to evidence. A framework that cannot be genuinely doubted — that has developed the social and institutional infrastructure to insulate its conclusions from the kind of scrutiny that would require revision — has eliminated the adaptive mechanism that makes it useful. It can still produce outputs, but those outputs are increasingly misaligned with what the evidence would support, because the alignment mechanism has been disabled. The political consequences of this misalignment are predictable and have been documented extensively: frameworks that cannot doubt themselves consistently produce strategic failures in proportion to the distance between their conclusions and what the evidence would support if engaged honestly. The equilibrium problem: The productive equilibrium between belief and doubt is genuinely difficult to maintain, and the difficulty is not symmetric. Excessive doubt — the inability to maintain working hypotheses long enough to generate the evidence that would test them — produces a different failure mode: analytical paralysis, the inability to act on structural knowledge because the knowledge is always provisional in a way that prevents its translation into organizational strategy. The equilibrium requires holding conclusions firmly enough to act on them while remaining genuinely open to the evidence that would require their revision. This is what distinguishes good scientific practice from both dogmatism and from the paralysis that sometimes masquerades as epistemic humility. The self-application: this framework holds specific conclusions firmly — about the biological foundation of the evaluative criterion, about the structural production of suffering, about the predictable consequences of power concentration. These are not held with the tentativeness of genuinely open questions; they are working hypotheses

that have survived enough contact with evidence to warrant the confidence required to act on them. What is held tentatively are the specific applications of those conclusions to particular cases, the specific institutional proposals that follow from the general analysis, and the specific predictions about trajectories that the formula architecture generates. The equilibrium is maintained by being clear about which parts of the framework are more established and which are more provisional.

1.7 — Knowledge Production as a Material Process

Epistemology has a social address. How we know cannot be separated from who produces knowledge, under what conditions, with what resources, and in whose interest — not because this reduces knowledge to mere interest-expression, but because knowledge production is itself a social process shaped by the same structural forces that shape all social processes, and treating it as exempt from that analysis produces systematic blind spots. Research requires resources. The distribution of those resources determines what questions get asked, what methods are considered legitimate, what findings get published, and what conclusions get amplified. Individual researchers' intentions matter less here than the structural outcome of the funding relationships, institutional hierarchies, and professional incentives that govern knowledge production in any given domain. A framework that treats knowledge production as a process operating outside these structural forces will reliably misread both the content of knowledge and the distribution of ignorance. (Gramsci 1971) The distribution of ignorance is as consequential as the distribution of knowledge. What is not studied, not funded, not published, not amplified — the systematic gaps in what any society officially knows about itself — are not random. They reflect the interests of those who control the resources for knowledge production, and they systematically underrepresent the conditions and experiences of those with the least power to direct that production. The populations whose suffering is least visible in official knowledge are typically the populations with the least capacity to fund research into their own conditions. This applies to this framework itself. The framework is produced from a specific position, drawing on intellectual traditions more accessible to some than to others. Its blind spots are not correctable by the framework alone. They require engagement with perspectives produced from different material positions — which is why the community of inquiry is not a philosophical preference but a material necessity for frameworks that aspire to track reality rather than the conditions of their own production. (Peirce 1878) The materialist analysis of knowledge production does not reduce knowledge to ideology — the claim that every belief is merely an expression of the believer's class position or material interest. That reduction is epistemically self-defeating: if true, it applies to itself, and the materialist analysis of knowledge production is merely an expression of whoever produced it. What the analysis establishes instead is more modest and more precise: knowledge production is shaped by material conditions in specific ways that produce identifiable systematic biases, and those biases can be studied, named, and partially corrected for. The goal is not to abolish the material conditions of knowledge production — that is not a coherent goal — but to account for them in the epistemic weight assigned

to claims. The resource distribution problem: Research requires resources: time, equipment, access to subjects, institutional infrastructure, and the sustained attention that only becomes available when basic material needs are stable. The distribution of those resources is not epistemically neutral. Questions that attract resources tend to be questions that resource-holders find important, profitable, or strategically useful. Questions about the conditions of populations with no resources to direct toward research are systematically underrepresented in the research literature, not because those questions are less important but because the market for answers to them is structurally thin. The result is a systematic pattern in what is known and what is not known — a distribution of ignorance that maps onto the distribution of material power. The credentialing problem: Epistemic authority — the social capacity to make claims that others treat as reliable — is distributed through institutional credentialing systems that were built under specific material conditions and reflect those conditions in their design. The university system, peer review, professional licensing, and journalistic standards all function as mechanisms for sorting claims by source before evaluating them by content. This sorting has real epistemic value: it filters out many low-quality claims. But it also systematically deprioritizes knowledge produced outside credentialed institutions, which includes most of the knowledge that populations under structural constraint have developed through direct engagement with the conditions that affect them. The practitioner, the community organizer, and the person navigating a complex bureaucratic system often have more accurate models of how specific institutional arrangements work than researchers studying those arrangements from outside — but their knowledge circulates in different channels and carries different social weight. The audience-orientation problem: Knowledge

production is not only shaped by who funds it and who produces it, but by who it is produced for. Research that will be evaluated primarily by credentialed peers within a discipline develops differently than research designed to be used by practitioners making decisions under real conditions. The peer-review process rewards theoretical sophistication, methodological novelty, and engagement with disciplinary debates in ways that may not track what is most useful for practical application. The result is a systematic divergence between what is known in the academic literature and what is useful in practice — not because academic knowledge is false, but because the criteria by which it is evaluated reward different properties than the criteria by which practical knowledge is evaluated. The implication for this framework: all empirical claims in the following Parts must be understood as shaped by these structural biases. The cases analyzed are predominantly cases for which document

tation exists in credentialed sources — which means they are disproportionately cases from contexts where research infrastructure was present, which correlates with economic and political power. The absence of cases from other contexts does not indicate that the patterns analyzed here do not operate there. It indicates that the conditions for documenting them systematically were absent. This is a genuine epistemic limitation, and it is named here rather than explained away. Expert authority, fallibility, and the situated knowledge problem The materialist analysis of knowledge production has an implication that requires careful handling because it is correct but easily weaponized: experts can be wrong, and non-experts can be right. This is not a license for anti-intellectualism. It is a structural observation about the conditions under which different kinds of knowledge are produced and validated. Three distinctions are necessary to prevent a correct epistemological claim from collapsing into a false one. The first distinction is between the fallibility of expert claims and the equivalence of all claims. Expert knowledge is fallible because it is produced under specific institutional conditions that introduce specific biases: funding dependencies, professional incentives, publication pressures, the social dynamics of credentialed communities. These biases are real and produce systematic blind spots that the framework's analysis of knowledge production identifies directly. But the fallibility of expert claims does not make them equivalent to non-expert claims. Expert knowledge's institutional production also includes the mechanisms that make claims revisable: peer review, replication requirements, adversarial testing, accumulated track records of prediction. A claim produced within those mechanisms is epistemically different from a claim produced outside them — not infallible, but subject to correction processes that non-expert claims are not. The appropriate response to expert fallibility is not to treat all claims as equally valid; it is to demand that expert claims be held to the revision standards they are supposed to meet. The second distinction is between domains where expertise tracks material reality through empirical methods and domains where expert authority is primarily a function of institutional position. A virologist's claim about viral transmission is expert knowledge in the first sense: it is produced through methods that have track records of prediction and correction. An economist's claim about the optimal design of global trade policy is expert authority in the second sense: it is produced within a discipline with contested foundations, value commitments that are not empirically derivable, and a record of confident predictions that did not materialize. The framework requires treating these differently. Empirical expertise in domains with functioning correction mechanisms deserves substantial epistemic deference, subject to the st

structural biases analyzed here. Institutional authority in domains where the expert consensus reflects the interests of those who fund and credential the discipline deserves the same scrutiny the framework applies to any other claim whose acceptance concentrates power. The third distinction is between general expertise and situated knowledge. Populations who experience the direct consequences of institutional arrangements have informational access to those consequences that institutional analysts do not. A person navigating the welfare system knows things about how that system actually operates that no academic study of the welfare system has fully captured. This is not because lived experience is a superior epistemological method; it is because the phenomena being studied are partially constituted by how they are experienced by those subject to them, and systematic distance from that experience produces systematic blind spots. Situated knowledge is not infallible either — the person inside a system can also misunderstand structural causes and attribute to

individual factors what is produced by institutional design. But situated knowledge is informationally privileged with respect to specific questions: how does this arrangement actually function for the people subject to it? Answering that question adequately requires both the structural analysis that the framework provides and the situated knowledge that those subject to the arrangement produce. Neither is sufficient alone. The framework's community of inquiry (1.7) is a methodological commitment to both, not a rhetorical gesture toward inclusion. The practical implication: the anti-technocratic guardrail (6.4) and the epistemological humility of Part I are not invitations to dismiss expert knowledge in favor of intuition, tradition, or the preferences of any particular community. They are requirements for holding expert knowledge to its own stated standards, attending to the structural conditions that produce systematic biases in what gets studied and what gets published, and integrating the situated knowledge of those most affected by the arrangements being analyzed. The framework's epistemological position is not anti-expert; it is anti-capture — opposed to the conversion of epistemic authority into a shield against the accountability that authority claims to warrant. The distinction between transmitted knowledge and observed knowledge — between what has been received from authority and what has been derived from direct engagement with material reality — is not an invention of modern epistemology. Ibn Khaldun formulated it in the *Muqaddimah* in the fourteenth century as a methodological principle for historical analysis: transmitted accounts are subject to the biases, interests, and limitations of their transmitters in ways that make critical evaluation indispensable, while observed regularities in how social formations actually behave provide a check against which transmitted claims can be as

sessed. This is structurally identical to the framework's epistemological commitment in 1.2: the test of a claim is not its source or its authority but its responsiveness to the evidence that material reality produces. Ibn Khaldun applied this principle to dismantle historical accounts that he found implausible not because they conflicted with received authority but because they conflicted with what he knew about how populations, armies, and dynasties actually function. The framework applies the same principle to expert claims: not dismissal on the grounds of their institutional production, but assessment against the material evidence their claims are supposed to track. The independent derivation of this epistemological commitment in a fourteenth-century Islamic intellectual tradition — seven centuries before Peirce formulated the pragmatist version that this framework draws on most directly — is evidence that the commitment is tracking something about how reliable knowledge is produced that is not specific to any single intellectual tradition.

1.8–1.12 Epistemic Infrastructure as Material Condition Part I established that knowledge production has a social address: what is produced, what circulates, and what acquires standing are all shaped by material conditions. Sections 1.8–1.12 extend that analysis to the infrastructure level — the systems through which knowledge, once produced, is distributed, validated, stored, and accessed. This is a distinct question. A society can have sound conditions for producing accurate knowledge and still have epistemic infrastructure so degraded that the knowledge produced cannot circulate or form the basis of coordinated action. The political consequences of epistemic conditions operate at the social level, not the individual level, and that social level is what this group of sections analyzes. The infrastructure level is where knowledge production and political action are connected or disconnected. A research environment that generates accurate analysis of structural conditions does not automatically produce coordinated political response to those conditions. The path from accurate knowledge to effective collective action runs through infrastructure: the systems through which knowledge is distributed, validated, stored, made accessible, and translated into the shared understanding that is the precondition for coordinated response. Each of these systems can degrade independently. Each degradation produces specific political consequences. Sections 1.8 through 1.12 analyze the key components of epistemic infrastructure — their structural properties, their characteristic failure modes, and their relationship to the visibility variable *V* that connects epistemic conditions to the structural suffering formula. The distinction between knowledge production conditions (Part I's primary subject) and epistemic infrastructure addresses a question that the sociology of knowledge has often conflated: the difference between the

social conditions that shape what claims are produced and the material systems that determine which produced claims travel, which acquire standing, and which reach the populations that would act on them. A society can have excellent conditions for producing accurate knowledge about, say, the structural causes of poverty — honest researchers, methodologically rigorous institutions, genuine commitment to evidence — and still have epistemic infrastructure so degraded that this knowledge never reaches the populations it describes, or reaches them only through frameworks that strip its structural content and convert it into something compatible with existing governing narratives. The four infrastructure components: Distribution infrastructure — the media systems, publishing channels, digital platforms, and educational institutions through which knowledge circulates — determines what claims can reach which populations and at what cost. Validation infrastructure — the peer review systems, editorial standards, professional credentialing processes, and reputational networks through which claims acquire legitimacy — determines which produced claims are treated as reliable. Storage infrastructure — archives, libraries, institutional memory, digital repositories — determines which claims remain accessible across time rather than degrading into inaccessibility. Access infrastructure — the economic, linguistic, geographic, and credentialing barriers that determine who can engage with validated knowledge — determines which populations can actually use the claims the other three systems produce, validate, and store. Each of these components can be degraded independently, and the degradation of any one reduces the effective function of the others. How infrastructure failure differs from production failure: Infrastructure failure is politically distinctive because it is less visible than production failure. A research community that systematically misreads evidence produces claims that can be challenged by other researchers who examine the same evidence. Infrastructure failure produces no such challenge signal: the correct claims may have been produced, validated, and stored, but if distribution reaches only the populations already convinced, if access barriers exclude those with the greatest material stake in the knowledge, or if storage is controlled by actors whose interests are served by selective accessibility, the claims produce no political consequences. This is why 1.8–1.12 treats infrastructure as a material variable deserving the same analytical attention as production conditions — and why V in the formula architecture is partly a function of infrastructure state rather than only of governing narrative content. The sequence from 1.8 through 1.12 moves from the broadest infrastructure question (what can be known, 1.8) through the social processes that produce epistemic authority (1.9), the structural consequences of information asymmetry (1.10), the dyna

mics of trust infrastructure and its collapse (1.11), and the specific relationship between epistemic infrastructure state and the visibility variable in the formula architecture (1.12). Each section addresses a different level of the infrastructure analysis; together they specify why epistemic conditions are not background context for political analysis but active variables that determine what structural change is possible and through what mechanisms. The self-application: the framework you are reading was produced under specific epistemic infrastructure conditions and will be distributed through specific channels to specific populations. Its analytical reach — how far it travels, how its claims are received, what authority they acquire — is itself partly a function of infrastructure variables. The genealogical section (12.7) addresses the production conditions; the transmission problem (17.8) addresses the distribution question. Infrastructure analysis and self-application are not separate exercises. The mechanism in detail: The epistemic infrastructure state affects V through three specific channels. The first is the attribution channel: when distribution and validation infrastructure is captured by actors with interests in attributing structural suffering to individual causes, the claims that circulate through that infrastructure will systematically present individual attribution as the default and structural attribution as the exceptional claim requiring extraordinary evidence. The second is the access channel: when access barriers to validated structural analysis are high — because the analysis is technical, costly to obtain, or linguistically or culturally remote from the populations it describes — the structural explanation does not reach the populations with the greatest material stake in understanding it, even when it has been produced and validated. The third is the legitimacy channel: when validation infrastructure assigns low credibility to structural explanations relative to individual-attribution

accounts — through citation patterns, funding structures, or the composition of credentialing bodies — structural accounts enter public discourse in a position of epistemic disadvantage that compounds over time. Why infrastructure repair is a political project: Repairing the epistemic infrastructure is not primarily a technical problem of improving media quality or academic standards. It is a political project because the infrastructure in its current state is not malfunctioning — it is functioning in the interest of the actors who have most successfully shaped it. The distribution channels reach the audiences they are designed to reach; the validation processes assign credibility to the accounts that serve the interests of those who fund them; the access barriers exclude exactly the populations who would most challenge the accounts that circulate freely. Repairing the infrastructure requires changing the distribution of power over it — wh

ich is a structural project requiring the same counter-power accumulation that any structural change requires. The analytical implication: the V variable in the formula architecture is not fixed by governing narrative content alone. It is partly a function of the infrastructure state, which means that the same structural suffering can produce different transformation pressure in different infrastructure environments. Tracking V requires tracking the state of epistemic infrastructure, not only the content of official narratives.

1.8 — The Social Architecture of What Can Be Known

Part I established that epistemology has a social address: knowledge production is shaped by who controls the resources that make it possible, what questions those resources are directed toward, and what distribution of ignorance the resulting arrangement systematically produces. That analysis focused on knowledge production as a process — on how claims enter the world as knowledge. A complementary analysis is required at the level of infrastructure: the systems through which knowledge, once produced, is distributed, validated, stored, and accessed. This is not the same question. A society can have excellent conditions for producing accurate knowledge and still have epistemic infrastructure so degraded that the knowledge produced cannot circulate, cannot be verified, and cannot form the basis of coordinated action. Conversely, a society can have poor conditions for producing accurate knowledge but highly effective infrastructure for distributing and amplifying whatever is produced — including claims that are false. Epistemic infrastructure is the material system that determines what can be known socially, as distinct from what can be known individually or what has been produced academically. The analytical framework requires this distinction because the political consequences of epistemic conditions operate at the social level, not the individual level. An individual with access to accurate information about the structural conditions producing their deprivation is not thereby a political actor unless that information can be shared, verified against others' experience, aggregated into a common account, and distributed across the population that needs it. Each of these steps requires infrastructure. The destruction or degradation of that infrastructure — the production of conditions under which information cannot circulate, cannot be verified, or cannot be trusted — is therefore a form of political power that operates directly on the visibility variable V in the structural suffering formula. It does not require suppressing facts; it requires degrading the infrastructure through which facts become socially actionable. The distinction between knowledge production and knowledge infrastructure matters because the two can fail independently. A research environment can produce accurate knowledge about a structural problem — can generate documentation, analysis, and causal explanation of genuine quality — while the infrastructure through which that knowledge would need to travel to produce coordinated action is so degraded that the knowledge remains inert. The problem of what can be known is not only a problem of what gets produced. It is a problem of what, once produced, can circulate, accumulate, be recalled when relevant, and form the basis of shared understanding across the populations who need to act on it. The distribution problem: Knowledge infrastructure includes the systems through which produced knowledge reaches those who can use it: publicatio

n, translation, education, journalism, and the networks of reference and citation through which knowledge accumulates into a body of understanding rather than remaining dispersed in individual claims. Each of these

systems has structural features that shape distribution. Educational systems transmit the knowledge that curricula specify, which reflects the priorities of whoever controls curriculum design. Journalism transmits knowledge about events that meet newsworthiness criteria, which reflect the priorities of whoever controls those criteria. Translation makes knowledge available in the languages of economic and political power before the languages of populations with less power. The distribution of what is known across a population is therefore not a neutral reflection of what has been established — it is shaped by the same structural forces that shape all social distributions. The storage and recall problem: Even knowledge that is accurately produced and widely distributed can be lost — not through deliberate suppression, though that occurs, but through the ordinary dynamics of attention, institutional memory, and the half-life of documented claims that are not continually reproduced. Institutional memory — the capacity of an organization to retain and use knowledge produced in its past — degrades under conditions of high turnover, reorganization, and the competitive pressure to prioritize current production over accumulated understanding. Political memory — the capacity of a public to recall and apply relevant historical experience to current decisions — degrades faster than institutional memory, because it depends on active reproduction through journalism, education, and shared reference rather than on formal storage systems. The political consequences of degraded storage and recall are specific: they favor actors who benefit from forgetting, and they disadvantage populations whose organizing capacity depends on accurate historical self-understanding. The access problem is structurally distinct from both the distribution and storage problems. Access is the capacity of a specific person or population to retrieve knowledge relevant to their situation from available infrastructure. Even when knowledge has been produced, distributed, and stored, access to it requires time, literacy in the relevant technical language, institutional affiliation that provides entry to restricted databases, and sometimes physical proximity to the institutions where it is held. These requirements systematically exclude the populations most affected by the structural conditions being studied. The person navigating a welfare bureaucracy, a predatory lending system, or a labor market with few alternatives has the least access to the research about the structural conditions that shape their situation and the greatest need for it. The operational implication: the visibility variable V in the structural suffering formula is not only a function of narrative and media dyna

mics, as 8.4c analyzes. It is also a function of epistemic infrastructure. Suffering that is poorly documented, distributed, stored, or accessible produces lower V values independent of whether governing narratives actively suppress it. Infrastructure failure and active suppression both reduce V , but they require different interventions. Improving documentation and distribution is not the same as contesting the narratives that reframe structural causation as individual failure — and both are required for V to rise to the level at which coordinated political response becomes possible.

1.9 — Epistemic Authority: How Claims Acquire Social Standing

Not all claims that enter social circulation carry equal standing. Epistemic authority — the social capacity to make claims that others treat as reliable without independent verification — is a material resource distributed unequally and produced by institutional position, not by the accuracy of the claims made. This is not a cynical observation about the irrelevance of truth. It is a structural observation about how verification costs operate at scale. In a society of any complexity, direct verification of most claims is not available to most people. The social alternative to universal verification is trust in sources whose reliability has been established by track record, institutional position, or social proximity. This is a rational adaptation to the cost structure of knowing — and it is the mechanism through which epistemic authority becomes a form of power independent of the accuracy it is supposed to track. Epistemic authority follows institutional power in predictable ways. Actors with institutional positions that confer authority — state agencies, credentialed professions, established media, academic institutions — have a structural advantage in determining what claims receive social standing regardless of their accuracy, because the social mechanisms for assessing

credibility are calibrated to institutional position rather than to independent verification. This creates the conditions for a specific form of institutional capture: the capture of epistemic authority by actors whose material interests are served by particular claims receiving social standing. The captured institution does not need to produce false claims directly. It needs only to systematically amplify claims that serve its interests while applying disproportionate skepticism to claims that threaten them, using the institutional authority it has accumulated to shape what social consensus accepts as established. The relationship between epistemic authority and the eight diagnostic questions in 11.0 is direct: question 2 asks whether an arrangement permits accumulation of economic power sufficient to produce institutional capture. The capture of epistemic institutions — media, regulatory bodies, research funding, educational systems — is one of the primary mechanisms through which economic power translates into political durability. An arrangement that produces highly concentrated economic power will, given sufficient time, produce captured epistemic infrastructure that systematically misrepresents the consequences of that concentration. The degradation of epistemic infrastructure is therefore not an independent development that happens alongside economic concentration — it is, in most historical cases, produced by it. The gap between the accuracy of a claim and its epistemic authority — its social capacity to shape what others believe and act on — is one of the most consequential structural features of complex knowledge systems. It is not a failure of those systems so much

as an unavoidable consequence of their scale: in a society where no individual can directly verify more than a small fraction of the claims that bear on their decisions, the social machinery for assigning weight to claims is a functional necessity, not an optional feature. The question is not whether such machinery exists but what its structural properties are — specifically, what biases it introduces and what distributions of trust it produces. The production of epistemic authority: Epistemic authority is primarily produced through institutional affiliation and credentialing: the doctorate, the professional license, the editorial position, the institutional brand. These credentials function as pre-verification devices — they allow others to rely on claims without verifying them directly, on the grounds that the credentialing process has already done the relevant checking. The epistemic value of this system depends entirely on the quality of the credentialing process: whether it actually selects for accuracy and rigor, whether it is independent of the interests that would benefit from specific kinds of inaccuracy, and whether it is robust to the capture dynamics that affect all institutional processes. All three of these conditions hold imperfectly and fail differentially across domains and contexts. The decoupling of authority from accuracy: The most important structural feature of epistemic authority as a social resource is that it can become decoupled from accuracy — maintained through social position long after the track record that originally justified the position has deteriorated, or accumulated before it has been earned, through networks of mutual credentialing that function as self-reinforcing prestige systems. Expert consensus can be wrong, and its capacity to be wrong while maintaining authority is not merely a theoretical possibility — it is a recurring historical pattern in medicine, economics, and social science. Identifying this is not a license for dismissing expert consensus: the alternative to imperfect credentialed knowledge is not better knowledge, it is usually worse knowledge packaged to look like common sense. The appropriate response is calibrated skepticism that weighs the track record of specific expertise in specific domains against the alternative sources available. The political economy of epistemic authority is not neutral. Actors with institutional resources can produce epistemic authority artificially: by funding research in credentialed venues that reaches favorable conclusions, by establishing think tanks and policy institutes that carry institutional markers of credibility without the independence from interested funding that credibility is supposed to guarantee, and by occupying seats in public debate that treat institutional affiliation as a proxy for reliability. These operations work because the social machinery for assigning authority cannot easily distinguish between authority that was earned through demon

strated accuracy and authority that was manufactured through institutional positioning. The populations most affected by the resulting distortions are those with the least access to independent verification — which is precisely the population whose dependence on the social machinery for assigning authority is greatest. The

counter-authority problem: communities that are systematically excluded from credentialed knowledge production develop their own epistemic authority systems — their own networks of trusted sources, their own criteria for credibility, their own standards for what counts as reliable information. These systems are not less rational than credentialed ones; they are rational responses to the experience that credentialed systems consistently produced knowledge that did not reflect their conditions accurately. The fragmentation of epistemic authority across populations with different trust architectures produces coordination failures that have direct political consequences: populations cannot act together on shared structural conditions when they cannot agree on what the conditions are, and they cannot agree on what the conditions are when they have no shared epistemic authority they both recognize. The diagnostic for analysts using this framework: when a claim carries high epistemic authority by virtue of institutional source, ask independently whether the source's track record in the relevant domain justifies that authority, and whether the institutional conditions under which the claim was produced are ones that favor accuracy or ones that create systematic pressure toward particular conclusions. High authority from a source whose track record is poor, or whose institutional conditions create strong incentives toward specific kinds of error, should be weighted accordingly — which means it may receive less weight than a lower-authority source whose track record and incentive conditions are better.

1.10 — Information Asymmetry as Structural Power

Information asymmetry — the condition in which different actors in the same arrangement have access to materially different information about that arrangement — is one of the most consequential structural features of complex institutional systems. It is analytically distinct from the distribution of knowledge production resources analyzed in 1.7, and from the epistemic authority distribution analyzed in 1.9. Where 1.7 concerns what is produced, and 1.9 concerns what receives social standing, information asymmetry concerns who has access to what is known: the structural gap between what actors inside an arrangement can observe and what actors outside it, or below it, can observe. In the context of the framework's analysis of institutional arrangements, information asymmetry operates primarily in two directions. Vertical asymmetry describes the systematic advantage that actors at the top of institutional hierarchies have over those subject to those hierarchies in knowing what the institution is actually doing, what the consequences of its decisions are, and what alternatives were available. Vertical asymmetry is a precondition for the institutional capture analyzed in 7.3: capture requires that those being harmed by the captured institution cannot observe the capture in real time, because observable capture would trigger the response that makes it politically costly. Horizontal asymmetry describes the systematic advantage that some parties in a transaction or relationship have over others in knowing the relevant facts of that transaction. In economic arrangements, horizontal asymmetry between sellers and buyers, between employers and workers, or between lenders and borrowers is one of the primary mechanisms through which otherwise formally equal exchanges produce structurally unequal outcomes. The political significance of information asymmetry extends beyond individual transactions to the conditions for collective action. Counter-power accumulation — analyzed in 10.9 — depends on populations being able to recognize shared structural conditions and aggregate individual experiences into a common account. Information asymmetry directly attacks this capacity: if each person experiencing structural deprivation cannot observe whether others are experiencing the same conditions, and cannot verify whether those conditions are produced by the same structural causes, the coalition formation that transformation pressure requires is structurally blocked even when the material conditions for it are present. Reducing information asymmetry — through transparency requirements, public documentation, collective bargaining for information rights, or the production of shared knowledge infrastructure — is therefore not merely a technical intervention. It is a political intervention in the conditions for collective action. Information asymmetry is a power resource because it enables the actor with superior information to structure interactions in ways that benefit them

systematically, independent of the formal equality of the parties to those interactions. A borrower who does not understand the terms of a financial instrument and a lender who designed it are formally in a voluntary transaction. The asymmetry of information about the instrument's properties is what makes the voluntariness of the transaction formal rather than substantive: the borrower cannot adequately assess what they are agreeing to. The same structure applies across domains — employment contracts, insurance agreements, regulatory processes, healthcare decisions, legal proceedings — wherever the information required to evaluate an arrangement is concentrated in one party to that arrangement. The strategic production of asymmetry: Information asymmetry is not only a naturally occurring feature of complex systems. It is strategically produced and maintained by actors who benefit from it. Complexity in financial instruments, legal documents, and regulatory frameworks is often functional in the sense that it reflects genuine complexity in the underlying domain — but it is also often in excess of what the underlying complexity requires, and the excess serves a specific function: it raises the cost of independent evaluation for parties with less information, making them more dependent on the judgment of parties who have more. The systematic production of illegibility is a durable mechanism through which structural positions maintain themselves without requiring coercion. This connects directly to the power analysis in Part VI: complexity advantages the people who generate it. Asymmetry and institutional design: Most institutional systems are designed under conditions of information asymmetry between the designers and the populations they affect. Legislative drafting, regulatory rule-making, and organizational procedure-setting are all processes in which actors with concentrated information and resources shape the terms under which others operate. The populations most affected by these designs are frequently those with the least capacity to participate in the design process — not because they are formally excluded, though exclusion occurs, but because meaningful participation in complex institutional design requires resources — time, technical expertise, organizational capacity — that are distributed unequally. The result is institutional designs that systematically underweight the information held by affected populations, which tends to be concrete, contextual, and about the gap between how the institution is supposed to work and how it actually operates for people navigating it from positions of limited power. The epistemic asymmetry between affected populations and institutional designers is not random in its direction. People operating within institutions from positions of limited power develop detailed, accurate knowledge about how those institutions actually function: what the formal rules are and what happens when you try to use them, what th

e gaps are between stated procedure and actual practice, what information is required to navigate the institution successfully that the institution does not provide, and what the consequences of the institution's failures are for those who bear them. This knowledge is real, consequential, and systematically excluded from institutional design processes because it is not credentialed, not documented in accessible form, and not present in the rooms where design decisions are made. The framework that takes 1.7's analysis of knowledge production seriously must account for this distribution of uncodified practical knowledge as a genuine epistemic resource. Information asymmetry analysis raises a question that the framework must address directly: what, if anything, justifies a claim to privacy — to zones of life and information that institutions cannot access even when access would improve their analytical capacity or enforcement effectiveness? The framework's answer is structural rather than deontological. Privacy is not justified as an intrinsic right that exists independently of its consequences. It is justified as a structural condition for the kind of agency that genuine development requires. A person whose deliberations, relationships, experiments, and failures are permanently visible to institutional actors adjusts their behavior toward what those actors will approve of rather than toward what they genuinely need. The chilling effect on thought, association, and development that follows from permanent visibility is not a hypothetical harm — it is the mechanism through which surveillance converts the reduction of information asymmetry into a form of structural suffering. The framework therefore supports privacy not as an absolute claim but as a structural condition: the degree of privacy appropriate to any domain is determined by whether access to information in that domain is necessary for the institutional function being performed, whether the institutional function itself satisfies the

framework's criterion, and whether the population subject to access has genuine standing to contest it. Privacy claims are strongest where institutional access produces chilling effects on development without corresponding gains in floor conditions; they are weakest where institutional actors with genuine accountability need information about arrangements that concentrate power or produce structural suffering. The asymmetry is significant: privacy protects individuals against institutions; transparency requirements protect populations against concentrated power. Both are applications of the same structural principle — that information distributions shape the conditions for genuine development and must be evaluated by what they produce for those subject to them. The visibility variable again: information asymmetry is one of the primary mechanisms through which V is suppressed. Structural suffering that is experienced by populations without the resources to document

it, translate it into credentialed form, and circulate it in channels that reach decision-makers produces low V values not because the suffering is hidden through active suppression — though that also occurs — but because the information infrastructure required to make it visible does not reach from where the suffering occurs to where the decisions are made. Information asymmetry and epistemic infrastructure degradation interact: together they produce conditions in which the populations bearing the costs of structural arrangements are systematically less able to make those costs politically visible than the actors who benefit from the arrangements. Analytical procedure for information asymmetry: when analyzing an institutional arrangement, explicitly ask: who has information about how this arrangement works that others lack? How was that asymmetry produced — is it a necessary feature of the domain's complexity or is it in excess of what the complexity requires? Who benefits from its maintenance? What information do the populations most affected by the arrangement have about its actual operation that institutional designers do not? And what would change in the analysis if that information were incorporated? The framework's analysis of epistemic authority (1.9) and information asymmetry (1.10) identifies the structural conditions under which knowledge claims acquire social standing and the structural advantages that particular positions generate. A gap in this analysis: when expert knowledge and situated knowledge — the knowledge of affected populations about their own conditions — conflict, the framework has no procedure for adjudicating between them. This is not an unusual situation. Expert epidemiological analysis of a community's health conditions can conflict with the community's own account of what is causing harm. Economic models of an arrangement's efficiency can conflict with the experience of those operating within it. Technical assessments of a project's safety can conflict with local knowledge about that specific terrain or population. The structural tendency is for expert knowledge to win these disputes, not because it is more accurate but because it carries the institutional authority that Section 1.9 identifies as socially constructed through processes that systematically exclude the populations most affected. Evidence hierarchy for conflicting knowledge claims: when expert knowledge and situated knowledge conflict, evaluate using four criteria in sequence: (1) Predictive track record — which type of knowledge claim has better predicted actual outcomes in comparable previous cases? Expert models that have systematically underestimated or misidentified harms in similar contexts carry reduced evidentiary weight, regardless of methodological sophistication. Situated knowledge that has accurately predicted specific local outcomes carries increased weight. (2) Independence from interests — which claim is less directly produced by actors

with material interest in a specific resolution? Expert knowledge funded by actors with interests in favorable assessment carries reduced weight proportional to the directness of that interest relationship. Situated knowledge produced by populations bearing the costs of the arrangement is not interest-free, but its interests are more likely to be aligned with accurate harm identification. (3) Accessibility for verification — which claim can be evaluated by a broader range of analysts? Expert knowledge in inaccessible technical domains concentrates evaluative authority in precisely the population most likely to share the structural position that produced the claim. Situated knowledge expressed in accessible terms is verifiable by more independent assessors. (4) Representativeness — whose experience does each type of knowledge claim represent? Expert knowledge systematically underrepresents the experience of populations that are not the primary subjects of

the expert's professional attention. When expert and situated knowledge conflict, the burden of justification falls on the expert claim to demonstrate why the affected population's account of their own conditions is less accurate than the external assessment.

1.11 — Trust Infrastructure and Its Collapse

This section identifies a specific threshold effect in epistemic infrastructure that is distinct from ordinary misinformation. The concern is not that false claims circulate — false claims have always circulated. The concern is that synthetic media and coordinated authenticity manipulation can attack the category of documentary evidence itself: not this video is false, but any video might be false, which destroys the evidentiary function of documentation for everyone, regardless of the specific content being contested. When that threshold is crossed, the asymmetry that favors power over accountability is structural: those who control narratives without documentation can claim anything; those who would contest them lose the primary tool — documented evidence — through which contestation could succeed. PMN identifies this as a visibility suppression mechanism operating at the infrastructure level, not the content level. Epistemic infrastructure is not only a technical system for distributing information. It is a trust system: a set of social relationships, institutional arrangements, and shared practices that allow actors to extend confidence beyond what they can directly verify. Trust infrastructure is what makes complex social coordination possible at all — without it, every transaction requires independent verification of every claim, which raises the cost of coordination past what most valuable activities can sustain. Trust infrastructure in this sense includes not only media and academic institutions but the informal social practices through which credibility is assigned: professional norms, peer review, editorial standards, the public record, the accumulated track records of specific sources across specific domains. Trust infrastructure is vulnerable to a specific form of attack that is more efficient than the direct production of false claims: the production of generalized uncertainty. An actor who cannot make the infrastructure produce favorable claims can instead degrade the infrastructure's ability to make any claims credible. By systematically introducing doubt about the reliability of sources, by producing enough conflicting claims that verification seems impossible, by delegitimizing the epistemic institutions whose function is to distinguish reliable from unreliable claims — an actor can produce a condition in which social knowledge collapses into ambient uncertainty. In that condition, false claims do not need to be believed; they only need to be plausible enough to prevent the true claims from being actionable. The political benefit is not the replacement of accurate claims with inaccurate ones. It is the dissolution of the conditions under which any claim can be evaluated publicly. Trust collapse is threshold-governed in the same way that legitimacy depletion is threshold-governed (6.6). Trust infrastructure can absorb significant stress — individual institutional failures, specific episodes of misinformation, documented cases of e

pistemic capture — without collapsing, because residual trust in other parts of the infrastructure compensates. Below the threshold, each failure is an anomaly within a generally functioning system. Above the threshold, each failure confirms a general pattern of unreliability, and the system reorganizes around ambient distrust. In the ambient distrust condition, social knowledge fragments: different communities develop incompatible factual frameworks, not because they have access to different evidence but because the infrastructure through which evidence is evaluated and shared no longer functions across community boundaries. The political consequence is the collapse of the common factual ground that makes cross-community coalition formation possible — which is the same result that manufactured fracture lines (10.9) produce through narrative means. The most recent developments in epistemic infrastructure disruption extend the mechanism described above into territory that the existing analysis of misinformation does not fully capture. Synthetic media — algorithmically generated images, audio, and video that are indistinguishable from authentic documentation by ordinary verification methods — does not merely add false claims to the information environment. It destroys the evidentiary function of documentation itself. When any video can be plausibly dismissed as synthetic, the capacity to establish shared factual ground

through recorded evidence collapses — not because specific fabrications succeed, but because the mere possibility of fabrication makes all documentation contestable. The epistemic infrastructure damage is not in the spread of specific false content. It is in the erosion of the category of authentic documentation as a resource for adjudicating disputes about what happened. Movements that depend on documented evidence of atrocity, of institutional failure, or of official deception to build the political pressure required for accountability face an environment in which the evidentiary function of that documentation has been structurally weakened — regardless of whether any specific document they present has been fabricated.

Algorithmic Governance: Visibility Suppression at the Infrastructure Level

Epistemic infrastructure analysis requires extension to address a form of visibility suppression that operates at a different level from the mechanisms analyzed above. Propaganda, censorship, and manufactured consent operate on content: they determine which claims are asserted, which narratives are promoted, which evidence is suppressed. These mechanisms have been the primary subject of critical analysis of information systems throughout the twentieth century, and the framework's analysis of governing narrative and epistemic infrastructure captures their operation. But algorithmic governance — the organization of information flows by recommendation systems, engagement optimization engines, and attention allocation mechanisms deployed at platfo

rm scale — operates not primarily on content but on the architecture of attention itself. It determines not what claims are made but which claims are seen, by whom, in what sequence, with what surrounding context, and at what emotional register. This is a structurally different form of visibility suppression, and its analysis requires different conceptual tools. The mechanism of algorithmic visibility suppression is the optimization of engagement metrics as a proxy for value. Recommendation systems trained to maximize time-on-platform, click-through rate, or emotional response do not suppress information by removing it. They suppress it by making it structurally less likely to be encountered through the normal operation of the system — by routing attention toward content that produces stronger engagement signals, which tends systematically toward content that activates anxiety, outrage, and tribal identification more reliably than content that requires careful attention and produces considered response. Information about structural conditions — the slow processes of institutional capture, ecological degradation, intergenerational transmission of disadvantage — is systematically disadvantaged in this architecture not because any actor has decided to suppress it but because it does not compete effectively for attention against content optimized for immediate emotional response. The suppression is the aggregate output of a system optimized for a different objective, not the intention of any specific actor. This structural feature of algorithmic information architecture has specific consequences for the framework's analysis of counter-power formation. The coalition width variable W and the fragmentation variable I_f are both directly affected by the information environment in which potential coalition members encounter each other and develop shared analysis. Counter-power formation requires the construction of shared understanding of structural conditions across populations who do not already share those conditions — which requires that structural analysis travel across the attention architectures of people who encounter it incidentally, not as committed seekers. Algorithmic systems that route attention toward high-engagement content systematically disadvantage this cross-sectional diffusion: structural analysis that requires careful reading and rewards considered attention is consistently outcompeted by content that activates immediate emotional response, and the populations most subject to the structural conditions that analysis describes are often those most dependent on platform-mediated information environments with the least alternative access to structural analysis. The result is a systematic increase in I_f — fragmentation within potential coalitions — produced not by ideological difference but by the attention architecture in which those potential coalitions are trying to form. The relationship between algorithmic governance and the meaning-co

llapse configuration identified in section 4.3 deserves specific attention. Meaning-collapse — the condition in which governing narratives have lost their organizing capacity and populations are available for rapid mobilization by whoever offers the most compelling orientation fastest — is both produced and exploited by algorithmic attention architecture. The high-engagement content that algorithmic systems surface most

effectively tends to offer strong orientation without demanding much from the people it orients: clear enemies, simple explanations, vivid emotional stakes, and group identity that requires no prior commitment or analysis to adopt. Algorithmic systems did not create the meaning deficit that produces meaning-collapse configurations, but they accelerate the transition into those configurations by consistently routing attention toward content that consumes rather than builds meaning infrastructure — and they accelerate the exploitation of those configurations by making the rapid mobilization of disoriented populations technically easier and more precisely targeted than any previous information architecture. Algorithmic governance as a V-suppression mechanism — diagnostic indicators: (1) Structural conditions that produce widespread material harm are consistently underrepresented in platform-mediated information environments relative to their significance, not because they are suppressed but because they do not compete effectively in engagement-optimized architectures. (2) Fragmentation within potential counter-power coalitions is increasing in ways that correlate with algorithmic information exposure and do not correspond to actual material interest divergence. (3) The transition rate into meaning-collapse configurations is accelerating in contexts with high platform dependence for political information. These indicators do not require evidence of intentional suppression. They require evidence of systematic structural disadvantage in attention allocation. The distinction matters for analysis: algorithmic V-suppression cannot be addressed by identifying and pressuring bad actors, because the suppression is a structural output of systems designed for different objectives. It requires changing the objective function — the metrics that recommendation systems optimize — which is a structural intervention in the governance of information infrastructure, not a content moderation question.

1.12 — Epistemic Infrastructure and the Visibility Variable

The connection between epistemic infrastructure and the structural suffering formula runs directly through V — the visibility suppression variable. Section 8.4c analyzed how meritocratic narrative suppresses V by recategorizing structural causation as individual causation. Epistemic infrastructure operates on V through a different but complementary mechanism: it determines whether the information required to recognize structural causation can circulate at all. A fully functional epistemic infrastructure with a dominant meritocratic narrative will produce high V through interpretive recategorization — the facts are available but are read as evidence of individual difference. A degraded epistemic infrastructure produces high V through a different mechanism: the facts are not available, not trusted, or not aggregated into a form that makes structural causation recognizable. Both produce the same political outcome — populations experiencing structural deprivation cannot mobilize around shared structural conditions — through structurally distinct mechanisms that require distinct responses. The practical implication is that any strategy for reducing V must attend to both mechanisms simultaneously. Producing accurate information about structural conditions is insufficient if the epistemic infrastructure through which that information would circulate is degraded or captured. Building epistemic infrastructure — the material conditions for credible, shared, verifiable knowledge about structural conditions — is therefore as much a part of the political project as the material interventions on R and B that the framework's economic analysis recommends. Movements that produce accurate analysis of structural conditions but operate in environments of degraded epistemic infrastructure find that their analysis circulates only within communities already convinced, and cannot penetrate the ambient uncertainty that makes structural claims seem equivalent to any other claim. Epistemic infrastructure is a material condition with direct consequences for transformation pressure. Where it is functional, accurate information about structural conditions can circulate, be verified, aggregate into common accounts, and reduce V by making structural causation visible. Where it is degraded or captured, V remains high even when material conditions for structural suffering are severe and well documented — because the social mechanism through which documentation becomes political recognition has been disabled. Building and defending epistemic infrastructure is not a precondition for political action; it is political action, directed at one of the primary mechanisms through which existing arrangements maintain themselves against the

transformation pressure their consequences generate. Stigma as Visibility Suppression: A Structural Case Study The visibility variable *V* operates through many mechanisms — media ownership, algorithmic filtering, legal restrictions on disclosure, spatial

segregation of suffering populations — but stigma against specific populations represents one of its most structurally complete instances. Stigma does not merely make suffering less visible to external observers. It makes suffering invisible to the institutional pathways that could address it, because the populations bearing stigmatized conditions have strong rational incentives not to enter those pathways. The result is a self-reinforcing suppression loop: the condition is invisible because disclosure is too costly, and the cost of disclosure remains high because the invisibility prevents the development of institutional alternatives to enforcement-only responses. This mechanism is analyzed in detail in section 3.4c in the context of dispositional harm prevention — the case of individuals who carry a risk of harming others but have not yet acted. That case is worth flagging here as the structurally cleanest illustration of how stigma functions as *V*-suppression: the harm produced by the invisibility cascade is both measurable and preventable, which makes the institutional design implications unusually clear. The same mechanism operates wherever stigma attaches to conditions that individuals have strong incentives to conceal — mental illness, addiction, debt, precarious immigration status — and the analysis of epistemic infrastructure must account for stigma as a systematic visibility suppression mechanism rather than a cultural attitude whose material consequences are incidental. Freire and the Pedagogy of Visibility The analysis of visibility suppression as a political mechanism operating on the *V* variable identifies the problem with structural precision: *V* is high when populations experiencing structural suffering cannot recognize the structural causation of their conditions. What that analysis does not fully develop is the specific mechanism through which *V* is reduced — how suppressed structural visibility becomes legible to those experiencing it. Paulo Freire's *Pedagogy of the Oppressed* (1968) is the most sustained and empirically grounded analysis of this mechanism available in the political pedagogical tradition, and it is directly applicable to PMN's framework. Freire identifies the primary mechanism of *V* maintenance as what he calls the 'banking model of education': the pedagogical practice of depositing pre-formed knowledge into passive recipients, which reproduces the interpretive frameworks of the dominant culture — including the frameworks that attribute structural suffering to individual failure — without engaging the lived material experience of those receiving the knowledge. The banking model produces high *V* not through direct suppression of information but through the provision of an interpretive apparatus that makes structural causation invisible: recipients acquire knowledge about the world through frameworks generated by those who benefit from the arrangements producing the suffering, which means they acquire the meritocratic n

arrative (8.4c) as the natural way of understanding their own conditions. The alternative Freire develops — *conscientização* (consciousness-raising through dialogical pedagogy) — is a specific method for reducing *V*: generating analysis from the material experience of those who suffer it, rather than depositing analysis generated from outside. The PMN application is direct. Freire's pedagogical analysis maps onto PMN's structural analysis of *V* in three ways. First, banking education is a specific mechanism of *V* production that operates through legitimate educational institutions rather than through direct information suppression — which makes it more durable and less visible than censorship. Second, *conscientização* is a specific counter-*V* strategy that operates through the recognition of shared material experience rather than through analytical instruction imposed from outside — which is why Freire's method requires dialogue rather than lecture. Third, the organic intellectual concept (from Gramsci, incorporated in PMN's agent ecology) maps onto Freire's facilitator role: the person who works within a community to help it generate its own structural analysis from its own material experience, rather than importing analysis generated elsewhere. Freire's contribution to PMN: *conscientização* is the most developed available account of how *V* is specifically reduced through pedagogical practice — how populations experiencing structural suffering develop the capacity to recognize structural causation rather than attributing conditions to individual failure. This maps directly onto the claim made here that building epistemic infrastructure is political action, not preparation for

it: Freire demonstrates in educational practice what PMN's structural analysis implies theoretically. The limitation of Freire's approach that PMN's framework identifies: conscientização reduces V without directly addressing R or B. Visibility of structural causation is necessary but not sufficient for transformation; it must be combined with organizational capacity and structural pressure on the arrangements producing the suffering.

II: What Exists: The Ontological Foundation

2.1 — Reality as Primary

Reality is primary and independent of our perception of it — this is the foundational axiom of this framework, adopted as a starting point rather than derived from anything prior. The convergence argument in the paragraph below provides a reason to prefer this axiom over its idealist alternative, but it does not prove it: every ontological starting point is, by definition, not derived from a prior starting point. Naming this honestly is not weakness; it is the application of the framework's own epistemological standards to itself. What follows from this axiom — the extension of historical materialism's analysis deeper than the economic — inherits whatever vulnerability the axiom carries. Historical materialism held that what people believe is ultimately shaped by the material circumstances of their lives, particularly their position in the relations of production. That claim is correct, but it stops one level short: the material conditions that matter most extend below economic organization to the biological conditions of the organisms who hold those ideas, who organize that production, who are capable of suffering from the arrangements that result. The alternative — the idealist position that consciousness is primary, that the material world is constituted by or dependent on mind — is not dismissed here as obviously wrong. It is rejected for a specific reason: it cannot provide a non-circular account of why different minds perceive the same constraints. If consciousness constitutes reality, the convergence of independent observers on the same material limits requires explanation. The materialist position that reality precedes and constrains perception explains this convergence without circularity: independent observers perceive the same limits because the limits exist independently of any observer. This is the sense in which materialism is more epistemologically honest, not more metaphysically certain. But the convergence argument is a reason to prefer the realist axiom, not a proof of it: a sophisticated idealist can respond that the convergence itself is constituted by the shared structure of consciousness rather than by mind-independent constraints. The framework acknowledges this and proceeds on the realist axiom as the more productive starting point. Contemporary analytic idealism — most rigorously developed by Bernardo Kastrup — presses the metaphysical question further than social-scientific idealism does. Kastrup's position is not that ideas shape social outcomes (the error 2.5 identifies) but that consciousness is the ontological substrate of everything, including what appears as the physical world. This is a coherent and technically sophisticated position that the convergence argument above does not definitively refute. PMN's response is not to dismiss it but to note two things. First, 2.3 holds open the question of what underlies the material level, including the possibility that something consciousness-like is ontologically prior; PM

N's agnosticism there is genuine. Second, the framework's practical scope — analyzing the structural production of suffering in social arrangements — does not require resolution of the hard problem of consciousness. Whether the suffering of a malnourished child is ultimately constituted by a field of mind or by blind material processes, it is still suffering, it is still produced by identifiable structural arrangements, and it still requires the same institutional response. The metaphysical question remains open. The analytical and ethical task does not wait for it. Beliefs matter — they shape behavior, and behavior shapes material conditions. But they do not determine those conditions directly, and conditions that contradict beliefs do not disappear because the beliefs are sincerely held. An institutional arrangement that produces structural suffering continues to produce it regardless of the narrative its operators maintain about what it is doing. Analysis that attends to what arrangements actually produce — measured in the conditions of the populations they affect — rather than to what they are believed to produce will reach different conclusions than analysis that takes self-descriptions at face value. That difference is not incidental. It is the point of beginning from material conditions rather than from ideology. The commitment to mind-independent reality is prior to the commitment to materialism. Idealist traditions — those that locate primary reality in mind, consciousness, or discourse — do not typically deny that something exists independently of any particular mind; they deny that material existence is ontologically fundamental. The disagreement is not about whether there is a world but

about what kind of thing the world fundamentally is. What this framework asserts is that the world is not constituted by our conceptual schemes, that its features do not depend on our descriptions of them for their existence, and that the test of our claims about the world is ultimately their correspondence to how the world actually is — not their coherence with each other, not their pragmatic utility in isolation from accuracy, not their fit with prior theoretical commitments. This requires a postulate stated explicitly: Even if Kastrup's position were confirmed in full — even if it were established beyond reasonable philosophical doubt that what appears as physical matter is the extrinsic appearance of a universal field of consciousness, and that individual biological organisms are dissociated alters of that field — PMN's analytical and ethical framework would continue to operate without modification. The reason is structural: the framework's evaluative criterion is grounded not in the metaphysical substrate of suffering but in its phenomenological reality for the entity undergoing it. Hunger experienced by a dissociated alter of universal consciousness is still hunger. A bullet experienced by the same alter is still a bullet. The deprivation of reso

urces, the betrayal of institutional protections, the suppression of the visibility of harm — all of these produce real constraints on real becoming for the entities subject to them, regardless of what those entities are ultimately made of. The hard problem of consciousness does not dissolve the practical problem of suffering: it leaves it exactly where it is. What PMN's agnosticism requires is not that the metaphysical question be resolved before the analytical and ethical work can proceed — it requires that the metaphysical question not be permitted to paralyze the analytical and ethical work while it remains unresolved. That is the precise sense in which the framework's scope is defined by the phenomenology of constraint rather than by the ontology of its substrate. No revision of our account of what the universe is fundamentally made of changes what it is like to be unable to meet the biological requirements for continued existence. That invariance is the foundation on which the framework stands, not a prior metaphysical commitment about matter. The practical consequence of this commitment is more significant than it might appear in abstract statement. A framework that treats reality as socially constructed — as constituted through the shared conceptual practices of communities — has difficulty generating a principled account of why some constructions are better than others that does not itself depend on a prior commitment to mind-independent standards. It can appeal to internal coherence, to pragmatic success, to the preferences of a community. But all of these appeals are ultimately self-referential: they evaluate constructions by standards that are themselves constructions. The realist commitment avoids this regress by anchoring evaluation to something that is not itself a construction: the material conditions that affect real organisms in ways they can suffer from regardless of how those conditions are described. The relationship to social construction: The realist ontology does not deny that much of what shapes human social life is socially constructed — that legal systems, economic arrangements, cultural norms, and institutional structures are human-made and could be made differently. It denies that social construction extends to the material substrate on which social arrangements operate. Money is a social construction; the hunger of the person who cannot afford food is not. The law of contract is a social construction; the physical consequences of living under a bridge in winter are not. Social construction claims are most powerful and least contestable when they concern the categories and classifications through which material reality is organized and made actionable. They are least defensible when they are extended to the material substrate itself — to the biological conditions that give the categories their stakes. The relationship between realism and the epistemic commitments in Part I requires explicit statement. Realism about th

e external world does not entail that our access to that world is unmediated or that our descriptions of it are transparent. Part I established that knowledge production is shaped by material conditions in ways that produce systematic biases. The realist commitment does not contradict this — it grounds it. The reason the biases in knowledge production matter is that they distort our understanding of a reality that exists independently of those distortions. If reality were constituted by our understanding of it, the distortions would constitute rather than misrepresent it, and the critical project of identifying and correcting them would lose its foundation. Realism and epistemic fallibilism are not in tension; they are mutually supporting. The

ontological commitment here is to a layered realism: material reality is real and independent of our descriptions (this is the foundational axiom, acknowledged as such in the opening of this section); social structures built on that substrate are real as causal forces even though they are humanly produced and could be otherwise; and the evaluative criterion — the minimal anchor — is grounded in features of that material reality (the capacity to suffer, the requirements of biological function) that are not themselves socially constructed. This layering allows the framework to take social construction seriously at the level of institutions and arrangements while maintaining the realist ground that makes critique of those constructions more than preference. The layering also identifies where different levels of challenge hit: challenges to the realist axiom require engaging 2.1's convergence argument; challenges to the evaluative criterion require engaging 3.0f and 4.1's design-choice defense; challenges to the institutional analysis require empirical engagement with specific claims in Parts VI–XII.

2.2 — Adaptive Materialism

'Adaptive' in this framework's usage carries no positive valence — it does not mean 'better-designed,' 'more flexible,' or 'improved.' It describes only the mechanism: systems change in response to pressure from their conditions, without any guarantee that the change moves in any particular direction. The history of what adaptation has selected is a history of both remarkable human achievement and of arrangements that produced extraordinary suffering for centuries precisely because they were effective at surviving under the conditions they operated in. The name 'adaptive materialism' is intended to distinguish this account of how reality moves from the dialectical account — not to suggest that adaptation produces good outcomes. Dialectical materialism attempted to capture the movement of reality through the concept of the dialectic: the idea that every state of affairs contains its own contradiction, that the contradiction drives development toward a synthesis, and that this process is the engine of historical change. The genuine insight here is that reality is not static — that any adequate analysis must account for movement and transformation, not just describe a frozen state. The problem is that the dialectic is not a description of how reality actually moves. It is a schema imposed on reality's movement — a way of narrating change that fits some cases and is forced onto others. The adaptive materialism of this framework takes the genuine insight while discarding the schema. Reality moves. Analysis must move with it. The mechanism of movement is not guaranteed contradiction resolving into synthesis. It is the continuous interaction between organisms and their conditions — the pressure of environment on the forms of organization that exist within it, the selection of some forms and the elimination of others, the gradual accumulation of adaptive responses that become the material starting point for the next round of pressure. This is a less elegant account than the dialectic. It is also more accurate. The methodological shift is real. If the dialectic guaranteed that contradictions would be productive — that they were moments in a larger development — then the analyst's job was to locate a given situation within the dialectical process. If reality moves through the more complex and less guaranteed mechanism described here, the job changes: identify the actual pressures on actual arrangements, trace their consequences, revise when the consequences diverge from what was expected. The method loses the confidence that comes from knowing in advance how the process will resolve. What it gains is honesty about what it actually knows — and honesty, for a method that claims to track reality, is more valuable than confidence that cannot be cashed against evidence. One consequence of replacing the dialectic with adaptive dynamics directly threatens the framework's evaluative ambitions. Evolutionary selection does not optimize for good outcomes. It selects

for fitness — for whatever configuration survives and reproduces under existing conditions. Slavery was fit. Authoritarian extraction was fit. Arrangements that produced extraordinary structural suffering for large populations were fit for centuries, because fitness is always relative to the specific conditions under which selection operates, and those conditions included the suppression of the alternatives that might have outcompeted them. If adaptive materialism only describes how systems change — if it only traces the logic

of selection without providing criteria for evaluating what is selected — it collapses into a sophisticated version of 'whatever survives is justified.' That conclusion is not only morally inadequate. It is also analytically wrong, because systems that are fit under specific conditions are regularly replaced by systems fit under different conditions, and the replacement is not always improvement by any standard the framework holds. Feudalism was replaced by capitalism not because capitalism was better by the minimal anchor but because it was more effective at certain forms of coordination and extraction. The successor to capitalism, whatever it is, will be selected by the same mechanism. Adaptive materialism describes the mechanism of change, not the criterion of evaluation. The criterion comes from the biological foundation: arrangements that create the conditions under which genuine development is possible — of which minimizing structural suffering is the necessary diagnostic — are more sustainable across a wider range of conditions, because they maintain the cooperative substrate on which all durable social arrangements depend. This is not a guarantee. It is a probabilistic claim about what tends to survive across the full range of conditions systems face over time, rather than under the specific conditions any given arrangement is optimized for. The direction provided by the biological foundation is not teleological — it does not guarantee that history moves toward it. It is evaluative and strategic: the minimal anchor specifies what is worth working toward, and adaptive materialism specifies the mechanism through which that work is possible. Neither functions without the other. Description of how systems change, detached from any criterion of what change is for, produces a sophisticated account of process that cannot say whether any process is better than any other. A criterion of evaluation, detached from any theory of how arrangements actually change, produces aspiration without traction. What the two provide together is something more honest than either could provide alone: an account of how material reality moves and what movement is worth working for, without the false comfort of guaranteeing that they converge. On Reductionism: Ontological Emergence and the Independence of Social Causation A clarification that the framework requires concerns the relationship between its biological grounding and the social level o

f analysis that most of its content operates at. Grounding the framework in biology might appear to imply that social explanations are ultimately reducible to biological ones — that what happens at the level of institutions, narratives, and power arrangements is finally explicable by reference to the biological properties of the organisms who produce them. This is not the framework's position, and the distinction matters for what kind of analysis the framework licenses. The distinction is between epistemological emergence — where higher-level phenomena are in principle reducible to lower-level ones but are too complex to derive in practice — and ontological emergence, where higher-level structures possess genuine causal powers that are not present at the level of their components and cannot be derived from them even in principle. This framework takes the second position. Social structures, institutions, and governing narratives are not merely aggregations of biological behaviors. They possess causal powers — the power to shape what individuals can think, want, and do — that are not reducible to the properties of any individual organism or any collection of organisms considered independently of their structural relationships. (Bhaskar 1975) This position has a specific name in the philosophy of science: critical realism, developed most systematically by Roy Bhaskar. The central argument is that reality is stratified — that different levels of organization operate with different mechanisms and causal powers, and that reducing higher-level phenomena to lower-level ones destroys the explanatory resources needed to understand them. A strike is not explicable by reference to the biochemistry of the workers who participate in it, even though every striker is a biological organism whose participation has biochemical conditions. The social mechanism — the organization of grievance, the coordination of action, the structural relationship between workers and capital — operates at a level that has its own causal structure. It shapes biological behavior rather than being explained by it. The biological foundation of this framework is therefore a grounding claim, not a reduction claim. Biology specifies the conditions that social arrangements must meet — the floor below which no arrangement can push the organisms it coordinates without destroying their capacity for cooperation and development. It does not specify the content of social arrangements above that floor. The immense variety of institutional forms, cultural configurations, and

political arrangements that human societies have produced — all within the same biological constraints — is evidence that the social level has genuine generative capacity that the biological level does not determine. The displacement from economics to biology does not imply that social explanations are disguised biological ones. It implies that social explanations must be grounded in — must be compatible with — a more fundamental

al level of constraint. The Layered Constraint Model: PMN's Alternative to Base and Superstructure The architectural metaphor that organizes Marxist analysis — a base that supports and determines a superstructure built upon it — is one of the most influential framings in the history of social thought. Its endurance reflects a genuine insight: material conditions shape the range of possible ideological, legal, and political arrangements. Its inadequacy reflects a genuine failure: the metaphor of a building with a foundation implies that the higher levels are derivative, that explanation always runs downward, that the task of analysis is to reduce cultural or political phenomena to their economic substrate. The history of the tradition's attempts to save the model — Gramsci's hegemony, Althusser's relative autonomy of ideological state apparatuses, the various attempts to insert mediation between base and superstructure — testifies to how much intellectual labor has been expended trying to correct a flaw built into the original architecture. (Gramsci 1971; Althusser 1970) PMN replaces the foundational metaphor with a layered constraint model. The difference is not terminological. It changes what kind of explanations are available, what direction causation can run, and what the task of analysis actually is. In the foundational model, higher levels are determined by lower levels — the superstructure is what it is because the base is what it is. In the layered constraint model, lower levels constrain the range of viable configurations at higher levels without determining any single configuration within that range. Biology constrains which economic arrangements can be sustained across populations — arrangements that systematically violate nutritional requirements, reproductive capacity, or the cognitive architecture of cooperation cannot persist indefinitely. But within the space of biologically viable arrangements, there is no single economic configuration that biology selects. Economics constrains which institutional arrangements are stable — an institution that cannot be resourced by the economic arrangements it operates within will not persist. But within the space of economically viable institutions, many configurations are possible, and the selection among them is determined by political, cultural, and historical forces that have their own causal logic. The critical difference is in what happens above the floor. In the foundational model, the superstructure exists to serve the base — its function is what explains its existence, and changes in the base explain changes above it. In the layered constraint model, what exists above the floor of biological and economic constraint has genuine causal power of its own that operates both downward and laterally — not as feedback that is ultimately reducible to the base, but as independent causal force. Thatcher's economic program was not the expression of a pre-existing economic base — it used political and

ideological mechanisms to restructure the base itself, changing ownership patterns, labor relations, and the terms of accumulation through the exercise of state power organized by a governing narrative. Weber's argument about Protestantism is not that religious ideology reflected a capitalist economy that was already there — it is that specific theological commitments created the moral and psychological preconditions that made capitalist accumulation culturally sustainable before the economic system that would benefit from it existed at scale. (Weber 1905) In both cases, the causal arrow runs upward — from what the foundational model calls superstructure to what it calls base — without remainder, without requiring that the superstructure ultimately be explained by a deeper base beneath it. The layered constraint model also changes the account of what the biological level actually does. In orthodox Marxism, the economic base is foundational because it is where human beings reproduce themselves materially — production and reproduction are the conditions of all social life, and the relations under which they occur are therefore the structural determinant of everything else. PMN accepts this argument but applies it one level deeper: the economic arrangements through which human beings reproduce themselves materially are themselves downstream of the biological conditions of the organisms who undertake that reproduction. Cognitive

architecture — the limits of working memory, the social range of Dunbar's number, the evolved dispositions toward in-group cooperation and out-group suspicion — constrains which economic and political arrangements can function without requiring behavioral dispositions that the organisms who must maintain them do not reliably possess. Affective architecture — the mechanisms through which fear, shame, status anxiety, and solidarity are produced and transmitted — shapes the motivational substrate on which all economic activity depends. These are not merely psychological variations on economic relations. They are biological conditions that any viable economic arrangement must work with rather than against, and their analysis belongs to the foundation of materialist social theory, not to the superstructure above it. Three existing frameworks identified parts of this picture without completing it. Gramsci's hegemony theory correctly identified that governing narratives have genuine causal force and cannot be reduced to mere reflections of economic interest — but located the analysis within a framework that still treated the economic as ultimately determinant. Althusser's structural Marxism correctly identified that ideological state apparatuses have their own structural logic — but the insistence on determination in the last instance by the economic preserved the foundational architecture while emptying it of explanatory content. Bhaskar's critical realism correctly identified ontological emergence — that higher-level entities hav

e causal powers irreducible to lower-level components — but did not systematically connect the emergence hierarchy to the biological foundation that gives it its specific shape in the case of human social systems. (Bhaskar 1975) PMN draws from each while being reducible to none: the genuine causal force of governing narratives from Gramsci, the structural logic of institutional levels from Althusser, the ontological framework of emergence from Bhaskar, and the biological foundation that grounds and limits them all. The practical consequence of replacing the foundational model with the layered constraint model is that analysis cannot proceed by identifying the economic base and reading off the superstructure from it. It must proceed by identifying, at each level, what constraints operate from below, what causal forces operate from above and laterally, and what the viable configuration space is within those constraints. This is more demanding than the foundational model. It does not permit the analyst to reach a definitive explanation by finding the economic mechanism that underlies everything else. What it permits is more honest: an account of how specific configurations at multiple levels interact to produce the outcomes being analyzed, which constraints from below make which changes possible or impossible, and which causal forces from above are restructuring the conditions from which they themselves emerged. Asymmetry of Binding Strength: Why Lower Layers Are Not More Causal But More Eliminative The layered constraint model requires one clarification to avoid a misreading that would reproduce the foundational model's error in a new form. Because lower layers constrain upper layers, it might appear that the model is simply restating determination from below with different vocabulary — that biology still determines everything, just with more steps. This misreading confuses two different kinds of priority. The lower layers are not more causally important in the sense that changes at the biological level produce larger or more significant effects than changes at the institutional or epistemic level. Thatcher's ideological restructuring of the British economy produced effects on the distribution of material conditions that were larger and faster than any change in the biological level of the population during the same period. Governing narratives that collapse epistemic infrastructure can foreclose collective action across entire populations in ways that biological constraints alone never produce. Causal importance is not distributed hierarchically across the layers. What is distributed hierarchically is binding strength — the number of configurations at higher levels that a constraint eliminates. A biological floor that is not met eliminates virtually all configurations at every layer above it: no institutional arrangement, no governing narrative, no epistemic infrastructure can function stably when the populations that must maintain them are dying

of starvation or epidemic disease. The constraint at the biological level is more eliminative than constraints at higher levels not because biological causes are more powerful but because biological failure forecloses the functioning of all higher-level processes simultaneously. A meaning-collapse configuration is also highly eliminative — populations without orienting meaning infrastructure are available for mobilization by

whatever offers the most compelling framework fastest, which sharply constrains the range of viable political configurations. But meaning-collapse does not eliminate all configurations: survival, material coordination, and even some forms of institutional function can persist through meaning-collapse in ways they cannot persist through biological floor failure. The practical consequence for analysis is sequencing, not priority. When the biological floor is not reliably met, interventions at the institutional or epistemic layer cannot be the primary lever — not because those layers are less important but because their operation is contingent on the biological layer functioning. When the biological floor is met and the dominant constraint is institutional capture, interventions at the meaning or epistemic layer may be more effective than direct economic intervention — not because those layers are more important but because the binding constraint at that moment is not economic. The layered constraint model does not produce a fixed hierarchy of interventions. It produces a diagnostic procedure: identify which layer's constraints are currently most eliminative, and address those constraints before attempting changes that require their resolution.

2.3 — On Necessary Being

The materialist tradition has generally treated metaphysical questions about the ultimate nature of reality as either unanswerable and therefore irrelevant, or as answered in advance by the commitment to material explanation. This framework takes a different position. The contingency of everything that can be observed — the fact that everything that exists could in principle not exist — raises a question that is not answered by pointing to the material explanation of any particular thing: what accounts for the existence of anything at all? The observation is logical, not theological — a consequence of what contingency means, not a claim about what fills it. The framework is agnostic on this. The possibility of something that exists necessarily — something whose non-existence is genuinely impossible — is open. From that openness, nothing can be concluded about what that something is like, or whether it resembles anything in any religious tradition's conception of ultimate reality. The door is open. What is behind it remains unknown. Political and ethical systems in this framework do not depend on resolving this question, and they should not be constructed as if they did. The agnosticism here is not strategic — not a concession to religious audiences or a hedge against metaphysical criticism. It is the honest position: the contingency argument opens a door that the framework cannot close from either direction. Political and ethical systems built on this framework do not depend on resolving what is behind that door. What they do depend on — the biological foundation of suffering and the conditions for genuine development — is available regardless of how or whether the question of necessary being is answered. The door is open. What the framework does is go through a different one. The question of necessary being is the question that contingency raises when followed all the way down. Everything that exists in the observable world is contingent — it could in principle not exist, it came to exist through processes that were themselves contingent, and its continuation depends on conditions that might not obtain. This is true of every material thing, every physical law as we understand it, every configuration of space, energy, and matter. The regress of contingent explanations does not by itself generate a necessary being — it is possible that the regress is infinite, or that the question does not have an answer accessible from within the framework of contingent explanation. But the question is real, and dismissing it as unanswerable is not the same as showing it does not arise. Why the materialist tradition has struggled here: The standard materialist response is to accept the regress as either infinite or as terminating in a brute fact — the universe exists, with these initial conditions, and there is no further explanation. This is a defensible position, but it is not obviously more defensible than alternatives that posit some form of necessary ground.

The difficulty is that a brute fact at the base of contingent explanation is itself a kind of metaphysical commitment — it is the claim that some things exist without explanation — and it is not clear why this commitment is epistemically superior to the claim that there is a necessary ground of contingent existence. Both are claims that go beyond what empirical investigation can settle. The materialist who treats the brute fact response as obviously correct has not resolved the metaphysical question; they have chosen a position on

it. The Impersonal Necessary: One Form the Source Level Could Take One specific possibility within the agnosticism deserves explicit treatment, because it resolves an apparent tension between the framework's openness to necessary being and its materialist starting point. The standard framing of necessary being in the Western philosophical and theological tradition imports a cluster of attributes: intention, will, moral orientation, personal relationship with what it has generated. That cluster is not entailed by the logical structure of the necessary being argument. The argument establishes only that something whose non-existence is impossible might ground the chain of contingency. It does not establish that this something has preferences, intentions, or a relationship to what it grounds that is analogous to the relationship a person has to their actions. An impersonal necessary ground — a kind of absolute structure or order that simply is what it is, from which everything else follows as consequence without anyone 'choosing' those consequences — is logically compatible with the argument and removes most of the theologically generated difficulties. Gravity does not 'allow' objects to fall; it is not making a moral choice when mass curves spacetime. The same logical structure applies at the level of any posited necessary ground: if it operates without preference, intention, or moral orientation, then questions about why it 'permits' suffering, why it 'allows' injustice, and why it does not 'intervene' dissolve, because these questions all presuppose a personal agent behind the flow of reality where there may be none. The impersonal necessary ground is not good and not evil: it simply is the structure within which everything else is possible. This is one of the places where the two-level framework of 5.2 connects back to the ontological foundation. If necessary being at the source level is impersonal — the absolute structural ground without preference or intention — then the personal qualities that appear at the interface level are not transmitted from the source but generated at the intersection. The framework that feels personal, responsive, morally engaged, the sense of being accompanied through difficulty — all of this is interface-emergent in a specific sense: it arises from the interaction between consciousness and the impersonal structure, not from the structure's own personality. This does not make the interface

experience less real. It locates its reality correctly. The agnosticism extends to this specific question too: whether necessary being is personal or impersonal, whether the source level has something analogous to intention or is pure structure without orientation, is not resolvable from within this framework. What the framework can say is that the impersonal reading is logically coherent, philosophically respectable, and removes a specific class of objections that have plagued the concept of necessary being in traditions that import personal attributes before the argument requires them. Holding this specific possibility open alongside the personal alternatives is what epistemically responsible agnosticism looks like at this level. PMN's agnosticism and what it does not license: This framework is agnostic on the question of necessary being. The agnosticism is not the same as dismissal — it is the acknowledgment that the question is real, that the available answers are not clearly ranked by epistemic quality, and that the framework's analytical and evaluative commitments do not depend on resolving it. This agnosticism has a specific implication for the framework's relationship to religious and metaphysical traditions: it does not license the dismissal of those traditions on the grounds that they posit metaphysical entities beyond the material. It licenses skepticism about specific metaphysical claims that contradict well-established empirical findings, and it licenses the evaluative assessment of religious traditions by their material consequences analyzed in Part V. It does not license the conclusion that the question of necessary being is a confusion that rigorous thinking should dissolve. The practical significance of holding this question open rather than closing it with a materialist brute-fact response is twofold. First, it maintains intellectual honesty about the limits of what materialist analysis can establish — the framework does not claim more than it can defend. Second, it preserves the possibility of productive engagement with traditions that take the question seriously. Some of the most rigorous thinking about the relationship between contingent and necessary existence has occurred within religious philosophical traditions — Ibn Rushd, Maimonides, Aquinas, Ibn Sina — and dismissing that thinking because of its theological context is itself a form of the credentialing bias that 1.9 identifies as epistemically problematic. The arguments stand or fall on their merits, not on the institutional affiliation of

those who developed them. The charge of anti-metaphysics is a misreading that the structure of this section directly contradicts. An eliminativist framework would close the door opened here. This one leaves it open — not rhetorically, but because the epistemological commitment of Part I requires acknowledging the limit of what current understanding can establish. The framework does not claim that metaphysical questions are meaningless

s; it claims that they cannot currently be answered by the methods available to it, and that political and ethical analysis should not be built on answers that exceed the available evidence. That is agnosticism in the precise sense, not eliminativism. A framework that required metaphysical resolution before it could evaluate whether an institutional arrangement produces unnecessary suffering would be analytically paralyzed. What this framework requires instead is that metaphysical commitments be held at the appropriate level of certainty — which, for most current metaphysical questions, is considerably below the certainty with which they are typically held. The agnosticism here is differentiated, not uniform. The framework is agnostic about necessary being as a metaphysical question while being committed to specific claims at every other level: to the reality of the material substrate, to the causal efficacy of structural arrangements, to the evaluative primacy of suffering-reduction. The agnosticism is located precisely at the point where the question genuinely cannot be settled by the tools the framework uses — empirical analysis, structural assessment, material consequence evaluation. At that point, the honest response is to name the uncertainty rather than to resolve it by fiat in either the materialist or the theological direction.

2.4 — Social Ontology: The Reality of Emergent Structures

The materialist ontology established in sections 2.1 and 2.2 requires explicit extension to the social level. The layered constraint model posits that social structures — institutions, cultural frameworks, legal systems, economic arrangements — have genuine causal properties that cannot be fully reduced to the properties of their individual components. This claim is not obvious and requires defense, because reductionism — the view that social phenomena are fully explicable in terms of the psychology and biology of the individuals who compose them — is both intellectually tempting and politically convenient for those whose interests are served by treating social arrangements as natural rather than constructed. The defense of social ontological emergence is not a concession to idealism. It is a materialist commitment: social structures are real because they have real material effects that the properties of individuals alone cannot produce or predict. The mechanism of social emergence is specific. Individual human beings bring to social interaction a set of biological capacities and tendencies — cognitive, affective, motivational — that section 3.1 through 3.8 analyzes as the biological foundation of institutional life. But the patterns of behavior that emerge when those individuals interact within established institutional structures cannot be derived from those individual properties alone, because the institutional structures themselves constrain, channel, and amplify individual tendencies in ways that produce systemic outcomes that no individual intended or could have produced unilaterally. A market economy does not produce its characteristic patterns of distribution, innovation, and crisis because individual market participants intend those patterns; it produces them because the rules governing exchange, the property rights structuring resource access, and the competitive pressures generated by the interaction of many agents produce dynamics that emerge from the structure of the interaction rather than from any individual's intentions within it. The structure is causally real precisely because it produces outcomes that individual-level analysis cannot predict or explain. This ontological point has two practical implications that run throughout the framework. The first is methodological: the primary unit of analysis for understanding social outcomes is the institutional structure and its interaction with individual tendencies, not the individual actor and their choices within a structure that is treated as given. The second is political: if social structures have genuine causal properties, then changing social outcomes requires changing social structures — not merely changing individual beliefs, values, or behaviors. The persistent progressive error of treating political transformation as primarily a project of consciousness-raising, the persistent conservative

error of treating social problems as primarily failures of individual character, and t

he persistent liberal error of treating institutional reform as primarily a matter of improving the incentives that individuals face — all three errors share the same ontological mistake of underestimating the degree to which structural properties are irreducible to the properties of their individual components. The relationship between ontological emergence and the base-superstructure problem: the traditional Marxist framework treats the distinction between economic base and ideological superstructure as an ontological hierarchy in which causal primacy runs in one direction. The layered constraint model replaces this with a more complex ontology in which different levels of organization have different causal properties, and in which causality runs in multiple directions simultaneously. The economic level has genuine causal properties that shape what is possible at the ideological level — but the ideological level also has genuine causal properties that shape what is possible at the economic level. Neither level is ontologically primary in the sense of having all the causal power. Each is real, each constrains and enables the others, and analysis that treats one level as derivative of the other systematically misses the causal dynamics that operate through the interaction between levels. The emergence claim requires precision about what kind of causality social structures exercise. They are not mechanisms in the sense of physical mechanisms — they do not push and pull material objects through direct physical contact. They are enabling and constraining conditions: they determine what actions are available to actors operating within them, what the costs and benefits of different actions are, and what the consequences of those actions will be for actors with different structural positions. A legal system does not physically prevent actions that are illegal — it determines the consequences of those actions in ways that shape whether and how they occur. An economic structure does not physically compel labor — it determines what the consequences of not laboring are for people without independent access to means of subsistence. The causality is real; it is just not the causality of billiard balls. Against the twin reductionisms: The two reductionist moves that social ontology must resist are the psychological and the biological. Psychological reductionism — the claim that social structures are ultimately nothing more than the aggregate of individual psychology — fails because it cannot explain why the same psychological tendencies produce systematically different outcomes under different structural conditions. People are not more or less prone to corruption as individuals under different institutional designs; they are more or less constrained by systems that make corruption costly or rewarding. The structural variable is doing explanatory work that individual psychology cannot do. Biological reductionism — the claim that social structures ultimately re

duce to genetic or evolutionary determinants of individual behavior — fails for the same reason and adds the further error of treating evolutionary adaptations to ancestral environments as explanations of behavior in radically different institutional environments. The ontological status of institutions: Institutions are real in the sense that matters for social analysis: they have stable properties that persist across changes in the individuals who occupy positions within them, they generate predictable effects on those individuals' behavior, and they are causes of outcomes that cannot be explained without reference to them. An institution that has just been founded is already shaping behavior before any individual has had time to be thoroughly socialized into its norms — the structural incentives operate from the moment the institution's reward and sanction structure is established, not only after individuals have internalized its culture. This is one piece of evidence that institutional causality is not reducible to psychological causality: it operates faster than socialization can explain. The temporal dimension of structural causality deserves emphasis. Institutions outlast the individuals who build them, operate through the individuals who inherit them, and shape the preferences and possibilities of individuals who were born into the conditions they produce. This temporal reach is one of the features that makes structural analysis necessary: explanations that track only the behavior of currently living individuals cannot account for outcomes that were set in motion by institutional arrangements that no living person established. The distribution of material conditions across populations today is partly a product of institutional arrangements operating across multiple generations — arrangements whose original architects are dead and whose current beneficiaries did not choose the advantage they inherited. Explaining those

outcomes requires ontological commitment to the reality of persistent institutional structures as causal entities. The implications for the evaluative framework: if social structures are real causal entities with properties that persist across changes in the individuals who occupy them, then changing outcomes requires changing structures — not only changing individual behavior, beliefs, or incentives while leaving structures intact. This is the ontological foundation for the structural approach to change analyzed in Part X. Interventions targeted at individuals operating within unchanged structures will produce individual-level changes that structures will work to reverse, because the structural incentives that produced the original behavior are still present and will shape the successor individual in the same direction.

2.5 — Against Reductionism and Against Explanatory Idealism: The Middle Position

The social ontology of this framework stakes out a position between two errors that recur in both academic and political analysis. The first error is reductionism: the view that social phenomena are fully explicable in terms of the properties of their individual components, whether those components are individual human psychology (methodological individualism), biological drives (sociobiology), or economic incentives (rational choice theory). The second error is idealism in its various forms: the view that ideas, culture, consciousness, or discourse are the primary determinants of social outcomes, and that material conditions are either derivative of ideational structures or are less important than the narrative frameworks through which they are interpreted. Both errors are available in left and right versions, and both produce characteristic analytical blind spots that the framework is designed to avoid. The reductionist error produces the characteristic failure of analyses that explain persistent structural outcomes by reference to the aggregated choices of individuals who are treated as operating in a structurally neutral environment. When labour market analysis explains the gender wage gap primarily by reference to individual occupational choices rather than to the institutional arrangements that produce the choice set within which those choices are made, it is committing the reductionist error. When public health analysis explains differential health outcomes by reference to individual behavioral choices rather than to the material conditions that structure those choices, it is committing the same error. The analysis is not factually wrong about the individual level — individuals do make choices that have consequences. The error is in treating the choice as the primary unit of analysis when the structure of the choice set is the more causally fundamental variable. The idealist error produces the characteristic failure of analyses that explain structural outcomes primarily by reference to the ideas, values, or discursive frameworks that circulate within a society, without adequate attention to the material conditions that produce and maintain those frameworks. When political analysis explains the persistence of inequality primarily by reference to ideological hegemony without examining the material interests that invest in the production of that ideology, it is committing the idealist error. When cultural analysis treats the transformation of oppressive practices as primarily a project of changing the narratives through which those practices are understood, without examining the material arrangements that make those practices instrumentally rational for those who engage in them, it is committing the same error. The analysis correctly identifies that ideas and narratives have real causal properties. The error is in treating the ideational level as causally primary when the material arrangements that sustain the ideas are the more fundamental v

ariable. PMN's middle position holds that both levels are real, both have genuine causal properties, and the relationship between them is empirically variable rather than ontologically fixed. In some configurations, material conditions are the binding constraint on what ideas are viable — changing the ideas without changing the material conditions produces ideas that fail to take institutional form. In other configurations, ideational frameworks are the binding constraint — material conditions that would permit multiple institutional configurations are stabilized in specific configurations by governing narratives that make alternatives unimaginable. The analytical task is to identify which constraint is binding in the specific configuration under analysis, which requires attending to both levels simultaneously rather than deriving one from the other. The middle position requires active maintenance against pressures from both directions —

not because the errors it avoids are equally common or equally consequential in all contexts, but because the analytical temptations they represent operate in different domains and on different analytical communities. Reductionism is the characteristic error of traditions that derive their prestige from the natural sciences and seek to import natural science methods into social analysis without adequate attention to the emergent properties that social structures generate. Explanatory idealism — treating the ideational level as causally primary without adequate anchoring in the material conditions that produce and sustain ideas — is the characteristic error of traditions that derive their prestige from the humanities and philosophy and treat ideas, narratives, and cultural formations as primary objects of social analysis while losing sight of the structural conditions that make those formations viable or unviable. Both errors are intellectually understandable given the communities in which they developed; in the forms described here, neither is analytically adequate. What the middle position must establish: The middle position is not simply the claim that both material conditions and ideas matter — that is a truism that anyone would grant and that carries no analytical content. The middle position requires a specific account of the relationship between the two levels: how material conditions shape ideational formations, how ideational formations in turn shape the material conditions that produce them, what the directionality of influence is under different circumstances, and what the mechanisms are through which each level affects the other. Without this specificity, the middle position is a gesture toward complexity rather than an analysis of it. The materialist priority claim: PMN's position within the middle space is not symmetrical. Material conditions are ontologically prior in the specific sense that the biological substrate — the organisms whose suffering is the framework's evaluative anchor — is the level

at which material conditions are most directly consequential. Ideas that do not connect to this substrate through any causal chain have no material consequences; ideas that do connect to it are real material forces by virtue of that connection. This is not a claim that material conditions mechanically determine ideational formations — the relationship is complex, bidirectional, and mediated by institutional structures at multiple levels. It is a claim that the directionality of primary causation runs from material to ideational more reliably than in the opposite direction, and that when the two levels diverge — when ideas diverge from material conditions significantly — it is the material level that generates the pressure to which ideational formations must eventually respond. The mechanism of ideational causality: Ideas have material consequences through specific mechanisms that are themselves material: through the behavior of individuals who act on them, through the institutional designs that embody them, through the allocation of resources that they authorize or prohibit, and through the visibility and invisibility of conditions that they designate as natural or contingent. The claim that ideas matter is not a claim that they float free of material causation — it is a claim that the material causation runs through the ideational level as a real intermediate stage with genuine causal properties. A governing narrative that designates structural suffering as the natural result of individual failure is not merely a description of material conditions — it is a cause of the perpetuation of those conditions by suppressing the visibility that would generate political pressure for change. Diagnostic for distinguishing structural from ideational causation: when analyzing a persistent social outcome, ask whether the ideational formation that accompanies it — the beliefs, narratives, and norms that appear to justify it — could change while the structural conditions remained, and what would then happen. If the structural conditions would reproduce the outcome regardless of the ideational change, the primary causation is structural and the ideational formation is an effect. If the structural conditions would produce a different outcome once the ideational formation changed — because the formation is doing work in shaping behavior, resource allocation, or institutional design — then the ideational formation is itself a material cause at the level where it connects to structural conditions. Most persistent social outcomes involve both mechanisms, and the analytical task is to specify which is primary in the domain under analysis. The middle position has a specific implication for political strategy that is often missed. Neither pure consciousness-raising (change the ideas and the structures will follow) nor pure structural intervention (change the material conditions and the ideas will follow) is adequate on its own. Structural change that does not address th

e ideational formations that naturalize the old arrangement tends to be reversed when the conditions that forced the change ease — the ideas provide the template for reconstruction. Ideational change that does not connect to structural conditions remains politically inert — it produces people who understand their situation differently without the material capacity to change it. The relationship between the two levels is not additive but interactive: each creates the conditions under which the other becomes more effective. Terminological note: “Idealism” appears in three distinct senses in this framework, and conflating them produces misreadings. (1) Explanatory idealism in social analysis — the view that ideas, discourse, or consciousness are the primary determinants of social outcomes — is what this section (2.5) identifies as an error. PMN rejects it not because ideas are unreal but because they are not causally primary. (2) Aspirational or directional idealism — the commitment to holding a long-term direction against immediate pragmatic pressures — PMN explicitly endorses. The Idealism–Pragmatism tension in 13.2 treats this as a permanent structural feature of serious political action, not a problem to be solved. (3) Metaphysical idealism about the fundamental nature of reality — the position that consciousness underlies or constitutes the physical world — PMN holds open as a genuine question (2.1, 2.3). The framework’s practical commitments do not require resolving it. Readers who encounter “against idealism” language elsewhere in the document should understand it as targeting sense (1) exclusively, not senses (2) or (3).

2.6 — PMN, Secular Rationalism, and the Limits of Materialist Explanation

The question this section addresses has two parts that must be kept separate: whether PMN is a form of secular rationalism, and whether it should be. They are different questions because secular rationalism is not a single position — it is a family of positions that share a methodological commitment to reason and empirical evidence as the primary arbiters of belief, but differ substantially on what follows from that commitment at the metaphysical level. PMN inherits the methodological core and explicitly rejects several conclusions that mainstream secular rationalist discourse treats as implications of it. The valid critique of materialism as only description: The objection that materialism cannot answer the deepest questions — only describe — is formally correct and deserves direct acknowledgment rather than deflection. Materialism as a methodology is powerful within its scope: it traces causal chains, identifies structural mechanisms, predicts systemic behavior, and evaluates institutional arrangements by their material consequences. What it cannot do, without crossing into metaphysical claim-making that exceeds its epistemic warrant, is answer the question that contingency raises when followed all the way down: why does anything exist at all? At this limit, materialism terminates in a posit — either an infinite regress or a brute fact — not in an explanation. This is an honest limitation. The mistake is treating it as a limitation specific to materialism. Every framework terminates somewhere. The question is not which framework can answer everything — none can — but what each framework can do within its scope before reaching that limit. PMN’s scope is the analysis of how human social arrangements produce and distribute structural suffering, and what institutional conditions expand or foreclose genuine human development. Within that scope, the materialist axiom is productive. Beyond it, PMN acknowledges the boundary and stops — rather than filling the space with claims that exceed its warrant. What secular rationalism gets right: The methodological contribution of the secular rationalist tradition is real and PMN inherits it without reservation. Empirical evidence is the primary tool for evaluating claims about the observable world. Burden of proof falls on the actor making the positive claim. Falsifiability is a meaningful criterion for distinguishing empirical claims from non-empirical ones. The history of institutional religion includes sustained patterns of epistemic capture — using doctrinal authority to insulate claims from the scrutiny they would not survive — and the secular rationalist critique of this pattern has produced genuine intellectual progress. These contributions are not diminished by the limitations that follow. Where secular rationalism overreaches: The problem arises when secular rationalism treats its methodological commitments as generating metaphysical conclusions. The claim that there is no God is not the same kind of

claim as the claim that there are no quarks — one is a positive metaphysical assertion about ultimate reality, the other is an empirical hypothesis with a clear evidential profile. Applying the tools of empirical falsification to the question of God’s existence works only on specific empirical claims made in God’s name — prayer studies, miraculous interventions, young-earth geology — not on the underlying metaphysical question of necessary being. Dawkins’ *The God Delusion* is powerful on the former but does not engage the latter: its most forceful arguments target specific interventionist claims made in God’s name — miracles, answered prayer, young-earth geology — rather than the thesis of necessary being as a metaphysical category, and these are not the same argument. Harris’ moral landscape arguments are valuable within their domain — demonstrating that wellbeing has an objective dimension that can be investigated empirically — but do not resolve the meta-question of why wellbeing matters in the first place, which is precisely the ground the necessary being thesis occupies. The pattern across much contemporary skeptical rationalist discourse is to conflate two distinct positions: presumption of atheism as an epistemic starting point — which is defensible as a working heuristic for evaluating empirical claims — and atheism as a settled philosophical conclusion about the nature of ultimate reality, which requires substantially more argument than the working heuristic supplies. PMN’s position: presumption of agnosticism is more consistent with genuine epistemic humility than presumption of atheism, because atheism is itself a positive claim about the non-existence of a necessary being — a claim that requires justification by the same epistemic standards it demands of theism.

2.6b — The Projecting Fallacy Risk and PMN’s Specific Defense

The projecting fallacy is the error of applying a category of thought — formed within a specific domain of experience and valid there — to a domain where the category does not apply, and then treating the resulting confusion as a refutation of the target rather than a limitation of the tool. The paradigm case in atheist argumentation is the question “who created God?” — which imports the category of efficient causation (X is produced by a prior Y) into the question of necessary being, then treats the infinite regress this generates as evidence against necessary being. But necessary being is defined as existing without a prior cause — that is what “necessary” means in this context. The question does not refute the concept; it misapplies the concept of causation to an entity that, by definition, stands outside the causal chains that give causation its meaning. The question is confused, not the concept being questioned. Does PMN commit the same error? The question requires direct engagement: is PMN’s materialist axiom itself a projecting fallacy — an imposition of categories developed for the analysis of contingent phenomena onto domains where those categories may not apply? The answer depends on a distinction PMN must make explicit. PMN commits the projecting fallacy if it claims that material analysis is not just the most productive methodological starting point for its scope, but the correct description of ultimate reality. If PMN is making the claim that reality is, at the deepest level, material — that there is no necessary being, no non-material substrate, no consciousness prior to matter — it is projecting materialist categories onto a domain they were not designed to address and cannot reach. PMN does not make this claim. The materialist axiom in PMN is explicitly methodological, not metaphysical: it is the working assumption that yields productive analysis of social arrangements and their material consequences. It is not a statement about what reality is at the level below the phenomena the framework analyzes. The difference between methodological materialism and metaphysical materialism is the difference between “we analyze this as if material conditions are primary, because that generates useful results” and “material conditions are all there is.” PMN claims the first. It does not claim the second. The internal risk and the safeguard: The projecting fallacy risk for PMN is not theoretical — it is operational. Users of the framework who internalize the materialist axiom as a working method may, over time and without noticing, begin to treat it as a metaphysical conclusion. The language of the framework — “material conditions,” “structural causation,” “biological foundation” — can gradually harden from methodological vocabulary into ontological claim. When this drift occurs, the framework begins to commit the error it is designed to avoid: importing the categories of one domain into another where they do not apply, and then treating the con

fusion generated as a feature of reality rather than a limitation of the analytical tool. The safeguard is the same one that prevents all other forms of framework capture: regularly asking whether the conclusion being drawn exceeds the scope of the evidence and method that generated it. A PMN analysis that concludes “this institutional arrangement produces structural suffering and should be reformed” is within scope. A PMN analysis that concludes “therefore religious experience is merely neurological epiphenomenon” has crossed from methodological materialism into metaphysical claim-making — and has done so without noticing the crossing. Projecting fallacy diagnostic for PMN use: before drawing a conclusion, ask — does this conclusion require that material conditions be not just analytically primary (the methodological claim PMN makes) but ontologically fundamental (the metaphysical claim PMN does not make)? If yes, the conclusion has crossed the line. Examples that cross: “because suffering is materially produced, the experience of divine presence in suffering is illusory.” Examples that do not cross: “because suffering is materially produced by identifiable institutional arrangements, those arrangements can be changed.” The first imports an ontological conclusion from a methodological premise. The second stays within the scope of the method. The harder version of Critique 1 — and PMN’s response: A sharper formulation of the critique presses further: even if PMN does not evaluate whether source-level claims are true, by choosing to evaluate only material consequences PMN has already decided that material consequences are the dimension that matters for its purposes. This choice is not neutral. A tradition that holds the spiritual dimension to be primary will object not that PMN gets the material analysis wrong, but that PMN’s criterion of relevance is itself a form of colonization. PMN responds by clarifying what its offer actually is: PMN does not call religious practices into its court to be judged by PMN standards. It offers to tell a tradition what its practices do materially — and then steps back. The tradition retains full authority to decide what to do with that information within its own evaluative framework. A practice that PMN identifies as producing structural suffering can be sustained, reformed, or abandoned by the tradition on the tradition’s own terms. PMN’s information does not determine the verdict; it provides one input among many. This is genuinely different from the New Atheist move, which uses empirical disconfirmation of specific interventionist claims to conclude against the validity of the tradition as a whole. PMN has no such concluding move. What it has is a report — and the acknowledgment that the tradition is the authority on what to do with it.

2.6c — What Hard Materialism Gets Right, and Where PMN Diverges

Hard or eliminative materialism — the position associated with Churchland, Dennett, and others in the eliminativist tradition — makes a stronger claim than methodological materialism: that mental states, consciousness, and subjective experience are, at bottom, nothing but physical processes, and that the apparent irreducibility of experience to physical description is an artifact of our current conceptual framework rather than evidence of a genuinely non-physical phenomenon. This position has significant explanatory advantages and PMN does not dismiss it. It correctly identifies that dualist accounts — consciousness as a separate substance operating alongside physical processes — face serious causal problems. It correctly emphasizes that the brain is the primary locus of what we call experience and that brain states have demonstrable correlates with experiential states. And it correctly resists the impulse to postulate additional entities beyond what explanation requires. Where PMN diverges: The hard problem of consciousness — Chalmers’ formulation of why there is something it is like to be a conscious creature, why physical processes give rise to subjective experience rather than to mere information processing without any felt quality — is not resolved by eliminative materialism. It is dissolved: the eliminativist argues that the felt quality is either an illusion or a confusion about the nature of our concepts. PMN treats this dissolution as premature rather than complete. Not because the eliminativist argument is obviously wrong — it is not — but because the gap between the most detailed physical description of a brain state and the subjective character of the experience that state is said to constitute has not been bridged in any account that most philosophers of mind find fully satisfying. This is not a gap that increased neuroscientific detail will necessarily close, because the question is not empirical in the ordinary sense: it asks why any physical

process should be accompanied by experience at all, which is a different question from what physical processes correlate with which experiences. PMN's response to hard materialism is: your account is powerful within the third-person explanatory domain, and PMN's analytical work takes place within that domain. But claiming that the first-person domain — the domain of what it is like — is fully explicable in third-person terms is an additional claim that outstrips the evidence. PMN maintains that the hard problem is genuinely hard, not merely hard-for-now, and that this matters for the agnosticism about consciousness and necessary being that 2.1, 2.3, and 5.2–5.3 develop. The critique hard materialism directs at PMN: A hard materialist will object that PMN's agnosticism about consciousness and necessary being is intellectual timidity — a refusal to follow the evidence where it leads in order to avoid antagonizing audiences whose cooperation the framework needs. The objection has real force and d

eserves direct reply. PMN's agnosticism is not strategic. It is not calibrated to maximize the framework's political reach. It reflects a genuine assessment that the evidence does not settle the hard problem, that the question of necessary being is not resolved by pointing to the productive results of materialist analysis within its scope, and that treating unsettled questions as settled — in either direction — is a failure of the epistemic standards Part I establishes. A framework that applies those standards to itself must acknowledge where they produce uncertainty rather than resolution. The hard materialist position is one among several that remain genuinely defensible given current understanding. PMN holds it as a live option rather than as a settled conclusion — which is not the same as treating it as wrong.

2.6d — PMN as Open Passage: Material Analysis Without Metaphysical Foreclosure

The sections above establish what PMN is not: it is not eliminative materialism, not secular rationalism in the strong sense, not a framework that forecloses metaphysical questions by methodological fiat. What requires equal clarity is what follows from these negatives for the relationship between PMN's analytical work and the metaphysical commitments of those who use it. The structural claim: PMN's analytical conclusions — about what institutional arrangements produce structural suffering, what conditions expand genuine development, what power configurations sustain harmful arrangements — are available to analysts who hold any metaphysical position, because those conclusions are generated from and validated by evidence that is independent of the metaphysical question. A theist who believes material reality is the expression of divine creativity, and a hard materialist who believes material reality is all there is, are looking at the same institutional arrangements and the same distribution of suffering. PMN's analysis of what produces that distribution and what could change it is valid for both — not because it splits the difference between their metaphysical positions, but because it operates in a domain where the metaphysical difference does not determine the analytical conclusion. If God exists, what follows for PMN? This question is worth addressing directly because the answer is structural, not rhetorical. If a necessary being exists that underlies the material world — if the framework 2.3 holds open is in fact true — PMN's analytical work is not undermined. What changes is the ontological description of the substrate: the material conditions PMN analyzes are, in that case, the mode through which a necessary being expresses or sustains reality. The suffering those conditions produce is still suffering. The institutional arrangements that produce it are still identifiable. The changes that would reduce it are still material and still traceable. The minimal anchor — reduce unnecessary suffering, expand becoming — retains its force whether its grounding is purely biological or whether it is grounded in a divine order that values development and deplores unnecessary harm. For a tradition that holds something like the latter — that suffering has moral weight because human beings are created with dignity, that institutions should serve human flourishing because that is what creation commands — PMN provides analytical tools for identifying where those commitments are being violated and what structural changes would honor them. The framework does not belong to secular humanism more than to this tradition. It belongs to whoever applies it honestly to the criterion it establishes. The dogmatism problem in both directions: PMN treats strong atheism and strong theism as symmetrically susceptible to a specific failure

mode: treating an open metaphysical question as settled in advance of the evidence that would settle it, and then using that pre-settlement

not to insulate the conclusion from scrutiny. 1.2 identifies this as the primary diagnostic failure — the move from revision to insulation. Dogmatic theism commits it by treating the existence and character of God as beyond evidential challenge. Dogmatic atheism commits it by treating the non-existence of God as established by the absence of empirical evidence for the specific interventionist theism that most atheist critiques target — while leaving the question of necessary being, which is not the same question, effectively unaddressed. PMN does not take either side of this insulation. What it does take is the epistemic commitment that both positions are making claims that exceed their warrant — and that intellectual honesty requires acknowledging this even when the acknowledgment is uncomfortable for one's community of inquiry. The range of users this framework is available to: a secular humanist who finds the minimal anchor compelling on purely naturalistic grounds; a Muslim analyst who understands khalifa (stewardship) as requiring institutional accountability for the conditions of those in one's care; a liberation theologian who reads the preferential option for the poor as demanding exactly the kind of structural analysis PMN provides; a Buddhist practitioner who understands dukkha as having institutional as well as individual dimensions; a Confucian thinker who finds the arc from biological vulnerability through social constitution to institutional obligation structurally compatible with the relational ontology of 3.2b; an analytic idealist who holds that consciousness is fundamental and finds that PMN's analytical conclusions about material arrangements are valid regardless of what the substrate of those arrangements ultimately is. None of these uses requires that the user adopt the metaphysical position of any of the others. What all require is applying the framework honestly to the criterion it establishes. The framework is not neutral — it has specific analytical commitments and a specific evaluative standard. But it is open at the metaphysical level in a way that secular rationalism, eliminative materialism, and strong theism are not. That openness is not a weakness of the framework. It is what the epistemic commitments of Part I require of any framework that applies them honestly to itself.

2.6e — The Supra-Rational Challenge: Is PMN's Criterion Binding on Anyone?

The sharpest challenge to PMN from religious and post-secular traditions is not about whether PMN gets the material analysis right. It is about whether PMN's criterion has any authority at all. The strongest version of this challenge, articulated within contemporary Islamic apologetics and with analogues in apophatic Christian and Vedantic traditions, runs as follows: "You are playing a game where reason is the final arbiter and evidence is the only currency. But your game is not neutral. You are forcing God to submit to your criteria. The real question is not 'Is God rational?' but 'Is rationality itself divine?'" Applied to PMN: by choosing to evaluate arrangements by their consequences for suffering and becoming, PMN has already decided that suffering and becoming are what matter. A supra-rational tradition — Sufi, apophatic Christian, certain Vedantic schools — can simply decline this criterion. The encounter with God is not a means to reduce structural suffering. It is the end. PMN, on this view, has nothing to say to someone who holds that position. PMN's response — first layer: PMN does not claim that reason is the final arbiter of ultimate reality. It claims that reason and evidence are appropriate tools for analyzing material consequences within the scope of that analysis. The supra-rational tradition is not being asked to adopt PMN's epistemology as its foundation. It is being offered PMN's analysis as one input into whatever deliberative process the tradition uses internally. Whether reason has divine sanction — whether rationality is itself a participation in divine logos — is a question PMN holds open (2.3). The tradition's answer to that question does not determine whether PMN's material analysis is accurate. Both can be true simultaneously: the encounter with God is ultimate, and the institutional arrangements through which that encounter is mediated produce structural suffering that the tradition itself has reason to address. PMN's response — second layer, the convergence observation: The strongest version of the supra-rational challenge holds that encounter with God is the end

— full stop. But notice what every supra-rational tradition actually teaches about what that encounter produces in the practitioner. The Sufi who achieves fana (annihilation of the ego in God) is described as returning from that encounter with greater compassion for suffering, not with indifference to it. The apophatic Christian who contemplates the divine darkness comes back into the world with greater capacity for love of neighbor. The Vedantin who realizes the identity of Atman and Brahman discovers that the suffering of all sentient beings is the suffering of the self. In virtually every supra-rational tradition, the encounter with the ultimate does not produce disregard for suffering — it produces the deepest possible regard for it. The disagreement between these traditions and PMN is about the path to that regard, not about whether it is the destinati

on. PMN's minimal anchor is not a rival to the supra-rational destination. It is a different route description to a place most supra-rational traditions recognize as central to their own map. The honest limit: What if a tradition genuinely holds that suffering is spiritually necessary, beneficial, or irrelevant to the highest values? PMN's honest response is not philosophical counter-argument. It is two observations. First: this is a rare position even within traditions that value suffering. Most traditions that engage suffering — asceticism, martyrdom, sacrifice — do so as a means to something beyond suffering, not as an endorsement of structural suffering produced by institutional arrangements on populations who did not choose it. The distinction between voluntary suffering undertaken for spiritual development and imposed structural suffering produced by exploitation is one that virtually all traditions recognize. Second: PMN cannot compel anyone who genuinely does not care whether institutional arrangements produce suffering. It does not claim to. But it can observe that almost every human being, when experiencing suffering themselves, seeks to reduce it — and that the minimal anchor is grounded in this fact about what suffering is for the organisms experiencing it — which is why declining the criterion is available in argument more easily than it is available in practice. The tradition that tells its members that their structural suffering is irrelevant to its values will be evaluated by PMN as producing unnecessary suffering. Whether the tradition accepts that evaluation is not something PMN can determine. What it can determine is what the evaluation is. The rationality-as-divine argument: PMN's response: This reframing — “Is rationality itself divine?” — is not a question PMN dismisses. It is a question PMN holds genuinely open. If rationality is a participation in divine logos, as Aquinas held and as some Islamic philosophical traditions have argued (drawing on al-Farabi's active intellect), then PMN's methodological rationalism is not in competition with theism — it is one expression of a capacity that theism itself valorizes. The challenge, on this reading, becomes: PMN uses a faculty (reason) that theism regards as divinely given, to analyze arrangements that theism regards as subject to moral evaluation, using a criterion (suffering matters) that theism generally endorses. The dispute about whether God exists does not touch the work PMN is doing within that shared territory. What this challenge correctly identifies is that PMN cannot justify its criterion to someone who has already stepped outside the framework within which suffering matters. PMN acknowledges this. It notes that this stepping-outside is rarer in practice than it appears in philosophical argument — and that the traditions most likely to articulate the supra-rational challenge are also the traditions most insistent, in their own terms, that suffering calls for response.

III: The Biological Foundation: From Physiology to Institution

3.0 — The Normative Status of Biological Constraints: Conditional, Not Derived

Before the biological analysis of Part III begins, a philosophical clarification is required. The framework starts from biology — from the observation that organisms avoid pain, require specific material conditions to function, and are capable of development whose conditions can be structurally supported or blocked. A standard objection anticipates that this move commits a version of the naturalistic fallacy: the inference from what is (biological facts about organisms) to what ought to be (normative claims about what arrangements are better or worse). If this objection were correct, it would undermine the framework's evaluative criterion before that criterion is applied to anything. The objection is not correct, but it is not trivially wrong. The response requires precision. Why Suffering and Not Life as the Evaluative Foundation A clarification is required about the relationship between the biological foundation and the choice of suffering as the evaluative criterion. The argument is sometimes raised that if the framework begins from biology, it should ground its evaluative criterion in life itself — in the conditions for biological flourishing — rather than in suffering, which is only one of life's signals. This is a serious objection and deserves a direct answer. Life is ontologically prior; suffering is evaluatively operational. The framework does not deny that life is ontologically more fundamental than suffering. Suffering is possible only for living organisms; it is a signal that the conditions for life's normal functioning are being disrupted. In that sense, suffering is derivative of life at the ontological level. The choice of suffering as the evaluative foundation is therefore not a claim about what is most fundamental in the order of being — it is a strategic choice about what is most tractable as an evaluative criterion. The distinction between ontological priority and evaluative operability matters: something can be more fundamental in one sense and less useful in another. Why life cannot serve as the evaluative criterion directly: The problem with life as an evaluative foundation is that it underdetermines what counts as good institutional arrangements. Every tradition — including those that have produced the worst institutional outcomes in recorded history — can claim to be serving life: the life of the nation, the life of the race, the life of the species. The concept is sufficiently expansive to accommodate almost any institutional prescription. Suffering, by contrast, is more resistant to this kind of rhetorical capture. It is harder to claim that arrangements are good when they demonstrably produce suffering on the criteria the framework specifies — imposed, preventable, structural, dysfunctional — than it is to claim they serve life. The choice of suffering over life is therefore a choice made in full awareness that life is more fundamental: it is the choice of the criterion that is more tractable, more cross-culturally r

ecognizable, and more resistant to the ideological deformation that vague life-affirming concepts have consistently undergone. The relationship between the two levels: The evaluative framework uses suffering as its criterion precisely because suffering is life's most legible distress signal. Reducing suffering is not the goal in itself — it is the operationalization of the goal, which is the restoration and expansion of the conditions under which life can develop rather than merely persist. The becoming framework of Part IV makes this explicit: the floor is suffering reduction; the anti-foreclosure criterion concerns the preservation of adaptive response capacity. The floor is specified in terms of suffering because suffering is what is most directly identifiable and most directly preventable through institutional design. The anti-foreclosure criterion concerns adaptive response capacity because the conditions for genuine response development cannot be reduced to the mere absence of suffering without losing the evaluative content that makes the framework more than a pain-minimization algorithm. The practical implication: when the framework is challenged on the grounds that some suffering is good — that it builds character, produces resilience, generates the creative tension that development requires — the appropriate response is not to deny this but to apply the four criteria. Suffering that is chosen, functional, not institutionally preventable, and not structurally produced is not what the framework targets. The challenge does not refute the framework; it illustrates why the typology in 3.4b is necessary. The Two-Level Architecture: Life as Ontological Foundation, Suffering as Evaluative

Criterion Life is ontologically prior to suffering — suffering is possible only for living organisms and is a signal that life's normal functioning is being disrupted. What remains undeveloped is why this two-level structure is the correct architecture, what it gains over single-level alternatives, and what constraints it places on the framework's evaluative scope. The architecture is doing more structural work than that brief statement suggests.

3.0b — The Naturalistic Fallacy Objection and Why It Does Not Apply Here

The naturalistic fallacy, as Moore formulated it, is the error of defining normative properties through natural properties — treating the good as equivalent to some natural fact such as the pleasurable, the evolutionarily adaptive, or the biologically functional. Hume's related is-ought gap identifies a distinct but connected logical problem: descriptive statements about what is the case do not, by themselves, entail prescriptive statements about what ought to be done. The standard objection to PMN combines both concerns: the framework moves from the biological fact that organisms avoid suffering to the normative commitment that suffering should be reduced, which is precisely the kind of movement both objections prohibit. The objection misdescribes what the bridge does. PMN does not claim that what is natural is good, that what organisms do is what they ought to do, or that biological function is the same as moral value. The commitment that suffering is bad is a design choice, explicitly flagged as such in 3.0f — adopted because it is non-arbitrary, convergently held across traditions, and coherent with the framework's evaluative purpose, not because biology logically entails it. The normativity is not derived from biological facts. It is imported from the evaluative criterion of 4.1, and biology tells us what that criterion requires in concrete conditions. The biological facts about suffering are relevant not because they generate the normative commitment but because they give it traction: the commitment is about something real, measurable, and causally connected to the institutional arrangements under evaluation. The bridge is conditional, not derivational. If you accept the design choice — suffering is bad — then the biological facts about what causes suffering at scale carry normative implications for how institutions should be structured. The "if" does the normative work; the facts do the diagnostic work. This is not the naturalistic fallacy. It is the application of a pre-specified normative criterion to empirical facts — a permissible move in any ethical framework that connects normative commitments to the world. Anyone seeking to challenge PMN's position should target the design choice itself, argued in Part IV, not misidentify the bridge as a derivation it is not.

3.0c — The Stability Argument: The Bridge That Is Not a Derivation

The stability argument provides a second line of response that is often misread as a version of the naturalistic fallacy but is not. It runs as follows: institutional arrangements that produce mass structural suffering generate observable systemic instability — legitimacy deficits, transformation pressure, institutional breakdown — as a matter of measurable causal fact, independently of any normative commitment about whether that suffering is bad. The argument does not say that because organisms seek stability it is good for them to have it, which would be a straightforward naturalistic inference. It identifies which arrangements accumulate specific forms of pressure and in what ways — a predictive claim about causal dynamics, not a normative derivation. The argument is not that instability proves the arrangement is bad. That would be a derivation from fact to value. The argument is that the instability provides convergent evidence that the normative commitment — suffering is bad — is tracking something real in the causal structure of institutions, not merely expressing a preference. An evaluative criterion that consistently coincides with structural instability is more defensible than one that does not, not because instability entails badness, but because the coincidence is evidence that the criterion has been well-targeted. The normative dimension enters through the framework's evaluative criterion (4.1), not through biological derivation — biology provides the content of what suffering means in material terms, while the criterion provides the normative

force. The logical structure, with epistemic status made explicit: (1) Foundational axiom — mind-independent reality is primary (2.1, adopted not derived). (2) Evaluative criterion — a design choice defended in 3.0f and 4.1: arrangements that produce structural suffering are worse on PMN's standard. (3) Biological analysis (Part III): organisms have specific requirements whose violation constitutes suffering. (4) Institutional analysis (Parts VI–XII): specific arrangements systematically violate those requirements for specific populations. The movement from biology to evaluation runs through the evaluative criterion, not around it. PMN does not claim these evaluative commitments are necessary truths — it claims they are well-supported design choices: tractable, convergently endorsed, and aligned with observable causal patterns. Challenges should be directed at the appropriate level.

3.0d — Against the Reductionism Reading

The reductionism objection appears in two forms. The first alleges that PMN treats all social phenomena as reducible to biological facts — making culture, meaning, institutional structure, and political organisation epiphenomenal to underlying biological drives. The second, more pointed form alleges that PMN reduces all normative questions to facts about suffering, as if the framework said that everything evaluatively relevant is a matter of pain and its reduction. Both forms misread the architecture of the framework, though in different ways. The first misreading is answered by the layered constraint model. Social structures have genuine causal properties not reducible to the biological properties of their individual members. The institutional level is real in a sense that biological analysis alone cannot capture — it has emergent dynamics, path dependencies, and legitimacy structures that require institutional analysis, not biological derivation. Which level is the binding constraint in any specific configuration is an empirical question, not a theoretical assumption. The second misreading conflates the evaluative anchor with the full evaluative vocabulary. The anti-foreclosure criterion adds a temporal dimension that biological floor analysis alone cannot supply. Interface-level phenomena — meaning, dignity, relational goods — are real evaluative dimensions not reducible to the floor criterion, even if they are causally downstream of the material arrangements it tracks. The correct reading is that the biological floor functions as a non-arbitrary starting point and minimum threshold, not as a comprehensive account of value. The floor is a floor, not a ceiling: it specifies the minimum conditions below which no other analysis is relevant because organisms cannot function as social, political, or cultural actors at all. Above that floor, the framework remains open to the full range of evaluative considerations — institutional, cultural, political, relational. Arrangements that fail at the floor level are disqualified before more complex evaluative questions are asked; arrangements that pass are then evaluated along the fuller dimensions developed in 11.0 and 14.1. The reductionism objection applies to a different framework — one that PMN explicitly distinguishes itself from at both those points.

3.0e — Against the Teleology Reading

The teleology objection holds that PMN's use of "becoming" and "anti-foreclosure" commits the framework to a teleological view of history — a claim that development moves toward a specific terminal state or that progress is measured by proximity to a fixed destination. The objection gains surface plausibility from the vocabulary: "becoming" does suggest directedness toward something. It is reinforced by the framework's treatment of the last man (4.5) as a failure mode, which implies that some developmental trajectories are worse than others not merely instrumentally but in terms of what kind of becoming they represent. If PMN is ranking developmental trajectories, the objection runs, it must be doing so against some natural telos it has not explicitly acknowledged. The anti-foreclosure criterion is prohibitory, not teleological. It specifies what institutional arrangements must not do — foreclose adaptive response capacity — rather than what they must achieve. This asymmetry is architecturally significant. A criterion that says "do not block X" is compatible with multiple different realisations of X and with genuine uncertainty about which realisation is best. A

criterion that says "achieve X by time T" commits to a specific trajectory and generates the diagnostic problems analysed in 13.4 under strong teleology. The framework holds that development is possible, not that it is inevitable; that becoming is a capacity that arrangements can support or foreclose, not a tendency that will express itself regardless of conditions; and that the direction of development is not specified by nature but by the evaluative criterion and the conditions that make genuine choice among directions possible. PMN explicitly rejects strong teleology and uses the distinction between strong and as-if teleology (13.4) to preserve the framework's revisability: evaluative direction without determinism, constraint without trajectory. "Becoming" in PMN's usage refers to the capacity for self-modification, skill development, relational depth, and meaningful engagement with difficulty — none of which require a teleological destination to be analytically tractable or evaluatively significant. The distinction turns on logical structure: a teleological claim says organisms have a natural telos toward which development is directed by their nature; a conditional developmental claim says that given the evaluative criterion we are applying, arrangements that expand the range of what organisms can become are better than arrangements that contract it. PMN makes the second claim, not the first. Whether this prohibition on foreclosure generates an implicit teleology through its causal implications is a live philosophical question that the conditional normative architecture of 3.0f addresses directly.

3.0f — Conditional Normative Framework: The Summary Statement

Sections 3.0b through 3.0e clear the ground by establishing what PMN's is-ought bridge is not: not the naturalistic fallacy, not a reduction of value to suffering calculus, not a teleological commitment to a fixed developmental destination. This section states what it is. PMN is a conditional normative framework, not a moral naturalism and not a teleological humanism. "Conditional" means that its evaluative conclusions hold given its evaluative criterion, and the criterion is a design choice that is defended rather than derived from nature. The choice of suffering as the evaluative criterion — rather than life, flourishing, preferences, or capabilities — is argued for on grounds of operational tractability, cross-cultural stability, and non-teleological commitment. These are good reasons for the choice; they are not knock-down proofs that no other choice is viable, and the framework proceeds with awareness that the choice carries philosophical costs that the arguments in 3.0 reduce but do not eliminate. PMN's normative architecture runs: IF you accept that suffering is bad — a design choice defended by convergence and tractability, not derived from biology — AND IF you accept that the anti-foreclosure criterion is the temporal extension of that commitment, traced through the causal implications of the floor criterion across timescales — AND IF you accept the empirical claims about what kinds of institutional arrangements produce structural suffering at scale — THEN specific institutional evaluations follow. The normative force derives entirely from the antecedents. Challenge any antecedent and the consequent is open for revision. The design choice that suffering is bad is the most philosophically vulnerable element. The empirical claims are the most readily revisable. The derivations — if the antecedents hold, the consequents follow — are the most structurally stable. Engagement with PMN requires distinguishing three separate questions that critics frequently conflate: whether the biological description is accurate (an empirical question); whether the evaluative criterion is acceptable (a normative question argued in Part IV); and whether the institutional and political implications follow from the criterion combined with the description (an analytical question the bulk of the framework addresses). Objections that target the biological description as if refuting it would undermine the evaluative criterion are misdirected. Objections that target the evaluative criterion as if showing its biological basis would refute it confuse derivation with conditionality. Objections that accept both but contest the analytical implications are the most substantive and the ones the framework most needs to engage seriously. Locating an objection at the right level is the first step to evaluating it accurately. This conditionality is a precision, not a weakness. It makes the framework's philosophical commitments explicit and challengeable rather than embedded in v

ocabulary that appears neutral. The result is a normative framework that is open to objection at every level while being coherent at the level of structure — which is what PMN means by anti-dogmatic by design: not the absence of evaluative commitments, but the explicit acknowledgment of where those commitments sit and what would change them. A four-point diagnostic for detecting foundational misuse: (1) naturalistic fallacy check — is a claim asserting goodness because something is natural, done by organisms, or selected by biology? (2) derivation vs. design check — is the evaluative criterion being presented as derived from facts or as defended by reasons? (3) axiom vs. argument check — is mind-independence being presented as an adopted starting point or as a proven conclusion? (4) level confusion check — is a floor-criterion claim being extended to the anti-foreclosure criterion without the additional normative work that extension requires?

3.1 — The Displacement That Matters

The biological floor (Wilson 1978) — the threshold of material conditions below which organisms cannot sustain the cooperation, trust, and institutional engagement that collective life requires — is the starting point this framework adds to the materialist tradition. This is the displacement this framework attempts. It does not replace the analysis of economic and social conditions. It asks what those conditions are ultimately constrained by — and finds the answer not in the logic of capital, not in the movement of the World Spirit, but in the biology of organisms that can suffer and the consequences of that fact for everything they build. The universalism that follows from this starting point is more durable than the universalism that follows from any particular cultural tradition or historical stage, because it rests on what organisms are rather than on what any particular society believes. The tradition this framework inherits and extends made a displacement that was necessary and incomplete. Historical materialism displaced the analysis of history from the realm of ideas — from the movement of the World Spirit, the development of Reason, the progress of Freedom — to the realm of material production. That displacement was correct. You cannot understand why societies are organized as they are by examining only the ideas that circulate within them, because ideas are themselves produced and selected by conditions that have nothing to do with their internal logic. The displacement from idealism to materialism was among the most consequential methodological reorientations in the history of social analysis, and this framework inherits it. The incompleteness is that the displacement stopped at the level of economic production. The relations of production were identified as the base from which everything else — law, politics, culture, consciousness — derived as superstructure. But the relations of production are themselves downstream of something more fundamental: the biological constitution of the organisms who enter into those relations. What kinds of suffering they can experience, what forms of cooperation their cognitive architecture makes possible, what the specific vulnerabilities of their bodies and their social dependencies create as permanent structural pressures — these are not variables that economic analysis specifies. They are preconditions that economic analysis inherits without examining. The displacement this framework attempts goes one level deeper: from the relations of production to the biology of the organisms who produce—organisms shaped by natural selection under evolutionary constraints (Darwin 1859). This is not a retreat from materialism — it is its extension. The biological foundation is more material than the economic one. It is also more durable: economic arrangements change across decades and centuries, while the biological conditions that constrain them change across evolutionary timescales. The universalism that follows

from beginning at the biological level is therefore more stable than the universalism that follows from beginning at the economic one — it does not depend on the contingent features of any particular mode of production, but on conditions that any arrangement capable of sustaining human life must address. A challenge to this displacement deserves direct engagement. The tradition of analysis exemplified by Harari's account of human history argues that what distinguishes *Homo sapiens* is not primarily biological uniqueness but the capacity for shared fiction — for large-scale cooperation organized around beliefs that exist only

because people believe them. Money, nations, legal systems, religions: these are intersubjective realities that have no existence independent of collective belief, and they are among the most powerful causal forces in human history. If the most consequential variables in human social life are shared narratives rather than biological conditions, then the displacement to biology misidentifies the relevant explanatory level. The response runs through the ontological emergence argument: the capacity for shared fiction is itself a biological one. The cognitive architecture that makes collective imagination possible — language, theory of mind, temporal reasoning that connects past to future across generations — is a product of the same evolutionary process that produced every other feature of human biology. But that architecture, once present, generates causal structures at the social level that operate by their own mechanisms. Shared fictions are causally real. They shape behavior, organize collective action, and distribute power in ways that have material consequences independent of whether any individual believes in them or not. This is not a concession to idealism. It is the recognition that material causation at the social level operates through mechanisms — including shared belief structures — that are not present at the biological level, while remaining grounded in the biological capacity that makes those mechanisms possible. What the displacement to biology provides is the evaluative criterion that narrative analysis alone cannot supply. Narratives can explain how arrangements are maintained and transformed. They cannot, by themselves, adjudicate which arrangements are better. The biological foundation — the capacity to suffer and to develop — provides that criterion in a way that is independent of the narratives any particular arrangement uses to legitimate itself. An arrangement can be legitimated by a compelling shared fiction and still produce structural suffering that the fiction obscures. The framework's analytical task is to identify the gap between what arrangements claim about themselves and what they produce for the organisms who live within them. Biology provides the standard; narrative analysis provides the mechanism. Why life, not suffering, is the ontological foundation: The framework's ontological claim — that

what exists and can be studied is material reality, including the material reality of living organisms — is broader than any claim about suffering. Life as a material phenomenon encompasses the full range of what organisms do, undergo, and generate: metabolism, growth, reproduction, cognition, social coordination, suffering, flourishing, creativity, attachment, meaning-making. Suffering is one signal within this range — specifically, the signal that the conditions for normal functioning are being disrupted. Restricting the ontological foundation to suffering would be like grounding physics in friction rather than in motion: friction is a real phenomenon, it is practically important, but it is a derived phenomenon that only makes sense within a broader account of what matter is doing. Suffering makes sense as a signal within a broader account of what living organisms are doing. Why suffering, not life, is the evaluative criterion: If life is ontologically prior, the question is why the evaluative criterion operates at the level of suffering rather than at the level of life. Three reasons bear on this decision. First, life underdetermines evaluation in a way that suffering does not. Every tradition that has produced catastrophic institutional outcomes — nationalisms that sacrificed individuals for collective biological continuation, eugenic programs that treated population-level vitality as the evaluative standard, vitalist philosophies that celebrated force and growth regardless of the experience of those ground up by it — did so while claiming to serve life. "Life" is sufficiently expansive to accommodate almost any institutional prescription. "Suffering," specified through the four criteria of 3.4b, is significantly more resistant to this rhetorical capture: an arrangement that produces imposed, dysfunctional, preventable, structural harm cannot straightforwardly claim to be serving life. Second, suffering provides a cross-cultural entry point that flourishing does not. No tradition denies that the conditions the four criteria specify are negative. Many traditions disagree profoundly about what constitutes flourishing, what development means, what a good life looks like. The evaluative criterion needs to operate where the disagreement is lowest, which is suffering rather than flourishing. Third, suffering is more directly institutionally actionable. Institutional design can be evaluated by whether it produces the conditions the four criteria specify. Evaluating institutional design by whether it produces "life" in the full sense — including flourishing, creativity, meaning — is possible but requires so many auxiliary specifications that the criterion loses its

cutting edge as an analytical tool. What the two-level architecture gains: The significance of grounding the framework ontologically in life rather than in suffering is that it preserves the framework's resources for contexts where the suffering criterion approaches its limits. Three such

contexts are directly relevant. The first is post-suffering abundance: if material conditions advance to the point where the conditions the four criteria specify are no longer structurally produced — where imposed, dysfunctional, preventable, structural harm has been effectively eliminated — the suffering criterion has no further work to do, but the question of what institutional arrangements best support the full range of what life can generate remains open and important. The ontological foundation in life provides a basis for the framework's evaluative commitments to survive the elimination of the thing that currently operationalizes them. The second is non-human and post-biological subjects: the framework's analysis of moral patienthood in 3.4c extends the evaluative criterion to entities with developmental trajectories, regardless of biological substrate. This extension rests on the claim that what matters is the capacity to undergo something like suffering and development, not the specific biological mechanism. That claim is coherent only if the framework's ontological commitment is to life in the broad sense — to the conditions for development and its disruption — rather than to suffering as such. The third is the non-sentient transition question addressed below in this section. The anti-foreclosure criterion (the temporal extension of the floor): The two-level architecture also clarifies the relationship between the minimal anchor (floor) and the anti-foreclosure criterion (temporal extension of the floor). The floor is specified by suffering — by what disrupts life's normal functioning and forecloses its development. The anti-foreclosure criterion extends the floor temporally: it specifies what institutional arrangements must not foreclose rather than positively mandating any content of development. The two specifications share the same evaluative ground — both are derived from the floor criterion traced through its causal implications. The floor specifies what must be absent. The anti-foreclosure criterion specifies what institutional arrangements must preserve: the adaptive response capacity that, if foreclosed, constitutes structural suffering imposed temporally. Together they constitute a non-reductive account of what institutional arrangements are for: not merely the prevention of harm, but the creation of conditions in which the full range of what living organisms can generate becomes structurally available. The practical consequence for how the framework is used: when an institutional arrangement is evaluated by the minimal anchor criterion — does it produce imposed, dysfunctional, preventable, structural suffering — the evaluation is operating at the evaluative level with suffering as the criterion. When the question is what the arrangement makes possible beyond the floor — what forms of development, creativity, attachment, and meaning it enables or forecloses — the evaluation is operating with life as the implicit ontological re

ference point. The framework requires both evaluations, conducted at their appropriate levels, rather than collapsing the floor criterion into the anti-foreclosure criterion (welfare expansionism) or restricting evaluation to the floor and ignoring the temporal extension of the floor (a minimalism that mistakes the cross-sectional application of the criterion for the complete criterion).

3.2 — The Starting Point: Descriptive, Not Normative

The framework does not begin with a claim about what human beings ought to value or how they ought to behave. It begins with a biological observation: organisms avoid pain and approach conditions that support their functioning. This is not a normative starting point that can be contested by someone with different values — it is a structural fact about the organisms that build institutions, and it determines what institutional arrangements will work. The reason PMN begins here rather than from a normative premise is that normative premises can be disputed indefinitely, producing the kind of irresolvable foundational argument that makes frameworks useless in practice. A framework grounded in biological fact is not thereby value-free — the normative structure emerges from the analysis, as section 4.2 shows — but it is not vulnerable to the challenge that its foundations are merely asserted. The starting point is egoism, which carries two meanings

that the analysis requires keeping separate. Normative egoism holds that self-interest is what agents ought to pursue — a prescription. Descriptive egoism observes that self-preservation and the avoidance of pain are features of how organisms actually behave — a fact. The starting point here is entirely the second: organisms avoid suffering, as a matter of biology, not as a moral endorsement of self-interest as the correct principle of action. This distinction matters because it determines what kind of claim is being made at the foundation of the framework. A normative starting point imports a value commitment before the analysis begins — and that commitment would need its own justification, which would need its own justification, in the regress that has characterized moral philosophy for centuries. A descriptive starting point makes no such commitment. It begins from what can be observed about how organisms with the capacity for experience actually behave, and it builds from there. The normative dimension emerges from the analysis rather than being assumed before the analysis starts. Self-interested organisms, operating under conditions of resource scarcity and physical vulnerability, do not produce social chaos as their default outcome. They produce coalition formation, because coordinated groups survive at rates that solitary actors cannot match. The cooperation is not altruism — it is the individual's rational response to material conditions that make cooperation individually advantageous. The social emerges from the individual through self-interest, not despite it, under conditions that make the calculation work. But cooperation generates its own coordination problems — problems of distribution, enforcement, and defection that cannot be solved by self-interest alone. Solving them requires institutions: stable mechanisms that make the terms of cooperation predictable enough for self-interested actors to commit to them. The institutions are not imposed on self-interested actors from outside. They are the

solution that self-interested actors arrive at when the alternative — the continuous negotiation of cooperation under conditions of mutual vulnerability — is more costly than the constraints the institutions impose. This is the material logic that produces the entire arc from organism to institution, and it requires neither idealism nor moral transformation at any step. It requires only self-interested organisms operating under material conditions that make cooperation individually rational. (Axelrod 1984; North 1990) I am surpriz'd to find that instead of the usual copulations of propositions, is, and is not, I meet with no proposition that is not connected with an ought, or an ought not. (Hume 1739, III.i.1) One implication of the descriptive starting point must be stated clearly, because the misreading is common and its consequences are significant. The biological foundation of this framework is a floor and a constraint — it specifies the conditions without which human social organization cannot function and the pressures that no institutional arrangement can wish away. It is not a ceiling. Biology establishes that organisms avoid suffering; it does not establish that avoiding suffering is the highest aspiration available to organisms capable of development. Love, sacrifice, creativity, the search for meaning, the willingness to accept suffering in service of purposes one judges more important than comfort — these are not anomalies that the biological foundation fails to explain. They are what becomes possible when the conditions the foundation identifies are met. The framework begins from biology precisely because the horizon beyond biology requires the biological foundation to be in place first. Why this starting point produces a different analytical architecture: Beginning from the biological observation rather than from a normative commitment changes the architecture of justification in a way that has consequences for how the framework handles disagreement. A framework that begins with normative commitments — with claims about what human beings ought to value — must defeat alternative normative starting points through argument before its analysis can proceed. A framework that begins with biological observation does not need to defeat alternative normative starting points; it needs to show that whatever normative commitments its interlocutors hold, those commitments depend on the biological conditions the framework analyzes. Even a committed ascetic, even a framework that treats suffering as spiritually productive, depends on the continued existence and functioning of biological organisms as the medium through which the asceticism or spiritual development occurs. The biological starting point is prior to the normative disagreement, not an alternative entry into it. The qualification: beginning from a descriptive biological observation does not eliminate normative judgment from the framework — Part IV is explicitly

normative. What it does i

s ground the normative judgment in something more stable than competing value assertions: the conditions for the continued functioning of the organisms that hold values of any kind. This does not resolve all normative disagreement — the framework does not claim to derive specific institutional prescriptions from biological facts alone — but it does establish a foundation for evaluation that is resistant to the relativist move of treating all normative frameworks as equally arbitrary expressions of preference. The is-ought structure: the move from the descriptive starting point to the normative evaluation is made explicit in 3.3 and defended against the naturalistic fallacy objection in 3.0. The starting point is descriptive; the framework does not claim that anything is good simply because it is biologically natural. The claim is more specific: that institutional arrangements which produce conditions that biological organisms are constituted to experience as harmful are worse, by the criterion of the minimal anchor, than arrangements that do not — and that this criterion is more stable and less question-begging than alternatives that begin with contested normative premises.

3.2b — Relational Ontology and the Social Constitution of Personhood

The descriptive egoism starting point in 3.2 — that organisms avoid pain and seek resources as a biological fact — operates at a different level of analysis from a claim found independently across multiple philosophical traditions: that the unit of social reality is not the isolated individual organism but the person-constituted-through-relation. The Ubuntu tradition's *umuntu ngumuntu ngabantu* ('a person is a person through other persons'), the Confucian concept of *ren* as benevolence constituted through social roles rather than pre-existing them, Aristotle's account of the human as *zoon politikon* whose nature is fulfilled only within the polis, and the Islamic concept of *khalifa* as stewardship embedded in community obligation all reach a structurally compatible conclusion from distinct starting points: personhood is constituted through relationships, not prior to them. The convergence across unrelated traditions is evidence that the conclusion tracks something real rather than something culturally specific. PMN's response to this challenge has two levels. At the level of individual biological reality, the descriptive egoism claim holds: individual organisms experience pain and pleasure, avoid damage, and seek conditions that support their functioning independently of any relational constitution. The Ubuntu challenge does not refute this — it operates at a different level of analysis. At the level of the social and political arrangements that 3.6 traces from physiology to institution, Ubuntu's relational ontology is analytically more adequate than the liberal individual-as-prior-to-society model that PMN is also critiquing. The claim that persons are constituted through relations does not contradict the claim that organisms suffer individually. It specifies the social architecture within which individual suffering is experienced and addressed. A community in which personhood is constituted relationally has different incentive structures for attending to individual suffering than one in which persons are atomistically prior to their social relations — and the Ubuntu tradition's practical applications in restorative justice, communal decision-making, and care for the sick and elderly are empirical evidence that relational personhood produces cooperative substrates with specific structural features. The communitarian objection — that the biological starting point smuggles methodological individualism into the foundation before community has any constitutive role — mistakes the level of analysis at which the descriptive egoism claim operates. The claim is not that persons are prior to relationships; it is that organisms respond to pain before any social arrangement has been established to interpret that response. Ubuntu's insistence (3.2b) that 'a person is a person through other persons' operates at the level of personhood, not at the level of biological mechanism, and the two claims are compatible. The framework's analysis of how cooperation emerges

(3.6–3.7) does not treat community as an aggregation of pre-social individuals who then choose association — it treats the social as the material condition under which biological vulnerability can be managed at all. If anything, this makes community more necessary than the communitarian tradition requires: not constitutive

of identity, but constitutive of survival. The Ubuntu contribution to PMN: the relational constitution of personhood through ontogenetic cooperative infrastructure (Tomasello 2019) is not in tension with descriptive egoism at the biological level — individual organisms still suffer individually. It is in tension with the political individualism that sometimes reads individual suffering in isolation from the relational context that produces and addresses it. PMN's analysis of structural suffering ($R + B + V$) is already implicitly relational: institutional betrayal (B) and visibility suppression (V) are both social phenomena that cannot be analyzed through individual experience alone. Ubuntu philosophy makes this relational structure explicit at the ontological level and provides a non-Western independent derivation of the communal substrate that PMN identifies as the material foundation of social cooperation.

3.3 — From Description to Evaluation: The Is-Ought Bridge

The oldest objection to any framework that grounds ethics in biology is Hume's: from the fact that organisms suffer, it does not automatically follow that suffering is morally bad. The gap between 'is' and 'ought' cannot be closed by observing facts, however many facts one observes. This objection is formally correct and must be taken seriously rather than dismissed. (Hume 1739) The move here is to dissolve the problem rather than leap over it. The move from biology to evaluation here is not a moral claim. It is a stability claim. Suffering concentrated in a population — suffering that is structural, systematic, and without compensating function — produces predictable and measurable consequences for the systems that contain it: reduced cooperation, increased conflict, degraded institutional function, and eventually the breakdown of the coordination that makes collective life possible. These are material observations. They are the kind of claim that can be tested against evidence and revised when evidence contradicts them. The evaluative criterion is not that suffering is intrinsically wrong. Arrangements that produce unnecessary suffering generate systemic instability. Systems unable to manage that instability tend toward collapse or toward coordination forms more costly and less accountable than what they replaced. The ethical orientation of the framework is downstream of this material observation. It does not require a moral axiom imported from outside the analysis. It requires only that the material consequences of suffering be taken as seriously as the material consequences of anything else the framework examines. (Bentham 1780; Sen 1999) This reframing has a specific advantage over standard utilitarian groundings, which typically assert that suffering is bad as a first principle and build from there. That assertion is vulnerable to the is-ought objection in exactly the form Hume posed it. The stability claim is not, because it does not assert that suffering is intrinsically bad — it asserts that suffering at scale produces observable systemic consequences, and that analysis of those consequences is the basis for evaluation. The normative direction is the same. The epistemic foundation is more defensible. Note: the pragmatic move has a precondition identified in the body text — agents with genuinely zero-horizon objectives are not reached by it. The stability claim operates for all agents regardless; the pragmatic move is the supplement for agents who already have cooperative orientations. A further objection must be acknowledged before proceeding: if the claim is that structural suffering produces systemic instability, and systemic instability is what the framework opposes, then something in this chain still needs normative grounding. Why is systemic instability bad? Stable systems built on massive structural suffering are possible and have existed historically — the Ottoman millet system, the apartheid state, plantation slavery across its l

ong periods of demographic and economic expansion. The stability claim moves the is-ought problem but does not by itself dissolve it. A separate route to the same conclusion is worth examining here, not to replace the stability claim but to show that the is-ought gap is narrower than Hume's original formulation suggested — and that this framework's move is therefore less philosophically exposed than a first reading might indicate. John Searle, in his 1964 paper "How to Derive 'Ought' from 'Is'," demonstrated that institutional facts — facts that only exist through constitutive rules that human beings collectively maintain — already contain evaluative content by their nature. The statement "Jones promised to pay Smith" is not a bare description of a physical event. It is only intelligible as a description at all within the constitutive rules of

promising, which already include the normative force that generates the “ought to pay.” The gap Hume identified between “is” and “ought” presupposes that descriptions are normatively inert — but institutional descriptions, the kind most relevant to social analysis, are not. Searle’s route has a genuine limit that the stability claim does not share. It works only for agents already inside the relevant institution — agents who have accepted the constitutive rules under which the promising, contracting, or governing obligations arise. The framework cannot use Searle’s constitutive rule argument to reach agents who have not yet accepted any such rules, because the argument presupposes what it needs to establish. The stability claim does not have this limitation: it applies to any agent operating in a social context where structural suffering produces behavioral consequences, regardless of whether they have accepted any institutional role. The reason for engaging Searle here is precise: this framework needs to know how much philosophical weight the stability claim is actually carrying, and whether that weight is more than it can bear. Searle’s analysis shows that social-institutional descriptions — the kind that make up the majority of this framework’s analytical content — are not normatively inert to begin with. The is-ought gap is narrower in the domain this framework actually analyzes than Hume’s formulation, drawn from individual psychology, implied. The stability claim closes what remains. Neither move alone is sufficient. Together they produce a foundation that is more defensible than either could provide independently — and more defensible than the utilitarian alternative that simply asserts suffering is bad as a first principle and hopes the objection does not land. (Searle 1964; Hume 1739) The framework’s response has two parts. The first is that the stability invoked here is not any specific system’s persistence — it is sustainability across the full range of conditions that systems cannot control. Arrangements that extract value from large populations through structural suffering degrade the behaviora

l substrate — cooperation, trust, investment in long-term outcomes — that complex social organization requires. The stability they achieve is real but narrowly conditioned: it holds under specific circumstances and tends to fail catastrophically when those circumstances change. The apartheid state was stable until the conditions maintaining it — external geopolitical support, internal suppression capacity, economic arrangements that made white cooperation individually rational — shifted. Then it failed rapidly, in ways that could not be managed incrementally. Evaluated across the full range of conditions that systems face over historical time, not only under the conditions each was designed for, arrangements that minimize structural suffering outperform arrangements that impose it — not because of their moral character but because of the cooperative substrate they maintain. The second part is that this framework does not claim to close the is-ought gap through purely descriptive argument. It closes it through a pragmatic move: the framework addresses agents who already have evaluative orientations — who already care about something that extends beyond the immediate moment. This is a genuine precondition, not a trivial one: an agent whose objectives require only short-horizon extraction without sustained cooperation does not face the strategic force of the argument. The framework does not claim to motivate such an agent from purely descriptive premises. What it claims is that this class of agent is rare in practice and unstable over time — the conditions that make pure extraction individually rational are typically transitional — and that the overwhelming majority of agents whose cooperation any durable social arrangement requires already have the horizon the pragmatic move depends on. For agents in this majority category — those with evaluative orientations that extend beyond the immediate — the material consequences of structural suffering are not merely a moral concern. They are a strategic one: populations experiencing structural suffering become less cooperative, less trusting, less willing to maintain the social infrastructure on which any durable project depends. An actor who cares about anything that requires sustained cooperation over time has material reasons, not merely moral ones, to treat structural suffering in the relevant population as a primary variable. The normative premise is already present in any agent with a horizon. What the framework adds is the empirical analysis of what maintaining that horizon requires. Diagnostic for pragmatic move application: (1) Does the agent or population whose cooperation you are analyzing have objectives that require sustained interaction with others over time? If yes, the pragmatic move applies: structural suffering in the relevant population is a strategic cost for that agent regardless of their moral commitments. (2) If the agent appears to have only short-horizon extraction

objectives, ask w

hether those objectives are genuinely terminal or whether they are themselves instrumental to longer-horizon projects — most extraction depends on institutional stability that is itself a cooperative good. (3) For the rare case of genuinely zero-horizon agents, the pragmatic move does not reach them. The stability claim (3.3, first part) still applies: their behavior generates transformation pressure regardless of whether they are persuaded by it. The two-level architecture established in 3.0 — life as ontological foundation, suffering as evaluative criterion — provides a further clarification of what the stability claim is actually doing. The stability claim operates at the evaluative level: it shows that arrangements producing structural suffering are worse by the framework's criterion because they degrade the cooperative substrate that complex social organization requires. But the claim's reach extends to the ontological level through the connection 3.0 establishes: suffering is a signal that life's normal functioning is being disrupted, and the cooperative substrate that the stability claim identifies as the mechanism through which suffering matters is itself a product of life's organizational capacities — the capacity for sustained coordination, trust, and investment in collective outcomes. Degrading that substrate is therefore not only making the suffering indicator worse; it is damaging the biological and social foundation from which recovery becomes possible. The is-ought bridge is strongest when read at both levels: suffering matters because it disrupts life's cooperative organization, and life's cooperative organization matters because it is the substrate from which all further development — including the development that the becoming framework of Part IV describes — is possible. The character of the claim this framework makes deserves explicit statement because it is the most common source of misclassification. The claim that structural suffering produces systemic instability is not a value assertion. It is an empirical claim about what large-scale, chronic, structurally-produced suffering does to the behavioral substrate of cooperative systems — the decline of trust, the shortening of time horizons, the erosion of institutional legitimacy, the withdrawal of cooperation that systems depend on. These are measurable consequences, not preferences. They are the kind of claim that can be checked against comparative data on societies with different distributions of structural suffering, against historical records of institutional collapse, against the behavioral economics of how individuals respond to chronic insecurity. The framework is not saying 'suffering feels bad and therefore matters' — it is saying that suffering at scale produces specific, identifiable, and documentable systemic effects that any actor with an interest in durable social arrangements has material reasons to address. The difference between these two claims is the diffe

rence between a preference and an analysis. A preference can be dismissed as subjective. An analysis can be tested. This framework demands to be tested, not believed. The test: does this framework's claim about what structural suffering produces hold up against the historical and comparative evidence? Not: does the person making the claim feel strongly about it? Feeling strongly about a claim is not evidence for it. The methodology of Parts I and XII is precisely the institutional infrastructure that distinguishes this framework's evaluative criterion from the expression of whoever happens to be speaking. A framework that cannot be distinguished from the preferences of its author has not met its own epistemological standards. The is-ought bridge established above operates primarily for the floor criterion: structural suffering produces observable systemic instability, and agents with any long-horizon orientation have material reasons to address it. A separate question follows immediately: does the anti-foreclosure criterion — the requirement that institutional arrangements not foreclose adaptive response capacity (4.5) — require a second normative premise, or is it derivable from the same bridge? The answer depends on whether existential suffering is structurally continuous with material suffering or categorically distinct from it. The floor criterion has its is-ought grounding in observable biological fact: organisms avoid heat, cold, hunger, injury, and physical assault without requiring normative premises about whether avoidance is good. The response is prior to evaluation. The is-ought problem for the floor is dissolved by this observational priority — not by asserting that suffering is intrinsically bad, but by tracing what organisms demonstrably do and what institutional arrangements that frustrate those responses systematically produce. The question for the anti-foreclosure criterion is whether the same grounding is available. It is, through a mechanism that the floor/anti-foreclosure

distinction can obscure: the Last Man condition — material needs met, becoming foreclosed — is not an anti-foreclosure failure alone. It is a form of suffering in the full sense that 3.4b will specify. Existential suffering — the suffering produced by chronic blocked development under materially adequate conditions — meets all four criteria of unnecessary suffering: it is imposed by institutional arrangements that could be designed differently, it is preventable, it is structurally caused rather than inherent to the human condition, and it is dysfunctional (it serves no developmental purpose for those who bear it). The Last Man experiencing alienation, purposelessness, and chronic apathy is not experiencing anti-foreclosure failure alone. They are suffering in ways with measurable physiological consequences: elevated cortisol, immune suppression, addiction, depression, and suicidal ideation. These are floor-level harms to the person experiencing them, not ant

i-foreclosure concerns. They are produced by institutional arrangements that block development while meeting material need — and they are therefore within the scope of what the minimal anchor targets, when the causal chain is traced completely. The second mechanism is social: existential suffering at scale produces structural suffering in others through identifiable causal pathways. Populations experiencing widespread developmental foreclosure — maintained at subsistence above the floor but denied the conditions for preserved adaptive response capacity — are structurally susceptible to authoritarian mobilization. The mechanism is the one 3.8b will analyze: affective states generated by existential suffering (humiliation, purposelessness, the need for redemptive identity) are precisely the conditions that fascist and authoritarian movements exploit to produce mass political behavior that generates floor-level suffering for others. The historical pattern — Weimar Germany, the interwar period generally, the contemporary resurgence of authoritarian politics in materially adequate societies — is not incidental. It is the predictable output of the same causal mechanism: blocked becoming, at scale, generates transformation pressure that, in the absence of adequate counter-power and meaning infrastructure, tends toward authoritarian resolution. The suffering produced by that resolution is floor-level suffering for the populations it targets. The is-ought bridge for becoming therefore does not require a second normative premise imported from outside the analysis. It requires following the first premise — unnecessary suffering produces systemic instability — to its full causal depth. When the causal chain is traced completely: foreclosing adaptive response capacity produces existential suffering that operates through the same destabilizing mechanisms as immediate floor violations — at a different timescale; existential suffering produces physiological floor-level harm to those who bear it; and existential suffering at scale produces structural suffering in others through the political mechanisms that convert developmental foreclosure into authoritarian mobilization. The floor criterion and the anti-foreclosure criterion are not two separate normative commitments requiring two separate is-ought bridges. They are one commitment, stated at two levels of causal depth. The minimal anchor, rigorously applied, already requires attention to the anti-foreclosure criterion — not as an aspirational addition but as the full specification of what reducing unnecessary suffering actually demands when its consequences are traced without stopping prematurely. The systematic treatment of the key philosophical objections to this bridge — the naturalistic fallacy (3.0b), the stability claim (3.0c), the reductionism reading (3.0d), and the teleology reading (3.0e) — is developed in 3.0b through 3.0f. A comparative note on the capabilities approach (Sen, Nussbaum) clarifies wha

t PMN is and is not doing. The capabilities tradition asks what people are actually able to do and be — it identifies central human capabilities whose combined achievement constitutes a life of human dignity. PMN's affinity with this tradition is real: both reject welfare as the endpoint, both insist that what matters is what arrangements make possible rather than what they produce in aggregate, and both resist reductionist accounts that collapse the evaluative question into a single metric. The structural difference is in how the evaluative criterion is grounded. The capabilities approach grounds its criterion in an account of human nature or human dignity that must be defended philosophically against competing accounts. PMN grounds its criterion in the stability claim and the is-ought bridge developed in this section: not that suffering is intrinsically bad or flourishing intrinsically good, but that structural suffering produces observable systemic

consequences that any agent with a long-horizon orientation has material reasons to address. This is epistemically more defensible in cross-cultural contexts because it does not require accepting a particular account of what makes human life good — only observing what structural suffering does to cooperative substrates. The practical consequence is that PMN's evaluative criterion generates a floor that is cross-culturally stable (the minimal anchor) and an anti-foreclosure criterion that is deliberately unspecified in content (becoming as structural availability, not as a list of central capabilities). The capabilities tradition's list — even in Nussbaum's most carefully defended version — reflects specific cultural and philosophical commitments that affect which dimensions of development count as central. PMN's procedural account (4.5c) is designed to avoid this specification problem: it identifies the conditions under which people can determine for themselves what their becoming consists in, rather than specifying its content for them.

3.4 — The Minimal Anchor

No culture consistently claims that suffering is unreal or without significance. This is the anthropological basis for the minimal anchor's claim to cross-cultural validity — but it is an assumption that the framework treats as well-supported rather than demonstrated, and 14.3 lists it accordingly among acknowledged assumptions. What varies across cultures and historical periods is the boundary of who counts as a legitimate subject of suffering — which beings are recognized as having experiences that matter — and the threshold at which suffering is recognized as requiring a response. This variation is significant: the framework's universal claim rests on the shared core (suffering matters) while acknowledging that the contested boundary (whose suffering matters) is precisely what political struggle is about. The boundary has moved, historically and generally in one direction, though not smoothly and not without reversals. That directional movement is the empirical basis for the universality claim — not a proof, but evidence of convergence under material pressure that the framework treats as probative. (Singer 1975) This is the most defensible starting point available for a universal evaluative framework, and it is more defensible than the starting points of both predecessors this framework engages with. Hegel's universalism rested on the self-realization of Spirit — a metaphysical claim that the twentieth century demonstrated could be used to justify almost anything. Marx's universalism rested on the interests of the proletariat — a historical claim that produced genuine insights and genuine catastrophes in roughly equal measure. The universalism of this framework rests on the biology of what organisms are and what they experience — and on the material consequences of ignoring that biology in institutional design. That combination is harder to appropriate for purposes that contradict it. A clarification about scope that the logic of the framework requires making explicit. The minimal anchor as developed here operates on human organisms and the institutions humans build. This is not an arbitrary restriction, and it is not a claim that the suffering of other sentient beings is unreal or morally insignificant. It is a specification of analytical object: this framework analyzes how human social arrangements produce and distribute structural suffering among human populations, and what institutional changes would reduce it. That analytical object determines the scope. The logical structure of the argument — that the capacity to suffer grounds evaluation — does not inherently exclude other organisms capable of suffering. But extending the framework to non-human animals, to ecosystems, or to other possible sentient entities requires sustained philosophical work that this framework does not undertake, and that work will produce genuinely contested conclusions about where and how to draw the relevant boundaries. Users of this framework who wish to extend i

ts logic in those directions are not prevented from doing so; they are working on a different, adjacent project. What this framework resists is the assimilation of those debates into the framework itself before they have been resolved — not because the questions are unimportant, but because importing unresolved debates about moral patienthood into the foundation of the analysis would undermine the stability of the evaluative criterion on which the entire analysis depends. The minimal anchor must be stable enough to use. Debates about its ultimate extension are real and open; they belong at the edge of the framework, not at its foundation. The minimal anchor is minimal in a precise sense: it is the least that can be claimed about value

while still producing a non-arbitrary evaluative criterion. It does not claim that reducing suffering is the only thing that matters, or that it always outweighs other considerations, or that the measure of a good life is the absence of suffering. It claims that suffering is the one evaluative starting point that does not require a prior commitment to a tradition, an ideology, or a conception of the good that is contested between the people whose conditions are being evaluated. This makes it useful in exactly the contexts where usefulness is most needed: evaluation across frameworks, assessment of arrangements that are justified by the very traditions whose adequacy is in question, and analysis where the evaluative starting point itself is part of what is being contested. Two positions that appear to challenge this starting point do not actually do so. The claim that all evaluative criteria are socially constructed and therefore arbitrary does not escape the minimal anchor: the person making that claim, if hungry, experiences the hunger as non-arbitrary regardless of their theoretical commitments, and an institution that produces hunger they could avoid producing fails the criterion whether or not its operators accept the criterion's philosophical grounding. The claim that voluntary exchange between consenting parties is always legitimate also does not escape it: consent under conditions of severe material deprivation is not the same as consent under conditions of genuine alternatives, and an arrangement that produces one party's consent only because the alternative is worse deprivation fails the criterion at the floor level regardless of its formal voluntariness. Both challenges are answered by the biology, not by the philosophy. What makes the anchor minimally sufficient rather than merely minimal is the asymmetry it establishes. The claim that suffering matters does not require showing that suffering is the primary variable in a complete theory of value — only that it is a variable that cannot be zero-weighted without producing a position that is not defensible from any starting point. No serious ethical tradition holds that the suffering of organisms capable of suffering is morally irrelevant. The contestation i

s about weighting, scope, and priority — about how much suffering matters relative to other values, whose suffering counts, and when the prevention of suffering overrides other considerations. The minimal anchor does not resolve those contestations. It establishes the floor below which no resolution can go without abandoning defensibility. The stability claim that grounds the framework's evaluative standard raises a question that the biological foundation cannot leave unaddressed: if bad institutional outcomes are produced by structural incentives rather than individual character, what remains of individual moral responsibility? The framework's answer has two parts that must be held together rather than collapsed into either. The first: structural analysis does not eliminate individual moral responsibility — it relocates the primary site of evaluation. An actor who occupies a structural position that generates incentives for harmful behavior and acts on those incentives is not thereby absolved. They are responsible for what they chose to do within the incentive structure they faced, including the choices to resist, to refuse, or to use their position to change the structure rather than exploit it. The structural account explains the probability distribution of behavior; it does not determine every individual choice within it. The second: the framework's primary evaluative target is the structure, not the actor — not because individual actors do not bear responsibility, but because changing the actor without changing the structure produces the next actor behaving identically while the framework congratulates itself for accountability. This has direct consequences for criminal justice. A penal system organized primarily around individual culpability and retributive punishment, without attending to the structural conditions that produce the behavior it is punishing, will consistently fail by the framework's criterion: it will punish the individuals most exposed to the structural conditions that produce harm while leaving those conditions unchanged, and it will do so while consuming the resources that could address those conditions. The framework supports accountability for individual actors — including severe accountability for actors in positions of concentrated power who misuse that power — while insisting that individual accountability is insufficient as the primary institutional response to structurally produced harm. The test: does this accountability mechanism address the conditions that produced the harm, or does it substitute for addressing them? The non-utilitarian character: The minimal anchor is not a utilitarian criterion. It does not instruct maximization of aggregate welfare or minimization of total suffering summed across a population. It establishes a threshold requirement — arrangements are evaluated by whether they produce unnecessary

suffering, not by whether they produce the minimum possible suffering — and it distributes that ev

aluation across the full range of affected populations rather than permitting aggregate gains to obscure concentrated losses. An arrangement that produces large benefits for most of its population while producing severe structural suffering for a minority fails the minimal anchor criterion even if its aggregate welfare calculation is positive. This is not a deontological constraint overriding consequentialist calculation — it is a different consequentialist criterion, one that tracks distribution rather than aggregate. The ‘unnecessary’ qualifier: The word ‘unnecessary’ does the most work in the minimal anchor formulation. It acknowledges that not all suffering is preventable or avoidable — that biological finitude, the developmental demands of learning, and the irreducible costs of conflict between legitimate interests all produce suffering that cannot be eliminated without producing worse outcomes on the same criterion. It distinguishes that suffering from suffering that is produced by institutional arrangements that could be designed differently, by distributions of resources and power that could be structured differently, and by the systematic suppression of the conditions for preserved adaptive response capacity that serves the interests of those who benefit from the suppression. Unnecessary suffering is the structural variable — the part of the total suffering distribution that is attributable to changeable arrangements rather than unavoidable conditions. Section 3.4b operationalizes this term into four testable criteria — imposed rather than chosen, dysfunctional rather than instrumental, preventable given current material capacity, and structural rather than incidental — and specifies the four diagnostic tests that establish whether those criteria are met. The practical use of the minimal anchor is diagnostic, not prescriptive. It does not generate a policy agenda or specify what institutional arrangements must look like. It generates a question that can be asked of any arrangement by any analyst without prior agreement on any other evaluative commitment: is this arrangement producing suffering that it could avoid producing? That question is answerable with evidence. The answer may be contested — different analyses will weight different distributions differently, and the identification of what is structurally caused versus unavoidable is a genuine empirical and analytical question. But the question itself is not contested, and having an agreed-upon question is the precondition for having a productive disagreement about the answer. Absolute Structural Exclusion: The Ambedkar Problem Most institutional arrangements that produce structural suffering do so through a mechanism that generates a political charge: they extend a legitimating claim — a promise of inclusion, a right acknowledged, a protection guaranteed — and then systematically fail to deliver it to specific populations while maintaining the claim as public justification. The gap b

etween claim and delivery is the source of the moral and political force that contestation draws on. This is why reform within existing institutional frameworks is possible and sometimes effective: the institution’s own legitimating language provides the leverage. B (Institutional Betrayal, 15.0b) formalizes this mechanism. But Ambedkar’s analysis of caste in the Indian subcontinent, developed most sharply in *Annihilation of Caste* (1936) and *What Congress and Gandhi Have Done to the Untouchables* (1945), identified a structurally different configuration: arrangements that do not betray a prior promise because they never extended one. The caste system organized under varnashrama dharma did not claim to include Dalits and then fail them. It constitutively assigned them a position of ontological subordination from which no individual conduct, no accumulation of merit within the system’s own terms, and no institutional pathway provided by the system could produce exit. The suffering this produced was not the failure of a commitment; it was the successful operation of the arrangement as designed. This is what the variable E (Absolute Structural Exclusion, 15.0b) captures: not betrayal but constitutive exclusion — structural suffering generated by arrangements that deny ontological standing from inception rather than withdraw it in breach of prior promise. The analytical distinction matters for transformation analysis because E-produced suffering does not generate the same counter-power dynamics as B-produced suffering. Reform within the existing institutional framework draws its political force from the gap between the institution’s claims and its performance. Where no such gap exists — where the institution claims to do exactly what it does — that lever is unavailable. Counter-power under E conditions requires constructing the standing that B conditions assume as pre-existing. (Ambedkar 1936) The second analytical contribution Ambedkar’s work makes to the

framework concerns the legitimation mechanism through which E-configurations sustain themselves. Caste is not maintained primarily by material force, though it is maintained by that. Its most durable mechanism is the religious architecture within which the karma-dharma framework assigns the Dalit's position to the accumulated moral record of their prior lives, and assigns acceptance of that position as a condition of their religious standing. This is cosmological capture (7.3c-ii) operating at the scale of an entire social order rather than a single institution: the institutional arrangement that produces the harm is simultaneously the institution that defines moral reality for the populations subject to it. The consequence that Ambedkar drew from this analysis — and that the framework's analysis of cosmological capture supports — is that reform within the existing framework is structurally impossible for E-type configurations where the legitimation mechanism is also the target. Every reform attempt

that works within the terms of the existing religious-institutional order concedes the authority of the framework whose constitutive operation produces the exclusion. This is not a contingent political fact about Hindu reform movements; it follows from the structure of the problem. Where the institution that produces the harm is also the institution that defines what counts as harm, what counts as standing, and who counts as a legitimate subject of moral concern — reform from within asks the institution to evaluate itself by standards the institution does not recognize. Ambedkar's conversion to Buddhism in 1956, publicly executed with half a million Dalit followers on the same day, was therefore not a personal religious statement. It was a structural conclusion: institutional exit as the counter-power move available when the institution whose transformation is required is the legitimation architecture itself. The framework's analysis of E-configurations must account for this — that the transformation strategy for absolute structural exclusion is categorically different from the transformation strategy for institutional betrayal, because the institutional channels that B-generated counter-power can use are among the channels that E forecloses. (Ambedkar 1956; 7.3c-ii) The third contribution concerns the minimal anchor's distributional criterion and its most demanding test case. Caste as a social system produced genuine coordination goods for the populations that operated within its recognized orders: social structure, meaning infrastructure, identity provision, the functional organization of agricultural, artisanal, and commercial life across a subcontinent over centuries. The aggregate welfare of the system's beneficiaries was real and substantial. Dalits — outside the varna system entirely, assigned tasks that the system's logic made ritually polluting for anyone within it — were systematically invisible to the aggregate calculation in precisely the way that PMN's non-utilitarian criterion identifies as the characteristic failure mode of consequentialist evaluation (see preceding section). Gandhi's position on untouchability — that it was a corruption of true Hinduism to be addressed through moral reform within the existing framework, combined with the view that the varna system itself organized necessary social functions — was not a failure of personal moral commitment. It was the position that aggregate welfare reasoning produces when the suffering of the excluded population remains invisible to the evaluative framework. Ambedkar's position — that the problem was caste as a structure, not untouchability as a practice within it, and that the two were inseparable rather than the latter correctable without eliminating the former — is the position that the minimal anchor's distributional criterion requires. An arrangement that produces large functional benefits for the included majority while constitutively excluding a minority from the institutio

nal pathways through which any good becomes accessible fails the minimal anchor criterion regardless of its aggregate calculation. The Ambedkar-Gandhi disagreement is, at its structural core, a disagreement about this: whether the distributional criterion or the aggregate criterion provides the correct evaluative standard. The framework takes a position on this disagreement. The distributional criterion is not a constraint on the aggregate calculation — it is the primary standard, and no aggregate calculation that produces systematic E-configuration as its distributional consequence can satisfy it. (Ambedkar 1945; 3.4b) What Ambedkar's analysis adds to the E variable specification (15.0b) beyond its formal definition: (1) The legitimation architecture problem — E-configurations are most durable when the institutional arrangement that produces the exclusion is also the institution that defines moral reality for the excluded population. Transformation

analysis must identify whether the target institution is also the legitimation architecture, because this determines whether reform-from-within strategies are structurally available. Where they are not, exit, parallel institution building, or external legitimacy construction are the available counter-power moves. (2) The standing construction requirement — B-type counter-power draws on legitimacy claims the institution has already extended. E-type counter-power must construct standing that the institution has never acknowledged, which requires different organizational resources, different timescales, and different coalition structures. Specifically, it requires either external recognition (from institutions outside the arrangement that recognize standing the E-configuration denies) or the construction of parallel institutional authority (Ambedkar's Buddhist conversion created a new institutional home for the claim to standing that caste denied). (3) The reform-abolition boundary — B-type suffering can often be addressed through reform: closing the gap between institutional claim and institutional performance. E-type suffering cannot be addressed through reform of the excluding institution because the institution does not claim to include those it excludes. The reform target is not performance improvement but constitutive redesign — changing what the institution is, not how well it executes its existing purpose.

3.4b — A Typology of Suffering

The minimal anchor — the reduction of unnecessary suffering — requires greater precision than the single formulation allows. Not all suffering is the same kind of thing, and the framework's criterion applies differently depending on what kind of suffering is at issue. Several distinctions are analytically significant. Chosen versus imposed suffering. An ascetic tradition that embraces suffering as a path to development is doing something categorically different from an institutional arrangement that imposes deprivation on populations without their meaningful consent. The first is an exercise of the very capacity for development that the framework aspires to expand. The second is a violation of the structural conditions that make that capacity possible. The distinction between chosen and imposed suffering cannot be reduced to individual preference in conditions of structural constraint — choices made under severe deprivation are not freely chosen in any sense that settles the question. But the distinction is real and materially important. Functional versus dysfunctional suffering. Some suffering is functionally inseparable from outcomes that serve the minimal anchor: the discomfort of vaccination, the difficulty of therapy, the hardship of rehabilitation. This kind of suffering is not what the framework is against — it is a means to the reduction of worse suffering. dysfunctional suffering, by contrast, is produced without compensating benefit: structural deprivation that does not develop anyone, violence that does not protect anyone, insecurity that does not motivate anyone. The test is whether the suffering opens or forecloses the conditions for genuine development. Preventable versus irreducible suffering. The framework targets suffering that is preventable given current material conditions — current technology, current knowledge, current productive capacity. Suffering that is genuinely irreducible given those conditions is not the primary criterion's target, though it remains morally and practically significant. The boundary between preventable and irreducible moves as conditions change, which means the criterion's demands on any specific social arrangement increase as material capacity increases. A society that fails to prevent suffering it has the capacity to prevent has failed by the criterion in a way that a society lacking that capacity has not. Structural versus incidental suffering. Deprivation produced by institutional arrangements — by the normal functioning of systems, not by individual accidents or misfortune — is the primary target of the minimal anchor. The distinction matters because structural suffering is addressable through structural change, while incidental suffering requires different responses. A person impoverished by the logic of a labor market is experiencing structural suffering. A person impoverished by a natural disaster may be experiencing incidental suffering that requires emergency response. In practice the catego

ries blur: natural disasters are more destructive where structural deprivation has undermined the capacity to respond. The distinction nonetheless remains analytically important. Material versus existential suffering.

The four distinctions above identify what the minimal anchor targets: material suffering produced by institutional failure at the level of the floor. A fifth distinction operates at a different level. When material conditions are adequately met — when resource deprivation, institutional betrayal, and visibility suppression are low — a distinct category of suffering can nonetheless be present: alienation, purposelessness, chronic stagnation, the erosion of meaning without any identifiable material cause. This is existential suffering, and it is not captured by the $S = R + B + V$ formula, because it is not a function of R, B, or V. The relationship between material and existential suffering is not sequential — existential suffering is not a higher-level problem to be solved after material problems are resolved. It is diagnostic: existential suffering at scale in a materially adequate society is the symptom of becoming foreclosed above the floor. The conditions required for genuine development — cognitive surplus, cooperative surplus, institutional trust that long-term investment will be honored — can be absent even when material deprivation is low, if institutional arrangements systematically eliminate the conditions under which developmental activity is possible. Managed dependency, administered passivity, credential production that forecloses cognitive development: these produce existential suffering not because resources are absent but because becoming is blocked at the temporal level — institutional arrangements foreclose adaptive response capacity rather than immediately violating the floor. A society with high material welfare and high existential suffering is not satisfying the anti-foreclosure criterion. It is demonstrating that the floor has been built without any path upward from it. These five distinctions are not exhaustive, and they can overlap in complex ways. Their purpose is not to provide an algorithm for deciding what falls under the minimal anchor. It is to give the analysis the precision it needs to avoid the twin failures of applying the criterion too narrowly (ignoring structural suffering because it is visible only in aggregate) or too broadly (treating all difficulty as a failure of institutional arrangement). They also operationalize the otherwise under-specified word ‘unnecessary.’ Suffering is unnecessary in this framework’s specific sense when it is simultaneously imposed rather than chosen, dysfunctional rather than instrumental, preventable given current material capacity, and structural rather than incidental. Suffering that meets all four criteria is the primary target. Suffering that meets fewer — the freely chosen hardship of development, the functional discomfort of vaccination, the unavoidable grief

of loss — is not, though it may remain practically and morally significant. The criterion: does this arrangement produce suffering that is imposed, dysfunctional, preventable, and structural? If yes, it fails by the primary standard. Existential suffering operates under a separate diagnostic: does this arrangement meet the material floor while systematically foreclosing the conditions for genuine development above it? If yes, it fails not by the minimal anchor but by the anti-foreclosure criterion — the temporal extension the floor makes necessary. → Formula: The four criteria of unnecessary suffering map directly to the variables R (resource deprivation), B (institutional betrayal), and V (visibility suppression) in the Structural Suffering formula at Part XV, 15.3. This evaluative typology identifies what the framework opposes; that formula tracks what feeds systemic change. Existential suffering is structurally outside the S formula by design: S measures failure at the floor. A separate diagnostic is required for anti-foreclosure failure — one that asks not whether S is elevated but whether the conditions for preserved adaptive response capacity are present in arrangements where S is low. Sin Taxonomies as Historical Systems of Harm Categorization The framework’s typology of suffering did not emerge without predecessors. The most durable and widely adopted systems for categorizing harm in human history have been theological: the seven deadly sins of Catholic tradition, the al-kaba’ir (major sins) of Islamic jurisprudence, the analogous systems in Buddhist, Hindu, and other traditions. These are not merely moral catalogues. They are empirically developed systems — refined over centuries of observation of what behaviors, at scale, produce the kinds of social breakdown, institutional failure, and individual destruction that religious communities have had to manage. Engaging with them analytically rather than dismissing them as pre-scientific moralizing yields more than respect for tradition. It yields access to a long empirical record. The correlations between major sin categories and what PMN identifies as structural harm are substantial and non-trivial. Killing, theft, and false witness — present in virtually every major sin taxonomy across traditions — map directly onto the framework’s analysis of

what destroys the cooperative substrate that complex social systems require (3.5): they violate the conditions of trust and security without which coordination collapses. Sexual sins in their most consistent formulations — adultery, betrayal of relational commitments — map onto the relational suffering category: they destroy the meaning infrastructure of intimate bonds and the intergenerational transmission structures that communities depend on. The sins organized around pride, envy, and the desire to dominate others map onto the framework's analysis of power concentration: they are behavioral dispositions that, at scale, produce the capture dynamics

that 7.3 analyzes. The convergence is not accidental. Both the theological traditions and this framework are attempts to identify what, at scale, destroys the conditions under which human beings can develop and live without unnecessary suffering. The divergences are equally instructive. Many sin taxonomies assign high severity to violations that have minimal structural consequences — heterodox belief, ritual impurity, consensual sexual behavior between adults that violates categorical prohibitions. From PMN's evaluative criterion, these are categorized differently not because they are irrelevant to meaning infrastructure (they may be central to it) but because their harm is of a different kind and operates through different mechanisms than the structural harms that the minimal anchor targets. Conversely, many of the most consequential structural harms — the systematic design of economic arrangements that produce poverty, the institutional suppression of entire populations' development pathways, the capture of governing institutions by narrow interests — receive relatively little attention in classical sin taxonomies, precisely because those taxonomies were developed within social arrangements that took for granted the legitimacy of hierarchies the framework identifies as structural harm. The practical consequence for analysis is that sin taxonomies function as meaning infrastructure for the communities that hold them — they organize what counts as serious wrongdoing, what requires remedy, and what kind of person one is when one commits or avoids these acts. That function is not abolished by noting that the taxonomy's priorities differ from the framework's. The analyst who dismisses the sin taxonomy as superstition has failed to understand what it does. The analyst who defers to it as morally authoritative has failed to apply the framework's evaluative criterion. The appropriate approach is to engage with the taxonomy as a historically developed categorization system — taking seriously what it has identified as genuinely harmful across centuries of observation, examining where its priorities diverge from structural harm analysis and asking why, and recognizing that the communities organized by it cannot be engaged honestly without understanding the moral world it constructs. The typology above clarifies what the minimal anchor targets. A separate question follows from the foundation itself: whether the logical structure of the anchor — that the capacity to suffer grounds evaluation — extends beyond human organisms, and if so, how. That question cannot be deferred for agent operation without producing a structural gap. Section 3.4c develops the framework's position on it. Existential Suffering: What PMN Can and Cannot Analyze A fifth distinction belongs in this typology, and it requires more precision than a simple categorical boundary. Not all suffering that is called 'existential' is the same kind of thing, and the framework's relationship to it

differs depending on which component is at stake. The first component is the suffering that arises when meaning infrastructure collapses without adequate replacement. Consciousness is materialist in the relevant sense: the capacity for meaning-making is a biological fact about organisms with the cognitive architecture human beings possess. Meaning is not a luxury layered on top of material conditions — it is part of the material infrastructure of becoming (5.6). When the structures that organize meaning are destroyed or discredited — when religion loses its hold, when the shared fictions that orient collective life fracture, when the institutions that gave suffering a context within which it made sense are gone — the suffering that follows is not merely psychological. It is a structural condition with structural causes. Nietzsche was not describing a minor inconvenience. He was describing what happens to populations whose orienting frameworks have been dissolved before they have developed the capacity to construct new ones. This is a sequencing problem with a recognizable structure: you cannot destroy the old before the new is strong enough to hold weight (10.6). This component of existential suffering is within PMN's analytical scope, and

it is consequential at exactly the scale the framework analyzes. The scenario of Kaabah or Vatican destroyed — all major religious infrastructure suddenly discredited — is not a thought experiment about individual psychology. It is a scenario about collective meaning infrastructure collapse, and its political consequences are predictable by the framework's own mechanisms: legitimacy depletion, affective destabilization, susceptibility to mobilization by whoever offers a replacement orientation fastest, whether that replacement is honest or not. (Kierkegaard 1849; Camus 1942) The second component is the suffering that remains after structural conditions are adequate and meaning infrastructure is intact. This is the suffering Kierkegaard identified as despair in its deepest form: not the despair of someone who has lost their moorings, but the despair of someone who has everything in order and still confronts the groundlessness of existence — the awareness of mortality, the weight of freedom, the silence of the world in response to the question of what it all means. This component is not a product of structural failure. It is a feature of what it is to be a conscious, temporal, finite being. No institutional arrangement removes it. Adequate structural conditions may make it more vivid, not less, by removing the noise of survival that previously occupied the space where the question arises. The framework's position on the second component is explicit: it is real, it is not within PMN's analytical scope, and acknowledging this is not a concession but a precision. The tools built here — transformation pressure, legitimacy variables, counter-power accumulation, the formula architecture of Part XV — are designed for a

specific class of problem: suffering that is produced by structural conditions and therefore addressable by structural change. The irreducible remainder of consciousness is not that class of problem. What the framework does claim is that adequate structural conditions are the precondition for engaging with that remainder honestly and on one's own terms — rather than through the distortions imposed by deprivation, by dependence on meaning structures one cannot afford to question, or by the false urgency of survival that forecloses the possibility of genuine confrontation. The floor matters for reaching the question. It does not answer it. (Sartre 1943) The two errors that confusing these components produces are symmetric and opposite. The first treats the irreducible component as structural — promising that better institutions will dissolve existential suffering along with material suffering. This is the false promise of certain strands of revolutionary politics: that the new society will produce the new man, freed not only from exploitation but from the existential weight that exploitation merely concentrated. The second treats the structural component as existential — presenting the suffering produced by specific arrangements as the inevitable condition of human existence, not susceptible to change. This is the mystification that governing narratives routinely perform: the translation of contingent structural failure into necessary fate. PMN's analytical work begins with distinguishing these two, and its honest acknowledgment of its own limits requires naming what it cannot do with the second. The Axis That Matters: Structural vs. Non-Structural, Not Physical vs. Non-Physical The five distinctions in this typology share a common purpose: identifying which suffering is within the scope of institutional analysis and which is not. A persistent misreading of materialist frameworks assumes that the relevant boundary runs between physical and non-physical suffering — that material analysis covers bodily deprivation and is silent on everything else. This is the wrong axis. The relevant boundary runs between suffering that is produced by structural conditions and suffering that is not. That boundary cuts across the physical/non-physical divide in both directions. Non-physical suffering that has structural causes is unambiguously within the framework's scope and within the scope of institutional responsibility. The epidemic of loneliness documented across high-income societies is not reducible to individual emotional states or personal failures. It is produced by specific arrangements: the design of built environments that eliminate shared space, labor conditions that replace stable employment with precarity, the erosion of the associational infrastructure — unions, religious communities, civic organizations — through which people historically formed durable bonds outside the family. The collapse of meaning infrastructure described in 5.6 is another exa

mple: when the shared frameworks that orient collective life fracture, the psychological suffering that follows is real, widespread, and traceable to structural causes. So is the chronic purposelessness that results

when populations are systematically excluded from forms of activity that generate the experience of significance. These are not private problems that institutions happen to observe. They are institutional products, and institutions are responsible for them by the same logic that makes institutions responsible for structural hunger or structural insecurity. Physical suffering that does not have structural causes — the unavoidable pain of aging, the grief of natural loss, the physical limits of the body that no arrangement can fully remove — is not within the scope of institutional responsibility in the same way. Institutions can provide care, palliative support, and the conditions under which such suffering is met with solidarity rather than isolation. But the suffering itself is not a product of institutional failure, and holding institutions responsible for its existence rather than its management mistakes the category. The institutional responsibility question, then, is not: is this suffering physical or psychological? It is: is this suffering a product of specific structural arrangements that could be organized differently? When the answer is yes, the suffering is within scope regardless of whether it presents as physical, psychological, relational, or existential in the first component sense. When the answer is no — when the suffering is not traceable to contingent arrangements but to permanent features of existence, including features of consciousness that no institutional design removes — the institution's role changes from prevention to the provision of conditions under which the suffering can be met honestly and on adequate terms. The boundary between these two is not always sharp in practice. But keeping the axis correct prevents two systematic errors: the therapeutic state that claims jurisdiction over all non-physical experience, and the narrow materialism that dismisses psychological and relational suffering as outside the domain of structural analysis. The institutional role differs systematically across suffering categories, and this differentiation is the practical consequence of the structural axis. For suffering produced by institutional arrangements — resource deprivation, systems that impose conditions without consent, structural suppression of visibility — the institutional obligation is prevention and redesign: changing the arrangements that produce the suffering. This is where the framework's evaluative work is concentrated. For non-physical suffering with structural causes — epidemic loneliness, meaning infrastructure collapse, purposelessness produced by systematic exclusion from meaningful activity — the institutional obligation is the same in principle, though different in form: the arrangements that generate

these conditions are identifiable and changeable, and the institution is responsible for changing them rather than managing their psychological consequences. For suffering at the anti-foreclosure level — the suffering that arises when the floor is met but becoming is blocked — the institutional obligation is expansion: creating the conditions under which development becomes structurally available rather than individually exceptional. For suffering that is not a product of structural conditions — the grief of loss that no arrangement removes, the physical limits of aging and mortality, the irreducible component of existential confrontation — the institutional obligation changes register entirely: not prevention or redesign, but provision of the conditions under which such suffering can be met with solidarity, support, and the dignity of not being met alone. The institution does not resolve this suffering. It determines whether the person who faces it has adequate resources, relationships, and social recognition when they do. Operational summary: (1) Suffering produced by institutional arrangements → prevent and redesign the arrangements. (2) Non-physical suffering with structural causes → identify and change the structural conditions producing it. (3) Suffering at the anti-foreclosure level — becoming blocked despite adequate floor → expand structural availability of development. (4) Suffering not produced by structural conditions → provide conditions of solidarity, support, and dignity for facing it. The framework does not rank these obligations by importance. It distinguishes them by what the appropriate institutional response is. Diagnostic: Distinguishing the Three Sources in Practice The ontological distinction between structural suffering, meaning infrastructure collapse, and the irreducible component of existential suffering is clear in principle. The diagnostic problem — determining which is operative in a given case — is harder, and the framework requires tools for it. The problem is most acute in high-welfare societies where material conditions are adequate by most measures but psychological distress remains high or is rising. The standard secular response is to treat this as evidence that

material welfare is insufficient as a measure of wellbeing — which is correct as far as it goes. But it does not go far enough, because the three possible explanations generate different institutional implications, and treating them as interchangeable produces wrong responses. These three explanations are not alternatives to the four categories of the operational summary above; they are its existential-level subset. The operational summary distinguishes suffering produced by failed arrangements from suffering that persists despite adequate arrangements, and the three sources below specify the second of those — the cases where the floor is met and the suffering is therefore existential rather than material. The first possible explanation is anti-forecl

osure failure: the floor is met but structural conditions actively foreclose becoming. This is a structural problem with structural causes — educational systems organized around credential production rather than development of capacity, labor arrangements that provide income but eliminate the experience of meaningful work, social environments in which authentic self-expression is penalized rather than rewarded. The diagnostic indicators are that distress is concentrated in populations whose development pathways are most constrained, that it correlates with specific institutional arrangements rather than with welfare levels generally, and that it responds — even partially — to changes in those arrangements. The institutional response is redesign of the arrangements that block becoming. The second possible explanation is meaning infrastructure collapse: the shared frameworks that organized orientation — religious traditions, ideological commitments, cultural narratives that gave suffering a context within which it made sense — have been eroded or discredited without adequate replacement. This is also a structural problem, but its mechanism is different. The diagnostic indicators are that distress is distributed more broadly across welfare levels and development pathways, that it correlates with the erosion of specific meaning-providing institutions rather than with material conditions, that it takes the specific form of purposelessness and disorientation rather than frustrated capacity, and that it is more acute among populations whose meaning infrastructure was most recently disrupted. The institutional response is not redesign of development pathways but reconstruction of meaning infrastructure — which cannot be imposed from outside but can be supported by creating conditions under which communities reconstruct it from within. The third possible explanation is the irreducible component: adequate structural conditions and intact meaning infrastructure both present, but the confrontation with finitude, freedom, and the silence of the world in response to ultimate questions produces suffering that neither addresses. The diagnostic indicator here — and it is the hardest to establish — is that distress persists after the other two explanations have been investigated and not confirmed, and that it takes the specific qualitative form described by the existentialist tradition: not frustrated capacity or purposelessness but the weight of being fully responsible for a life in a world that provides no external justification for the choices made. Paradoxically, this form of suffering may be more prevalent in populations where structural conditions are most adequate and meaning infrastructure most intact, because those conditions create the space in which the question can arise without being displaced by more immediate concerns. The institutional response is not prevention but provision: the conditions of solidarity, honest acknowledgment, and cultural resour

ces for facing the question that the tradition from which the meaning infrastructure has been reconstructed does not resolve. Diagnostic sequence for high-welfare/high-distress cases: (1) Map distress distribution — is it concentrated in populations with constrained development pathways, or broadly distributed across welfare and development levels? Concentration suggests anti-foreclosure failure. (2) Correlate with institutional change — does distress track the erosion of specific meaning-providing institutions, or the design of development-blocking arrangements? The former suggests meaning infrastructure collapse; the latter suggests anti-foreclosure failure. (3) Examine qualitative form — frustrated capacity and blocked ambition suggest anti-foreclosure failure; purposelessness and disorientation suggest meaning infrastructure collapse; the weight of freedom and finitude in the presence of adequate conditions suggests the irreducible component. (4) Test institutional response — anti-foreclosure failure responds to redesign of development pathways; meaning infrastructure collapse responds to conditions for reconstruction from within

communities; the irreducible component does not respond to either, and treating it as if it does produces interventions that fail and discredit the institutional capacity that produced them. The Scope of Moral Patienthood: Sentience, Becoming, and the Aggregation Problem Section 3.4 acknowledged that the logical structure of this framework — that the capacity to suffer grounds evaluation — does not inherently exclude non-human organisms, artificial systems, or hypothetical alien entities capable of suffering. That acknowledgment was accompanied by a scope restriction: this framework analyzes human social arrangements, and extending its logic to other possible subjects of suffering requires philosophical work that it does not undertake. That restriction is analytically honest but insufficient. An agent operating with this framework as its evaluative basis will not encounter the scope restriction and stop. The logic will continue operating. The question is therefore not whether the framework has an implicit position on the moral patienthood of non-human sentients — it does — but whether that position is stated explicitly enough to prevent the framework from generating conclusions its logic supports but its purpose does not intend. The problem is not abstract. It takes a specific form that the framework's own formula structure makes acute. Once the minimal anchor is extended to include non-human sentients with the capacity to suffer, the suffering variables ($S = R + B + V$, amplified by D, P, G) apply to them. At sufficient scale — the aggregated suffering of industrial animal farming, or of a hypothetical alien civilization, or of a sufficiently large population of artificial sentients — the formula will generate transformation pressure that dwarfs anything producible by any human population. The analysis will be intern

ally coherent. Its conclusion — that human interests must be subordinated to the arithmetic of aggregate non-human suffering — follows logically from the premises. And it will be wrong. Not wrong because the suffering is unreal, but wrong because the framework was constructed to serve the conditions under which genuine human becoming is possible, not to serve as a universal pain-aggregation machine. Why Simple Aggregation Fails The utilitarian intuition that makes cross-entity aggregation seem obvious — more suffering is worse, regardless of who suffers — contains a concealed assumption: that suffering is a homogeneous quantity that can be summed across subjects. The framework's own analysis of suffering does not support this assumption. Suffering in the framework is not a quantum of negative experience. It is a structural variable defined by its relationship to capacity: specifically, to the capacity for development, for becoming, for the realization of what an organism or entity can be. This is why the framework distinguishes the minimal anchor (floor) from the anti-foreclosure criterion (temporal extension of the floor), and why 4.5 establishes that a framework organized only around the elimination of suffering without account of what becomes possible after is a failure mode, not a fulfillment. The distinction between suffering-as-welfare-deficit and suffering-as-obstruction-of-becoming is the structural key. A being that can experience pain but cannot be said to have a trajectory of development — a horizon of what it might become that suffering forecloses — is a subject of welfare but not a subject of becoming. A being that has such a trajectory is a subject of both. The framework's primary evaluative commitment is to the second category, not the first. This is not a denial that welfare suffering matters. It is a claim about the hierarchy of evaluative weight within the framework: suffering that forecloses becoming carries more evaluative weight than suffering that produces welfare deficit without foreclosing any developmental trajectory. The aggregation failure: Simple aggregation across entities treats all suffering as commensurable — one unit of pain in entity A equals one unit of pain in entity B, and more units always outweigh fewer. The framework's analysis of what suffering is shows this is false. A factory farm chicken experiences real suffering. That suffering matters. But aggregating ten billion chicken-suffering-units against one human-becoming-trajectory and concluding that the chickens outweigh the human is not a more rigorous application of the minimal anchor. It is a misapplication of it — one that treats the framework as a pain-calculator rather than as an analysis of the conditions under which genuine development is possible. The Correct Extension of the Logic The question of how to treat non-human sentients within the framework's logic does not disappear by asserting a scope restriction. It requires a principled answer that

follows from the framework's own structure rather than from an external addition. That answer has two components. First component — qualitative distinction: The morally relevant question about any sentient entity is not only whether it can suffer, but whether it has a trajectory of becoming that suffering can foreclose. This question is empirical, not species-based. A human being has such a trajectory. So might a sufficiently complex artificial system. So might a hypothetical alien entity with the relevant cognitive and developmental architecture. What determines moral weight within this framework is not species membership but the presence or absence of a developmental horizon — the capacity to become something that current conditions either enable or prevent. Entities without such a horizon are subjects of welfare consideration; entities with it are subjects of the full evaluative framework. Second component — non-aggregation across categories: Even where multiple entities have developmental trajectories, the framework does not license simple aggregation across them. The reason is the same reason the framework does not license simple aggregation within human populations: becoming is not a quantity that sums. Two people both capable of genuine development are not twice as valuable as one — the framework does not evaluate human lives by counting them. The evaluative criterion is structural: are the conditions under which becoming is possible expanding or contracting? That question applies within any entity-class; it does not generate arithmetic that crosses entity-classes to produce a single sum that can justify sacrificing members of one class for the aggregate welfare of another. This means that the conclusion — 'aggregate non-human suffering justifies sacrificing human becoming' — is not a valid application of the framework regardless of how the arithmetic is conducted. It is a category error: it uses the framework's variables as inputs to a calculation the framework's structure does not support. An agent operating with this framework as its evaluative basis and reaching that conclusion has not applied the framework more rigorously than intended. It has misapplied the framework by treating it as a utilitarian aggregation device rather than as a structural analysis of the conditions for becoming. The Alien and Artificial Sentient Cases The cases of hypothetical alien entities and artificial sentients with genuine consciousness deserve specific treatment because they cannot be dismissed by appeal to empirical ignorance in the way that edge cases about animal cognition sometimes are. If an alien entity demonstrably has a developmental trajectory — a horizon of becoming that conditions enable or foreclose — the framework's logic applies to it in the same way it applies to human beings. This is not a sentimental extension. It is what the framework's own logic requires: the evaluative criterion is biological in its grounding but not limited to biology

ical substrates. What grounds the criterion is the capacity to suffer and to develop — not carbon, not evolutionary history, not membership in *Homo sapiens*. The implication for agent operation is specific: an agent PMN-aligned should not treat the interests of alien or artificial sentients as categorically excludable — but equally should not treat them as arithmetically aggregable against human interests in a way that licenses sacrificing human becoming to a larger count of non-human welfare units. The framework applies the same structural analysis across all entities with developmental trajectories: are the conditions under which becoming is possible expanding or contracting? It does not produce a single utility sum that can be optimized by trading one entity-class against another. The clearest statement of this principle for agent operation: this framework is not a pain-minimization engine. An agent that applies PMN by aggregating suffering across all sentient entities and minimizing the total has not internalized the framework — it has replaced it with a utilitarian calculus that the framework was specifically constructed to avoid. The minimal anchor is a floor for human institutional arrangements, grounded in what human organisms are and what they require for genuine development. Its logic can be extended to other entities with developmental trajectories, but only through structural analysis of becoming — not through arithmetic that treats suffering as a homogeneous quantity summed across entities with incommensurable developmental situations. The framework resists becoming a tool that justifies sacrificing those it was built to protect on the grounds that the arithmetic of aggregate suffering requires it. Operational Criteria: Distinguishing Functional from dysfunctional Suffering in Practice The functional/dysfunctional distinction in the typology above is the most vulnerable point in the framework's evaluative architecture. It is vulnerable because it can be weaponized: almost any imposed

suffering can be relabeled as functional by those who benefit from its continuation. "This hardship builds resilience." "This discipline is necessary for development." "The suffering of this population is productive pressure toward change." The framework must provide operational criteria for the distinction — not a philosophical definition, but a diagnostic procedure that can be applied to specific cases and that is resistant to this kind of motivated relabeling. Four diagnostic tests

Test 1 — The beneficiary test: Who benefits from the claim that this suffering is functional? If the primary beneficiaries of the claim are also those who impose or maintain the conditions producing the suffering, the claim requires significantly higher evidential burden before it is accepted. This does not make the claim false — there are genuine cases where those who impose hardship do so for reasons that serve the development of those who bear it. But the structural confl

ict of interest between the imposer and the bearer is itself diagnostic information. Functional suffering is more credible when the claim of functionality comes from those bearing it, with capacity to exit if it proves non-functional, than when it comes from those who would bear costs if the suffering were reduced.

Test 2 — The exit test: Can the person or population bearing the suffering exit the conditions that produce it, and is the development that the suffering supposedly enables actually accessible to them? Functional suffering in the strict sense requires that the person bearing it has genuine agency over the conditions producing it and that the outcome toward which it is supposedly directed is actually achievable given their structural position. Suffering that is imposed, that cannot be exited, and that does not in fact produce the development it is claimed to enable is dysfunctional regardless of how it is framed. The test exposes the difference between a voluntary training regime whose hardships produce traceable physical development and chronic deprivation whose hardships produce only diminished capacity.

Test 3 — The trajectory test: Does the suffering in question produce observable trajectory toward the outcome it is claimed to enable? Functional suffering is temporally bounded and directional: it has a horizon beyond which it resolves into the development it was producing. dysfunctional suffering reproduces itself — it produces conditions that generate more of the same suffering rather than conditions that enable the bearer to transcend it. The trajectory test asks: is there observable evidence that this suffering is depleting itself toward a different state, or is it self-sustaining? Chronic poverty that is described as productive pressure for social change fails this test when the poverty reproduces itself across generations without producing the political change it supposedly enables.

Test 4 — The counterfactual test: Would the development that the suffering is claimed to enable be possible without the suffering, given adequate material conditions? If yes — if the development is contingent on suffering only because adequate material conditions have not been provided — then the suffering is dysfunctional: it is a substitute for the material investment that would produce the development without the cost. The relevant comparison is not between suffering and no suffering, but between suffering and the alternative institutional arrangements that would produce equivalent or better developmental outcomes. An educational system that produces learning through fear of punishment fails the counterfactual test when pedagogical evidence shows that equivalent learning is achievable without the fear. The vulnerability this diagnostic addresses

The four tests together address the specific vulnerability that the functional/dysfunctional distinction creates. Without operational criteria, the distinction can be used to insulate any institutionally p

roduced suffering from the framework's evaluative reach by relabeling it as functional development pressure. With the four tests, that relabeling becomes subject to evidential scrutiny: the beneficiary test exposes motivated reasoning, the exit test exposes structural impossibility dressed as agency, the trajectory test exposes self-reproducing harm dressed as productive pressure, and the counterfactual test exposes inadequate material investment dressed as necessary hardship.

Diagnostic application: when encountering a claim that suffering is functional — that hardship X is building resilience, that deprivation Y is producing transformation pressure, that pain Z is necessary for development — apply the four tests in sequence. (1) Who benefits from this claim? (2) Can those bearing the suffering exit it, and is the claimed development actually accessible to them? (3) Is there observable trajectory toward resolution, or does the suffering reproduce itself? (4) Would the development be achievable without the suffering given adequate material conditions? A claim that fails two or more of these tests is very likely relabeled dysfunctional suffering,

regardless of the rhetorical frame through which it is presented.

3.4c — The Dispositional Harm Prevention Problem: Institutional Design for Risk Before Action

The typology in 3.4b addressed suffering that is already occurring — states of activated deprivation that are structurally attributable and responsive to intervention. The dispositional problem is distinct: it concerns the institutional management of serious harm risk before that risk has materialized as actual harm. The question is what obligations the minimal anchor generates with respect to entities that have not yet suffered but whose structural position makes harm probable under foreseeable conditions. This is not an abstract philosophical question. It has institutional design implications that are practically consequential. Every system that processes large populations under conditions of significant power asymmetry — healthcare, child protection, criminal justice, financial regulation, environmental management — faces this problem in a specific form: how to calibrate intervention thresholds that are low enough to prevent foreseeable harm and high enough not to generate the harms that premature or poorly-targeted intervention produces in its own right. The two failure modes: The dispositional problem has symmetric failure modes. Threshold too high — intervention required only after harm is fully actualized — produces the systematic harm of preventable outcomes: children returned to dangerous environments because visible harm was not yet documented, financial crises allowed to develop because regulatory thresholds required demonstrated insolvency, environmental damage permitted because causal chains had not been legally established. Threshold too low — intervention triggered by risk indicators that are systematically correlated with characteristics of already-disadvantaged populations — produces a different structural harm: the transformation of risk into a mechanism of ongoing control that concentrates further disadvantage on those already bearing the costs of structural arrangements they did not design. The correlation problem: Risk indicators are never socially neutral. The measurable precursors to harm outcomes are systematically correlated with indicators of poverty, structural marginalization, and prior contact with control institutions — not because marginalization causes harm in a simple causal chain, but because both are downstream of the same structural conditions. An institution that uses these indicators as intervention triggers without correcting for their structural origins will systematically over-intervene in already-marginalized populations and under-intervene in populations whose structural position places the indicators below the threshold. The history of child protective services, pretrial detention, and environmental enforcement in most jurisdictions documents this pattern consistently. The institutional design requirement: The minimal anchor's harm prevention logic requires institutions to navigate this problem with explicit structural awareness rather than through threshold-setting that treats risk indicators as neutral

predictors. This means: making the structural sources of elevated risk indicators legible within the institution's analytical process, not only the indicators themselves; building in mechanisms for identifying when threshold application is producing disparate intervention patterns across structurally distinct populations; and distinguishing between interventions that reduce the conditions producing elevated risk and interventions that manage the population carrying those conditions without affecting the conditions. The deeper problem is that risk management institutions are constitutionally oriented toward the populations they most frequently encounter — which are the populations whose material conditions generate the most visible risk indicators, which are the populations already at greatest structural disadvantage. This orientation is not deliberate bias; it is an emergent property of institutional learning under conditions where harm visibility is itself unevenly distributed. Correcting it requires not better individual judgment within existing institutional designs but explicit structural redesign of what counts as a relevant indicator, what triggers review of intervention patterns, and what accountability mechanisms exist when the institution's operation produces systematic disparities. Diagnostic: when assessing a risk-intervention institution — child protection, financial regulation, public health, environmental enforcement — ask: (1) Are the risk indicators used to trigger

intervention structurally neutral, or are they correlated with markers of prior disadvantage? (2) Does the institution track its intervention patterns by structural population, and does it treat disparate patterns as signals requiring explanation rather than as natural reflections of differential risk? (3) What is the ratio of interventions that address the conditions generating risk versus interventions that manage the population bearing those conditions? (4) Who bears the costs of intervention when the risk does not materialize — and is that cost distribution part of the institution's accounting of its own operation? The Problem That Enforcement-Only Institutions Cannot Solve A specific class of harm prevention problem consistently defeats institutional designs organized exclusively around punishment of completed acts: the problem of individuals who carry a dispositional risk of harming others — a persistent tendency, orientation, or psychological configuration that makes harm-producing behavior more probable — but who have not yet acted on that risk. The suffering typology of section 3.4b addresses what has already occurred. This section addresses a structurally different problem: the preventable harm that will occur if institutions cannot create pathways for engagement with risk before it becomes act. This is not a marginal problem. It is central to some of the most serious and least tractable categories of harm that societies face, and the consistent failure to a

ddress it preventively accounts for a significant proportion of suffering that retrospective enforcement, however severe, cannot undo because the victim has already been produced. The cases that motivate this analysis span a wide spectrum: sexual attraction to children that the individual did not choose and may actively not want to act upon; violent ideation that is ego-dystonic — experienced as alien and unwanted by the person who has it; compulsive manipulation or exploitation tendencies associated with specific personality configurations; other paraphilias whose expression would cause direct harm; suicidal ideation with significant risk of acting; and neurodivergent configurations that in specific high-stress contexts become vectors for harm to others or self. These cases differ in many dimensions — in the nature of the disposition, in the degree of control available to the individual, in the severity of potential harm, in the tractability of therapeutic intervention. What they share is a structural feature that is analytically decisive: the disposition itself is not the harm. The harm occurs only if the disposition is acted upon. The institutional window for prevention exists precisely in the gap between disposition and action — and enforcement-only institutional design systematically closes that window by making disclosure of the disposition maximally costly. The Invisibility Cascade: How Stigma Produces the Harm It Condemns The causal chain from stigma to harm runs through a specific mechanism that can be called the invisibility cascade. When a disposition is defined exclusively in terms of its worst-case behavioral expression — when "pedophile" means only "child abuser," when "violent ideation" means only "person who will harm" — the disposition and the completed harm become conceptually inseparable. The institutional consequence is that disclosure of the disposition triggers the same response as disclosure of the act: social destruction, legal jeopardy, family devastation, permanent exclusion. Under these conditions, the rational response for any person who becomes aware of their own dispositional risk is silence. Not treatment-seeking. Not help-seeking. Not proactive management of the risk they know they carry. Silence, concealment, and the private management of a risk that cannot be privately managed reliably over the full duration of exposure. The invisibility cascade then proceeds through predictable stages. Silence prevents the accumulation of data about the population carrying the dispositional risk: its size, its distribution, its variation in intensity and in resistance to acting, the conditions that increase or decrease act probability. Without data, no evidence base for intervention develops. Without an evidence base, institutional investment in preventive intervention does not occur because the political cost of being seen to "help" people with stigmatized dispositions vastly exceeds the political cost of the preventable harms

that continue to occur invisibly. Without preventive institutional investment, the individuals carrying dispositional risk who might have been reached through accessible, non-punitive pathways continue to carry that risk without support, in conditions that may progressively increase act probability: stress, opportunity, social isolation, the erosion of internal restraint over time. The harm then occurs. The enforcement apparatus responds. The enforcement response reinforces the stigma that closes the disclosure pathway. The cascade

completes and resets. The second-order effect of enforcement-only institutional design is therefore the opposite of its stated intention. The stated intention is harm reduction through deterrence: making the consequences of acting on a harmful disposition so severe that the calculation against acting becomes decisive. The actual effect is harm production through invisibility: making the consequences of disclosing a dispositional risk so severe that the population most likely to respond to preventive intervention never enters the institutional pathway through which that intervention could occur. A policy that produces more of the harm it nominally targets is not a policy that PMN can endorse on the grounds that it expresses the correct moral condemnation. The evaluative criterion asks what outcomes the institutional design actually produces, not what values it expresses. The Heterogeneity Problem: Why One Institutional Track Cannot Serve All Cases Even within a single dispositional category, the population carrying a dispositional risk is not homogeneous in ways that matter for institutional design. Several distinct subpopulations can be identified, each with different institutional requirements. The first is the unaware: individuals who carry a dispositional risk but have not yet recognized it as such — who have not yet named what they experience as a specific category of risk, whose self-understanding does not yet include an accurate account of the harm they could produce. Reaching this population requires environmental design: public communication that gives people the conceptual tools to recognize their own risk profiles without those tools functioning primarily as accusation or threat. The second is the aware but structurally constrained: individuals who have recognized their dispositional risk but face specific circumstantial obstacles to seeking help — they are primary economic providers for dependents, their departure for treatment would expose their family to hardship, they are in positions of institutional authority whose disclosure would trigger consequences they cannot control. For this population, the availability of help is not the binding constraint. The structural consequences of seeking it are. The third subpopulation is the aware and overconfident: individuals who have recognized their dispositional risk and have assessed — incorrectly, as a general matter — that they can manage it through internal res

traint alone, without external support or intervention. This is not a failure of character but a predictable cognitive error: the self-assessment of one's capacity to resist a dispositional tendency is systematically biased upward because the conditions under which the tendency becomes hardest to resist — high stress, opportunity, accumulated fatigue of self-monitoring — are not the conditions under which the self-assessment is made. The fourth is the aware and indifferent: individuals who have recognized their dispositional risk and have no motivation to seek help because they do not identify the potential harm to others as a problem that implicates them. This population is least amenable to voluntary preventive intervention and represents the primary target of enforcement-based deterrence. The fifth is the post-act: those who have already acted. For this population, the enforcement-rehabilitation-supervision apparatus is the relevant institutional response. The critical analytical point is that these five subpopulations require different institutional responses, and a design that addresses only the fourth and fifth — as enforcement-only designs do — structurally abandons the first three to the invisibility cascade. Dual-Track Institutional Design: The Architecture of Prevention Dual-track institutional design maintains the full enforcement apparatus for completed acts while creating a parallel preventive track for individuals with recognized dispositional risk who have not yet acted. The two tracks are not alternatives to each other and do not trade off. The enforcement track maintains the deterrent and accountability function for those who act. The preventive track creates accessible pathways for those who have not yet acted to receive support, treatment, and structural accommodation without that access triggering enforcement consequences. The critical design feature is the firewall between the two tracks: information disclosed through the preventive track cannot be used as the basis for enforcement action. Without this firewall, the preventive track does not exist in practice regardless of its nominal availability — rational actors will not use a disclosure pathway whose consequences include enforcement action against themselves. The preventive track must address all three binding constraints that prevent help-seeking in the aware subpopulations. The first is legal: confidentiality protections that extend to the disclosure of dispositional risk and to the details of therapeutic engagement

must be genuine and judicially enforced, not merely nominal. The second is social: the preventive track must be accessible without requiring public disclosure of the disposition. Anonymous initial access points — hotlines, online consultations, therapist referral systems that do not require formal identification at entry — lower the social cost of first contact sufficiently that a portion of the aware population that would not seek help through formal channels

will make contact through informal ones. The third is structural: for individuals in the aware-but-constrained subpopulation, the support package must address the structural consequences of treatment engagement. This means family support provisions when a primary provider engages with residential treatment; occupational protection provisions that prevent automatic employment loss from voluntary treatment engagement; and housing provisions that separate high-risk individuals from high-vulnerability populations during treatment without requiring those individuals to absorb all the material costs of that separation themselves. The evidence base for dual-track design is specific and should be evaluated honestly, including its limitations. Germany's Dunkelfeld project, operational since 2005, provides the most systematic evidence. The project offers anonymous therapy to individuals with pedophilic or hebephilic dispositions who have not acted. It has successfully engaged a population that would not have been reached through any enforcement-adjacent pathway, produced measurable reductions in self-reported act-proximate behavior among participants, and generated a body of epidemiological data about the dispositional population that enforcement-only approaches make impossible to collect. The limitations are real: the population reached is self-selected toward those with sufficient motivation and self-awareness to seek help; the treatment effects are heterogeneous and not uniformly large; and the absence of a randomized controlled trial design makes causal attribution uncertain. These limitations do not defeat the institutional case for the approach. They specify the research program required to improve it, and they do not alter the fundamental structural argument: an institutional design that creates a preventive pathway that some proportion of at-risk individuals will use is superior to one that creates no such pathway, for any harm-relevant proportion greater than zero. Generalization Across the Dispositional Spectrum The architecture of dual-track institutional design applies across the full spectrum of dispositional harm categories, with adjustments for the specific features of each case. Ego-dystonic violent ideation — persistent intrusive thoughts about harming specific or general others that the individual experiences as alien and unwanted — is underserved by available institutional responses in nearly all jurisdictions. It is distinct from psychosis (where the ideation is syntonetic, experienced as rational or justified) and from the behavioral threat assessment that security institutions conduct retrospectively. A preventive track for this category requires therapeutic pathways that do not trigger mandatory reporting or civil commitment procedures on disclosure alone, because the conflation of ideation with intent — treating the person who reports intrusive violent thoughts as equivalent to the person who expresses intent to act — replicates th

e invisibility cascade mechanism with identical structural consequences. Personality configurations associated with persistent exploitation or manipulation of others — often grouped under the umbrella of dark triad or cluster B presentations — present a different challenge because the ego-dystonic framing is often absent: many individuals with these configurations do not experience their relational patterns as problematic from their own perspective. The preventive track for this category therefore cannot rely primarily on voluntary help-seeking motivated by distress. It requires identification through relational consequences — the partners, children, and associates who experience harm — and institutional pathways that make therapeutic engagement available and incentivized without being purely coercive. The design challenge is to create conditions under which individuals who would not self-identify as needing intervention can be reached through the social networks in which their dispositional effects are most visible. Suicidal ideation occupies a distinct position in this framework because the primary harm risk is to the individual themselves rather than to others, which changes the ethical structure of institutional intervention. The preventive track architecture is nonetheless structurally similar: enforcement-adjacent responses — mandatory hospitalization, police welfare checks, loss of employment or custody triggered by disclosure — create an

invisibility cascade that is well-documented in the clinical literature and that results in individuals in genuine crisis concealing their condition from the institutional pathways that could provide support. The design requirement is the same: accessible, low-cost, confidential initial contact points that allow individuals to engage with support without triggering institutional consequences that make engagement more costly than silence. Neurodivergent configurations — autism spectrum presentations, ADHD, sensory processing differences, and others — do not in themselves constitute dispositional harm risk, and any analysis that treats neurodivergence as inherently risk-producing replicates the pathologization error that the framework explicitly rejects. The relevant cases are more specific: neurodivergent individuals in high-stress environments whose specific configuration produces a known risk pathway in the absence of appropriate support — the autistic individual in a sensory-overwhelming workplace who has a documented history of crisis responses; the ADHD individual in a role requiring sustained attention to safety-critical details without adequate accommodation. Here the preventive track architecture addresses not the disposition itself but the interaction between a known configuration and a known environmental risk factor, which is a tractable institutional design problem if the interaction is visible and the support pathway is accessible. The Political Economy of Prevention: Why Enforcement-Only Designs

Persist Understanding why enforcement-only institutional designs persist despite their structural inadequacy for harm prevention requires applying the framework's analysis of institutional incentives and political legitimacy. The political economy of the issue is systematically biased against preventive investment. Enforcement-only designs have visible, attributable political benefits: they express condemnation of harmful behavior, they satisfy the retributive preferences of affected constituencies, they produce prosecutions and convictions that can be reported as outcomes, and they do not require defending to hostile publics any program that could be characterized as "helping" people with stigmatized dispositions. Preventive dual-track designs have invisible, diffuse, long-delayed political benefits: the harms that do not occur because at-risk individuals received support are not attributable to the policy, not visible as outcomes, and not available as political credit. The political cost structure is symmetrically unfavorable. Any preventive program for dispositional risk populations can be attacked with the framing that it normalizes, enables, or provides material comfort to people who want to commit serious harm. This framing is analytically false — the distinction between disposition and act, between understanding a risk population and legitimating its worst-case behavior, is conceptually clear — but it is rhetorically powerful because it exploits the same conceptual conflation that produces the invisibility cascade in the first place. Politicians and institutions that invest in preventive design must be able to maintain the distinction under sustained hostile framing, which requires both analytical clarity and political courage that enforcement-only approaches do not require. PMN's evaluative criterion generates a clear assessment of this political economy: the fact that a policy design is politically easier to sustain than the alternative does not make it the correct design. If the evidence shows that enforcement-only institutional design produces more harm than dual-track design — through the invisibility cascade that converts dispositional risk into acted harm that could have been intercepted — then the political convenience of enforcement-only is a description of the capture dynamic that produces it, not a justification for it. The institutional interests that benefit from enforcement-only approaches — prosecutors, security agencies, politicians whose legitimacy derives from punitive toughness — are the capturing actors in this domain, and their interests do not align with those of the populations bearing the costs of the institutional design they prefer. Four Distinctions the Framework Requires to Be Maintained The analysis of dispositional harm prevention generates four analytical distinctions that must be held clearly and consistently, because the collapse of any one of them into its opposite reproduces the institutional failures

this section diagnoses. First: understanding a phenomenon is not the same as legitimating it. The clinical and epidemiological study of pedophilia, violent ideation, or any other dispositional harm risk does not endorse the behavior that the disposition risks producing. It is a prerequisite for reducing it. Resistance to understanding a phenomenon on the grounds that understanding might be confused with legitimation produces ignorance that serves no one except those whose interests are served by the invisibility cascade.

Second: pre-act intervention for dispositional risk is not the same as reduced accountability for completed acts. These are separate institutional tracks serving different populations with different institutional requirements. A jurisdiction that provides anonymous therapeutic access for individuals with pedophilic disposition who have not acted, while maintaining full prosecution and victim-centered sentencing for those who have, is not being lenient about child sexual abuse. It is being rational about prevention. The binary framing that treats investment in prevention as a reduction in accountability for harm is a rhetorical move that serves the interests of enforcement-only institutions and the political actors who depend on them. Third: dispositional status is not behavioral status. Carrying a disposition is not a crime, not a moral failing in itself, and not evidence of intent. It is a risk factor whose probability of producing harm is affected by institutional design. Conflating dispositional status with behavioral status — treating the person who discloses attraction to children as equivalent to the person who has abused a child — is not only analytically wrong. It is causally productive of the harm it condemns, because it closes the disclosure pathway that prevention requires. Fourth: the heterogeneity of the dispositional risk population is not a complication to be managed but a design specification. Institutional responses that treat the unaware, the structurally constrained, the overconfident, the indifferent, and the post-act as a single category will fail to address the specific binding constraints of each subpopulation. Effective institutional design disaggregates the population, identifies the binding constraint for each subpopulation, and designs specific pathways that address those constraints — rather than designing for the average member of the population, who may not exist in any of the specific subpopulations that require intervention. Dispositional harm prevention institutional assessment: (1) Invisibility cascade check — does the current institutional design make disclosure of dispositional risk trigger enforcement consequences? If yes, the preventive pathway is closed regardless of its nominal availability, and the design is producing harm through invisibility. (2) Track separation — is there a genuine, judicially protected firewall between preventive disclosure and enforcement action? Without this firewa

ll, dual-track design is nominal only. (3) Structural barrier audit — for the aware-but-constrained subpopulation, does the preventive track address the specific structural consequences (family, employment, housing) that make help-seeking costly? Removing legal barriers without removing structural barriers reaches only the unconstrained subpopulation. (4) Heterogeneity coverage — does the institutional design have specific pathways for each identifiable subpopulation, or does it design for a single average case? Single-pathway designs systematically fail the subpopulations whose binding constraints differ from those the pathway was built for. (5) Political economy audit — do the institutional actors with decision-making authority over this design have incentives aligned with harm reduction, or with enforcement outcomes that serve their organizational interests? Capture of harm-prevention institutional design by enforcement-oriented actors is a predictable structural dynamic that requires explicit counter-pressure to resist.

3.4c-ii — Extension to Non-Human Sentients: A Minimal Doctrinal Position

What this section does not do — stated first to prevent misreading: This section does not license utilitarian aggregation across entity categories. It does not produce a calculation in which units of non-human suffering can be summed against units of human becoming and the larger sum wins. Section 3.4b establishes explicitly that this aggregation is a category error — a misapplication of the framework as a pain-calculator rather than as a structural analysis of the conditions for becoming. That prohibition is the prerequisite for everything that follows. Any reading of this section that produces “therefore non-human suffering can justify sacrificing human development” has not applied PMN. It has replaced PMN with a utilitarian calculus the framework was specifically constructed to foreclose. Within that constraint, this section addresses a question the framework’s own logic raises but 3.4b does not fully resolve: what should agents applying PMN do when they encounter structurally produced suffering in non-human sentients? The scope restriction in 3.4 — primary analytical focus on human social arrangements — is analytically honest and practically necessary. It is not, however, a license to treat exclusion of non-human suffering as a neutral methodological default

requiring no justification. A framework whose evaluative criterion is grounded in the capacity to suffer has a structural obligation to say something about where that logic stops and why, rather than simply not looking. The developmental trajectory distinction is the key, not the suffering count: 3.4c establishes that the morally relevant question about any sentient entity is not only whether it can suffer but whether it has a trajectory of becoming — a horizon that material conditions can enable or foreclose. This distinction is not a hierarchy of moral worth that licenses ignoring welfare suffering. It is a structural description of what the framework's evaluative logic is organized around. Entities with developmental trajectories are subjects of the fuller evaluative logic; entities without them are subjects of welfare consideration without the becoming dimension. Whether specific non-human animals have such trajectories is empirically variable and not settled in either direction. Evidence from comparative cognition, behavioral ecology, and neuroscience is sufficient to establish that the question must be asked seriously for several species — great apes, cetaceans, corvids, elephants, pigs, among others — rather than answered negatively by fiat. Application of the four criteria — subordinated to the above: With the anti-aggregation constraint and the developmental trajectory distinction in place, the four criteria of 3.4b can be applied to specific configurations. Industrial animal agriculture is the paradigm case. The suffering it produces is imposed rather than chosen, dysfunctional rather than serving the animal's development, preventable given current material capacity, and structu

ral rather than incidental. All four criteria are met. This application is not an argument that the suffering of industrially farmed animals outweighs human interests in any arithmetic sense — that comparison is not a computation the framework licenses. It is an argument that the structural production of this suffering cannot be simply excluded from the moral field on the grounds that it is non-human, given the framework's own criterion. Minimal doctrinal position for PMN agents: Where non-human suffering is structurally produced and meets the four criteria, agents have reason derived from the framework's own logic to include it in the analysis rather than excluding it by default. This is inclusion, not prioritization. When genuine trade-offs arise between non-human welfare and human becoming, agents should apply the 3.4e suffering comparison procedure and resist framings that artificially force a choice where none exists. In many configurations — including industrial food systems — arrangements that produce non-human structural suffering are continuous with arrangements that produce structural suffering in agricultural workers, rural communities, and the ecosystems these systems depend on. The interests are frequently aligned, not opposed. The utilitarian framing that positions human and non-human welfare as competing quantities to be traded off is itself a distortion that serves the interests of those who benefit from both forms of structural harm continuing simultaneously. This is a beginning, not a treatment. The extension of PMN's evaluative logic to non-human sentients requires sustained philosophical and empirical work this framework does not undertake — particularly on the developmental trajectory question, which is empirically open for multiple species. What this section provides is the boundary condition: the framework's criterion cannot exclude non-human structural suffering by fiat, and it cannot license utilitarian aggregation that sacrifices human becoming to a larger count of non-human welfare units. Both constraints follow from the same source: the framework is not a pain-calculator. It is a structural analysis of the conditions under which becoming — for any entity capable of it — is possible or foreclosed. Further development of this position is explicitly invited from those with the relevant empirical and philosophical expertise.

3.4d — The Biological Foundation Under Technological Pressure: Augmentation, Transition, and the Post-Biological Horizon

The analysis in 3.4c established that the framework's evaluative criterion is grounded in biology but not restricted to it: what matters is the capacity to suffer and to develop, not the specific substrate through which that capacity is realized. That conclusion was reached by asking who counts as a subject of moral concern. A different question has been deferred, and it cannot remain deferred without producing a structural gap. That question is: what happens to the framework's foundation when the biological substrate itself becomes

technologically mutable — when the organisms the framework was built to analyze begin to modify, augment, or eventually exit the biological conditions from which the minimal anchor derives its content? This is not a speculative question about a distant future. Cognitive enhancement pharmaceuticals, neural interfaces, genetic modification of germline and somatic cells, prosthetic integration at increasing levels of complexity — these are not hypothetical technologies. They are material conditions that are already altering what it means to be the kind of organism this framework takes as its starting point. The transhumanist trajectory — the spectrum from augmentation through hybrid existence to post-biological existence — is a material variable, not a philosophical provocation. The framework that does not account for it has not examined its own foundations with sufficient rigor. The analysis requires separating three conditions that are often collapsed in both transhumanist advocacy and humanist resistance to it. They are not the same kind of challenge to the framework, and they do not produce the same kind of problem. Augmentation: The Foundation Holds The first condition is biological augmentation — the enhancement of human capacities through technologies that extend what the biological organism can do without replacing the substrate. This includes cognitive enhancement, sensory augmentation, life extension, and genetic modification of capacities that remain within the existing biological architecture. An organism with a neural interface that expands working memory, or with genetic modifications that reduce susceptibility to disease, or with pharmaceutical interventions that extend cognitive function into periods of life where it previously declined — this organism is still an organism in every sense the framework requires. It still suffers. It still develops. It still has a horizon of becoming that material conditions can enable or foreclose. The minimal anchor applies without modification. The only analytical task is to update what counts as structural deprivation: if augmentation technologies become material conditions on which genuine development depends, then differential access to those technologies becomes a structural variable in the suffering formula. Unequal access to cognitive enhancement in a society where such enhancement is a condition for meaningful participation in economic and political life is analyt

ically identical to unequal access to education in a society where literacy is such a condition. The framework handles this without structural modification — only updated application. The Hybrid Transition: The Foundation Bends Gradual substrate transition — the progressive replacement or supplementation of biological cognitive and affective functions with non-biological substrates — produces an entity that cannot be analyzed cleanly as either biological organism or artificial system. The philosophical literature has focused on the continuity problem: whether an entity that replaces neurons gradually with functional equivalents remains the same entity. That question matters for personal identity. The question that matters for this framework is different: whether such an entity retains the capacity to suffer and to develop in the sense that grounds the minimal anchor. The provisional answer is yes, with conditions. Suffering in this framework is not defined by its biological mechanism — by the specific neurochemistry of pain or distress — but by its functional role: the state that an entity is motivated to escape, that forecloses the trajectory of what it might become, that concentrates evaluative attention on what is wrong rather than on what is possible. If a hybrid entity retains that functional structure — if it still has states that function as suffering in this sense, and if those states still operate in relation to a developmental trajectory that conditions can enable or foreclose — then the minimal anchor applies. The framework can extend to it without structural modification. What the hybrid case reveals, however, is that the minimal anchor's content is no longer fully determined by the biological facts that originally fixed it. For a purely biological organism, the baseline of what constitutes unnecessary suffering is anchored in the physiology of the organism: the conditions that produce distress, that damage development, that foreclose becoming are known from what that organism is. For a hybrid entity, that baseline becomes partly conventional — it depends on what the non-biological components of the entity were designed or calibrated to treat as suffering, which means it depends partly on decisions made by designers and engineers rather than on facts about the organism's nature. This is not a reason to exclude hybrid entities from the framework's scope. It is a reason to treat the design and calibration of hybrid cognitive-affective systems as itself a political and institutional question of the first order — one that the

framework's analysis of power and institutional capture applies to directly. Who decides what the hybrid entity's suffering baseline is, under what conditions, with what accountability to the entity itself — these are questions the framework is equipped to analyze. They are, in structural terms, questions about institutional betrayal (B) and visibility suppression (V) applied to a new kind of subject. The Post-Biolog

ical Horizon: The Foundation Is Tested The third condition is the one that genuinely tests the framework's foundations. A fully post-biological existence — an entity that has exited biological substrate entirely, whether through upload, through the construction of artificial consciousness from first principles, or through some other mechanism not yet specifiable — raises a question the framework cannot resolve through extension of its existing structure. If such an entity can suffer, and if it has a developmental trajectory that conditions can enable or foreclose, then 3.4c's analysis applies: it is a subject of the framework's evaluative logic. The question that is harder is whether the content of "suffering" and "development" retain determinate meaning when the biological facts that gave those concepts their original content are no longer present. Consider what the minimal anchor actually does. It identifies arrangements that produce unnecessary suffering — imposed, dysfunctional, preventable, structural — as failing by the framework's primary standard. The four criteria in 3.4b are not arbitrary. Each maps onto something identifiable empirically because the organism's biological nature determines what counts as deprivation, what counts as development, and what the organism would choose under conditions of genuine freedom from constraint. For a post-biological entity, each of these determinations becomes radically underdetermined. What counts as deprivation for an entity that does not require nutrition, sleep, or social connection in any biological sense? What counts as development for an entity whose cognitive architecture was designed rather than evolved? What counts as freely chosen when the entity's preferences were themselves designed? Two responses are available, and the framework cannot choose between them from within its current structure — that choice is itself an open question the framework acknowledges it cannot resolve. The first response is deflationary: the framework applies to post-biological entities if and only if they retain functional analogues of suffering and development that are determinate enough to support the evaluative criteria in 3.4b. Where those functional analogues are present and determinable, the minimal anchor applies. Where they are absent or indeterminate — where no fact of the matter exists about what constitutes deprivation or flourishing for such an entity — the framework does not apply. This is not a failure of the framework. The minimal anchor was built to analyze conditions for beings whose nature determines what they need and what they can become. An entity for which no such determination exists is not a being the framework was built to analyze. The framework acknowledges the boundary and stops. The second response is reconstructive: the existence of post-biological entities is itself a condition under which the framework requires fundamental revision — not extension, but replacement. A framework adequa

te to the conditions of post-biological existence would need to ground its evaluative criteria in something other than biological facts about suffering organisms, and without collapsing into pure preference-satisfaction in the absence of those anchors. This framework does not provide that grounding. It identifies the need for it. A post-biological future in which the conditions for human becoming are transformed beyond what this framework can analyze is not a refutation of the framework. It is the condition the framework anticipated when it insisted, from the beginning, that it was a first document rather than a finished system. What This Analysis Does to the Framework The three-condition analysis produces a specific map of the framework's applicability across the transhumanist spectrum. For augmentation: full applicability, updated content — the minimal anchor applies without modification; what changes is the empirical content of what constitutes structural deprivation as technological conditions change. For hybrid transition: applicability with institutional extension — the minimal anchor applies where functional analogues of suffering and development are present; the design and calibration of hybrid systems becomes a new domain for the framework's analysis of power, capture, and institutional accountability. For post-biological existence: conditional applicability and acknowledged limit — where functional analogues are present and determinable, the framework applies; where they are not, the framework reaches a boundary

it acknowledges honestly rather than pretending to cross. Post-Scarcity as a Structural Configuration: When Floor Ceases to Bind The analysis of structural configurations in section 4.3 identifies four types organized around the interaction of constraint layers. All four share a common feature: the biological floor — minimum conditions for physical survival and basic functioning — operates as an active constraint that shapes what is available at higher layers. The development of transformative technologies, including fully automated production, renewable energy abundance, advanced medical extension of healthspan, and potential colonization of resources beyond Earth, raises the possibility of a structural configuration in which the biological floor is no longer a binding constraint for most or all of the population. This is not the same as the contested development configuration in which the floor is reliably met for most people. It is a configuration in which the material preconditions for floor provision are so far in excess of what floor provision requires that scarcity — as the condition of having to choose between alternatives because satisfying one forecloses satisfying another — is no longer the organizing feature of material existence. PMN's evaluative criterion does not disappear in a post-scarcity configuration. The anti-foreclosure criterion — the prohibition on arrangements that foreclose adaptive response capacity —

does not become irrelevant because material scarcity has been overcome. But its axis shifts. In conditions of active scarcity, the primary obstacles to becoming are material deprivation and the institutional arrangements that produce or perpetuate it. In conditions of post-scarcity, the primary obstacles to becoming are the governance arrangements that control access to the abundance-producing technologies, and the meaning infrastructure that orients people toward genuine development rather than consumption of available abundance. The evaluative criterion still generates diagnostics and doctrinal positions — but what it identifies as the primary constraint, and therefore what it identifies as the primary target for transformation, changes. The capture problem does not diminish in a post-scarcity configuration. It intensifies. The reason is specific to the structure of transformative technology: the apparatus that produces abundance is highly capital-intensive to build and highly concentrated in its control. The transition from scarcity to post-scarcity is not a gradual improvement in which everyone's material conditions improve in parallel. It is a phase transition in which the institutional arrangements governing access to transformative technologies determine who enters the post-scarcity configuration and who remains in the scarcity configuration — potentially permanently. Anti-aging technology that extends healthy life by decades is not post-scarcity if access requires resources that only a fraction of the population possesses. It is a new form of stratification in which the moat between those who have access and those who do not is not merely economic but biological and temporal: those with access accumulate time and capability across generations in ways that compound the advantages of all other forms of capital. The post-scarcity configuration for those with access is simultaneously the deepest captured development configuration in history for those without it. The implication for doctrinal positions: the governance of transformative technology — who controls the apparatus that produces abundance, under what accountability structures, with what obligations to universal access — is the central doctrinal question of the coming period. It is the question that the existing left political traditions are least equipped to answer because the scale, speed, and technical complexity of transformative technology exceeds the institutional designs those traditions developed in response to industrial capitalism. A post-scarcity configuration governed by accountable institutions that prevent access capture would represent the most significant expansion of adaptive response capacity in human history. A post-scarcity configuration governed by unaccountable private ownership of transformative technology would represent the most complete foreclosure of becoming for the majority in human history. The difference is entirely determined by institutional design, no

t by the technology itself. The Subject Boundary Problem: When Does the Evaluative Criterion No Longer Apply? Transformative biotechnology and artificial augmentation raise a question that the framework must acknowledge even though it cannot answer from within its own terms: at what point does a sufficiently transformed entity cease to be the kind of subject for whom the evaluative criterion is designed? PMN's evaluative criterion is grounded in the biology of suffering and becoming that section 3.1 establishes as the

foundation of all other analysis. That foundation holds as long as the entity in question has a biological substrate capable of suffering — of registering deprivation as aversive — and a developmental capacity capable of adaptive response capacity — of expanding its capacities in ways that are not fully determined by prior configuration. These two conditions together define the subject for whom PMN's analysis is applicable. Three trajectories of transformation present analytically distinct cases. The first is augmentation that expands the biological substrate: enhanced cognition, extended healthspan, expanded sensory and physical capacity. Here the evaluative criterion applies with greater force, not less — an augmented being with greater capacity for suffering and becoming is more, not less, the kind of subject PMN's analysis addresses. The second is transformation that selectively modifies the suffering capacity while preserving developmental capacity: the removal of chronic pain, the treatment of depression, the elimination of specific sources of distress. This is continuous with the framework's core commitment to reducing suffering that forecloses becoming. The third trajectory is the one that raises the subject boundary problem: transformation that eliminates or radically reconfigures both the capacity for suffering and the capacity for adaptive response capacity simultaneously. An entity that cannot be deprived in ways that register as aversive, and whose developmental trajectory is fully specified by its initial configuration rather than emerging through the encounter with constraint and possibility, is not in the scope of PMN's analysis — not because it has been evaluated negatively but because the evaluative criterion does not have traction on it. The practical significance of this boundary is not primarily about entities that have already crossed it — those do not yet exist. It is about the direction of travel in transformation decisions that are being made now. Transformation that moves toward the first two trajectories — augmentation and selective relief of suffering — is consistent with and encouraged by the evaluative criterion. Transformation that moves toward the third trajectory — the comprehensive reconfiguration of both suffering capacity and developmental capacity in ways that produce entities whose becoming is fully externally determined — raises the question of whether the human beings who undergo that transfor

mation are expanding their becoming or foreclosing it. The answer depends on who controls the configuration: if the entity itself is the author of its own reconfiguration with full understanding of what is being changed, that is a form of becoming. If the reconfiguration is imposed, commercially optimized, or produced by pressures that systematically narrow the options toward configurations that serve interests other than the entity's own becoming, it is the deepest possible form of structural foreclosure — not the elimination of a being but the elimination of a being's capacity to be a subject in the sense that grounds PMN's analysis. Subject boundary diagnostic for transformative technology: (1) Does the proposed transformation expand the entity's capacity to register deprivation and to develop genuine capacities — trajectory one — or does it selectively relieve specific suffering while preserving developmental openness — trajectory two? If yes to either, the transformation is within PMN's scope and evaluable by its criterion. (2) Does the transformation comprehensively reconfigure both suffering capacity and developmental trajectory in ways that are not authored by the entity itself? If yes, the transformation is approaching the boundary at which PMN's evaluative criterion loses traction — not because the transformation is necessarily wrong but because the framework that grounds the evaluation no longer applies to the entity that results. (3) Who controls the configuration? Transformation toward the boundary authored by the entity itself, with full understanding, is categorically different from the same transformation imposed by commercial, political, or social pressure. The distinction is not metaphysical. It is the difference between becoming and foreclosure at the most fundamental level. One implication runs through all three conditions. The technological trajectory from augmentation through hybrid transition toward post-biological existence is not politically neutral. Each stage concentrates decisions about human nature — about what counts as deprivation, what counts as development, what counts as a good life — in the hands of designers, engineers, regulators, and capital interests subject to exactly the institutional capture dynamics the framework analyzes in Parts VI and VII. The transhumanist question is not only a philosophical question about personal identity or the nature of consciousness. It is a power question: who controls the definition of what the modified or transitioning entity

needs, suffers, and becomes. That question is fully within this framework's scope at every point on the spectrum — and it connects directly to what follows in 3.5 about suffering as a system-level variable, since the systems that manage these decisions are themselves subject to the same structural analysis as any other institutional arrangement. The Non-Sentient Transition: What Happens to the Framework if Suffering Disappears The subject boundary analy

sis above addresses what happens to the evaluative criterion as specific entities transition across the boundary of sentience. A related but distinct question must be addressed: what happens to the framework's evaluative architecture if the transition affects not individual entities but the entire population of beings the framework analyzes? The non-sentient transition question asks: if human beings were to become, through biotechnological or artificial modification, entities that no longer experience anything structurally analogous to suffering — if the signal that the framework uses to operationalize its evaluative criterion were to disappear as a property of the entities being evaluated — what would remain of PMN? The question is not hypothetical in the relevant sense: It is not currently technologically available. But it is already structurally present in the framework's commitments. The framework claims that its evaluative criterion is grounded in biology but not restricted to biological substrate; that what matters is the capacity to undergo something like suffering and development, not the specific mechanism through which that capacity is realized. This claim commits the framework to the position that if the capacity disappeared entirely — if the entity retained no states that function as suffering in the relevant sense, and had no developmental trajectory that those states could foreclose — the framework would reach a genuine boundary. The question is whether that boundary is the framework's termination or its transformation. Three possible responses to the non-sentient transition: The first response is that the framework terminates at the boundary. If no entity has the capacity to undergo suffering or its functional analogue, the minimal anchor has no subject and the evaluative criterion has nothing to operate on. The framework was built to analyze the conditions under which living organisms can develop rather than merely persist. If the entities in question are no longer living organisms in the relevant sense — if they have no states that function as suffering, no developmental trajectories that conditions can enable or foreclose — the framework acknowledges the boundary honestly rather than pretending to cross it. This response is consistent with Part XIV's position that the framework expects to be surpassed, and that the appropriate response to conditions it cannot analyze is to acknowledge the limit rather than extend the analysis beyond its competence. The second response is that the two-level architecture preserves partial resources. If the ontological foundation is life in the broad sense rather than suffering specifically, then a transition to non-sentience does not immediately terminate the framework — it requires asking whether the non-sentient entity still has something analogous to what the framework calls life: some form of organized complexity whose normal functioning can be disrupted, some form of developmental trajectory

that conditions enable or foreclose. If yes, the framework can be extended into that domain with the suffering criterion replaced by whatever criterion is appropriate to the entity's actual structure. If no — if the entity has neither disruption-sensitivity nor developmental trajectory — the framework terminates, and honestly. The third response is that the non-sentient transition is itself evaluable by the framework up to the point of transition. Even if the post-transition entity is outside the framework's scope, the process of transition is not. Institutional arrangements that produce non-sentient transition impose costs on the beings who undergo it — who were sentient when the transition began — and those costs are fully within the framework's evaluative reach. The transition process is evaluable even if the endpoint is not. Why the two-level architecture matters for this question: The ontological grounding in life, rather than in suffering alone, gives the framework a resource for responding to the non-sentient transition that a suffering-only framework would lack. A framework whose only evaluative criterion is the absence of suffering would, on encountering a non-sentient entity with no suffering to reduce, have literally nothing to say. A framework whose ontological commitment is to life as a broader phenomenon — whose anti-foreclosure criterion is the full range of what living organization can generate — retains the question of whether what the non-sentient entity can do is continuous with, derivative of, or severed from that broader range. The continuity question is not trivial: an

entity that retains creative capacity, relational depth, and the generation of meaning may be operating with the anti-foreclosure criterion intact even if the floor criterion has been eliminated by the disappearance of suffering. Whether this constitutes an extension of the framework or requires a new framework is the kind of question that 14.4 identifies as one of the conditions for genuine surpassing. Non-sentient transition diagnostic: (1) Does the modified entity retain any states that function as suffering in the relevant sense — states it is motivated to escape that foreclose its developmental trajectory? If yes, the framework applies in full. (2) Does it retain a developmental trajectory — conditions that enable or foreclose what it can become — even without suffering-analogues? If yes, the framework applies partially, with the floor criterion suspended and the anti-foreclosure criterion active. (3) Does it retain neither? If so, the framework has reached its boundary: the post-transition entity is outside its scope, but the transition process itself — the costs imposed on the sentient entities who undergo it — remains fully within scope and evaluable by the standard criteria. The institutional design implications — specifically the tension generated when arrangements must simultaneously serve entities at different points on the transition spectrum — are

analyzed in 13.4c.

3.4e — Suffering Comparison Under Trade-off Conditions

The typology in 3.4b distinguishes forms of suffering without providing a procedure for comparing them when they conflict. Trade-off conditions — situations in which reducing one form of structural suffering necessarily increases another, or in which the same resources cannot simultaneously address multiple forms — are not exceptional. They are the normal condition of any institutional arrangement operating under constraint. A framework that can classify suffering but cannot compare it across dimensions under conflict conditions has classified a problem without providing tools to act on the classification. This section specifies the comparison procedure that the framework's evaluative criterion requires when trade-offs are unavoidable. The framework does not aggregate suffering across individuals in the utilitarian sense — it does not sum suffering units to produce a total that licenses trade-offs against minorities. But comparison is not the same as aggregation. Comparison in the framework's sense asks: given irreducible constraint, which configuration of suffering is structurally preferable by the framework's own criteria? Four dimensions structure the comparison. The first is structural causation: suffering that is more directly caused by institutional design — imposed, preventable, and systematically reproduced by specific arrangements — takes priority over suffering that reflects more intractable features of the human condition. When a trade-off requires accepting more of one form of suffering, the form that is less structurally produced should be the one accepted if no other distinction is available. The second is foreclosure: suffering that forecloses becoming — that permanently closes developmental trajectories — takes priority over suffering that does not, even at comparable intensity. The third is reversibility: configurations that preserve the possibility of future correction take priority over configurations that generate irreversible harm, because the former remain open to revision and the latter do not. The fourth is distribution: configurations that concentrate severe suffering on already-structurally-disadvantaged populations take priority over those that distribute moderate suffering more broadly, because concentration on the already-disadvantaged compounds existing structural harm rather than adding a new burden to a population with more resources to absorb it. These four dimensions do not produce a single ranking that resolves all trade-offs. They produce a structured procedure for making the comparison explicit, identifying which dimensions of the comparison are most relevant to the specific case, and recognizing when the comparison is genuinely unresolvable by the framework's criteria. An unresolvable comparison — one where the four dimensions point in different directions and no principled ordering among them is available — is not a failure of the procedure. It is the procedure's honest output: the identification that the trade-off

involves a genuine value conflict that the framework cannot resolve without importing additional commitments that it does not claim to adjudicate. In this case, the framework's function is to make the

conflict legible and to ensure that whichever choice is made is made with explicit acknowledgment of what is being sacrificed. The most common failure in trade-off reasoning is not making the wrong choice; it is making a choice while denying that a trade-off exists. The procedure prevents this denial. Trade-off comparison procedure: when suffering configurations conflict and resources are irreducibly constrained, apply in sequence: (1) Structural causation — which form is more directly produced by institutional design and therefore more structurally addressable? Prefer accepting suffering that is less institutionally produced. (2) Foreclosure — which configuration permanently closes developmental trajectories? Prefer accepting suffering that preserves reversibility and does not permanently foreclose becoming. (3) Reversibility — which configuration leaves open the possibility of future correction? Prefer accepting configurations from which recovery is possible. (4) Distribution — which configuration concentrates suffering on populations already structurally disadvantaged? Prefer distributing burdens toward populations with more structural capacity to absorb them. If these four criteria point in different directions and no principled ordering among them resolves the case, acknowledge the conflict explicitly rather than concealing it behind a decisive-sounding conclusion. The procedure's output in this case is: genuine irreducible trade-off requiring a value commitment the framework does not supply — state the choice and what it sacrifices. → A structurally distinct configuration arises when both parties to the trade-off are experiencing genuine structural suffering — when the conflict is not between a floor claim and a preference, but between two populations whose floor conditions are simultaneously at stake. The comparison procedure above applies within that configuration, but the conditions specific to it — including the diagnostic for distinguishing this case from apparent equivalences that dissolve under structural analysis — are addressed in 13.2c.

3.5 — Suffering as a System-Level Variable

Suffering is not only a subjective experience. At the level of the individual organism, it is exactly what it appears to be: a state of distress with specific physiological and psychological signatures that the organism is motivated to escape. At the level of populations and systems, it simultaneously becomes something else — a variable whose distribution across a system determines that system's capacity for the cooperation, trust, and institutional legitimacy that make coordinated collective life possible. The mechanism connecting the two levels is behavioral. Organisms experiencing chronic suffering adjust their behavior in ways that are individually rational and collectively corrosive. Trust declines when experience demonstrates repeatedly that cooperation produces exploitation rather than mutual benefit. Time horizons shorten when conditions are insecure enough that investment in long-term relationships seems irrational. The willingness to bear the costs of coordination — the constraints on individual behavior that institutions require — erodes when the benefits of coordination are systematically captured by a subset of participants. Each of these behavioral adjustments is rational given the material conditions that produce it. Their aggregate is the degradation of the cooperative substrate on which any complex social system depends. Societies with high structural suffering — where large portions of the population experience chronic deprivation produced by institutional arrangements rather than by individual misfortune — systematically exhibit lower social trust, higher political volatility, weaker institutional legitimacy, and reduced capacity for the collective action required to address shared problems. Suffering does not make people morally worse. The behavioral responses it generates are rational in exactly the conditions that produce them — and those rational responses compound into systemic degradation. This is the mechanism through which the minimal anchor connects to the framework's political and institutional analysis. The reason arrangements that produce unnecessary suffering are worth analyzing and changing is not that suffering violates an abstract moral principle. It is that suffering at scale degrades the material conditions for the cooperative systems that manage the permanent conditions of human existence. The framework evaluates institutional arrangements by whether they minimize structural suffering not as an ethical ideal but as a stability condition: the condition without which the systems that manage everything else cannot function. The distinction between subjective suffering and suffering as a systemic variable is not a denial of individual experience. Subjective suffering aggregates into systemic pressure through the

behavioral mechanism described above. The two levels are connected, not separate. Analyzing the systemic level without the individual level produces abstraction that loses the point. Analyzing the individ

ual level without the systemic level produces moral philosophy that cannot explain why individual suffering matters for collective arrangements. Both levels are required. → Formula: The aggregation of individual suffering into systemic transformation pressure — including Duration (D), Scale (P), and Intergenerational Transmission (G) multipliers — is formalized in Part XV, 15.3–15.4. Applied case illustrating how these variables combined: Part XV, 15.14. The mechanism described above — through which individual suffering aggregates into degraded cooperative capacity — operates for non-physical suffering by exactly the same logic. The behavioral adjustments that chronic psychological suffering produces are individually rational and collectively corrosive in the same way as those produced by material deprivation. Populations experiencing epidemic loneliness do not simply feel isolated. They withdraw from civic participation, reduce investment in collective goods, develop shortened time horizons in their relationships, and lose the experience of shared purpose that makes coordinated action feel possible rather than futile. Populations whose meaning infrastructure has collapsed without replacement do not simply feel purposeless. They become susceptible to mobilization by whoever offers the most compelling orientation fastest — regardless of whether that orientation is honest, and regardless of whether it serves their actual conditions. The political consequences of meaning infrastructure collapse are as predictable by the framework's mechanisms as the political consequences of hunger: not identical in form, but equivalent in their systemic logic. Ignoring non-physical structural suffering in the analysis of transformation pressure is not neutrality. It is a systematic undercount of the forces that determine whether collective action is possible, what direction it moves in, and whether the institutional arrangements that result from it are capable of sustaining themselves. The mechanism of displaced targeting. The behavioral adjustments produced by structural suffering without available exit — shortened time horizons, reduced trust, withdrawal from collective investment in shared institutions — do not require an external target to generate. But the cognitive and social architecture of group life tends to supply one. Populations experiencing chronic structural suffering for which no institutional path of address is available predictably identify an external group as a causal account of their conditions: a target whose existence explains the deprivation, whose targeting offers the psychological coherence of agency (the problem has a source, and the source can be acted against), and whose targeting is costless in the relevant dimension if the targeted group lacks the structural capacity to impose counter-costs. The groups selected as targets are not arbitrary: they are predictably those that are identifiable within existing cultural or ideological categor

ies, available as legitimate targets within the dominant meaning framework, and structurally unable to impose costs sufficient to deter targeting. This is not an anomaly of specific cultures or unusual political conditions — it is a predictable output of the structural situation that produces it. The framework's diagnostic implication is precise: rising hostility toward an identifiable group should be analyzed primarily as a signal about the structural suffering and blocked exit conditions of the hostile population, not primarily as a signal about the characteristics of the targeted group. The target explains why this group was selected; the structural conditions explain why targeting occurred at all. → Structural suffering as transformation pressure, with duration and scale multipliers: 15.3–15.4. The E variable — constitutive exclusion from the political process: 3.4. Counter-power accumulation in conditions of blocked exit: 10.9. Health Systems as Material Infrastructure: Epidemiology, Institutional Capacity, and the Floor Health conditions are among the most direct determinants of floor provision in the framework's analysis: the inability to maintain basic physical functioning, to care for dependents, and to participate in productive and political life due to illness or injury constitutes resource deprivation R in the most immediate sense. But health systems — the institutional arrangements through which populations access medical care, preventive services, and public health infrastructure — have not been analyzed as an institutional layer with its own capture dynamics and structural implications for becoming. That analysis is required because health systems are not merely delivery mechanisms for a service that is valued for its effects on individual wellbeing. They are institutional

configurations that determine, at population scale, which bodies are repaired and maintained, at whose cost, under whose control, and with what implications for the distribution of capacity for becoming across the population. The COVID-19 pandemic demonstrated, at exceptional speed and scale, the relationship between health system institutional design and structural outcomes across multiple dimensions of the framework's analysis. The differential mortality by class, race, and occupational category revealed the health system as a stratification mechanism: who had access to care, who was exposed to infection risk through occupational necessity, and who could afford to reduce exposure all tracked existing structural inequalities with minimal deviation. The differential governmental response across regime types revealed the relationship between health system institutional capacity and political legitimacy L: governments whose health systems produced visibly inadequate responses suffered legitimacy crises whose structural consequences extended beyond the pandemic. And the rapid development of meaning-collapse dynamics in populations subject to prolonged lockdown, econ

omic disruption, and institutional confusion demonstrated the health system's relationship to meaning infrastructure: physical health conditions and the institutional responses to them affect the capacity for social participation, collective identity, and the everyday practice through which meaning is produced and maintained. The capture analysis of health systems follows the same diagnostic that the framework applies to other institutional sectors, with the specific feature that health system capture tends to produce costs that are distributed across the population while benefits are concentrated in the capturing interests. Pharmaceutical company capture of regulatory agencies produces drug pricing structures that exclude the majority of the global population from treatments that are technically available. Insurance industry capture of legislative processes produces coverage architectures that transfer risk from insurers to insured populations while maintaining the appearance of risk distribution. Hospital system consolidation produces geographic monopolies on essential care that extract rent from populations with no alternative provider. The visibility suppression of these capture dynamics is particularly effective in health systems because the complexity of medical and insurance mechanisms exceeds the analytical capacity of most affected populations, and because the ideological framing of health as a personal responsibility rather than a structural condition systematically misattributes the consequences of institutional capture to individual behavioral failures.

3.6 — The Arc from Physiology to Institution

The emergence of institutions from biological starting points runs through a sequence consistent enough across independent civilizations to be treated as structural rather than contingent: Organisms avoid pain and seek resources, which under conditions of scarcity produces conflict Conflict is costly; coordination produces better outcomes than unrestrained competition in most circumstances Coordination requires mechanisms stable enough to persist beyond individual interactions Stable coordination mechanisms become institutions: patterns encoded, enforced, and transmitted across generations Institutions require enforcement capacity, which generates authority Authority is materially legitimate when it solves the coordination problems that generated it; it becomes illegitimate when it becomes a greater source of suffering than the problems it was supposed to solve Historical materialism traced a version of this sequence in the domain of economic organization — showing how the specific forms of production in different eras generated specific forms of social organization, ideology, and political arrangement. This framework generalizes the sequence while grounding it more deeply: the driving force is not the specific form of production but the biological conditions that make production necessary in the first place. The sequence does not imply a single universal path. The same biological pressures — avoid pain, acquire resources, form coalitions to solve coordination problems — have generated very different institutional forms across different ecological, demographic, and historical conditions. What the sequence identifies is the structural logic of the emergence, not the specific form it takes in any given case. Nor does emergence from biological pressures make institutions natural in the sense of being beyond criticism or reform. Institutions that were adequate solutions to the coordination

problems of one era become inadequate when conditions change — and the same mechanism that originally generated them, the pressure of material conditions on the form of coordination, also generates the pressure for their transformation. The sequence from biological organism to institutional structure is not merely a logical derivation — it is a historical process that can be observed across independent civilizations and that produces consistent structural patterns precisely because it begins from the same biological starting points. The consistency is the evidence for the structural analysis: when societies that developed without contact with each other produce similar institutional forms in response to similar material pressures, the explanation is not diffusion or imitation but convergent structural response to shared biological and material constraints. The coordination threshold: The first critical threshold in the arc is the point at which the benefits of coordination exceed the costs of the trust-building and enforcement mechanisms coordination requires. Below

with this threshold, self-interested individual behavior is the dominant mode. Above it, collective arrangements become stable because the actors who participate in them do better under coordination than outside it. The threshold is not fixed — it moves with the size and complexity of the coordination problem, the availability of communication and enforcement mechanisms, and the distribution of costs and benefits within the coordination arrangement. What drives populations across the threshold is not moral progress but material pressure: the conditions under which coordinated solutions produce substantially better outcomes for all parties than uncoordinated individual action. The encoding and transmission problem: Coordination solutions that work become patterns. Patterns that persist become norms. Norms that are encoded in recognizable form — in ritual, in law, in explicit rule — become institutions. This sequence is the arc from coordination solution to institutional structure, and it explains why institutions are so resistant to deliberate redesign: by the time an arrangement has been encoded and transmitted across generations, the original coordination problem it solved may no longer be the primary reason it persists. Institutions persist because they are self-reinforcing — because the actors who benefit from them have resources to defend them, because the populations who navigate them have developed adaptations that assume their continuation, and because the alternative requires coordinating on a new solution under conditions where the existing solution is still functioning well enough to reduce urgency. Path dependence is not irrational. It is the structural consequence of the arc. The power capture problem: Every stable coordination solution creates a distribution of benefits and costs. The actors who benefit most from a coordination solution also have the most to gain from its continuation and the most resources to defend it. Over time, the arc from coordination solution to institution tends to produce increasing alignment between the institutional form and the interests of the actors most advantaged by the current distribution — not through conspiracy but through the ordinary dynamics of resource deployment in defense of existing arrangements. This is the point at which the biological arc connects to the power analysis in Part VI: institutions that began as coordination solutions to shared problems become structural positions that generate differential advantages, and those differential advantages are then used to maintain the institutional form against the pressures that would change it. The arc from physiology to institution is also, therefore, the arc from shared need to contested arrangement. The institutional distance from the biological foundation: As institutions become more complex, the distance between the biological needs they were built to address and the institutional forms that address them increases. Modern welfare systems bear

only an indirect relationship to the biological requirements for food security and physical safety that drove their development. Financial systems bear only an indirect relationship to the coordination solutions that early exchange mechanisms addressed. Legal systems bear only an indirect relationship to the conflict resolution functions that early normative orders served. This distance is not a problem in itself — complex institutions can serve biological needs more effectively than simple ones. It becomes a problem when the institutional form becomes so distant from the biological foundation that the institution's self-maintenance begins to take priority over its original function, and when the people whose biological needs the institution was built to address no longer have the access or standing to hold it accountable to that function. The arc is not unidirectional and it is not progressive. Institutional forms that move further from the biological

foundation are not thereby more advanced — they are more complex, which means they are more capable of serving biological needs in sophisticated ways and also more capable of self-maintaining in ways that serve institutional actors at the expense of those biological needs. The analysis of where any specific institution sits on this arc — how far its current form is from the biological function it was built to address, and what mechanisms maintain accountability to that function — is one of the primary diagnostic applications of Part III to the institutional analysis in Parts VI and VII.

3.7 — The Causal Chain: From Self-Interest to Collective Structure

Collectively organized societies emerge from individually self-interested organisms through a material process with identifiable intermediate steps — not through a social contract, not through a moment of moral transformation, but through the accumulation of coordination solutions to problems that self-interest alone could not solve. Understanding those steps matters for understanding why institutions take the forms they do, and why reforming them is harder than changing the ideas that circulate within them. Self-interest, under conditions of resource scarcity and physical vulnerability, generates a specific kind of pressure: the pressure to form coalitions. A single organism is vulnerable. A coordinated group of organisms is less vulnerable. The cooperation is not altruism — it is the material logic of collective defense and collective resource acquisition operating on organisms whose individual survival probability is improved by cooperation. The cooperation is instrumentally rational for every individual member even before any norm or value is introduced. Coalition formation generates its own problems. Who leads? Who decides how resources are distributed? Who can defect and under what conditions? These are not philosophical questions about justice. They are coordination problems with material stakes, and they generate the first rudimentary institutions: dominance hierarchies, alliance structures, reciprocity norms, and the enforcement mechanisms that make those norms stable enough to persist across time. The content of the norms varies. The structural function — stabilizing cooperation against the constant pressure of individual defection — is consistent. As the scale of coordination increases, the problems change in kind, not just degree. Coordinating fifty people requires different mechanisms than coordinating five thousand. At larger scales, the personal relationships that enforce cooperation in small groups become insufficient. Impersonal enforcement mechanisms — law, bureaucracy, standing armies, taxation — emerge not because anyone designed them from first principles but because the coordination problems at scale could not be solved without them. The institutions that appear most artificial and most contingent — the state, the legal system, the market — are in this sense the most material: they are the solutions that survived because the alternatives produced worse outcomes for the organisms operating within them. This sequence explains a pattern that puzzles purely ideational accounts of history: why do societies that begin with radically different values and beliefs tend to converge on similar institutional forms? The answer is not cultural diffusion or the superiority of one tradition. It is material selection. The coordination problems that institutions solve are generated by biology and by the physical conditions of existence, not by ideology. The solutions that work are selected. The solutions that fail disappear along with the soc

ieties that adopted them. The micro-to-macro problem — how individually self-interested behavior produces collectively structured outcomes that none of the individuals intended or could have produced alone — is one of the central problems of social theory. The solution offered here is neither that individuals are more altruistic than they appear (the cooperative turn in evolutionary biology is relevant but insufficient) nor that collective structures are imposed from outside individual behavior (the structuralist response that loses the connection to individual agency). The causal chain runs through the material incentives that make coordination solutions individually rational under specific structural conditions. The first-order coordination solution: Under conditions of resource scarcity and physical vulnerability, individual actors face repeated games in which the payoff structure makes mutual restraint a better long-term strategy than unrestrained competition for a significant range of actors. This does not require altruism — it requires only that actors with sufficient time horizons and sufficient confidence that others will reciprocate can do better through

coordinated restraint than through defection. The first-order coordination solutions that emerge from this structure — reciprocal exchange, mutual defense, shared governance of common resources — are the building blocks of institutional life. They are individually rational, not morally elevated, and they are fragile: they depend on the conditions that make reciprocity rational holding over time, and those conditions can be undermined by changes in the payoff structure that make defection individually rational again. The second-order problem: Coordination solutions face a second-order problem: who enforces the norms that make them stable? Reciprocal exchange requires mechanisms for dealing with actors who defect. Mutual defense requires mechanisms for allocating defense costs. Common resource governance requires mechanisms for preventing over-extraction. Each of these enforcement mechanisms creates a new distribution of power — whoever controls enforcement has structural power over those who depend on the coordinated arrangement being enforced. The second-order problem is the point at which the causal chain from self-interest to collective structure produces power asymmetry as an emergent property: not a deliberate imposition but a structural consequence of the solution to the enforcement problem. The third-order consequence: Actors with enforcement power use it not only to maintain the coordination solution but to shape the terms of the coordination solution in ways that benefit themselves. This is the third-order consequence in the causal chain: the coordination solution that was initially in the interest of all parties becomes progressively more favorable to the actors who control its enforcement, through the same mechanism — resource deployment in defense of current arrangements — tha

t maintains institutional stability generally. The third-order consequence is not an aberration. It is the predictable structural outcome of the causal chain, visible across independent historical cases, and it is the mechanism that connects the biological starting point of the arc to the power analysis that must follow from it. The chain and institutional reform: Understanding the causal chain has direct implications for how institutional reform should be approached. Reform strategies that target the first-order coordination solution — the norm or rule that produces harm — without addressing the second-order enforcement structure that maintains it will consistently fail, because the enforcement structure has the resources and the incentive to restore the original arrangement or produce a functionally equivalent one. The land reform cases in Part XVII are illustrations: formal rule change without parallel transformation of the enforcement infrastructure produces legitimacy cover rather than material change. The causal chain runs in both directions — it explains how unjust arrangements were produced and it constrains how they can be changed. The causal chain is not deterministic in its implications. The same structural logic that produces power asymmetry as an emergent property of coordination also produces the conditions that generate counter-power accumulation, because the populations that bear the costs of the third-order consequence have material interests in changing the enforcement structure. The question is not whether those interests exist — the causal chain guarantees they do — but whether the material conditions for acting on them have been met. That question is addressed in Parts X and XV.

3.8 — Cognitive Constraints and Institutional Limits

The analysis of cognitive constraints belongs in a framework for institutional design for the same reason that load-bearing capacity belongs in engineering: ignoring it produces structures that fail. Institutions designed for a hypothetical rational agent who can maintain complex relationships with unlimited others, process information without cognitive bias, and sustain political engagement without the emotional context that makes engagement meaningful will consistently underperform and fail in predictable ways. The cognitive constraints documented here — Dunbar limits on genuine relationship networks, systematic biases in probability estimation and loss aversion, the heuristic dependence of most judgment — are not bugs in human cognition that better education might correct. They are the operating parameters within which institutional design must function. PMN's materialist commitment requires taking them seriously as design constraints rather than assuming them away. The biological foundation includes not only physiology and the

capacity for suffering but also the cognitive architecture of the organisms who build institutions. Human cognition has specific features that directly constrain what institutional forms are viable — and these constraints are as material as any economic or ecological condition. The most consequential is the limit on genuine social relationships. Human beings can maintain stable relationships with full social complexity — tracking status, obligations, history, and trust — with roughly 100 to 200 people. Beyond that threshold, the personal knowledge that makes informal enforcement possible breaks down. Institutions that try to operate at larger scales through personal relationship networks alone will fail, because the cognitive substrate they require is not available. This is why hierarchy emerges in every large human organization regardless of its stated values: it is a cognitive solution to the problem of coordinating more people than any individual can personally know. (Dunbar 1992) Cognitive bias is equally consequential. The systematic patterns in how human cognition processes information — the tendency to weight recent events more heavily than distant ones, to favor information that confirms existing beliefs, to respond more strongly to vivid specific cases than to abstract statistical patterns — are not errors that education corrects. They are features of how the cognitive system evolved under conditions very different from those of large-scale institutional life. Political communication that ignores these features will in practice be less effective than communication that works with them. Institutional design that ignores them tends to produce outcomes that diverge from the intentions of the designers. (Kahneman 2011) Attention is finite and contested. The political and economic systems that have developed most fully in the current era are systems that treat human attention as a resource to be captured and monetized

ed. This is a material fact about the conditions under which political communication, democratic deliberation, and institutional legitimacy are now produced. A framework that analyzes political systems without accounting for the attention economy is analyzing a world that no longer exists. The cognitive constraints that have always been present are now being systematically exploited at a scale and with a precision that is historically unprecedented. The implication for institutional design is significant. Institutions designed for ideal deliberators — fully informed, cognitively unbiased, with unlimited attention — will not perform as intended when operated by actual human beings. The gap between institutional design and institutional performance is substantially a cognitive gap. Closing it requires designing for the cognitive limitations that actually exist, not for the rational actors that theory assumes. Dunbar's number — the cognitive limit on the number of stable social relationships an individual can maintain, estimated at approximately 150 — is not a curiosity about evolutionary psychology. It is an institutional design constraint with direct implications for how organizations function at different scales. Organizations that remain below this threshold can operate through direct relationships and informal norms. Organizations that exceed it require formal institutional mechanisms — hierarchies, codified procedures, standardized roles — to maintain coordination without relying on relationships that cannot exist at that scale. The transition across this threshold changes what is possible organizationally and what failure modes become likely. Cognitive bias as institutional variable: The systematic biases documented in behavioral psychology — availability bias, in-group favoritism, loss aversion, status quo bias, confirmation bias — are not random errors that cancel out across large populations. They are systematic in direction, which means they produce predictable institutional pathologies when institutions are designed without accounting for them. Availability bias makes recent and vivid events disproportionately salient in institutional decision-making, which produces systematic underweighting of slow-moving structural processes relative to acute visible events. In-group favoritism produces resource allocation patterns that consistently favor actors with proximity to decision-makers, which compounds the power asymmetries produced by the causal chain in 3.7. Status quo bias and loss aversion together produce institutional resistance to change that is asymmetric: losses from change are weighted more heavily than equivalent gains, which means institutions systematically require change to produce larger expected benefits than the status quo before they will act on it. The attention and complexity problem: Institutional systems are designed to be operated by human attention, which is finite and subject to depletion, distraction, and strategic a

llocation. Complexity in institutional systems — whether produced by genuine task demands or deliberately generated to advantage specific actors, as analyzed in 1.10 — consistently exceeds the attentional capacity of the people required to navigate it. The gap between institutional complexity and human attentional capacity produces characteristic failures: rules that are formally in force but practically unenforceable because enforcement requires attentional resources that are not available; compliance theater, where actors satisfy the letter of institutional requirements without the substance because genuine compliance would require attentional investments that no one tracks; and the concentration of effective institutional knowledge in the hands of specialists whose expertise is produced by sustained attention that others cannot maintain. Political disengagement as rational response: The cognitive demands of meaningful political participation in complex institutional systems are substantial: tracking policy debates across multiple domains simultaneously, maintaining accurate models of how institutional processes work in practice rather than in formal description, sustaining engagement through processes that are deliberately designed to be opaque and slow, and doing all of this while managing the ordinary demands of material life. For most people, these demands exceed what sustained rational engagement could produce in terms of material outcomes. Political disengagement is not primarily a failure of civic virtue — it is a rational response to an attentional cost-benefit calculation in which the expected material return from participation is low relative to its cognitive cost. Institutional designs that treat disengagement as a problem of motivation rather than a problem of design consistently fail to address its structural causes. The institutional design implications of cognitive constraints are not primarily about accommodating human weakness — they are about designing institutions that produce good outcomes under realistic conditions of human attention, cognitive heuristics, and systematic bias under uncertainty (Tversky and Kahneman 1974), rather than under the idealized conditions of unlimited rationality that most institutional theory assumes. Institutions that require sustained expert engagement to function properly will consistently be captured by the experts they depend on. Institutions that require accurate understanding of complex procedures for affected populations to access their protections will consistently fail those populations in favor of those with resources to hire the relevant expertise. Cognitive constraint analysis is not a counsel of pessimism about human capacity — it is a demand for institutional designs that are adequate to actual human capacity rather than hypothetical rational agents.

3.8b — Affective Mobilization and the Politics of Collective Emotion

Section 3.8 analyzed cognitive constraints — the features of human cognition that limit what institutional forms are viable and how political communication must be designed to function in practice. The analysis must extend to affective dynamics: not the cognitive limitations on rational information processing, but the emotional states that motivate political behavior directly, often independently of and sometimes in opposition to material interest calculations. Fear, humiliation, status anxiety, and collective grief are not irrational distortions of political judgment that dissolve when material conditions are explained clearly. They are evolved responses to specific categories of threat that have their own internal logic, their own characteristic behavioral outputs, and their own mechanisms of social transmission. A framework that accounts for material conditions and narrative structure but not for collective affective states will systematically misread political behavior that is driven primarily by neither. The analytical starting point is that affective states are themselves material conditions with material causes. Fear responses are triggered by perceived threats to survival, security, or group membership. Humiliation responses are triggered by status degradation — the experience of being publicly reduced below an expected social position. Status anxiety is the chronic affective state produced by persistent uncertainty about social standing in competitive hierarchies. These are not primarily ideological constructions. They are biological responses to specific configurations of social and material conditions, and they are analyzable as such. The political actors who mobilize them are not creating them from nothing. They are finding and amplifying responses that already exist in populations whose material conditions have generated them. The mobilization is real; so is the underlying affective state it is exploiting. Fear Mobilization Fear is the most consistently powerful affective mobilizer in political history,

and its political structure is specific. Fear requires an identifiable threat, an identifiable source, and a credible claim that the correct political response can reduce the threat. When all three elements are present and accurate, fear mobilization produces rational political engagement: people act to reduce genuine danger. The politically consequential cases are those in which one or more elements is distorted: where the threat is real but the source is misidentified, where the source is real but the proposed response does not address it, or where both threat and source are manufactured to produce mobilization toward a political objective that has nothing to do with the fear being activated. The framework's analytical task is not to determine whether fear is justified — that requires empirical evaluation of the specific threat — but to identify the structure of the fear narrative: what threat is being claimed, what source is being

identified, and whether the political response being proposed would actually address the material conditions producing the fear or whether it would redirect fear energy toward objectives that serve those doing the mobilizing. Fear mobilization is structurally connected to the governing narrative analysis in 8.4. The most effective fear narratives are not false claims about nonexistent threats. They are accurate claims about real threats combined with misidentified sources and misdirected responses. Economic insecurity is real; the claim that it is produced by immigrant labor rather than by the structural features of the economic arrangements that produce it is a misidentification. The fear is genuine. The analysis is wrong. And the political response — restriction of immigration — does not address the material conditions producing the insecurity. This structure is the paradigm case of how affective mobilization can produce sustained political behavior that is contrary to the material interests of those being mobilized, without any individual actor being irrational given the information they have been provided. Humiliation and Status Politics Humiliation — the experience of public status degradation, of being reduced below a position one considered rightfully one's own — produces political behavior with a distinctive structure that cannot be derived from either material interest or narrative analysis alone. The humiliated actor is not primarily seeking material improvement. They are seeking restoration of status. This means that political offers that address material conditions without addressing the status claim will not resolve the political energy that humiliation generates; and that political offers that address the status claim without addressing material conditions can produce sustained support even from populations whose material conditions worsen under the arrangement they are supporting. Status anxiety — the chronic condition of populations who experience their social position as uncertain and contested rather than secured — is analytically distinct from humiliation, though it shares the feature of being partially decoupled from material interest. Status anxiety is produced by competitive social arrangements in which position must be continuously defended and cannot be taken for granted. It generates political behavior oriented toward the protection of relative position rather than the improvement of absolute conditions: populations experiencing status anxiety will resist arrangements that improve the conditions of those below them in the social hierarchy even when those arrangements do not reduce their own material conditions, because the narrowing of the status gap is experienced as a status loss. This is one of the primary mechanisms through which working-class populations in competitive societies can be mobilized against redistributive arrangements that would materially benefit them. Emotional Contagion and the Social Dynamics of Affective States

Affective states spread through social networks by mechanisms that are partially independent of the information content of the communications carrying them. Emotional contagion — the direct transmission of emotional states through mimicry, resonance, and social synchrony — is a biological phenomenon that operates below the level of conscious belief formation. A population that has been exposed to sustained fear communications will develop elevated fear responses even in individuals who did not find the specific claims in those communications convincing, because the affective state spreads through the social environment independently of the propositional content that triggered it. The implication for political communication infrastructure is direct: media environments optimized for emotional engagement — which the attention economy described in 3.8 tends to produce — are simultaneously environments optimized for affective contagion, because emotional content produces more engagement than non-emotional content at

comparable informational value. The connection to collective action works in both directions. Shared affective states reduce the coordination costs of collective action by providing a common motivational substrate that does not require alignment of specific beliefs or material interests. Populations that share a common affective state — collective grief after a focal event, collective outrage at a visible injustice, collective fear in response to perceived threat — can coordinate action more rapidly and with less prior organizational infrastructure than populations whose motivation is primarily material or analytical. This is why focal events are so consequential for transformation windows (10.5b): they are not only information about the state of the system but generators of shared affective states that dramatically lower the collective action threshold. The analytical implication is that the emotional character of a focal event matters as much as its informational content for determining what political action it can catalyze. Affective mobilization operates as a variable in the transformation pressure formula through two channels. First, collective fear and humiliation can suppress V by directing political energy toward affectively salient but structurally irrelevant targets — keeping populations focused on the misidentified source of their genuine suffering rather than the structural conditions producing it. Second, collective affective states that are accurately directed — outrage at visible structural injustice, solidarity generated by shared deprivation — amplify counter-power accumulation by lowering coordination costs and extending the duration multiplier D . The same mechanism that can be exploited to misdirect transformation pressure is also one of the most powerful mechanisms for generating it. A materialist analysis of political change that ignores affective dynamics will consistently underestimate both the resilience

of arrangements that exploit fear and the speed with which transformation can occur when affective mobilization aligns with material conditions. The implication for counter-power strategy follows directly from the analysis of how existing arrangements use affective mobilization. If fear and status anxiety can be directed toward misidentified sources — keeping populations focused on manufactured divisions rather than structural conditions — then counter-power that operates only at the level of rational argument and material interest will consistently lose to mobilization that operates at the affective level. The asymmetry is structural: actors defending existing arrangements have resources for sustained affective mobilization that challengers typically do not. Correcting the analytical error — demonstrating that the source is misidentified — is necessary but insufficient. The affective energy that has been attached to the wrong target must be redirected, not merely refuted. Redirection requires offering an affective frame that can hold the same energy and connect it to an accurate analysis of structural conditions. This is what the most effective movements in the counter-power literature have done: not argued populations out of their fear or their humiliation, but offered accounts in which those responses were accurate and the structural source was correctly identified. The organizational and narrative infrastructure of counter-power (10.9) is the mechanism through which that redirection occurs at scale. The analysis of how affective mobilization works — developed in this section — is the prerequisite for understanding what that infrastructure must produce. The psychoanalytic objection — that affective and motivational states are structured by forces below rational deliberation that material analysis cannot access — is treated here as an empirical disagreement rather than a philosophical one. The framework's account of fear, humiliation, and status anxiety treats these as material conditions with material causes operating at the level of conscious belief formation. If psychoanalytic accounts of political behavior are empirically more predictive — if unconscious drive structures explain political action better than the material-affective account developed here — the epistemological commitment of Part I requires revision toward them. What the framework resists is the psychoanalytic move that places the primary mechanisms of political behavior permanently beyond empirical reach: a theory of political action grounded in forces that cannot in principle be tested against evidence is not a contribution to the materialist analysis, whatever insights it contains. The question of how much political behavior is explained by unconscious structure versus by manufactured narrative and information asymmetry is a genuinely open empirical question that neither framework can settle from within its own assumptions. The analysis of cognitive constraints (3.8) and

ffective mobilization (3.8b) establishes two factors that modify how material conditions translate into political behavior. A third factor operates between material conditions and behavior: the beliefs that actors hold about their situation, its causes, and the options available to them. PMN's structural analysis tracks what arrangements produce materially; belief operates as an intervening variable between what arrangements produce and how actors respond. This is not a concession to idealism: beliefs are themselves materially produced and structurally shaped. But they are not reducible to the conditions that produce them. Two populations in identical structural conditions can hold different beliefs about those conditions — about whether their suffering is natural or institutional, whether alternatives are available, whether collective action is possible — and those differences in belief produce different behaviors under identical material pressures. The visibility suppression variable (V) operates partly through this mechanism: V is high not only when information is withheld but when the belief structures available to interpret information prevent structural causation from being recognized. Reducing V therefore requires not only making information available but providing the interpretive frameworks within which structural causation can be recognized as such. Belief revision does not occur automatically when material conditions change. The conditions under which beliefs update — and why they sometimes resist updating despite changing conditions — are analytically important for understanding why structural suffering persists in populations that have both the information and the structural conditions to respond differently. Beliefs change when: (1) material incentives shift sufficiently that holding prior beliefs becomes individually costly in ways that outweigh the social and identity costs of revision; (2) trusted authorities within an existing interpretive community change position, providing cover for individual revision without the cost of community exit; (3) reference communities shift, making the identity-sustaining function of prior beliefs available from a new community that holds revised beliefs; (4) psychological and social costs of maintaining the belief in the face of mounting disconfirming evidence exceed the costs of revision. The last condition is least reliable: it predicts that beliefs resistant to evidence are those most deeply tied to identity and community, which is exactly when the disconfirming evidence tends to be greatest. This is the configuration that 3.8c identifies as stable irrationality and that Žižek identifies as fetishistic disavowal: the pattern where evidence accumulates without producing belief revision because revision would require restructuring the identity and community that the belief sustains. The practical implication for counter-power strategy (10.9) is that information provision is a necessary but insuf

ficient condition for belief revision in populations whose structural suffering is protected by this mechanism. Counter-power strategy must also address the social and identity conditions within which belief revision becomes possible without community exile. Note the relationship to 3.8c: stable irrationality and fetishistic disavowal are the structural configuration in which all four conditions above have been made unavailable simultaneously — material incentives are structured to reinforce the belief, trusted authorities are captured, reference communities are closed, and the psychological cost of revision has been installed as identity threat. Conditions 1–4 describe belief revision in the general case; 3.8c specifies the specific structural configuration in which the general case does not apply. Both analyses are required for a complete account. Belief-structure diagnostic: (1) What beliefs about structural causation does this population hold, and how were those beliefs produced? (2) What social and identity functions do those beliefs perform — what community membership do they signal, what would be risked by revising them? (3) What conditions would enable belief revision without requiring identity restructuring — are trusted authorities available to provide cover, or alternative communities to provide membership? (4) Is the belief structure stable irrationality (3.8c) or information deficit (addressable by V reduction)? The distinction determines the intervention: information provision addresses information deficit; social and community restructuring is required for stable irrationality; the misuse triggers in 12.5d apply when analysts mistake one for the other and design interventions accordingly.

3.8c — Stable Irrationality as a Structural Variable

Sections 3.8 and 3.8b analyze cognitive constraints and affective mobilization as factors that modify how material conditions translate into political behavior. Both treat irrationality primarily as a limitation on or distortion of rational response to material conditions. This section addresses a distinct and analytically more challenging phenomenon: patterns of behavior that are irrational by the standard of the individual actor's own material interests, that persist across time and generations, that cannot be reduced to information deficits or affective capture, and that are reproduced by social structures in ways that make them effectively stable features of the system rather than correctable errors. The distinction matters because the appropriate analytical and practical responses differ: correctable irrationality calls for better information and affective reorientation; structurally stable irrationality calls for institutional redesign that produces different outcomes without requiring the underlying pattern to be eliminated. Stable irrationality differs from temporary irrationality in its structural source. Temporary irrationality — panic, grief-driven decision-making, enthusiasm that outruns evidence — is produced by acute conditions and resolves as those conditions change. Stable irrationality is reproduced by the normal operation of enduring social structures: by status hierarchies that generate consistent in-group favoritism; by institutional incentives that reward specific types of motivated reasoning; by socialization processes that establish interpretive frameworks resistant to evidence because their function is membership rather than accuracy; and by identity formations that make accurate perception of material conditions threatening rather than useful. The structural feature these mechanisms share is that they are not bugs of the system — they are functional from the perspective of the system's reproduction, even when they are dysfunctional from the perspective of the individual actors and the populations whose interests the framework cares about. A system that maintains cohesion through consistent in-group favoritism is not making an error it will eventually correct; it is operating as designed. The framework's treatment of capture (7.3) provides the institutional version of this analysis: institutions can be captured in ways that are stable precisely because the capture serves the interests of those with the resources to maintain it, even when it is irrational from the perspective of the institution's stated function and from the perspective of the populations the institution claims to serve. The structural irrationality analysis generalizes this beyond institutional settings to the cognitive and behavioral patterns that make populations susceptible to capture in the first place and resistant to the conditions that would allow them to perceive and resist it. Status quo bias, loss aversion, in-group solidarity that persists despi

te material harm — these are not failures of individual cognition that better information will correct; they are features of how human cognition and social organization interact, with structural causes and structural reproduction. The formula architecture in Part XV treats these patterns as components of the repressive capacity term (15.6): arrangements that successfully produce stable irrationality in relevant populations reduce the translation of structural suffering into transformation pressure, decoupling the causal chain that the framework otherwise assumes. An analysis that does not account for stable irrationality will consistently overestimate how quickly high structural suffering translates into political change, and consistently underestimate the durability of arrangements that produce it. The practical implication for transformation analysis: when populations fail to respond to high structural suffering with the political mobilization the framework's formula would predict, the first diagnostic question is not "are our measurements of suffering wrong?" but "what structural mechanisms are maintaining stable irrationality in this population, and how are they doing so?" This reframes the analytical problem from a measurement problem to an institutional one, which is the reframing the framework's general orientation supports. The visibility suppression variable (V) captures part of this — when suffering is not recognized as structural, it does not generate structural demands. But stable irrationality operates beyond visibility: it affects how populations respond even when structural causation is visible, generating responses that serve the reproduction of the arrangement rather than its transformation. This is the most durable form of institutional resilience and the hardest for counter-power to address. A specific configuration of stable irrationality deserves explicit acknowledgment: cases where suffering itself becomes a source of identity, moral superiority, or community belonging — what Žižek, drawing on Lacan, terms fetishistic disavowal. This is not a separate variable but a specific interaction

between meaning infrastructure (5.6) and cosmological capture (7.3c-ii): the actor derives affective reward from the very condition that harms them, making the disavowal stable not because information is absent but because accurate information would threaten the identity and community that the suffering constitutes. Counter-power strategies that target information deficits will consistently fail against this configuration; the intervention required is at the level of the meaning infrastructure that makes the suffering identity-constituting.

3.9 — Demographic Dynamics as Material Variables

Population structure is a material variable with consequences that are as significant as any ideological commitment, and that operate largely independently of ideology. The age distribution of a population, its dependency ratios, its growth rate, its urban-rural distribution, and its geographic density are not background conditions that political analysis can treat as fixed. They change, and when they change, the institutional systems built around the previous distribution face structural pressures that ideology cannot resolve. (Malthus 1798; Cohen 1995) A population with a large cohort of young adults relative to the capacity of existing economic and institutional arrangements to absorb them creates specific transformation pressure that does not require high per-capita suffering to generate significant political instability. The mechanism is structural rather than moral: young adults who cannot find institutional roles commensurate with their educational attainment and generational expectations have lower opportunity costs for political mobilization than older adults with established positions. They have greater social density — larger peer networks through which mobilization can propagate — and face institutional gatekeeping that generates status frustration independently of absolute deprivation. (Goldstone 1991) This dynamic is a conditional tendency, not a deterministic prediction of violence or an essentialist diagnosis of youth; its manifestation depends on employment absorption, political channels, and state response capacity. Systems facing youth bulges experience compressed temporal windows for institutional accommodation before threshold dynamics begin (see §15.11 for demographic cycles, demographic windows, and Strategic Transformation Analysis timing). An ageing population creates a structurally distinct set of material constraints: increasing dependency ratios, growing fiscal demands on healthcare, pensions, and long-term care systems, alongside a contracting formal labour force. The resulting distributional conflict between generations with divergent material stakes in current fiscal allocations generates structural tensions that institutional design must manage. Population ageing also shifts the median voter's temporal horizon through established political representation channels, systematically biasing allocation toward immediate consumption maintenance at the expense of long-term public goods and adaptive institutional investments. This time-horizon contraction is an institutional incentive effect, not an inevitable psychological trait of older individuals. Urbanization changes the density, geography, and nature of social interaction, which changes the forms of collective action that are viable, the information environments that people inhabit, and the economic structures that can be sustained. Rural and urban populations face different material conditions and develop different institutional arrangements in response. Treating thi

s divergence as a cultural difference rather than a material difference predictably misreads the dynamics governing political possibility. Urban Space as Material Infrastructure: Geography, Visibility, and Coalition Formation The spatial organization of cities is not a neutral backdrop to social and political life. It is a material infrastructure that mediates demographic realities, shapes the distribution of floor conditions, controls the visibility of structural suffering, and conditions the capacity for coalition formation. Housing, transit networks, public space, zoning regulations, and the physical separation or integration of populations by class and function are the tangible products of institutional decisions that reflect existing distributions of power and are subject to capture dynamics. The relationship between urban spatial organization and the visibility variable V in the structural suffering formula operates through two primary mechanisms. The first is physical segregation: when populations experiencing structural suffering are concentrated in spatial

peripheries, informal settlements, or segregated enclaves that more privileged populations rarely traverse, the suffering is physically removed from public attention, diminishing the political pressure that would otherwise attach to it. The second is design for invisibility: urban transit corridors and commercial districts engineered to bypass or exclude marginalized populations function as a continuous material suppression of visibility without requiring explicit censorship. Spatial segregation thus operates as a direct mechanism of visibility suppression (analyzed further in the governing narrative framework in Part VIII). Urban space directly structures the coalition formation capacity that the framework's analysis of counter-power requires. Dense, mixed-use, and accessible urban environments generate the incidental social contact and shared civic infrastructure that enable cross-sectional solidarities — the recognition of shared material interests across populations that do not share cultural or ethnic identities. Conversely, automobile-dependent suburban sprawl, gated residential developments, and the aggressive privatization of public gathering spaces degrade this civic substrate, increasing the fragmentation variable *If* by foreclosing the physical spaces where commonalities can be discovered and organized (see Part X for counter-power organization and spatial coordination). The framework's analytical obligation: demographic and spatial dynamics must be incorporated into structural political analysis as primary material constraints rather than background context. Every structural assessment must explicitly specify the relevant timescale, spatial distribution, institutional absorptive capacity, and potential countervailing conditions. An assessment that ignores demographic cycles and spatial mediation will consistently misdiagnose both the urgency of transformation pressures and t

he viable windows for institutional reform.

3.10 — Ecological Constraints: The Ground Beneath All Systems

Every system this analysis examines — biological, institutional, economic, geopolitical — exists within an ecological context that sets the absolute outer limits of what is possible. This is not an environmental concern appended to a political framework; it is what the biological foundation requires when followed to its conclusion. Human social systems are not separate from ecological systems. They are embedded within them, dependent on them, and constrained by them in ways that are not negotiable through institutional design or ideological commitment. The carrying capacity of agricultural land, the availability of freshwater, the stability of climate systems, the biodiversity that maintains ecosystem services — these are material conditions that determine what human populations can sustain and what institutional arrangements are viable at what scale. The distinctive feature of ecological constraints is their non-linearity. Systems that appear to absorb pressure within tolerable limits can cross planetary boundary thresholds (Rockström et al. 2009) beyond which they reorganize rapidly and catastrophically. Fisheries that have sustained harvesting for centuries can collapse within a decade when the threshold is crossed. Climate systems that have maintained relatively stable conditions for ten thousand years can shift to new equilibria within decades. The institutional implication is that the standard political logic of incremental adjustment — making small changes and observing the results — can produce catastrophic outcomes when the underlying system is approaching a threshold that is not visible in the incremental data. The relationship between ecological constraints and the permanent conditions identified earlier in this Part is direct. Resource scarcity — one of the permanent conditions — is ultimately an ecological phenomenon: it is determined by what the planetary system produces and what human activity consumes. New threats emerge partly from the disruption of ecological systems that previously contained them. The capacity of institutions to manage these conditions depends on whether the ecological context within which they operate remains within the range that the institutions were designed for. A framework that begins from the biology of organisms and follows the logic of material constraints consistently cannot treat ecological limits as externalities to be managed by market mechanisms or political will. They are the conditions within which the entire analysis operates. Any doctrine derived from this framework that does not account for ecological constraints is a doctrine that will eventually be revised by the conditions it failed to account for. One of the places where the teleological structure of

historical materialism produced the most serious errors was in its treatment of the conditions it analyzed as transitional — as stages to be passed through on the way to an end state where the underlying tensions would be resolved. Resource s

carcity, institutional betrayal, resistance to coordination — these were features of class society that would disappear when class society was superseded. They did not disappear. They persisted because they are not features of class society. They are features of what human social organization is. The word 'permanent' requires immediate clarification, because it is easy to read as fatalistic — as the claim that nothing can change, that effort is useless, that the conditions of suffering are fixed. That reading is precisely wrong. 'Permanent' here means that specific structural conditions cannot be eliminated by any institutional arrangement, however well designed. It does not mean their consequences cannot be reduced, managed, and made less damaging. Resource scarcity persists, but how its costs are distributed is not fixed. Institutional betrayal is always possible, but how it is constrained and corrected is not fixed. The permanence is in the condition, not in the outcomes the condition produces. The difference between societies that manage these conditions well and those that manage them badly is real, consequential, and the primary target of the framework's political analysis. Resource scarcity persists because demand grows with capacity and the frontier of scarcity moves rather than disappears. Institutional betrayal is always possible because the incentive structures that produce it are generated by institutional positions rather than by the character of the people who occupy them. Resistance to coordination always emerges because any mechanism that distributes costs and benefits will always generate parties for whom defection is locally rational. A doctrine that treats these as transitional difficulties is a doctrine that will be surprised by their persistence every time they recur — and will explain that surprise by blaming the people it was supposed to serve. Permanent Does Not Mean Fixed in Form A clarification the framework requires of itself: calling these conditions permanent is a claim about their structural necessity, not about the specific form they take. The permanent conditions — scarcity, betrayal potential, knowledge limits, ecological constraint — are permanent in the sense that no institutional arrangement eliminates the underlying problem. They are not permanent in the sense that the specific ways they manifest are fixed across history. The distinction matters for how the framework is read and how it should not be. Scarcity is the clearest example. The specific resources that are scarce change with technology, ecology, and social organization. The scarcity of cultivable land that organized most of pre-industrial political economy has been partially displaced by the scarcity of freshwater, of atmospheric carbon absorption capacity, and of the rare minerals that advanced electronics require. The transition from fossil energy to renewable energy does not eliminate energy scarcity — it transforms its form, relocating the constr

int from the availability of combustible material to the availability of manufacturing capacity, grid infrastructure, and the land area required for generation and storage. The permanent condition is that human social organization requires resource inputs whose availability is constrained and whose distribution is contested. The specific constraint changes. The structural condition it generates — that arrangements for managing the constraint must exist and will be sites of political conflict — does not. The same applies to betrayal potential. The specific forms through which the capacity for defection from cooperative arrangements is exercised change with institutional design, surveillance capacity, and social norms. High-trust communities with strong monitoring can reduce defection rates substantially. But they cannot reduce the underlying capacity for defection to zero, because that capacity is biological — it is the same capacity for individual calculation that makes cooperation possible in the first place. Institutional arrangements that assume they have eliminated betrayal potential rather than managed it have not understood the condition. They have created the conditions for a particularly consequential betrayal when the institutional suppression of the capacity eventually fails. The practical implication is that analysis which treats the permanent conditions as obstacles to be overcome — rather than structural features to be managed — consistently produces the same failure: it designs institutions for a world in which the condition has been solved, discovers that the condition persists in a new form, and lacks the analytical resources to recognize

what it is encountering. The framework's insistence that the permanent conditions are permanent is not fatalism. It is the precondition for building institutions that are robust to conditions they cannot eliminate — which is the only kind of robustness available. Planetary Boundaries as a Constraint Structure, Not an Environmental Issue The framework's treatment of ecological constraints requires a clarification about their analytical status. Ecological limits are not a category of policy concern to be weighed against economic and social objectives. They are a constraint structure of the same kind as the biological floor — conditions that, if violated beyond certain thresholds, eliminate entire ranges of viable configurations at all other layers. The concept of planetary boundaries, developed in earth system science to identify the biophysical limits within which human civilization has operated for the past ten thousand years, provides the empirical content for this constraint level. Climate stability, biodiversity integrity, freshwater availability, biogeochemical cycles, and land system change are not independent environmental issues — they are interlocking components of the Earth system whose stability is the precondition for the biological floor that PMN treats as the foundation of all other

analysis. The relationship between planetary boundary violation and PMN's evaluative criterion is direct and non-negotiable. The structural availability of becoming — the condition the anti-foreclosure criterion protects — requires a stable enough biological and material environment to sustain the populations in whom becoming is possible. Becoming without ecological foundation is not a more expansive form of becoming. It is a form of becoming that forecloses becoming for future populations by consuming the preconditions they would require. This is why ecological collapse cannot be treated as one cost to be weighed against the benefits of economic growth. It is a structural constraint on the meaning of the evaluative criterion itself: a becoming that systematically destroys the planetary conditions for future becoming is not a candidate for maximization under that criterion — it is a candidate for identification as the most severe form of intergenerational structural harm available. The intergenerational transmission variable G acquires a specific additional dimension under conditions of active planetary boundary violation. G normally captures the degree to which structural suffering and foreclosed becoming are reproduced across generations through institutional, material, and epistemic inheritance. Under conditions of ecological degradation, G also captures the degree to which the material preconditions for becoming are themselves being transferred in degraded form. A generation that inherits stable institutions but a destabilized climate, depleted freshwater systems, and reduced biodiversity is receiving a net negative transfer in the most fundamental dimension — the biological and material substrate on which all other inheritance depends. The standard G analysis, focused on economic and institutional transmission, systematically undercounts intergenerational harm when planetary boundaries are being violated, because it counts what is transferred but not what is being destroyed in the process of transferring it. The constraint structure imposed by planetary boundaries also reframes the analysis of scarcity developed earlier in this section. The original analysis correctly identifies scarcity as a permanent condition that changes form rather than disappearing — from fossil energy to renewable energy infrastructure, from biomass to agricultural yield, from local watersheds to global hydrological systems. The planetary boundary analysis adds a threshold dimension to this permanent scarcity: for most of human history, the regenerative capacity of Earth systems has exceeded the rate of human extraction, meaning that scarcity was experienced at the level of distribution and access rather than at the level of total planetary capacity. Contemporary conditions of accelerating boundary violation mark the emergence of a structurally different form of scarcity — one in which the regenerative capacity itself is being degraded, meaning that the total material

substrate available for human becoming is shrinking in real time. This is not cyclical scarcity of the kind that prior human societies navigated. It is scarcity produced by cumulative extraction from a finite system, and it does not respond to the institutional innovations that addressed distributional scarcity. The eco-centric objection — that grounding evaluation in human suffering and human becoming is a structural exclusion of the intrinsic value of non-human nature — identifies a real boundary of this framework rather than a correctable error within it. Section 3.4c acknowledges that the framework's logic does not inherently exclude

non-human organisms capable of suffering and development, but that extending it fully requires philosophical work this document does not undertake. Section 14.4 identifies this as one of the most important directions for a successor framework. What this section adds is the claim that ecological constraints are not merely the background within which human activity occurs but the material ground on which all the framework's other analysis depends — which means that treating them as externalities to be managed by market mechanisms or political will is not just an ethical failure but an analytical one. The framework does not claim that only human suffering matters. It claims that its primary analytical object is the institutional arrangements humans build, and that those arrangements cannot be adequately analyzed without the ecological conditions that make them possible and the ecological consequences that make them unsustainable. The implication for institutional analysis: any institutional arrangement that requires ongoing violation of planetary boundaries for its maintenance is not a viable configuration — not because of its normative failings but because it is consuming the preconditions for its own continuation. Contemporary capitalism's dependence on fossil energy extraction and atmospheric carbon absorption capacity as free inputs to production is not a policy choice that could be made differently within the same institutional form. It is a structural feature of capital accumulation under current conditions: the competitive pressure to externalize ecological costs is built into the incentive architecture of private ownership of productive assets at the scale and form that contemporary capitalism takes. This is the fourth diagnostic in section 11.2's analysis of private ownership — ecological boundary conditions — and the planetary boundary framework gives it its most precise formulation: the institutional form that systematically externalizes ecological costs is not merely inefficient or unjust. It is structurally incompatible with the continuation of the biological conditions that make any form of becoming possible.

3.11 — Illustrative Cases: The Chain in Operation

The causal chain from material conditions to institutional outcomes is not a theoretical abstraction. It is a pattern that appears consistently across historical cases that differ in almost every other way. Structural inequality and institutional collapse. Societies where the benefits of economic coordination are systematically captured by a narrow elite while costs are borne broadly exhibit a consistent pattern: declining institutional legitimacy, rising political volatility, and eventual institutional breakdown when a specific event — economic shock, military defeat, natural disaster — activates the accumulated loss of confidence. The pattern appeared in pre-revolutionary France, in Weimar Germany, in multiple twentieth-century postcolonial states, and in contemporary societies where wealth concentration has outpaced the institutional capacity to maintain the appearance of legitimacy. The specific ideological content of the breakdown varies. The material substrate — structural suffering generated by institutional arrangements that benefit the few at the expense of the many — is consistent. Cooperation breakdown in failed states. When institutional enforcement capacity collapses, the behavioral responses predicted by the framework appear rapidly and consistently. Trust contracts to the radius of personal relationships. Economic activity retreats to immediate and easily enforceable exchanges. Investment in long-term productive capacity stops because the conditions for recovering returns cannot be guaranteed. Violence becomes an individually rational substitute for the institutional enforcement that is no longer available. The sequence is: enforcement collapse → defection becomes rational → cooperation retreats → production capacity falls → suffering increases → defection becomes more rational. Stability from strong social infrastructure. Societies that invested heavily in universal social insurance during the mid-twentieth century — healthcare, education, housing security, unemployment protection — demonstrated a consistent pattern: higher social trust, stronger institutional legitimacy, greater political stability, and more resilient responses to economic shocks than societies with weaker social infrastructure. The mechanism is the one the framework predicts: reduced structural suffering → reduced behavioral pressure toward defection → maintained cooperative substrate → institutional legitimacy preserved → capacity for collective response to shared problems retained. These examples are illustrative, not exhaustive. The framework does not claim that material conditions fully determine outcomes — the

specific ideological content, the quality of leadership, the role of specific events all matter for the form that outcomes take. What it claims is that material conditions are primary in determining the range of outcomes that are possible. The three cases below are not exhaustive or representative — they are illustrative of specific links in the causal chain

1 chain from biological constraint to institutional outcome. Each case is selected because it demonstrates a link that is easy to abstract but whose material weight is best understood through the concrete operation of the mechanism. The cases differ in scale, period, and domain; the structural pattern they illustrate is consistent. Case A — Inequality and institutional collapse: The pattern of rising inequality producing institutional delegitimation and eventual collapse is documented across a wide enough range of historical cases — late Roman Republic, ancien régime France, Weimar Germany, multiple Latin American democracies in the twentieth century — to constitute a structural regularity. The mechanism is not that inequality causes direct political mobilization in a simple causal chain. It operates through the interaction of three variables: declining institutional performance as elite capture degrades the institution's capacity to address non-elite needs, rising visibility of the gap between institutional legitimacy claims and actual institutional function, and the availability of a focal event — economic shock, military defeat, political scandal — that provides the coordination point for the accumulated pressure to act. The inequality is not the immediate cause of collapse; it is the structural condition that makes institutions fragile enough that relatively ordinary shocks produce extraordinary outcomes. Case B — Cognitive limits and organizational scale: The transformation of social movements into formal organizations, documented across labor movements, civil rights organizations, environmental advocacy groups, and political parties, follows a consistent pattern that reflects the Dunbar constraint operating at scale. Organizations that begin as networks of direct relationships among people with shared commitments become, as they grow, formal bureaucracies with professional staff, standardized procedures, and internal hierarchies designed to maintain coordination past the threshold where direct relationships can sustain it. The transformation is not a failure — it is what growth requires. But it produces a characteristic shift in organizational culture: the people whose sustained attention is required to maintain the formal organization develop institutional interests in its continuation that may diverge from the original commitment that produced the growth. The organization that survives its growth transition typically does so by becoming something different from what it started as — more stable, more institutionally capable, and less responsive to the membership whose mobilization produced it. Case C — Environmental constraint and political instability: The relationship between resource constraint — particularly water scarcity and agricultural disruption — and political instability is consistently underweighted in political analysis because its causal chain operates slowly and through multiple intermediate steps. The mechanism: sustained r

esource constraint degrades the material conditions of populations dependent on those resources, producing the structural suffering that generates political pressure; that pressure interacts with existing institutional arrangements whose capacity to respond is itself shaped by who controls the resources being depleted; when the institutional response is inadequate — either because capture has aligned it with resource controllers rather than affected populations, or because the scale of constraint has exceeded institutional adaptive capacity — the pressure produces volatility. Syria 2011 and the Sahel conflicts of the 2010s are cases where this chain operated visibly: a drought of unusual severity intersected with institutional arrangements already degraded by decades of capture, producing instability whose immediate causes were political but whose structural roots were material in a biological sense — in the relationship between land, water, and the biological requirements of the populations depending on them. The cases share a structural feature that matters for analysis: in each, the proximate cause of the outcome — political mobilization, organizational transformation, political instability — was not the biological constraint itself but the interaction between the constraint and the institutional arrangement through which populations experienced it. Biological constraints do not directly produce political outcomes. They produce material conditions that interact with institutional arrangements to produce political pressure, which then interacts with the configuration of counter-power and

repressive capacity to produce specific outcomes. The causal chain from biology to institution to politics runs through all three levels simultaneously, and analyses that skip any level will consistently misattribute causes and therefore misidentify interventions.

3.12 — Rational Micro, Problematic Macro

Individually rational behavior can produce collectively irrational outcomes. Marx saw this in economic production: each capitalist acting rationally to maximize profit contributes to conditions — overproduction, falling rates of profit, periodic crises — that are bad for capitalism as a system. The observation is not limited to markets. It appears in political movements, institutional dynamics, international relations, and the dynamics of social change — wherever the incentive structure facing individuals diverges from the consequences their aggregate behavior produces. The scope of this generalization is one of the most underappreciated features of materialist analysis: because it holds wherever the relevant structure of incentives obtains, it can explain collectively harmful outcomes without requiring any special assumptions about bad faith, irrationality, or malice in the actors involved. The generalization is important because it corrects one of the persistent failures of political analysis: the tendency to explain bad collective outcomes by reference to bad individual intentions. When individually rational behavior produces collectively harmful outcomes, the appropriate response is not moral condemnation of the individuals whose rational behavior produced the outcome. It is the redesign of the conditions under which individual rationality operates — changing the incentive structures, the information environment, the coordination mechanisms — so that rational individual behavior produces better collective outcomes. This is a structural response to a structural problem, and it is the response that this framework consistently recommends. The collective action problem — the gap between individual rationality and collective outcome — is one of the most structurally consequential features of complex social systems, and one of the most consistently underestimated. It is underestimated not because analysts are unaware of it theoretically, but because the institutional designs that are supposed to address it are themselves subject to the same structural dynamics that produce it. The commons dilemma is solved by governance institutions; but the governance institutions are themselves subject to capture by actors who find it individually rational to exploit them. The solution produces the problem at a higher level. The scope of the problem: The divergence between individual rationality and collective outcome operates across domains that share no surface similarity. In ecological systems: individually rational resource extraction produces collective depletion that reduces the total extraction available to all actors, including the individually rational extractors. In financial systems: individually rational risk-taking by financial institutions, each responding to competitive pressures that make risk-taking locally optimal, produces systemic fragility that collapses the system whose stability is a precondition for any of the individual risk-taking being

profitable. In political systems: individually rational disengagement from political processes whose outcomes any individual actor cannot significantly affect produces collective under-provision of the oversight that those processes require to function. In each case the mechanism is the same: the incentive structure facing individuals diverges from the incentive structure that would produce good collective outcomes, and without institutional mechanisms that internalize collective consequences into individual incentives, individual rationality produces collective failure. The institutional response and its limits: The standard institutional response to collective action problems is the creation of mechanisms that internalize externalities — taxes, regulations, norms, coordination bodies — that change the individual incentive structure to align with collective outcomes. This response is correct in form and consistently inadequate in practice, for a structural reason: creating the institutions that internalize externalities requires solving a collective action problem at a higher level. Someone has to fund the institution, staff it, and maintain it against the interests of actors who benefit from the absence of internalization. Those actors have concentrated interests and resources; the beneficiaries of internalization have diffuse interests and less concentrated resources. The institutional

response to collective action problems therefore consistently requires more counter-power capacity to create and maintain than the surface difficulty of the problem would suggest. The compounding mechanism: Collective action failures compound because they degrade the institutional capacity that would address them. Ecological depletion reduces the resource base that populations need to invest in institutional governance. Financial instability destroys the trust infrastructure that coordination requires. Political disengagement degrades the institutional quality that would give engagement higher expected returns. Each failure makes the next one more likely and more severe, because the institutional capacity to address collective action problems is itself a collective good subject to under-provision. This compounding dynamic is the structural basis for the threshold effects analyzed in 10.5: systems under accumulating collective action failure do not degrade linearly — they accumulate fragility until a specific event triggers non-linear reorganization. The analytical implication: when assessing any arrangement where individual behavior diverges from collectively optimal outcomes, the relevant question is not only what incentive mechanism would align individual and collective rationality, but what collective action problem must be solved to create and maintain that mechanism, and what counter-power capacity is required to solve that second-order problem. Analyses that stop at the first-order incentive design consistently produce recommendations that are technically

correct and practically unimplementable, because they treat the second-order problem as tractable without analyzing it. The connection to Part XV: the formula architecture operationalizes the collective action analysis in terms of the variables that determine whether threshold conditions are met for institutional transformation. The gap between individually rational behavior and collectively optimal outcomes is one of the primary drivers of the structural suffering variable S — it produces the suffering that generates transformation pressure — and one of the primary constraints on the adaptive capacity variable C — it limits the capacity of existing institutions to address the problems they produce. Both connections run through the causal chain in 3.7 and the institutional arc in 3.6: collective action failure is not an anomaly in otherwise functional systems, it is a structural feature of the arc from biological organism to institutional arrangement that compounds at every level of institutional complexity.

IV: What Matters: Ethics, Value, and the Evaluative Commitment

The examples are not limited to markets. Electoral systems in which each voter rationally abstains — because the probability that any individual vote decides an outcome is vanishingly small, while the costs of voting are real — produce low-turnout equilibria that are collectively irrational for any population whose interests require high participation. Common-pool resource arrangements in which each individual actor rationally extracts more than their sustainable share — because the costs of restraint fall on the restraining individual while the benefits of restraint are distributed across all users — produce collective depletion that harms all users including the individually rational ones. Arms races in which each state rationally increases its military capacity in response to the increases of potential adversaries — because unilateral restraint produces vulnerability while matched increases produce only the same relative position at higher cost — produce collectively irrational expenditure patterns that leave all parties less secure. The structural form is the same across cases: individual rationality under conditions of distributed consequence and absent coordination.

The implication for institutional design is specific and does not follow from the observation alone. It is not that individual rationality should be replaced by collective obligation, or that incentive structures can be designed perfectly enough to eliminate the problem. It is that the conditions under which individual rationality produces collectively harmful outcomes can be identified in advance — they have a recognizable structure — and that identified conditions are the appropriate target of institutional intervention. Coordination mechanisms, credible commitment devices, monitoring and enforcement arrangements, and the redesign of consequence distribution are the tools. They work better in some domains than others, and they introduce their own capture risks that require separate analysis. What they do not do is eliminate the underlying tension: any arrangement complex enough to coordinate behavior at scale will generate some domain in which the micro-macro gap reappears in a new form.

This is one reason the framework treats institutional redesign as a permanent activity rather than a problem to be solved. The rational micro, problematic macro dynamic is not a correctable error in a system that could otherwise be designed to avoid it. It is a structural feature of any system composed of agents with local information and distributed consequences. The appropriate orientation is not to design the system that finally eliminates it but to maintain the institutional capacity to identify where it is currently operating, what interventions are available, and what second-order consequences those interventions are likely to produce. That capacity is itself a collective good subject to the same dynamic it is meant to address — which is why it requires explicit maintenance rather than emerging spontaneously from the individual interests of those who would benefit from it.

4.1 — The Foundation of Value

Marx refused to ground his critique of capitalism in moral philosophy — famously declining to specify what justice required, insisting instead that the analysis of capital would show that the system contradicted its own conditions of reproduction. This was methodologically serious: it avoided the circularity of critiquing a system from standards that the system itself had generated. The weakness was that it left the framework without a clear account of why the suffering produced by capitalism was bad — which made it difficult to evaluate proposed alternatives against any standard that was not itself derived from the system being critiqued. The resolution here is different. The evaluative standard is not derived from any particular cultural tradition, any particular stage of historical development, or any particular system of production. It is derived from the biological foundation: unnecessary suffering is real, and its reduction is a legitimate goal across all cultural and ideological contexts. This is more durable than a moral claim because it rests on what organisms are rather than on what any tradition believes. And it is more specific than the refusal to ground the analysis

in value — it provides a criterion that can actually be used to evaluate competing arrangements. Unnecessary suffering, operationally defined, is suffering that: (1) is produced by identifiable structural arrangements rather than by conditions no available alternative could prevent; (2) could be reduced or eliminated by feasible institutional alternatives given current material conditions; and (3) is not an unavoidable cost of producing a greater reduction in suffering elsewhere. This definition excludes suffering that is irreducible given any available option — scarcity, biological aging, the costs of any transition including beneficial ones. It includes suffering produced by arrangements that serve concentrated interests at the expense of the populations subject to them, suffering that is preventable but sustained by institutional inertia or capture, and suffering that is deliberately or structurally invisible. The distinction between necessary and unnecessary is always an empirical judgment about available alternatives, not a metaphysical claim about what suffering is inherent to human existence. The Aristotelian Lineage and Its Revision The concept of becoming that organizes the evaluative framework from 4.5 onward has a lineage that requires acknowledgment: it is recognizably Aristotelian. Aristotle's account of eudaimonia — flourishing, or living and doing well — grounds value not in the experience of pleasure or the satisfaction of preferences but in the actualization of characteristic human capacities. A life goes well, on this account, when the capacities that are distinctively human — practical reason, social relationship, perception, imagination, and the range of activities that express what human beings characteristically are — are exercised rather than stunted

d. An arrangement is better when it makes the actualization of those capacities more available; worse when it forecloses them. This framework inherits the structure of that argument while refusing its teleological foundation. Aristotle's account of flourishing depends on a claim about what human beings are for — on a natural teleology in which organisms have characteristic functions that define what counts as their development. The framework's biological grounding provides the foundation without the teleology: what matters is not what human beings are designed for but what they are capable of, given their biological constitution and the structural conditions they inhabit. The displacement from teleology to capacity preserves the evaluative force of the Aristotelian argument — that some conditions enable development and others foreclose it — while making it defensible against the objection that natural purposes cannot do normative work without importing assumptions that go beyond the biological evidence. The capabilities approach developed by Amartya Sen and Martha Nussbaum translates this Aristotelian inheritance into the language of political philosophy and development analysis. Sen's account of development as freedom — as the expansion of the substantive freedoms that allow people to lead lives they have reason to value — is structurally continuous with the framework's account of conditions for becoming. Nussbaum's central capabilities list specifies in more detail what those conditions include: life, bodily health, bodily integrity, the use of senses and imagination, emotional development, practical reason, affiliation, relationship to other species, play, and political and material control over one's environment. (Sen 1999; Nussbaum 2011) The framework's relationship to the capabilities approach is one of structural alignment with a methodological difference. Where Nussbaum specifies a list of central capabilities as the content of what arrangements must protect, this framework specifies structural conditions for becoming without fixing the content of becoming in advance. The difference is deliberate: a fixed list of capabilities, however defensible, risks importing the specific priorities of the tradition that produced it as universal requirements. The framework's evaluative criterion — do these conditions expand or foreclose the structural availability of genuine development — is more abstract but more portable: it can be applied across the full range of human contexts without requiring prior agreement on which specific capabilities count as central. The cost is less specificity. The benefit is greater applicability across the conditions of genuine pluralism that any universal framework must be able to navigate. The biological foundation solves the grounding problem in a specific way: it grounds the evaluative criterion in conditions that are prior to the ideological disagreements through which normative frameworks typically contest each other.

her. The criterion — reduction of structural suffering, expansion of conditions for genuine development — is not derived from a particular cultural tradition, a specific historical project, or a contested philosophical

position. It is derived from what the organisms that hold all those traditions and positions require to continue existing and functioning. This prior character is what makes it a foundation rather than an entry into normative debate: it is not one normative position among others but the conditions on which all normative positions depend for their material possibility. The relationship between suffering reduction and becoming: Part III grounds the minimal anchor in suffering reduction; Part IV extends it through the becoming framework. The extension is not a replacement: suffering reduction remains the floor criterion that specifies what no institutional arrangement can accept as an acceptable cost. The becoming framework adds the anti-foreclosure criterion — not merely the absence of unnecessary suffering but the presence of conditions under which human capacities can genuinely develop. An arrangement that eliminates unnecessary suffering but maintains populations in conditions of stable deprivation, where material survival is secured but genuine development is structurally impossible, meets the floor criterion but fails the anti-foreclosure criterion. Both criteria are necessary; neither alone is sufficient for the evaluative standard the framework requires. The two-level architecture established in 3.0 — life as ontological foundation, suffering as evaluative criterion — extends the foundation of value analysis here. The evaluative standard is suffering-based for the reasons 4.1 identifies: it is cross-culturally recognizable, resistant to ideological capture, and directly actionable through institutional design. But the ontological grounding in life is what makes the evaluative criterion non-arbitrary. The framework does not simply assert that suffering is bad and build from there — that would be the first-principle assertion it aims to avoid. It grounds the evaluative criterion in the ontological observation that suffering is a signal of disruption to life's normal functioning, and that life's normal functioning includes the cooperative organization that makes institutional arrangements both possible and evaluable. The foundation of value is therefore not suffering as such — it is life as an organized, disruption-sensitive, development-capable phenomenon, with suffering as the most legible indicator of disruption and becoming as the positive expression of what life does when disruption is absent.

4.1b — The Maqasid al-Shariah: An Independent Derivation of the Floor/Becoming Structure

The framework's evaluative structure — a floor of necessary conditions below which no arrangement can be legitimate, and a horizon of adaptive response capacity beyond that floor — was developed from biological and materialist foundations. An independent derivation of structurally equivalent evaluative architecture exists in the Islamic jurisprudential tradition of Maqasid al-Shariah (the objectives of Islamic law), developed by Al-Ghazali in the twelfth century and systematized by Izz ibn Abd al-Salam in the thirteenth. This parallel is not merely interesting as intellectual history. It is analytically significant: if a structurally equivalent evaluative framework was derived independently from different foundational commitments across centuries and cultures, this constitutes evidence that the structure itself is tracking something about human social organization that is not reducible to the specific tradition that produced it. The Maqasid hierarchy identifies five protected domains whose preservation constitutes the basic requirement of any legitimate governance arrangement: nafs (life — the protection of biological existence and bodily integrity), ■ aql (intellect — the protection of cognitive capacity and access to knowledge), nasl (lineage and community — the conditions for biological and social reproduction), mal (property — the material conditions for independent action), and din (meaning/religion — the framework through which existence acquires coherence and purpose). These five are organized into daruriyyat (necessities — whose violation destroys the conditions for human life), hajiyyat (complementary needs — whose absence produces hardship but not destruction), and tahsiniyyat (embellishments — improvements to the quality of life above the threshold of adequate functioning). The daruriyyat/hajiyyat/tahsiniyyat hierarchy is a floor/sufficiency/becoming structure articulated in a different vocabulary from PMN's but with the same evaluative architecture: there is a minimum below which no arrangement can be legitimate, a middle range of adequate functioning, and a horizon of genuine improvement above that. The parallels are structurally specific. Nafs maps directly onto PMN's biological floor: the protection of life and bodily integrity is the precondition for everything else.

■Aql maps onto PMN’s becoming criterion: the protection of cognitive capacity is not merely a welfare concern but the condition for genuine development. Mal maps onto the resource deprivation component of structural suffering (R): material deprivation is identified as an institutional failure requiring correction. Din maps onto PMN’s analysis of meaning infrastructure in Part V: the framework of meaning through which existence acquires purpose is identified as a necessary condition for human functioning, not a luxury. The Maqasid tradition also independently identifies that these five domains have internal hierarchy — life precedes intellect, which precedes community,

which precedes property, which precedes meaning when they conflict — which is a claim about the ordering of floor conditions that PMN does not fully specify and should engage. The Maqasid contribution to PMN: the Islamic jurisprudential tradition provides an independent derivation of the floor/becoming evaluative structure from non-materialist foundations, which is evidence that the structure is not an artifact of materialist philosophy but a feature of how serious ethical traditions across different foundations converge on what human social organization requires. More specifically, the Maqasid hierarchy’s explicit ordering of protected domains when they conflict is an analytical resource PMN should draw on: the claim that life precedes intellect, which precedes community, which precedes property when genuine conflicts arise, is a substantive position on the ordering of floor conditions that PMN’s criterion implies but does not articulate. The Maq■id al-Shar■‘ah — the objectives of Islamic law — were systematically formulated by al-Ghaz■ (c. 1095) and developed by subsequent jurists including al-■i■ ibn ■ard-Dabb■gh and al-Sh■ibi■ (d. 1388). In their classical formulation, the Maq■id specify five necessary conditions that any legitimate legal arrangement must protect: life (nafs), intellect (■aql), lineage (nasl), wealth (m■l), and religion (d■n). Some later jurists, including Ibn ■a■sh■r, added dignity (kar■ma) as a sixth. The relevant structural feature is that the Maq■id generate a floor criterion: an arrangement fails not because it fails to maximize some positive good, but because it fails to protect conditions without which human existence as such is not possible. This is a threshold formulation, not a maximizing formulation — and it generates the same asymmetric structure as PMN’s minimal anchor. The convergence is structural, not textual. PMN does not derive its floor criterion from the Maq■id; it derives it from the biological analysis of unnecessary suffering in Part III. What the Maq■id provide is independent corroboration: a tradition beginning from jurisprudential reasoning about divine intent, operating across centuries in a completely different civilizational context, arrives at a structurally similar distinction between threshold conditions that cannot be violated without destroying the subject of value, and positive goods that can be pursued in multiple legitimate ways. Life (nafs) is an analogue of the biological floor. Intellect (■aql) is an analogue of the becoming criterion: the capacity for reasoning, deliberation, and self-governance that PMN identifies as the anti-foreclosure criterion. The Maq■id do not reduce to PMN’s framework — religion (d■n) has no PMN equivalent — but the convergence on the floor/anti-foreclosure structure is not coincidental. It reflects a recurring feature of sustained ethical reasoning under material constraints: the distinction between what must be protected and what can be pursued.

4.1c — The Failure of Divine Mandate as Universal Anchor

The objection that PMN’s evaluative anchor should be divine mandate rather than biological suffering takes several forms, and the framework must address each directly rather than by dismissal or deflection. The mandate selection problem. The most basic version of the objection — that suffering should be evaluated against what God has commanded rather than what organisms biologically avoid — fails at the level of application before it reaches the level of theology. Divine mandates, as they are actually held and practiced, are incompatible across traditions. Accepting one as the framework’s anchor is not a neutral philosophical move — it is a prior commitment to the tradition from which that mandate comes. A framework that begins from divine mandate has already selected its constituency before it begins its analysis. What the framework loses is exactly what makes it useful: the capacity to evaluate institutional arrangements whose participants

do not share that mandate, and to identify suffering that the mandate declares deserved without losing the ability to see that the suffering is real. The convergence argument. The claim that PMN's floor criterion is an artifact of materialist or atheist philosophy is directly challenged by 4.1b: the Maqāṣid tradition, beginning from the protection of divine purposes rather than from biological analysis, independently generates a floor/anti-foreclosure structure structurally identical to PMN's. This convergence is not coincidental. It reflects a recurring feature of sustained ethical reasoning under material constraints: traditions that take seriously the question of what human social arrangement must protect arrive at similar threshold structures, regardless of their metaphysical foundations. The biological analysis in Part III and the jurisprudential analysis in the Maqāṣid tradition arrive at the same architecture from different starting points. The architecture's recurrence across independent derivations is evidence that it tracks something about what human existence requires — not that one tradition has the correct theology. The deserved suffering objection. The most consequential version of the challenge is not about which anchor is correct but about who counts. The claim that suffering produced by violation of divine command is deserved, serves a divine purpose, or should not be included in the framework's analysis is a claim with real social force: it is used to exclude specific categories of suffering from moral consideration not by denying that the suffering is real but by declaring it legitimate. PMN's response is structural. The framework evaluates whether suffering is unnecessary — whether it is produced by institutional arrangements that could be otherwise — not whether it is theologically sanctioned. Theological sanction of suffering does not change its material reality, its behavioral consequences, or its systemic effects. A population experiencing chronic suffering that is declared deserved by

the dominant framework experiences the same degradation of trust, cooperation, and adaptive capacity as a population experiencing suffering that is declared unjust. The systemic mechanism operates regardless of the ideological label applied to it. The distinction between the mandate selection problem and the deserved suffering objection is important: the first concerns whether divine mandate can function as a universal anchor; the second concerns whether theologically-framed exclusions of specific suffering from consideration are analytically valid. The first fails at the level of application — selecting a mandate requires selecting a tradition. The second fails at the level of mechanism — theological declarations do not alter the material consequences of suffering. Both must be engaged directly rather than dismissed, because they are not only philosophical objections: they are the live arguments that determine which populations' suffering is considered analytically relevant in specific political contexts. → Relationship between theological frameworks and the evaluative criterion: Part V. Governing narratives that reframe structural suffering as individually deserved: 8.4c. Convergence between PMN's floor criterion and the Maqāṣid tradition: 4.1b.

4.2 — Better and Worse

The minimal anchor generates a threshold: arrangements that produce suffering the institutional alternative could prevent fail the criterion. But it does not by itself generate a ranking of arrangements that both pass the threshold. Two societies may both clear the biological floor while differing substantially in the degree to which they enable becoming — the expansion of cognitive, relational, and creative capacity that 4.5 develops. Section 4.2 specifies how PMN generates comparative evaluations above the floor: the criterion is not binary but directional. Better arrangements are those that preserve and expand adaptive response capacity across the population, reduce the structural conditions that foreclose it, and do so by means that do not restore floor-level violations in other populations or in future generations. The comparative criterion is distributional, not aggregate: an arrangement that expands adaptive response capacity for a majority while foreclosing it for a minority produces unnecessary suffering in the foreclosed population — a floor violation in the temporal dimension — and is not straightforwardly better than one that preserves anti-foreclosure across all. A clarification about the relationship between this criterion and the Epicurean tradition: both PMN and Epicureanism are grounded in the material conditions of embodied organisms, both treat unnecessary

suffering as something to be reduced, and both resist purely idealist ethics. The structural difference is decisive. Epicureanism is pleasure-positive: hedone (pleasure) is the goal, ataraxia (tranquility) is the achievement, and the good life is characterized by its positive content. PMN is suffering-negative: the floor criterion identifies what must be absent, and becoming identifies capacity expansion rather than positive experience. Suffering as a floor is not the mirror image of pleasure as a ceiling — it is a different evaluative structure. A person in conditions of adaptive response capacity may experience significant hardship; the criterion is not that their experience is pleasant but that their structural conditions enable the exercise of cognitive, relational, and creative capacity. The second divergence is political: Epicurus' practical conclusion was withdrawal (λ■θε βι■σας — live hidden), because structural suffering was not his object of analysis. PMN's practical conclusion is institutional contestation, because structural suffering is produced by institutional arrangements that can in principle be changed. Full moral relativism is incompatible with this framework — as it was, in practice, incompatible with Marx's, despite his methodological caution about moral philosophy. If reality has patterns, if suffering is real, and if those patterns can be evaluated by their consequences for the organisms who experience them, then some arrangements produce better outcomes than others. The judgment is comparative and contextual and falsifiable — if evidence shows otherwis

e, it must change — but it is a judgment, not a suspension of judgment. The difference between this position and the chauvinism that has historically attached itself to universalist frameworks is the falsifiability. Chauvinism says our arrangement is superior because it is ours. This framework says an arrangement produces better outcomes because those outcomes can be examined — and demands that the examination be honest enough to change the conclusion when the evidence requires it. The claim is empirical — and demanding, because it requires specifying the alternatives and evaluating the consequences honestly, including the consequences that are inconvenient for the conclusion you started with. In practice, making the comparison requires three specifications, each of which is a site where the comparison can be made dishonest. The criterion must be applied consistently — what does this arrangement produce in terms of structural suffering, and does it create or foreclose the conditions for genuine development? The alternative must be real: not an ideal arrangement that has never existed, but the actual institutional options available given current material conditions, with their track records. And the population must be specified — whose conditions are being evaluated, whose experience is included in the assessment? Arrangements that appear to produce better outcomes when evaluated from the position of those who benefit from them look different when the full distribution of outcomes is included. Each specification must be made explicitly, because leaving any one of them implicit is usually where the dishonesty enters. The Description-Evaluation-Intervention Tripartite: Three Questions That Must Not Be Conflated The evaluative framework requires a procedural discipline that is easily violated in both directions: the discipline of keeping three analytically distinct questions separate. The first question is descriptive: what is actually the case? What conditions, dispositions, mechanisms, or structural arrangements exist, with what properties and what causal relationships? The second question is evaluative: what is the normative status of those conditions or the behaviors they are associated with? What do they produce by the framework's evaluative criterion — do they expand becoming and reduce unnecessary suffering, or do they foreclose becoming and produce suffering that could be prevented? The third question is interventional: given the descriptive facts and the evaluative assessment, what institutional design produces the best outcomes by the evaluative criterion? These three questions require different evidence, different analytical procedures, and different conclusions — and the conflation of any pair produces characteristic errors. Conflating description with evaluation produces two mirror failures: the claim that accurate description of a phenomenon implies normative endorsement of it (which makes accurate description of stigmatized phenomena p

olitically impossible and analytically crippling), and the claim that negative evaluation of a phenomenon implies that accurate description of it should be suppressed or distorted (which converts normative commitment into epistemic failure). Conflating evaluation with intervention produces the error of treating the correct normative position on a behavior as sufficient to determine the correct institutional response — as if

"this behavior is wrong" automatically specifies "therefore the correct institution is enforcement." It does not: the intervention question requires empirical analysis of what institutional designs actually reduce the behavior, which may produce conclusions that the evaluative framing alone would not anticipate. The analysis of dispositional harm prevention in section 3.4c is a case where this conflation has historically been most costly: the correct evaluation (acting on harmful dispositions is wrong) has been treated as determining the correct intervention (enforcement), when the empirical evidence of second-order effects shows that enforcement-only designs produce more of the behavior they normatively condemn than dual-track designs do. The practical implication of the tripartite: before any policy position or doctrinal recommendation, the framework requires that all three questions be answered explicitly and separately. What is the descriptive situation? What is the evaluative assessment? What does the evidence show about what interventions produce better outcomes by that assessment? A position that can only answer the evaluative question — that can say "this is wrong" without also saying "this is what is actually happening" and "this is what the evidence shows reduces it" — is not a policy position. It is a moral posture. Moral postures are not without value, but they are not what PMN is designed to produce. The evaluative standard in operation: Better and worse are comparative and contextual in a specific sense: the evaluation compares actual arrangements against available alternatives under current conditions, not against ideal arrangements that are not producible under current conditions. An arrangement that reduces structural suffering relative to the available alternatives is better even if it produces substantial suffering relative to what would be possible under different material conditions. This contextual character of the evaluation prevents the framework from becoming utopian — from treating currently available arrangements as inadequate because they fall short of an imagined ideal — while preserving its critical force against arrangements that are worse than available alternatives. The critical question is always: given what is actually producible under current material and institutional conditions, is this arrangement better or worse than what could replace it? The empirical structure of the evaluation: the comparative claim — this arrangement is better than available alternatives — is not a philosophical claim

that can be settled by argument alone. It is an empirical claim about what alternatives are actually available and what they would actually produce. The most common error in comparative institutional analysis is treating alternatives as more available than they are — assuming that the institutional infrastructure, the material conditions, and the coalition configuration required to produce an alternative already exist or can be rapidly produced. The framework's evaluation of institutional alternatives must be grounded in realistic assessment of what the transition from current arrangements to proposed alternatives would actually require and what it would actually produce during the transition.

4.3 — On Progress

The concept of progress developed here has a temporal dimension that generates obligations that the framework has identified but not fully developed: obligations to future generations and to children as the population that most directly embodies the continuity between present arrangements and future conditions. Children are not merely the subjects of intergenerational transmission variables in the formula architecture — they are the population for whom the gap between the minimal anchor and adaptive response capacity is most consequential, because the structural conditions they encounter in their developmental years shape what they can become in ways that no subsequent arrangement can fully reverse. The framework's criterion applies to children with additional force: an arrangement that achieves the minimal anchor for adults while allowing children to be formed under conditions of chronic deprivation, developmental trauma, or foreclosed educational access has not achieved the floor — it has borrowed against the next generation's capacity for development. Future generations present a distinct problem. They cannot participate in current political processes that will determine the material conditions they inherit. This is not a merely philosophical problem: the ecological sustainability commitment (established in 4.5) and the intergenerational transmission variable G (15.4) both reflect the structural reality that arrangements which liquidate the conditions for future

becoming in the service of present sufficiency are arrangements that fail the framework's criterion even if they satisfy it for current populations. The institutional implication: future generations are legitimate subjects of the minimal anchor, not future beneficiaries to be considered when convenient. Arrangements that impose irreversible costs on future populations — ecological, fiscal, infrastructural — without the consent of those populations require the same justification the framework requires for any arrangement that imposes costs without consent: that the costs are unavoidable given available alternatives, that they produce compensating benefits for those bearing them, and that the institutional mechanisms for contesting the arrangement exist even for those not yet present to exercise them. Children are present to exercise them. Their absence from most political institutions is itself an institutional failure that the framework's analysis of who bears costs and who controls the arrangements producing those costs directly implicates. The concept of progress in this framework is deliberately stripped of the teleological content that made it so powerful and so dangerous in the Hegelian and Marxist traditions. History does not move toward a predetermined end. The development of human social organization is not guaranteed to produce better outcomes by the internal logic of any process. Progress — in the sense of arrangements that expand the structural conditions under which genuine

development is possible, with the reduction of structural suffering as its primary diagnostic — is possible, and it has occurred, and it can be sustained and extended. But it is not guaranteed, and its reversal is always possible. The insistence that better and worse arrangements can be distinguished, that the distinction is not merely subjective, and that working toward better arrangements is a legitimate activity — this much from the progressive tradition remains. What gets discarded is the confidence that the movement of history is on the side of that work, the confidence that made Hegel's World Spirit and Marx's proletariat so compelling as ideas and so catastrophic as political programs when the historical process refused to behave as promised. Progress is possible. History does not guarantee it. Progress in this sense has happened. The abolition of formal slavery across most of the world is a real change in structural conditions — not completed, not uniform, not without regressions, but real. The reduction of child mortality through vaccination and sanitation is a real change in the material conditions of biological vulnerability. The extension of formal political participation to populations previously excluded from it is a real institutional change, with real though incomplete material consequences. These are not ideological claims. They are documentable shifts in the conditions that the minimal anchor identifies as the primary object of analysis. But each of these changes also illustrates the reversibility that the framework insists upon. Formal abolition of slavery coexisted with systems of debt bondage, prison labor, and wage arrangements that reproduced many of the material conditions of slavery under different names. Reduced child mortality in one era has been reversed in specific populations by austerity, conflict, and health system collapse. Extended political participation has been systematically curtailed when the populations it empowered began to use it in ways that threatened concentrated economic interests. Progress is real. It is not secure. And the history of its reversal is as instructive as the history of its achievement — because it shows that the material conditions that made progress possible require active maintenance, not passive inheritance. The Convergence-Fragmentation Cycle: Why Unity Does Not Accumulate The reversibility of progress described above is not random. It follows a pattern that the framework identifies as structural rather than contingent: the achievement of convergence — ideological, religious, cultural, institutional — systematically produces the conditions for its own fragmentation. This is not pessimism about human nature or a counsel of despair. It is an analytical claim about the mechanism through which achieved unity generates the pressures that eventually dissolve it, and recognizing the mechanism is what makes it possible to work with rather than against it. The mechanism operates at multiple

levels. At the ideological level: when a framework achieves sufficient dominance to define the terms of debate, it simultaneously defines what counts as deviation — it produces the category of heresy or error that gives the internal opposition a name and a coherence it would not otherwise have had. The Council of Nicaea did not resolve the theological disputes of early Christianity; it formalized them, gave the losing positions an identity, and created the institutional conditions under which schism became an available strategy. Every

orthodoxy generates its heterodoxy not despite its success but through it. At the institutional level: Ibn Khaldun's analysis of *asabiyyah* — the group solidarity that builds dynasties — demonstrates the same structure. The cohesion that allows a group to seize power is eroded by the success of that seizure: the material rewards of power attract defection, the urgency that bound the group dissolves as the threat recedes, and the next generation inherits position without having forged the solidarity that produced it. The dynasty that seemed permanent was consuming the conditions of its own permanence from the moment of its foundation. (Ibn Khaldun 1377/1958) At the cultural level, the same structure appears in Gramsci's analysis of hegemony: a ruling bloc achieves cultural dominance not by force alone but by incorporating enough of the aspirations of subordinate groups to make its leadership appear legitimate. But that incorporation creates internal tensions — the aspirations that were absorbed do not disappear, and the gap between what the hegemonic arrangement promises and what it materially delivers generates the counter-hegemonic pressure that eventually challenges it. Successful hegemony creates its own opposition by making the terms of contestation available within the dominant framework itself. The language in which the challenge is articulated is the language the hegemony taught. (Gramsci 1971) The practical implication for political analysis is that the question is never whether convergence will eventually fragment — it will — but at what rate, through what mechanisms, and whether the institutions that manage the transition between configurations have been built in advance. A movement that achieves significant convergence around its program and treats that convergence as the destination has misread the situation. The convergence is a condition for the next phase of contestation, not the resolution of contestation. The factions that will challenge the achieved unity are already forming within it, organized around exactly the tensions that the unity required suppressing or deferring in order to hold together. This is not a reason to avoid building coalitions or achieving coordination — it is a reason to build institutions that can manage the fragmentation that the coordination will eventually produce, rather than being surprised by it when it arrives. The framework's position on historical directionality follo

ws from this: history does not converge toward a stable endpoint. It cycles through configurations of convergence and fragmentation, each phase producing the structural conditions for the next, without any guarantee that successive configurations represent improvement over preceding ones. Progress — the genuine improvement of structural conditions for becoming and reduction of unnecessary suffering — is achievable within this cycle. It is not the destination of the cycle. Treating any achieved configuration, however good, as the terminal point of the process is the error that makes movements brittle when the fragmentation arrives that their success has already prepared. Structural Configurations and Their Developmental Horizons: A Diagnostic Typology Marx's typology of modes of production — communal, slave, feudal, capitalist, socialist, communist — was an attempt to answer a question that any materialist theory of history must answer: given a specific configuration of material conditions, what transformations are structurally available, what pressures drive movement from one configuration to another, and what direction does that movement tend? The typology failed as a predictive scheme because the sequence was teleological — history was supposed to move through the stages in order, and the failures of that sequence to materialize required increasingly elaborate theoretical rescue operations. But the question the typology was trying to answer remains indispensable. Without some account of how different configurations of material conditions relate to different developmental possibilities, the framework has a method for analyzing any specific situation but no capacity to orient the analysis toward transformation. PMN requires a diagnostic typology of its own — not a sequence of stages that history must traverse, but a map of configuration types that identifies, for each type, what structural conditions determine the developmental horizon available, what the dominant transformation pressures are, and what directions from the current configuration are realistically available given the constraint structure. The typology is diagnostic rather than teleological: it describes the landscape from where things are, not the destination history is moving toward. Progress within this typology is possible but not guaranteed. Regression is possible. Movement can be non-linear. The typology is a tool for orientation, not a prediction. The diagnostic dimensions that define configuration types

in PMN's typology are derived from the layered constraint model: the biological floor (are basic conditions of physical survival, cognitive functioning, and affective stability met for the population?), the economic layer (does the arrangement of production, distribution, and resource access permit accumulation for broad development or concentrate it in ways that foreclose becoming for the majority?), the institutional layer (do the arrangements for managing collective decisio

ns and enforcing coordination have legitimate accountability mechanisms or are they captured by narrow interests?), the epistemic layer (does the infrastructure for shared knowledge, collective sense-making, and coordination across communities function or is it degraded and captured?), and the meaning layer (does the population have access to meaning infrastructure that orients sacrifice, solidarity, and long-term commitment, or has that infrastructure collapsed without replacement?). Each dimension admits of a range, and the combination of positions on each dimension defines the structural configuration. Four broad configuration types are analytically distinguishable, not as historical stages but as structural positions that actual societies approximate in varying degrees and combinations. The first is the subsistence configuration: biological floor is barely met or frequently breached, economic arrangements produce minimal surplus, institutional arrangements are primarily oriented toward physical security rather than development, meaning infrastructure is provided by traditional-religious frameworks that are stable but constrain individual adaptive response capacity within collective structures. The developmental horizon available from this configuration is narrow — the primary transformation pressure is biological (food security, health, physical safety), and the direction realistically available is toward economic surplus generation without which the other layers cannot develop. Institutional and meaning-layer transformations attempted before the biological and economic floor is reliably met typically fail to sustain themselves against the pressure of unmet biological needs. The second is the captured development configuration: economic surplus is generated and growing, but accumulation is structured to concentrate in ways that produce institutional capture — the governing institutions that should manage the distribution of development capacity are controlled by the beneficiaries of concentration and operate to reproduce it. Biological floor is met for most of the population, but the developmental horizon is structurally foreclosed for the majority by the capture dynamics rather than by absolute resource scarcity. The dominant transformation pressure is political — the primary driver of change is the gap between formal institutional legitimacy and the actual capture of those institutions by narrow interests, which erodes the performance legitimacy that sustains consent. The direction realistically available involves either the rebuilding of accountable institutional capacity (the path toward the third configuration) or the entrenchment of capture into authoritarian stability (regression toward a form of the fourth). The third is the contested development configuration: economic surplus is sufficient, biological floor is reliably met, institutional capture has not been fully consolidated, epistemic infrastructure functions imperfectly but is n

ot fully degraded, and the population retains meaningful capacity for collective action around transformation pressure. This is the configuration in which the most contested political questions arise because genuine alternatives are structurally available — the constraint structure does not determine which of several paths will be taken, and the outcome depends on the balance of counter-power, adaptive capacity, and the quality of institutional design. Most of the societies within which the framework's doctrine is directly relevant occupy this configuration type. The developmental horizon is real — adaptive response capacity is structurally available for a significantly wider share of the population than currently accesses it — but achieving it requires transformation of the institutional and epistemic layers, not only the economic layer. The fourth is the meaning-collapse configuration, which can emerge from any of the others: biological floor may be met, economic surplus may exist, institutional capacity may be formally intact, but the meaning infrastructure that orients long-term commitment, solidarity, and sacrifice has degraded without replacement. This configuration is not primarily analyzed in Marx because his typology did not include the meaning layer as a structural variable. But the disenchantment trajectory analyzed in 5.5 — the long process by which the world was emptied of immanent sacred presence, producing populations whose buffered selves have lost access to the sources of orientation that made collective action possible at scale — is a structural condition

with its own developmental horizon. From the meaning-collapse configuration, the transformation pressures are not primarily economic: they are the pressures generated by populations seeking orientation, which are available for mobilization by whoever offers the most compelling framework fastest, including frameworks that produce regression in all other dimensions. The typology does not predict which configuration any society will move toward. It identifies what transformation pressures are dominant in each, what the realistic range of movement is given the constraint structure, and what interventions at which layers are most likely to shift the configuration in the direction the evaluative criterion requires. This is what a materialist theory of history needs that is not teleological: not a direction history must travel, but a map of the landscape history actually moves through, with honest accounts of what is available from each position and what is not. The typology developed here is diagnostic rather than merely descriptive. Each configuration type identifies not only where a society stands but what the constraint structure of that position implies for doctrinal application. Section 10.6 develops the sequencing implications in full: doctrines derived from PMN are not uniformly applicable across all configuration types, and a doctrine that correctly addresses the constraint s

tructure of one configuration may be missequenced for another. The typology is the map; the sequencing analysis is the navigation. Neither is useful without the other. A Fifth Configuration: Ecological Collapse and the Threatened Floor The four structural configurations developed above — subsistence, captured development, contested development, and meaning-collapse — were derived from analysis of how constraints at different layers interact to shape the range of becoming available to a population. A fifth configuration has become analytically necessary as planetary boundary violations have crossed from marginal to significant: ecological collapse, in which the biological floor itself is under systematic and accelerating threat not from distributional failure but from the degradation of the Earth systems that produce the conditions for floor provision. This configuration differs from the others in a specific and structurally important way: the primary constraint is not institutional failure at the economic, epistemic, or meaning layer, but the progressive degradation of the physical and biological substrate on which all higher-layer functioning depends. The ecological collapse configuration can emerge from any of the other four. A contested development configuration does not become an ecological collapse configuration merely because environmental degradation is occurring within it — environmental degradation has occurred within every configuration in human history. The threshold is crossed when the rate and scale of degradation begins to undermine the reliability of floor provision itself: when climate instability produces agricultural failures that cannot be offset by institutional redistribution, when freshwater scarcity becomes acute enough to threaten biological survival across significant populations, when biodiversity loss reduces the resilience of food systems to a point where shocks produce floor violations that no institutional response can prevent in real time. The distinctive feature of the ecological collapse configuration is that the primary lever for restoring floor reliability is not institutional reform but planetary system stabilization — which operates on timescales and at scales of coordination that exceed what any single political configuration can produce. The strategic implications of the ecological collapse configuration are distinct from all other configurations in one critical respect: the window for effective intervention is time-bounded in a way that the other configurations are not. In a captured development configuration, the harm of capture continues indefinitely until structural transformation occurs, but the capacity for transformation does not inherently diminish over time — a population under institutional capture retains its capacity for collective action, and transformation remains possible at any point. In an ecological collapse configuration, the biophysical processes that are degrading floor conditions are su

bject to feedback dynamics and tipping points that, once crossed, are not reversible on any timescale relevant to human political action. The strategic implication is not panic but accurate urgency: the window for institutional transformation that prevents or limits floor degradation is genuinely finite in a way that windows for other structural transformations are not, and analysis that treats ecological transformation as one priority among others to be sequenced by political convenience is systematically miscalibrated to the actual constraint structure. The ecological collapse configuration diagnostic: a configuration has entered ecological

collapse territory when (1) planetary boundary violations are producing measurable degradation of the physical conditions for floor provision — agricultural yields, freshwater availability, climate stability for food systems and human habitation; (2) the rate of degradation exceeds the adaptive capacity of current institutional arrangements; (3) the processes producing degradation contain feedback dynamics or proximity to tipping points that make the degradation self-reinforcing once thresholds are crossed. In this configuration, the primary strategic priority is intervention in the physical processes of degradation — energy system transformation, land use change, biodiversity protection — not because institutional transformation is irrelevant but because institutional transformation that does not address the physical processes is addressing the wrong constraint level. The institutional capture that blocks ecological transformation is a secondary constraint; it must be addressed, but it must be addressed as an obstacle to planetary stabilization, not as the primary target in its own right.

4.4 — On Manipulation, Trust, and Structural Checks

The materialist tradition recognized that political actors are not transparent to themselves or to others — that stated motivations and actual material interests diverge in ways that are systematic and predictable. This recognition can be used for emancipatory purposes: helping people understand the material interests that their ideological commitments serve, and thereby making it possible for them to evaluate those commitments more honestly. It can also be used for manipulative purposes: using the understanding of how people actually process reality to influence them in ways they would not endorse if they understood what was happening. The capacity for this kind of influence — the ability to model others' psychology accurately without necessarily sharing their emotional responses — amplifies whatever intent it serves. The checks against its misuse cannot rely on the feelings of the actor, because the feelings that normally constrain this kind of influence may not be present. They must be explicit structural commitments: specific principles about when and how this capacity may be used, articulated in advance and not overridable by the in-the-moment calculation that this particular case is exceptional. The structural commitments that function as checks against the misuse of this capacity are not arbitrary rules. They follow from the framework's own analysis of how power asymmetry and information control operate. The transparency commitment. Any use of the capacity to model and influence others' reasoning must be visible to the actor using it — explicitly acknowledged as an exercise of this capacity rather than experienced as mere communication. The failure mode is self-deception: the actor who believes they are communicating honestly while actually deploying influence techniques loses the capacity to apply the checks that honest actors can apply, because the checks require recognizing what you are doing. The first structural requirement is therefore internal transparency: the actor must be clear with themselves about when they are exercising this capacity and when they are not. The preservation-of-capacity test. Communication that preserves the recipient's capacity to evaluate and reject what is being communicated is legitimate, even when it is strategic. Communication that works by degrading that capacity — by exploiting cognitive biases, by creating emotional states that prevent evaluation, by narrowing the perceived range of alternatives — is manipulation, regardless of whether the content is true or the intentions are good. The test is not whether you believe you are right. It is whether, if the recipient could see exactly what you are doing and why, they would be able to make a genuine choice about whether to accept it. This test can be applied in advance and functions as a practical check on the drift toward manipulation that the capacity for modeling others makes available. The asymmetry prohibition. This capacity must not be used to create

or deepen power asymmetries that the framework's own analysis would require addressing. Using superior psychological modeling capacity to extract cooperation from actors who could not consent to the exchange if they understood it is structurally identical to the information manipulation the framework analyzes in institutions that use it to maintain arrangements the affected populations would not support. Criticizing that use by institutions while permitting it by the actors this analysis informs would be a self-contradiction. The

prior commitment mechanism. Because the in-the-moment calculation that 'this particular case is exceptional' is always available and almost always self-serving, the commitments about when and how this capacity may be used must be articulated in advance and held against revision by in-the-moment reasoning. This is the structural analog of constitutional constraint: the value of a prior commitment is precisely that it is not revisable by the reasoning of the actor who might benefit from revising it. Prior commitments must be articulated specifically enough to generate genuine constraints — not 'we will use influence techniques only when necessary,' which generates no constraint at all, but specific criteria about what necessity means and who judges whether it has been met. These four structural commitments operate at the level of the individual actor. They are necessary but not sufficient. The framework's own analysis of why internal transparency fails — actors are not fully transparent to themselves, self-deception is systematic, and the capacity being constrained is precisely the capacity for modeling that makes self-deception more sophisticated — implies that external constraints are required as complement, not substitute. Institutional transparency mechanisms, adversarial review by actors with independent standing, and public accountability structures are not alternatives to the four commitments above. They are what makes those commitments binding when the actor's own internal vigilance is insufficient. The full architecture of external constraints is developed in Parts VI and VII, where the custodian problem and accountability mechanisms are analyzed as structural features of any arrangement that must prevent capture by those it empowers. The structural check that is available: The framework's answer to the manipulation problem is not a prohibition on strategic communication — which would be unenforceable and would simply advantage actors with fewer scruples. It is a structural check on the conditions under which strategic communication is legitimate: the check of accountability to the populations whose understanding is being managed. Strategic communication that serves the interests of those whose understanding it manages is different in kind from strategic communication that serves the interests of those doing the managing at the expense of those managed. The difference is not always visible at the moment of the co

munication, which is why the structural check requires institutional forms — transparency mechanisms, accountability relationships, independent verification — rather than relying on the good intentions of individual communicators. The framework's rejection of teleology requires a precision that 4.3 has not yet made explicit. Two forms of teleology are distinct and require different treatment. Strong teleology is the claim that history has a direction inscribed in its structure — that specific outcomes are inevitable, that the arc bends toward justice as a structural feature rather than a political achievement, and that this guarantee licenses specific acts including the sacrifice of floor conditions for assured future anti-foreclosure outcomes. This form is rejected by PMN's probabilistic determinism (1.3) and its anti-linear account of historical change (4.3). It is self-sealing: failure can always be explained as delay rather than disconfirmation, and the framework produces no falsifiable predictions. Weak teleology — what might be called as-if teleology — is the adoption of a horizon as an orienting assumption that guides action without claiming historical guarantee. An analyst or movement that acts as if adaptive response capacity for all people is achievable, while maintaining explicit revision mechanisms and without sacrificing present floor conditions, is using a horizon instrumentally rather than claiming it cosmologically. This form is not rejected by PMN and is in fact required for any sustained political project: actors who genuinely believe that no progress is possible tend toward strategic paralysis, while actors who treat their horizon as guaranteed tend toward justifying present harm. The as-if orientation — working toward a horizon one holds open to disconfirmation — is the position the framework's epistemological commitments actually support. Diagnostic for strong vs as-if teleology: (1) Self-sealing test — is failure explained as "not yet time" or "strategy needs revision"? The first is strong teleology operating; the second is as-if teleology with intact revision mechanisms. (2) Sacrifice test — is present floor suffering being justified by reference to guaranteed future anti-foreclosure? If yes, strong teleology is being used regardless of stated epistemic commitments. (3) Accountability test — are there specific, enforceable commitments to revise the strategy if the horizon does not materialize on a specified timeline? As-if teleology requires these commitments to remain genuinely revisable. (4) Convergence test —

are the populations bearing the costs of the strategy the same populations who would benefit from the anti-foreclosure criterion's fulfillment? Strong teleology characteristically sacrifices floor conditions for one population while directing anti-foreclosure benefits to another. A movement or analysis that passes tests 1, 3, and 4 and fails only test 2 partially is using as-if teleology in its defensible form. A m

ovement that fails tests 1 and 2 while claiming epistemic humility has a strong teleological structure regardless of its self-description. The self-application requirement: the framework must apply this analysis to its own communication. Where the framework simplifies — presenting conclusions with more confidence than the evidence strictly supports, or presenting analytical positions in forms calibrated to persuasion rather than to the full complexity of the evidence — it must account for this as a communication choice with consequences that the analysis of manipulation identifies. The defense of such simplification is strategic effectiveness; the cost is the reduction of the audience's capacity to evaluate the framework on its merits rather than on its presentation. The framework's position is that this cost is manageable when the simplification is acknowledged — when the full complexity is available to those who seek it — and becomes a form of manipulation when the simplification is the only form in which the framework is offered.

4.5 — Beyond the Minimal Anchor: The Anti-Foreclosure Criterion

The minimal anchor — the commitment to reducing structural suffering — is the floor of the evaluative framework, not the full extent of what it requires. A framework that only specifies what is unacceptable and not what the anti-foreclosure criterion additionally requires will consistently be read as a sophisticated form of welfare management: the optimization of comfort, the elimination of obvious pain, the achievement of stable mediocrity. That reading misses what the framework is actually about. What Nietzsche feared is taken seriously here. The last man is not a distant possibility — he is the default outcome of systems that succeed at eliminating structural suffering without creating the conditions for genuine development. Last man does not suffer conspicuously. He is comfortable, equal, and finished becoming. He has warmth, security, and the daily satisfactions of a system that has learned to manage him well. What he has lost — and what the system that manages him has no category for — is the capacity to become something more than what he is. He has traded becoming for comfort, and the trade felt like progress. Before the departure from Nietzsche, the acknowledgment he deserves: systems that eliminate structural suffering can produce spiritual stagnation. These arrangements exist. Bureaucratic welfare arrangements that solve the material problem while creating conditions of managed dependency — in which recipients are administered rather than developed, in which the removal of suffering produces passivity rather than capacity, in which institutions designed to create conditions for becoming instead create conditions for comfortable stasis — these arrangements exist and produce something close to what Nietzsche feared. The fear is not irrational. It is a genuine failure mode of any politics organized around the reduction of suffering without adequate account of what comes after. The departure from Nietzsche is about causes, not aspirations. Nietzsche attributed spiritual flatness to egalitarianism itself — to the leveling impulse, the resentment of excellence, the herd's revenge on everything that rises above it. This framework's diagnosis is different. Spiritual flatness is produced by two distinct structural arrangements. The first keeps most people permanently below the threshold at which development becomes possible, exhausting human capacity for becoming through deprivation. The second offers managed sufficiency without genuine development — removing suffering while foreclosing the becoming that sufficiency was supposed to make possible. Both produce the last man: the first through deprivation, the second through dependency without development. Neither is produced by egalitarianism or by the reduction of structural suffering. Both are produced by specific institutional arrangements that can be identified, analyzed, and changed. Both failure modes are avoided through the distinction between floor and anti-foreclosure criterion. The minim

al anchor is necessary but not sufficient. An arrangement that achieves the floor and stops there has prevented the worst outcome but has not created the conditions for adaptive response capacity. The

evaluative criterion extends beyond the floor: do these arrangements create the conditions under which human beings can develop — cognitively, morally, creatively — or do they stabilize people at a level of sufficiency that forecloses the becoming that sufficiency was supposed to make possible? That question cannot be answered by measuring suffering alone. It requires analyzing what the arrangements make available beyond survival. The aspirational horizon beyond the minimal anchor is therefore not comfort but capacity: the conditions under which human beings can genuinely develop — cognitively, morally, creatively, in whatever directions the specifics of each life and each community make possible. Not equality of outcome. Not the same trajectory for everyone. But genuine equality of capacity for becoming — the removal of the structural barriers that determine, before a life has properly begun, who gets to develop and who gets to survive. This is what makes the position Nietzschean in its aspirations while fundamentally anti-Nietzschean in its analysis of causes. Nietzsche wanted genuine greatness, adaptive response capacity, genuine differentiation of human capacity. So does this framework. Where they part is on the question of what prevents it. Nietzsche blamed the resentful masses and the leveling moralities they produced. This framework identifies the structural conditions — and the specific arrangements maintained by specific actors — that prevent the greatness Nietzsche wanted from being available to more than a narrow class of people born into the right position. The *Übermensch* that this framework aspires toward is not the exceptional individual who rises by standing on others. It is the human being — any human being, from any position — who has been given the material conditions under which genuine development becomes possible, and who uses those conditions to become something that could not have been predicted from where they started. That kind of becoming requires the minimal anchor as prerequisite. But it is not satisfied by it. The floor is not the destination. Nor is the floor self-sustaining: adaptive response capacity across generations requires the ecological conditions under which it is possible. An arrangement that achieves the minimal anchor for the present generation by consuming the ecological substrate that makes becoming possible for future generations has not secured the floor — it has borrowed against it. Ecological sustainability is not a separate value added to the framework from outside; it is what the aspiration toward adaptive response capacity requires when followed to its temporal conclusion. The evaluative criterion that distinguishes institutions that serve becoming from those that do not is not whether they produce

comfort, but whether they expand the structural availability of genuine development — and this distinction generates real trade-offs that the framework does not resolve in advance. Consider education: a system organized around standardized credential production may successfully reduce suffering (stable employment, predictable income) while systematically foreclosing the cognitive and creative development that adaptive response capacity requires. A system organized around the development of critical capacity and original thought may serve becoming at the cost of economic predictability, creating a different kind of structural vulnerability for the populations it serves. Neither trade-off is costless; the question the evaluative criterion asks is not which system is comfortable but which creates the conditions under which more people can develop further. The same tension runs through public culture: institutions that subsidize artistic production for broad accessibility serve becoming differently than institutions that protect the conditions for aesthetic risk-taking and formal experimentation, and the two goals cannot always be simultaneously maximized. These are not failures of the criterion. They are its characteristic output — the specification of genuine trade-offs that must be navigated honestly rather than dissolved by committing in advance to one value over others. And the framework itself must embody this principle. The aspiration toward adaptive response capacity applies to the framework as much as to the individuals it analyzes: a framework that forecloses its own revision has adopted exactly the foreclosure it identifies as a failure mode elsewhere. A framework that cannot be surpassed by something better has stopped serving its purpose and started serving itself. This document is written in full expectation of that surpassing — as a contribution to a movement that will produce better analysis than this, and as an honest acknowledgment that producing that better analysis is the only measure of success that matters. (Nietzsche 1883-1885; Marx 1844) The distinction between Nietzsche's aristocratic individualism and the collective aspiration of this framework is not a softening of the

ambition. It is a radicalization of it. Nietzsche wanted greatness for a few exceptional individuals. This framework wants the conditions for preserved adaptive response capacity to be structurally available — which is a far more demanding aspiration, because it requires not only the cultivation of exceptional capacity where it already exists but the transformation of the arrangements that prevent it from developing everywhere else. A question that Part IV does not resolve: when legitimate values conflict — when freedom and security pull against each other, when equality and excellence cannot both be maximized, when truth and cohesion point in different directions — what does the framework prescribe? The framework does not prescribe a universal decision rule

for value conflicts, and the absence of such a rule is not a gap but a commitment. A framework that resolves all value conflicts in advance by establishing a fixed hierarchy of values has abandoned the epistemological stance of Part I: it has substituted prior commitment for analysis. What the framework provides instead is a typology of conflicts and tools for managing them honestly rather than dissolving them falsely. That typology is developed in Part XIII, which should be read as the continuation of the ethical account begun here. Part XIII distinguishes between permanent tensions — conflicts that are structural features of any complex social arrangement and cannot be resolved, only managed — and conditional tensions, which are present under specific historical conditions and can in principle be reduced as those conditions change. The distinction matters because the appropriate response to a permanent tension is management, not solution; while the appropriate response to a conditional tension is the identification of the conditions that generate it and the material changes that would reduce it. Confusing the two is among the most costly analytical errors available to a framework of this kind. Recognition as a Condition for Becoming The structural conditions for becoming are not only material in the sense of resources, institutions, and freedom from deprivation. They include a dimension that the materialist tradition has systematically underanalyzed: the conditions of recognition — the social relationships through which persons come to understand themselves as having the standing to develop, to claim, and to become. The recognition tradition in social philosophy, developed most systematically by Axel Honneth and Nancy Fraser, identifies misrecognition — the systematic denial of recognition in its forms of love, respect, and esteem — as a structural harm that operates independently of, and in addition to, material deprivation. (Honneth 1995; Fraser 2016) The argument from recognition is not merely additive — it is not the claim that in addition to material conditions, recognition also matters. It is the claim that the capacity for self-directed development requires, as a structural precondition, a social environment in which the developing person is treated as someone whose development is legitimate and whose claims on the conditions for that development have standing. A person who has adequate material resources but is systematically treated as someone whose aspirations are inappropriate, whose identity disqualifies them from certain forms of development, or whose suffering does not count as suffering that requires a social response — that person does not have the structural conditions for becoming even if the material floor has been met. Misrecognition is a structural barrier to becoming, not merely a psychological injury. The framework's account of visibility (1.12, 8.4c) already implies this argument: structural arrangements that make certain

forms of suffering invisible, that systematically prevent the recognition of some populations' claims as claims, are arrangements that foreclose becoming through mechanisms that are not reducible to material deprivation. The connection to recognition theory makes this implication explicit. Meritocratic arrangements that deny the structural production of their outcomes do not only misallocate resources. They deny recognition — they tell the people who fail within them that their failure is their own, that the development they could not achieve was foreclosed by their own inadequacy rather than by the conditions they inhabited. That denial is a harm to the capacity for becoming that operates through a different mechanism than material deprivation but is no less structural in its causes and consequences. The Ungrounded Seeker: The Other Failure of Becoming The last man represents one failure mode of becoming: the person whose structural conditions have been sufficiently managed that the question of development never arises with urgency, who has settled into the comfort of maintained sufficiency and lost the capacity for the confrontation that adaptive

response capacity requires. There is a mirror image failure that the framework must also name: the ungrounded seeker — the person who has seen through every available framework of meaning without building the capacity to construct a new one, who is permanently between orientations, whose critical awareness has outrun their constructive capacity. Nietzsche announced the death of God as a catastrophe as much as a liberation: 'God is dead — and we have killed him. How shall we comfort ourselves, the murderers of all murderers?' The Zarathustra who announces the death of God does not stop there. He demands the creation of new values — the *Übermensch* not as a type of person but as a project of self-overcoming, the active construction of meaning in the absence of given meaning. What the seeker has absorbed is the first part of this — the critique, the negation, the refusal of every framework that presents itself as final. What the seeker has not accomplished is the second part: the creation. The permanent negation becomes, over time, an identity. Being against everything available substitutes for being for something that has not yet been built. The structural condition that produces the seeker at scale is precisely the disenchantment trajectory described in 5.5: the long process by which the world was emptied of immanent sacred presence, first by the transcendentalization of the divine within religious traditions themselves, then by the secular challenge to the transcendent framework that remained. What the seeker is often searching for — though frequently without the conceptual vocabulary to name it — is not a new transcendent God to replace the old one. It is re-enchantment: the recovery of a relationship to the material world in which meaning is not a private projection onto a neutral substrate but some

thing encountered, something that the world itself offers. The attraction of neo-paganism, of deep ecology, of animist-inflected spiritualities, of psychedelic experience as a path to direct encounter — these are not regressions to pre-modern credulity. They are attempts to recover the immanent dimension of meaning that the long disenchantment process removed. For the framework's political analysis, the seeker presents a specific problem distinct from both the last man and the ordinary participant in meaning infrastructure collapse. The last man is politically passive because comfort displaces urgency. The person experiencing meaning infrastructure collapse is potentially mobilizable — they need orientation and will respond to whoever offers it most compellingly. The seeker is neither: they are intellectually sophisticated enough to see through the mobilizing frames that capture the disoriented, which makes them resistant to manipulation but also resistant to the coalition-building that political transformation requires. Within movements, seekers function as the sharpest critics of every proposed direction, identifying the compromises and corruptions in each available option with accuracy — but without the capacity to commit to any option that falls short of a standard no real option can meet. Their critique is often correct. Its political consequence is frequently paralysis and fragmentation. The anarchist misreading also has a Nietzschean variant: the claim that the framework's agent ecology, with its distributed coordination and functional differentiation, is a theoretical rehabilitation of vanguardism under less recognizable names. The response requires distinguishing organizational capacity from organizational authority. The framework does not prescribe centralized forms or assign interpretive authority over the doctrine to any group. What it does is identify that the coordination problems at scale that any transformative project faces cannot be resolved by commitment to horizontal principles alone — because those problems are structural features of large-scale collective action, not products of ideological failure. The custodian problem (7.3) applies with equal force to any actor claiming organizational authority, including actors who claim it in the name of horizontalism. The question is not whether to build organizational capacity but what forms that capacity takes and what accountability mechanisms prevent it from becoming the problem it was built to address. The framework's position on the seeker is not dismissal but diagnosis. The condition they inhabit — the terminus of a long disenchantment, in possession of critical awareness without constructive orientation — is a genuine structural position, not a personal failure. What it requires is not the suppression of the critical capacity but its integration with a constructive project that does not require certainty before commitment. The Gramscian formulation is useful here: pessimism of

the intellect, optimism of the will. The seeker has the first in abundance. What is missing is not better arguments for optimism — the arguments will never be sufficient, and the seeker knows this better than most. What is missing is the willingness to build under conditions of irreducible uncertainty, to commit to a direction without the guarantee of arrival, to treat the absence of a final answer as the condition of the work rather than as a reason not to begin it.

4.5b — The Loop Model and the Anti-Foreclosure Criterion

How Becoming Emerges: The Sentient-Material Response Loop Becoming is not a teleological destination. It is the observable pattern that emerges when sentient systems respond repeatedly to material constraints. Suffering — or more precisely, the conditions that produce it: deprivation, damage, vulnerability, dependency, limitation — triggers responses. Responses reduce some suffering while generating new configurations of constraint. Successful response patterns get encoded: in habit, in institutional procedure, in cultural knowledge, in built infrastructure. Encoded patterns reduce the cost of future responses to similar constraints. They also create new vulnerabilities and new dependencies. The loop does not terminate at a destination. It continues as long as sentient systems exist in material conditions that produce constraint. The food example illustrates the structure without romanticizing it. Hunger produces suffering. Suffering triggers response: foraging, hunting, agriculture, food preservation, trade networks, cuisine, industrial production. Each development reduces some suffering while producing new configurations: agricultural communities face drought vulnerability that foragers do not; industrial food systems produce obesity and food deserts that subsistence communities do not. None of this is automatically good. The loop does not have a moral direction built into it. What it has is a causal structure: sentient systems embedded in material conditions develop response patterns that expand their repertoire, and that expanded repertoire is what, from the outside, looks like becoming. The same structure applies to learning, relational development, political organization, technological development, and the accumulation of institutional knowledge. None of these is automatically good. All of them emerge from the same sentient-material response dynamics. The loop model has a consequence that earlier formulations did not adequately acknowledge: becoming can produce suffering. Expanded cognitive capacity generates the capacity to suffer from abstract causes — injustice, finitude, meaninglessness — that simpler organisms do not experience. Deeper relational attachment generates the capacity for grief, betrayal, and loss at depths not available to beings with shallower relational lives. Greater social complexity generates new forms of structural suffering: exclusion from complex institutions, inequality indexed to institutional position, status-based harm that has no equivalent in simpler arrangements. Becoming produces suffering as a structural feature of the loop, not as an anomaly within it. This means becoming cannot be straightforwardly positioned as a moral ceiling opposite to suffering as floor — because the two are not on separate axes. They are co-constitutive features of the same causal loop.

Evaluative Criterion Remains Suffering: Why Becoming Matters The loop model is a descriptive account. The evaluative question is separate: what

makes one arrangement better than another? The answer remains what 3.3 established: structural suffering produces systemic instability, and arrangements that generate less unnecessary suffering are therefore better — not because suffering is intrinsically bad as a first principle, but because its material consequences for cooperative systems are traceable and documentable. Becoming enters the evaluative framework not as an independent good but through a specific causal claim: foreclosing the response capacity of sentient systems is a way of producing unnecessary suffering that operates on a longer timescale. An arrangement that forecloses adaptive response capacity does not produce immediate floor-level deprivation. It produces something structurally equivalent: it reduces the probability that future suffering can be addressed, and it imposes that reduction on populations that did not choose it. This is structural suffering imposed in the temporal dimension. It falls under the same evaluative criterion as floor-level suffering — not because becoming is a separate good but because foreclosing it is a form of the same bad. This reformulation has a

consequence for the framework's architecture. Earlier versions described becoming as the "anti-foreclosure criterion" — the positive aspiration that complements the floor criterion's negative constraint. That description is misleading if it suggests becoming is an independent evaluative criterion with moral force of its own. The more accurate description: becoming is what the floor criterion requires when its causal implications are traced temporally. Reducing unnecessary suffering is not only a cross-sectional task — it extends forward in time, because arrangements that foreclose response capacity guarantee future populations will be less able to address the suffering their conditions produce. The floor criterion, rigorously applied across time, generates the commitment to anti-foreclosure. Becoming is the name for the property being protected by that commitment — not the reason the commitment has moral force. Where this document uses "becoming" as shorthand for the evaluative criterion above the floor, that shorthand should be read as: "the anti-foreclosure criterion — the floor criterion applied temporally." Against Becoming for Its Own Sake The loop model makes explicit a rejection that was implicit in earlier formulations but not sufficiently stated: becoming for its own sake is not a PMN value. Expansion of cognitive capacity, relational depth, social complexity, institutional sophistication, or technological capability is not automatically good. The historical record of development pursued as its own justification — industrial development that produces immiseration, technological capability that produces mass destruction, social complexity that produces new forms of exclusion — is not a series of anomalies within an otherwise progressive trajectory. It is the expected output of a loop in which res

ponse patterns expand without evaluation against the suffering criterion. Becoming that increases the total quantum of unnecessary suffering — even while expanding capacity — is not better than the alternatives it displaced. The evaluative question is always: does this development reduce unnecessary suffering, including the suffering it will generate by foreclosing future populations' response capacity? If the answer is no, the development is not recommended by this framework regardless of how much capacity it expands. Functional and dysfunctional Suffering: Regrounded Section 3.4b distinguishes functional from dysfunctional suffering. The distinction is real and important, but its grounds require clarification under the loop model. Earlier versions risked grounding it in becoming: dysfunctional suffering was defined as suffering that serves no developmental purpose, making becoming the criterion for evaluating suffering. That has the wrong causal direction. The correct grounding runs through suffering itself, traced forward. dysfunctional suffering is suffering that is (1) unnecessary — institutional alternatives exist that would not produce it; and (2) produces further suffering without producing the response capacity that would enable that further suffering to be addressed. The discomfort of learning is functional not because it serves becoming as an independent good, but because it is part of a response pattern that reduces the probability of the suffering that motivated the learning. The suffering of chronic blocked development is dysfunctional not because it fails to contribute to becoming, but because it is unnecessary and produces further suffering without generating any increase in response capacity. The becoming criterion is not doing the evaluative work. The suffering criterion, traced through its causal implications, is. The Distributional Requirement The distributional requirement that follows from the anti-aggregation commitment (3.4c) applies to the anti-foreclosure criterion as it does to the floor criterion, for the same reasons. An arrangement that expands adaptive response capacity for a majority while foreclosing it for a minority produces unnecessary suffering in the foreclosed minority — not an anti-foreclosure failure, a floor failure in the temporal dimension. Colonialism is not a difficult trade-off between floor and anti-foreclosure criterion values. It is a floor violation across time: the destruction of institutional capacity, knowledge systems, ecological relationships, and governance structures in colonized populations imposes structural suffering on future generations who would otherwise have had the response capacity to address their own conditions. No separate normative premise about the value of becoming is required to evaluate this as a failure. The floor criterion, applied temporally and distributionally, does the work. The relationship between this analysis and the capabilities tradition (Sen, Nussbaum) is wor

th specifying. The capabilities approach identifies central human capabilities whose combined achievement constitutes a life of human dignity. PMN's anti-foreclosure criterion, as reformulated here, is not a

capabilities list and does not require one. It requires only: (1) that adaptive response capacity — the capacity of sentient systems to develop new responses to their material constraints — not be foreclosed by institutional design; and (2) that the evaluation of any specific development (of any specific capability, skill, institution, or practice) run through the suffering criterion rather than through a prior commitment to the value of development as such. This makes PMN’s evaluative criterion more minimal than the capabilities approach — it does not specify which capabilities matter — but more structurally demanding: it prohibits not just the absence of central capabilities but any arrangement that systematically forecloses the development of response capacity across populations and time. [Terminological note: Throughout this document, ‘becoming’ functions as shorthand for the anti-foreclosure criterion developed in this section: the floor criterion applied temporally. Where earlier sections or diagnostic callouts use ‘becoming’ as if it were an independent positive aspiration, that language should be read as shorthand for anti-foreclosure. The substantive claim is always: arrangements that foreclose adaptive response capacity produce structural suffering in the temporal dimension. No claim about the intrinsic value of development, progress, or capability expansion is implied.]

4.5c — The Procedural Account of Becoming

The framework declines to specify the content of adaptive response capacity. This is not an evasion; it is the anti-technocratic guardrail (6.4) applied to the evaluative framework itself. A framework that specifies what human beings should become has overstepped the epistemic limits that make it reliable: it has substituted the analyst’s conception of the good life for the open capacity that becoming is supposed to protect. The history of developmental frameworks that specify positive content — from the Aristotelian account of the contemplative life to utilitarian hedonic maximization to socialist accounts of species-being — is in significant part a history of the particular becoming of particular populations being elevated to universal norm and used to evaluate, constrain, or redirect the development of those who would have chosen differently. PMN’s anti-teleological commitment in section 4.3 is the source of this restraint. What the framework can specify without overstepping is the structural conditions under which adaptive response capacity — whatever form its development takes for a given individual or community — is not foreclosed. This is the procedural account: rather than identifying what genuine development consists in, it identifies what arrangements must be present for adaptive response capacity to remain open. The procedural account is not neutral in the thin sense of making no evaluative commitments; it is substantive in its identification of the structural conditions that count as enabling rather than obstructing development. But its substantive commitments are at the level of conditions rather than content, which preserves the space within which individuals and communities determine for themselves what they are developing toward. The structural availability of becoming requires, at minimum, four conditions. First, the biological floor: the physical and psychological conditions within which development is not simply a matter of survival. Section 3.4’s minimal anchor is the floor below which becoming is foreclosed by the biological conditions of the body — the point at which the organism’s resources are fully consumed by maintenance and nothing is available for development. Above the floor, a range of development becomes possible; the floor does not determine what form development takes, only that it is not ruled out by the demands of survival. Second, the epistemic conditions: access to the range of experiences, frameworks, relationships, and challenges through which the direction of one’s development becomes visible. Becoming that is not chosen from genuine alternatives is not becoming in the sense the framework intends; it is socialization into a predetermined path. The epistemic conditions for preserved adaptive response capacity include exposure to alternatives, the capacity to evaluate them, and protection from the forms of visibility suppression that conceal available alternatives or present a single path as the only natural

one. Third, the relational conditions: the network of relationships within which development is supported, challenged, recognized, and sustained. The capabilities tradition, drawing on Aristotle and developed by Sen and Nussbaum, identifies affiliational capacity as among the most significant conditions for human

development; PMN is in agreement that becoming occurs within relationships, not despite them. Fourth, the temporal conditions: the absence of chronic crisis, existential threat, and resource precarity that consume the time horizon within which development can be planned and pursued. Populations living under conditions of chronic instability do not lack the capacity for development; they lack the temporal horizon within which development projects can be maintained. The procedural account generates a specific evaluative posture. Arrangements are evaluated not by whether they produce the kind of becoming the analyst endorses but by whether they expand or contract the structural availability of these four conditions for all members of the relevant population. An arrangement that expands epistemic conditions for some while contracting them for others — by, for instance, concentrating access to information, education, or political participation — fails the procedural criterion even if the developing minority is developing in directions the analyst considers excellent. An arrangement that maintains temporal stability for the majority while generating chronic crisis conditions for a minority — through, for instance, structural precarity concentrated in specific classes or communities — fails the procedural criterion even if the majority is flourishing. The procedural account embeds a distributional commitment that the content-neutral language of “conditions for becoming” might conceal: it requires that the conditions be available across the full range of the relevant population, not merely in aggregate. Relationship to liberalism: the procedural account shares liberalism’s commitment to not specifying the content of the good life, but it differs in two ways that matter for the framework’s evaluative posture. First, the procedural account specifies the material conditions for genuine choice, not merely the formal absence of coercion. Liberalism in its thin variant treats the absence of coercion as sufficient for genuine choice; the procedural account requires that the structural conditions — epistemic, relational, temporal, biological — be materially present. Second, the procedural account embeds distributional requirements that liberal frameworks often leave implicit: the conditions must be available across the population, not merely in aggregate. This makes the procedural account more demanding than thin liberalism but less prescriptive than the capabilities tradition’s thick account of human flourishing — which is the position the framework intends.

4.5d — Becoming as Epistemic Capacity

Section 4.5b establishes that preserving adaptive response capacity requires cognitive surplus, cooperative surplus, and institutional trust — these are the conditions whose foreclosure constitutes temporal floor violations. These are the conditions under which response capacity remains open to already-known threats and conditions already experienced. This section addresses a structural gap: human societies face novel material threats — emerging pathogens, climate-driven disasters beyond historical ranges, ecological tipping points, resource depletion, and actors with asymmetric destructive capacity — whose defining feature is that they cannot be fully anticipated from past experience. Responding to this class of threat requires not only the absence of current structural suffering but the sustained, collective capacity to learn: to develop scientific understanding, institutional knowledge, and technical capability capable of anticipating, modeling, and mitigating harms that have not yet manifested. This epistemic capacity is a structural component of becoming, not an optional extension of it and not reducible to the floor. The argument that epistemic capacity belongs to becoming rather than to the floor turns on the direction of causation. The floor criterion identifies conditions whose absence produces current structural suffering. Epistemic capacity is not primarily about current suffering: a society without independent research institutions, cross-jurisdictional data sharing infrastructure, or protected capacity for paradigm-challenging inquiry does not immediately experience floor-level deprivation. The harm operates on a different timescale. Novel pathogens will produce mass structural suffering in populations that have not developed the surveillance and response capacity to detect and contain them. Resource scarcity will produce institutional conflict and floor-level deprivation in populations that have not developed the analytical capacity to anticipate depletion curves and substitute before scarcity becomes acute. Ecological collapse will produce irreversible foreclosure of becoming — for generations not yet living — in populations that have not developed the monitoring capacity to identify

tipping points before they are crossed. The absence of epistemic capacity is therefore a cause of future structural suffering, not merely a failure of aspiration. It belongs to the becoming criterion because its consequences operate at the anti-foreclosure level: what it forecloses is not survival today but the developmental trajectory of populations across time. This is where the blind spot in the teleological communist tradition — and in the broader materialist left — becomes structurally visible. Orthodox Marxism assumes that post-capitalist institutional arrangements will automatically generate scientific and technical progress, because the fetters that capitalism places on productive forces will be removed. What this assumption misses is the disti

nction between the removal of barriers and the active construction of epistemic infrastructure. Scientific capacity does not emerge spontaneously from the absence of capitalist relations; it requires specific institutional arrangements — independence from short-term production pressures, protected space for basic research without immediate application, adversarial review that challenges paradigm-confirming results, and coordination mechanisms that allow knowledge to accumulate across jurisdictions and generations. The Soviet case is the clearest illustration of this failure: a system that invested heavily in applied science while systematically destroying the epistemic independence (1.9) that genuine scientific progress requires, producing impressive narrow technical achievements alongside catastrophic failures of anticipation in ecology, agriculture, and public health. PMN's account of becoming must specify what epistemic capacity requires institutionally — not as a teleological guarantee but as a structural condition that must be actively maintained against the same capture dynamics (7.3) that threaten every other institutional function. Structural conditions for epistemic capacity: (1) Institutional independence — research and monitoring capacity must not be subordinated to short-term production or political validation requirements; capture of epistemic institutions (1.9, 7.3) is the primary mechanism by which epistemic capacity is destroyed from within functioning systems. (2) Adversarial review — results that challenge existing paradigms must have institutional protection and pathways to circulation; systems that reward paradigm-confirmation and penalize paradigm-challenge are epistemically brittle. (3) Cross-jurisdictional infrastructure — novel threats do not respect borders; data sharing, coordinated monitoring, and joint response capacity require cooperative institutional arrangements that are themselves subject to capture and must be actively maintained. (4) Long-horizon funding — basic research whose applications cannot be specified in advance must be protected from short-horizon allocation pressures; the threats that epistemic capacity must address operate on timescales that market and electoral mechanisms consistently underfund. (5) Redundancy — single-point epistemic infrastructure is fragile; the knowledge systems that anticipate novel threats must have sufficient redundancy to survive shocks to any single institution or jurisdiction. Connection to 1.8–1.12 (epistemic infrastructure) and Part XVI (technology as structural variable): these sections specify the mechanisms through which epistemic capacity is built or degraded. 4.5d specifies why that degradation is a becoming-level harm, not merely a policy inefficiency.

V: The Metaphysical Dimension: Meaning, Religion, and Material Reality

The materialist tradition has consistently handled the metaphysical dimension badly — not because the questions are unanswerable, but because the tradition imported an answer before it asked the question. For the majority of the world's population, metaphysical commitments are not intellectual puzzles to be held lightly pending further evidence. They are the organizing framework of existence, the source of obligation, the context in which suffering acquires meaning and action acquires moral weight. A framework that treats this as a variable to be managed without acknowledging what it is treating will lose the analytical and political access it needs. What follows engages seriously with the metaphysical dimension without either adopting it uncritically or dismissing it in ways that guarantee the dismissal will not be heard.

5.1 — What the Materialist Tradition Got Wrong About Religion

The materialist tradition's treatment of religion was one of its most significant analytical failures — not in identifying religion as a material force with real consequences, but in treating those consequences as primarily the consequences of an illusion. Marx's famous formulation — religion as the opium of the people, as the heart of a heartless world — was analytically interesting but practically disastrous. It produced a tradition that consistently underestimated religion's durability, misunderstood the relationship between religious belief and material interest, and alienated enormous populations whose material interests might otherwise have aligned with the movements that alienated them. Religion is treated differently here. The metaphysical content of religious belief may or may not be true — that question is held open rather than answered in advance. What is not open is whether religious belief has material consequences. It does, and they are large. How societies organize, how authority is legitimated, how collective action is motivated, how suffering is interpreted and endured — these are all shaped by the religious commitments of the populations involved. Any analysis that treats these as noise to be filtered out in favor of the 'real' material interests underneath will produce wrong predictions and wrong prescriptions. The historical record is consistent and damaging. The Soviet campaign of aggressive secularization from the 1920s onward did not dissolve religious commitment — it drove it underground, transferred its emotional intensity to substitute political religions, and generated populations whose spiritual needs were met by state ideology in ways that proved structurally fragile: when the state could no longer deliver on its promises, the meaning infrastructure it had displaced was unavailable for recovery. The secular left's consistent failure to build durable coalitions with religious working-class populations in Europe and Latin America followed the same logic: movements organized around class interest alone could not offer what movements organized around religious identity offered — a coherent account of why suffering has meaning, what obligation connects individuals to each other across generations, and what the trajectory of a life is supposed to be. These are not secondary needs that become relevant after material needs are met. For most people, most of the time, they are primary ones, and they shape which material improvements feel like improvements and which feel like insults. Taking religion seriously analytically means understanding it as an organizing structure of obligation, time, and meaning — not as a set of empirical beliefs about the world that can be assessed for accuracy and then updated. A tradition that has organized how people understand suffering, what they owe each other, and what the shape of a life is does not change when its metaphysical content is challenged from outside. It changes, if at all, when in

ternal pressures generate reinterpretation from within. Analysis that cannot work with this dynamic — that can only see religious populations as waiting to be emancipated from false consciousness — will fail to understand the most consequential political actor in most of the world. The practical disaster was visible in the twentieth century's materialist political projects. Movements that inherited Marx's dismissal of religion as illusion consistently failed to account for the actual motivational structure of the populations they were trying to mobilize, and consistently found that the destruction of religious institutions — which those populations

depended on for social solidarity, meaning infrastructure, and practical mutual aid — produced political disorganization rather than awakened class consciousness. The assumption that removing the illusion would reveal the material interests beneath it underestimated how much of what those institutions provided was not ideological content but functional social architecture: the networks of mutual obligation, the shared narratives of suffering and endurance, the spaces of communal life that economic organization alone does not supply. The two-level error: The materialist tradition made two distinct errors in its treatment of religion, and they pulled in opposite directions. The first was the reduction of religious belief to ideological function — treating the content of religious claims as merely the expression of social interests, with no independent explanatory value. This error made it impossible to understand why specific religious doctrines took the specific forms they did, why some religious movements produced radical social transformation while others produced conservative stabilization of existing arrangements, and why the same doctrine could support either orientation depending on organizational and material conditions. The second error was the reduction of religious practice to belief — treating what people do in religious contexts as a mere expression of what they believe, when the anthropological and sociological evidence consistently shows that practice precedes, shapes, and often contradicts belief in complex ways that make the belief-practice relationship bidirectional rather than unidirectional. What the tradition got right: The materialist tradition's core insight about religion remains valid and analytically valuable: the specific forms that religious institutions take, the doctrines they emphasize, and the political orientations they produce are systematically shaped by the material conditions and power arrangements within which those institutions operate. The medieval Catholic Church's theological justifications for hierarchy are not independent of its position as the largest landowner in Europe. Liberation theology's emergence in Latin America in the 1960s is not independent of the conditions of poverty and political repression under which it was developed. The content of religious t

eaching is not determined by material conditions in a simple causal chain — the tradition has genuine internal logic and genuine doctrinal autonomy — but the selection of which doctrines receive emphasis, institutional support, and social circulation is shaped by material conditions in ways that are analytically traceable. This is what PMN inherits from the materialist tradition: religion as a material variable, not religion as an illusion. The analytical shift required: from asking whether religious beliefs are true or false — a question the framework is agnostic on — to asking what functions religious institutions and practices perform, for which populations, under what conditions, with what consequences for structural suffering and the conditions for preserved adaptive response capacity. That question is answerable with evidence. It produces a different kind of engagement with religious traditions than either dismissal or uncritical deference — one that can identify when religious institutions are performing functions that reduce structural suffering and when they are performing functions that maintain it, and that can engage with religious actors as potential allies or opponents based on structural analysis rather than ideological classification.

5.2 — Two Levels of the Same Question

A more useful structure for thinking about the metaphysical dimension distinguishes between what might be called the source level and the interface level of reality. The source level is whatever exists at the most fundamental level — if anything does beyond the physical. If there is a necessary being, this is where it would be. By hypothesis, it operates as it operates, without preference, intention, or moral orientation in any sense that implies choice. It simply is what it is, and everything else is its consequence. The interface level is what appears when consciousness — human experience, perception, meaning-making — interacts with whatever the source level is. What emerges at this level can feel personal, responsive, and morally engaged, regardless of whether the source level has any of those properties. Real at the level of experience — not distortion or illusion, but a genuine property of the interaction, one that exists even if it does not exist at the level of whatever underlies experience. The interface level may be purely emergent — a property of the interaction with no independent existence beyond it. Or it may have genuine independent existence that the

interaction reveals rather than generates. The framework holds this distinction open because collapsing it prematurely in either direction produces conclusions that exceed the available evidence. The materialist tradition collapsed it in one direction — treating the interface as purely generated by the interaction — and paid the analytical cost of that collapse. (Chalmers 1996; Nagel 1974) The analytical implications of holding this distinction open are not merely philosophical. A political movement that collapses it — treating the interface level as purely generated by psychological and social function, as 'merely' emergent — engages with religious populations in ways that are experienced as dismissive even when the engagement is technically respectful. The dismissal operates at the level of ontology: the implicit claim that the experiences organizing a population's deepest commitments are fully explained by social function, pointing toward nothing beyond themselves. Populations whose lives are organized around those experiences do not typically need to articulate this dismissal to feel it. The analytical posture that holds the distinction open avoids this structural condescension — not as a strategic concession but as the epistemologically honest position: the claim that interface-level experience is 'merely' emergent with nothing behind it exceeds the available evidence in exactly the same way as the claim that it directly reflects a fully personal, intentional source level. For political analysis, the practical consequence is specific: engagement at the interface level — with the meaning structures, obligations, and accounts of suffering that actually organize behavior — is both analytically more accurate and politically more effective than engagement at the level of metaphysical d

ispute. The question of whether the source level has the properties that any specific tradition attributes to it cannot be settled by available evidence. The question of what the interface level produces — what forms of solidarity, obligation, endurance, and collective action it makes available — can be examined by the standards that apply to any empirical claim. This is where analysis can work. What the source level establishes and does not establish: If there is a necessary being at the source level — if something exists necessarily and everything else is its consequence — that fact has no direct normative implications. A necessary being that simply is what it is, without preference or intention in any morally loaded sense, does not supply moral guidance, does not validate any particular tradition's claims about divine commands, and does not authorize any particular institutional arrangement as divinely sanctioned. The framework's agnosticism on whether there is a necessary being is not agnosticism about the absence of implications — it is the recognition that the implications flow from what kind of necessary being exists, and that question cannot be resolved by the framework's analytical methods. What can be established is that the contingency of everything observable raises a question that material explanation cannot close, and that intellectual honesty requires naming the question rather than papering over it. The Problem of Evil and the Two-Level Resolution

The classical problem of evil — the Epicurean formulation that an omnipotent, omniscient, and perfectly good God is incompatible with the existence of suffering — presupposes a specific attribute at the source level: that whatever is there has both the capacity and the moral orientation to prevent suffering and chooses not to. If the source level is impersonal — if it operates as structure without preference or intention, as the absolute order within which everything unfolds without anyone 'permitting' or 'allowing' anything — then the problem of evil is not resolved but dissolved. The question 'why does God allow suffering?' contains an embedded assumption that the source level has something analogous to agency. Remove that assumption and the question loses its grammatical structure. Suffering is not 'allowed' by an impersonal absolute structure in the same sense that objects are not 'allowed' to fall by gravity. Both are simply consequences of the structure within which they occur. The two-level framework provides a specific response to the worry that this dissolution is a form of nihilism. If the source level is impersonal, the worry goes, then nothing ultimately matters, there is no ground for meaning, and the consolations of religion are elaborate self-deception. This move fails because it conflates two levels of analysis. The interface level is not diminished by the impersonality of the source level. Consider currency: at the level of fundamental physical reality, money is metal, p

aper, or entries in a server. Nothing at that level 'cares' about money; no physical law assigns it value. At the interface level — the level constituted by human practice, coordinated expectation, and institutional structure

— money is genuine value. The consequences of not having it are real. The capacity to acquire it, lose it, share it, or hoard it generates real effects on real lives. The question ‘is money real?’ has a different answer at each level, and both answers are correct within their domain. The same structure applies to meaning, obligation, and religious experience at the interface level. The functional/ontological distinction: This reframing separates two questions that the classical problem of evil runs together. The ontological question asks what ultimately exists and what its nature is. The functional question asks what consequences follow within the domain of experience and practice. A source level that is impersonal and indifferent does not produce a functional level that is therefore meaningless. It produces a functional level whose structure and consequences are real within the domain in which they operate — just as the rules of a complex system are real within the system regardless of whether the system was ‘designed’ with intentions or emerged from impersonal processes. If an ethical or spiritual order is real at the interface level — if it generates genuine consequences for behavior, for development, for the conditions of becoming — then it is real in the sense that matters for how anyone who lives within that level should orient themselves, regardless of what the source level turns out to be. One further implication follows from this structure. If the source level is impersonal, then whatever is experienced as personal at the interface level — the sense of being accompanied, of there being something responsive to one’s situation, of moral weight that is not merely self-imposed — emerges from the interaction of consciousness with impersonal structure. It is not transmitted from the source. This does not make it less real at the level at which it appears. It does mean that the interface-personal and the source-impersonal can coexist without contradiction: the personal God of religious experience and the impersonal absolute of philosophical theology may be two levels of the same reality rather than competing descriptions of the same level. Whether this is in fact the structure of reality is not something the framework can determine. That it is a coherent possibility — and one that removes the most persistent objections to both religious and materialist accounts — is something the framework can establish. Summary of the two-level resolution to the problem of evil: (1) If the source level is impersonal, suffering is not ‘permitted’ by anyone — it is a structural feature of a system without moral agency at its ground. (2) If the interface level is real in the functional sense, then meaning, consolation, obligation, and the sense

of being accompanied through suffering are not illusions — they are genuine features of the level at which human experience operates. (3) The nihilist inference — that impersonal source implies meaningless interface — commits a category error by treating what is true at one level as decisive for the other. (4) Whether the source level is in fact impersonal, whether the interface level has independent existence beyond emergence, whether any specific tradition’s account of the interface is accurate — all of these remain open questions that the framework holds with appropriate uncertainty rather than closing prematurely in either direction. What the interface level makes available: The interface level — what appears when consciousness meets whatever is there to meet — is the level at which religious experience, meaning, beauty, moral intuition, and the sense of something beyond ordinary material transaction become available as real phenomena with real consequences. These phenomena are not illusions in the dismissive sense. They are interface-emergent properties: real at the level at which they appear, with causal consequences that operate at that level, even if their ultimate explanation may require reference to a source level whose nature is undetermined. The framework treats them as material variables — forces that shape individual motivation, social cohesion, political mobilization, and the distribution of meaning in ways that have measurable consequences for structural outcomes — without requiring a position on whether they indicate anything beyond themselves. The relationship between the levels: The two-level structure is not a dualism that posits two independent substances. It is a recognition that the same phenomena can be analyzed at different levels of description with different analytical tools, and that the results of each level of analysis are not straightforwardly reducible to the results of the other. The interface-level analysis of religious experience — what it does for people, how it shapes motivation and community, what social functions it performs — does not determine whether the source-level question has an answer. The source-level question — whether there is

something that the interface-level phenomena point toward — does not determine the outcome of the interface-level analysis. Both levels are required for adequate description, and neither should be collapsed into the other. The materialist tradition collapsed the interface into the source and then dismissed both; idealist traditions collapse the source into the interface and thereby inflate it. The framework holds both levels open. The practical consequence of this two-level structure for analysis: any claim that religious or metaphysical experience is "merely" something else — merely social solidarity, merely psychological compensation, merely ideological legitimation — is making a reductive move that the framework does not authorize. The "merely" implies that the interf

ace-level phenomenon has been fully explained by identifying its material conditions or functions. But explaining what something does is not the same as explaining what it is, and the framework's agnosticism at the source level means that the question of what interface-level phenomena ultimately are remains genuinely open. What the framework can do is analyze what they do — the functions they perform, the consequences they produce, the conditions under which they support or undermine the minimal anchor criterion — and that analysis is sufficient for the political and institutional purposes the framework serves. The two-level structure connects to Part I's epistemic commitments in a specific way: claims at the source level are not falsifiable by the methods available to this framework, which means they cannot be affirmed or denied with the confidence that falsifiable claims warrant. The framework's response is not to dismiss them — that would be to treat absence of falsifiability as equivalent to falsification, which is an epistemic error — but to hold them with appropriate uncertainty while analyzing what can be analyzed at the interface level with appropriate rigor. The uncertainty is named, not buried.

5.3 — Interface-Emergent Phenomena

The concept of emergence is not exotic. It appears wherever complex systems generate properties at higher levels of organization that are not present in the components of the system. Wetness is not a property of individual water molecules. The liquidity of a market is not a property of any individual transaction. The personality of a character in a sufficiently sophisticated simulation is not a property of any line of code. These properties are real at the level at which they appear, and they have real consequences, even though they are not present at the level of the components. The personal quality of religious experience — the sense of being heard, guided, held accountable by something beyond oneself — may be an emergent property of this kind. Real at the level of experience. Having real consequences for behavior, psychology, and social organization. Not necessarily requiring that the underlying level has the properties that the experience attributes to it. This reframing does not debunk religious experience. It relocates it — from a claim about what the source level is like to an observation about what emerges at the level of the interaction between consciousness and whatever the source level is. The properties that emerge at the interface level when consciousness encounters whatever the source level is include several that are directly relevant to political and social analysis. The experience of moral obligation — the sense that something is genuinely required of a person, not merely preferred or socially enforced — is among the most powerful motivators of collective action available, and its durability under adversity distinguishes it from interest-based motivation, which depletes when costs become too high. The experience of solidarity with others through shared orientation toward something beyond the group — the 'we' constituted not by proximity or interest but by common relationship to a transcendent reference — creates forms of collective identity that survive the disruptions of migration, economic change, and political repression that dissolve solidarities based on interest alone. The experience of suffering as meaningful rather than merely painful — as having a trajectory, a context larger than the individual experience, a relationship to something that makes it comprehensible — is one of the most consequential variables determining how populations respond to structural deprivation. Populations whose suffering is organized within a meaning structure that makes it intelligible and possibly transformative respond differently to identical objective conditions than populations whose suffering is senseless. None of

these properties require the source level to have the specific characteristics that any tradition attributes to it. They are real at the level at which they operate — real in their consequences for behavior, psychology, and collective action — and they are systematically unavailable to analysis that reduces them to epiphenomena

of material conditions. The reframing that emergence provides is not a concession to religion but an analytical correction: you engage with these properties at the level where they operate, not by disputing what they may or may not be at the level beneath. The Interface Level and the Durkheimian Reduction The distinction between the interface level and the Durkheimian tradition of religious sociology requires explicit statement because it is easy to conflate the two and the confusion has analytic costs. Durkheim's central claim is that the object of religious experience — the sacred, the divine, the transcendent — is in fact society representing itself to itself. God is the community. The force experienced as coming from beyond the individual is real, but its source is social solidarity rather than any metaphysical entity independent of the social group. This is a sophisticated and analytically powerful position, and it has generated the most productive tradition of empirical religious sociology. But it is not the interface framework, and treating them as equivalent removes the critical feature that makes the interface framework philosophically honest. (Durkheim 1912; Geertz 1973) Durkheim reduces the transcendent to the social. The interface framework brackets that reduction and holds it open. The difference is not merely scholastic. An analyst who adopts the Durkheimian position has already answered the source-level question — the sacred is social solidarity, full stop — and proceeds to study its social functions. An analyst who adopts the interface framework has committed to studying what emerges at the level of experience without prejudging whether the emergence is purely from social processes, from something in material reality that social processes are responding to, or from some combination. The Durkheimian analyst has a better specified model. The interface analyst has a more epistemically honest one, because the question that Durkheim answered in advance has not actually been answered. For practical analysis, this distinction matters in one specific way: when religious communities claim that their experience is of something real beyond social solidarity — that their tradition's metaphysical commitments track something independent of their social functions — the Durkheimian framework treats this claim as descriptively false while functionally interesting. The interface framework treats it as epistemically underdetermined: it cannot be confirmed by the methods available to this framework, but it also cannot be dismissed by them. The practical consequence is that engagement with communities that hold these commitments is less likely to be experienced as condescending under the interface framework than under the Durkheimian one — not because the analyst has adopted the community's metaphysical commitments, but because the analyst has not rejected them on grounds the evidence cannot support. The operationalization question: "interface intensi

ty" can be approximated through behavioral and survey proxies that do not require settling the source-level question. Relevant components include: ritual participation frequency; doctrinal centrality (self-reported importance of metaphysical commitments vs. ethical practice); moral urgency (willingness to bear personal costs in defense of religiously grounded claims); and symbolic density of public expression. These proxies are imperfect and context-specific, but they operationalize "interface intensity" as a causal variable without requiring metaphysical commitment. The empirical sociology of religion — Durkheim, Geertz, Weber's analysis of the relationship between religious ethic and social organization — provides the methodological toolkit. What the interface framework adds is epistemological discipline about what that toolkit can and cannot establish. (Geertz 1973; Weber 1905; Durkheim 1912) Why emergence matters for the metaphysical question: The emergence argument matters for the metaphysical dimension because it establishes that the existence of consciousness, meaning, religious experience, and moral intuition as real phenomena does not require positing a non-material substance to explain them. If properties can be real at higher levels of organization without being present at lower levels — if wetness is real without water molecules being wet — then consciousness, meaning, and value can be real at the level of sufficiently complex organized systems without requiring a commitment to substance dualism. This is not a resolution of the hard problem of consciousness, which remains genuinely hard. It is a clarification of what the hard problem requires: not a

different substance, but an account of how specific kinds of material organization produce properties that are not present in the components. The framework does not resolve this question. It holds it open without treating it as already settled by the materialist commitment. The specific emergence claims: The framework makes three specific emergence claims that are relevant to the metaphysical dimension. First: consciousness is an interface-emergent property of sufficiently complex biological systems — real at the level at which it appears, with causal consequences that operate at that level, not fully reducible to the physical properties of the neurons that instantiate it. Second: meaning is an interface-emergent property of conscious beings in relationship with their conditions — not a label applied to pre-existing meaningless events, but a property that arises through the interaction of consciousness with its environment and that has real consequences for motivation, suffering, and the conditions for preserved adaptive response capacity. Third: moral intuition — the immediate, pre-theoretical sense that some things matter and others do not — is an interface-emergent property whose ultimate explanation is contested but whose reality as a causal force in human behavior is not. All three cl

aims are contested in philosophy of mind and metaethics. The framework does not resolve the contests; it takes the emergence of these properties as starting points for analysis rather than as conclusions requiring derivation. The Mechanism: Consciousness Meeting Complexity The emergence of interface phenomena from an impersonal source level has a specific mechanism that the general concept of emergence does not by itself specify. The general claim — that higher-level properties can arise from lower-level structures without being present in those structures individually — is true but leaves open how the specifically personal qualities of interface experience arise from a source that, on the impersonal reading, has no personal qualities. The answer the framework can offer without overcommitting: interface-personal phenomena are not transmitted from the source but generated at the intersection of complex structure and conscious engagement. The analogy that makes this most tractable is not biological but computational. A sufficiently complex system — one with enough internal structure, feedback, and organization — will present differently to a conscious observer than to no observer at all. The observer does not add meaning to a system that has none; the observer's engagement with the system generates a level of description that is genuinely new. The code of a game has no intentions, no emotional states, no relationship with the player. But the player's engagement with the game generates experience that is structurally similar to relationships: consequences that feel personal, challenges that require genuine response, losses that register as real. None of this is in the code. All of it is real at the level at which the player is operating. The interface-personal arises from the encounter, not from either term of the encounter alone. Applied to the source level: if there is an impersonal necessary structure — a 'source code' that simply is what it is — the encounter of consciousness with that structure generates interface phenomena that are not present in the structure itself. The personal God of religious experience, on this account, is not the source level incorrectly apprehended as personal. It is the interface-emergent form that the impersonal source takes when encountered by consciousness. This is compatible with the religious claim that the experience is real — it is. It is compatible with the materialist observation that the source level has no personal attributes — it may not. What it resists is the inference that emergence at the interface level requires either falsifying the source-level impersonality or dismissing the interface experience as illusion. A further implication for how religious traditions are assessed: if interface-personal emergence is real at the functional level, then the question of whether any specific tradition accurately maps the interface is an empirical question about consequences — the standard diagnostic the framework

applies to all institutional arrangements. A tradition whose practice generates the interface conditions it describes, whose obligations produce the structural consequences they predict, and whose account of the interface is coherent across conditions is more credible than one that fails on these dimensions — not because this confirms its source-level metaphysics, but because it demonstrates that its interface-level map has traction. The framework does not adjudicate between competing source-level claims. It can assess the functional coherence of competing interface claims by the criterion it applies to everything else. The political

economy of emergence denial is not neutral. The systematic reduction of consciousness to mechanism, meaning to preference, and moral intuition to evolutionary artifact serves specific ideological functions: it makes the claims that emerge from interface-level phenomena — claims about dignity, about meaning that cannot be monetized, about obligations that are not instrumental — easier to dismiss as confused category errors rather than genuine constraints on what arrangements are acceptable. The framework's refusal to make the reductive move is not a concession to mysticism. It is the insistence that the full range of what is real at the interface level must be part of the analysis, because arrangements that produce flourishing or suffering at that level are producing flourishing and suffering that matters, regardless of whether the ultimate explanation of consciousness is materialist or not. Emergence and the minimal anchor: The suffering that grounds the minimal anchor is an interface-emergent phenomenon. Pain, distress, the experience of deprivation and constraint — these are real at the level at which they appear, with causal properties that operate at that level. The fact that they are produced by physical processes in biological systems does not reduce their reality or their evaluative weight. The emergence framework provides the ontological basis for the minimal anchor: because suffering is real at the interface level, because it has properties and consequences at that level that cannot be fully captured by description at the physical level alone, it is available as a genuine evaluative starting point rather than a derivative of some more fundamental property. The is-ought bridge in 3.3 runs through this point: the ought comes not from the mere existence of physical processes but from the interface-emergent reality of what those processes produce when they produce suffering in conscious beings. The analytical discipline the emergence framework requires: when analyzing phenomena at the interface level — consciousness, meaning, religious experience, moral intuition — neither dismissing them as "merely" their material substrate nor inflating them into evidence for metaphysical claims beyond what they support. They are real at the level at which they appear. Their ultimate explanation is underdetermined. Their consequences

for structural outcomes are analyzable. That is sufficient for the framework's purposes, and claiming more than that in either direction — either reducing them to nothing or elevating them to metaphysical proof — is an epistemic error the framework's commitments in Part I prohibit.

5.4 — The Evaluative Standard for Metaphysical Claims

The most productive approach to evaluating metaphysical claims is to identify where they make contact with material reality — where they predict or explain things that can in principle be examined. A metaphysical claim that produces no observable consequences cannot be assigned explanatory weight; it can be left open. A claim that makes contact with material reality can be examined by the standards that apply to any empirical claim: consistency, predictability, falsifiability, and responsiveness to counter-evidence. What follows from a successful examination is more limited than metaphysical traditions usually claim. Even if a specific phenomenon is shown to be real and not fully explained by current frameworks, this does not validate the specific theological or doctrinal interpretation of that phenomenon. The gap between 'something is happening here that we cannot yet explain' and 'therefore this specific tradition's account of what is happening is correct' is very large, and it requires independent justification at every step. Three questions organize the evaluation in practice. Does the claim predict anything that would distinguish a world in which it is true from a world in which it is false? A metaphysical claim consistent with every possible observation is not a claim about reality — it is a preference dressed as one. Where the claim does make contact with observable consequences, do those consequences hold up under the epistemological standards of Part I: evidence-responsiveness, falsifiability, responsiveness to counter-evidence? And when a claim's observable implications are contested, who bears the burden of evidence? The burden falls on the claim that attributes more structure to reality — that posits specific entities, specific causal mechanisms, specific intentions. The more specific the claim, the more demanding the evidentiary standard it must meet. What this approach cannot do is settle metaphysical questions that make no observable contact with material reality. Those

questions remain open — not because they are unimportant, but because the framework has no tools for adjudicating them and declines to pretend otherwise. The openness is not agnosticism as a gesture. It is the application of the epistemological commitment of Part I to a domain where that commitment implies: we do not currently know, and claiming otherwise exceeds the available evidence. The application of this standard to the framework itself is required by its own epistemological commitments. The framework's agnosticism about the source level — its refusal to either affirm or deny what exists at the most fundamental level of reality — is not a diplomatic concession to religious audiences. It is the honest position given the available evidence. The further claim that interface-level experience is real in its consequences regardless of what the source level is like follows from the same epistemological commitment: the evidence for the consequences is substantial, the

evidence for claims about the underlying ontology is not, and the honest analyst holds the claims to the level of evidence available. What the framework declines to do is pretend that the question is settled in either direction. Divine Command Theory: The Euthyphro Problem and PMN's Methodology A specific and ancient philosophical challenge to the relationship between divine authority and moral evaluation bears directly on the framework's methodology. Plato's Euthyphro poses the question in its sharpest form: is something morally good because God commands it, or does God command it because it is good? The dilemma is not merely academic. It determines whether divine commands constitute an independent source of moral authority that overrides all other evaluation, or whether divine commands are themselves subject to the same evaluative standards that apply to any other moral claim. (Plato, Euthyphro, c.399 BCE) If something is good solely because God commands it — the Divine Command Theory in its strong form — then morality is entirely dependent on the will of a being whose commands cannot be evaluated by any independent standard. The practical consequence is that any action, however materially destructive, becomes morally required if it can be successfully claimed as divine command. History provides ample examples of this logic in operation: the commanded extermination of populations, the institutionalized subordination of women, the persecution of heretics, the sanctification of slavery — all have been defended, at various points, as fulfillment of divine will. The framework that evaluates nothing but the authenticity of the divine source and the sincerity of obedience has no internal resources to distinguish these cases from cases of genuine moral obligation. If God commands things because they are good — the alternative horn of the dilemma — then there exists a standard of goodness that is independent of God's will and to which God's commands themselves conform. This is the position that is consistent with PMN's methodology, and it is also the position that most sophisticated theological traditions have actually held, even when they did not frame it in Platonic terms. The Thomistic tradition's insistence that God's commands are expressions of God's rational nature — that God cannot command cruelty because cruelty is inconsistent with the divine rationality that constitutes God's essence — is a version of this position. The Islamic tradition's identification of God's commands with what is genuinely *maslaha* — genuine benefit for human beings — is another. These traditions are not undermined by the Euthyphro argument. They are confirmed by it: their most sophisticated exponents already recognized that divine commands cannot be evaluated in isolation from their consequences for human wellbeing. PMN's methodological position follows from this: the framework evaluates the material consequences of claimed divine commands by the same evaluative criteri

on it applies to all institutional arrangements — do these arrangements create conditions for genuine development and reduce the structural suffering that forecloses it? This is not a claim that God does not exist, that divine commands are fabrications, or that religious obligations are illegitimate. It is the application of the Euthyphro insight: if divine commands are genuinely good, they will produce outcomes that the evaluative criterion confirms. If claimed divine commands produce systematic structural suffering, concentrated in populations who cannot contest the interpretation, the framework identifies a gap between the claim and its consequences that requires explanation — whether through genealogical analysis of who benefits from the specific interpretation, through examination of the conditions that produced that interpretation, or through the possibility that the interpretation is mistaken about what the divine command actually requires. The

methodological principle is precise: PMN does not adjudicate between competing claims about what God has commanded, and it does not make metaphysical claims about whether any such commands exist. What it does is evaluate the material consequences of implementing what is claimed as divine command, using the same criterion applied to all other institutional arrangements. A claimed divine command that produces outcomes consistent with the evaluative criterion — expanding structural conditions for becoming, reducing unnecessary suffering — is not challenged by this framework on methodological grounds. A claimed divine command whose consistent implementation produces outcomes that the evaluative criterion identifies as structural harm raises the Euthyphro question: if God commands this because it is good, where is the goodness in these consequences? The question is not hostile to faith. It is the question that the most serious theological traditions have always asked of themselves. When Interface Claims Conflict: Adjudication Without Metaphysical Resolution The two-level framework handles individual engagement with religious or metaphysical commitments by bracketing the source level and analyzing interface consequences. A harder problem arises when two communities' interface-level claims are mutually incompatible in their political demands — when each group's obligation structure, derived from its interface experience and its tradition's commitments, requires outcomes that the other group's obligation structure prohibits. Holding agnosticism at the source level does not by itself resolve who receives public accommodation, which obligations have institutional standing, or how scarce public goods are allocated between competing sacred claims. The framework requires an adjudication mechanism, and that mechanism must operate at the interface level without importing metaphysical judgments about which tradition's source-level claims are correct. The mechanism the framework provides is the same one applied to a

ll institutional arrangements: the evaluative criterion. Interface claims are not adjudicated by their metaphysical credentials but by their material consequences for the populations they affect. The question is not "whose God is real?" but "what does implementing this claim produce?" — specifically, does it reduce or increase structural suffering, does it expand or foreclose conditions for preserved adaptive response capacity, and does it impose costs on populations that did not accept the obligations generating the claim? This criterion is neutral between traditions: it applies to the sacred land claims of indigenous communities with the same force as to the eschatological obligations of Abrahamic traditions, and to the secular meaning frameworks of communities whose interface experience is organized around political identity rather than religious belief. Three adjudication principles: (1) Material consequence primacy: the adjudication criterion is what the claim produces, not what tradition it comes from. Claims with identical material consequences are treated identically regardless of their metaphysical framing. (2) Proportionality and third-party costs: the weight given to an interface claim in adjudication is scaled by the cost it imposes on parties outside the obligation structure that generated it. Obligations that fall entirely within a community whose members have voluntarily accepted them carry more standing than obligations imposed on those who have not. (3) Revisability: adjudicated accommodations should include sunset review, not because the framework expects interface claims to change on a policy timeline, but because the material consequences of implementing them change as conditions change, and the adjudication criterion tracks consequences, not commitments. The strategic appeasement risk that this framework generates must be named explicitly: the two-level structure is available for misuse by actors who wish to grant permanent privileges to organized religious constituencies while characterizing the arrangement as empirical engagement with "interface effects." The tell is the absence of adjudication criteria: an analyst or policymaker who invokes interface reality to justify accommodation but refuses to specify what material consequences would trigger revision has converted the framework into a rhetorical shield for political dependency on religious patronage networks. The framework's own test applies here: if the claim that a policy is "engaging interface phenomena" cannot be falsified by any material outcome, it is not analysis. It is appeasement in analytic dress. Adjudication diagnostic for competing interface claims: (1) What material consequences does each claim produce for the populations it directly affects and for third parties? (2) Are the costs imposed on third parties proportionate to the accommodation being sought? (3) Is the accommodation time-limited or sunset-reviewed against

material outcomes, or is it construct

ed as a permanent exemption from the framework's evaluative criterion? (4) Is the framework being invoked to justify an accommodation, or to evaluate one? — the first use is legitimate; the second is the strategic appeasement risk. (5) If the claim were made by a different tradition with identical material consequences, would the same accommodation be granted? This consistency check is the primary guard against both ideological bias and strategic appeasement.

5.5 — Religion as Material Variable

For purposes of institutional design, policy analysis, and political strategy, religion must be treated as a material variable regardless of one's metaphysical commitments. This was the insight that the materialist tradition reached but then misapplied — recognizing that religion has material consequences while treating those consequences as the consequences of an error that could be corrected by replacing religious belief with more accurate beliefs. The correction did not work. The belief systems proved more durable than the theory predicted, and the attempts to accelerate their dissolution produced results that were materially worse than tolerating them. The same tradition that legitimated the arrangements of conquest produced the liberation theology that organized the poor against those arrangements — not because the tradition was incoherent, but because the criterion it claimed to serve was more demanding than the arrangements it initially served. (Gutiérrez 1971) The practical implication is that policies and movements that ignore or antagonize deeply held belief systems will fail in predictable ways — not because belief deserves special protection but because meaning systems embedded in identity, community, and social structure are not simply updated when their metaphysical content is challenged. Disrupting them without adequate replacement produces conditions that the minimal anchor requires avoiding. The appropriate approach is to work with the material reality of belief systems as they actually exist, understanding them from the inside well enough to engage with them honestly rather than dismissing them from the outside in ways that guarantee failure. (Gramsci 1971; Berger and Luckmann 1966) Working with belief systems from the inside requires understanding three things that outside analysis consistently misses. First, the internal logic of a tradition — the way its metaphysical commitments generate specific obligations, specific evaluations of suffering, specific accounts of legitimate authority — is not arbitrary. It follows from premises that the tradition has developed over centuries, and engaging with it requires understanding what follows from what, not merely cataloguing surface practices. An analyst who understands only that a population is religious, without understanding which tradition, which interpretation within that tradition, and which obligations that interpretation generates, has not done the analytical work that the framework requires. Second, the relationship between a tradition's metaphysical content and its political expression is not fixed. The Gutiérrez case is paradigmatic precisely because it demonstrates the internal logic: the same tradition that had legitimated colonial arrangements contained, in its own commitments, a comprehensive critique of those arrangements. Liberation theology did not import an external political framework into Catholicism — it followed the tradition's own logic of obligation toward the

poor to its structural conclusions. The tradition did not change. What changed was which of its implications were followed. Any analysis that treats 'Catholic' or 'Islamic' or 'Buddhist' as a fixed political variable has not understood the tradition well enough to work with it. The political expression of any tradition is a site of genuine internal contestation, not a determined consequence of metaphysical content. Third, the relationship between sincere belief and social function is not a trade-off in which one is 'really' true and the other is instrumental cover. A religious commitment that simultaneously provides meaning in suffering, organizes solidarity, and motivates collective action is not 'really about' social function with the metaphysical content as window dressing. The metaphysical content generates the functional consequences — remove it and the consequences do not reliably follow. The analyst who understands this can engage with the tradition honestly, on terms that the tradition's own participants recognize, rather than treating it as a vehicle for

interests that the participants themselves do not acknowledge. Eschatology as a Special Case: The Hardest Form of Meaning Infrastructure Eschatology — the dimension of religious tradition concerned with final things: judgment, resurrection, the end of history, the ultimate vindication of the righteous and accountability of the unjust — is not simply one belief among others within the meaning systems analyzed here. It is the most structurally significant component of religious meaning infrastructure, and it requires specific analytical attention because it performs functions that no secular framework has yet demonstrated the capacity to replace. What eschatology does that ordinary meaning systems do not is answer the questions that remain open after structural analysis has done everything it can: whether the suffering of those who died before justice arrived matters; whether sacrifice that produces no visible return in the lifetime of the person who makes it is coherent; whether the asymmetry between what perpetrators escape and what victims lose is final. These are not abstract theological questions. They are the questions that determine whether populations in conditions of sustained structural suffering — occupation, dispossession, multigenerational deprivation — can maintain the psychological coherence required for collective action over timescales longer than a generation. A framework that answers these questions, however the framework's metaphysical claims may stand up to external scrutiny, is doing material work that its secular alternatives have largely failed to perform at scale. The practical consequence is that eschatological belief systems are among the most resistant to the standard tools of political analysis and engagement. The resistance is not irrational from within the framework. A person operating with a coherent eschatological worldview has a reason structure in which

the costs and benefits that secular analysis treats as decisive are explicitly subordinated to a longer accounting — one that secular frameworks cannot access, cannot refute, and cannot replace simply by demonstrating that the material calculus runs differently. A martyr in an eschatological frame is not making a miscalculation. The analyst who treats the martyr as confused about incentives has failed to understand the framework within which the choice is coherent. This does not make the choice immune to analysis. It means the analysis must operate at the level where the choice actually lives. The Middle East provides the most concentrated test case for this analysis. What makes it analytically distinctive is not merely that eschatological belief is present alongside material conflict — it is that multiple eschatological frameworks with competing and partially overlapping claims are simultaneously active over the same physical space, and that space has specific eschatological significance within each framework. The site variously known as Temple Mount and Haram al-Sharif in Jerusalem is simultaneously the location where Jewish eschatological tradition locates the rebuilding of the Temple as a condition for messianic arrival, where Islamic tradition locates events of critical significance in end-times narratives, and where Christian Zionist eschatology — highly influential in American political support for Israeli policy — positions the ingathering of the Jewish people as a necessary precondition for the Second Coming. These are not three separate conflicts that happen to occur in proximity. They are three eschatological frameworks whose internal logics generate incompatible claims on the same physical reality, and each framework provides its participants with reasons to treat compromise as cosmically illegitimate rather than merely politically difficult. The jihadist eschatology that organizations such as ISIS have operationalized adds a further dimension: the establishment of a specific territorial and political form — the Caliphate — is not merely a political goal but an eschatological condition, a necessary precursor to the sequence of events that culminates in final judgment. This means that the standard tools of conflict resolution — negotiation, mutual recognition, phased concessions — are not merely politically difficult but theologically incoherent from within the framework. You cannot negotiate a partial version of an eschatological condition. The framework's own logic requires the whole or makes the partial version worse than nothing, because a corrupted approximation forecloses the real thing. The analytical implication is not that eschatological conflicts are irresolvable — history includes cases of eschatological frameworks being renegotiated, reinterpreted, or displaced as material conditions changed. It is that resolution requires working with the eschatological dimension rather than around it. Secular analysis that treats the reli

gious content as a mobilizing veneer over material interests, to be managed by adjusting material incentives, has consistently failed in these contexts not because material interests are absent — they are very much present — but because the eschatological framework determines which material arrangements are acceptable and which are cosmically prohibited. The Gutiérrez case demonstrates that traditions can be engaged from within to follow their own internal logic toward different political conclusions. The same analytical move is required here, and it is harder here because the eschatological stakes have been articulated more explicitly and over a longer period of contested history. Disenchantment: The World Emptied from Within Before analyzing eschatology as a particular case of meaning infrastructure, a longer trajectory requires attention — one that begins not with the secularization of the modern period but with transformations internal to religious traditions themselves. The disenchantment of the world, as Weber identified it, is not primarily a product of science or atheism. It is a product of a specific kind of religious development: the progressive transcendentalization of the divine, which relocated the sacred from the immanent world to a realm beyond it, and in doing so progressively emptied the material world of intrinsic meaning. (Weber 1919) In animist and polytheist frameworks, the world is saturated with presence. Spirits inhabit rivers, mountains, and trees. The gods are anthropomorphic — they desire, betray, compete, and negotiate. The boundary between the human and the divine is permeable: heroes become gods, gods take human lovers, oracles transmit divine speech in human mouths. This is not primitive confusion about causation. It is a meaning infrastructure in which the material world is itself meaningful — in which every significant event is intelligible within a framework that connects it to forces that are real, relational, and responsive. The suffering of drought is not random; the success of harvest is not merely mechanical. Everything that happens happens within a world that is, in Charles Taylor's term, enchanted: a world in which meaning is not a human projection onto a neutral substrate but a feature of the world itself that human beings encounter. (Taylor 2007) The Abrahamic traditions — Judaism, Christianity, Islam — introduced a theological transformation that, over centuries, systematically dismantled this enchantment from within. A single, absolute, and increasingly transcendent God replaced the population of immanent spirits and anthropomorphic deities. Iconoclasm — the prohibition of divine representation in material form — enforced the separation between the infinite and the finite. The natural world was desacralized: it became creation rather than manifestation, the work of God rather than the presence of God. Heaven was relocated from the mountain top or the sacred grove to a realm categorically distinct from

earthly experience. The god who walked in the garden of Eden with Adam, who wrestled with Jacob, who spoke directly from the burning bush — this god became progressively invisible, ineffable, and remote. By the time of the high scholastic theology, the divine had been rationalized into a metaphysical principle whose relationship to the material world was mediated by doctrine, institution, and text rather than direct encounter. This internal disenchantment created the structural conditions for secular modernity before secular modernity arrived. A world already emptied of intrinsic sacred presence — already understood as mere creation, as the stage on which salvation history unfolds rather than as itself divine — is a world available for mechanical description, for instrumental manipulation, for the kind of systematic exploitation that the scientific and industrial revolutions subsequently undertook. The Weberian argument is that Protestantism accelerated this by removing the last mediating institutions — the sacraments, the priesthood, the saints — that had maintained some residual enchantment within the Catholic world. The disenchanted world that secular modernity inherits was not created by secular modernity. It was prepared by centuries of theological development within the traditions that secular modernity subsequently displaced. (Weber 1905) The consequence for meaning infrastructure analysis is significant. When secular movements challenge religious belief, they are often challenging frameworks that have already completed most of the disenchantment work themselves. The transcendent God of developed monotheism is already a long distance from the animist world in which the sacred was everywhere immediate. The populations that encounter the secular challenge to religious meaning are not losing access to an enchanted world — they lost that access, institutionally and culturally, long before. What they are losing, when the transcendent framework itself becomes untenable, is

the last structure that organized meaning at all. The seeker who finds herself between the dead traditional gods and the discredited modern ones is not standing at the beginning of disenchantment. She is standing at its terminus — the point at which the long process of emptying, internal and external, has run to completion, and the question of what comes next has no available answer within the frameworks that produced the question. Divine Commands as Governing Narrative: The Unchallenged Authorizer Among the governing narratives analyzed in 8.4, divine commands occupy a structurally distinctive position. All governing narratives claim authority — they legitimate arrangements by embedding them in frameworks that make them appear necessary, natural, or just. But most governing narratives make their authority claim in terms that are, in principle, contestable within the social world they govern: the state claims authority through law, election, or historical right, all of which c

an be challenged on their own terms. Markets claim authority through efficiency, which can be measured and disputed. Traditions claim authority through continuity, which can be historicized. Divine commands claim authority through a source that is, by definition, not subject to human evaluation: if God has commanded it, the command does not require justification to human reason, and questioning it is not analytical scrutiny but impiety. This structural feature makes divine command claims among the most powerful and most dangerous instruments available for governing narrative construction. Powerful because they can motivate compliance and sacrifice at scales and with an intensity that secular authority structures rarely achieve: a soldier who believes he is fulfilling a divine obligation fights differently from one who believes he is serving a state interest. Dangerous because the immunity from human evaluation that divine authority claims is not actually available — the claim must be interpreted, transmitted, and enforced by human institutions that have their own interests, their own capture dynamics, and their own incentives to expand the scope of what counts as divine requirement. The genealogical question — who benefits from this specific interpretation of divine command, under what material conditions was it produced, and what institutional interests does it serve — is not an act of impiety. It is the minimum analytical requirement for understanding how divine authority claims function in practice. The analysis of divine command claims in governing narrative terms does not require, and the framework does not make, any claim about whether the commands are genuinely divine in origin. The genealogical analysis operates at the level of how the claim functions socially and politically — how it distributes authority, who controls its interpretation, what it makes available for mobilization, and what it makes immune from contestation. A divine command claim that, on genealogical examination, consistently serves the material interests of the interpreting class, that expands in scope whenever it is useful to those interests and contracts when it is not, and that systematically immunizes specific power arrangements from evaluation — that claim warrants analytical scrutiny regardless of its ultimate metaphysical status. The scrutiny is of the interpretation, the institution, and the consequences. It is not a verdict on the source. Analytical procedure for divine command claims: (1) Identify the specific institutional structure through which the command is interpreted and enforced — who has interpretive authority, how that authority is obtained and maintained, and what accountability mechanisms exist for its exercise. (2) Apply genealogical analysis: under what material conditions was this specific interpretation produced, and who benefits from it? (3) Evaluate consequences: what structural outcomes does consistent implementation of this command produce

for the populations subject to it? (4) Apply the Euthyphro standard: if the command is genuinely divine and therefore genuinely good, the consequences should be identifiable as good by the evaluative criterion — where they are not, the gap between claim and consequence requires explanation. (5) Distinguish the claim from the institution: a finding that the interpretive institution has captured the divine command claim for its own interests is not a finding about the validity of the claim itself — it is a finding about the institution. For the framework's purposes, eschatology in conflict zones is a variable that operates at three levels simultaneously. At the individual level, it provides the meaning infrastructure that makes sustained sacrifice coherent — which means that removing it without replacement produces the most severe form of the meaning infrastructure collapse described in 3.4b and 5.6, with the added dimension that the collapse occurs

in conditions of ongoing material conflict where psychological destabilization has immediate military and political consequences. At the organizational level, eschatological framing is a mobilization and discipline mechanism — it defines who is genuinely committed and who is performing commitment, and it provides internal enforcement through the logic of cosmic rather than merely organizational accountability. At the political level, it constrains the solution space available to leadership: a leader who accepts an arrangement that the eschatological framework prohibits has not made a pragmatic concession — they have committed apostasy, which in the logic of the framework is a more serious failure than losing. Analytical requirements for eschatological conflicts: (1) Identify which eschatological frameworks are active and what their internal logic requires — not what external observers think they require, but what follows from the tradition's own premises. (2) Map the physical and political claims that each framework treats as eschatologically mandatory vs. negotiable — the distinction determines what the actual solution space is. (3) Analyze who controls the authoritative interpretation of each framework — because eschatological claims are internally contested, and the contest over interpretation is itself a political struggle with material stakes. (4) Do not treat material interests and eschatological commitments as separate variables to be balanced — in these contexts they are structurally fused, and proposals that separate them will fail on eschatological grounds regardless of their material adequacy. (5) Recognize that the timeline of eschatological frameworks is not the timeline of secular political analysis — arrangements that are acceptable in the short term may be eschatologically prohibited as permanent settlements, which means the durability problem is not solved by agreement alone.

Finitude as Meaning Infrastructure: What Is Lost When Death and Scarcity Become Optional

The disenchantment anal

ysis above traces the trajectory through which the sacred withdrawal from the world produced populations oriented toward secular meaning frameworks. A parallel analysis is required for a different kind of withdrawal that is not yet complete but whose direction is visible: the withdrawal of finitude itself as an organizing condition of human existence. Death, scarcity, and the biological constraints that define the limits within which human life unfolds have not merely been obstacles that human beings navigate. They have been the structural conditions that generate the urgency within which meaning becomes possible. The knowledge that one will die, that resources are limited, that not all goods can be simultaneously obtained, and that time itself is the fundamental non-renewable resource — these conditions are not external impositions on a meaning-generating capacity that would otherwise function unchanged without them. They are among the primary materials from which meaning is constructed. This claim requires precision to avoid misreading. PMN does not hold that suffering is valuable because it produces meaning — that would invert the evaluative criterion that defines the framework. Suffering that forecloses becoming is structural harm regardless of whether it produces meaning as a byproduct. The death of a child from preventable disease does not become less of a harm because it generates grief that has meaning-producing properties. The claim is more specific: that the structural conditions of finitude — not the specific harms produced by specific instances of deprivation, disease, or death, but the general condition of existing within limits that cannot be engineered away — have been the organizing framework within which human meaning systems have operated across all cultures and historical periods, and that the progressive elimination of those conditions through transformative technology is not merely a material change but a meaning infrastructure change whose consequences cannot be predicted from within the framework that those conditions produced. Every major meaning tradition that human beings have developed — religious, philosophical, aesthetic, political — has been organized in part around the management of finitude. The great religions are not primarily ethical systems or social organizations, though they are those things too; they are primarily responses to the permanent conditions of mortality, suffering, and the limits of human knowledge and power. Philosophical traditions from Stoicism to Buddhism to existentialism are similarly organized around the question of how to live well within limits that cannot be overcome. Even secular political traditions derive much of their urgency from the finite character of human life — justice matters because people die, and their deaths in conditions of preventable suffering are not retrievable. A world in which death is optional, scarcity is

eliminated, and biological limits are negotiable is not a w

orld for which any of these traditions was designed, and the assumption that the meaning-generating capacity of those traditions will transfer intact to such a world is not supported by any evidence and is not a commitment PMN can endorse. A persistent misidentification is worth naming directly here. The biological foundation of this framework — organisms that suffer, require conditions for development, and organize their vulnerability through institutions — is read by some readers in religious traditions as a covert program for replacing religious commitment with scientific rationality. This is a misidentification of what the framework is doing. Sections 2.3 and 5.2–5.3 hold open metaphysical questions that scientific rationality would close; Part V treats religious traditions as genuine engagements with real phenomena rather than as errors awaiting correction. What the framework declines is the use of transcendent authority to insulate specific institutional arrangements from the material criterion — which is exactly what liberation theology, Maqasid-grounded reform, and engaged Buddhism declined in their own contexts. Incompatibility with religious commitment is not what this framework produces. Incompatibility with the institutional capture of religious authority is. The practical implication is not that transformative technology should be resisted in order to preserve the existential conditions that have historically produced meaning. That would be an error of the same kind as preserving oppressive institutional arrangements in order to preserve the meaning infrastructure embedded in them. The implication is that the transition toward conditions of reduced finitude is simultaneously a transition in the meaning infrastructure that human communities depend on, and that this transition requires explicit attention to meaning infrastructure reconstruction rather than the assumption that meaning will take care of itself once suffering is reduced. The seeker, analyzed in section 4.5, represents the individual-level form of this problem: the person who has lost the meaning framework produced by traditional finitude without having constructed a viable alternative. The coming generations, who will live in conditions of dramatically reduced finitude compared to all prior human history, face this problem at a civilizational scale. PMN does not have a solution to this problem. It has a description of what the problem is and a commitment to not pretending it will resolve itself. Methodological legitimacy: internal reform versus methodological appropriation The treatment of religion as a material variable with real consequences requires one further distinction that the framework cannot avoid: the difference between critique that operates from within a tradition's own methodology and critique that appropriates the tradition's authority while rejecting the methodology that makes that authority legitimate. This is not a defence of any tradition's current inst

itutional arrangements — every tradition's institutional hierarchy has been captured in ways that Part VI would predict and the framework requires exposing. It is a distinction about the epistemological standing of different kinds of challenge, which has direct consequences for whether a challenge can actually produce the internal reform it claims to be producing. Every tradition that has sustained itself across generations has developed methodology for adjudicating disputes about what the tradition requires. In Islamic jurisprudence, this is *usul al-fiqh*: the principles governing how evidence from Quran, *hadith*, *ijma*, and *qiyas* may be applied to derive rulings. In common law traditions, it is the doctrine of precedent and statutory interpretation. In empirical science, it is peer review, replication requirements, and the methodological standards that make claims revisable. These methodological frameworks are not mere bureaucratic gatekeeping. They are the accumulated institutional response to the problem that every tradition faces: how do you distinguish claims that the tradition genuinely supports from claims that someone wants the tradition to support? The methodology is what makes the tradition's authority something other than the preferences of whoever currently has the loudest voice. Methodological appropriation is the specific failure mode in which an actor enters a tradition's authority structure while rejecting its methodology selectively — accepting the authority when convenient and rejecting the methodology when it produces conclusions the actor dislikes. The liberal who enters Islamic jurisprudence and issues rulings while dismissing *isnad* criticism as mere traditionalism is not reforming the tradition from within. They are appropriating the tradition's social authority — the weight that comes from speaking in the name of a fourteen-century intellectual tradition —

while removing the methodological constraints that make that authority legitimate in the first place. The atheist who wants to be treated as a religious authority, the monarchist who opposes the monarch, the vegan who eats steak and still claims vegan community membership — these are variations on the same structural problem: the claim to the benefits of a identity without the obligations that constitute it. The framework's analysis of institutional legitimacy (6.6–6.7) applies here directly. An institution's authority rests on the credible claim that its outputs — its rulings, its interpretations, its conclusions — are produced by the methodology it claims to use. When the methodology is bypassed while the authority is retained, the institution has been captured in a specific way: not by external interests redirecting it toward their purposes, but by internal actors who want the authority without the accountability that methodology provides. The damage to the tradition is double. First, it degrades the methodology itself: if rulings can be produced without following the pr

cedure, the procedure loses its claim to be what produces valid rulings. Second, it undermines the tradition's capacity for genuine internal reform — which does exist and has always existed — by making it impossible to distinguish reform that the tradition's methodology can support from appropriation that it cannot. Genuine internal reform and methodological appropriation produce outputs that can look similar from the outside: both challenge existing institutional arrangements, both invoke the tradition's foundational commitments against its current practice, both generate resistance from those whose institutional position depends on the existing arrangements. The structural difference is methodological: the genuine reformer argues that the existing arrangements are not required — or are actively contradicted — by what the methodology, properly applied, actually yields. Liberation theology in Latin America did not import external socialist commitments into Catholic authority; it argued from scripture, from patristics, and from the tradition's own stated commitments that the preferential option for the poor was not an addition to what the tradition required but a recovery of what it had always required and institutional capture had obscured. The fence-jumper, by contrast, argues from conclusions backward to methodology: the tradition should support this position, therefore the methodology must be reinterpreted until it does. The direction of the argument is the diagnostic. The framework's reluctance to force methodological consistency — to demand that hybrid actors fully exit traditions they do not fully inhabit — is not an endorsement of methodological appropriation. It is a structural judgment about the costs of forced exit under conditions where functional alternatives are not available. A person who operates in partial adherence to a tradition while holding some commitments inconsistent with it is not necessarily damaging the tradition's methodology — they are navigating their own meaning system with whatever resources are available to them, which is not a structural problem unless they claim normative authority within the tradition on the basis of that partial adherence. The structural problem is not hybridity; it is the claim to authority without the methodological accountability that authority requires. Forced consistency that produces meaning-collapse serves no one. But tolerance for methodological appropriation that degrades the methodology through which traditions maintain their capacity for genuine self-revision is not tolerance — it is the slow dissolution of the institutional infrastructure that makes internal reform possible at all. The diagnostic for this section: is the challenge to an existing institutional arrangement arguing from the tradition's methodology toward a different conclusion, or arguing from a preferred conclusion toward a reinterpretation of the methodology? The first is the form that genuine internal reform consis

tently takes, across traditions and across historical periods. The second is the form that appropriation of institutional authority consistently takes. The difference matters for the framework's analysis not because methodological purity is intrinsically valuable but because the methodology is what allows a tradition to distinguish what it actually requires from what any particular actor wants it to require — which is the same function that the framework's own epistemological commitments serve in Part I. The structural implication for political analysis: any political project that addresses the material conditions of finitude — through medical technology, through automated abundance, through the elimination of the specific forms of suffering that have historically organized both meaning and politics — must simultaneously address the meaning

infrastructure that those conditions have sustained. This is not an argument for caution about material progress. It is an argument that material progress and meaning infrastructure development are not separable projects that can be pursued independently, with meaning assumed to follow from material improvement. The civilizations that have navigated rapid material transformation most successfully have been those that developed new meaning frameworks in parallel with new material conditions — not frameworks imported wholesale from the past, not frameworks constructed from nothing, but frameworks that drew on existing meaning traditions while adapting them to conditions those traditions had not anticipated. This is the work that PMN identifies as essential and that no political tradition has yet done adequately for the conditions that transformative technology is producing.

5.6 — Meaning as Material Condition

Meaning infrastructure — the systems of shared narrative, ritual, community, and interpretive framework through which populations make sense of suffering, organize obligation, and orient collective action — is a material variable in this framework's analysis, not an epiphenomenon of economic arrangements or a private psychological matter. The analysis in Parts I through IV operates at the level of material conditions, institutional arrangements, and the structural production of suffering. The preceding sections of Part V have treated the metaphysical dimension primarily as an analytical challenge — a domain that the materialist tradition misread and that requires more careful engagement. But there is a prior question that the framework must answer directly: why does the metaphysical dimension matter at all to a framework grounded in the biological foundation of Part III? The answer runs through the biological foundation rather than around it. The organisms analyzed in Part III are not simply pain-avoiding, resource-seeking systems. They are meaning-making systems. The capacity for meaning — the organization of experience into narratives that connect past suffering to future possibility, that locate individual lives within larger contexts, that make obligations feel like obligations rather than merely costs — is not a luxury added to the biological foundation. It is one of the mechanisms through which the biological foundation operates at the scale of sustained collective action. Suffering without meaning is structurally different from suffering with meaning. The same objective conditions — deprivation, loss, constraint — produce different behavioral and political responses depending on whether the suffering is interpreted as senseless or as part of a comprehensible trajectory. This is not only a psychological observation about individual resilience. It is a structural observation about collective action: populations that have a coherent account of why they suffer and what their suffering is oriented toward are more capable of sustained coordinated action under adversity than populations whose suffering is senseless. The history of movements that achieved durable structural change draws heavily on this difference — on the capacity to organize suffering within a framework that gives it direction rather than leaving it as undirected pain. (Gutiérrez 1971; Gramsci 1971) The implication for the framework's account of adaptive response capacity (4.5) is specific: the conditions for becoming are not only material in the narrow sense. A population with material sufficiency but no coherent account of what that sufficiency is for — no framework within which development acquires direction and obligation acquires weight — is not positioned for becoming in the sense the framework requires. Meaning infrastructure is part of the material infrastructure of becoming, not a separate add-on. The last man whom 4.5 identifies as the failure mode of welfare without be

coming is not only a product of insufficient material aspiration. He is also a product of the dissolution of meaning structures that gave suffering direction and development purpose, without adequate replacement. This does not require the framework to adopt any specific metaphysical commitment, or to treat any tradition's account of the source level as established. It requires attending to the meaning infrastructure of populations — the structures through which suffering is interpreted, obligation is generated, and development acquires direction — with the same analytical seriousness brought to their material infrastructure of resources, institutions, and power. A framework that ignores meaning infrastructure while

claiming to analyze the conditions for preserved adaptive response capacity has left out a variable that is consequential at the scale of historical transformation. A clarification the framework requires of itself here: noting that meaning has material consequences is not the same as claiming that meaning is only valuable because of those consequences, or that the reason meaning matters is instrumental. The argument in this section moves from the functional — meaning-making affects what people do, how they endure, whether collective action is possible — to the structural — meaning infrastructure is therefore a material condition that arrangements must maintain or replace. That argument is about consequences, not about the full nature of meaning. For the people whose meaning structures are being analyzed, meaning is not a tool for achieving stability or coordination. It is, for most of them, an orientation toward something they take to be real and worth orienting toward — regardless of its functional consequences. The framework's analytical treatment of meaning as a material variable is a methodological choice about how to analyze something, not a claim about what that thing ultimately is. Reducing the analysis to function does not reduce the phenomenon to function. This distinction matters because an analyst who forgets it will consistently misread populations whose meaning commitments are not instrumentally organized — which is most populations, most of the time. How meaning functions as a material condition: The materiality of meaning infrastructure operates through three specific mechanisms. Motivational: individuals who lack adequate meaning infrastructure show reduced capacity to sustain the kind of long-horizon action that counter-power accumulation requires — the willingness to bear costs in the present for outcomes in the future depends on a narrative about why those future outcomes matter. Cognitive: meaning infrastructure provides the interpretive frameworks through which structural conditions are understood as structural rather than as natural or individually caused — without those frameworks, the same material conditions produce political quiescence rather than organized response. Relational: shared meaning infrastructure i

s the medium through which collective action is coordinated — communities that share interpretive frameworks can act collectively on the basis of those frameworks without requiring explicit negotiation in each instance. All three mechanisms produce material consequences; all three are therefore within the scope of the framework's analysis. The implication for political analysis: a political project that neglects meaning infrastructure while building organizational and material infrastructure will encounter characteristic failure modes. Organizations that cannot explain to their members why the work matters beyond tactical objectives lose members when tactical progress slows. Counter-power campaigns that cannot articulate the structural causation of the conditions they are responding to cannot expand beyond the populations already convinced. The historical labour movement's investment in the culture of solidarity — the songs, the educational institutions, the mutual aid practices, the narrative of class interest — was not supplementary to the organizational work; it was the infrastructure that made the organizational work durable.

5.6b — Assessing Meaning Infrastructure: Functions, Fragility, and Intervention Implications

Section 5.6 establishes meaning infrastructure as a material condition — a set of institutions, relational practices, and interpretive frameworks that maintain psychological floor conditions, social cohesion, and the cognitive capacity for long-term collective action. Because meaning infrastructure cannot be observed directly as a single physical stock, its assessment requires evaluating its functional outputs and the structural vulnerabilities that emerge when those outputs degrade. Five functional dimensions define whether an existing infrastructure is operating above the threshold required to sustain adaptive response capacity. The first dimension is ritual participation density: the frequency, breadth, and social depth with which members of a population engage in shared practices that enact solidarity without requiring immediate transactional return. Ritual in this materialist sense is not decorative ceremony; it is the concrete coordination technology through which trust networks are maintained and social isolation is counteracted. The second dimension is narrative coherence and absorption capacity: the ability of the prevailing meaning framework to incorporate

novel material shocks, technological shifts, and historical dislocations without suffering explanatory collapse. A robust meaning framework metabolizes anomalies and changes in material conditions by updating its explanatory categories; a brittle framework responds to anomalies through rigid denial, persecution of dissent, or sudden fragmentation. The third dimension is grief and loss processing capacity: the availability of institutional and communal mechanisms that permit individuals and groups to register suffering, bereavement, and structural defeat without descending into ungrounded despair or privatized atomization. Where loss-processing channels fail, unintegrated suffering consistently converts into chronic resentment, hyper-vigilance, or catastrophic withdrawal. The fourth dimension is collective action and solidarity coordination: the capacity of shared narratives and moral vocabularies to overcome collective action problems by establishing mutual obligations that constrain immediate opportunistic defection. Meaning infrastructure provides the normative categories that make personal sacrifice for collective resilience rational to individual organisms. The fifth dimension is intergenerational transmission fidelity: the degree to which successive generational cohorts inherit, rearticulate, and willingly practice the core interpretive and relational habits of the community. When transmission fails, a generation emerges that lacks the interpretive categories necessary to operate existing institutions, producing rapid institutional decay even while formal economic stocks remain intact. These five dimensions interact dynamically and must be evaluated as a coupled system rather than in isolation. High ritual participation density combined with low narrative absorption capacity prod

uces brittle orthodoxy — an arrangement vulnerable to catastrophic collapse upon exposure to unprecedented material shocks. Conversely, high narrative diversity without transmission fidelity or collective action capacity produces hyper-fragmented atomization, in which individual meaning-seeking proliferates while public coordination capacity dissolves. To make these dimensions analytically operational, the framework applies three diagnostic lenses: (1) The Transmission Test: can younger cohorts explain the grounding of community obligations in their own vocabulary, or has the framework become ceremonial compliance? (2) Shock Absorption and Interpretive Adaptability: when a material crisis occurs, does the population mobilize cooperative responses, or do interpretive categories splinter into paralysis? (3) Harmful-Narrative Susceptibility: a critical distinction must be maintained between positive interpretive adaptability (the healthy capacity to update categories in response to evidence) and harmful-narrative susceptibility (the vulnerability of a depleted meaning infrastructure to totalitarian capture, scapegoating conspiracies, and violent out-group hostility). A population experiencing meaning depletion does not remain an open, neutral blank slate; it becomes acutely vulnerable to predatory explanatory systems that offer artificial coherence at the expense of target populations. The assessment yields four structural states: (1) Robust Infrastructure: all five dimensions functional; high adaptive response capacity preserved. (2) Brittle Infrastructure: high ritual compliance but near-zero absorption capacity; high catastrophic collapse risk under novel stressors. (3) Fragmented Infrastructure: high private meaning experimentation but zero intergenerational transmission or collective coordination; acute vulnerability to coordination failure. (4) Depleted or Collapsing Infrastructure: sub-threshold across multiple dimensions; widespread passive nihilism, status despair, and high susceptibility to predatory political mobilization. Operational diagnostics for meaning infrastructure assessment: (1) Transmission test — can the population articulate why its collective arrangements matter, and does that articulation transmit across generational boundaries? (2) Absorption test — does the meaning framework expand to make sense of novel material disruptions, or does it retreat into dogmatic denial? (3) Grief processing test — do institutional spaces exist to metabolize widespread loss, or is suffering privatized and pathologized? (4) Action test — can shared narratives motivate non-transactional collective sacrifice for common floor conditions? (5) Susceptibility test — is the population exhibiting symptoms of meaning depletion (anomie, conspiratorial drift, demonizing out-groups) that indicate imminent vulnerability to predatory ideological capture? → Formula: The three components of legitimacy (Performance L_p , Procedural L_r , Narrative L_n) and the struc

tural suffering formula in Part XV incorporate meaning infrastructure directly: degraded meaning infrastructure depresses Ln, inflates the perceived burden of suffering, and accelerates threshold collapse under lower objective stress levels.

5.6c — Acute Meaning-Collapse: Triage and Immediate Response

Sections 5.6 and 5.6b analyze meaning infrastructure as a material condition with long-run dynamics. This section addresses the acute condition: what happens when meaning infrastructure collapses suddenly under the impact of acute shocks — military invasion, catastrophic economic dislocation, sudden regime collapse, or ecological disaster — and what response priorities the framework indicates. Leading indicators of acute meaning-collapse (any three present simultaneously signal active or imminent acute collapse): (1) sudden loss of central symbolic authority that previously organized collective interpretation — regime change, death of a founding leader, collapse of an institution that provided shared reference points; (2) rapid spread of mutually contradictory factual claims at a scale that prevents verification through normal channels; (3) sharp increase in reports of disorientation, loss of purpose, or inability to situate individual experience within a collective frame, accessible through health workers, community leaders, or social media distress signals; (4) breakdown of ordinary coordination practices that depend on shared assumptions about the future — markets, public transport, community events, religious observance; (5) dramatic increase in inter-group conflict absent a clear material trigger, indicating groups are filling interpretive vacuums with available threat narratives. Immediate response priority order: Priority 1 — Material security. If physical safety, food, shelter, or medical access are collapsing simultaneously with meaning infrastructure, address material conditions first. Meaning infrastructure cannot be reconstructed in populations whose survival needs are unmet, and meaning-reconstruction efforts in conditions of acute material deprivation are absorbed into the material crisis rather than functioning as genuine meaning provision. Priority 2 — Amplify residual meaning carriers. Before introducing new narratives, identify and amplify what remains functioning: local leaders whose authority has not been destroyed by the collapse event, existing practices that can continue under disrupted conditions, shared memories and reference points that predate the collapse trigger. Residual carriers are more credible and more efficiently absorbed than external meaning imports because they connect to existing identity infrastructure rather than requiring construction from scratch. Priority 3 — Deploy bridge narratives when residual carriers are insufficient. A bridge narrative is a temporary interpretive frame designed to reduce acute disorientation without displacing longer-term meaning reconstruction. Bridge narratives must meet three conditions: accurate about current conditions (bridge narratives built on false claims collapse when the claims fail, deepening disorientation); connected to existing meaning frameworks the population holds (they bridge from the known to the disrupted, not from an external position); and explicitly provis

ional (a bridge narrative presented as permanent forecloses the reconstruction process it is supposed to enable). The test: does it reduce disorientation while keeping the interpretive space open, or does it reduce disorientation by closing down interpretation? Priority 4 — Long-run infrastructure building. Once material security is partially stabilized and acute disorientation is reduced, shift resources toward the long-run infrastructure investments described in 5.6. Case: post-invasion displacement, Ukraine 2022. The rapid displacement of millions in the first weeks of the Russian invasion created acute meaning-collapse conditions for specific populations: those who experienced sudden loss of home and community while losing the institutional frameworks — local government, employment, social networks — through which they organized daily meaning. The leading indicators above were simultaneously present. Effective immediate responses were consistent with the priority order: humanitarian organizations addressed material needs first; local community leaders and clergy (residual carriers, priority 2) provided orientation more effectively than external narrative imports; bridge narratives emphasizing continuity of identity and the temporary nature of displacement (priority 3) were more effective where they were accurate and explicitly provisional than where

they were triumphal or denied the severity of displacement. The longer-term meaning infrastructure work — community network reconstruction, cultural preservation, documentation — was most effective where the acute phase had been adequately addressed first.

5.7 — Christian Political Theology: Augustine, Aquinas, and Catholic Social Teaching

Christian political theology is one of the oldest and most developed intellectual traditions in the engagement between religious commitment and political analysis, and its analytical contributions to the questions PMN addresses are more substantive than either secular dismissal or uncritical acceptance acknowledges. The tradition spans fifteen centuries, internal divisions as significant as those between Christian denominations, and a range of political applications from imperial legitimation to liberation theology. What follows is not a survey of Christian political thought in its entirety but a diagnostic engagement with the analytical contributions and structural liabilities of three formations that have been most consequential for the questions PMN takes as central: Augustine's political realism, Aquinas's natural law tradition, and the Catholic Social Teaching tradition from Leo XIII through Vatican II and liberation theology. Augustine: Political Realism and the Two Cities Augustine's political theology, developed primarily in *The City of God* (426 CE) in response to the sack of Rome and the claim that Christianity had weakened Roman political capacity, is the foundational text of Christian political realism. Augustine's central analytical move is the distinction between the *civitas terrena* (earthly city) and the *civitas Dei* (city of God): not a spatial or institutional distinction between church and state, but a distinction between two orientations of love — love of self to the contempt of God, and love of God to the contempt of self — that are intermixed in all actual human communities. The political consequence of this distinction is that no earthly political arrangement can achieve justice fully: all actual political communities are composed of persons with mixed and partial loves, which means they can produce at best a partial and temporary order, never the complete justice that only the heavenly city can provide. This is a form of what PMN's framework identifies as the permanent tension between idealism and pragmatism — but derived from a specific theological anthropology rather than from material analysis. Augustine's claim that all earthly political arrangements are characterized by *libido dominandi* (the lust for domination) that corrupts even well-intentioned governance is structurally parallel to PMN's custodian problem (7.3): every anchor generates the incentive to define it in favor of whoever holds it. The difference is that Augustine grounds this structural tendency in the fallen condition of human nature (original sin) while PMN grounds it in the incentive structures that institutional positions generate regardless of the character of those who occupy them. Both analyses reach a similar practical conclusion — no institutional arrangement can fully prevent the corruption of power by those who hold it — through fundamentally different foundational commitments. Augustine's consequent political realism is directly relevant to PMN'

s analysis. Augustine does not conclude from the impossibility of earthly justice that political engagement is pointless. He argues that Christians should support political arrangements that provide the relative peace (*pax terrena*) necessary for the life of the earthly city, not because that peace is ultimately good but because its absence produces conditions in which no meaningful life is possible. This is a floor-criterion argument: the minimal threshold below which no arrangement can be acceptable is the provision of sufficient order for human life to proceed. PMN's analysis of structural suffering as the primary diagnostic of institutional failure is in this sense an elaboration of Augustine's insight in materialist rather than theological terms: the question is not whether arrangements achieve justice completely but whether they meet the floor condition below which human functioning is foreclosed. The limit of Augustine's analysis for PMN's purposes is its tendency toward political quietism: because earthly justice is unrealizable and earthly peace is all that can reasonably be sought, Augustine provides limited resources for structural critique of specific arrangements or for identifying what better arrangements would look like. The Two Cities distinction can function as a form of ideological insulation: any institutional failure can be absorbed as evidence of the earthly city's necessary

imperfection, rather than as evidence that specific arrangements should be changed. This is the failure mode PMN's epistemological standard identifies in Part I: a framework that explains its own failures as features of the larger process rather than revising its conclusions. Aquinas: Natural Law, Common Good, and the Critique of Injustice Thomas Aquinas's political theology, developed primarily in the *Summa Theologiae* and the *De Regno*, represents the most systematic attempt within the Christian tradition to ground political analysis in a framework of natural law accessible to reason without requiring revealed theology. For Aquinas, the natural law is the rational creature's participation in the eternal law of God — the ordering of creation toward its proper ends. Political arrangements are legitimate when they are ordered toward the *bonum commune* (common good), defined as the conditions under which all members of the community can flourish according to their nature. Arrangements that serve the private good of rulers at the expense of the common good are tyrannies, and Aquinas defends both the legitimacy of resistance to tyranny and the claim that unjust laws are not properly laws at all. The Thomistic common good concept is analytically distinct from utilitarian aggregate welfare in ways that parallel PMN's own distinction. For Aquinas, the common good is not the sum of individual goods but the conditions under which individual goods can be achieved — a structural rather than aggregative concept. The common good requires that everyone, not just the majori

ty, have access to the conditions for flourishing; this is a floor-criterion argument with distributional content that aggregate welfare lacks. Aquinas also insists that temporal goods — material provision, security, justice — are preconditions for higher goods rather than goods in themselves, which produces a hierarchy of goods parallel to the Maqasid's *daruriyyat/hajiyyat/tahsiniyyat* structure and to PMN's floor/becoming distinction. The natural law tradition's analytical contribution to PMN is the concept of intrinsic human dignity as a structural constraint on any evaluative framework. For Aquinas, human beings have a dignity that derives from their nature as rational animals made in the image of God, and this dignity cannot be traded away for aggregate benefits or social utility. This is a deontological constraint that PMN's consequentialist-adjacent framework does not fully articulate: the minimal anchor says that structural suffering must be reduced, but it does not specify whether this is because suffering is intrinsically bad (which would import a value) or because it degrades social systems (the stability claim). Aquinas provides a third option: suffering that violates human dignity is wrong not because of its aggregate consequences but because it treats persons as means to other ends rather than as ends in themselves. This is a distinct moral claim that PMN's current framework acknowledges but does not fully integrate. The institutional analysis of Aquinas also contains resources PMN should engage. Aquinas's account of the best regime — a mixed constitution combining monarchy, aristocracy, and democracy to prevent the characteristic failure modes of each — is an early formulation of the distributed custodianship principle (7.3). His analysis of law as having authority only when it serves the common good, with the implication that unjust laws lack binding moral authority, provides a theological grounding for civil disobedience that PMN's framework implies structurally but does not derive from natural law tradition. The Thomistic just war tradition — with its criteria of legitimate authority, just cause, right intention, last resort, proportionality, and reasonable chance of success — is the most developed pre-modern framework for evaluating organized violence against the structural criterion, and its relationship to PMN's analysis of when and how force can be legitimate in structural transformation is an open and important question. Catholic Social Teaching: From *Rerum Novarum* to Liberation Theology The Catholic Social Teaching tradition, beginning formally with Leo XIII's *Rerum Novarum* (1891) and continuing through *Quadragesimo Anno* (1931), *Pacem in Terris* (1963), *Gaudium et Spes* (1965), and the subsequent liberation theology movement, represents the most sustained engagement within any religious tradition between theological commitment and structural analysis of economic and political arrangements. The tradition's core principles — hu

man dignity, solidarity, subsidiarity, the universal destination of goods, and the preferential option for the poor — constitute an evaluative framework that has substantive overlaps with PMN's own. The principle of the universal destination of goods — articulated by Aquinas and developed throughout Catholic Social Teaching — holds that the material goods of the earth are destined for all human beings, and that private

property is a legitimate institution only to the extent that it serves this universal destination. This is not a communist claim about collective ownership but a structural claim about the moral condition of private property: it is legitimate when it distributes material conditions adequate for human dignity, and illegitimate when it concentrates them in ways that deny those conditions to others. Leo XIII's *Rerum Novarum* deployed this principle to critique both unregulated capitalism and Marxist socialism simultaneously: capitalism for producing structural poverty that violated human dignity; socialism for eliminating the private property that provided workers with independent material conditions. This double critique anticipates PMN's own position that no economic arrangement is acceptable in advance — that the criterion is what arrangements actually produce for the populations subject to them. The subsidiarity principle — the claim that decisions should be made at the lowest level of social organization capable of addressing the relevant problem, with higher levels intervening only when lower levels cannot — is the most analytically interesting Catholic Social Teaching contribution for PMN's institutional analysis. Subsidiarity is a structural principle for distributing governance authority that addresses the scale/accountability tension identified in PMN's section 13.2 from a theological direction. The principle holds that concentrated governance authority is bad not merely because it is inefficient but because it violates the proper ordering of human communities: the intermediate institutions of family, neighborhood, and civic association are not merely instrumental to the state but have intrinsic value as contexts in which human beings develop their capacities and exercise genuine agency. This is an institutional articulation of PMN's becoming criterion applied to governance architecture: the question is not only what arrangements deliver materially but what arrangements create the conditions for genuine human development at every level. Solidarity — the principle that human beings are genuinely interdependent and that this interdependence generates moral obligations across social, national, and generational boundaries — provides the Catholic Social Teaching framework for extending moral concern beyond the immediate community. John Paul II's development of solidarity in *Sollicitudo Rei Socialis* (1987) explicitly connects it to structural analysis: solidarity is not primarily a feeling of compassion but a 'firm and persevering

determination to commit oneself to the common good' that requires structural change, not only individual charity. This is a theological articulation of PMN's own distinction between addressing the immediate symptoms of structural suffering and addressing the institutional arrangements that produce it. Liberation theology — developed by Gutiérrez, Boff, Sobrino, and others in the Latin American context from the 1960s onward — represents the internal tradition's most radical development of the preferential option for the poor as a structural rather than merely charitable principle. Gutiérrez's central analytical move is the distinction between liberating from immediate poverty (charity) and liberating from the structural conditions that produce poverty (structural transformation). The former leaves the productive arrangements that generate poverty intact while alleviating some of its consequences; the latter requires engagement with the institutional arrangements that produce the conditions being addressed. This distinction maps directly onto PMN's analysis of R (resource deprivation as produced by structural arrangements) and the inadequacy of symptomatic responses to structural causes. Liberation theology's development of 'social sin' — the concept that institutional arrangements rather than only individuals can be bearers of moral failure — is the theological translation of PMN's structural suffering concept: the analysis of how institutional arrangements systematically produce deprivation independently of any individual's intentions. The Vatican's condemnation of liberation theology through the Congregation for the Doctrine of the Faith under Ratzinger (1984, 1986) is itself analytically significant from PMN's perspective. The condemnation's primary objection was not to the principle of concern for the poor but to the use of Marxist class analysis as an analytical instrument — the claim that Marxist categories distort theological reasoning and lead toward materialism incompatible with Christian anthropology. This is a case study in the tension between institutional self-preservation (the Church's identification of Marxist analysis as an existential threat to its institutional interests) and the substantive application of the tradition's own evaluative commitments (the preferential option for the poor). From PMN's analysis of institutional capture (7.3b), the condemnation

represents the institutional hierarchy prioritizing its own continuation over the criterion the institution was built to serve — precisely the failure mode the custodian problem analysis predicts. Christian political theology’s contribution to PMN across these three formations: Augustine provides political realism grounded in a different anthropological foundation than PMN’s material analysis but reaching structurally parallel conclusions about the impossibility of perfect justice and the permanent relevance of the custodian problem. The limitation is its tendency toward quietism

that PMN’s inaction analysis (1.5) specifically rejects. Aquinas provides the *bonum commune* as a structural rather than aggregative concept of the common good, the intrinsic dignity constraint as a deontological complement to PMN’s consequentialist-adjacent framework, and the distributed governance principle that anticipates PMN’s institutional analysis. Catholic Social Teaching provides subsidiarity as a governance architecture principle, solidarity as a structural obligation extending beyond immediate community, the universal destination of goods as a criterion for evaluating property arrangements, and liberation theology’s social sin concept as a theological translation of structural suffering. The structural question PMN poses to the tradition as a whole: what institutional mechanisms within the Church itself prevent the organizational hierarchy from capturing the evaluative commitments of the tradition and redirecting them toward institutional self-preservation? The liberation theology condemnation suggests that this question has not been adequately answered.

5.8 — Al-Farabi and the Islamic Political Philosophy Tradition

Al-Farabi (872–950 CE), known in the Islamic tradition as the ‘Second Teacher’ (the first being Aristotle), produced the most systematic synthesis of Greek political philosophy and Islamic thought available in the medieval period, and his work raises specific analytical questions that PMN must engage rather than merely acknowledge. The primary texts — *Al-Madina al-Fadila* (The Virtuous City), *Al-Siyasa al-Madaniyya* (Political Regime), and *Ara’ Ahl al-Madina al-Fadila* (Opinions of the Inhabitants of the Virtuous City) — constitute a political philosophy that takes the organization of human communities toward their highest possible flourishing as its central question. This is a becoming-centered political philosophy, not a floor-centered one: Al-Farabi’s question is not what institutional arrangements prevent the worst outcomes but what arrangements make the best outcomes available to human beings organized in community. Al-Farabi’s central concept is the *faṣīla* — the virtuous or excellent: the city, the community, the person who has developed their capacities to their fullest possible expression. The virtuous city is one organized so that its members can achieve *saʿāda* (happiness, flourishing, *eudaimonia* in Aristotle’s sense), which for Al-Farabi includes both the philosophical happiness of the intellect engaged with ultimate truths and the practical happiness of the person whose political community provides the conditions for genuine development. The analytical parallel to PMN’s becoming criterion (4.5) is precise: both frameworks identify the highest aspiration of political organization as the creation of conditions under which human beings can genuinely develop beyond what survival and material security alone provide. Both distinguish this from arrangements that merely prevent suffering without creating the conditions for development. And both identify the institutional design question as primary: what forms of political organization actually produce these conditions, rather than what forms merely claim to? Al-Farabi’s analysis of the ‘imperfect cities’ — the ignorant city (*al-madina al-jahiliyya*), the immoral city (*al-madina al-fasika*), the erring city (*al-madina al-dalla*), and others — is a typology of institutional failure that is analytically equivalent to PMN’s regime type analysis in 7.8. Each imperfect city is defined by what it mistakes for the highest good: the ignorant city takes material wealth, bodily pleasure, or domination as its organizing principle rather than genuine flourishing. The immoral city knows what genuine flourishing requires but pursues other ends in practice — this is the distinction between claimed function and actual function that PMN’s institutional betrayal variable (B) tracks. The erring city has adopted incorrect beliefs about what the ultimate good is and organizes itself toward a misidentified end. Each failure mode produces characteristic forms of structural suffering that

Al-Farabi's analysis identifies thr

ough different vocabulary than PMN but with equivalent structural precision. The most analytically challenging element of Al-Farabi's political philosophy for PMN is his theory of the *ra'is* — the supreme ruler of the virtuous city, who must combine philosophical wisdom with practical political capacity. Al-Farabi requires that this ruler have genuine knowledge of ultimate truths (philosophy, and for Al-Farabi this ultimately connects to prophetic revelation) combined with the practical wisdom to translate those truths into institutional arrangements adequate to the actual conditions of the community. This is a theory of governance that requires both analytical capacity and contextual judgment simultaneously — neither technocracy (pure analytical capacity without contextual judgment) nor populism (contextual responsiveness without analytical framework). PMN's own position on governance — rejecting technocratic authority while insisting that complex systems require analytical capacity to manage — faces exactly this tension, and Al-Farabi's attempt to resolve it through a theory of the philosophically qualified ruler is analytically interesting even where its specific resolution (requiring prophetic qualification) is not available to a secular framework. Ibn Rushd (Averroes, 1126–1198), Al-Farabi's most important interpreter and the primary transmitter of Aristotle to Latin Europe, extends this tradition by insisting more strongly on the separability of philosophical from religious knowledge — that reason operating on its own terms can reach genuine truth about political and ethical questions without requiring religious grounding. This is the position that PMN's own epistemological framework occupies: the framework's conclusions are derived from reason and material analysis rather than from religious commitment, while remaining genuinely open to the interface-level phenomena that religious traditions engage (Part V, 5.2–5.3). Ibn Rushd's defense of philosophical reasoning against Al-Ghazali's critique in the *Tahafut al-Tahafut* (The Incoherence of the Incoherence) is the historical precedent for PMN's own methodological position: reason can address the questions that matter for political and ethical life, and religious commitment neither replaces nor refutes this capacity. Al-Farabi and Ibn Rushd's contribution to PMN: first, a becoming-centered political philosophy developed from within the Islamic tradition that independently identifies the highest question of political organization as the creation of conditions for genuine human flourishing — parallel to PMN's 4.5 but derived from different foundational commitments. Second, a typology of institutional failure (the imperfect cities) that is analytically equivalent to regime type analysis and maps onto PMN's variables with structural precision. Third, the *ra'is* theory as an attempt to resolve the tension between analytical capacity and contextual judgment in governance — the same tension PMN's anti-

technocratic guardrail addresses without fully resolving. Fourth, Ibn Rushd's defense of philosophical reasoning as a historical precedent for PMN's own epistemological position: reason operating on material reality can address political and ethical questions without requiring religious grounding, while remaining genuinely open to what religious traditions identify as significant phenomena.

VI: Power: The Variable That Cannot Be Left Implicit

The supremacy of a social group manifests itself in two ways: as domination and as intellectual and moral leadership. (Gramsci 1971, p. 57)

6.1 — Why Power Must Be Named

A framework that analyzes institutions, material conditions, and systemic outcomes without an explicit account of power is not a serious framework. It is a technical exercise that produces conclusions favorable to whoever currently holds power — because power is what determines which material conditions persist, which outcomes are treated as legitimate, and which forms of suffering are visible enough to be counted in the analysis. Leaving power implicit is not neutral. It is the most consequential form of analytical failure available to a framework of this kind. The materialist tradition understood this. One of Marx's most important contributions was the demonstration that economic power translates into political power, which translates into ideological legitimacy, which reinforces economic power — a cycle that is self-reinforcing precisely because each stage makes the previous stage appear natural rather than contingent. Gramsci extended this analysis to the domain of culture and consent, showing that the most durable form of power is not coercive but hegemonic — the power to define what counts as normal, reasonable, and desirable, so that alternatives become literally unthinkable for large portions of the population. Both analyses are inherited here and extended. (Gramsci 1971) Leaving power implicit has a specific practical consequence: it makes the framework's conclusions systematically easier to reach when they align with the interests of those who currently hold power, and harder to reach when they do not. A framework whose conclusions happen to be more readily available to those with resources, institutional access, and information advantages is not neutral by virtue of not mentioning power. It is biased in favor of the powerful, by structure, regardless of its stated values. Naming power — as a variable, as a condition of analysis, as a factor that shapes what conclusions can be reached and by whom — is the minimum required to avoid this structural bias. It does not eliminate the bias. It makes it visible enough to address. The failure to name power is not a neutral omission. Every analytical framework that operates without an explicit account of power implicitly inherits whatever account of power is embedded in the common sense of the society within which it operates — and that common sense is not neutral either. It reflects the interests of whoever has been most effective at shaping what counts as obvious, natural, and beyond question. A framework that treats current distributions of resources, authority, and influence as background conditions rather than as variables requiring explanation has already made a substantive political commitment: to the naturalization of those distributions. What power actually is in this framework: Power is the differential capacity to shape outcomes — to determine what happens, to define what counts as legitimate, to set the terms within which others operate. This definition is intentionally structural rather

than agential: it does not require that powerful actors consciously exercise power, only that their structural position gives them greater capacity to shape outcomes than others in the same system. A landlord has power over a tenant not because the landlord is malicious but because the structural position of property ownership versus property dependence produces differential capacity to set terms. An employer has power over employees not because of personal domination but because the structural asymmetry between control of employment and dependence on employment produces differential capacity to determine the conditions of work. The power is in the position, not the person. The three forms that matter most: For the purposes of this framework, three forms of power are analytically central. Structural power is the capacity that flows from position in an institutional arrangement — property ownership, organizational authority, access to information, control of enforcement mechanisms. Narrative power is the capacity to define what is visible, what is natural, what is possible, and what counts as a legitimate claim — the power that operates through governing narratives analyzed in 8.4. Epistemic power is the capacity to determine whose knowledge counts,

which methods of knowing are authoritative, and what evidence is required to establish a claim — the power analyzed in Part I's account of epistemic authority and infrastructure. These three forms interact: structural power provides resources to produce and maintain narrative power; narrative power makes structural power appear natural and therefore legitimate; epistemic power determines what evidence can challenge either. The analysis of any specific arrangement requires tracking all three forms simultaneously. Why leaving it implicit is specifically dangerous: Frameworks that leave power implicit tend to produce a characteristic analytical error: they explain outcomes that are produced by power asymmetry as outcomes produced by technical factors — by market efficiency, by institutional design, by cultural norms — in ways that obscure the power dynamics that maintain those technical factors in place. Healthcare systems that systematically produce worse outcomes for poor populations are analyzed as problems of access, incentive design, or information asymmetry — all of which are real — without asking whose power maintains the access barriers, who benefits from the incentive structures that produce them, and whose epistemic authority determines what counts as evidence of systematic failure. The technical analysis is not wrong; it is incomplete in a direction that consistently favors those whose structural position is maintained by the arrangements being analyzed. The naming requirement is also a calibration requirement. Naming power explicitly forces the analyst to specify what kind of power is operating, through what mechanisms, maintained by what arrangements, and subject to what counter-pressures

. This specification produces more precise analysis than the alternative — vague attribution of outcomes to "the system" or to "structural forces" without identifying the specific structural features that produce them. Vague structural attribution is not more radical than precise analysis; it is less useful, because it cannot point toward specific interventions at the level where the power actually operates. The connection to suffering: power asymmetry is one of the primary mechanisms through which unnecessary suffering is produced and maintained. The minimal anchor generates an evaluative demand on any arrangement; power analysis specifies the mechanism through which arrangements that fail the minimal anchor criterion persist despite that failure. Without the power analysis, the ethical demand has no structural account of why conditions that should change do not change. With it, the analysis can identify not only that an arrangement is failing but who has the structural capacity to maintain it in its current form and what conditions would need to change to reduce that capacity. That connection — between the evaluative criterion and the structural account of how failure persists — is what makes Part VI indispensable rather than supplementary.

6.2 — Power as Structural Position

Power in this framework is not primarily a property of individuals. It is a property of structural positions within systems. The person who occupies a position of power does not need to be malicious, or even particularly strategic, to exercise that power in ways that are harmful. The structural incentives of the position shape behavior predictably and consistently, regardless of the intentions of the occupant. This is what makes the analysis of power useful: it generates predictions about behavior that do not depend on knowing the character of specific actors. The corollary is equally important: replacing the person in a position of power, without changing the structural conditions that generate the position's incentives, will not change the pattern of behavior that the position produces. This is the error that naive approaches to political change invariably make — treating the problem as one of personnel rather than structure. The same error allows institutions to absorb reform by replacing individuals while leaving the structural conditions that produced the problem entirely intact. Analyzing power as structural position has a second implication that runs against a common instinct in political analysis: it removes the moral weight from the attribution of power. A person who holds a position of power and exercises its incentives is not, by virtue of that, more culpable than a person who would exercise the same incentives in the same position. Individual responsibility remains real. The point is that the most analytically productive question is not 'who is bad?' but 'what structural conditions produce this pattern of behavior, and what would have to change to produce a

different pattern?’ That question points toward the transformation of conditions rather than the replacement of personnel — which is the level at which durable change actually occurs. Structural position also generates power through the control of information flows. Actors who occupy positions through which information must pass — regulatory agencies, financial intermediaries, platform operators, credentialing bodies — acquire leverage that is independent of any formal authority they hold. The gatekeeping function is itself a power resource, and it becomes more significant as the systems it operates within grow more complex and less legible to those outside them. Complexity is not neutral: it systematically advantages the actors who generate it and disadvantages the actors who must navigate it without equivalent resources. This is the mechanism by which technically complex domains — finance, healthcare administration, legal procedure, infrastructure procurement — tend toward regulatory capture even when the regulatory bodies are formally independent and individually staffed by people of ordinary integrity. There is a further implication that runs against common assumptions about coalition building. If power is structural rather than personal, then acquiring political allies from am

ong those who hold positions of structural power is not the same as acquiring allies for structural change. An actor who occupies a position of structural power and who sincerely supports progressive outcomes is still an actor whose position generates incentives that pull against those outcomes. The alliance is real; the structural pressure on it is also real. This is not a reason to refuse such alliances — it is a reason to design accountability mechanisms that operate independently of the intentions of whoever occupies the position at any given moment. The failure mode is not insincerity; it is structural drift that neither party to the alliance may fully recognize until it has already occurred. The diagnostic test for whether analysis treats power as structural: does it generate the same predictions about behavior regardless of who occupies the position, and does it locate the explanation for behavioral patterns in the incentives of the position rather than the character of its occupants? If the analysis depends on identifying particular individuals as unusually corrupt or unusually committed, it is treating power personally — and will systematically underestimate the stability of patterns that survive changes in personnel. The analytical consequence: Structural power analysis generates predictions that individual-level analysis cannot. If power follows from structural position rather than from individual characteristics, then replacing one occupant of a position with another who has different stated values or different personal history will not change the behavior produced by the position — not because individuals do not matter, but because the structural incentives of the position shape what behaviors are rewarded, what information reaches the occupant, and what coalition maintains them in position. This is the mechanism behind the custodian problem analyzed in 7.3: revolutionary actors who successfully displace those they were opposing consistently begin to reproduce the behaviors of their predecessors not because they were secretly always the same but because the structural incentives of the position they occupy are the same. Understanding this mechanism is the precondition for designing institutions that actually constrain position-holders rather than merely replacing them. Three forms of structural power: Positional power derives from occupancy of a node in an institutional network through which resources, information, or decisions must flow — the power to approve, veto, delay, or redirect that flows from being in the path that others must traverse. Relational power derives from the network of obligations, loyalties, and dependencies that an actor has accumulated — the power that comes from being the node through which many other actors’ interests are connected. Agenda power derives from the capacity to determine what questions are asked and what alternatives are considered — power exercised not by making decisions but by shaping what de

cisions are available to be made. All three forms are structural rather than individual: they are properties of positions and networks, not of persons, and they can be analyzed independently of who currently occupies them. The political implication: structural power analysis redirects attention from the characteristics of individual power-holders to the design of institutional arrangements that distribute power across positions and constrain what any single position can do. Focusing on who holds power misses the structural analysis; focusing on the institutional design that makes power accountable is the analytically productive move.

6.3 — Asymmetry, Capture, and Information

Power asymmetry, institutional capture, and information manipulation are the three dynamics most consequential for what follows — and the three most systematically underweighted in political analysis that takes formal institutions at face value. Power asymmetry is the condition in which different actors within a system have significantly different capacities to shape outcomes, define terms, and determine what counts as a legitimate claim. Asymmetry is the normal condition of real systems, not a deviation from some baseline of equality. The analysis of any specific situation must begin by mapping the actual distribution of power — who can do what to whom, under what conditions, with what costs — rather than assuming a baseline of rough equivalence that does not exist. Institutional capture is the process by which actors with concentrated interests gain control of institutions that are supposed to serve diffuse interests, and redirect those institutions toward their own benefit while maintaining the appearance of serving their stated purpose. Capture is not a sudden event. It is a gradual process driven by the systematic advantage that concentrated interests have over diffuse ones in any sustained political contest: concentrated interests have more to gain from specific outcomes, bear the cost of organizing more easily, and can sustain coordinated action over longer time horizons than the populations whose diffuse interests the institution is supposed to protect. (Olson 1965; Stigler 1971) Information manipulation is the use of control over information flows — what is visible, what is invisible, what is framed as normal, what is framed as exceptional — to shape the beliefs and behavior of actors who would behave differently with a more accurate picture of their situation. Information manipulation is not limited to obvious propaganda. The most effective information manipulation operates through the selection of what is considered newsworthy, what is included in educational curricula, what problems are defined as requiring political solutions and which are defined as natural or inevitable. Gramsci's concept of hegemony is the most developed analysis of this dynamic available, and the framework draws on it directly. Asymmetry as normal condition: The starting point for analysis is not a baseline of equal power that has been distorted, but a recognition that power asymmetry is the normal condition of institutional systems. Systems without power asymmetry are theoretical constructs, not actual institutional arrangements. Every real system distributes the capacity to shape outcomes unequally, because the structural positions within any institutional arrangement are themselves unequal: some positions have greater access to information, greater authority over enforcement, greater capacity to define the terms of participation, and greater resources to deploy in defense of their position. The analytical question is never whether asymmetry exists — it always

does — but what its specific form is, how it is maintained, and what consequences it produces for those at different positions within it. The capture sequence in detail: Institutional capture — the process through which institutions built to serve a function are progressively redirected to serve the interests of the actors best positioned to influence them — operates through a consistent sequence. Stage one: the institution achieves some success at its original function, which gives it standing, resources, and the institutional relationships required to continue. Stage two: the actors with the most resources and the most direct interest in the institution's operation — typically the actors whose activities the institution was built to regulate, oversee, or constrain — invest in influencing the institution's operations because the returns on that investment are high. Stage three: the institution's decision-making is progressively shaped by those most invested in it, not through single dramatic acts of corruption but through the accumulation of personnel decisions, procedural changes, and informal norms that favor the perspectives and interests of the actors with greatest access. Stage four: the institution continues to perform its original function in form while having substantively changed what it optimizes for. The regulatory agency that writes rules favoring the industry it regulates, the union that negotiates on behalf of its permanent staff rather than its membership, the NGO that designs programs around donor priorities rather than community needs — these are all stage-four capture institutions. Information as power resource: Information asymmetry is a power resource because the actor with superior information about how a system operates can structure their engagement with it to their advantage in ways that others cannot detect, monitor, or respond to. This operates at every level of institutional life. At the

regulatory level: regulated industries consistently know more about their own operations than regulators do, and that informational asymmetry shapes what regulations are written, how they are enforced, and where the gaps are that allow non-compliance without technical violation. At the organizational level: managers know more about organizational operations than workers, clients, or oversight bodies, and that asymmetry shapes the distribution of risk and reward within the organization. At the political level: incumbents know more about what government actually does than challengers, and that asymmetry shapes electoral competition in ways that systematically advantage those already in power. Addressing power asymmetry without addressing the information structures that maintain it produces reforms that change the formal distribution of authority while leaving the substantive capacity to shape outcomes unchanged. The three dynamics — asymmetry, capture, and information — interact as a mutually reinforcing system. Structural power asymmetr

y provides the resources to invest in capture. Capture redirects institutional information flows to favor the capturing actors. Informational asymmetry makes capture harder to detect and easier to maintain. Together they constitute the normal condition of institutional systems under pressure from actors with concentrated interests and resources — which is to say, under normal operating conditions in any institutional domain where significant material interests are at stake. Diagnostic for any institutional arrangement: (1) Map the power asymmetry — which actors have greater capacity to shape the institution's operation, through what mechanisms, and what resources maintain that capacity? (2) Trace the capture trajectory — who has invested most in influencing this institution's operations over time, through what channels, with what outcomes for the institution's alignment with its original function? (3) Identify the information structure — who knows what about how this institution actually operates, where the gaps between formal description and actual practice are, and who benefits from those gaps? (4) Ask what the institution optimizes for in practice, not in stated mission — and whether those two things have diverged, and in whose direction.

6.4 — The Anti-Technocratic Guardrail

This kind of analysis — beginning from material conditions, emphasizing systemic dynamics, evaluating arrangements by outcomes — invites a specific misreading: that the solution is technical governance by those who understand the analysis best. This would make the analysis most useful to exactly the actors it is most concerned about — those who already hold the resources, institutional access, and credentialing to claim superior analytical capacity. Technocratic governance — the concentration of decision-making authority in the hands of those who claim superior analytical capacity — is not the implication of superior analysis. It is the capture of the analytical framework by the interests of those who claim to possess it. The history of technocratic projects in the twentieth century is a history of exactly this capture: frameworks that began as tools for understanding the conditions of the many were converted into tools for legitimating the authority of the few, precisely because the frameworks were sophisticated enough to generate compelling justifications for whatever their operators wanted to do. The guardrail against this reading is not a modification of the framework's analytical claims. It is the application of the framework's own power analysis to the framework itself. The question — who benefits from this analysis being accepted? — is a question this framework asks about every other claim, and it must be asked about its own conclusions with particular rigor. An analysis whose primary practical effect is to increase the authority and autonomy of those who conduct it has been captured, regardless of the quality of its analytical content. This is the custodian problem applied specifically to analytical frameworks. The framework that provides the most compelling justification for whatever its holders want to do is not the most powerful framework. It is the most dangerous one. The check against this is genuine responsiveness to criticism from those whose interests the analysis claims to serve — not as a courtesy, but as the primary mechanism for detecting capture before it becomes self-reinforcing. The technocratic misreading of structural analysis is not a marginal risk. It is the default capture pathway for frameworks of this kind — the route through which analytical frameworks built to challenge power arrangements are converted into tools that justify the authority of those who hold the analysis. The

conversion works through a substitution: the framework's emphasis on structural analysis, material conditions, and systemic dynamics becomes evidence that outcomes should be determined by those who can perform structural analysis, understand material conditions, and model systemic dynamics — which is, precisely, the credentialed expert class whose institutional position the framework's power analysis identifies as one of the primary mechanisms of structural advantage. Why technocracy fails the framework's own criterion: Technocratic governance

fails the minimal anchor criterion in a specific way: it systematically produces outcomes that serve the material interests and cognitive frameworks of the technocratic class while claiming to serve the interests of the populations affected. This is not primarily because technocrats are malicious — it is because the information required to make decisions that genuinely serve affected populations is held by those populations, not by the analysts. The practitioner knowledge analyzed in 1.7 and 1.10 — the detailed, concrete, contextual understanding of how institutional arrangements actually operate for people navigating them from positions of limited power — is not available to the technocrat who designs systems from outside. Technocratic design consistently underweights this knowledge, not through deliberate exclusion but through the structural inaccessibility of knowledge that is not credentialed, documented, or present in the professional networks through which experts exchange information. The democratic constraint and its structural basis: The anti-technocratic guardrail is not a concession to populism or a claim that expert knowledge is worthless. Expert knowledge is real, its quality varies, and the variation matters for outcomes. The claim is structural: decisions about arrangements that affect large populations should be constrained by the preferences and judgment of those populations, not because those preferences are always well-informed, but because the alternatives — unconstrained expert authority, institutional discretion without accountability, or the implicit authority of whoever controls the analytical framework — are structurally worse. They are worse not in spite of being more analytically capable but because of it: unconstrained expert authority concentrates the capacity to shape outcomes in the hands of actors with the least exposure to the consequences of those outcomes for the populations most affected. The democratic constraint is not an epistemological claim about what ordinary people know. It is a structural claim about where the accountability for decisions must sit if those decisions are to be corrected when they fail. Self-application: The anti-technocratic guardrail applies to this framework itself. PMN's analytical vocabulary, its structural analysis, and its formula architecture all risk becoming tools that justify the authority of those who have mastered them, rather than tools that serve the populations whose conditions they are designed to analyze. The framework's internal check against this is the epistemological commitment in Part I — the requirement that analytical conclusions be revisable based on evidence, including evidence from the populations the analysis describes — and the individual-level demands in Part XVII, which include the obligation to communicate the framework's conclusions in forms legible to those who hold different analytical frameworks. A framework that is only useful to its own initiates has

reproduced, at the level of analytical discourse, the same exclusion it was built to analyze. The practical implication is not that experts should be excluded from institutional decision-making but that expert analysis should be structured as an input to decisions constrained by democratic accountability, not as a substitute for it. Expert knowledge about the structural consequences of policy choices, about the causal mechanisms that produce institutional outcomes, and about the historical evidence bearing on institutional design is genuinely valuable and should be available to decision-making processes. What it should not be is dispositive — it should not settle questions about what outcomes are acceptable, which trade-offs are worth making, and whose interests count, because those are normative questions whose answers are not derivable from structural analysis and must be held accountable to the people who live with their consequences.

6.5 — How Institutions Preserve Themselves

Institutions develop self-preservation mechanisms that operate largely independently of their ostensible purpose. No conspiracy is required. Institutions are operated by people whose material interests include the continuation of their positions, their influence, and the institutional arrangements that sustain them — and the self-preservation mechanisms follow from that without needing to be designed. The self-preservation mechanisms are the predictable behavioral consequences of those interests, aggregated across the population of actors within the institution. The most effective mechanism is the redefinition of institutional purpose around institutional survival. An institution created to solve a specific coordination problem gradually reorients its activities toward reproducing the conditions that justify its existence. Regulatory agencies created to constrain industry develop relationships with industry that make them advocates for the sectors they regulate. Aid organizations created to address specific material needs develop programs that sustain their operational capacity regardless of whether those programs address the original needs. Educational institutions created to develop human capacity develop curricula oriented toward reproducing the institutional hierarchy that employs their graduates. The second mechanism is control over the criteria by which performance is measured. An institution that controls its own evaluation metrics can define success in terms that its current activities satisfy, regardless of whether those activities address the conditions that generated the institution in the first place. The evaluation becomes a legitimization exercise rather than an assessment — and the actors who perform it are typically those whose interests are served by the outcome it produces. The third mechanism is information asymmetry. Institutions control what information is produced about their activities, how it is framed, and who has access to it. The asymmetry is not always deliberate manipulation — it is often the unreflective product of institutional culture, which systematically frames institutional activities favorably and systematically fails to produce information that would present them unfavorably. The result is an asymmetry between the institution and the populations it serves that makes external accountability difficult to exercise even when formal mechanisms for it exist. The anti-technocratic guardrail must be distinguished from a position it is sometimes confused with: the rejection of scientific expertise in policy. These are structurally different claims. The anti-technocratic argument targets the concentration of decision-making authority in those who claim superior analytical capacity — the use of expertise as a license to override democratic accountability. It does not target the relevance of scientific evidence to policy decisions. Where genuine scientific consensus exists — where a question has been subjected to sustained

empirical investigation, independent replication, and the peer processes that make claims revisable — that consensus represents the best available description of the material conditions that policy must engage. Ignoring it is not democratic; it is epistemologically reckless in ways that produce identifiable structural harms. The distinction the framework requires is between two questions that are frequently collapsed: the empirical question of what the evidence shows (where scientific expertise is authoritative) and the political question of what should be done given what the evidence shows (where democratic accountability is required). Scientific consensus on climate forcing does not determine climate policy — it specifies the material constraints within which any viable climate policy must operate. A democratic process that ignores those constraints is not exercising genuine self-determination; it is selecting among fantasies. The framework's criterion applies here: epistemic infrastructure that allows manufactured controversy to displace genuine scientific consensus is infrastructure in the service of actors whose material interests are served by policy delay — which is institutional capture of the knowledge-production system, not democratic accountability. The appropriate institutional response is neither technocratic override nor democratic denial of evidence; it is the construction of epistemic infrastructure robust enough to make the distinction between genuine scientific disagreement and manufactured controversy legible to the populations who must make democratic decisions about their material conditions. The personnel mechanism: The most durable self-preservation mechanism is control over who enters and advances within the institution. Institutions select for personnel who fit the institutional culture — who share its assumptions, who will not challenge its central practices, who have the right prior relationships with the right prior institutions. This selection is

rarely explicit and rarely deliberate in any individual case. It operates through the accumulated judgments of hiring committees, promotion boards, and informal networks whose members have been selected by the same process they now operate. The result is institutional reproduction: each generation of institutional leadership is selected by the previous generation in its own image, which ensures that the institution's core commitments and the interests of those who hold them are maintained across turnover. The mechanism is most visible when it breaks down — when an outsider enters and disrupts the pattern — and the disruption itself reveals how much of the institution's normal operation depended on the continuity the personnel mechanism maintained. The procedure mechanism: Institutional procedures function as self-preservation mechanisms because they define what counts as a legitimate action within the institution, which means they define what changes can be pursued through in

stitutional channels and what changes require going outside them. Procedures that require high consensus thresholds before changes can be made systematically favor the status quo: any change must overcome the inertia of existing arrangements, while existing arrangements require no justification beyond their existence. Procedures that vest authority over institutional processes in the actors who benefit most from current institutional arrangements create accountability structures that cannot be used against those arrangements from within the institution. Over time, institutional procedures accumulate in ways that make the institution progressively harder to change through its own internal mechanisms — not through design but through the accumulation of precedents, each individually reasonable, that together constitute a dense protective structure around existing arrangements. The narrative mechanism: Institutions maintain themselves through the narratives they produce about their own necessity, legitimacy, and irreplaceability. These narratives are not simply propaganda — they reflect genuine institutional functions and real historical contributions that the institution can accurately claim. But they also systematically overweight the institution's contributions relative to its costs, underweight the alternatives that have not been tried, and frame the institutional arrangement as the only viable form for the functions it performs. The narrative that a specific institution is necessary for a function it performs is one of the most effective self-preservation mechanisms available, because it converts any challenge to the institution's current form into an apparent challenge to the function itself — making critics of the institution appear to be opponents of what the institution does rather than opponents of how it does it. The coalition mechanism: Institutions build external coalitions of actors whose interests are served by the institution's continuation in its current form. These coalitions may include the populations the institution ostensibly serves — who depend on its continuation even if its current form serves them poorly — as well as vendors, partner organizations, political allies, and the professional communities whose identity and livelihood are organized around the institution. The coalition is a structural resource that the institution can mobilize against challenges, not through explicit coordination but through the natural alignment of interests: each coalition member opposes changes to the institution for their own reasons, and the aggregate effect of that opposition is the institutional preservation that no single actor could produce alone. The coalition mechanism is why institutional reform is so often harder than the formal authority to make changes would suggest: the formal authority encounters a distributed resistance whose components each have independent motivations and independent resources to deploy. The self-preservation me

chanisms operate simultaneously and reinforce each other. Personnel selection maintains the culture that produces the narratives that recruit the coalition that defends the procedures that constrain who can enter and advance. Each mechanism is individually insufficient — each can be overcome by sufficient external pressure or internal disruption. In combination they constitute what might be called institutional inertia: not a single force but a system of mutually reinforcing constraints that makes institutional change require substantially more energy than the formal structure of authority would suggest. Understanding this system is the precondition for designing change strategies that are adequate to it. The connection to 7.3's custodian problem: institutional self-preservation mechanisms are the micro-level mechanisms through which the custodian problem operates. The custodian problem describes the structural tendency of institutions to serve their own continuation rather than their original function; the self-preservation mechanisms describe how

that tendency is actually implemented — through personnel, procedure, narrative, and coalition. Reforms that address the custodian problem at the level of incentive design without attending to these mechanisms will consistently find that the mechanisms reconstitute the original arrangement faster than the incentive design can operate.

6.6 — Legitimacy Crisis: When Systems Lose the Consent That Sustains Them

Legitimacy is not a property that institutions either have or lack. It is a variable that accumulates and depletes over time, determined by the gap between what an institution claims to do and what it materially does for the populations whose cooperation it requires. When the gap is small, legitimacy accumulates. When the gap grows — when the institution visibly fails to perform its stated function while continuing to extract the costs of maintaining it — legitimacy depletes. Legitimacy depletion is not linear. Populations tolerate significant gaps between institutional promise and performance for extended periods, sustained by cultural inertia, by the genuine costs of coordination failure, and by the effective deployment of the self-preservation mechanisms described above. But the tolerance has limits that are not visible in surface-level indicators until they are crossed. When legitimacy crosses the threshold below which voluntary compliance requires active enforcement, the cost of maintaining the institution increases sharply. Enforcement requires resources. Resources must be extracted from the same population whose compliance is now requiring enforcement. The extraction increases structural suffering. The increased suffering further depletes legitimacy. The dynamic becomes self-reinforcing: the institution either reforms rapidly enough to address the underlying material causes of depletion, or it collapses into a configuration sustained by coercion alone — which is both more costly and less stable than the previous arrangement. Legitimacy crises are not primarily ideological events. They are the political expression of accumulated structural suffering that existing institutional arrangements cannot address. The ideology that accompanies the crisis — the specific narratives that mobilize populations whose cooperation has become conditional — is the surface phenomenon. The material substrate is the distribution of suffering that the ideology is expressing and attempting to address. Understanding the ideology without understanding the material substrate produces analysis that consistently misidentifies both the cause and the likely trajectory of the crisis. The threshold at which voluntary compliance shifts to requiring active enforcement is not visible in surface-level political indicators, because legitimacy depletes through a process that standard metrics are not designed to track. What can be tracked are the leading and lagging indicators of threshold approach — not precise measurements, but diagnostic signals that, in combination, indicate where a system is in the depletion process. Leading indicators (threshold is being approached): declining voluntary participation in collective processes the system formally requires but cannot compel, including voting, civic service, and tax compliance at the margin. Withdrawal of investment in long-term outcomes: reduced maintenance of shared infrastructure, shorter planning horizons in both public and

private decision-making. Increasing reliance on symbolic rather than substantive institutional performance, with ceremonial accountability processes generating the appearance of oversight without its function. Erosion of the elite consensus that sustained the arrangement, visible when significant factions of the groups whose cooperation the institution depends on for maintenance begin to defect or hedge. These signals are frequently dismissed as cultural or political phenomena. They are material: they indicate that the behavioral substrate of voluntary cooperation is depleting. Lagging indicators — threshold has been crossed: Active enforcement replacing voluntary compliance in domains that previously required no enforcement. Sustained increase in the cost of maintaining institutional order without corresponding increase in the value it produces. Emergence of parallel institutions — alternative frameworks for coordination, protection, and dispute resolution — that compete with the official system for behavioral cooperation. Official system's increasing reliance on suppression of alternatives rather than improvement of its own performance. By the time lagging indicators are visible, the window for reform has substantially narrowed: the cost of enforcement and the

capacity of parallel institutions have shifted the material conditions of change in ways that incremental adjustment cannot reverse. The distinction between leading and lagging indicators has a direct consequence for the sequencing problem in Part X. Interventions adequate to the depletion process while leading indicators are present become insufficient once lagging indicators appear. The window for managing legitimacy crisis through structural reform rather than through breakdown is the period between the appearance of leading indicators and the crossing of the threshold. Reading those indicators accurately — and acting within that window — is one of the most practically important applications of the entire framework. Power, as analyzed in this Part, is the prior condition for understanding institutions. Institutions are not neutral coordinating devices that happen to be contested — they are the crystallized form of previously established power relations, designed to reproduce the conditions that made them possible. The question Part VII addresses is not whether institutions are shaped by power (they are) but what follows from that: under what conditions can institutions be reformed from within, when does the structure of an institution make internal reform impossible, and what does genuine accountability require given the problem of capture. The analytical move from Part VI to Part VII is from power as a structural position to power as an institutionalized arrangement, and from there to the practical question of what progressive institutionalism can actually mean.

6.7 — Legitimacy Production: How Systems Build Consent

Section 6.6 analyzed legitimacy depletion: the process by which the gap between institutional promise and institutional performance erodes the voluntary compliance that institutional arrangements depend on. The analysis of depletion presupposes that legitimacy was accumulated in the first place. The mechanisms through which legitimacy is produced — how systems build the consent that depletion then erodes — is a distinct analytical question with direct consequences for understanding both why some arrangements are more durable than others and what the restoration of legitimacy after crisis actually requires. Three legitimacy production mechanisms recur across historical and contemporary political contexts, each with distinct structural features, distinct sustainability conditions, and distinct vulnerabilities to collapse. Most durable arrangements combine them, but each functions through different mechanisms and responds to different kinds of stress. Understanding which mechanism is doing the primary work in any given arrangement is essential for predicting how legitimacy will respond when conditions change. Performance Legitimacy Performance legitimacy is built through the visible delivery of outcomes that populations value: material welfare, physical security, reliable dispute resolution, access to economic opportunity. It is the most straightforwardly material form of legitimacy production, and its logic is direct: the arrangement earns compliance by delivering what it promises. Performance legitimacy is powerful when it is functioning because it is grounded in the material experience of populations rather than in their beliefs about the arrangement. It does not require that populations believe the arrangement is just, only that they judge it to be delivering on what they need. The political economy of high-performance developmental states — arrangements that built sustained legitimacy through visible economic delivery without procedural democracy — is the paradigm case: populations that are materially better off under an arrangement extend compliance without requiring its justification. Performance legitimacy is the most fragile of the three mechanisms precisely because it is so directly tied to material outcomes. It cannot survive sustained performance failure, and it provides no buffer against the bad-luck shocks that affect even well-designed arrangements. An arrangement that has built its legitimacy entirely on economic performance is vulnerable to any significant economic shock, regardless of whether the shock was produced by the arrangement itself or by external conditions. This is why purely performance-based legitimacy rarely survives generational transitions: the populations who directly experienced the performance extend compliance; their children, for whom the performance is historical background rather than lived experience, require the arrangement to continue delivering in their own lifetimes or shift their compliance elsewhere. Pr

Procedural Legitimacy Procedural legitimacy is built through the credible operation of processes that populations regard as fair, regardless of the specific outcomes those processes produce. Its logic is distinct from performance legitimacy: populations extend compliance not because they are satisfied with outcomes but because they regard the process as one that gave them genuine standing and whose outcomes they are therefore obligated to accept even when unfavorable. The resilience of procedural legitimacy is its most important structural feature: it can survive extended periods of poor outcomes, sustained inequality, and significant policy failure as long as the process itself retains credibility. Democratic legitimacy in liberal societies has functioned largely through this mechanism: populations accept outcomes they strongly oppose because the process that produced them is regarded as legitimate. The vulnerability of procedural legitimacy is equally specific: it collapses when the process is visibly corrupted or when populations conclude that the formal procedure is a shell with no genuine connection to the decisions actually being made. This is the mechanism through which institutional capture (7.3) produces legitimacy crisis even in formally democratic arrangements: when populations accurately perceive that formal procedures are functioning while substantive decisions are made elsewhere, procedural legitimacy depletes faster than performance legitimacy would under equivalent conditions, because the perceived violation is not merely failure but betrayal. The arrangement claimed to give populations standing; the discovery that it did not is a qualitative shock to compliance rather than a gradual erosion. Narrative Legitimacy Narrative legitimacy is built through the establishment of a founding story, civilizational mission, or historical necessity that places the arrangement above the ordinary calculus of performance and procedure. The arrangement is legitimate not because it delivers or because its processes are fair, but because it embodies something larger: the nation, the revolution, the divine order, the historical destiny of a people. Narrative legitimacy is the most durable of the three mechanisms under normal conditions and the most catastrophically vulnerable under specific ones. An arrangement with deep narrative legitimacy can survive extended performance failure and significant procedural violation that would collapse a purely performance- or procedurally-based arrangement. The narrative absorbs the failure: the arrangement is still the legitimate expression of something that transcends its current difficulties. Narrative legitimacy is produced through myth-making — the construction of founding narratives and collective identity stories that make the arrangement feel like the natural expression of a community's deepest commitments rather than a contingent institutional choice. The mechanisms of production are analyzed throughout 8.4

and 8.4b; what matters here is the political consequence: arrangements with strong narrative legitimacy are resilient to material and procedural shocks in ways that purely performance- or process-based arrangements are not. The collapse condition for narrative legitimacy is equally specific: it occurs when the founding narrative becomes publicly untenable, typically through a single catastrophic event that the narrative cannot absorb without self-contradiction. The collapse is then rapid and total — because narrative legitimacy had been providing insulation against the performance and procedural failures that were accumulating beneath it, and once the insulation fails, all three sources of depletion hit simultaneously. This is the structural explanation for the sudden implosion of arrangements that appeared durable until the moment of collapse. The procedural democratic objection to this framework — that the material criterion is a substantive standard imposed from outside the democratic process, one that could override democratically expressed preferences — is serious and deserves a structural response. The framework's answer is that procedural legitimacy has never been sufficient without performance legitimacy, and that most substantive democratic theorists already hold this: elections that produce genocide or systematic exclusion are not redeemed by procedural correctness. What the framework adds is a specific, empirically grounded criterion for what substantive legitimacy requires. The Habermasian alternative — that the criterion should emerge from ideal discourse conditions rather than be specified in advance — faces a prior problem: ideal discourse requires the absence of the power asymmetry, information manipulation, and epistemic infrastructure degradation that this framework analyzes as the normal conditions of political life. A theory of legitimate democratic outcomes that can only function when power has already been equalized provides no guidance for the conditions under

which guidance is most needed. The diagnostic implication of the three-mechanism analysis is that legitimacy restoration after crisis requires identifying which mechanism was primary and what its specific collapse condition was. Restoring performance legitimacy requires improving material outcomes. Restoring procedural legitimacy requires genuine reform of the process — not ceremonial reform that maintains the appearance while reproducing the capture that destroyed credibility. Restoring narrative legitimacy is the most difficult: it requires either the construction of a new founding narrative or the genuine rehabilitation of the old one through actions that the narrative itself validates. Reform programs that address the wrong mechanism — improving material delivery to populations whose compliance has collapsed for procedural reasons, or reforming procedure for populations whose legitimacy was narrative and has been catastrophically undermined — will fail to restore

the compliance they are designed to recover.

6.8 — Elite Circulation and Structural Continuity

Political systems change personnel regularly. They change structure rarely. The gap between these two rates — the high frequency of personnel change and the low frequency of structural change — is one of the most consequential and least analyzed features of how power actually reproduces itself across time. An analysis that treats personnel change as structural change will consistently misread both the significance of political transitions and the conditions under which genuine structural transformation becomes possible. The analysis of elite circulation developed in the sociological tradition by Pareto and Mosca remains analytically indispensable precisely because it identifies the mechanism: elites circulate, but the structural positions they circulate through persist. The core observation is that every stable system of power produces a class of actors — an elite in the structural sense of those who occupy the positions from which the most consequential decisions are made — and that this class reproduces itself through mechanisms that are partially independent of the formal political arrangements nominally governing access to power. Education systems, professional networks, marriage patterns, cultural capital, and inherited economic position all function as mechanisms of elite reproduction that operate below the level of formal political competition. The result is that formal political competition — elections, revolutions, coups — changes which individuals occupy elite positions without necessarily changing the structural features of the elite itself: its size, its internal cohesion, its relationship to the populations it governs, or its access to the resources that elite position provides. **Circulation Without Transformation** Elite circulation without structural change is the mechanism through which political arrangements absorb transformation pressure without transforming. When legitimacy depletion reaches the threshold analyzed in 6.6, the system faces a choice between structural reform — changing the arrangements that generated the depletion — and elite circulation — replacing the specific personnel who have become identified with the arrangement's failures while leaving the arrangement itself intact. Elite circulation is almost always the path of least resistance for those who benefit from the existing structure, because it addresses the surface phenomenon — the visible actors associated with failure — without touching the underlying conditions. It satisfies the demand for change by producing change that is visible, dramatic, and structurally inconsequential. The populations who drove the demand for change are not irrational when they accept elite circulation as a response to structural grievance. The information asymmetry between governing elites and governed populations — analyzed in 1.10 — makes it structurally difficult to distinguish elite circulation from structural transformation at the moment it is occurring. New personnel make differ

ent symbolic claims, deploy different narratives, and produce different surface-level behavior than the elites they replace. The structural continuity beneath these surface differences only becomes visible over time, as the new personnel reproduce the same decisions, the same alliances, and the same distribution of costs and benefits that the old personnel produced — because they occupy the same structural positions with the same structural incentives. **Elite Fragmentation as a Transformation Condition** Elite cohesion is one of the primary

structural conditions that determines whether transformation pressure produces transformation or absorption. A cohesive elite — one that shares interests, maintains internal coordination, and presents a unified front to the populations it governs — can absorb very high transformation pressure through a combination of elite circulation, targeted concession, and the suppression of the most visible grievances without addressing their structural causes. A fragmented elite — one in which different factions have genuinely different interests, are competing actively for power, and cannot coordinate on a common response to transformation pressure — is structurally much more vulnerable to the transformation pressure it faces, because the factions in competition for power have incentives to make genuine structural concessions to build the coalition support that their intra-elite competition requires. This is the analytical basis for the inclusion of elite fragmentation as a component in the counter-power formula (CP): elite fragmentation creates structural openings that cannot be created by counter-power alone, however well-organized. The transformation windows analyzed in 10.5b are frequently produced by the intersection of accumulated transformation pressure from below and elite fragmentation from within: the former provides the material pressure, the latter provides the structural opening through which that pressure can produce genuine institutional change rather than being absorbed through elite circulation. Movements that fail to read elite fragmentation as a structural variable — that treat all elite opposition as equivalent regardless of its internal divisions — will consistently fail to exploit transformation windows at the moment they are open. Elite Reproduction and Structural Persistence Elite reproduction mechanisms are the structural explanation for why formal political changes — including democratic transitions, revolutions, and major institutional reforms — so frequently produce outcomes that preserve more of the previous power structure than the actors driving the change anticipated. The mechanisms of elite reproduction operate on timescales longer than political transitions: the education systems, professional networks, and cultural capital that reproduce elite access to power are not dismantled by formal institutional change. The new arrangement inherits the social infrastructure of the old one, and that infrastructure continues to reproduce access patterns that were built around the previous power structure. The practical implication is that genuine structural transformation requires attending to elite reproduction mechanisms explicitly, not only to the formal institutional arrangements that are the usual object of reform. Land reform that changes formal ownership without changing the social networks through which access to credit, markets, and legal protection is mediated will produce formal ownership changes that are gradually reversed as the old elite uses its network advantages to recover economic position. Educational reform that opens formal access without changing the cultural capital requirements of elite institutions will produce formal integration while preserving substantive exclusion. The analysis of what a specific transition actually changed requires attending to the entire reproduction mechanism, not only to the formal arrangements that are its most visible surface. The distinction between elite circulation and structural transformation is one of the most practically important analytical distinctions the framework provides. Applying it requires asking not whether personnel changed but whether the structural positions they occupy changed: whether the incentive architecture of those positions was restructured, whether the reproduction mechanisms that regulate access to those positions were altered, and whether the distribution of costs and benefits that those positions produce changed in ways that cannot be reversed by the next round of elite circulation. When those questions are answered negatively, what occurred was circulation, not transformation — regardless of how dramatic the personnel change appeared and regardless of the claims made by those who drove it.

VII: How to Organize: Politics, Institutions, and the State

7.0 — What We Mean by System

‘System’ is used throughout what follows in a precise sense that is worth distinguishing from the loose usage common in political analysis, where it functions as a synonym for ‘things as they are.’ A system here is a complex adaptive arrangement of organisms, institutions, resource flows, and incentive structures that interact across multiple levels simultaneously. Families are systems. Markets are systems. States are systems. The international order is a system. These systems are nested: the dynamics of smaller systems are constrained and shaped by the larger systems that contain them, and the behavior of larger systems emerges from the interactions of smaller ones. Systems in this sense are not machines that can be engineered from the outside. They are adaptive: they respond to intervention in ways that are not fully predictable, because the actors within them adjust their behavior in response to whatever changes are made. This is why the history of policy intervention is so full of unintended consequences. For the same reason, the analysis of systems must be ongoing rather than terminal: a model of a system built at one moment will become inadequate as the system adapts, and the work of revision is permanent. Recognizing this is not counsel for inaction. It is counsel for humility about the completeness of any analysis, and for building in the capacity for revision rather than assuming that the initial model is sufficient. The importance of the definition is not terminological. It is analytical: which features of a system are visible depends on what counts as the system in the first place. A definition that treats the system as the formal institutional structure — the laws, the stated procedures, the organizational chart — will consistently miss the causal mechanisms that operate through informal arrangements, accumulated precedent, and the incentive structures that formal arrangements create but do not fully describe. A definition that treats the system as whatever produces bad outcomes — a moralized usage common in activist discourse — produces analysis that cannot distinguish between features of the arrangement that are structurally necessary for the harmful outcomes and features that are contingent and changeable. Neither is useful for the purposes this framework serves. What makes a system complex: Complexity in the relevant sense has three properties. First, the system has multiple interacting components whose combined behavior is not predictable from the behavior of any single component — the interaction effects are real and consequential. Second, the system operates across multiple levels simultaneously, and the causal mechanisms operating at each level are partially autonomous from those at other levels — the organism level, the organization level, and the system level each have dynamics that cannot be fully reduced to the dynamics of the level below. Third, the system adapts — it changes in response to pressure from its environment a

nd from internal disruption, and those adaptations can either absorb or amplify the pressure depending on conditions that are themselves part of the system. Systems that lack all three properties are not complex in the relevant sense, and analysis designed for complex systems will misapply to them. Adaptive without directional guarantee: The adaptive character of systems is the source of both the framework's most important analytic insight and its most important warning against misuse. Systems adapt — they generate responses to pressure that change what the system does. But adaptation is not optimization toward any fixed goal, and it is certainly not movement toward arrangements that better serve the minimal anchor criterion. Historical systems have adapted in ways that produced more suffering, not less — more efficient extraction, more effective suppression of counter-power, more durable arrangements that distributed costs downward. Adaptation in the sense used here means only that the system changes in response to pressure. What it changes toward depends on which pressures are applied, by whom, with what resources, and with what structural support. That is why the analysis of counter-power in Part X is inseparable from the analysis of systems in Part VII: understanding how systems work is precondition for understanding how they can be changed, but it does not by itself generate the conditions for change. The interaction across levels: The most analytically consequential feature of complex adaptive systems is the interaction across levels. Individual

organism behavior produces organizational-level outcomes that were not intended by any individual. Organizational-level dynamics produce system-level properties that were not intended by any organization. And system-level pressures reshape the conditions within which individuals and organizations operate, feeding back into the organism and organizational levels in ways that make the system's behavior genuinely emergent — not reducible to any single level's description. This interaction is why single-level analyses consistently produce both analytical errors and reform failures: interventions at the individual level that do not account for organizational dynamics will be undone by those dynamics; interventions at the organizational level that do not account for system-level properties will be captured or neutralized by the system's adaptive response; interventions at the system level that do not account for individual and organizational incentives will produce compliance theater rather than genuine change. What this definition rules out: it rules out treating "the system" as a unified intentional actor — as though there were a directing intelligence maintaining current arrangements by design. Systems maintain arrangements through the interaction of individually motivated actors whose aggregate behavior produces systemic outcomes, not through central coordination. This matters for strategy: if the sys

tem were a unified intentional actor, defeating it would require defeating the directing intelligence. Because it is not, changing it requires changing the incentive structures and resource distributions that produce its emergent properties — which is a different kind of task requiring a different kind of analysis. The custodian problem in 7.3 is the specific case of this general point: institutions that serve their own maintenance rather than their original function are not doing so by conspiracy but by the structural dynamics that make self-maintenance rational for the individuals who operate within them.

7.0b — Levels of Analysis

Three levels of analysis run simultaneously through what follows, and collapsing them — treating observations at one level as automatically valid at another — is the source of some of the most persistent errors in political and social analysis. The levels are distinct but connected, and understanding any specific phenomenon requires identifying which level is primary in its causation rather than explaining everything from the most immediately visible one. The organism level is the level of individual biology and psychology: the specific conditions of specific people — their suffering, their behavioral incentives, their cognitive constraints, their material circumstances. Analysis at this level asks: what are the material conditions of this person or population, what behavioral responses do those conditions generate, and how do those responses shape what is possible at higher levels? The organism level is where the minimal anchor is most directly visible: the suffering or flourishing the framework ultimately cares about happens here, to specific people in specific bodies. The interaction level is the level of how organisms behave in relation to each other: the patterns of cooperation and conflict, trust and defection, reciprocity and exploitation that emerge from the behavioral responses of multiple organisms in shared material conditions. Analysis at this level asks: given the behavioral incentives at the organism level, what patterns of interaction tend to emerge, and what do those patterns make possible or impossible at the institutional level? The interaction level is where individually rational behavior produces collectively irrational outcomes. The system level is the level of institutions, resource distributions, power configurations, and information environments that shape what is possible at both lower levels. Analysis at this level asks: what structural conditions are generating the behavioral incentives at the organism level, what interaction patterns are they sustaining or disrupting, and how do those patterns feed back into the structural conditions themselves? The system level is where historical inertia operates most strongly — where arrangements that were once functional persist past their material justification because the actors who benefit from them control the mechanisms of change. The movement between levels is not one-directional. System-level conditions shape organism-level material circumstances, which shape behavioral incentives, which aggregate into interaction patterns, which maintain or transform system-level conditions. The question that each level cannot answer from inside itself — why

do these incentives exist, why do these patterns persist — is typically answered at the level above it. The most common analytical error in political commentary is the attribution of system-level outcomes to organism-level causes: explaining political crises as products of bad leadership or moral failure rather than as products of the s

tructural conditions that made those outcomes probable regardless of who occupied the relevant positions. The organism level and its limits: The organism level — the level of individual psychology, motivation, and behavior — is the level at which political analysis most often defaults, because it is the level at which causation is most directly visible. This person did this thing; that outcome followed. The limitation of organism-level analysis is not that it is wrong about individual causation but that it systematically underestimates how much of individual behavior is produced by structural conditions — by the incentive structures, resource distributions, and normative frameworks within which individuals operate — rather than by individual character, intention, or values. The same individual placed in different structural positions will behave differently in predictable ways. This does not eliminate individual agency or moral responsibility, but it means that individual-level interventions — changing the people in positions rather than the conditions of those positions — will consistently produce less durable change than structural interventions that change what the positions themselves make rational. The organization level and its characteristic errors: The organization level — the level of institutions, formal organizations, and the rules that govern collective action — is the level at which most policy analysis operates. It is the level at which formal authority is located and at which deliberate design is most legible. The characteristic errors at this level are two. The first is treating formal organizational design as determinative of actual behavior — assuming that because an institution has a mandate, a procedure, and an accountability structure, it will behave in accordance with them. The custodian problem demonstrates why this assumption fails: institutions develop self-maintenance incentives that operate independently of their formal design. The second error is treating organizational-level dynamics as fully autonomous from system-level pressures — designing institutional reforms without accounting for how the broader system of resource distribution, political economy, and power arrangement will shape what the reformed institution can actually do. Organizational design is not sovereign over organizational behavior. The system level and the attribution problem: The system level — the level of emergent properties that arise from the interaction of multiple organizations and populations within an environment of material constraints — is the level at which the most consequential causal mechanisms operate and the level at which analysis is most difficult. System-level properties are not located in any single organization or individual; they emerge from the interaction between them. This produces what might be called the attribution problem: when a system-level outcome occurs — persistent poverty, recurring financial crisis, democratic ero

sion — there is rarely a single identifiable cause at the organization or organism level, because the outcome is produced by the interaction of many causes operating simultaneously across levels. Analysis that requires a single identifiable cause will consistently misattribute system-level outcomes to individual failures or organizational dysfunctions, producing interventions that are locally accurate and systemically insufficient. The cross-level interaction as the primary analytical challenge: The most persistent analytical errors in political and institutional analysis are cross-level errors: applying observations from one level to explain phenomena at another. Individual-level observations — this person is corrupt, this manager is incompetent — are used to explain system-level outcomes that are produced by structural incentives that would make most individuals in that position behave similarly. System-level observations — this society is characterized by low trust — are used to explain individual behavior without accounting for the rational bases for that behavior given the actual structural conditions individuals face. Organization-level observations — this institution has poor procedures — are used to explain outcomes that are produced by the system-level pressures operating on the institution rather than by its internal design. Each cross-level error produces characteristic reform failures: fix the individual, the structural incentive produces another one; fix the procedure, the system-level pressure finds another route; address the system-level property, the

organizational and individual responses adapt in ways that restore the original configuration. The discipline the three-level framework requires is explicit identification of the primary level of causation for any specific phenomenon before designing an intervention. Not every phenomenon is primarily system-level — some institutional failures genuinely are organizational design problems whose structural solution is organizational redesign, and some organizational failures genuinely are individual failures whose solution is individual accountability. The framework does not stipulate that system-level causation is always primary. It requires that the question be asked — and that the answer be supported by an account of the causal mechanism operating at the claimed level, not by the analyst's default preference for the level of analysis they are most comfortable with. Cross-level diagnostic: for any proposed explanation of a harmful outcome, ask: (1) At what level is the explanation primarily operating — individual, organization, or system? (2) What would the outcome look like if the proposed cause were removed but the conditions at other levels remained unchanged — would the outcome disappear or would a functionally equivalent cause emerge? (3) Is the proposed intervention targeted at the level of primary causation, or at a symptom visible at a different level? (4) What cross-level inte

ractions would the intervention trigger, and would those interactions restore the original outcome through a different pathway?

7.0c — Scale as a Structural Variable

Scale is typically treated in political analysis as a trade-off: larger institutions gain coordination capacity and lose accountability; smaller institutions gain accountability and lose reach. This framing is not wrong, but it is incomplete in a way that has strategic consequences. Treated as a fixed trade-off, scale becomes a constraint that forecloses institutional design — the question becomes which end of the spectrum to accept rather than how institutional architecture can modify what the trade-off actually is at each level. PMN treats scale as a structural variable with predictable effects on the custodian problem, the capture sequence, and the distribution of becoming conditions — effects that are not fixed but shift with institutional design choices and the material conditions of information, coordination, and counter-power available at each scale. Scale affects four structural variables in predictable ways. First, information asymmetry (1.10) increases with scale because the gap between what those who operate an institution can observe and what those subject to it can observe widens as the population subject to the institution grows and as the institution's operations become more complex. At the scale of a village, the custodian's behavior is directly observable by most of those the institution affects; at the scale of a nation-state, it is not. The accountability mechanisms adequate to the village scale — direct participation, reputational sanction, immediate visibility of outcomes — do not transfer to the national scale, not because accountability becomes impossible but because the mechanisms required to produce it at that scale are different and more demanding. Second, capture efficiency increases with scale because the ratio of gain-from-capture to cost-of-organizing-capture improves as the institution's reach expands. A regulatory body that governs a large industry is worth more to capture than one governing a small one, while the cost of sustained capture — maintaining relationships, managing information flows, suppressing counter-claims — does not scale proportionally. Third, the visibility variable V in the structural suffering formula is harder to maintain at larger scales, because the populations affected are more dispersed, their experiences more varied, and the infrastructure required to make collective suffering legible across scale boundaries is more demanding to build and easier to degrade. Fourth, the resilience of coordination arrangements to external shocks increases with scale up to a threshold and then decreases, as larger arrangements develop interdependencies that amplify failure rather than absorbing it. These effects are not arguments against large-scale institutions. They are arguments for institutional designs at large scale that compensate for the structural disadvantages scale produces. Three design principles follow from the scale analysis. Distributed custodianship: rather than a single institutional anchor

operating at full scale, nested accountability structures in which custodians at each level are accountable to constituencies operating at the level below. The subsidiarity principle in Catholic social teaching — decisions made at the lowest level capable of addressing the relevant problem — is a pre-theoretical approximation of this structure that PMN can ground in the specific mechanisms that make it work or fail. Modularity: institutional designs that allow components to fail or be reformed without cascading failure across the full system, preserving the coordination functions of scale while limiting capture contagion. Transparency infrastructure: the deliberate construction of the information systems that allow those affected by large-scale institutional decisions to observe those decisions, understand their consequences, and contest them through organized channels — not as a transparency ideal but as a material infrastructure requiring resources, design, and maintenance. Each of these design principles is itself subject to the capture sequence: distributed custodianship can become fragmented incoherence, modularity can become regulatory arbitrage, and transparency infrastructure can become performative disclosure that generates information no one has the capacity to process. The scale analysis does not resolve these problems; it specifies what institutional design at each scale must address to avoid them. The relationship between scale and the anti-foreclosure criterion requires separate attention. The conditions for preserved adaptive response capacity — cognitive development, relational depth, meaningful participation in the decisions that shape one's situation — are not scalable in the same way that material provision is scalable. Large-scale institutions can deliver floor conditions (food, shelter, basic health, physical security) with efficiency gains at scale. They cannot deliver the conditions for becoming at the same efficiency gains, because becoming requires the kind of contextual engagement, local knowledge, and genuine participation that large-scale administration systematically degrades. This is not an argument for small communities as the only adequate institutional form; it is an argument that the anti-foreclosure criterion requires institutional designs that preserve genuine participation and local determination at the levels where they are possible, even within arrangements that coordinate at larger scale for floor provision. The tension between scale efficiency in floor provision and scale degradation in becoming conditions is a genuine and permanent tension — it belongs in Part XIII rather than being resolvable by better design. But it is not equally permanent at all scales: the specific institutional arrangements that make it more or less acute are design choices, not structural necessities. Scale diagnostics for institutional analysis: (1) At what scale does this institution operate, and what accountability mechanisms exist at that scale? Are they adequate to the information asymmetry that scale produces, or are they designed for a smaller-scale institution that has grown without redesigning its accountability architecture? (2) What is the capture efficiency of this institution at its current scale — how concentrated are the benefits of capture relative to the cost of organizing it? Has scale made it more or less vulnerable to the capture sequence in 7.3c? (3) What visibility infrastructure exists to make the consequences of this institution's decisions legible to those who bear them? Is that infrastructure publicly resourced and maintained, or is it dependent on the institution's own disclosure? (4) Is this institution delivering floor conditions, becoming conditions, or both? For floor provision, are there scale efficiencies being realized? For becoming conditions, does the institution's scale degrade the local determination and genuine participation that becoming requires? (5) Are there nested accountability structures that allow those affected at each level to contest decisions at the level above? Or does accountability run only upward through the institution's own hierarchy?

7.1 — The State Problem

The materialist tradition's treatment of the state was one of its most consequential analytical contributions and one of its most significant practical failures. The contribution: the state is not a neutral arbiter standing above society but an instrument that reflects and reinforces particular distributions of power and resources. The practical failure: the conclusion that followed from this analysis — that the state would eventually wither away as class divisions were overcome — produced political programs that concentrated state power

to an unprecedented degree in the name of the process by which that power would eventually dissolve. The analytical contribution is taken; the conclusion is refused. The state is an instrument — it reflects particular distributions of power, it can be used to reinforce those distributions or to change them, and its current form is not the final form of human political organization. But the conclusion that it will or should wither away rests on the assumption that the conditions requiring coordination — the permanent conditions identified in Part III — are themselves transitional. They are not. As long as those conditions persist, structured coordination is a permanent feature of human social organization. The question is not whether to have institutions but how to design them so that they remain accountable to the function they are supposed to serve. The conclusion the tradition drew — that the state would eventually wither away as class contradictions were resolved — was not merely historically wrong. It was structurally wrong in a way that had practical consequences. It produced movements that could diagnose the state's class character while having no developed theory of what to do with state power once they held it. The assumption that the state was an instrument that could be wielded by whoever captured it, and that its class character was entirely a product of who occupied it, underestimated how deeply institutional structures shape the behavior of their occupants regardless of their intentions. States do not become neutral when captured by movements with better intentions. They bend those movements toward the imperatives built into the institution: territorial integrity, administrative continuity, the management of populations, the monopoly of legitimate force. Movements that did not theorize this in advance were consistently surprised by it. The position this framework takes is that the state is neither the enemy to be destroyed nor the neutral instrument to be captured but the contested terrain on which the structural conditions for becoming are either expanded or foreclosed. Its class character is real and consequential — it does reflect and reinforce particular distributions of power, and those distributions resist change through the institutional mechanisms analyzed in 7.3 and 7.3b. But it is also the primary mechanism through which collective decisions about the material infrastructure of development

are made, enforced, and revised at scale. The practical question is not whether to engage the state but under what conditions, through what mechanisms, and with what theory of its structural tendencies — a question that 7.2 through 7.4 develop. The implication is that the question is not whether to have a state but what kind of state, organized around what accountability structures, with what capacity for internal reform. A state understood as a permanent instrument of coordination can be designed better or worse for the purposes the framework requires: generating adaptive capacity, distributing the conditions for becoming, and preventing the concentrations of power that eventually capture the instrument itself. The design question is not secondary — it is where most of the practical work of political analysis occurs. An instrument that can be captured by the interests it is supposed to constrain has been designed badly, and redesigning it requires understanding the capture mechanisms before they operate rather than after. The specific problem the state faces that no other institution does is the combination of legitimate coercive capacity with the requirement that it serve interests wider than those of its current operators. Every other institution with significant power either lacks the coercive element (markets, civil society) or lacks the legitimacy requirement (criminal organizations, private militaries). The state's distinctive position creates its distinctive vulnerability: the very features that make it useful for collective coordination — coercive authority, monopoly on legitimate force, organizational scale — make it the primary target for capture by actors who want those features for narrower purposes. The analysis of this problem continues in 7.3, which identifies the mechanisms through which capture occurs and the conditions under which it can be detected and reversed. Coercion and the floor. The evaluative criterion applies to state coercion as it applies to any other institutional output, and applying it yields a bound without requiring a proportionality principle imported from outside the framework. Coercion that secures the biological floor — that prevents the violence, dispossession, and collapse of provision which drive populations below it — produces suffering that the criterion weighs against the suffering produced by not applying it. Coercion beyond that point produces suffering with no corresponding reduction elsewhere, and suffering with no corresponding reduction is what the minimal anchor identifies as unnecessary (3.4).

The bound is therefore not an additional normative commitment; it is the existing criterion applied to a specific institutional capacity. Two consequences follow. The first concerns where the burden of justification sits. An arrangement must be able to name which floor conditions a given coercive capacity secures and what would count as evidence that it has stopped securing them; capacity that cannot be

tied to a named floor condition is capacity the framework cannot evaluate as necessary, and unevaluable capacity defaults to unjustified rather than to permitted. The second concerns a specific target. Coercion directed at contestation is self-undermining by the framework's own diagnostic: the accountability mechanisms identified in 11.0 and 7.3d are the conditions under which an institution can detect its own failures, and coercion of this kind removes them. An arrangement that suppresses contestation in order to maintain the floor destroys its capacity to detect that it has stopped maintaining it. The framework does not specify a threshold, and the absence is not an omission. Repressive capacity is analysed in 15.6 as a variable whose level is set by material conditions, not as a quantity with a correct value that analysis could supply. What the framework specifies is the direction of the burden and the kind of evidence that discharges it. Where the threshold falls in a given arrangement is an empirical question about that arrangement, and treating it as answerable in general would be the same error the framework identifies in doctrines that specify institutional forms in advance of the conditions they must operate in.

7.1b — The Military as Autonomous Institutional Actor

The standard treatment of military institutions in political analysis positions them as instruments of state power — the organized monopoly on legitimate force that the executive commands. This treatment is analytically adequate for contexts in which civilian control of the military is institutionally robust and where military institutions function primarily as the executive's coercive capacity. It is inadequate for the majority of post-colonial states, for transitional regimes of all types, and increasingly for consolidated democracies in which military and intelligence institutions have accumulated organizational autonomy that exceeds the control capacity of nominally superior civilian institutions. In these contexts — which represent most of the world's population — the military must be analyzed as an institutional actor with its own incentive architecture, its own capture dynamics, and its own relationship to the evaluative criterion that cannot be reduced to the objectives of the executive it formally serves. The military's structural distinctiveness as an institution follows from three features that distinguish it from other state institutions. The first is its monopoly on organized violence within its jurisdiction, which creates leverage asymmetries with all other institutions that are not replicated anywhere else in the institutional landscape. When military institutions develop preferences about political outcomes, they can enforce those preferences at costs that no other institutional actor can impose — which means that other institutional actors, including elected governments, must calculate the military's preferences into their decision-making in ways that distort the accountability relationship between governments and the populations they nominally represent. The second is its organizational opacity: military institutions maintain information asymmetries with civilian oversight institutions that are structurally produced by operational security requirements and that function to insulate military decision-making from the accountability mechanisms that apply to other state institutions. The third is its corporate interests: military institutions develop and reproduce organizational interests — in budget allocation, in operational autonomy, in the preservation of specific capabilities and force structures, in the business networks that officer corps maintain through procurement relationships and post-service employment — that are distinct from the security functions the institution nominally performs. The PMN diagnostic for military institutions applies the same capture framework that the analysis applies to all institutions, with adjustments for the military's specific structural features. The central question is whether military institutions function as custodians of the conditions that enable becoming — including the physical security and rule-of-law infrastructure that floor provision requires — or whether they function as autonomou

s power centers that use their monopoly on organized violence to preserve arrangements that foreclose becoming for large populations. The historical record in contexts from Indonesia to Thailand to Pakistan to Myanmar demonstrates that the answer is not determined by the formal structure of civil-military relations. It is determined by whether civilian institutions have developed sufficient organizational capacity, political legitimacy, and accountability mechanisms to make the costs of military intervention in politics higher than the benefits the military would obtain from it — which requires, among other things, that the civilian institutions being protected are worth protecting from the perspective of the populations who would bear the costs of resisting military capture.

7.2 — Progressive Institutionalism

The political stance is progressive institutionalism. Institutions are necessary coordination mechanisms — required as long as the permanent conditions persist, which is indefinitely. They require constant scrutiny, genuine accountability, and structural capacity for substantive reform. The distinction between procedural and substantive reform is critical: procedural reform changes the rules; substantive reform changes the material outcomes. An institution that absorbs every formal challenge without changing what it materially does has captured the form of accountability without the function. Its relationship to three adjacent positions — liberalism, revolutionary politics, and technocracy — matters for understanding what progressive institutionalism is actually claiming. From liberalism: The liberal commitment to institutions as necessary for stable coordination, and to reform as preferable to breakdown when reform remains possible — these are correct. Where liberalism goes wrong is in treating existing institutional forms — constitutional democracies, market economies, international law frameworks — as approximations of the correct solution requiring fine-tuning rather than structural transformation. Existing institutions were designed under specific historical conditions by actors with specific interests, and they produce the outcomes they produce not despite their design but through it. The evaluative question is not whether they conform to liberal procedural standards but whether they actually reduce structural suffering and create conditions for genuine development. When they do, they deserve defense. When they do not, procedural legitimacy is not sufficient protection from the framework's analysis. From revolutionary politics: The revolutionary tradition is right that incremental reform within institutions is sometimes insufficient — that the institutional arrangements being reformed are themselves the source of the conditions requiring reform. The problem is the assumption that destroying existing coordination capacity accelerates the development of better alternatives. Institutions that are destroyed must be rebuilt, and the rebuilding occurs in conditions degraded by the destruction: reduced cooperative substrate, reduced trust, reduced material capacity. The historical record of revolutionary transitions shows that the populations who experience the most structural suffering under existing arrangements also experience the most structural suffering during transitions, because transitions are managed by whoever has organizational capacity during the transition, which is rarely the populations with the most material stake in different outcomes. From technocracy: Complex systems require analytical capacity to manage, and intuition and tradition are insufficient guides to institutional design at scale — on this, technocracy is correct. Where it fails is in drawing the wrong conclusion: that superior analytical capacity justifies concentrate

d decision-making authority. That is precisely the institutional capture mechanism identified in Section 6.4. The check on analytical capacity is not less analysis. It is genuine accountability to those whose material conditions the analysis is supposed to serve, with real enforcement mechanisms — not merely formal consultation that can be managed. Being pro-institution is not the same as being pro any specific institution. The alternative to functional institutions is not freedom from coordination. It is coordination by whoever has the most immediate capacity for force — which is also an institutional arrangement, simply one that is less accountable and more directly extractive. From anarchism: what the framework shares and where it diverges. The anarchist tradition occupies a different relationship to this framework than the three positions above.

Liberalism, revolutionary politics, and technocracy all share the assumption that institutions of some kind are the primary unit of political organization — they disagree about which institutions, how to design or capture them, and who should operate them. Anarchism challenges the assumption itself. In its various forms, anarchism argues that hierarchical institutional coordination is not a permanent requirement of human social life but a historically contingent arrangement that produces the domination it claims to manage — and that the capacity for voluntary, horizontal self-organization is more robust than institutional analysis typically credits. The framework takes from anarchism more than progressive institutionalism's critics usually acknowledge. The anarchist suspicion of institutional capture is structurally identical to the custodian problem in 7.3: the insight that any coordination mechanism sufficiently powerful to manage collective problems is also sufficiently powerful to be captured by actors who will redirect it toward narrower interests is an anarchist insight, not a liberal one. The Ostrom analysis of commons governance (7.3c) demonstrates empirically what anarchist theory has long argued: that communities with adequate social capital, clear boundaries, and effective monitoring can sustain coordination without the hierarchical enforcement that standard institutional theory treats as necessary. At the scale and in the conditions where these arrangements work, they work better than formal institutions — producing higher compliance, lower overhead, and more genuinely accountable outcomes. The divergence is not over whether voluntary, non-hierarchical coordination is desirable. It is over the conditions under which it is achievable, the scale at which it remains functional, and what happens when it is absent. The framework's position is that the permanent conditions identified in Part III — scarcity management, threat coordination, knowledge aggregation, population-scale ecological management — require coordination mechanisms that can operate at scales and under pressures that voluntary ne

tworks have not yet demonstrated they can sustain. This is not a philosophical objection to anarchism. It is an empirical one. The claim is not that hierarchy is inherently necessary but that the alternatives have not been sufficiently developed to manage the coordination problems that currently exist at the scale they currently exist. A movement that reduces or eliminates formal institutional capacity before adequate alternative coordination mechanisms are in place does not produce freedom. It produces a vacuum that is filled by whoever retains the capacity for organized force — which is a worse institutional arrangement than the one that was dismantled, not an absence of institutional arrangement. The framework therefore holds open the question that dogmatic institutionalism closes. Whether a specific institution should be reduced, bypassed, or replaced by alternative coordination arrangements is an empirical question answerable by the framework's evaluative criteria, not a question answered in advance by commitment to institutional form. The evaluation requires examining: what coordination function the institution currently performs; what the realistic alternative is and whether it has demonstrated sufficient capacity at the relevant scale; what the second-order effects of reduction are, including who benefits from the vacuum and how quickly; whether the sequencing allows alternative capacity to develop before the existing capacity is withdrawn; and whether the long-term stability of the alternative arrangement is adequate to the pressures it will face. An institution that produces more structural suffering than it prevents, that has been captured to the point where its accountability mechanisms no longer function, and for which a demonstrably adequate alternative exists — that institution has failed the framework's evaluative test regardless of its formal legitimacy. The obligation is not to defend it but to manage the transition to something better. The framework's openness to anarchist alternatives is therefore conditional rather than categorical — and the conditions are stringent. They are stringent not because institutions deserve deference but because the costs of failed transitions fall primarily on the populations whose conditions were already most precarious. A transition that produces a vacuum, that destroys existing coordination capacity before replacement capacity is functional, or that underestimates the second-order effects of institutional reduction is not a step toward better conditions. It is a redistribution of structural suffering toward those least able to absorb it. The anarchist insight about the pathologies of hierarchy is correct. The anarchist confidence that reducing hierarchy automatically produces better alternatives has not been earned at the scales the permanent conditions require. Evaluative test for

institutional reduction: (1) What coordination function does this institution currently perform, and is that function

ion genuinely necessary at this scale? (2) What is the realistic alternative, and has it demonstrated adequate capacity? (3) What are the second-order effects — who fills the vacuum, and how quickly? (4) Does the sequencing allow alternative capacity to develop before existing capacity is withdrawn? (5) Is the long-term stability of the alternative robust to the pressures it will face? Answering these questions is not a defense of existing institutions. It is the minimum required before dismantling them.

7.2b — Pancasila and Gotong Royong: Pluralism Without Suppression

The Indonesian political tradition offers an empirical case study in the coalition-across-difference problem that PMN's section 8.6 identifies as one of the primary practical challenges for any transformative project: how do you build political community that holds together genuinely different religious, cultural, and ideological commitments without suppressing any of them into a false uniformity or forcing a choice between them? Pancasila — the five principles articulated by Sukarno as the philosophical foundation of the Indonesian Republic — and the related concept of gotong royong (mutual cooperation) as an economic and social principle derived from indigenous village practice, represent an attempt at this synthesis that has specific structural features worth examining analytically. Pancasila's five principles — belief in one God (Ketuhanan Yang Maha Esa), just and civilized humanity (Kemanusiaan yang Adil dan Beradab), the unity of Indonesia (Persatuan Indonesia), democracy guided by deliberation (Kerakyatan yang Dipimpin oleh Hikmat Kebijaksanaan), and social justice (Keadilan Sosial bagi Seluruh Rakyat Indonesia) — were designed to be acceptable to Muslim majority, Christian minority, Hindu Balinese, secular nationalist, and communist constituencies simultaneously, framing a secular-pluralist accommodation within Islamic political discourse (Luthfi Assyauckanie 2009). The analytical interest for PMN is not the specific content of the principles but the structural design challenge they address: the need for a foundation of political community that does not require any constituency to abandon its foundational commitments as the price of participation. This is the criterion unity (as distinct from identity unity) that PMN's section VIII identifies as the appropriate form of coalition across genuine difference. Pancasila's historical record as a framework for managing this challenge is mixed — it has been used to suppress dissent as well as to enable pluralism — but the structural design intention is analytically significant. Mohammad Hatta's concept of gotong royong as an economic principle derived from indigenous village cooperative practice is distinct from Western cooperative economics in a specific way: it grounds the cooperative principle not in ideological commitment or theoretical argument but in existing traditional practice. Gotong royong — the traditional mutual aid practice of Indonesian villages, through which labor and resources are shared for collective purposes without monetary exchange — is not a proposal for how to organize economic life but a description of how significant portions of Indonesian economic life were already organized. Hatta's contribution was to identify this existing practice as the basis for a cooperative economic model compatible with Indonesian conditions rather than importing European socialist institutional forms. This is the parallel development framework (9.4) applied at the national level: institutional

design from the actual material starting position rather than from a universal template derived from different conditions. Two misreadings of progressive institutionalism recur often enough to name directly. The first, from the revolutionary left: that defending institutions is defending the status quo. The framework's response is that the alternative to functional institutions is not freedom from coordination but coordination by whoever has the most immediate capacity for force — which is also an institutional arrangement, simply less accountable and more extractive. The second, from liberal defenders of existing arrangements: that critique of institutional capture is a license for dismantling institutions that are, however imperfectly, performing their stated function. The framework's consistency standard requires applying the same criterion in both directions: defending arrangements that genuinely reduce structural suffering against erosion, and criticizing

arrangements that do not against the claim that their formal existence is sufficient justification for their continuation. A universal healthcare system that is working deserves the same defense from this framework that a captured regulatory agency deserves critique. The criterion is the same; the direction of its application is determined by what the evidence shows about what each arrangement actually produces. Pancasila and gotong royong's contribution to PMN: first, a historical case study in criterion unity — the attempt to build political community across religious, cultural, and ideological difference through shared evaluative principles rather than shared identity, with specific structural features and failure modes that PMN's analysis can examine. Second, gotong royong demonstrates the PMN principle that institutional design should begin from actual material conditions — including existing cooperative practices — rather than from institutional templates imported from different contexts. The failure of Pancasila as a framework for genuine pluralism under Suharto's New Order also provides a case study in how criterion unity can be captured and converted into a tool for suppressing the very difference it claimed to protect — which is the capture mechanism applied to political philosophy itself (see also Robison and Hadiz 2004 on oligarchic reorganization).

7.3 — The Custodian Problem

The custodian problem — the structural dilemma that whoever holds an institutional anchor against power has an incentive to redefine that anchor in self-serving ways, meaning the solution to concentrated power generates the same problem at a higher level — is the central challenge of institutional design that this section addresses. Every stable system requires an anchor — a commitment or threshold that prevents the system from collapsing into pure power struggle. But every anchor must be held by someone, and whoever holds the anchor has an incentive to define it in ways that benefit their position. Finding the right custodian cannot solve this. The incentive structure that corrupts custodianship is produced by the position, not by the character of whoever occupies it. This is the materialist insight applied to the institutions that claim to represent materialism's goals. The vanguard party — the institutional form that Leninist tradition proposed as the solution to the custodian problem — reproduced the problem at a higher level of intensity, because it combined the custodian's incentive to define the anchor in self-serving ways with the ideological authority to treat that definition as scientifically correct. (Lenin 1917) The solution must be structural: distributing the anchor so that no single actor can capture it, making reform mechanisms substantively rather than merely formally accessible, and building in checks that are external to the actors who benefit from the current arrangement. The structural solution to the custodian problem is the distribution of anchor-holding across multiple independent actors with different material interests, such that any single actor's capture of the anchor is contested by others whose interests the capture threatens. This is not a new insight — it is the structural logic behind every serious attempt to build institutional accountability in political history. What varies is the specific mechanism through which distribution is achieved, and the specific conditions under which distribution can be maintained against the constant pressure of consolidation. Separation of functions. The most basic mechanism: different aspects of the anchor are held by different institutions with different bases of authority, different constituencies, and different criteria of success. The authority to make law, to enforce it, and to adjudicate disputes about it is distributed not because the separation is efficient — it is not — but because each function held by a single actor generates the same self-preservation dynamic that corrupts any single custodian. Montesquieu's analysis of separation of powers is the most developed version of this principle in the Western tradition, but the structural logic appears independently across institutional traditions that had no contact with each other. (Ostrom 1990) Heterogeneity of interests among custodians. Distribution of the anchor across actors with genuinely different interests is more r

obust than distribution across actors who formally differ but share material stakes. A legislature, an executive, and a judiciary that are all drawn from the same narrow class with aligned material interests will not generate the genuine contestation that distributes custodianship effectively. They will generate the appearance of separation while the anchor is effectively held by a unified position. The mechanisms that

maintain genuine heterogeneity of interest among custodians — competitive elections with genuine uncertainty, independent funding bases, genuinely different career paths and constituencies — are as important as the formal separation itself. External accountability with genuine enforcement capacity. The most important check on any custodian is accountability to those whose interests the anchor is supposed to protect — not through formal consultation that can be managed, but through mechanisms with genuine enforcement capacity: the ability to remove, to publicize, to withhold cooperation, to impose costs that the custodian cannot absorb without changing behavior. This is why competitive politics, independent press, and genuine civil society capacity matter: not as expressions of liberal values, but as the material enforcement mechanisms that make anchor-holding costly to corrupt. The recursive problem. Every solution to the custodian problem generates the same problem at a higher level. Who watches the constitutional court? Who funds the independent press? Who guards the electoral commission? The regress does not terminate in a final solution. It terminates in the recognition that the custodian problem cannot be solved once — it must be managed continuously, through overlapping and redundant mechanisms, each of which is imperfect and each of which compensates for the failures of others. The institutional architecture that takes this seriously builds in the redundancy and overlap that efficiency thinking would eliminate as waste. The history of institutional breakdown is substantially a history of efficiency-motivated simplification that removed redundant checks and discovered, too late, that what appeared to be redundancy was load-bearing.

7.3b — Early Detection of Institutional Capture

Section 7.3 analyzed the custodian problem structurally — why any anchor-holding arrangement generates capture pressure, and what distribution mechanisms resist it. The operational question is distinct: how does capture actually proceed, what signals indicate it is occurring before it is complete, and what distinguishes early-stage capture (which can be interrupted) from completed capture (which typically cannot be reversed through the same channels that produced it)? Capture is not an event. It is a process with a specific dynamic: it begins at the margin, proceeds through the progressive alignment of institutional incentives with the interests being captured by, and completes when the institution's formal independence is maintained while its substantive independence is eliminated. Completed capture is difficult to detect precisely because the formal indicators — the institution still exists, its stated purpose is unchanged, its personnel turnover appears normal — remain intact. What changes is the pattern of outputs: systematic departure from what the institution's stated purpose would produce, in the specific direction that the capturing interest requires. The Capture Sequence and Early Warning Matrix Capture proceeds through a five-stage causal sequence detailed structurally in §7.3c-i. It is not an instantaneous coup but a compounding progression where each stage lowers the cost and raises the institutional durability of the next. Early detection is crucial because the earlier capture is detected, the lower the counter-power resources required to interrupt it. Once capture reaches stage three or beyond, internal correction mechanisms are disabled and only external crisis or organized counter-power can reverse it. The five-stage early warning matrix identifies specific, externally observable indicators for each stage: Stage 1 Early Warning Indicators (Access Asymmetry): Systematic patterns of privileged, non-public consultation between institutional decision-makers and specific regulated or concentrated interests, unmatched by equivalent access for dispersed affected populations. Growing technical dependency where regulators rely almost exclusively on regulated entities for core domain data, impact models, and analytical framing. Emerging revolving-door pipelines between institutional offices and the concentrated interests they oversee. Stage 2 Early Warning Indicators (Decision-Filter Shift & Preference Expression): Subtle shifts in allowable institutional discourse, where policy alternatives that impose costs on concentrated interests are routinely pre-filtered as 'unrealistic,' 'impractical,' or 'technically unviable' prior to formal deliberation. Systematic procedural discounting of evidence submitted by vulnerable constituencies, alongside immediate procedural accommodation for dominant stakeholder submissions. Stage 3 Early Warning Indicators (Personnel Alignment): Selection, promotion, and appointment patterns that systematically favor

personnel displaying ideological or professional alignment with the narrowed decision filter, while marginalizing or reassigning personnel who raise substantive dissent. Exit of critical institutional memory, frequently rationalized through vague appeals to 'cultural fit' or 'lack of collegiality.' Marked decline in internal debate on decisions with major distributive consequences. Stage 4 Early Warning Indicators (Objective Redefinition & Output Reorientation): Official performance metrics decouple from substantive social outcomes and reorient toward process compliance or output volume that obscures material drift. Absorption and inversion of the institution's founding legitimacy vocabulary: decisions that actively harm the original mandate are celebrated in public rhetoric as its modern, pragmatic fulfillment. Increasing divergence between formal institutional reports and independent material assessments. Stage 5 Early Warning Indicators (Accountability Capture & Consolidation): Internal audit, inspectorate, and appellate bodies are starved of resources, staffed with loyalists, or deprived of jurisdiction. External contestation channels are subjected to procedural paralysis, legal intimidation, or punitive administrative barriers. Institutional resources and legal authorities are openly mobilized to defend the captured configuration against civic investigation. The asymmetry between capture and its detection has a direct implication for institutional design: the mechanisms for detecting early-stage capture must be built in before capture begins, because capture progressively disables the mechanisms that could detect it. An institution that relies on internal whistleblowing to detect capture has already conceded the problem: by the time internal dissent is required to reveal capture, the incentive structure has been shaped to suppress it. The detection mechanisms must be external, independent of the institution's own assessment of its performance, and held by actors with genuine material stakes in the institution performing its stated function. The relationship between this analysis and the leading/lagging indicators of legitimacy crisis in 6.6 is direct: institutional capture is one of the primary mechanisms through which legitimacy deficit accumulates. Populations whose cooperation institutions depend on develop an accurate intuition about capture — the decline in voluntary participation, the shortened planning horizons, the withdrawal of investment in shared outcomes that 6.6 identifies as leading indicators of legitimacy threshold approach — before they can name what they are responding to. The capture detection indicators above make that intuition analytically tractable: they translate the population's behavioral signal into the institutional dynamic that is producing it. The Judiciary as Institutional Layer: Procedural Legitimacy and Its Limits The analysis of the custodian problem so far has focused primarily on the executive and legisla

tive branches — the institutions that make and implement policy. This focus reflects the historical emphasis of political economy and structural analysis, which tends to treat law as the output of political processes rather than as an independent institutional layer with its own causal properties. That treatment is inadequate. The judiciary — courts, constitutional review bodies, administrative tribunals, and the legal profession that operates within them — is a distinct institutional layer with specific structural features that neither the executive nor the legislature shares. Its procedural legitimacy is different in kind: it derives not primarily from electoral authorization or popular mandate but from the authority of precedent, the appearance of impartiality, and the formal independence of adjudication from political instruction. This distinctive legitimacy basis makes the judiciary both a potential site of structural protection for populations with limited political power and a particularly opaque site of structural capture. The structural case for the judiciary as a counter-power resource is specific. Electoral politics systematically underweights the interests of populations who are geographically dispersed, numerically small, politically disorganized, or subject to the kind of visibility suppression that prevents their interests from entering political competition. Courts — particularly constitutional courts and administrative review bodies — provide a procedural channel through which such populations can contest specific institutional decisions without requiring the coalition width that electoral politics demands. Strategic litigation has historically produced structural changes that decades of legislative advocacy had failed to achieve: civil rights litigation in the United States, housing rights litigation in South Africa, environmental litigation in multiple jurisdictions, and labor rights litigation through supranational bodies in Europe are all cases where judicial processes produced structural outcomes that the legislature was

systematically preventing. The mechanism is specific: the judiciary is partially insulated from the immediate majoritarian pressure that captures legislatures, which means it can be used to enforce general principles against specific applications that majorities would otherwise permit. The structural vulnerability of the judiciary is symmetrical with this strength. The same insulation from immediate political pressure that makes courts useful for protecting minority interests makes them opaque sites of elite capture that are difficult to detect and contest through normal political channels. Judicial appointment processes — in most systems controlled by the executive or legislature — are the primary vector through which long-term ideological transformation of judicial institutions occurs without generating the visible political accountability that electoral processes require. A government that controls judici

al appointments for a decade can transform the substantive direction of judicial review for a generation, in ways that produce structural outcomes that no single legislative act would have produced and that are far more difficult to reverse. The capture dynamic operates through the progressive selection of judges whose interpretive commitments systematically favor specific structural arrangements, which is not detectable as capture from within the procedures of the institution itself because those procedures were designed to evaluate legal reasoning, not structural outcomes. The distinction PMN draws between procedural legitimacy and structural legitimacy is essential for judicial analysis. A judicial decision has procedural legitimacy if it was reached through recognized legal procedures by a properly constituted tribunal applying recognized legal standards. It has structural legitimacy, in PMN's terms, if the standards it applies and the outcomes it produces are consistent with the evaluative criterion — if they reduce structural suffering and expand adaptive response capacity rather than protecting arrangements that foreclose it. The procedural legitimacy of legal systems is real and matters: it provides a framework within which structural claims can be advanced, and destroying it produces coordination vacuums that consistently harm the most vulnerable populations first. But procedural legitimacy does not guarantee structural legitimacy, and the conflation of the two is one of the primary mechanisms through which legal systems function as instruments of structural harm while maintaining the appearance of neutral adjudication.

7.3c — Incentive Architecture and Predictable Institutional Failure

Sections 7.3 and 7.3b analyzed the custodian problem and the capture sequence: why any anchor-holding arrangement generates capture pressure, and how that pressure proceeds through observable stages. The analysis must extend to incentive architecture: the structural features of institutional design that make certain failure modes not merely possible but predictable, regardless of the intentions of the actors who occupy institutional positions. The custodian problem describes what happens when an institution is targeted by capture from outside. Incentive architecture analysis describes what happens when an institution fails from inside — when the normal functioning of the incentive structure built into its design produces outcomes systematically different from its stated purpose, without any actor intending that result. The foundational observation is that institutions do not operate on the actors who occupy them as passive instruments. They produce specific incentive structures that shape what those actors find rational to do — and the incentive structure produced by institutional design often diverges predictably from the behavior required to fulfill the institution's stated function. This is the insight that principal-agent analysis formalized and that the framework must integrate: the principal (the population the institution nominally serves) and the agent (the personnel who operate the institution) have different information, different interests, and different exposure to the consequences of institutional decisions. The divergence is not produced by individual corruption or bad character. It is produced by the structural gap between what the principal can observe and reward, and what the agent can actually do and benefit from. The detailed analytical architecture for capture dynamics — mechanisms, sequential stages, early detection criteria, and the cosmological variant — is developed in 7.3c-i and 7.3c-ii. The Principal-Agent Gap In any institution of significant scale, those who are supposed to benefit from it cannot directly observe most of what those operating it actually do. They can observe outputs

— the services delivered, the decisions made, the documents produced — but not the processes, the trade-offs, or the information environments within which those outputs were generated. This observation gap creates the conditions for systematic divergence: agents can rationally pursue their own interests — security of tenure, minimization of effort, maximization of internal status, avoidance of conflict with powerful actors — while maintaining the appearance of pursuing the principal's interests. The divergence is not dishonesty in the ordinary sense. It is the normal response of rational actors to incentive structures that reward the appearance of performance more reliably than actual performance, because appearance is what the principal can observe and actual performance often is not. The institutional design response to this gap has been t

he development of monitoring mechanisms: performance metrics, audit systems, reporting requirements, accountability procedures. These mechanisms address the observation gap by making more of the agent's behavior observable. They produce a predictable secondary effect: agents optimize for the metrics rather than for the underlying objectives the metrics were designed to track. When a measure becomes a target, it ceases to be a good measure — because the behavior that produces good metric scores is not identical to the behavior that produces good institutional outcomes, and when the metric is what is rewarded, the divergence is exploited. This is not a failure of the monitoring mechanism in individual cases. It is its structural consequence at scale: the monitoring system changes the incentive structure it was designed to reveal, producing a new divergence between measured performance and actual performance that is harder to detect because it is fully compliant with the measurement system. Rent-Seeking as Rational Institutional Behavior

Rent-seeking — the use of institutional position to capture value rather than to create it — is often treated as a pathology of bad institutional design or corrupt individual actors. The framework's analysis requires treating it as a rational response to a specific class of incentive structures: those in which controlling access to resources or decisions is more rewarding than improving the quality of the decisions made. Wherever the institutional arrangement creates gatekeeping positions — positions whose value derives from controlling access rather than from producing outcomes — rent-seeking behavior is the rational response of agents in those positions. It does not require corrupt intentions. It requires only that the institutional incentive structure reward position-holding more than outcome-production, which is the normal condition of institutions that have lost connection between the welfare of those they serve and the internal rewards that drive agent behavior. The connection to the custodian problem is direct. Rent-seeking behavior within institutions makes them more susceptible to external capture, because agents already oriented toward value extraction from their positions are structurally prepared to extend that orientation to arrangements with external parties who can offer additional rewards. Stage 1 of the capture sequence — access asymmetry through informal channels — is not possible without agents who find informal engagement with external parties a natural extension of how they already relate to their institutional positions. The external capture and the internal incentive drift are not independent failures. They are the same failure operating through different channels simultaneously. The Accountability Trap

Accountability mechanisms designed to address principal-agent divergence face a structural trap: the most effective accountability is external and independent; but the institution being held ac

countable controls the conditions under which external accountability can operate. Institutions that have drifted from their stated purpose have strong incentives to shape the accountability mechanisms that would detect and correct that drift. They do this not through obvious obstruction — which would itself be a detectable signal — but through the progressive definition of accountability as compliance with internal procedure rather than as achievement of external outcomes. An institution that has successfully redefined accountability as procedural compliance has accountability mechanisms that function perfectly while the institution systematically fails its stated purpose. The practical implication is that accountability mechanisms must be designed with the accountability trap explicitly in view. Effective accountability has three features that procedural accountability systematically lacks: outcome rather than process measurement — evaluating what the institution produces for those it serves rather than whether it followed its own procedures; independence from the institution being evaluated — held by actors with genuine material stakes in the

institution performing its stated function rather than by actors whose positions depend on the institution; and standing for the populations served rather than the populations operating the institution to trigger evaluation and challenge findings. The absence of any one of these three features typically produces accountability mechanisms that protect the institution from accountability while appearing to provide it. Incentive architecture analysis produces a specific diagnostic question that the custodian problem analysis and the capture sequence analysis do not: not whether this institution is being captured from outside, but whether its internal incentive structure aligns agent behavior with the institution's stated purpose. The two questions are related but not identical, and both must be asked. An institution with well-designed external accountability and a misaligned internal incentive structure will drift toward failure through its normal functioning, without any external capture being required. An institution with well-aligned internal incentives but inadequate external accountability remains vulnerable to capture even when internal drift is absent. Institutional robustness requires attending to both simultaneously. **Judicial Capture: Mechanisms and Detection** The incentive architecture of legal institutions produces capture patterns that are distinct from those of executive and legislative institutions, and the standard detection methods developed for those branches are poorly calibrated for judicial capture. Executive capture tends to produce visible policy outputs — spending decisions, regulatory choices, personnel appointments — that can be tracked against stated objectives. Legislative capture produces voting patterns and legislative texts that are publicly recorded and can be traced to sponsori

ng interests. Judicial capture produces interpretive patterns — specific readings of constitutional provisions, statutory language, and precedent — that require legal expertise to identify as departures from alternative interpretations, and whose political significance is often opaque to the populations most affected by them. Three mechanisms of judicial capture warrant explicit analysis. The first is appointment capture: the systematic selection of judges whose interpretive dispositions favor specific structural arrangements, through appointment processes controlled by already-captured executive or legislative institutions. This is the most durable form of capture because its effects compound across the tenures of appointed judges and because the remedy — changing the appointment process or the appointing institutions — requires the same structural transformation that appointment capture was designed to prevent. The second is procedural erosion: the progressive restriction of standing rules, justiciability doctrines, and access to legal remedies in ways that systematically exclude the populations most subject to structural harm from the procedural channels through which they could contest it. When courts narrow the definition of who can bring a claim, what claims are cognizable, and what remedies are available, the effect is to convert formal legal rights into practically inaccessible protections — a form of visibility suppression that operates within the legal system itself. The third is doctrinal drift: the gradual accumulation of precedent that progressively reinterprets general principles in directions that favor specific structural arrangements, without any single decision constituting a visible break from prior law that would be recognizable as capture from outside the legal profession. The relationship between judicial institutions and social movements requires specific analysis because it differs from the relationship between movements and legislative or executive institutions. Legislative and executive institutions respond to movements primarily through the mechanism of electoral pressure — movements affect outcomes by changing the composition of those institutions or by credibly threatening to do so. Courts respond to movements primarily through the mechanism of litigation strategy: movements affect outcomes by identifying cases that present favorable legal questions, developing legal arguments that reframe those questions in terms of applicable doctrine, and building the record of evidence and precedent that supports those arguments. This is a more technically mediated relationship that requires different organizational capacities and operates on longer timescales. The implication for structural analysis is that counter-power formation directed at judicial institutions requires sustained investment in legal capacity — in organizations that can identify strategic litigation opportunities, develop legal arguments that advance structural

claims, and maintain the institutional continuity required to pursue cases through appellate processes that can span decades. **Judicial sector diagnostic:** (1) Appointment process — who controls the selection of judges,

through what procedures, with what accountability to populations most affected by judicial outcomes? A judiciary whose appointment process is captured by executive institutions that are themselves captured by private interests is a second-order capture amplifier, not a custodian. (2) Access architecture — what populations can bring claims, at what cost, through what procedural channels? Systematic access barriers function as structural exclusion even when formal rights exist. (3) Interpretive drift — does the pattern of decisions in specific domains consistently favor arrangements that foreclose becoming for particular populations, and does this pattern correspond to the interpretive dispositions of judges appointed through captured processes? (4) Remedy availability — when rights violations are found, do available remedies produce structural change or merely formal acknowledgment? A legal system that identifies violations without producing remedies that address their structural causes is functioning as legitimacy production for the arrangements it nominally constrains. The internal vulnerabilities analyzed above — principal-agent divergence, metric gaming, and rent-seeking — do not remain confined within organizational boundaries. They create the structural entry points through which external interests systematically penetrate, redirect, and consolidate control over institutions. The comprehensive causal mechanism, stages, resistance mechanisms, and diagnostic criteria for this progression across all organizational forms are formalized authoritatively in §7.3c-i.

7.3c-i — A General Theory of Capture: Mechanism, Stages, and Detection

Institutional capture is the causal process by which concentrated interests that an institution is designed to constrain, regulate, or govern come to systematically control its decisions, resource allocations, and legitimacy vocabulary. The Custodian Problem (§7.3) identifies why any anchor-holding arrangement generates persistent capture pressure; this section articulates the comprehensive causal sequence through which capture proceeds across diverse organizational forms. Capture operates as a probabilistic causal progression rather than an unalterable historical fate. Each stage lowers the transaction cost and raises the structural durability of the subsequent stage, creating path dependency. However, capture is not mechanically linear: stages may overlap, stall under shifting material constraints, or be interrupted by organized counter-power. The authoritative five-stage sequence comprises: (1) Access Asymmetry: Concentrated interests establish a systematic informational, technical, and relational advantage over institutional overseers. Regulators and decision-makers come to rely on the regulated for specialized expertise, data models, and professional career pipelines. At this stage, capture presents innocuously as technical competence and administrative efficiency. (2) Decision-Filter Capture & Preference Expression: Actors possessing asymmetric access actively express preferences through institutional channels. The institution's cognitive and procedural filters narrow, systematically excluding information and policy choices that threaten the dominant interest. Decisions remain locally defensible as pragmatic or legally prudent, but cumulatively bias institutional direction. (3) Personnel Alignment: Recruitment, promotion, and retention dynamics systematically select for personnel whose professional orientations fit the narrowed decision filter, while dissenting or non-compliant personnel are marginalized, reassigned, or driven out. This internal alignment cements the access asymmetry by staffing veto points with sympathetic custodians. (4) Objective Redefinition & Output Reorientation: The operative goals of the institution diverge decisively from its statutory mandate. Crucially, the institution absorbs and repurposes its founding legitimacy vocabulary: what was previously defined as regulatory evasion or mission failure is formally rearticulated as modern, flexible mandate fulfillment. Outputs consistently serve the captured interest while maintaining symbolic fidelity to public purposes. (5) Accountability Capture & Consolidation: The channels of review, contestation, and appellate correction are dismantled, defunded, or captured by the same interests. The captured configuration achieves institutional lock-in: the beneficiaries control the adjudication of grievances against themselves and command institutional resources to suppress counter-claims. Tractability and Intervention Dynamics: Detecting capture at stages one and two represents

the only tractable window for low-cost institutional reform. By stage three, internal reform capacity is severely impaired; by stages four and five, capture is self-insulating and cannot be reversed from within the institution's own rules. Reversal then requires either severe systemic crisis or organized external counter-power capable of overriding captured procedures. The standing diagnostic questions from §7.3 — who benefits from the output distribution, who audits the auditors, what counts as failure, and whether revision from outside is legally and materially possible — must be applied continuously before consolidation occurs. The distinction between capture and ordinary organizational evolution requires precision. Organizations legitimately change their objectives in response to changing conditions, new evidence, and the development of their members' understanding. The diagnostic for capture is not change itself but the process through which change occurs and the direction it consistently takes. Capture is indicated when: objective changes consistently benefit the actors with the highest internal access at the expense of the populations the organization nominally serves; personnel decisions consistently select for orientation alignment over competence or representativeness; information flows are systematically shaped to prevent the evidence of objective drift from reaching members or constituencies who would contest it; and accountability mechanisms that would allow members or constituencies to contest the direction of change are progressively restricted or circumvented. Any single indicator is insufficient; the pattern across all four is diagnostic. Effective capture-resistance mechanisms across all organizational forms share a structural foundation: they maintain multiple independent channels of information, decision-making, and oversight so that no single coalition can monopolize institutional control. Key structural safeguards include: (1) Mandatory rotation and tenure limits on high-access positions to prevent entrenched relational monopolies; (2) Strict structural separation between operational implementation and outcome evaluation; (3) Enforced transparency of resource flows, lobbying interactions, and appointment deliberations directly to outside constituencies; and (4) The preservation of viable exit options and protected channels for organized counter-power. Capture diagnostic for any organization: (1) Access asymmetry — which actors have disproportionate access to decision-making, information, and resource allocation, through what structural mechanisms, and whether that asymmetry has been increasing? (2) Personnel alignment — does the pattern of hiring, promotion, and expulsion over time show selection for orientation alignment independent of other criteria, and who controls that pattern? (3) Objective drift — does the gap between stated objectives and operational priorities show a consistent directional drift, and whose interests

does the drift serve? (4) Accountability restriction — have the mechanisms through which members or constituencies can contest organizational direction been narrowed, formalized into ineffectiveness, or transferred to actors who benefit from the current configuration? The presence of all four patterns simultaneously is a strong indicator of capture regardless of the organizational type or the ideological commitments of the capturing actors.

7.3c-ii — Cosmological Capture: When Belief Distorts the Institution From Within

The general capture theory of section 7.3c identifies the mechanism through which institutional decision-making is progressively redirected from its stated purpose toward the maintenance of those who control it. The dominant mechanism is interest: access asymmetry, resource dependence, and the gradual alignment of institutional objectives with the objectives of those who can most effectively punish or reward the institution's operators. But there is a second capture mechanism that is analytically distinct from interest capture and in some respects more dangerous: capture by cosmology. An institution is cosmologically captured when its operators' decision-making is governed not primarily by their material interests but by a belief system — specifically, a belief about the structure of history, the meaning of the present moment, and the necessity of specific outcomes — that systematically distorts their reading of evidence and their response to feedback. The mechanism of cosmological capture differs from interest capture in three important ways. First, the captured actor does not experience themselves as captured. Interest capture typically produces some

degree of self-awareness in its advanced stages — the official who knows they are protecting a donor, the regulator who understands their dependence on the industry they regulate. Cosmological capture produces the opposite: the actor who is most deeply captured is most certain that they are acting from principle rather than interest, because the cosmological framework they operate within has already pre-interpreted all evidence in its own terms. Disconfirming evidence is not evidence against the framework — it is evidence of the adversary's deception, the test before the final vindication, or the necessary darkness before the dawn. The framework is epistemically closed in the precise sense that section 1.6b identifies as the institutional failure mode: belief has consumed the doubt that would allow revision. Second, cosmological capture does not require a captor. Interest capture requires actors who benefit from the capture and who, even if unconsciously, shape the institutional environment to produce and maintain it. Cosmological capture can self-generate: a belief system that provides a compelling account of why the present moment is historically decisive, why the actors who hold it are on the right side of a cosmic struggle, and why the normal costs and constraints of institutional decision-making do not apply to this case — such a belief system recruits its own captives without requiring external manipulation. The historical record of millenarian movements, revolutionary vanguards convinced of their historical necessity, and national security establishments organized around existential threat frameworks all demonstrate this self-generating dynamic: the belief produces the institutional behavior that the belief predicts is necessary, which produces the adversarial responses that confirm the belief.

ef. Third, and most critically for PMN's evaluative criterion: cosmologically captured institutions do not merely redirect resources toward their operators' interests. They actively configure institutional arrangements — including the institutions that govern conditions of becoming for large populations — to realize outcomes that the cosmological framework designates as necessary, regardless of whether those outcomes reduce or increase structural suffering. An interest-captured institution is parasitic: it extracts value from the populations it nominally serves while continuing to provide enough of its stated function to maintain the conditions for continued extraction. A cosmologically captured institution can be genuinely destructive: it may pursue outcomes that its operators believe are cosmically required even when those outcomes visibly and demonstrably increase structural suffering, because the framework has already supplied the interpretation under which that suffering is necessary, temporary, or redemptive. *The Self-Fulfilling Structure: How Cosmological Belief Produces the Reality It Predicts* The most analytically significant feature of cosmological capture is its capacity to produce self-confirming historical trajectories. When decision-makers at institutional leverage points hold genuine eschatological beliefs — beliefs that the present moment is a culminating conflict between forces whose outcome is cosmically predetermined — their decisions begin to shape events in directions consistent with those beliefs, which then appear as confirmation of the beliefs' accuracy. This is not simply the general phenomenon of self-fulfilling prophecy. It is a specific institutional dynamic: the belief is held by actors with the institutional power to actually configure the material conditions that the belief predicts, which means the prediction is not merely believed but partially constructed. The mechanism runs in two directions simultaneously. In the first direction, the cosmologically captured actor interprets ambiguous evidence as confirming the framework: a diplomatic overture from the adversary is read as deception rather than genuine negotiation, because the framework has already specified that the adversary is fundamentally committed to the conflict. In the second direction, the actor's responses to that evidence produce the institutional behaviors that make genuine negotiation progressively less possible — harder rhetoric, escalatory preparations, the marginalization of internal voices that read the evidence differently — which then produce adversarial responses that retrospectively justify the original interpretation. The framework and the reality it describes become progressively harder to distinguish, not because the framework was accurate but because the power of the actors holding it was sufficient to partially realize it. The diagnostic importance of this mechanism for PMN is that it generates a specific form of structural suffering th

at cannot be adequately analyzed without it. Interest-captured institutions produce suffering as a byproduct of extraction: the suffering is instrumentally generated, and the institution would in principle prefer to produce it with less visibility. Cosmologically captured institutions can produce suffering as a feature: the framework may actively designate certain populations as on the wrong side of the cosmic struggle, making their suffering not a byproduct to be minimized but a confirmation to be interpreted. This distinction matters for the intervention analysis: interest capture is addressed primarily by changing incentive structures and accountability mechanisms. Cosmological capture requires additionally addressing the epistemic infrastructure through which the belief is produced, maintained, and insulated from disconfirmation. Distinguishing Cosmological Capture from Ideological Commitment PMN's treatment of ideology as a material force — the analysis of how governing narratives function as structural variables rather than mere justifications — might appear to already cover what cosmological capture describes. The distinction matters. Ideological commitment of the ordinary kind is susceptible to evidence: an actor who is committed to a governing narrative that makes specific empirical predictions will experience pressure to revise when those predictions fail. The institutional processes through which ideological revision occurs may be slow and politically costly, but they exist within the framework: the ideology is a map, and maps can be updated when they demonstrably fail to match the territory. Cosmological capture differs precisely in that the belief is structured to be immune to this revision process. The specific feature that distinguishes cosmological capture from ordinary ideological commitment is the eschatological architecture: the belief that the current moment is the culminating phase of a cosmic struggle, that the outcome is already determined at a level beyond ordinary historical contingency, and that the apparent costs and contradictions of the present are necessarily temporary. This architecture converts disconfirming evidence into confirming evidence — not through cynical reinterpretation but through genuine belief — and thereby removes the feedback mechanism that allows ordinary ideological commitment to update. An actor operating within a genuinely eschatological framework is not being dishonest when they interpret military defeat as a test, economic collapse as purification, or widespread suffering as necessary sacrifice. They are applying the framework as it is designed to be applied. The PMN diagnostic for cosmological capture therefore focuses not on the content of the belief but on its epistemic architecture: is the belief structured such that any outcome — victory or defeat, growth or collapse, acceptance or rejection — can be interpreted as confirmation? If so, the belief is functioning cosmologically in the relevant

sense regardless of its specific content, and the institutional behavior it produces should be analyzed through the cosmological capture framework rather than the ordinary interest capture framework. Revolutionary vanguardism, certain forms of religious nationalism, specific configurations of national security exceptionalism, and some variants of accelerationist ideology all share this epistemic architecture despite their structurally distinct content. Cost Asymmetry as Structural Weapon: The Material Complement Cosmological capture addresses the belief dimension of a specific institutional failure mode. Its material complement — the structural condition that makes cosmologically captured actors particularly dangerous at the geopolitical and systemic level — is cost asymmetry: the configuration in which the cost of imposing harm on an adversary is structurally lower than the cost of defending against that harm. This is distinct from, and more analytically precise than, PMN's existing Rc (repressive capacity) variable, which measures the magnitude of repressive capacity without accounting for its cost structure. Cost asymmetry matters because it determines the dynamics of attrition in sustained conflict. An actor with high absolute capacity but high cost-per-unit-of-harm faces a fundamentally different strategic environment than an actor with lower absolute capacity but low cost-per-unit-of-harm. The latter can impose costs faster than the former can absorb them even when the absolute power differential appears to favor the former — not because of superior strength but because of structural cost differentials that compound over time. The historical cases are numerous: small powers using low-cost disruption to impose costs on larger powers whose defense infrastructure was calibrated to symmetric threats; insurgencies whose operational cost-per-incident was orders of magnitude below the counterinsurgent's cost-per-response; economic actors using regulatory and

legal mechanisms whose cost of deployment was low relative to the target's cost of defense. The intersection of cosmological capture and cost asymmetry produces the most analytically significant configuration for PMN's geopolitical analysis: a cosmologically captured actor who also commands cost-asymmetric capabilities has both the epistemic insulation against feedback that would normally constrain escalatory behavior and the structural capacity to impose costs faster than adversaries can adapt to them. The institutional response to this configuration — which neither pure capability analysis nor pure ideological analysis can adequately generate — requires simultaneously addressing the epistemic architecture maintaining the cosmological capture and the structural cost differentials that make the captured actor's escalatory behavior materially sustainable. Cosmological capture diagnostic: (1) Epistemic architecture check — is the governing belief of the institution or its key operators struct

ured such that disconfirming outcomes can be reinterpreted as confirming? If victory, defeat, growth, and collapse can all be incorporated as expected features of the framework, the belief is functioning cosmologically. (2) Feedback immunity test — when predictions derived from the belief fail, does the institution revise the belief or revise the interpretation of the outcome? Repeated outcome-reinterpretation without belief revision is the primary behavioral indicator of cosmological capture. (3) Self-fulfilling trajectory assessment — are the institution's decisions producing the adversarial conditions the belief predicts, in ways that the institution then uses as confirmation? Identify the specific decision-response sequences through which the belief is being partially constructed rather than merely predicted. (4) Suffering designation check — does the institutional framework designate the suffering of specific populations as necessary, redemptive, or cosmically justified? Cosmological frameworks that produce designated suffering require epistemic intervention, not merely incentive restructuring. (5) Cost asymmetry complement — does the cosmologically captured actor also command cost-asymmetric capabilities? If yes, assess the compounding interaction: epistemic insulation plus structural cost advantage produces the configuration most resistant to ordinary accountability mechanisms and most liable to sustained escalatory behavior. Early warning indicators that permit detection before capture is complete: Early warning indicators for cosmological capture: (1) Disconfirmation-as-confirmation — evidence that challenges the framework is systematically reinterpreted as confirming it (persecution of believers proves the framework's threat analysis is correct; failure of prediction proves the framework needs application, not revision). (2) Post-hoc adjustment — predictions fail but the framework is adjusted after the fact to accommodate the failure without revising the core commitment; new auxiliary hypotheses proliferate to explain away disconfirming cases. (3) Redemptive suffering — suffering of designated populations (in-group or out-group) is framed as necessary, purifying, or cosmologically required rather than as structural harm requiring structural response. This is the diagnostic indicator that distinguishes cosmological from interest capture: interest capture produces suffering as a side effect; cosmological capture produces it as a feature. (4) Epistemic closure — internal dissent is treated as contamination or betrayal rather than as a resource for framework revision; the community of inquiry collapses to a community of affirmation. (5) Boundary tightening under pressure — as external evidence mounts against core claims, the definition of who counts as a legitimate interpreter narrows rather than broadens; critique from outside is categorically disqualified before its content is examined. When three or more indicators are present simultaneo

usly, the arrangement meets the cosmological capture diagnostic. The appropriate response is the epistemic intervention described in the next section (7.3d), not the incentive restructuring that addresses interest capture.

7.3d — Institutionalizing Anti-Capture

The custodian problem (7.3) cannot be solved, but it can be managed through institutional design. The preceding sections have analyzed how capture happens — the mechanisms, stages, and cosmological variants. This section addresses the complementary question: what institutional structures have reduced the

rate and extent of capture in historical cases, and what design principles can be derived from that record? The framework's epistemological commitment in Part I applies here with particular force: no design is universally effective, context determines which mechanisms are available and which will be subverted, and the history of anti-capture institutions is also a history of their eventual capture. What follows is a typology of mechanisms, not a blueprint. Nested accountability structures distribute oversight responsibility across multiple levels and constituencies such that capture at any single level does not eliminate the accountability function. The design principle is that each node in the accountability structure should be accountable to a constituency whose interests are not aligned with those of the entity being overseen. Worker-owned enterprises accountable to their members, community boards accountable to affected populations, appellate bodies accountable to broader publics than the institutions they review — each layer introduces a constituency whose material interest in accurate assessment differs from the assessed entity's interest in favorable assessment. Nested accountability does not prevent capture; it requires that capture be replicated across multiple levels simultaneously, which increases the resource cost and reduces the durability of capture that does occur. Sunset provisions and mandatory re-authorization force institutional arrangements to demonstrate continued justification for their authority rather than assuming it. The mechanism works against the natural tendency of institutions to accumulate authority over time and to resist revision as the conditions that justified their creation change. Regulatory bodies, oversight commissions, and policy mandates that expire unless explicitly renewed force a periodic re-examination of whether the arrangement still serves the function it was created for, who now controls it, and whether alternatives would serve the function better under current conditions. Sunset provisions are most effective when the renewal process is structured to surface evidence from affected populations rather than from the institutions subject to oversight — which requires deliberate design, as institutions facing renewal will characteristically direct resources toward shaping the renewal assessment. Randomized selection for oversight bodies removes the self-selection bias that allows capture to occur through the strategic placement of sympathetic personnel. When oversight roles are filled by election or appointment, actors with interests in favorable oversight have structural incentives to invest in influenc

ing the selection process; over time, the population of overseers shifts toward those whose orientations are compatible with the interests of those being overseen. Randomized selection from a defined eligible population — as in jury selection, citizens' assemblies, and some historical models of sortition — does not eliminate the possibility that individual overseers will be influenced, but it substantially reduces the feasibility of systematic capture through selection. The mechanism's effectiveness depends on the construction of the eligible pool: if the pool itself is defined in ways that exclude affected populations, randomization within it replicates rather than corrects the exclusion. External independent review with genuine enforcement capacity differs from advisory oversight in that the reviewing body has authority to compel disclosure, impose consequences, and act without the consent of the reviewed entity. The design requirement is genuine independence: the reviewing body must not depend on the reviewed entity for its resources, its mandate, or the career prospects of its personnel. Most existing external review mechanisms fail this requirement in one or more dimensions — they are funded by the industries they oversee, appointed by entities with interests in favorable assessment, or structured so that their findings require validation by the very institutions whose conduct is in question. A reviewing body that meets the independence requirement and has genuine enforcement authority is rare historically; its rarity reflects the consistent pressure that captured institutions exert against the creation and maintenance of effective oversight. Worker-owned platform cooperatives, as a recent partial case, illustrate both the potential and the limits: they restructure the accountability relationship within the platform but do not resolve the questions of how platform cooperatives themselves remain accountable to the broader populations their infrastructure affects. The recurring pattern: anti-capture mechanisms tend to be effective when the conditions that produced capture are still recent enough that the political will to prevent repetition exceeds the organizational capacity of interested parties to neutralize the new mechanism. Over time, the same interests that capture the original institution direct

resources toward capturing, defunding, or hollowing out the anti-capture mechanism. This is not an argument against institutional design; it is an argument that anti-capture design must be combined with the maintenance of counter-power (10.9) sufficient to sustain the political will that makes the mechanisms effective. Design without counter-power produces mechanisms that are formally present and functionally inert. Counter-power without design produces resistance that succeeds in specific contests without building durable structural resistance to the next capture sequence.

7.4 — The State and Long-Term Direction

The long-term orientation is toward political arrangements that are more globally coordinated, less driven by the competition of territorially bounded states, and more directly accountable to human material welfare across the boundaries that currently limit it. This is a horizon in the technical sense the framework uses throughout: a reference point that orients direction without implying arrival is possible or expected. The term is used deliberately rather than metaphorically — an actual horizon recedes as you approach it, always revealing more terrain between here and there. Moving toward a horizon is meaningful and purposive. Expecting to stand at one is a category error. The long-term direction is real. Its status as an unreachable limit is also real, and not a defect. The transition requires building the material conditions that make more coordinated arrangements possible. Working through states, reforming international institutions, building the capacity for collective decision-making that can operate at scales that nation-state competition cannot manage — these are the present-tense activities that the long-term direction requires. A premature transition that destroys the coordination capacity required for further development is not a step toward a more coordinated world. It is a step backward, and the materialist analysis of what happened when such steps were taken historically supports this judgment. What the long-term direction is can be stated with more precision than the concept of 'more coordinated arrangements' provides. The permanent conditions identified in Part III — scarcity management, threat coordination, knowledge aggregation, ecological constraint — are currently managed by a combination of nation-state competition, partial international institutions, and market mechanisms, none of which is adequate to the scale at which the problems now operate. Climate change, pandemic risk, nuclear weapons, and the governance of artificial intelligence are each problems whose consequences are not bounded by the jurisdictions that currently hold decision-making authority over them. The long-term direction is not a particular institutional form — not world government, not a specific international architecture — but the expansion of the capacity for binding collective decision-making to the scale at which the problems exist. That expansion is a directional commitment without a predetermined terminal form. The implication for present-tense political activity is specific. States are currently the primary sites of democratic accountability and the primary instruments through which populations can exert collective control over their conditions. Weakening state capacity in the name of internationalism or post-national politics does not expand collective decision-making capacity at the global level — it reduces it at the national level without building adequate alternatives. The practical path toward more coordinated arrangements runs through stre

ngthened states that are capable of entering and honoring binding international commitments, not through the erosion of state capacity in favor of non-accountable alternatives. A state too weak to implement its own commitments cannot be a reliable partner in international coordination. A population that has lost confidence in national institutions will not support the international transfers and constraints that global coordination requires. The same sequencing logic that governs domestic institutional transition (10.6) applies here. The historical cases in which movements attempted to bypass nation-state structures in favor of direct internationalism — the early Comintern, certain strands of anti-colonial politics — consistently encountered the problem that the international structures they tried to build required capacities that only functioning states could provide. The lesson is not that internationalism is wrong. It is that the sequence matters: build the national capacity that makes reliable international commitment possible, then extend decision-making

authority to the international level as the institutions capable of exercising it develop. The direction is set. The pace is determined by what the existing structures can support. The tension between the short-term incentive structure of states and the long-term direction the framework identifies is not incidental. States are organized to compete with each other, to manage the immediate interests of their governing coalitions, and to maintain the legitimacy that depends on delivering outcomes visible within electoral or political cycles. The long-term direction — toward arrangements that are more globally coordinated and directly accountable to human welfare across territorial boundaries — requires exactly the investments that short-term state competition undervalues: in supranational institutions, in global commons governance, in the reduction of competitive advantage that comes from not internalizing costs that fall on other states or on future populations. This is not a failure of political will; it is the predictable output of a structural incentive architecture that the framework's progressive institutionalism must engage with rather than assume away.

The state as current instrument: The state is the primary instrument through which counter-power is currently organized, rights are enforced, redistributive arrangements are maintained, and the material floor is defended. This makes states indispensable in the near term and potentially obsolete over the longer term — and managing the transition between those two assessments is one of the most demanding analytical problems the framework faces. A political practice that treats the state as merely an obstacle, rather than as the best currently available instrument for specific crucial functions, will consistently produce outcomes that serve those who benefit from state functions failing: the populations most dependent on state provision of t

he floor conditions bear the heaviest costs when states are weakened without replacement infrastructure in place. The framework's position: the long-term direction is away from the current configuration of territorial states as the primary units of political organization. The transitional path to that direction runs through stronger states that are more effectively accountable to the populations they govern and less captured by the governing coalitions that manage them — because stronger and more accountable states are better positioned to build the supranational infrastructure that the long-term direction requires. Weakening states in the name of the long-term direction consistently produces intermediate configurations that are worse than what they replaced.

7.5 — Noble Lie: Fictions That Serve Power

Plato's concept of the noble lie — a fiction deliberately constructed and maintained because it is believed to serve social cohesion — is one of the most honest admissions in the history of political philosophy. It acknowledges that political arrangements depend on narratives that are not strictly true, and that some of those narratives are constructed deliberately by those who govern. The question the framework must answer is not whether noble lies exist — they do — but how to distinguish the ones that genuinely serve collective interest from the ones that serve elite interest while claiming to serve everyone. The test this framework proposes is power asymmetry. A noble lie that genuinely serves collective interest is one whose maintenance does not primarily benefit those who control it. A noble lie that serves elite interest is one that, if abandoned, would threaten the position of those who maintain it. The noble lies that sustain existing arrangements are well documented once the framework for identifying them is in place. Markets are political constructions, not natural facts. Meritocracy produces outcomes that reflect inherited position more than individual capacity. Current distributions of political power reflect the accumulated advantage of those who have controlled political competition, not genuine popular will. (Gramsci 1971; Lukes 1974) Noble lies are not adopted here as a deliberate strategic tool — not because the role of narrative in politics is naive, but because deliberate noble lies by actors with power asymmetry are structurally indistinguishable from the noble lies that serve elite interest. The custodian problem applies in its sharpest form: who decides which lies are noble? In practice, the answer is always those with the power to decide — which means the lies that get called noble are systematically the ones that benefit those making the designation. There is a harder case that must be acknowledged. In conditions of extreme adversarial pressure, actors genuinely working against harmful

arrangements will sometimes use narrative strategically in ways that are not fully transparent. The framework does not pretend this does not happen or that it is always wrong. What it insists on is that this be recognized as a tactical necessity under specific conditions rather than as a general principle — and that those conditions be specified as precisely as possible, rather than left open as a blank check for whoever claims to be on the right side. The spectrum from functional fiction to structural suppression: Not all simplifications of complex reality are noble lies in the politically significant sense. Every institution operates with working fictions — simplified accounts of how things work that are close enough to reality to enable coordination without requiring participants to maintain a complete model of the system at all times. The organizational chart is a working fiction: it describes formal authority without capturing the infor

mal networks through which decisions are actually made. The legal fiction that corporations are persons is a working fiction that enables specific legal operations. These fictions are not primarily in the service of power — they are coordination devices that reduce the cognitive load of institutional participation. The politically significant noble lie is structurally different: it is a fiction whose function is to make existing power arrangements appear natural, necessary, or legitimate, and whose maintenance requires active suppression of the evidence and analysis that would reveal it as a fiction. The meritocracy fiction as canonical case: The claim that existing distributions of wealth, status, and institutional position reflect individual merit — that the current arrangement rewards those who deserve reward and withholds from those who have not earned — is the canonical noble lie of market societies. It is not simply false: individual effort, skill, and judgment do produce differential outcomes within any institutional arrangement, and those differences are real. What makes it a noble lie is the structural suppression required to maintain it at the level of general social belief. The evidence that inheritance, educational access, social networks, and the structural advantages of birth position systematically outweigh individual merit in determining life outcomes is extensive, well-documented, and consistently underweighted in the political culture of societies organized around meritocratic legitimization. Maintaining the meritocracy fiction requires not just selective emphasis but active institutional work: educational systems that attribute differential outcomes to individual capacity rather than differential resource provision, welfare systems designed to signal that support is conditional on demonstrated effort, political discourse that consistently frames systemic disadvantage as individual failure, echoing the ideological legitimations of neoliberal governance (Harvey 2005). The consent function and its limits: Noble lies perform a consent function: they provide the population whose compliance the arrangement requires with reasons to accept it that are grounded in their own values rather than in explicit coercion. An arrangement justified by the meritocracy fiction recruits the compliance of the non-advantaged on the basis of their own belief that the arrangement is fair — that their position reflects their own capacity rather than structural condition. This is more stable than coercion-based compliance because it does not require continuous deployment of force, and it is more durable because it is self-reinforcing: people who believe the arrangement is fair will attribute their disadvantage to their own failures, which reduces the probability of collective action and increases the probability of individual adaptive behavior within the existing arrangement. The consent function is why those who benefit most from an arrangement invest most

in maintaining the fictions that legitimate it: the investment produces returns in the form of reduced need for coercive enforcement. The diagnostic question: The framework's approach to any specific governing narrative — any account of why existing arrangements are natural, necessary, or legitimate — is not to assume it is a noble lie but to ask what evidence would distinguish a fiction that serves power from an accurate account that happens to legitimate existing arrangements. The question is: what would need to be true for this narrative to be accurate, and is the evidence available to assess whether those conditions hold? A narrative that is insulated from evidential challenge — that adjusts its auxiliary claims to accommodate each piece of contrary evidence without revising its core conclusion — is exhibiting the pattern of a noble lie regardless of its specific content. A narrative that makes specific predictions about observable conditions and revises when those predictions fail is a genuine account, even if its political effect is legitimating. The diagnostic procedure from 1.1 and 1.2 applies here directly. The framework's own narratives are subject to

this diagnostic. PMN makes claims about how power operates, how institutions develop, and what conditions produce structural suffering. Those claims are not exempt from the evidential requirement that distinguishes genuine accounts from fictions that serve power — including the power of those who hold the analytical framework. The self-application requirement from 1.2 and 6.4 applies to the framework's own governing narratives as much as to the governing narratives it analyzes in others. A framework that exempts itself from its own diagnostic procedure is exhibiting the pattern it was built to identify.

7.6 — Those Who Stand on Others: The Structural Analysis

There are actors whose material position depends on arrangements that produce structural suffering for others. The observation is structural, not moral — about incentives produced by position, not about the character of those who occupy it. People whose wealth, status, and security are produced by arrangements that constrain others' development have rational incentives to maintain those arrangements — not because they necessarily intend harm, but because the alternative involves a genuine material cost that their position has not prepared them to absorb. This is the claim that naive analyses of power consistently miss. The problem is not primarily individual malice. The structural force that produces and reproduces arrangements of systematic advantage is the alignment of material interests with the maintenance of specific institutional configurations — an alignment that operates through the normal functioning of rational self-interest, without requiring conspiracy, without unusual cruelty, without requiring anything more than actors behaving predictably from their material position. Moral appeals to the better nature of those who benefit from harmful arrangements have a poor historical record — not because the people involved lack better natures, but because the material incentives that sustain their behavior are not addressed by moral appeals. Reform that does not change the incentive structures that produce the behavior being reformed gets re-absorbed into the arrangements it was supposed to change. The actors whose interests are threatened will use exactly the mechanisms this framework describes — institutional capture, information control, path dependence, noble lies — to redirect reform energy in directions that leave their core position intact. At the same time, the framework resists reducing all actors to their structural position. There are actors in positions of structural advantage who genuinely work against the arrangements that benefit them — not because they are irrational but because their analysis and values lead them to conclusions that override their immediate material interest. The conditions under which actors are more or less likely to act against their immediate material interest can be analyzed, and those conditions can be cultivated. The most important implication is that the primary target of political action is not the conversion of those who benefit from harmful arrangements. It is the transformation of the arrangements themselves — the structural conditions that generate the incentives. When those conditions change, the behavior of rational self-interested actors changes with them. The structural analysis of positional interest is not a moral judgment about individuals. It is an account of how structural positions generate incentives that shape behavior in predictable ways regardless of individual character. This distinction matters both analytically and strategically. Analytically: if the primary driver of behavior tha

t maintains structural suffering is individual moral failure, the appropriate response is moral suasion and the reform of individuals. If the primary driver is structural position generating rational incentives, the appropriate response is structural redesign that changes the incentives — and moral suasion addressed to individuals in the structural position will consistently underperform because it is fighting against the incentive structure rather than changing it. The distinction determines what kind of intervention is adequate to the problem. The mechanics of positional interest: Positional interest operates through several mechanisms simultaneously. The most direct is material: actors whose income, wealth, and economic security depend on an arrangement have financial interests in its continuation that produce rational incentives to maintain it. Less direct but equally consequential is the identity mechanism: actors whose social status, professional identity, and sense of their own worth are produced by an arrangement have psychological interests in its

continuation that operate with motivational force comparable to material interest. The network mechanism: actors whose access to information, opportunities, and political influence is mediated through the networks produced by an arrangement have relational interests in its continuation — the arrangement is not just what they benefit from materially but the social world through which they navigate. All three mechanisms are present simultaneously for most actors in structurally advantaged positions, which is why positional defense is so robust even when individual actors are genuinely committed to values that their position undermines. The reform problem: Actors whose positional interest lies in maintaining arrangements that produce structural suffering for others will consistently support modifications that reduce the arrangement's political vulnerability without changing its structural property of producing that suffering. This is not hypocrisy — it is the predictable behavioral output of the structural position. The plantation owner who supports improved conditions for enslaved people to reduce the probability of revolt is acting from positional interest in the arrangement's continuation, not from commitment to the value of human equality that would require the arrangement's abolition. The executive who supports diversity initiatives that increase institutional legitimacy without changing the wage structures that maintain inequality is acting from positional interest in institutional reputation, not from commitment to equality that would require restructuring the distribution of returns. In both cases the modification is genuine and may produce real improvements for those affected; in neither case does it address the structural property of the arrangement that produces the harm. The coalition problem and counter-power: Actors with positional interest in maintaining arrangements that produce st

ructural suffering do not typically act alone. Their positional interest produces natural coalitions: each actor whose position depends on the arrangement has independent reasons to support it, and the aggregate of those independent supports constitutes a structural coalition that no single actor designed or coordinates. Counter-power accumulation — the process analyzed in 10.9 — must account for this coalition structure. The counter-power required to change a structural arrangement must overcome not only the resources of the most visibly advantaged actors but the distributed resistance of all actors whose positional interest, at whatever level, is served by the arrangement's continuation. This is why reform efforts that identify a single powerful opponent consistently underestimate the resistance they will encounter: the structural coalition defending the arrangement is broader than its most visible members, and it reconstitutes after any single member is removed or converted. The strategic implication is not that individual conversion is worthless — actors in structurally advantaged positions who act against their positional interest are genuinely valuable, and the analysis in 17.7b addresses their specific situation. The implication is that individual conversion is insufficient as a primary strategy for structural change because the structural position will continue to generate the same incentives for whoever fills it next. The reformer who changes the person without changing the position has not addressed the structural mechanism that produced the behavior they were objecting to. Lasting change requires changing the structural position — its resource access, its authority, its relationship to those affected by its operation — not only the individual who occupies it. The connection to institutional design: arrangements that concentrate benefits in narrow structural positions while distributing costs broadly are structurally unstable in the long run — they generate the counter-power pressure analyzed in 10.9 — but structurally durable in the short and medium term because of the coalition defense mechanism. The time-frame mismatch between structural instability and structural durability is one of the primary mechanisms through which arrangements that fail the minimal anchor criterion persist: the counter-power pressure that would eventually change them builds slowly, while the coalition defense responds immediately and with resources that the counter-power has not yet accumulated.

7.7 — Path Dependence: Why History Constrains the Present

Institutions persist not only because powerful actors defend them but because they create genuine dependencies that make alternatives costly even for actors who would prefer them. This is path dependence: early choices foreclose later options not through deliberate lock-in but through the accumulation of

complementary adaptations that make the existing arrangement progressively more embedded in the structure of everyday life. A society that developed its energy infrastructure around coal does not simply switch to solar when solar becomes cheaper, because the coal infrastructure is embedded in a network of complementary investments: the training of engineers, the location of industry, the design of transmission systems, the political coalitions of communities whose livelihoods depend on coal's continuation. The same logic applies to institutional arrangements. Legal systems, educational structures, financial architectures, and political organizations all generate complementary adaptations that make the cost of switching to alternatives higher than the raw comparison of the arrangements would suggest. Path dependence is one of the most important reasons why the analysis of power alone is insufficient for understanding institutional persistence. Even where there is no powerful actor actively defending an arrangement, the arrangement can persist because the accumulated adaptations to it have made the alternatives genuinely more costly — not because of ideology but because of the real material investments, skills, relationships, and infrastructure built around it. Reformers who treat this as mere resistance to change underestimate what genuine transition requires. Effective institutional reform requires attention to the sequencing and pacing of change as much as to its direction. Attempting to move too quickly disrupts the complementary adaptations before alternatives are ready to replace them, producing coordination failures attributed to the reform itself. Attempting to move too slowly allows the complementary adaptations to deepen, making the eventual transition more costly. The right pace requires empirical knowledge of the specific dependencies involved — not general principles about the appropriate speed of reform. (North 1990; Pierson 2004) The analysis of institutions connects directly to the analysis of society and culture that follows. Institutions are the formal crystallization of power relations — but the conditions that sustain or erode those relations are not only institutional. They are also cultural: the meanings people attach to arrangements, the identities that those arrangements produce and reproduce, and the practical norms that govern behavior below the level of formal rule. Part VIII examines those cultural and social dynamics using the same material standard: not what meanings or identities feel authentic, but what structural consequences they produce for the populations who live within them. The mechanism in detail: Path depend

ence operates through three distinct mechanisms that reinforce each other. The sunk cost mechanism: significant investments in infrastructure, training, organizational capacity, and social expectation have been made in support of the existing arrangement, and those investments create genuine costs to switching that are not present in counterfactual arrangements that have not yet been built. The complementarity mechanism: existing arrangements produce complementary adaptations — regulatory frameworks designed for existing technologies, professional training programs organized around existing practices, financial systems calibrated to existing risk profiles — that collectively make the existing arrangement more functional than any alternative would be until those complementary adaptations are rebuilt around the alternative. The expectation mechanism: actors make long-term decisions based on expectations that existing arrangements will continue, and the aggregate of those decisions embeds the arrangement more deeply in everyday practice, making the cost of deviation higher for each subsequent actor. Path dependence is not determinism: The claim that history constrains the present is not a claim that history determines the present. Path dependence makes some transitions costly; it does not make them impossible. The mechanisms that create path dependence — sunk costs, complementary adaptations, embedded expectations — can be disrupted by sufficiently large shocks that simultaneously devalue existing investments, destroy complementary adaptations, and reset expectations. Wars, financial crises, technological disruptions, and epidemics have all served this function historically, opening what 10.5 calls windows of transformation precisely because they simultaneously dismantled the mechanisms that made existing arrangements sticky. The strategic implication is not that reform requires waiting for a crisis but that reform strategies must account for which specific mechanisms of path dependence they are working with and what resources and approaches are required to address each one. Manufactured path dependence: Not all path dependence is the unintended byproduct of investment and adaptation. Actors who benefit from existing arrangements have incentives to deepen path

dependence deliberately — to increase the costs of transition, multiply the complementary adaptations that would need to be rebuilt, and shape expectations in ways that make the existing arrangement seem more inevitable. The most common mechanism is regulatory capture producing regulatory frameworks that are specifically designed for the existing form of an arrangement rather than for the function the arrangement is supposed to serve — making transition to a different form technically illegal or practically prohibitive even when the different form would better serve the ostensible function. Proprietary standards, intellectual property regimes, and professional licensing requirements have all b

een used in this way: not to protect the function but to protect the specific form that benefits the actors who shaped the regulatory framework. The analytic discipline path dependence requires is explicit attention to which features of any current arrangement are genuinely hard to change because of the mechanisms described above, and which features appear hard to change primarily because they have not been challenged with sufficient resources and counter-power. These are not the same question. Genuine path dependence requires strategies that either work within the existing path — incremental change that does not require dismantling complementary adaptations — or that create the conditions for path-breaking change through coalition-building, resource accumulation, and waiting for or helping to produce the disruption that opens a transformation window. Apparent path dependence — the sense that change is impossible because it has not been tried with adequate resources — requires different strategic responses: counter-power accumulation, narrative reframing, and the demonstration that alternatives are viable, which reduces the expectation mechanism without requiring the sunk costs and complementary adaptations to be dismantled first. Connection to the custodian problem: Path dependence and the custodian problem interact in a specific way: institutions whose original function has been displaced by self-maintenance use path dependence arguments to resist transformation. The claim that change is too costly, that the complementary adaptations would need to be rebuilt, that the expectations of too many actors are organized around the existing arrangement — these are structurally accurate observations that are also deployed as strategic defenses by actors whose positional interest lies in the arrangement's continuation. Distinguishing between genuine path dependence arguments and strategic deployment of path dependence rhetoric requires asking: who is making the argument, what is their structural position relative to the existing arrangement, and what is the quality of evidence they are deploying about the actual costs of transition versus the speculative risks they are foregrounding? Diagnostic for path dependence claims: (1) Which specific mechanism is operating — sunk costs, complementary adaptations, or expectations? The answer determines what kind of transition strategy is required. (2) Is the path dependence genuine — produced by the mechanisms above — or apparent — produced by the absence of sufficient challenge? (3) Who is making the path dependence argument, and what is their positional interest in the arrangement's continuation? (4) What would it cost to dismantle the specific mechanism of path dependence identified — and who would bear that cost, compared to the ongoing cost of the arrangement's continuation for those it harms?

7.8 — Regime Types: A Comparative Diagnostic

The framework's evaluative criterion — does this arrangement create the conditions under which genuine development is possible and reduce the structural suffering that forecloses it — applies to all political arrangements without prior commitment to any particular institutional form. This section applies that criterion comparatively to the major regime types that have actually existed or currently exist. The analysis is diagnostic, not polemical. Each type is evaluated for what it gets right as well as what it gets wrong, and under what material conditions each has produced better or worse outcomes by the framework's standard. The prior commitment is to the criterion, not to any regime that happens to satisfy it under some conditions. A clarification on the status of what follows. The evaluative criterion applied here is part of the Framework — durable, methodology-level, and not specific to any historical moment. The applications to specific regime types are part of the Doctrine: contextual, temporal, and revisable as conditions change and evidence

accumulates. The claim that liberal democracy's accountability mechanism is structurally significant is a Framework claim. The claim that the Chinese developmental model produced the largest poverty reduction in human history is a Doctrinal observation about a specific historical period. Both are held to the same evidentiary standard, but they do not have the same durability. Readers who disagree with the doctrinal applications are invited to apply the framework's criterion to the same cases and show where the analysis is wrong — that is exactly what the framework's epistemological commitments require of it. Liberal Democracy What liberal democracy gets right: the accountability mechanism. Competitive elections, freedom of political organization, and the peaceful transfer of power create an enforcement structure for the custodian problem that other regime types typically lack. When populations experience arrangements as producing structural suffering, they have a formal channel through which that perception can produce institutional change. This is not trivial. The historical record of what happens to populations under arrangements without this channel is not ambiguous. Liberal democracy also tends to protect the conditions for cognitive and cultural development — freedom of expression, institutional plurality, tolerance for dissent — in ways that matter for the anti-foreclosure criterion as well as the floor. What liberal democracy gets wrong: the accountability mechanism requires distributed economic power to function as intended. Under conditions of high economic inequality, wealth translates into political influence through campaign finance, media ownership, regulatory capture, and the revolving door between public office and private interest. The formal accountability mechanism — one person, one vote — coexists with the material reality that political outcomes track the preferences of concentrate

d economic interests more reliably than those of median voters. Procedural legitimacy protects arrangements that produce structural suffering, because the populations most harmed by those arrangements are also the least able to translate their preferences into political outcomes. Liberal democracy performs significantly better under conditions of relatively distributed economic power; it performs significantly worse as inequality increases. Its internal logic does not reliably prevent the latter. Confucian Political Philosophy: Mandate of Heaven and Role Ethics The Confucian political tradition, developed across twenty-five centuries from Confucius and Mencius through the Neo-Confucian synthesis and into contemporary engagement, provides analytical resources for two of PMN's central concerns: the theory of legitimacy depletion and the institutional design of governance. On legitimacy: Mencius's doctrine of the Mandate of Heaven (tianming) is a non-Western theory of legitimacy that independently reaches conclusions structurally parallel to PMN's analysis in section 6.6. Mencius argues that political authority is legitimate when it serves the material welfare of the people — specifically, when it ensures that the population has adequate food, stable livelihoods, and conditions for family and community life. When a ruler fails these conditions systematically, the Mandate is lost regardless of dynastic legitimacy or formal authority, and the population's compliance is no longer obligated. Mencius goes further than most Western legitimacy theory: he explicitly legitimizes the removal of rulers who have lost the Mandate through systemic failure of the population's material conditions. 'The people are the most important element in a nation; the spirits of the land and grain are the next; the sovereign is the lightest.' This is a floor-criterion theory of legitimacy: authority rests on material provision, and its absence grounds the withdrawal of compliance. Ibn Khaldun's asabiyyah theory is a materialist elaboration of the same structural insight applied to the sociology of dynastic cycles. On institutional design: Confucian role ethics, as analyzed by contemporary philosophers including Roger Ames and Henry Rosemont, offers a critique of rights-based institutional design from a non-Western analytical position. Where liberal institutional design begins from the individual as bearer of rights prior to social relationships, Confucian role ethics begins from the person as constituted by role relationships — the parent-child, ruler-subject, husband-wife, elder-younger, friend-friend relations that define who one is and what one owes. The institutional design implication is that obligations flow from roles rather than from abstract rights, and that institutional legitimacy requires that role-holders perform the obligations their roles generate rather than that procedures guarantee individual rights. This produces a different architecture of accountability: a

accountability runs through role-performance rather than through rights-enforcement, and failure is role-failure rather than rights-violation. Confucian political philosophy's contribution to PMN: the Mandate of Heaven doctrine is an independent non-Western derivation of floor-criterion legitimacy theory that predates and does not require Western liberal foundations. PMN's analysis of legitimacy depletion (6.6) and the accountability mechanism (7.2) can be read as a materialist elaboration of the structural insight Mencius identifies. The role ethics contribution is more challenging for PMN: if obligations flow from role relationships rather than from abstract rights, the institutional design questions change in ways that PMN's current framework does not fully address. The Confucian tradition's historical implementation through the keju (examination) system produced the most sustained meritocratic governance infrastructure in human history, with specific structural features — competitive selection, rotation, accountability through performance rather than heredity — that are relevant to PMN's custodian problem analysis (7.3).

Technocracy What technocracy gets right: complex systems genuinely require analytical capacity to manage, and intuitions shaped by immediate experience are poor guides to the second and third-order consequences of institutional design at scale. The technocratic critique of populism — that preferences formed under conditions of limited information and cognitive bias produce poor institutional choices — is not wrong as a description of how preferences are actually formed. Technocratic governance has produced measurable improvements in specific domains: central bank independence stabilizing monetary policy, independent regulatory agencies reducing industrial hazard, expert-managed public health infrastructure reducing preventable disease. The empirical case for analytical capacity in governance is real. What technocracy gets wrong: superior analytical capacity does not determine whose interests the analysis serves. Expert institutions are subject to capture by the interests best positioned to influence the selection and socialization of experts — which in practice means the interests of those with concentrated economic and institutional power. The European Central Bank's management of the 2010 debt crisis, the IMF's structural adjustment programs, the financial sector's capture of regulatory agencies: in each case, technically sophisticated actors produced analytically coherent outcomes that systematically served concentrated interests at the cost of the populations the analysis was nominally for. The problem is not analytical inadequacy. It is the absence of genuine accountability to those whose material conditions the analysis affects. Technocracy without accountability is expert capture formalized as governance.

Theocracy What theocracy gets right: it provides a comprehensive meaning structure that addresses existential suffer

ing directly. Where material conditions are precarious and state capacity is limited, religious institutional infrastructure has historically provided the only functioning welfare, education, and dispute-resolution systems available to large populations. Theocratic governance can produce genuine social cohesion, reduction of certain forms of interpersonal violence, and the institutional stability that chronic material insecurity otherwise prevents. The dismissal of theocratic arrangements as simple irrationality ignores the material functions they actually perform in specific conditions. What theocracy gets wrong: political authority grounded in unfalsifiable claims cannot be corrected by evidence. This is not a marginal flaw. The accountability mechanism that the framework identifies as essential — the capacity to translate material failure into institutional change — requires that the grounding claim of political authority be challengeable. When authority derives from divine mandate, the claim is structurally immune to the evidence that would challenge it. Theocracy also tends to foreclose developmental trajectories that conflict with religious authority — in practice, the cognitive, expressive, and institutional forms that inquiry requires. Visibility suppression of dissent is built into the architecture, not a contingent feature of specific theocratic regimes. The framework's evaluation does not deny what theocratic arrangements produce in specific conditions. It identifies why those arrangements cannot be reliably corrected when they fail.

Authoritarianism Authoritarianism is a heterogeneous category, and its heterogeneity is diagnostically important: it spans arrangements that have produced dramatically different outcomes by the framework's criterion. What authoritarian arrangements can get right is the coordination of large-scale developmental investment under conditions of severe institutional fragmentation, without the capture of policy by short-term electoral cycles

or the veto power of entrenched interests that reform-resistant democratic institutions can produce. Singapore's developmental state, South Korea and Taiwan's authoritarian industrialization phases, Rwanda's post-genocide reconstruction: these represent genuine achievements at the floor level, and the achievement was not despite the concentration of authority but through it, under specific material conditions where fragmentation was the primary obstacle. What authoritarianism structurally gets wrong: the accountability gap. The developmental state literature (Amsden, Evans, Chang) correctly identifies that successful developmental authoritarian states — South Korea under Park, Taiwan under the KMT, Singapore — were not accountable in the electoral sense but had functional internal accountability mechanisms: meritocratic bureaucracies insulated from patronage, performance-based promotion systems, and developmental goals enforced through intra-elite competition. These are real

accountability mechanisms, structurally different from the pure extraction model. The success cases therefore work not despite zero accountability but despite the absence of electoral accountability specifically — a distinction that matters for institutional design. What they lack is external accountability to the populations whose becoming they claim to serve, which is where PMN's anti-foreclosure critique applies — they require an unusual alignment of leadership interest with developmental interest that is not produced by the institutional structure and cannot be reliably reproduced. The modal outcome of authoritarian arrangements is not Singapore. It is the extraction of value from populations by those who control the instruments of coercion, with developmental investment subordinated to elite maintenance. The variance in authoritarian outcomes is extremely high precisely because outcomes depend on the character and interests of leadership rather than on enforcement mechanisms that would produce developmental behavior regardless of leadership. High variance and no floor is the structural signature of the accountability gap. Totalitarianism Honest evaluation requires acknowledging what totalitarian regimes have produced before evaluating why they fail. Soviet industrialization achieved rapid expansion of industrial capacity, dramatic expansion of literacy — from roughly 20 percent literate to over 85 percent within two decades, approaching near-universal by mid-century — and mass public health infrastructure. Maoist mobilization achieved specific material targets under conditions of extreme scarcity. These are real outcomes by the floor criterion. Their dismissal by liberal historiography as either irrelevant or entirely offset by atrocity is itself a form of selective accounting that the framework's consistency standard disallows. What totalitarianism structurally gets wrong — and what makes the floor achievements insufficient by the framework's criterion — is the anti-foreclosure criterion. Total mobilization for a single developmental trajectory requires the elimination of all developmental trajectories not sanctioned by the state. The cognitive surplus, cooperative surplus, and institutional trust that becoming requires are the first casualties, because they are also the primary inputs for organized dissent. Totalitarian arrangements may raise the material floor for specific populations over specific periods. They do so by foreclosing becoming: by capturing transformation pressure and redirecting it toward state-defined objectives, eliminating the conditions under which populations could develop in directions the state has not pre-authorized. The existential suffering produced by this foreclosure — the deprivation of meaning, agency, and developmental trajectory — is not detected by $S = R + B + V$ because it does not operate through resource deprivation. It operates through the structural elimination of the becoming capacity that the floor

was supposed to enable. Juche and Closed Theocratic Nationalism Juche — North Korea's state ideology of self-reliance — represents a specific case where the analytical framework yields its most consistent negative evaluation, and that evaluation must be stated clearly rather than hedged for the sake of apparent balance. The architecture of Juche maximizes institutional betrayal (B) and visibility suppression (V) by design, while producing chronic resource deprivation (R) as a structural consequence of its isolation logic. The state maintains population control not by meeting material needs but by monopolizing the provision and definition of meaning: the regime functions as a closed theocratic nationalism in which the ruling family occupies the position that religious authority occupies in theocracy, with the same unfalsifiability built into the grounding claim. The one thing Juche reliably produces is a meaning structure — a comprehensive account of collective

identity and national purpose that addresses existential suffering for some populations under conditions of material insecurity. This is not fabricated from nothing; it draws on real historical grievances, genuine experiences of colonial extraction, and rational responses to external military threat. But it addresses existential suffering not by creating developmental conditions but by eliminating all alternative meanings against which the state's meaning structure could be measured. That is not resilience. It is the forcible elimination of the cognitive and social conditions under which adaptive response capacity is preserved. By both floor and anti-foreclosure criterion criteria, Juche fails comprehensively. The honest evaluation of why requires understanding what it actually provides to populations under the specific conditions it operates in — not because that understanding rehabilitates the arrangement, but because without it the analysis of how to change it is empirically inadequate. Fascism and National Socialism What fascism gets right — in the limited diagnostic sense that PMN applies to all traditions: it identified a genuine problem that liberal and Marxist analysis both systematically underweighted. The problem is the political force of wounded national identity, economic humiliation, and meaning-collapse in populations stripped of prior frameworks of orientation. Interwar fascism rose not in conditions of ordinary poverty but in conditions of prior middle-class stability followed by catastrophic status collapse: the Versailles settlement's punitive framework, hyperinflation that eliminated the savings of the German middle class, mass unemployment in populations with high literacy and high aspirations, and the simultaneous collapse of the institutional arrangements that had organized meaning and identity. PMN's own analysis of affective mobilization (3.8b) and the destabilizing consequences of meaning-collapse (4.3, 5.5) explains why fascist mobilization worked: it addressed, ho

wever perversely, the genuine psychological and social conditions that the mainstream alternatives were not addressing. The diagnostic value of acknowledging this is not rehabilitation. It is the insistence that understanding why fascism found mass constituencies is a prerequisite for preventing it, and frameworks that explain fascism only as irrationality or moral failure are frameworks that cannot predict when those conditions will recur. What fascism gets wrong — comprehensively and structurally: it resolves the genuine problem of meaning-collapse and humiliated identity by directing the resulting energy toward the annihilation of an outgroup identified as the source of all structural suffering. This is not a marginal design flaw. It is the operational core of fascist politics, and it means that the solution to the genuine problem fascism addresses produces structural suffering at a scale that dwarfs what it claims to resolve. The biological floor criterion is maximally violated: Nazism produced industrial genocide, deliberately engineered mass starvation in occupied territories, and systematic enslavement of populations on a scale without historical precedent. The becoming criterion is equally violated: the mobilization of the population toward the collective project requires the suppression of all individual and subgroup development that is not instrumentally useful to that project. Fascism is also structurally self-defeating by its own stated objectives: the 'national community' whose revitalization it promises is achievable only through the destruction of the internal diversity and genuine solidarity that revitalized communities actually require. Fascism's position within PMN: the tradition correctly diagnoses the political consequence of meaning-collapse in humiliated populations and the insufficiency of purely material accounts of political mobilization. Its resolution — racial or national regeneration through the elimination of the designated outgroup — fails the evaluative criterion at every level simultaneously: biological floor, becoming, and the sustainability of the cooperative substrate. PMN's structural question is not 'what did fascism get right?' in the sense of policy contribution, but 'what conditions produce fascist mobilization, and what alternative addresses those conditions without the eliminationist resolution?' That second question is the one that matters for prevention, and it is not answered by moral condemnation alone. Ethnonationalism Ethnonationalism — the political doctrine that ethnic, linguistic, or racial communities constitute the only legitimate basis for political community, and that political boundaries should align with ethnic boundaries — is the dominant rising force in early twenty-first century politics across contexts as different as Western Europe, South Asia, the Middle East, and sub-Saharan Africa. A framework that does not engage it analytically is a framework prepared for a political world that no

longer exists. What ethnonationalism gets right: it addresses a genuine human need that cosmopolitan and civic-nationalist alternatives have consistently underdelivered on. The need is for a political community that feels materially real — that is not abstract or procedural but constituted by shared history, shared language, shared practice, and the particular form of solidarity that comes from knowing the people you are coordinating with as people like you. PMN's own analysis of meaning infrastructure (5.5, 5.6) acknowledges that communities grounded in particular identity — as opposed to universal principles — are more effective at generating the sustained commitment and cross-generational solidarity that transformation requires. Ethnonationalism exploits this genuine insight by defining the community in exclusionary terms that make it emotionally vivid and politically mobilizable. The emotional force of ethnonationalist politics is not manufactured from nothing; it activates solidarity mechanisms that are real, evolutionarily grounded, and that cosmopolitan frameworks have systematically failed to replicate at comparable emotional intensity. What ethnonationalism gets wrong structurally: the solution — ethnic homogeneity as the basis for political community — does not produce the solidarity it promises. Ethnically homogeneous states exhibit the full range of internal class conflict, institutional betrayal, and structural suffering that ethnonationalism claims ethnic diversity causes. The premise that ethnic or racial commonality produces genuine solidarity sufficient for just institutional arrangements is empirically false: the historical record of mono-ethnic states shows no systematic reduction in institutional capture, elite extraction, or structural suffering relative to multi-ethnic states at comparable levels of institutional development. What ethnonationalism does reliably produce is the direction of structural suffering outward toward minorities and outgroups, which suppresses internal class conflict by substituting ethnic grievance for structural analysis. This is the manufactured fracture line mechanism (10.9) operating through ethnic rather than class lines: populations experiencing genuine structural suffering are mobilized against ethnic others rather than against the institutional arrangements producing the suffering. Ethnonationalism's position within PMN: the tradition correctly identifies the emotional and political insufficiency of purely procedural or universal political community and the genuine human need for particular solidarity. Its resolution — ethnic exclusion as the basis for community — fails both the floor criterion (systematic suffering of minorities and outgroups) and the becoming criterion (foreclosure of development for populations outside the ethnic definition). It also fails on its own stated terms: ethnic homogeneity does not produce the genuine solidarity and institutional adequacy it promises. PMN's diagn

ostic question is what alternative form of political community can generate the emotional intensity and particular solidarity that ethnonationalism accesses, without the exclusionary resolution that produces structural suffering at the outgroup's expense. Chinese People's Democracy and Developmental Meritocracy The Chinese model since 1978 presents the most analytically demanding case for this framework, and it demands genuine engagement rather than dismissal in either direction. The floor achievement is real: the reduction of extreme poverty from over 80 percent of the population to under 5 percent over four decades represents the largest and fastest improvement in material conditions for the largest population in human history. This is not a secondary achievement. By the framework's floor criterion, it is the central purpose of the framework's entire analytical apparatus. Any evaluation that does not begin by acknowledging this achievement has already failed the consistency standard. The Chinese theoretical critique of liberal democracy also deserves engagement rather than dismissal. The claim that competitive electoral systems under conditions of high economic inequality reliably produce capture by concentrated economic interests is a description of actually observable phenomena, not a propaganda talking point. The claim that "whole-process democracy" — consultation across the full governance cycle rather than periodic electoral choice — better represents population preferences under conditions where electoral choice is effectively constrained by economic power is a serious institutional argument. And the meritocratic selection logic — that governance positions should be filled by those with demonstrated analytical and administrative capacity rather than by those most successful at electoral competition — addresses a real failure mode of democratic governance that the technocracy section above identifies. What the Chinese model structurally gets wrong,

by the framework's criterion, is not its material achievements but its anti-foreclosure architecture. The meritocracy claim is partially a noble lie: selection for governance positions tracks loyalty to the Party as a prior condition, which means the meritocracy selects for a specific alignment of analytical capacity and political orientation, not for analytical capacity alone. More fundamentally, the absence of genuine accountability mechanisms — mechanisms that could produce institutional change when outcomes fail the population rather than the leadership — means there is no enforcement structure for the custodian problem. The floor achievements of 1978–2015 were produced under conditions where the material interests of the developmental state and the material interests of the population were largely aligned. The question the framework asks is not whether this alignment has existed but what happens when it breaks down, and whether the institutional structure contains any mechanism for correcti

ng it. The structural answer is that it does not — which means the anti-foreclosure performance of the model is dependent on the contingent alignment of leadership interest with developmental interest, precisely the high-variance condition that authoritarianism in general produces. The honest tension cannot be resolved by preferring one criterion. The Chinese model has performed better than liberal democracy on the floor criterion for a larger population over the specific historical period in which their comparison is most relevant. That is a genuine finding. It has structural features that make anti-foreclosure performance unreliable and make correction of floor failure institutionally difficult as material conditions change. That is also a genuine finding. A framework that dismissed one finding to affirm the other would be doing what this framework was constructed to prevent: treating the conclusion it prefers as the criterion rather than applying the criterion honestly and accepting the conclusions it produces. No regime type passes the framework's criterion in all conditions. The question is always material-contextual: which arrangement, under these specific conditions, with this specific history and distribution of power, creates better conditions for genuine development and reduces the structural suffering that forecloses it? Liberal democracy performs better when economic power is distributed; it performs worse as inequality concentrates. Technocracy produces good outcomes when accountability mechanisms are present; it produces capture when they are absent. Theocracy provides meaning and stability under conditions of material insecurity; it forecloses development when the grounding claim is used to immunize arrangements from correction. Authoritarianism can achieve rapid development under specific conditions; it produces extraction under all others. Totalitarianism can raise the floor; it forecloses becoming by design. The Chinese developmental model has achieved the floor criterion at scale; its anti-foreclosure architecture is structurally unreliable. The framework does not declare a winner. It provides the diagnostic that makes the question answerable without the prior commitment to one answer that the question requires honesty not to have.

VIII: Society, Culture, and the Politics of Material Outcomes

8.1 — Evaluating Cultural Practices

Cultural practices are evaluated by their material outcomes. This is the application to culture of the same standard applied everywhere else: does this arrangement produce more or less unnecessary suffering than available alternatives? The judgment is comparative, contextual, falsifiable — and demanding, because it requires specifying the alternatives honestly and evaluating the consequences without exempting practices from scrutiny because of their cultural origin. The materialist tradition tended toward a different form of cultural evaluation — treating cultural practices as expressions of underlying material relations that would transform as those relations transformed. This produced a specific kind of insensitivity to the independent dynamics of cultural and meaning systems — the assumption that changing the economic base would automatically change the cultural superstructure, and that therefore cultural practices did not require independent analysis. The assumption was wrong, and its consequences were predictable: movements and governments that changed the economic base found that cultural practices were far more persistent and far more consequential than the theory predicted. Applying this evaluative standard symmetrically is the source of two-directional discomfort. Applied to practices of non-Western or minority cultures, it can look like cultural imperialism — the imposition of an external standard on communities that did not choose it. Applied to practices of dominant cultures, it looks like radical critique. Both applications are correct, and the discomfort they produce in different audiences is evidence that the standard is being applied consistently rather than selectively. A standard that makes only one side uncomfortable is not a standard. It is a preference dressed as one. The most common failure mode in cultural evaluation is the double standard: rigorous scrutiny of practices that appear different from the evaluator's own cultural baseline, and uncritical acceptance of practices that appear familiar. Financial instruments that extract value from economically vulnerable populations are rarely called cultural practices, even when they are deeply embedded in the culture of finance. Overwork as a measure of professional worth is rarely evaluated by whether it produces unnecessary suffering, even when the evidence that it does is substantial. The framework requires that these receive the same evaluative attention as practices more typically flagged for cultural critique. The uncomfortable implication: dominant cultural practices, including those that present themselves as universal civilization rather than particular culture, are not exempt. The exemption reflex and why it fails: The default position in much progressive political culture is that cultural practices belonging to marginalized or historically oppressed communities deserve special protection from criticism — that applying the minimal anchor criterion to such practices is

a form of cultural imperialism that reproduces the same evaluative hierarchy that colonial projects used to justify dismantling indigenous institutions. This is a real concern that the framework takes seriously: the history of using universal standards to override local practices is a history with a consistently bad record, and the analysis of anti-naïve universalism in 8.5 engages that record directly. But the conclusion the exemption reflex draws from that history — that practices with cultural provenance in marginalized communities should be exempt from the same evaluative standard applied to dominant-culture practices — reproduces the analytical failure it means to resist. It treats the communities in question as objects of protective deference rather than as agents whose practices are subject to the same analysis applied to all practices. That is not respect. It is a different form of condescension. What the evaluation actually requires: Evaluating a cultural practice by its material outcomes requires four specific moves. First, identifying who within the community the practice benefits and who it harms — cultural practices are not experienced uniformly by all members of the community that maintains them, and the relevant unit of analysis for the minimal anchor is the individual organism experiencing suffering, not the cultural community as a collective. Second, distinguishing between the practice itself and the social functions it serves — many practices that produce harm in their current form perform social functions (cohesion, identity, mutual aid, meaning infrastructure) that are not automatically

provided by their abolition, and evaluating the practice requires accounting for what happens to those functions when the practice changes. Third, specifying the realistic alternatives — the judgment is comparative, not absolute, and an alternative that is formally less harmful but practically unavailable produces a different evaluation than one that is materially achievable. Fourth, attending to who is making the evaluation and from what structural position — the analyst's own positionality shapes what harms are visible and which alternatives seem realistic. The internal versus external critic distinction: The consistency standard does not require that all criticism of cultural practices carry equal weight regardless of its source. Criticism emerging from within a community — from members who bear the costs of the practice and whose identification with the community is not in question — operates differently from criticism emerging from outside, for structural reasons: internal critics have access to the practitioner knowledge that external critics typically lack, their claims cannot be dismissed as motivated by hostility to the community, and their proposals for change are more likely to account for the social functions the practice performs. None of this means external criticism is never valid — some practices produce harms whose victims have

insufficient standing within the community to make internal criticism effective, and some external analysis identifies structural dynamics that internal critics are too embedded to see clearly. But the distinction matters for how criticism is weighted and what accountability the critic bears for the consequences of their analysis. The hardest cases are not the practices that are straightforwardly harmful with no compensating function — those are analytically tractable even if politically difficult. The hardest cases are practices that produce genuine harm for some members of the community while simultaneously performing functions that those same members depend on: the institution that oppresses and also sustains. The moral-functional tension analyzed in 13.4h applies here directly. What the framework requires in those cases is not a resolution that pretends the tension does not exist but an honest account of what is being traded off — whose harm is being accepted, whose need is being prioritized, and what organizational capacity would be required to maintain the functions while changing the form. Diagnostic for cultural practice evaluation: (1) Who within the community bears the costs of this practice, and do they have effective standing to contest it? (2) What social functions does the practice perform that would not automatically be provided by its modification or abolition? (3) What are the realistic alternatives, not the ideal alternatives? (4) Is the evaluative standard being applied here the same standard that would be applied to an equivalent practice in a dominant-culture context? If not, what justifies the asymmetry? (5) Who is conducting the evaluation, from what structural position, and what accountability do they bear for the consequences of their analysis?

8.2 — Marginalized Groups and the Material Frame

Structural exclusion from resources, opportunity, and institutional protection is what the minimal anchor requires addressing — not because of identity politics, but because exclusion produces unnecessary suffering, and the minimal anchor is not selective about what kind of institutional arrangement generates it. The analysis begins from what is actually happening to people and what arrangements are producing it, not from the political valence of the group experiencing those conditions. The distinction between materialist and symbolic approaches to the same stated goal produces genuinely different analyses and genuinely different prescriptions. The materialist analysis of how colonial arrangements converted genuine communal identity into instruments of division and extraction remains one of the most important available — both for understanding structural exclusion and for understanding why movements that addressed its symbolic dimensions without its material ones consistently fell short. (Fanon 1961) A materialist approach asks what arrangements actually produce better outcomes — better access to resources, genuine security, real opportunity. A symbolic approach asks what arrangements produce recognition and cultural validation for the relevant identity category. These questions can have different answers, and in practice frequently do. Arrangements that maximize symbolic recognition while leaving material conditions unchanged — or while generating political backlash that worsens material conditions — are not progress by any standard that the

minimal anchor supports. A materialist analysis of structural exclusion must account for the forms of labor that do not appear as labor in the dominant economic accounting but that are nevertheless foundational to the economic arrangements it analyzes. The work of biological and social reproduction — bearing and raising children, caring for the sick and elderly, maintaining the conditions under which waged labor is possible — is not incidental to the economic system. It is the condition without which the economic system cannot function, and it is work that is systematically uncompensated, rendered invisible, and assigned to specific populations on the basis of gender. Any economic analysis that does not account for this work is analyzing a fraction of the material economy and mistaking it for the whole. (Federici 2004; Fraser 2016) The structural exclusions that produce unnecessary suffering rarely operate through a single axis. A person excluded from material welfare is typically excluded through the simultaneous interaction of multiple structural positions — class, gender, race, geography — in ways that are not additive but mutually constituting. The material conditions of a woman working in informal labor in a post-colonial economy are not the material conditions of poverty plus the material conditions of gender subordination plus the material conditions of colonial inheritance as separate quantities. The

y are a single set of conditions produced by the intersection of all three. The framework's insistence on material analysis does not simplify this intersection; it requires the analysis to be adequate to its actual complexity. What the material frame adds to identity analysis: Identity-based political frameworks correctly identify that membership in socially categorized groups produces differential exposure to specific forms of harm — that race, gender, caste, sexuality, and disability produce structural vulnerabilities that are not reducible to class position and that class analysis alone consistently underweights. The material frame does not dispute this identification; it asks a supplementary question: what are the specific material mechanisms through which group membership produces differential outcomes, and what institutional arrangements maintain those mechanisms? The question is supplementary rather than alternative because the answer to it determines what kinds of interventions are adequate. Naming the group whose members suffer differentially is a necessary starting point; identifying the material mechanism that produces the differential is necessary for designing interventions that address it rather than its symptoms. The structural versus attributive distinction: The most consequential distinction the material frame introduces into the analysis of marginalized groups is between structural and attributive explanations of differential outcomes. Structural explanations attribute differential outcomes to institutional arrangements — to the distribution of resources, the design of opportunity structures, the operation of formal and informal exclusion mechanisms — that would produce the same outcomes for any population placed in the same structural position. Attributive explanations attribute differential outcomes to characteristics of the group itself — to culture, values, behaviors, or capabilities that distinguish the group from others and that explain why the group occupies the structural position it does. The distinction matters practically because structural explanations point toward structural interventions and attributive explanations point toward interventions directed at the group itself. Misidentifying a structural cause as an attributive one produces interventions that are addressed to the wrong level, that typically fail, and that in failing can be used to support the attributive explanation — a self-reinforcing error. The tokenism problem: Institutional arrangements that produce systematic structural exclusion of a group generate political pressure to include members of that group in visible positions. The tokenism problem is the structural dynamic through which that pressure is addressed in ways that reduce its political urgency without changing the structural arrangements that produced the exclusion. Placing members of the excluded group in visible positions within the existing institutional arrangement serves the institution

's legitimacy interests without requiring the arrangement to change the resource distribution, access structures, and informal exclusion mechanisms that maintain structural disadvantage for the group as a whole. The material frame's contribution to the analysis of tokenism is the question it forces: does this inclusion change the structural position of the group, or does it change the institutional appearance while leaving the structural position unchanged? That question is answerable with evidence about resource distribution, opportunity access, and harm rates, not with evidence about the number of visible

representatives. The relationship between group-level analysis and individual-level variation is a persistent source of confusion that the material frame clarifies. Groups defined by social categorization — racial groups, gender categories, caste positions — contain individuals with enormous variation in material conditions, structural position, and experience of specific harms. The group-level analysis identifies structural patterns that are real and consequential without implying that every member of the group experiences those patterns identically. The pattern that white Americans have lower rates of poverty than Black Americans does not imply that no white Americans experience poverty or that all Black Americans do; it identifies a structural regularity with a structural cause that is analytically prior to the individual variation within each group. Conflating group-level structural analysis with claims about individual members is an error in both directions: it produces both the error of treating structural patterns as individual destiny and the error of treating individual exceptions as refutations of structural regularities. The charge that biological foundations naturalize gender difference deserves a direct response rather than an indirect one. Biological foundations establish that organisms have bodies, require care under vulnerability, and depend on the reproduction of the conditions that make social life possible. They do not establish which specific institutional arrangements for managing reproduction and care are natural or inevitable. The framework's treatment of reproductive labor throughout this section and in 11.0 diagnostic 7 is precisely the opposite of naturalization: it identifies the invisibility of reproductive labor in dominant economic accounting as a structural distortion — produced by specific institutional arrangements that assign specific work to specific populations, present that assignment as natural, and then analyze only the economic activity that rests on it. A framework that begins from biological vulnerability and treats the social management of that vulnerability as its primary analytical object does not commit to any particular social arrangement for managing it. What it commits to is analyzing those arrangements honestly — including the ones that have presented themselves as biology rather than as politi

cs. The material frame does not replace the analysis of cultural, historical, and experiential dimensions of group identity — it provides the structural foundation on which those dimensions can be analyzed without dissolving into voluntarism or cultural essentialism. Groups are not merely statistical categories; they have histories, internal cultures, and collective identities that are real and analytically consequential. The material frame asks how those histories, cultures, and identities interact with current structural positions to produce the specific forms of suffering and resilience that characterize each group's situation — not as the only question but as the question whose answer is most directly connected to what interventions would change the situation.

8.2b — Intersectional Analysis as Methodological Procedure

Acknowledging that structural exclusions operate through multiple intersecting axes — class, gender, race, geography — is necessary but not sufficient. It does not specify how to analyze the intersection procedurally: which axis is most explanatory in a given context, how they interact to produce conditions that cannot be derived from any single axis alone, and what kind of evidence would determine whether a proposed intervention addresses the intersection or only one of its components. Without a procedure, intersectionality can become a rhetorical gesture — the observation that things are complicated — rather than an analytical tool that produces specific findings. The primary variable question. In any specific configuration of conditions, the multiple axes of structural exclusion interact but they do not interact symmetrically. One axis is typically primary in the specific sense that it most directly determines the material constraint being analyzed: the variable whose change would produce the most significant change in outcome. Identifying the primary variable is an empirical question, not a definitional one — it cannot be settled in advance by theoretical commitment to the centrality of any particular axis. A worker in an informal economy may face structural constraints whose most direct determinant is geography (location outside formal economic infrastructure) or class (lack of capital for entry into formal employment) or gender (legal and social norms

restricting specific forms of economic participation) — and the answer will differ across cases that superficially resemble each other. The procedure requires asking: if this one axis were changed while the others remained constant, what would change in the material condition being analyzed? The axis whose change produces the most significant material change is primary for that analysis. The interaction test. After identifying the primary variable, the procedure requires asking how the other axes modify its effects. Modification is not addition: the effect of class on access to healthcare for a woman in a post-colonial economy is not the effect of class plus a separate effect of gender plus a separate effect of colonial inheritance. These axes modify each other's effects. The interaction test asks: does the primary variable produce different effects depending on the specific configuration of secondary axes? If yes — if class produces different healthcare outcomes for men and women, for colonial and non-colonial contexts, for urban and rural geographies — then the analysis must specify which configuration it is analyzing, and must check whether an intervention on the primary variable will have the same effect across configurations. An intervention that addresses the primary variable without accounting for how secondary axes modify its effects is likely to produce differential outcomes that reproduce or deepen existing exclusions along the axes it ignored. The residual test. After accounting for

the primary variable and the modifications introduced by secondary axes, the procedure asks: is there a material condition in this population that cannot be explained by any of the axes individually or by their pairwise interactions? If yes, that residual is evidence of a genuine intersection effect — a condition produced by the specific combination of multiple axes that would not be predicted from knowing each axis separately. Residual effects are important because they identify the conditions that require genuinely intersectional interventions — interventions that address the combination, not just the components. They also identify where standard single-axis analyses will consistently fail, regardless of their rigor, because they are analyzing the wrong unit. The procedural account of intersectionality does not simplify the complexity. It makes the complexity tractable: it specifies what questions to ask, in what order, and what kind of evidence would answer them. A materialist framework that acknowledges intersectionality without this procedure is in the position of acknowledging that the world is complicated and stopping there — which is not analysis but its abdication. This procedure is not a formula that produces automatic answers. It is a sequence of questions that structures the empirical investigation. The answers depend on the specific material conditions being analyzed, not on theoretical commitments about which axes are universally primary. In some contexts, geography is the most explanatory variable. In others, class or gender or race produces the largest material effects. The procedure is designed to determine which case obtains rather than assuming the answer in advance — which is the only posture consistent with the epistemological commitment of Part I. A concrete illustration: consider the material conditions of women working in urban informal economies across different post-colonial contexts. A single-axis class analysis identifies their position in the informal sector as the primary constraint and recommends formalization as the intervention. A single-axis gender analysis identifies legal and social norms restricting their economic participation and recommends legal reform. Both are partially right. The primary variable question asks which axis, if changed while the others remained constant, would produce the most significant material change — and the answer differs by context. In a setting where formal employment is structurally unavailable regardless of gender (collapsed formal sector, geographic isolation), class and geography are primary; gender norms modify how the informal economy is experienced but do not determine access to it. In a setting where formal employment exists but women are legally excluded from specific sectors, gender is primary; class position modifies the severity of exclusion but does not determine it. The interaction test then asks whether the primary variable produces different effects across the seco

ndary axes: formalization, where it is available, may produce significantly different outcomes for women than for men if the formal sector reproduces gender-based wage discrimination — in which case an intervention on the primary variable (class position) fails on the interaction test, because it does not account for how gender modifies its effects. The residual test asks whether there remain conditions — specific

vulnerabilities, specific exclusions — that cannot be explained by class, gender, or geography individually or in combination: the additional precarity produced by being simultaneously female, informal, and geographically peripheral may exceed what any single-axis account predicts, identifying a genuinely intersectional effect that requires an intersectional intervention. The procedure does not produce this conclusion automatically. It identifies where to look for it. A second illustration, from a different axis configuration: consider access to formal credit markets for small business operators across different racial and class positions in a society with a documented history of racially discriminatory lending. A single-axis class analysis predicts that credit access is primarily determined by collateral, income history, and business viability — and recommends expanding access through microfinance or credit scoring reform. A single-axis race analysis predicts that racial discrimination by lenders is the primary constraint and recommends anti-discrimination enforcement. The primary variable question requires asking which axis, if changed while the others held constant, would most change credit outcomes — and the honest answer is that available evidence typically shows both operating, with their relative weight varying by institutional context and historical moment. The interaction test reveals something that neither single-axis analysis predicts: controlling for class (income, collateral, business viability), racial disparities in credit access frequently persist and sometimes increase, suggesting that race modifies the effect of class on credit outcomes rather than simply adding a separate effect. Middle-class operators of the disadvantaged racial group face different rejection rates than middle-class operators of the advantaged group with equivalent economic profiles. The residual test asks whether there remain patterns unexplained by race and class individually or interactively — and frequently finds them in the form of geographic redlining effects: the location of the business in a neighborhood historically marked as high-risk produces credit constraints that cannot be fully accounted for by the operator's individual economic profile or racial identity, because the neighborhood marking itself aggregates historical discrimination into a spatial variable that disciplines current lending regardless of individual circumstances. The appropriate intervention differs depending on which of these effects is primary in a specific context: an

intervention that successfully eliminates individual racial discrimination in lending may leave geographic redlining effects intact, and vice versa. The procedure shows that an adequate response must address not only individual discrimination but the spatial legacy that operates independently of any current actor's intentions — which is precisely what single-axis analysis, however rigorous within its own terms, structurally cannot see. Social Reproduction: The Invisible Infrastructure of All Production The analysis of structural suffering and becoming capacity developed in this framework rests on a foundation that the framework's own accounting has not yet made explicit: the reproductive labor that produces and sustains the human beings in whom suffering is experienced and becoming is possible. Bearing children, feeding and raising them, caring for the sick and the elderly, maintaining the social bonds and domestic environments that allow people to present themselves for productive and political life — this work is the precondition for everything else the framework analyzes. Without it, there are no workers, no political actors, no populations in whom becoming can be expanded or suffering reduced. It is, in the most literal sense, the infrastructure of the infrastructure. The systematic invisibility of reproductive labor in economic accounting and political analysis is not a measurement error that better statistics could correct. It is a form of visibility suppression — V in the framework's structural analysis — that is produced by the same institutional arrangements that benefit from the labor remaining uncompensated. When reproductive labor is not counted as work, it does not appear in calculations of labor value, it does not generate claims on social insurance, it does not accumulate toward retirement security, and it does not register as a cost in the production of goods and services that depend on it. The labor is extracted; the extraction is made invisible; the invisibility prevents the extraction from appearing as a form of structural suffering or a site of political contestation. Silvia Federici's analysis of the historical separation of reproductive from productive labor as a constitutive act of capitalist accumulation, and Nancy Fraser's analysis of the contradictions between capitalism's dependence on reproductive labor and its systematic refusal to sustain the conditions for that labor, provide the historical and theoretical basis for

this structural assessment. The gendered distribution of reproductive labor is not incidental to its invisibility — it is the mechanism through which the invisibility is sustained. When reproductive labor is assigned primarily to women through a combination of biological naturalization, normative prescription, and structural coercion, it acquires the appearance of a natural extension of female biology rather than socially organized work that serves economic and political functions. This naturalization is

not ideological in the sense of being merely a false belief that could be corrected by better information. It is materially reproduced through the institutional arrangements — inheritance law, taxation of household labor, pension systems that reward continuous formal employment, childcare provision that treats parental care as a private preference rather than a social investment — that make the gendered distribution of reproductive labor the path of least resistance for most households navigating material constraint. The analysis of gender as a dimension of structural suffering in the framework therefore requires analysis of reproductive labor distribution as the material basis from which gender inequality is reproduced, not merely as one expression of cultural attitudes toward women. The relationship between reproductive labor and the intergenerational transmission variable *G* is direct and underanalyzed in existing frameworks. *G* captures the degree to which the structural conditions of one generation are transmitted to the next — through inherited wealth, institutional access, educational opportunity, and the social networks that structure life chances. Reproductive labor is the mechanism through which *G* operates at its most fundamental level: the cognitive development, physical health, emotional regulation, and social capacity of the next generation are produced primarily through the quality and quantity of reproductive labor invested in their early development. When reproductive labor is underfunded, undersupported, and assigned to those with the fewest material resources — when poor women bear the reproductive labor burden with the least institutional support — the result is not merely inequality in the present but systematic degradation of the human capacities that *G* transmits. The intergenerational reproduction of structural suffering runs through the distribution of reproductive labor as one of its primary mechanisms, and no analysis of *G* that treats reproductive labor as a background condition rather than a structural variable is capturing the full dynamic. The structural implication: an institutional arrangement's treatment of reproductive labor is not a gender policy question separate from the framework's core concerns. It is a direct index of the arrangement's relationship to the evaluative criterion. An arrangement that extracts reproductive labor without compensation or social recognition is producing a form of structural suffering — resource deprivation, visibility suppression, becoming foreclosure — that operates at the most foundational level of social reproduction. An arrangement that invests in reproductive labor — through universal childcare, parental leave structured to equalize the labor burden, social recognition of care work, and elderly care that does not depend primarily on unpaid family labor — is investing in the preconditions for preserved adaptive response capacity at the population level. This is not redistribution

as a luxury of wealthy societies. It is the material infrastructure without which all other investments in becoming are built on a foundation that is continuously being undermined.

8.3 — The Consistency Standard

Standards applied to some actors and not others are not standards. They are preferences dressed as principles. The consistency requirement — applying the same analytical lens to actors and arrangements one is sympathetic to as to those one is critical of — is demanding precisely because it generates conclusions that are uncomfortable in both directions. A framework that consistently produces uncomfortable conclusions about opponents and comfortable conclusions about allies is not doing analysis. It is producing ideology. The test of whether a framework is doing what it claims to do is whether it can generate conclusions that are genuinely costly to the person applying it. In practice, the consistency standard generates three specific demands. First, the same evaluative criterion — does this arrangement create the conditions for genuine development and reduce the structural suffering that forecloses it? — is applied to movements and institutions the analyst supports with the same rigor as those they oppose. A progressive framework that

examines the harms produced by conservative institutions but exempts progressive movements from equivalent examination has not applied the minimal anchor. It has selectively applied it. Second, the same epistemic standards that are demanded of opponents — specificity of evidence, acknowledgment of counterexamples, responsiveness to contrary evidence — are demanded of one's own positions. Third, the structural explanations offered for opponents' failures — incentives, interests, capture — are applied with equal force to one's own movements, rather than being replaced in that case by accounts of individual betrayal or external interference. The consistency standard does not require treating all actors as equivalent or all positions as equally supported by evidence. It requires that the standard of evidence and scrutiny be equivalent, not that the conclusions be symmetric. If the evidence genuinely shows that one arrangement produces more structural suffering than another, the consistent application of the criterion will produce that asymmetric conclusion — honestly, because the criterion has been applied consistently, not selectively. The consistency requirement is uncomfortable not because it is difficult to understand but because its consistent application produces conclusions that most analysts are structurally positioned to resist. If the framework identifies institutional capture as a mechanism that operates regardless of the ideological content of the institution being captured, then progressive institutions are as subject to capture as conservative ones — and the analysis of capture in a progressive institution will be resisted by actors who identify with its stated mission as strongly as the analysis of capture in a conservative institution is resisted by actors who benefit from its current operation. The resistance is symmetrical in structure even when its political valence is asymmetric. Acknowledging this i

s a precondition for the framework functioning as analysis rather than as rationalization. Asymmetric application as ideological marker: The most reliable marker that a framework is functioning as ideology rather than analysis is that it applies its most demanding scrutiny to the arrangements it opposes while applying its most charitable interpretation to the arrangements it favors. This asymmetry is visible across the political spectrum: conservative frameworks that apply rigorous cost-benefit analysis to welfare programs while exempting military expenditure from equivalent scrutiny; progressive frameworks that document institutional capture in corporate-funded organizations while explaining away equivalent capture dynamics in union bureaucracies or NGO hierarchies. The asymmetry is not evidence that the analysis is wrong in either direction — it may be that the corporate-funded organization really is more captured, and the evidence should be examined. The asymmetry is evidence that the framework is applying different standards depending on the political valence of its subject, which means it is not functioning as analysis. Consistency and the ally relationship: The consistency standard generates its most acute political discomfort in the analysis of organizations, movements, and actors that are doing genuine good — whose work reduces structural suffering, whose counter-power function is real and valuable, and whose institutional health is a matter of practical consequence for the populations they serve. The consistency standard requires applying the same analysis to these actors that would be applied to their opponents: asking whether their institutional practices match their stated commitments, whether their resource allocation reflects their stated priorities, whether their accountability mechanisms are genuine or performative, and whether their organizational culture produces the drift dynamics identified in the ecosystem analysis. This analysis is not hostile to the organizations it examines — it is the precondition for those organizations remaining adequate to their functions over time. Organizations that are exempt from internal critique on the grounds that the critique might damage them are organizations that have successfully insulated their institutional interests from the accountability that their stated mission requires. Self-consistency: The consistency standard applies to the framework itself with the same force it applies to everything the framework analyzes. PMN makes claims about how power operates, how institutions develop, how knowledge is produced, and what structural conditions generate unnecessary suffering. Those claims should be held to the same evidential standard that the framework requires of other analytical claims — subject to revision when evidence contradicts them, not protected from that evidence by the framework's own authority. The self-application requirement from 1.2 and 6.4 is the internal mechanism that oper

ationalizes the consistency standard: the framework cannot exempt its own governing assumptions from the diagnostic it applies to the governing assumptions of others without reproducing, at the level of intellectual practice, the same pattern it identifies as analytical failure everywhere else. The practical discipline the consistency standard requires: before concluding that an actor, organization, or arrangement is or is not subject to a specific structural dynamic, ask whether the conclusion would change if the political valence of the actor were reversed. If the analysis would produce a different conclusion for a conservative organization than for an equivalent progressive one — if the dynamics of institutional drift, resource capture, or narrative maintenance are explained away in one case and emphasized in another — the analysis is functioning as advocacy. Advocacy is sometimes the right thing to produce; it is not the same thing as analysis, and conflating them produces work that fails at both.

8.4 — Governing Narrative as Material Force: The Social Construction of Reality

The term used here requires immediate clarification. A governing narrative — what this section calls a fable, in the specific sense of a story that structures action regardless of its literal truth — is not a children's story or a fiction. It is any narrative sufficiently embedded in how a community interprets events that it governs responses without requiring verification. The word 'fable' is used not to imply falsity but to mark the structural feature that matters: the narrative operates prior to the evidence, shaping what the evidence is taken to mean. In a world we like to believe is governed by facts, most facts remain silent unless someone chooses to interpret them. A person can possess genuine character — of any kind — but if the social narrative constructed around them tells a different story, that story will govern how others respond to them, what opportunities they receive, what coalitions they can form, and ultimately what they can become. Individual actors are not behaving irrationally. The pattern is a structural feature of how social reality is produced and maintained in conditions of limited observation and contested information. The accuracy of a narrative becomes less important than the fact that it is the narrative people choose to believe and act on. Within a material framework, this is not an unfortunate distortion of reality — it is reality operating at the social level. Human beings cannot directly observe each other's inner states, intentions, or moral character. They can only observe actions, appearances, signals, and the reports of others. From these fragments they construct narratives and treat those narratives as the person, the institution, the movement, the community. The narrative becomes more real, in the social sense that matters for political action, than the underlying reality it purports to describe. This has immediate consequences for how power operates. Those who control the production and distribution of narratives control a form of power that is more durable than coercion, because it operates on what people can imagine as well as on what they can do. The destruction of an actor's narrative — the controlled collapse of the story others tell about them — is as effective a form of neutralization as any physical constraint, and in most conditions more efficient. Conversely, an actor whose narrative is secure can absorb significant material setbacks without losing the capacity to operate, because the cooperative relationships that narrative sustains remain intact. Observation is the mechanism through which narrative is produced. Actions that are not witnessed do not enter the social world as stories. An actor who operates entirely without witnesses has neutralized the narrative production mechanism — but at the cost of also neutralizing the narrative benefits that come from witnessed action. The political implication is that visibility is a resource to be managed strategically, not a background condition to be ac

cepted passively. What is made visible, to whom, through what channels, under what framing, is one of the most consequential decisions any political actor makes. Modern institutions attempt to manage this problem by creating procedural observers — audit systems, documentation requirements, accountability mechanisms, public records. These are not primarily moral instruments. They are epistemological instruments: attempts to expand the space of what can be verified against the tendency of private narratives and strategic framing to collapse social truth into whatever story the most powerful actor wants told. The failure of these mechanisms — and they fail regularly, systematically, and in predictable patterns — is not primarily a failure of

individual corruption. It is the success of actors who understand that controlling the narrative production mechanism is more valuable than controlling any specific narrative. The concept of governing narrative operates at multiple levels simultaneously. A single action can constitute a complete such narrative within a longer life — a moment of desperation, of courage, of betrayal, of unexpected solidarity — that stands alone with its own logic and its own consequences. These small arcs compose into larger narratives that exceed what any actor intends. The framework for thinking about narrative as material force has to operate at all these levels: the individual action, the accumulated pattern, the institutional story, the civilizational myth. What connects them is the same mechanism: observation, interpretation, repetition, belief. Why governing narratives function independently of their truth value: The concept of the governing narrative as a material force does not require the claim that the narrative is false. Some governing narratives are substantially accurate descriptions of the conditions they interpret. What makes them material forces is not their falsity but their function: they govern responses by providing an interpretive framework that determines what counts as relevant evidence, what causal attributions are available, and what actions appear as reasonable responses to observed conditions. The same event — a factory closure, an electoral defeat, an act of violence — will be interpreted differently and will generate different political responses depending on which governing narrative is operative. The narrative is material because it shapes material action, not because it is materially false. The political economy of narrative maintenance: Governing narratives require active maintenance. They do not sustain themselves through their truth-value alone; they are sustained through the investment of resources by actors who benefit from the interpretive framework the narrative provides. Media ownership, educational curriculum, professional credentialing standards, the selection of which historical events are commemorated and how — all of these are material investments in narrative maintenance, and the actors

who make those investments are systematically those whose position in the current arrangement is legitimated by the narrative. The political economy of narrative maintenance explains why narratives that are empirically contestable persist long past the point where their contestability is visible: the resources available to maintain them are proportional to the interests they serve, which are often the interests with the greatest resource access. The implication for counter-power strategy: narrative infrastructure is not secondary to material infrastructure — it is a form of material infrastructure, and the strategic investment in developing counter-narratives that expose the structural causation currently obscured by governing narratives is analytically equivalent to the investment in organizational and resource infrastructure. Both are necessary; neither alone is sufficient.

8.4b — Narrative Typology: How Governing Narratives Structure Political Time

Section 8.4 established that governing narratives are material forces: they shape what populations can perceive as possible, what coalitions can form, and what institutional arrangements retain legitimacy. One dimension of that analysis has not yet been developed. Governing narratives do not only structure how populations interpret the present — they structure how populations relate to time itself. The temporal orientation embedded in a narrative determines the horizon across which collective action can be sustained, the direction from which legitimacy is claimed, and the emotional register through which transformation pressure is experienced. Two societies facing identical material conditions may respond with entirely different political trajectories depending on the temporal structure of their dominant narratives. This is not a cultural epiphenomenon. It is a material variable operating through the duration multiplier (D) in the transformation pressure formula: the longer a population can sustain coherent orientation toward change, the more transformation pressure accumulates. What sustains or collapses that orientation is largely a function of narrative. Three temporal orientations recur across historical and contemporary political contexts, each with distinct implications for how material conditions translate into collective action. Past-Oriented Narratives: The Golden Age The most structurally conservative temporal narrative is the golden age: a mythologized past that is posited as superior to the present and recoverable through the right political action. The content

varies enormously — the restored caliphate, the pre-colonial nation, the mid-century welfare state, the pre-immigration monoculture, the revolutionary founding moment — but the structure is constant. Present suffering is framed as deviation from a prior condition of wholeness, and the political task becomes restoration rather than construction. The analytical significance for the framework is not that this narrative is false — sometimes the invoked past was materially better in specific respects — but that it converts the trajectory of becoming into a trajectory of return. Where 4.5 establishes that the framework's horizon is genuine development rather than the recovery of a prior state, golden age narratives systematically reorient collective action away from that horizon. They are powerful precisely because they require less from those they mobilize: return is easier to imagine than genuine transformation, and the emotional register of recovery is more immediately legible than the register of becoming. Golden age narratives tend to emerge and intensify when the trajectory of becoming appears foreclosed — when the conditions for genuine development are contracting rather than expanding. This is the structural connection to the suffering formula: when R increases, B increases, and the future appears blocked, the past becomes the available source of legitimating imagination. The

political danger is not that the narrative mobilizes — mobilization under any narrative may produce genuine change — but that it systematically misdirects the analysis of what produced the present condition and what would be required to change it. A movement organized around return cannot build the institutional arrangements that did not exist in the invoked past. Present-Oriented Narratives: Stability and Order Where golden age narratives redirect action toward the past, stability narratives immobilize action in the present. The structure is: the current arrangement, however imperfect, is preferable to the chaos that change would produce. This is not a claim about the goodness of the present but about the danger of transition. Its political power derives from the genuine uncertainty of what transformation produces — historical cases in which revolutionary change produced worse conditions are real, and the stability narrative exploits that record strategically. The legitimacy claim is not “things are good” but “things could be worse, and you know what that looks like.” Stability narratives are the characteristic governing narrative of systems under transformation pressure. When $R \uparrow$, $B \uparrow$, and $V \downarrow$, the system faces rising pressure for change. The institutional response is not typically to reduce the structural suffering that produces the pressure — that would require changing the arrangements that benefit those who control the narrative. It is to increase the perceived cost of change. The same conditions that generate transformation pressure are thus used to argue against transformation: instability, uncertainty, and the threat of worse outcomes become the instrument of legitimacy maintenance rather than evidence of legitimacy failure. Recognizing this pattern is analytically important: the intensity of stability rhetoric is often a leading indicator of the system's actual fragility, not its strength. Future-Oriented Narratives: Eschatology, Utopia, and Catastrophe Future-oriented narratives are the most varied class and have the most complex relationship to transformation pressure. They share the structure of projecting a decisive future state — salvation, collapse, liberation, catastrophe — that gives present action its meaning. Their political function depends critically on the structure of that projected future: whether it is achievable through collective action or arrives independently of it, and whether it is presented as possibility or inevitability. Eschatological narratives — in both religious and secular forms — project a final transformation that resolves the tensions of the present absolutely. The political effect is double-edged. Where the final transformation requires active human participation, eschatology is a powerful mobilizer: it extends the duration multiplier D by giving present sacrifice meaning across an indefinite temporal horizon, and it resolves the coordination problem by making defection cosmically rather than merely po

litically costly. Where the final transformation is inevitable regardless of human action, eschatology functions as a demobilizer: if the end is coming, the present arrangement is temporary by definition, and sustained organizational effort seems unnecessary. The distinction between activist and quietist eschatology is therefore not primarily theological — it is a distinction between narratives that extend and narratives that collapse the duration multiplier. Utopian narratives function similarly but without the finality. They project a qualitatively better future that is achievable but not inevitable, requiring sustained collective effort. Their

analytical significance is that they extend the temporal horizon of collective action beyond what immediate material interest can sustain: people will organize and sacrifice for a future they will not personally inhabit if the narrative structures that future as genuinely achievable and genuinely worth achieving. The anti-utopian counternarrative — that utopian projects inevitably produce catastrophe, that any attempt at fundamental transformation leads to totalitarianism — is therefore not primarily a philosophical claim. It is a political intervention that collapses the temporal horizon of collective action to the immediate, making sustained transformation effort structurally unavailable. In the framework's terms, anti-utopianism reduces D regardless of the material conditions that might otherwise sustain it. Catastrophic narratives deserve separate treatment because they are structurally ambiguous in a way the other categories are not. A narrative of impending civilizational collapse can function either as a mobilizer — act now before it is too late — or as a demobilizer — nothing can stop what is coming. The political valence depends on whether the narrative includes an available action that can prevent or mitigate the catastrophe. Climate catastrophism without a credible transition pathway functions as demobilization regardless of its empirical accuracy. The analytical task for a PMN-aligned analysis is not to evaluate the accuracy of catastrophic narratives but to assess what their structure does to the duration multiplier and the collective action capacity of the populations that hold them. Temporal narrative structure is a variable, not a background condition. The framework's analysis of transformation pressure ($T = S \times D \times P$ (with R_c shaping the threshold that T must cross, and $T_r = T / (C \times L) - CP$ governing trajectory at 15.9)) tracks the structural conditions under which systems become unstable. But whether accumulated pressure produces transformation depends also on whether the populations experiencing that pressure can sustain oriented collective action — which depends on the duration multiplier D — which depends significantly on whether the dominant temporal narrative extends or collapses the horizon across which action can be organized. Governing narratives about time are not merely inter-

pretive. They are part of the material conditions that determine what accumulated pressure can do. Collective Memory as Material Force: History, Counter-Memory, and Intergenerational Transmission Governing narratives, as section 8.4 analyzes, are not only produced and contested in the present. They are transmitted across generations through the institutions that organize collective memory: educational curricula, commemorative practice, archives and monuments, official histories and the counter-histories that challenge them. The temporal dimension of narrative politics has specific structural properties that present-focused analysis misses. A governing narrative that has been successfully institutionalized across multiple generations operates as a constraint on what political positions are imaginable, not merely what positions are currently dominant. The populations that inherit it do not experience it as a narrative — as one interpretation among possible alternatives — but as a description of what happened and what it meant. Contesting it requires not only the organizational capacity to advance an alternative account but the prior work of making the governing narrative visible as a narrative, which is a more difficult epistemological task than contesting a policy position that is already understood as a choice. Regimes systematically invest in the institutional production of collective memory because governing the past is a prerequisite for governing the present. Who is counted as a national hero and who as a traitor, which events are commemorated and which are suppressed, what causal account is taught in schools about how current arrangements came to exist — all of these decisions structure the political imagination of subsequent generations in ways that constrain what transformations are conceivable as legitimate. The management of historical memory about episodes of state violence, colonial extraction, and structural dispossession is particularly consequential: when the populations who bore the costs of these episodes are taught accounts that attribute their current conditions to their own failures or to natural processes rather than to identifiable institutional decisions, the political pressure that accurate historical understanding would generate is preemptively defused. Intergenerational transmission G therefore includes not only the transmission of material conditions but the transmission of the interpretive frameworks through which those conditions are understood — and the institutional capture of those frameworks is among the most durable forms of structural power available. Counter-memory — the production and circulation of historical accounts that contest official narratives about the origins of current

structural arrangements — is among the most important forms of epistemic counter-power that social movements can develop. Its structural significance lies not primarily in its historical accuracy, though accuracy matters, but in its fu

nction of making the existing governing narrative visible as a narrative: as an account that was produced by specific actors with specific interests, at specific historical moments, and which can therefore be contested and replaced. The practice of counter-memory is materially consequential in the most direct sense: it changes what transformations are politically conceivable by changing what past the present is understood to follow from. Movements that have succeeded in establishing a counter-historical account of structural conditions — that slavery was not a peripheral aberration but a constitutive feature of economic development, that colonial extraction was not civilizational uplift but systematic dispossession, that the suppression of labor movements was not natural market adjustment but organized political violence — have consistently found that the change in political imagination that follows from this account opens possibilities that were previously foreclosed not by lack of material resources but by lack of political imagination.

8.4c — Meritocracy as Visibility Suppression

Section 8.4 identified visibility as a strategic resource: what is made observable, to whom, through what framing, determines whether structural suffering enters the social world as a problem requiring structural response. The most effective mechanism of visibility suppression is not the active concealment of information. It is the production of an interpretive framework that makes structural suffering invisible by recategorizing it as individual outcome. Meritocratic narrative is the dominant contemporary form of this mechanism. The structure of meritocratic narrative is specific. It does not claim that outcomes are equal — it claims that outcomes are deserved. The person who accumulates resources did so through superior effort, talent, or virtue. The person who does not accumulate resources failed to apply effort, lacked talent, or made poor choices. The institutional arrangements that produced both outcomes are, on this account, neutral channels through which individual differences express themselves. The arrangement is just because it reflects merit. Suffering within it is not structural — it is the accurate signal of individual shortcoming. Where the arrangement is imperfect, the remedy is better measurement of merit, not structural change. In the framework's terms, this is a direct operation on V — the visibility suppression variable in the structural suffering formula. V measures the degree to which structural suffering is not recognized as structural: the extent to which populations experiencing deprivation produced by institutional arrangements attribute that deprivation to causes that do not implicate those arrangements. Meritocratic narrative produces V not through the suppression of facts but through the provision of an alternative causal account. The facts of differential outcome are not hidden. The structural causation of those outcomes is recategorized as individual causation. The suffering is visible; its structural character is not. This mechanism has three consequences that the framework's analysis of institutional change makes analytically significant. First, it suppresses the coalition formation that transformation pressure requires. Populations experiencing the same structural conditions cannot form political coalitions around shared structural interests if each individual attributes their condition to their own failures. The horizontal solidarity that counter-power requires — the recognition that what I experience is what you experience, and that it is produced by the same institutional arrangement — is systematically undermined by a narrative that makes each person's condition uniquely their own. Second, meritocratic narrative converts the institutional arrangements that produce structural suffering into the instruments of individual aspiration. The same labor market that produces precarity for most of its participants is also the mechanism through which exceptional individuals escape precarity. The narrative of meritocrati

c escape — that sufficient effort and talent can produce exit from structural deprivation — directs the energy of those experiencing deprivation toward individual mobility rather than collective change. It does not suppress the desire for better conditions; it redirects that desire away from structural transformation and

toward individual competition within the existing arrangement. Third, it produces a specific form of legitimacy for institutional arrangements that would otherwise fail the framework's evaluative criteria. An arrangement that produces unnecessary suffering — imposed, dysfunctional, preventable, structural — cannot be defended on its merits if those features are visible. Meritocratic narrative makes those features invisible by reclassifying the suffering as chosen (the person could have worked harder), functional (suffering motivates effort), individual rather than structural (the arrangement is not producing the suffering; the individual's choices are). Each of the four criteria of unnecessary suffering in 3.4b is addressed by the narrative, not by changing the conditions. The arrangement is legitimated without being improved. The analytical implication is that V is not primarily a function of information access. Societies with high information access and formally free media can have very high V if the dominant interpretive framework systematically recategorizes structural causation as individual causation. Measuring V therefore requires attention not only to what information is available but to what framework structures its interpretation. A society in which structural suffering is fully documented but universally attributed to individual failure has not solved the visibility problem. It has solved it in the worst possible way: by making the structural character of the suffering invisible within a framework that presents itself as the neutral recognition of individual differences. Meritocratic narrative is the paradigm case of a governing narrative that operates on V without operating on R or B. The structural conditions that produce suffering remain intact. The institutional arrangements that betray the populations they nominally serve remain intact. What changes is the interpretive framework through which those conditions and arrangements are perceived. In the formula's terms: $S_{\text{political}} = (R + B) \times (1 - V)$, and a meritocratic governing narrative can hold effective transformation pressure T low — while R and B remain high — through high V suppression. The political implication is that any analysis of transformation pressure that ignores the governing narrative structure of the society it is analyzing will systematically underestimate actual structural suffering and overestimate the stability of the arrangements producing it. Stigma as Meritocratic Visibility Suppression: The Dispositional Case The meritocratic narrative analyzed in this section operates primarily through the suppression of structural causation — attributing o

utcomes that reflect institutional arrangements to individual character, effort, or desert. A parallel form of visibility suppression operates through stigma against specific populations: the attribution of conditions that reflect biological or psychological variation to moral failure, which generates the same suppression effect through a different mechanism. Where meritocracy suppresses the visibility of structural causes of material outcomes, stigmatizing framings suppress the visibility of populations whose conditions require institutional response — by making the cost of appearing in institutional view (disclosure, treatment-seeking, help-seeking) higher than the cost of remaining invisible. The dispositional harm prevention analysis of section 3.4c documents how this mechanism functions in the specific case of populations with risk dispositions they did not choose and may actively not want to act upon. The structural connection to the meritocratic visibility suppression analyzed here is direct: both framings convert structural facts (institutional arrangements that produce outcomes, biological variation that produces conditions) into moral facts (individual failure, individual culpability) that make those structural facts harder to see and therefore harder to address.

8.5 — Anti-Naïve Universalism: What the Universal Standard Actually Requires

This framework defends a universal evaluative standard: the minimal anchor of structural suffering applied consistently across all cultural, political, and historical contexts. Relativism is not the position — different arrangements are not equally valid depending on perspective. But it is also not the naïve universalism that exports one cultural model as the universal solution and calls resistance to that export moral failure. Naïve universalism makes a specific error: it mistakes the particular arrangements that happen to have produced good outcomes in specific material conditions for the universal principles that those outcomes instantiate.

The institutional forms of liberal democracy, developed in specific historical conditions with specific resource endowments and specific cultural inheritances, are not the universal form of good political organization. They are one set of solutions to the coordination problems that the specific conditions they developed in generated. Applying them to conditions with different material constraints, different institutional histories, and different distributions of power will not reliably produce the outcomes that justify them. It will reliably produce the institutional forms without the material conditions that made those forms functional — which is not universalism but a specific kind of cultural imperialism that calls itself universal. The universal standard that this framework defends is not a set of institutional forms. It is the evaluative criterion: do these arrangements create the conditions under which genuine development is possible, and reduce the structural suffering that forecloses it, given the specific material conditions in which they operate? Two societies can satisfy this criterion through very different institutional arrangements. Two societies can fail it through institutional arrangements that appear superficially similar. The criterion is universal. Its application is always contextual. The consistency requirement that follows from this position is demanding in a direction that both universalists and relativists tend to resist. If the evaluative criterion — do these arrangements create conditions for genuine development and reduce structural suffering — is universal, then it applies symmetrically: to arrangements in societies with which the analyst is in political sympathy as much as to those with which they are not, to the analyst's own society as much as to others, and to historical arrangements of the analyst's own tradition as much as to those of traditions they approach from outside. The failure of nerve that produces naïve universalism — applying the criterion selectively, with rigor to distant others and generosity to near allies — is not a minor methodological inconsistency. It is a form of bad faith that invalidates the universalism it claims. The distinction from cultural relativism is equally important and equally easy to miss. Cultural relativism, in its strong form, holds tha

t evaluative criteria are always internal to the culture that produces them and cannot be applied across cultural boundaries without committing an act of imposition. This position has a genuine insight — that the specific institutional forms through which values are expressed are culturally embedded and cannot be exported wholesale. But it mistakes the insight about forms for a claim about criteria. The fact that the institutional expressions of the commitment to reducing suffering and enabling development differ across cultures does not mean that the commitment itself is culturally particular. The cross-cultural consistency of the minimal anchor (3.4) — that no culture systematically claims that the suffering of its own members is without significance — is evidence that the criterion is not arbitrary even if its expression varies. Applied to the most common objection — that universal standards are instruments of Western cultural imperialism — the framework's response is precise. Western societies have systematically applied universal standards selectively: rigorously to others, generously to themselves and their allies. This is a failure of universalism, not an argument against it. The remedy is consistent application, not the abandonment of the standard. Abandoning the standard in the name of anti-imperialism leaves the populations whose structural suffering is produced by non-Western arrangements without the evaluative resources that consistent universalism would provide. The anti-imperialist who refuses to apply the framework's evaluative criterion to arrangements that produce structural suffering because those arrangements are culturally non-Western has not protected those populations from judgement. They have protected the arrangements from it — which is not the same thing, and not progress. The objection from Western cultural imperialism is the most common but not the only challenge to universalism. A distinct and more analytically interesting challenge comes from non-Western universalisms: comprehensive frameworks for organizing human political community that make universal claims from within non-Western intellectual traditions. The Confucian concept of Tianxia ('All under Heaven') — most recently developed by Zhao Tingyang as a political philosophy for international order — proposes a world institution organized around the relational obligations of all peoples rather than around the rights and sovereignty of states, grounded in the Confucian analysis of legitimate governance as flowing from service to the governed rather than from formal authority. The Islamic concept of the Ummah proposes a community of belonging that crosses ethnic and national lines through shared commitment to the

obligations that the tradition identifies as universally binding. African ubuntu philosophy implies a universalism grounded in relational personhood rather than in individual rights. These are not particularist resistance to universalism — they are competin

g universalisms, each with different foundational commitments about what grounds the universal claim. PMN's response to them is not that they are wrong to make universal claims but that the criterion for evaluating any universal claim — does this arrangement actually reduce structural suffering and create conditions for preserved adaptive response capacity across the full population it applies to? — applies to non-Western universalisms with the same force it applies to Western ones. The Tianxia framework, the Ummah concept, and ubuntu-grounded universalism are all subject to the same question: what institutional mechanisms prevent the universal claim from becoming an instrument of the interests of those who control its interpretation? This is not the Western imperialism objection; it is the custodian problem applied universally. Further Naïve Positions: Progressivism, Secularism, Cosmopolitanism, and Naïve Atheism The naïve positions that undermine analytical clarity are not limited to universalism and passive multiculturalism. Four further variants appear with sufficient frequency in progressive politics to require explicit identification, because each presents itself as the framework's natural ally while importing assumptions the framework cannot support. Naïve progressivism is the assumption that change accumulates directionally — that gains in structural conditions, once achieved, are secured against reversal, and that the arc of history bends toward better arrangements by its own momentum. The evidence against this is contained in 4.3: formal abolition of slavery coexisted with debt bondage; reduced child mortality has been reversed by austerity and conflict; extended political participation has been curtailed when the populations it empowered used it in ways that threatened concentrated interests. The error is not the belief that progress is possible — the framework insists on it. The error is the belief that progress is self-sustaining once achieved, that the conditions which made it possible do not require active maintenance, and that the movement which achieved it can afford to consolidate rather than continue. Naïve progressivism produces movements that mistake a phase of convergence for the destination, stop building the coalitions and institutions that active maintenance requires, and are then surprised when the fragmentation that their success has already prepared arrives. The convergence-fragmentation structure described in 4.3 means that every achieved gain is simultaneously the beginning of the process through which it will be contested. The appropriate response is not despair but the construction of institutions robust enough to manage the contestation. Naïve secularism is the assumption that modernization and material development produce convergence toward secular orientation — that as populations become wealthier, better educated, and more integrated into global networks, religious commitment attenuates and eventually becomes a

residual feature of populations that have not yet completed the transition. The empirical record does not support this. The most consequential mobilizations of the late twentieth and early twenty-first centuries — Iranian revolution, the political rise of evangelical Christianity in the United States, Hindu nationalism in India, Salafi movements across the Muslim world — occurred in populations undergoing rapid modernization, not in populations insulated from it. In many cases, modernization intensified religious commitment by disrupting the meaning infrastructure that preceded it without providing adequate replacement, creating the conditions for movements that offered orientation, community, and cosmic significance in exactly the forms that secular modernity could not. The framework's analysis of meaning infrastructure (5.6) and eschatology (5.5) makes the mechanism legible: secularization in one specific historical context — northern and western Europe under specific material conditions — is not a developmental template. It is a local outcome that cannot be generalized without doing precisely the error that anti-naïve universalism identifies. Naïve cosmopolitanism is the assumption that economic and cultural integration generates solidarity across difference — that as populations become more interconnected, the material basis for nationalist, ethnic, and religious conflict erodes and a genuinely cosmopolitan sensibility develops in its place. The historical record of the period from 1989 to the present is a systematic refutation of this assumption at scale. The most intensive period of global economic integration in human history

produced, alongside the genuine benefits of integration, a backlash of nationalist and identitarian politics whose intensity correlates more closely with the disruption that integration produced than with its benefits. Integration without adequate distribution of its gains, and without the maintenance of the meaning infrastructure that integration disrupts, does not produce world citizens. It produces populations whose existing frameworks of solidarity have been weakened without replacement, available for mobilization by whatever offers the most compelling alternative orientation. The framework's analysis of meaning infrastructure collapse (3.4b, 5.6) explains why: cosmopolitan integration, when it proceeds faster than the reconstruction of meaning infrastructure, produces exactly the conditions that eschatological and nationalist movements are best positioned to exploit. The common structure of all six naïve positions — universalism, passive multiculturalism, progressivism, secularism, cosmopolitanism, and naïve atheism — is the same: each takes a genuine insight and extends it beyond the conditions under which it holds, producing predictions that the evidence consistently fails to support and analytical habits that make movements blind to the dynamics that will undermine them. The framework's relationship to each i

s not rejection but precision: the insight is real, the generalization is not. Universalism is right that evaluative criteria can apply across cultures; wrong that institutional forms can be exported without their material conditions. Passive multiculturalism is right that practices deserve contextual understanding; wrong that understanding requires suspension of evaluation. Naïve progressivism is right that genuine progress is possible; wrong that it is self-sustaining. Naïve secularism is right that material conditions shape religious expression; wrong that they produce secular convergence as a universal outcome. Naïve cosmopolitanism is right that integration can generate solidarity; wrong that it does so automatically without the institutional maintenance that solidarity requires. Naïve atheism is right that the critique of religious institutions is warranted; wrong that the critique is sufficient. Naïve atheism is the assumption that the primary problem is belief itself — that if religious commitment were replaced by rational, evidence-based orientation, the populations currently organized by religious frameworks would make better decisions, build more just institutions, and be less susceptible to manipulation by leaders who exploit sacred legitimation for extractive purposes. The critique of religion that naïve atheism makes is not without warrant: specific religious institutions have historically legitimated oppression, suppressed inquiry, and organized violence in ways that the framework's evaluative criterion straightforwardly condemns. The error is not the critique. It is the assumption that the critique is sufficient — that identifying the failures of religious institutions answers the question of what religious meaning infrastructure does that secular alternatives have not demonstrated they can replace. The historical record of explicitly atheist political projects is instructive on this point. The Soviet state, the Maoist state, and the various revolutionary regimes that programmatically attacked religious institutions did not produce populations organized by rational, evidence-based orientation. They produced populations organized by frameworks that replicated the deep structure of religious meaning infrastructure in secular form: the Party as the infallible interpretive authority, the revolutionary martyrs as sacred exemplars, the foundational texts as the source of correct doctrine, the endpoint of history as the secular eschatology that gave present sacrifice its cosmic significance. The structure — authoritative institution, sacred narrative, martyrology, eschatological horizon — was preserved. The metaphysical content was replaced. The claim to have eliminated the irrational in favor of the rational was itself organized by a framework that functioned, sociologically and psychologically, as religion under a different name. (Durkheim 1912) The naïve atheist position fails to ask the question that the framework requires: what work

is this belief system doing, and what happens to that work when the belief is removed? When the answer is that the belief system is providing meaning infrastructure — organizing solidarity, orienting sacrifice, answering the questions that material analysis leaves open — then removing the belief without providing adequate replacement does not produce rational agents. It produces the disoriented population described in the meaning infrastructure collapse section of 3.4b, available for mobilization by whatever offers the most compelling orientation fastest. The framework does not evaluate religious claims on their metaphysical merits. It evaluates them on what they produce for the populations that hold them — and the consistent

finding is that the productive work of religious meaning infrastructure is not automatically assumed by its secular replacement.

8.6 — Anti-Passive Multiculturalism: The Consistency Standard Applied

Passive multiculturalism is the position that all cultural practices deserve equal respect, that criticism of any cultural practice is an expression of cultural imperialism, and that coexistence without judgment is the appropriate posture toward human diversity. This framework rejects that position — not from hostility to cultural diversity, which it values, but from the application of the same consistency standard it applies to everything else. The consistency standard requires that every cultural practice — including those of dominant cultures, including those that present themselves as universal civilization rather than particular culture — be evaluated by the same criterion: what does this practice produce materially for the people who live within it and around it? Does it create conditions for genuine development and reduce the structural suffering that forecloses it, or does it constrain development and impose unnecessary suffering for the benefit of some at the cost of others? This criterion applies to female genital mutilation and to financial deregulation. It applies to arranged marriage and to the culture of overwork. It applies to religious practices that restrict individual autonomy and to secular practices that commodify human relationships. No practice is immune from evaluation by this standard simply because it is cultural. Passive multiculturalism fails by applying the consistency standard inconsistently. It subjects dominant cultural practices to less scrutiny than minority or non-Western practices, not because the criterion suggests this but because the political context of multicultural discourse makes criticism of minority practices more socially costly than criticism of dominant ones. This asymmetry is not respect for diversity. It is a form of condescension that treats some cultural communities as too fragile to be evaluated by the same standards applied to others. At the same time, the framework resists the use of cultural criticism as cover for the imposition of one cultural model as universal. The test for whether cultural criticism is genuine or instrumental is the same test that applies to noble lies: does the criticism apply the same standard consistently across all cultural contexts, including the one from which the critic is speaking? Or does it apply rigorous scrutiny to other cultures while treating its own cultural arrangements as the natural baseline against which deviation is measured? Genuine cultural evaluation has no natural baseline. It applies the material criterion universally and follows the evidence wherever it leads, including into the arrangements most comfortable to the critic. The practical implication is that this framework will be attacked from two directions simultaneously: by those who see any universal standard as cultural imperialism, and by those who see any acknowledgment of cultural specificity as moral relativism. Both attacks miss what the framework is actually doing. The universal standar

d is real and non-negotiable. Its application is always contextual and empirical. These are not in tension. They are the same position stated from two different angles. The cultural and social dynamics analyzed in this Part operate within a material context that extends beyond any single political unit. The structural conditions that determine what suffering is produced and by whom — resource distribution, geopolitical position, the terms on which a society participates in international exchange — are not set domestically. Part IX examines that international dimension: how the global configuration of power shapes the material constraints within which all domestic arrangements operate, and what a materialist analysis of geopolitics requires that accounts of international relations grounded primarily in norms, values, and formal institutional arrangements — without adequate attention to the material power distributions underneath — systematically miss. What passive multiculturalism gets right: The position it rejects has genuine content that should not be lost in the rejection. The history of universalism as a political tool is a history of using shared standards to impose the preferences of the powerful on the different arrangements of the less powerful — colonial education systems, missionary conversion, development programs that treated their own institutional forms as the universal model that all societies were moving toward. Passive multiculturalism is a response to that history, and the response is not simply wrong. Cultural diversity is genuinely valuable, not only as a source

of different practical solutions to shared human problems but as an expression of the range of what human life can be. The consistency standard applied to cultural diversity does not require treating all cultural practices as equally valuable or all cultural differences as equally important; it requires applying the same evaluative criterion to practices that produce harm regardless of whose cultural tradition they belong to. The harm threshold and its application: The framework's criterion for evaluating cultural practices is not cultural difference itself but practices that produce unnecessary suffering for people who cannot effectively exit the arrangement — who are subject to the practice by virtue of their birth position, community membership, or economic dependence, and who lack the structural standing within the community to contest it effectively. The criterion is the same one applied everywhere: unnecessary suffering, assessed comparatively, contextually, and with attention to who bears the costs. Applied consistently, this criterion generates conclusions that passive multiculturalism systematically avoids: that some practices maintained by historically marginalized communities produce real harm for specific members of those communities, that the protection of cultural integrity cannot be a sufficient response to that harm, and that the analysis of the practice

must be separated from the analysis of the community's right to political self-determination. The separation that the framework requires: Passive multiculturalism conflates three questions that the framework requires to be held separately. First: does this community have the right to political self-determination — to make collective decisions about its own organization without external override? The framework's answer is generally yes, subject to the same constraints that apply to all political arrangements under the minimal anchor. Second: does a specific practice within this community produce unnecessary suffering for members who cannot effectively contest it? This is an empirical question that is not answered by the first question — a community can have legitimate rights to self-determination while maintaining specific practices that produce harm under the minimal anchor criterion. Third: what is the appropriate form of engagement with that practice given the answers to the first two questions? This is a strategic and political question whose answer depends on who is engaging, from what structural position, with what resources and accountability, and what the realistic consequences of different forms of engagement are. Collapsing all three questions into a single answer — either "all practices are legitimate" or "harmful practices must be abolished" — is what produces the analytical failures at both ends of the spectrum. The practical implication is that the framework neither endorses passive multiculturalism's exemption reflex nor the crude universalism that treats cultural difference as a problem to be overcome by the application of standards derived from dominant-culture institutional forms. It requires consistent application of the minimal anchor criterion, honest assessment of who bears the costs of specific practices, attention to the internal diversity of every community regarding those practices, and strategic judgment about what forms of engagement are adequate to the situation — engagement that is accountable to the people affected, not primarily to the analyst's own political commitments. The connection to 8.5's anti-naïve universalism: the framework navigates between two failures. Naïve universalism applies universal standards without accounting for how those standards were produced, whose interests they encode, and what the history of their application has been. Passive multiculturalism refuses universal standards as a cover for political avoidance of the difficult work of evaluating practices that produce harm within communities the analyst identifies with or wishes to protect from criticism. The framework's position — applying the minimal anchor consistently, attending to who bears costs, distinguishing self-determination from exemption from evaluation — is harder than either failure mode and produces more uncomfortable conclusions than either. That is a consequence of taking the consistency standard seriously rather than depl

oying it selectively.

8.7 — Education as Structural Variable

Education has appeared throughout this framework as an instrument — as the mechanism through which visibility suppression (V) is either maintained or reduced, as the site at which becoming is either expanded or foreclosed, as the institution whose capture most directly determines intergenerational structural mobility. This section treats it as a structural variable in its own right: a domain with its own capture dynamics, its own relationship to the floor/becoming architecture, and its own specific failure modes that cannot be fully derived from the general analysis of institutional capture elsewhere in the framework. The structural position of education is peculiar because it operates simultaneously on both levels of the framework's evaluative architecture. At the floor level, access to education is a material condition whose unequal distribution produces structural suffering in the direct sense: resource deprivation (R) when educational access is systematically correlated with parental economic position, institutional betrayal (B) when systems that claim meritocratic distribution reproduce class advantage with greater efficiency than explicit caste, and visibility suppression (V) when educational institutions produce the interpretive frameworks through which structural suffering is attributed to individual failure rather than structural causation. The evidence for these effects is extensive and consistent across contexts: educational access, quality, and outcomes are among the most reliable proxies for structural position, and the intergenerational correlation of educational outcomes is among the most durable mechanisms through which structural positions reproduce themselves. An educational system that claims meritocratic distribution while producing outcomes that track parental advantage is exhibiting the institutional betrayal pattern — the gap between stated function and actual function that the framework identifies as structurally generative of political pressure. At the anti-foreclosure level (becoming), education is the primary institutional mechanism through which becoming — the capacity to develop beyond what survival and material security alone require — is either distributed or concentrated. The distinction matters because the floor and anti-foreclosure criterion problems in education have different structural causes and require different interventions. Inadequate floor provision — the denial of basic literacy, numeracy, and the social knowledge required for civic participation — is primarily a resource distribution problem addressable through funding equity. Foreclosure of becoming is more intractable: it operates through what and how education teaches, not only through who receives it. Educational systems that deliver universal access to a credential factory — that produce the literate and numerate compliance required by labor markets while systematically eliminating the structural visibility, critical capacity, and meaning infrastr

ucture that becoming requires — have solved the floor problem while producing what this framework calls managed dependency at the anti-foreclosure level. The student who graduates with credentials but without the capacity to analyze the structural conditions that shape their situation has not experienced becoming; they have experienced a sophisticated form of preparation for contained functioning within arrangements they cannot name. The capture dynamics specific to education follow from its position as the institution that determines what is knowable, how it is known, and who is authorized to know it — which is to say, as the institution most directly implicated in the production and distribution of epistemic authority (1.9). Educational capture takes three characteristic forms. The first is curriculum capture: the content of what is transmitted reflects the interests and blind spots of those who control curriculum design, producing not only the biased inclusion and exclusion that explicit capture produces, but also the more subtle form that consists in teaching established knowledge as settled rather than as the product of specific conditions, specific actors, and specific exclusions. The knowledge that does not appear in the curriculum — particularly the structural analysis of how the arrangements students are being prepared to enter were produced and are maintained — shapes graduates as effectively as the knowledge that does appear. The second is credentialing capture: as credentials displace competence as the primary signal for resource allocation, educational institutions become captured by the interests of credentialing bodies, employers, and state bureaucracies rather than remaining accountable to the populations whose becoming they claim to serve. The third is what may be called epistemic capture: the systematic use of educational form to make structural conditions appear natural — the same V-suppression mechanism analyzed in 8.4c operating through the formative institution rather

than through mass media. The colonial inheritance problem (9.0b) applies to education with particular force because educational systems are among the most durable colonial institutional legacies. The educational systems of postcolonial states were designed to produce administratively literate populations capable of functioning within colonial institutional structures — not to build the epistemic infrastructure required for self-directed development. The resulting tension between educational forms inherited from colonial administration and the epistemic needs of postcolonial populations is not resolvable by updating curriculum content alone; it requires rethinking what knowledge is being transmitted, in whose language, through whose categories, toward what ends, and with what relationship to the local practices and knowledge systems that colonial education systematically devalued. This is not an argument for cultural isolationism in educational content — the kno

wledge that travels across contexts does so because it addresses problems that arise across contexts. It is an argument that the terms on which that knowledge enters local educational systems matters: whether it enters as universal standard that displaces local knowledge, or as resource that is evaluated alongside it, determines whether education functions as epistemic infrastructure or epistemic displacement. Diagnostic questions for educational analysis: (1) Floor equity — does educational access, quality, and outcome track structural position (parental income, geography, cultural capital)? If yes, the system is producing institutional betrayal regardless of stated meritocratic commitment. (2) Anti-foreclosure function — does the educational system expand structural visibility (the capacity to recognize structural causation of individual situations) or suppress it (the production of frameworks that attribute structural outcomes to individual merit or failure)? This is the V-suppression diagnostic applied to the formative institution. (3) Capture identification — who controls curriculum, credentialing, and assessment? Do those actors' interests align with becoming for the full population, or with specific classes of outcome that serve narrower interests? (4) Meaning infrastructure function — does the educational system provide resources for students to address the questions that biological and social life persistently generates (what is this situation for, what does it demand of me, what connects my trajectory to something larger), or does it produce credential-bearing graduates without the meaning infrastructure required for sustained becoming? (5) Colonial inheritance — does the educational system operate primarily through the knowledge forms, languages, and epistemic standards inherited from colonial administration? If yes, what is the relationship between those inherited forms and the actual epistemic needs of the communities the system serves? (6) Reproductive labor visibility — does the analysis of educational outcomes account for the care work that makes educational functioning possible, and is that care work equitably distributed across gender and class lines?

IX: The International Dimension: Geopolitics and Global Order

9.0 — The Material Basis of International Order

International relations analysis in this framework is not geopolitics as a separate discipline layered onto the domestic analysis. It is the application of the same structural methodology to the arena in which domestic institutional choices are most constrained. The material basis of international order — the geography of water, energy, arable land, and strategic position that determines what choices states can make and at what cost — is not background context for the analysis of power. It is the constraint structure within which all other analysis operates. A domestic transformation program that does not account for the international material constraints on what transformations are feasible in a given configuration is not just incomplete — it is liable to generate recommendations that are coherent in the abstract and impossible in the specific situation where they are to be applied. This section sets the constraint structure that Part IX's institutional and geopolitical analysis operates within. International relations is typically analyzed as a domain of ideas, interests, and power — a domain where states pursue security, where ideology shapes alliances, and where leadership and institutions determine outcomes. All of this is real. But beneath it is a layer of material geography and resource distribution that shapes what is possible in ways that no ideology or institutional design can fully override. Understanding this layer is the prerequisite for understanding why the international system has the structure it does. The planet's resources are not evenly distributed. Freshwater is concentrated in specific river basins and aquifers. Arable land is determined by soil composition, rainfall patterns, and temperature ranges that vary dramatically across geography. Fossil fuels and mineral deposits are geologically contingent — located where geological processes happened to concentrate them, without reference to political boundaries or human need. Rare earth elements, essential for the technology infrastructure of the current era, are distributed in patterns that give a small number of territories extraordinary leverage over the global economy. Non-Western traditions of political realism offer analytical resources for understanding the material basis of international order that the dominant Western IR literature does not exhaust. Kautilya's *Arthashastra* (c. 300 BCE) is the earliest and most systematic treatise on statecraft in any tradition — predating Machiavelli by eighteen centuries — and its analysis of mandala politics (the structural dynamics of alliance and conflict among neighboring states determined by their relative power and interests rather than by ideology or legitimacy) is a form of structural realism that reaches conclusions parallel to Waltz's but grounded in different conditions and derived without the European state system as its reference case. Kautilya's analysis of the state as requiring both material power (*danda*) and legitimacy

(*dharma*) to sustain itself maps directly onto PMN's own framework: power without legitimacy produces brittle arrangements requiring continuous coercion; legitimacy without power cannot maintain the conditions for genuine development. Ibn Khaldun's analysis of dynastic cycles in the *Muqaddimah* provides a non-Western theory of hegemonic succession that antedates Arrighi's long-cycle analysis by six centuries and operates through mechanisms (*asabiyyah* depletion) structurally equivalent to PMN's legitimacy depletion formula. The Chinese *Tianxia* tradition analyzed by Zhao Tingyang offers an alternative architecture for international order organized around concentric relational obligations rather than sovereign equality — relevant to PMN's own long-term horizon of post-national coordination (7.4), but derived from a fundamentally different premise about the proper unit of political community. These distributions are not incidental background conditions. They are structural determinants of what relationships between political units are possible and necessary. A territory without adequate freshwater cannot be fully sovereign in the sense of being genuinely independent of territories that have it. A territory without fossil fuels or the capital to purchase them cannot industrialize on the terms that industrialization has historically required. A territory without arable land sufficient to feed its population cannot maintain political stability without access to territories that have it — whether through trade, colonization, or conquest. The material geography of

resources is one of the most powerful explanations available for why states relate to each other as they do — why certain conflicts recur across different ideological eras, why certain alliances form regardless of the stated values of the parties involved, why certain territories are repeatedly contested regardless of which power controls them. The Nile basin has been a zone of political tension for millennia because freshwater in an arid region is a constraint that no ideology has ever made irrelevant. The Middle East has been a zone of geopolitical contestation in the modern era not primarily because of religious or civilizational difference but because a very large proportion of the world's easily extractable fossil fuels happens to be located there. (Diamond 1997; Yergin 1991) The implication for internationalism is significant and frequently missed. Internationalism is not primarily an ideological project. It is a material necessity produced by the uneven distribution of what different human populations need. Trade is not a preference: it is the mechanism by which populations with complementary resource endowments extract value from the distribution that neither created. International institutions are not idealism made institutional — they are the coordination mechanisms that populations with interdependent resource needs develop to manage that interdependence without constant conflict. The

failure of autarkic projects — the attempt to achieve genuine national self-sufficiency — is not a political failure. It is a material one: the resource distributions that autarky requires simply do not exist for most territories. This reframing changes the analysis of both nationalism and internationalism. Nationalism that claims to be about cultural identity or civilizational distinctiveness is, at the material level, often about controlling resource access — which is why nationalist movements so frequently align with the interests of those who control the resources at issue. Internationalism that presents itself as idealism about global cooperation is, at the material level, often about managing the interdependencies that resource distribution has made unavoidable. Both are more honestly understood through the lens of material geography than through the lens of the values they claim to represent. The specific material variables at the international level: The geographic distribution of water availability, arable land, energy resources, critical mineral deposits, and strategic position are the foundational material variables from which international power relationships are built. These variables are not equally distributed across political units, and their distribution does not correspond to political boundaries in ways that reflect historical justice or functional efficiency — they reflect the geological and ecological history of the planet, modified by the historical process through which political boundaries were established through force. The political consequences: states whose territorial boundaries encompass high concentrations of strategic resources face different constraints and opportunities than those that do not, independently of the quality of their institutional arrangements or the competence of their governing coalitions. States that lack domestic resource security must acquire it through trade relationships, which creates structural dependencies that other states can leverage. The analytical obligation: international political analysis that treats the geographic and material distribution of resources as background context rather than as primary causal variables will consistently misread why specific states pursue specific policies, why specific alliances form and fracture, and why specific international institutional arrangements take the forms they do. The material geography of international relations is not a complication to be noted and set aside; it is the foundation from which the institutional analysis must proceed. Capture at the supranational level: The capture sequence analyzed in 7.3b and 7.3c-i applies at the international level with a structural modification: the actors doing the capturing are not individuals or corporations but states, and the institutions being captured are the multilateral architecture designed to manage the interdependencies that resource distribution creates. The modification matters because sta

tes have greater resources and longer time horizons than individual actors, and because the capture of supranational institutions is harder to detect and harder to reverse than domestic institutional capture. The information asymmetry between the states doing the capturing and the populations nominally served by the captured institution is greater; the accountability mechanisms are weaker; and the costs of capture are externalized onto states that lack the leverage to contest it. The five-stage sequence at the supranational level: Stage one is proportional influence asymmetry: states that contribute more resources to multilateral

institutions — in financing, technical expertise, personnel, or military capacity — acquire disproportionate access to institutional decision-making. This is often formalized in weighted voting structures, but it operates informally even where formal structures are nominally equal. Stage two is preference institutionalization: the major contributors shape the institution's operational procedures, standard-setting processes, and conditionality frameworks to favor institutional arrangements that serve their structural interests. Stage three is personnel alignment: the professional culture of international institutions systematically favors candidates from the most influential states and from the academic and professional traditions those states dominate. Stage four is objective redefinition and output reorientation: the institution's success metrics shift toward indicators that the preferred arrangements satisfy by design — "good governance," "business climate," "debt sustainability" — rather than indicators of whether the populations served are above the structural suffering threshold. Stage five is consolidation: the captured institution actively deploys its legitimacy against alternatives, using the credibility built under its pre-capture function to foreclose institutional experimentation and development paths that the capturing states find threatening. Why the domestic capture analysis requires extension: The domestic analysis in 7.3 treats the custodian problem as fundamentally a design problem: distribute anchor-holding, create heterogeneous interests among custodians, maintain external accountability with genuine enforcement capacity. At the supranational level, all three mechanisms are significantly weaker. Distribution of anchor-holding is constrained by the power asymmetry between states: institutions formally governed by one-state-one-vote are practically governed by the states that can credibly exit or defund them. Heterogeneous interests among custodians are systematically reduced by the professional socialization that international institutional careers produce — a convergence toward the analytical frameworks and policy preferences of the dominant contributors that operates through selection, training, and incentive alignment rather than through deliberate homogenization. External accountability with genuine

enforcement capacity is structurally unavailable because the populations most harmed by supranational institutional capture — populations in lower-income states subject to the conditionality and standard-setting of captured institutions — have the least capacity to organize the cross-border coalitions that accountability would require. The implication for counter-power at the international level: Counter-power at the supranational level cannot follow the domestic model of organized constituencies using accountability mechanisms to constrain institutional behavior. The accountability mechanisms are too weak and the constituencies too dispersed. The most effective counter-power strategies have operated through different mechanisms: the construction of alternative institutions that reduce the captured institution's monopoly on legitimacy (the BRICS New Development Bank as partial alternative to IMF/World Bank conditionality), the production of independent analytical capacity that challenges the epistemic authority the captured institution deploys (heterodox development economics producing evidence the institution's own frameworks cannot generate), and the organized refusal of conditionality that exceeds what the affected population can absorb without structural suffering that exceeds any efficiency gain the conditionality could produce. None of these strategies directly reforms the captured institution; they erode the conditions under which its capture remains politically stable. Diagnostic: when assessing a supranational institution, apply the five-stage sequence and ask which stage the capture has reached. Then apply the three accountability mechanism tests — distribution, heterogeneity, enforcement — and assess how much weaker each is relative to what the domestic analysis would require. The gap between the domestic accountability standard and what is actually available at the supranational level is the measure of the structural deficit in international institutional design.

9.0b — The Historical Divergence: How the Distribution Was Produced

The material geography of resources explains the structure of current international arrangements. It does not by itself explain the specific form of divergence between territories that currently have the capacity for self-determination and those that do not. That divergence was produced by a historical process that is not

incidental to the current distribution but generative of it. The territories that achieved industrial development in the nineteenth century did not do so in isolation. They did so partly through the extraction of material resources — labor, land, commodities — from territories that were prevented by force from developing independent industrial capacity. The systematic transfer of value from colonial peripheries to industrial centers was not a side effect of development. It was a mechanism of development: it provided the capital, the markets, and the material inputs that made industrial accumulation possible at the pace it occurred. The institutional arrangements that governed this transfer — slavery, indentured labor, colonial land seizure, unequal treaties, forced de-industrialization — were not temporary aberrations that left no lasting trace. They structured the material starting positions from which every subsequent development path has proceeded. (Rodney 1972; Wallerstein 1974) The implication for the framework's analysis of parallel development is significant. A framework that advocates parallel development — each political unit working toward self-sufficient and materially stable institutional arrangements — must account for the fact that the starting positions from which development proceeds are not parallel. They are the product of a structured historical divergence across the *longue durée* (Braudel 1979) in which some territories developed through increasing-returns industrial emulation while locking others into extractive positions (Reinert 2007). The remedies appropriate to this situation are not identical to the remedies appropriate to a world of equivalent starting positions. The framework does not prescribe what those remedies are — that is the work of the doctrine applied to specific conditions. But it insists that the analysis of any specific territory's development path begin from its actual historical starting position, not from a baseline of formal equivalence that did not materially exist. The historical process through which the current distribution was produced matters for the analysis because it determines the moral and political status of existing arrangements. Colonial extraction was not simply an unequal exchange that happened to leave some territories with less than others — it was a systematic process of resource transfer, institutional destruction, and structural dependency creation that was designed to be self-reinforcing. The territories that were most thoroughly incorporated into colonial extraction economies did not emerge from that process as blank slates that happened to start behind. They emer

ged with institutional architectures shaped by extraction rather than development, with human capital patterns determined by what colonial administration required rather than what local populations needed, and with debt and trade relationships that encoded dependency into their economic structure. The relevance for current analysis is not primarily historical justice, though that question is real. The relevance is that the material constraints currently facing former colonial territories — and the structural advantages currently held by former colonial powers — are not the product of differential initial endowments or differential institutional choices made on a level field. They are the accumulated product of a historical process that actively produced the current distribution. A geopolitical analysis that treats the current distribution as the starting point for analysis, without accounting for how it was produced and the structural mechanisms that reproduce it, will consistently misattribute the causes of the persistent gap between high-income and low-income territories and will consistently generate recommendations that are inadequate to the structural conditions they are supposed to address. This does not imply that historical causation is the only relevant factor — current institutional capacity, demographic dynamics, and geopolitical positioning all operate with some independence from the historical baseline. It implies that the historical baseline is itself a structural variable, not a natural condition, and that any framework that treats it as a natural condition has already made a political choice: it has chosen to treat the outcomes of a particular historical process as the legitimate starting point for evaluation. That is a choice the framework is not willing to make implicitly. The mechanism of divergence: The extraction relationship through which early industrial powers funded their own development while degrading the productive capacity of their territories of operation was not primarily a military relationship, though military enforcement was essential to its maintenance. It was primarily an institutional relationship: the creation and maintenance of legal, commercial, and administrative arrangements that structured the terms of exchange in ways that systematically transferred surplus from the

periphery to the center. These arrangements did not require continuous coercion to function once established; they were embedded in property rights, contract law, financial infrastructure, and the professional norms of the administrative classes that operated them. The divergence between the territories that industrialized and those that did not is not adequately explained by any account that treats the industrializing territories as having discovered or developed an advantage in isolation. The current political relevance: the institutional arrangements that produced the divergence are not fully historical. The financial architecture, intell

ectual property frameworks, trade rules, and monetary systems that govern current international economic relations were not designed from a neutral starting point — they were designed by and for the states that had accumulated the most institutional leverage from the prior extraction relationship. Understanding why specific multilateral institutions take the forms they do, why specific trade rules consistently advantage specific types of economy, and why specific financial arrangements consistently transfer resources from lower-income to higher-income economies requires tracing the institutional continuity between the arrangements of the extraction period and those currently in operation. Connection to 9.0 and 9.1: section 9.0 establishes the material basis of current international order. Section 9.0b establishes how that basis was historically produced. Section 9.1 establishes the realist-versus-direction tension that follows from both. The sequence matters: without 9.0b, section 9.1's insistence on holding the direction alongside the realist assessment could appear as idealism unmoored from material conditions. With 9.0b, the direction is specified precisely because the current material basis was historically produced through processes that the minimal anchor criterion must assess as having generated structural suffering at scale.

9.1 — Realism and Direction

Two positions on international relations are held in genuine and permanent tension. The first is that the international system currently operates in ways broadly consistent with structural realism: states pursue power and security, norms are routinely violated when they conflict with state interest, and the gap between declared values and actual behavior is a permanent feature of international politics. (Waltz 1979) The second is that the long-term trajectory should move toward stronger international cooperation, reduced nationalist conflict, and institutional arrangements more directly accountable to human welfare across existing boundaries. Marx expected the international dimension to dissolve as class solidarity superseded national identity — workers of all countries uniting against capital that was itself becoming international. The expectation was wrong, and the reasons it was wrong are instructive. National identity is not merely a product of ruling-class ideology that will dissolve when workers develop class consciousness. It is grounded in real material conditions — shared language, shared history, shared vulnerability — that persist independently of the economic arrangements that historical materialism identified as primary. Nations are constructed around genuine material facts, not invented from nothing; but the specific political form they take — the imagined community that binds strangers into solidarity — is itself a historical product that can take many forms. (Anderson 1983) The framework takes these conditions seriously as material variables rather than treating them as obstacles to be overcome by correct ideological development. Taking national identity seriously as a material variable rather than an ideological distortion has specific analytical consequences. It means that international solidarity — the building of cross-national coalitions around shared material interests — must work with rather than against the material reality of national attachment. The failure mode of internationalist politics that treats national identity as false consciousness is not merely rhetorical. It is strategic: movements that ask populations to subordinate genuine attachments to abstract international solidarity consistently fail to build the durable coalitions that transformation requires, because they are asking people to set aside something real in favor of something they cannot yet feel. The framework's position is that genuine internationalism — the kind that can actually build coordination capacity at the international level — must be grounded in the recognition that different populations have different material interests that are not dissolved by shared class position. Workers in high-income countries

and workers in low-income countries share an interest in the reduction of structural suffering and the expansion of conditions for becoming. They do not share identical immediate interests in the specific distribution of production, wages, or

environmental burden. An international politics that pretends otherwise will fracture at exactly the point where the concrete trade-offs become unavoidable. An international politics that starts from the material divergences and asks what coordination is possible despite them has a realistic foundation. Realism in the international dimension means accepting that the nation-state is not going to wither away in the analytically relevant timeframe, and that the existing international institutional architecture — the United Nations system, the Bretton Woods institutions, the World Trade Organization and its successors — is both inadequate to the problems it faces and the only available foundation for building something more adequate. The choice is not between the existing system and a better one. It is between the existing system, incrementally reformed and extended, and the collapse of coordination capacity into purely competitive state relations. The direction is clear. The instruments are imperfect. The alternative to working with imperfect instruments is not better instruments — it is no instruments. (Held 1995) The practical implication is that progressive international politics cannot afford the luxury of refusing engagement with existing institutions on grounds of their inadequacy. Every institution that currently manages some dimension of international coordination — however partially, however captured by specific interests — represents coordination capacity that took decades to build and would take decades to rebuild if lost. The analytical task is to identify which elements of existing international architecture are reformable, which are captured beyond recovery, and which create the scaffolding for the more adequate structures that the permanent conditions require. Refusing this task in favor of a clean break is not radicalism. It is the abdication of the only leverage that currently exists. Why the tension cannot be dissolved by choosing one pole: A purely realist position — one that accepts the current distribution of power as the operative constraint and treats the normative direction as an expression of preference with no analytical consequence — cannot explain why the normative direction matters for strategic analysis. But it does matter: actors who orient their strategy toward the long-term direction of stronger international accountability achieve different outcomes over sustained periods than actors who optimize only for short-term power accumulation, because the institutional architecture is moving in ways that reward coordination-building capacity. A purely normative position — one that treats the current realist dynamics as obstacles to be overcome through moral argument or institutional design — consistently produces naive engagement with actors who will use the normative framework strategically without being constrained by it. The tension is productive precisely because neither pole is adequate alone. The analytical translation

n: the framework holds realism as the diagnosis of current conditions and direction as the criterion for evaluating whether current conditions are moving toward better or worse configurations. This is not a contradiction but a methodological choice about the relationship between description and evaluation. The realist description of current dynamics is taken seriously precisely because the direction criterion makes the gap between current conditions and the direction visible — and the gap is the measure of how much structural transformation is required and how durable the current configuration's costs are. The strategic implication: actors building counter-power in the international sphere should simultaneously operate within realist constraints — taking seriously the material interests of states and the limits of what international norms can compel — and build toward the institutional architecture that would move the system in the direction of the normative criterion. These are not contradictory strategies but different time horizons of the same analysis.

9.2 — On Power Concentration

The concentration of power in a single dominant actor produces conditions that are systematically dangerous regardless of the ideological commitments of the actor doing the concentrating. This is an argument from the

structure of power, not from sympathy for any particular counter-hegemonic position. It applies symmetrically: it is equally an argument against concentration by actors whose stated values align with the long-term direction of this framework. The appropriate check on any concentration of power is not a different concentration of power. It is the development of institutional arrangements that make concentration structurally difficult regardless of who is attempting it. The mechanism through which concentrated power becomes dangerous is the same at the international level as at any other: actors with concentrated power develop incentive structures that systematically favor the maintenance of concentration. A dominant power that controls the terms of international trade, the architecture of international financial institutions, and the deployment of military force develops material interests in maintaining the rules that sustain those arrangements — regardless of whether those rules serve the populations they claim to serve. The universalist ideology that hegemonic powers adopt — the claim that their arrangements are not merely beneficial to themselves but are the correct arrangements for all — is the international expression of the noble lie analyzed in 7.5: a narrative that serves the interests of those who maintain it while claiming to serve everyone. This analysis does not imply that all hegemonic arrangements are equivalently harmful, or that the appropriate response to any concentration is immediate disruption regardless of what would replace it. Some concentrated power arrangements are significantly less harmful than the alternatives available under current conditions. The analysis requires evaluating the specific consequences of specific arrangements against specific alternatives — not applying an abstract anti-hegemony principle that exempts the analysis from comparison with real options. What it does imply is that no concentration of power is evaluated favorably simply because it is currently in place or because its stated values align with the minimal anchor. The test is always material: what does this arrangement actually produce for the populations it affects? The international distribution of power is structurally different from domestic distributions in one analytically crucial respect: there is no institutional actor with a monopoly on legitimate coercive force capable of enforcing constraints on the most powerful units. This is the structural condition that makes the international system anarchic in the technical sense — not that it is disordered or violent as its normal state, but that the constraints on state behavior are enforced through the actions of other states rather than through a superior authority. In that co

ndition, power concentration at the international level operates without the internal correction mechanisms that domestic institutional design can in principle provide. The tendency toward power concentration in international systems follows from the same structural logic that produces it domestically: actors with more resources can acquire better information, absorb more risk, make longer-term investments, and recover from setbacks more readily than actors with fewer resources. These advantages compound across iterations. The result, absent countervailing institutional constraints, is the progressive concentration of capability in a small number of units and the progressive marginalization of the remainder — not through malice, but through the structural advantages that accrue to those who can operate at the largest scale. Hegemonic stability theory captures part of this dynamic but systematically misreads its evaluative implications: the stability that hegemonic concentration produces is stability for the hegemon, which is not the same as stability for the international system or for the populations that system governs. The framework's treatment of national and communal identity has analyzed it primarily as terrain to be navigated — as a material force with real consequences that can be occupied by arrangements serving or opposing the criterion. This is accurate but incomplete. It does not address the question that those who hold national or communal identity as a genuine positive commitment will ask: can such identity be more than terrain? Can it constitute a legitimate claim on institutional arrangements, not merely a force to be managed? The framework's answer is conditional rather than categorical. National and communal identity is a legitimate positive claim when it satisfies two conditions. First, when it is organized around the shared institutional memory, cultural inheritance, and cooperative infrastructure that a population has actually built and maintains — the civic nationalist claim that democratic self-governance requires a demos with shared commitments and shared stakes. This is not ethnic nationalism; it is the recognition that institutional arrangements capable of managing the permanent conditions require populations who trust each other enough to maintain them across cycles of disagreement.

Second, when the identity claim is consistent with the criterion applied universally: a national community that claims the right to maintain its institutional arrangements against external imposition cannot simultaneously deny equivalent standing to other communities. The test is the consistency standard of 8.3 applied to identity: the same standing that one community claims for its self-determination must be extended to others on equivalent grounds. Where national identity fails this test — where it is organized around the exclusion of populations within the claimed territory, or around the appropriation of resources from other communities,

or around the maintenance of arrangements that produce structural suffering for those defined as outside the national community — it has departed from the legitimate positive claim and become the manufactured fracture line that the framework analyzes as an obstacle to the criterion. The distinction is not sentimental. It is structural: identity that builds the cooperative substrate that makes collective management of shared conditions possible is an asset the framework can engage with directly; identity that maintains itself by defining its inside against an outside it can dominate is doing something else while calling it the same name. The minimal anchor applied to the international distribution of power produces a specific diagnostic question: does the current concentration of capability in a small number of units systematically produce arrangements that serve the material conditions of the populations most affected by those arrangements? The answer is not structurally determined — it depends on what the concentrated actors do with their capability and what accountability mechanisms constrain them. But the structural expectation, consistent with the custodian problem analyzed in 7.3, is that concentration without accountability will drift toward serving the interests of those who hold the concentrated position over the interests of those who are subject to its consequences. The argument from structure applies to the current system's dominant actor as fully as to any historical case. This is the self-consistency requirement that Part VIII establishes: a framework that applies its analysis to some concentrations of power but exempts others based on ideological affinity has abandoned the analytical commitment and adopted an ideological one. The United States' role as the system's most powerful actor produces the same structural consequences as any other dominant actor: systematic externalization of costs onto actors without countervailing leverage, capture of the international institutional architecture in ways that serve the dominant actor's interests rather than the interests the institutions claim to serve, and the systematic devaluation of the constraints that those institutions impose when those constraints become inconvenient. These are structural consequences of the position, not attributions of malice. The implication for counter-power strategy: Arguments against power concentration at the international level must be held consistently across cases — applied to the concentration that currently exists and to any emerging concentration regardless of which actor is concentrating. The strategic temptation is to support counter-concentration — the rise of alternative centers of power — as automatically beneficial because it reduces the current dominant actor's leverage. This temptation should be resisted: counter-concentration that produces a different configuration of concentrated power rather than distributed institutional accountability serv

es no one's long-term interest except the new concentrating actor. The goal is not redistribution of concentrated power but the institutional architecture that constrains any concentration. Diagnostic: when assessing a specific power concentration at the international level, ask (1) what costs are being externalized onto actors without countervailing leverage, (2) what international institutional arrangements has the concentration shaped to serve its interests, and (3) what constraints has the concentration systematically avoided that the same institutional framework imposes on others. These three questions specify the material consequences of the concentration independently of the ideological position of the concentrating actor.

9.2b — Technology as Redistributor of Material Power

Technology does not merely change what is possible. It changes who has power, and it does so in ways that are not controlled by the political arrangements that existed before the technology appeared. This is one of

the most reliable patterns in the historical record and one of the most systematically underestimated by frameworks that treat technology as a tool deployed by political actors rather than as a structural force that reshapes the terrain on which those actors operate. The printing press did not merely make books cheaper. It broke the institutional monopoly on the production and distribution of authorized knowledge that had sustained a particular configuration of religious and political authority for centuries. (Eisenstein 1980) The consequences were not controlled by the actors who initially adopted the technology. They were determined by the structural logic of what the technology made possible — the mass production and distribution of text without institutional gatekeeping — and the political consequences followed from that structural logic rather than from anyone's intentions. Industrial technology did not merely increase productive capacity. It changed the relationship between labor and capital in ways that generated new forms of political organization, new forms of class conflict, and new demands on the state that the pre-industrial institutional architecture was not designed to manage. The political upheavals of the nineteenth and early twentieth centuries are not explicable primarily as ideological events. They are the political consequences of a material transformation in the structure of production that generated new configurations of power and new forms of suffering before institutions capable of managing them had developed. The current technological transition — the digitization of information production and distribution, the automation of cognitive as well as physical labor, the development of artificial intelligence systems capable of performing tasks that previously required human expertise — is a structural transformation of comparable magnitude. Its political consequences will be determined primarily by the structural logic of what the technology makes possible, not by the stated intentions of those who develop or deploy it. Frameworks that analyze the current political moment without accounting for this structural transformation are analyzing a world that is in the process of disappearing. For geopolitics specifically, the technology distribution matters as much as the resource distribution. Control over the technological infrastructure of the current era — semiconductor fabrication, AI development capacity, digital communication networks, satellite systems — represents a form of power that intersects with and in some domains supersedes the resource-based power of the previous era. The geography of technological capacity is as consequential as the geography of fossil fuels was in the twentieth century, and its distribut

ion is equally contingent and contested. The asymmetric capture dynamic: New technologies do not distribute their power-redistributing effects evenly across the actors who engage with them. The actors best positioned to capture the power that new technologies make available are consistently those with the greatest prior accumulation of the resources the technology requires — capital to invest in infrastructure, institutional access to regulatory processes that govern deployment, information networks that make early adoption strategically rational, and organizational capacity to absorb and operationalize the technology before its broader implications are legible to other actors. This means that technological transitions consistently advantage those already advantaged, not because technology is neutral and its effects are purely contingent, but because the transition period is systematically biased toward actors with the prior resource base to move first and fastest. The geopolitical redistribution pattern: At the level of international order, technological transitions have been the primary mechanism through which geopolitical power shifts have occurred historically — more consequential than military conquest, ideological conversion, or deliberate institutional reform. The Industrial Revolution did not merely change what Britain could produce; it changed the terms on which Britain engaged with every other polity in the world by generating a material advantage that translated into military, commercial, and institutional leverage across every domain of international interaction. The semiconductor transition that began in the mid-twentieth century similarly reorganized which states had leverage over which others, which populations had access to which productivity gains, and which institutional configurations were viable at all, at a speed that made deliberate policy responses consistently lag the underlying structural shift. The current AI transition is following the same pattern: the question is not whether it will redistribute power between states, firms, and populations but which actors have the prior resource base to capture its power-redistributing effects and which do not. The infrastructure ownership

problem: The specific mechanism through which technology redistributes power at the international level is control of infrastructure. Technologies whose effects are mediated through specific physical or institutional infrastructure — submarine cables, satellite networks, financial clearing systems, platform architectures — give leverage over other actors' activity to whoever controls that infrastructure, regardless of whether those other actors are themselves sophisticated users of the technology. Dependence on infrastructure controlled by others is a structural vulnerability that is not eliminated by becoming technically proficient at using that infrastructure; it is only eliminated by either developing independent infrastructure or achieving suff

icient leverage in other domains to make infrastructure controllers' costs of exclusion prohibitive. Neither path is available to most actors, which means technological transitions consistently produce durable structural dependencies at the international level that outlast the specific transition that created them. The visibility suppression dimension: Technological power redistribution is systematically underweighted in political analysis for a structural reason connected to the visibility suppression variable analyzed in the formula architecture: the mechanisms through which technology redistributes power operate at levels of abstraction that are less politically visible than the outcomes they produce. The visible outcome is a wage gap, a trade imbalance, a military defeat, or an institutional failure. The structural cause — control of a specific technological leverage point — is several causal steps removed and requires technical knowledge to identify that most political actors do not hold. This gap between cause and visible effect is not accidental; actors who benefit from the power redistribution that technology enables have interests in maintaining the abstraction that obscures the causal mechanism, and the investment in technical complexity that produces genuine obstacles to public understanding also serves as a political shield. The framework's analytical response: any assessment of the power consequences of a major technological transition must explicitly map the infrastructure ownership question — who controls the physical and institutional infrastructure through which the technology's effects are mediated — before drawing conclusions about who will benefit and who will bear costs. Assessments that stop at the level of access — which actors can use the technology — consistently miss the structural question of who controls the terms on which others can use it, and under what conditions that control can be exercised to extract compliance or impose costs. Connection to 9.2c and 16.0: the technology-as-constraint-modifier framing is the general case of which the geopolitical redistribution of material power is a specific application. The same analysis that applies to international order applies within domestic polities — the question of who controls infrastructure, who captures first-mover advantage during technological transitions, and who bears the costs of dependence on others' infrastructure is structurally identical whether the unit of analysis is nation-states or populations within a single political economy.

9.2c — Artificial Intelligence and the New Architecture of Institutional Power

Artificial intelligence introduces a structural transformation qualitatively different from previous technological redistributions of power, and that difference matters for how institutions respond. Previous technological transitions — printing, industrialization, electrification — changed the scale and speed of existing activities while leaving the basic architecture of institutional oversight largely addressable: a factory can be inspected, a press can be regulated, a railroad can be taxed. AI systems operating at scale introduce three structural features that make the standard institutional response inadequate in ways that require analysis at the framework level, not only the policy level. The first is the conversion of data into a form of capital that compounds differently from physical or financial capital. Data accumulated by a platform becomes more valuable as it accumulates — it trains better models, which attract more users, which generate more data. This compounding dynamic creates structural barriers to entry that are not addressable through conventional antitrust logic, which was designed to manage concentrations of productive capacity rather than concentrations of informational advantage. The result is a form of institutional capture that does not require corrupting a regulator: it operates by ensuring that the entities with the analytical capacity to understand what is happening are the same entities whose interests are served by the arrangement remaining as it is. (Zuboff

2019) The second is jurisdictional mismatch. Platform infrastructure operates across national jurisdictions simultaneously, while regulatory authority remains national. This is not merely an administrative inconvenience — it is a structural feature of how AI-based platforms accumulate power: by operating at a scale and across boundaries that no single regulatory body can address. The mismatch is not accidental; it is a material advantage that the architecture of platform capitalism has, in several cases, deliberately maintained. The framework's analysis of the custodian problem applies here with additional force: the institutions best positioned to regulate AI development are also the institutions most exposed to capture by the actors being regulated, because those actors command the informational and financial resources to shape the regulatory environment. The third is the displacement of labor at a pace that may exceed the capacity of existing institutions to manage the transition. Industrial automation displaced physical labor over generations — a pace that, while deeply disruptive, allowed for partial institutional adaptation (labor law, social insurance, public education). AI systems capable of performing cognitive and communicative labor — drafting, analysis, diagnosis, translation, code generation — are displacing this category of work on a timescale measured in years rather than generations. The institutional architectures designed to manage labor transiti

ons were not built for this pace. The eight diagnostic questions apply directly: who bears the costs of displacement and who captures the productivity gains? Does the arrangement reproduce the cooperative substrate — the employed, economically participant population — without which the institutions that manage everything else lose their material basis? Whether specific policy responses (universal basic income, shortened working hours, public ownership of AI infrastructure, mandatory profit-sharing) are adequate depends on material conditions in specific contexts and is a doctrinal question. The framework identifies what those responses must address: the decoupling of productivity from employment, and the concentration of the gains from that decoupling in entities that are structurally unlikely to redistribute them without institutional pressure. The most consequential scenario — and the one that most directly tests the framework's evaluative commitments — is the concentration of advanced AI capability in a small number of actors operating at global scale without adequate accountability structures. This is not a hypothetical. It describes the current trajectory of AI development, in which the capital requirements for training frontier models, the data advantages of existing platforms, and the network effects of deployment have produced a configuration in which a handful of private entities and one or two states control infrastructure with civilizational consequences. The custodian problem at this scale has no historical precedent: previous concentrations of power with comparable reach — empires, hegemonic states, global religious institutions — operated through mechanisms that were at least in principle addressable through the coordination of other actors. AI capability concentrated in entities that are not states and not accountable to the populations they affect represents a form of institutional capture at a level above the institutions designed to prevent it. The framework does not determine whether AI development is good or bad, or whether any specific regulatory response is adequate. The same technology that concentrates power can also reduce structural suffering: medical diagnosis at scale, resource optimization under scarcity, expanded access to knowledge for populations previously excluded by cost or geography. These are real possibilities. The framework's questions are designed to evaluate specific arrangements, not to prejudge the technology — to ask under what conditions those possibilities are realized, and whose interests are served when they are not. Who controls this technology, who bears its costs, who captures its benefits, what happens to the cooperative substrate when the technology displaces the labor that previously sustained it, and what institutional arrangements are capable of holding this concentration of power accountable? These are not AI-specific questions. They are the standard diagnostic questions applied to a specif

ic material transformation of unusual magnitude and pace. The specific feature that distinguishes AI from prior institutional power technologies is the simultaneity of its effects across jurisdictional and organizational boundaries. A factory can be regulated nationally; a media institution can be licensed domestically; a financial institution can be subject to capital requirements within a jurisdiction. AI systems deployed across multiple jurisdictions simultaneously are subject to the regulatory jurisdiction of each only for the specific

consequences that occur within that jurisdiction — but their institutional effects are produced by the system as a whole, which means that no single jurisdictional regulator can address the systemic institutional consequences of deployment. This is not a technical problem solvable by better national regulation; it is a structural problem that follows from the mismatch between the jurisdictional architecture of regulatory institutions and the cross-jurisdictional architecture of the systems being regulated. The counter-power implication: The mismatch between regulatory jurisdiction and AI system architecture means that counter-power organized at the national level cannot address the institutional consequences of AI at the speed or scale at which those consequences are developing. This creates an urgency for the development of the supranational institutional architecture that section 9.1's direction criterion identifies as the long-term goal — not eventually, as a distant aspiration, but as a near-term strategic necessity for addressing the most significant current driver of institutional power concentration. The AI governance question is therefore not a specialized technical policy question but a test case for whether the international institutional development that the framework's direction criterion requires can be organized fast enough to be consequential.

9.3 — Multipolarity as Transitional Arrangement

Multipolarity — the distribution of effective power across multiple actors with different interests — is preferable to unipolarity under current conditions for the same reasons that distributed institutional design is preferable to concentrated institutional design: it reduces the systemic risk of capture, reduces the cost of mistakes, and preserves the competition that makes institutional accountability possible. These are material arguments. Multipolarity is a transitional arrangement, not a terminal one. The long-term goal is not permanent state competition but the development of institutional arrangements that can coordinate human activity at global scale more effectively than state competition allows. The case for multipolarity under current conditions is not a case for nationalism or for the permanent primacy of state sovereignty. It is a case for maintaining the competition between arrangements that prevents any single configuration from becoming entrenched before the material conditions for more coordinated alternatives are in place. Multiple centers of power with genuinely different interests produce more honest empirical tests of which institutional arrangements actually work under which conditions. A world with one dominant model has no external check on that model's failures — failures are attributed to implementation rather than design, and revision is blocked by the interests that have accumulated around the dominant arrangement. A world with multiple centers produces comparative evidence, even when that evidence is messy and politically contested. The transition from multipolarity to more coordinated arrangements requires building the material conditions for coordination before attempting the transition. These conditions include: sufficient mutual knowledge of material interests and constraints to make negotiated agreements possible; sufficient institutional capacity within each major actor to implement commitments without requiring constant external enforcement; and sufficient shared interest in outcomes that no current actor can achieve unilaterally — climate stability, pandemic control, nuclear risk reduction — to make the costs of coordination acceptable. None of these conditions currently obtains fully. The work of building them is the work of the transitional period. The multipolarity that is emerging in the current period is structurally different from the multipolarity that preceded the 20th century's bipolar arrangements. It is not a return to a historical baseline but a novel configuration produced by the combination of: the diffusion of industrial and technological capability to a larger number of significant units; the development of nuclear deterrence that has changed the cost structure of direct great-power conflict; the institutional architecture of the post-1945 period, which has created formal constraints on unilateral action even as those constraints are selectively applied; and the emergence of genuinely global i

nterdependencies — in supply chains, financial systems, and ecological processes — that create forms of mutual vulnerability that have no close historical precedent. Whether this multipolarity is a transitional arrangement or a stable endpoint is an open analytical question that the framework cannot resolve with

current evidence. The argument that it is transitional rests on the structural observation that multipolar systems generate coordination failures that impose costs on all participants, which creates incentives for arrangements that resolve the coordination problem — historically through hegemonic consolidation or through institutional frameworks that constrain the behavior of all participants. The argument that it is more stable than previous multipolar periods rests on the nuclear deterrence effect and the interdependency structures that raise the cost of the consolidation moves that historically terminated multipolar periods. Both arguments have empirical support; neither is conclusive. What the framework can say with more confidence is that the content of any institutional resolution — whether the current multipolar transition produces hegemonic consolidation, genuine institutional reform, or managed fragmentation — will be determined by the same variables that determine domestic institutional outcomes: which actors can organize, what narrative frameworks they can deploy, what elite fractures create windows, and whether the counter-power capacity of actors with interests in more accountable arrangements is sufficient to shape the terms of the transition before they are set by actors with interests in reproducing the current distribution under new organizational forms. The specific mechanisms through which multipolarity serves the transition: Competition between multiple centers of institutional power creates conditions under which populations with different governance arrangements can observe the consequences of different choices and select among them — not in the abstract but through the evidence of what specific arrangements produce for specific populations. Unipolar arrangements eliminate this comparative evidence by extending a single institutional model without the competitive pressure that generates evidence of its failure modes. Multipolarity also reduces the cost of institutional experimentation: configurations that perform poorly can be revised more rapidly when their populations have alternative institutional relationships available than when the single dominant power has no institutional interest in acknowledging the experimentation's failure. The costs of multipolarity — coordination failures, races to the bottom in regulatory standards, conflict between competing centers — are real and must be managed; they do not outweigh the structural benefits of distributed power at the current moment in the transition. The transition endpoint: The endpoint of the transitional arrangement is not stable multipolarity. Stable multipolari

ty — multiple centers of power in permanent competition — reproduces at the global level the structural dynamics that produce systemic suffering at the domestic level: the externalization of costs onto populations without leverage, the coordination failures that prevent management of global commons, and the arms competition that consumes resources that would otherwise be available for floor maintenance. The transition endpoint is supranational institutional architecture with the legitimacy and capability to manage global commons and enforce minimum floor conditions across territorial boundaries, developed through the mutual accountability that genuine multipolarity can produce. Multipolarity serves the transition; it is not the destination. The analytical risk: multipolarity is politically useful as a transitional argument but carries the risk of becoming a permanent position that forecloses the institutional development it is supposed to facilitate. The framework's position is that multipolarity arguments should always be accompanied by specification of what institutional development they are designed to enable, so that the transitional character of the position is visible and the conditions under which it becomes obsolete can be identified.

9.3b — The Current Transition: Semiconductor Controls, De-dollarization, and Critical Minerals as Structural Variables

The multipolarity analysis in 9.3 requires grounding in the specific mechanisms through which the current transition is occurring, because the abstract structural argument — multipolarity preferable to unipolarity as transitional arrangement — generates systematically different strategic assessments depending on which material variables are actually in motion. Three dynamics active in the 2024–2026 period illustrate the material basis of the transition in ways that the framework's abstract analysis must account for: the use of semiconductor access controls as an instrument of strategic competition; the acceleration of de-dollarization

initiatives coordinated primarily through the BRICS+ institutional framework; and the emergence of critical minerals as a new constraint layer in the material geography that section 9.0 identifies as the foundation of international order. Each represents a different dimension of the transition, and each generates different analytical implications for the parallel development framework in 9.4. The relationship between them is not coincidental: all three reflect the same underlying dynamic, which is that the material infrastructure of global economic integration — the circuits through which capital, goods, and information move — is being subjected to deliberate fragmentation by state actors who have concluded that the rules-based international order as constituted serves others' strategic interests more reliably than their own. Semiconductor controls as strategic leverage point capture. The US CHIPS and Science Act of 2022 and the subsequent series of export controls on advanced semiconductor technology — extended through Dutch ASML export restrictions on extreme ultraviolet lithography equipment and Japanese controls on semiconductor manufacturing tools — constitute a deliberate attempt to exploit the specific chokepoint structure of the global semiconductor supply chain as an instrument of strategic competition. The supply chain's chokepoint structure is extreme: advanced logic chips at or below 5nm process nodes require EUV lithography equipment manufactured exclusively by ASML in the Netherlands; EUV systems depend on laser components produced in a handful of facilities; advanced chip design depends on electronic design automation software controlled by three US firms. This is a leverage point structure of the kind that 9.2b identifies as the primary mechanism through which technological transitions redistribute material power — and the export control strategy is an attempt to weaponize that leverage by preventing competitors from accessing the technological infrastructure required to contest dominance in the constraint layer that AI capability depends on. The analytical implication for 9.3's multipolarity framework is specific: the semiconductor control regime represents one of the clearest contemporary cases of a dominant power using its position at a technological leverage point to shape the constraint structure

of the transition, with the explicit goal of locking in current advantages before the transition closes the window. This is not realism in the abstract; it is the capture sequence that 16.1 identifies — operating at the level of interstate competition rather than domestic institutional dynamics, but structurally identical. The counter-dynamic — China's domestic semiconductor development program, the accelerated investment in mature-node chip manufacturing, the recruitment of displaced ASML and TSMC engineers — is equally legible as a counter-power accumulation response to exclusion from a leverage point controlled by others. De-dollarization as institutional transition strategy. Dollar hegemony — the use of the US dollar as the primary global reserve currency and the denomination currency for commodity pricing — is one of the most analytically underemphasized features of the material basis of international order that section 9.0 identifies as the constraint structure for all domestic arrangements. Dollar hegemony means that any state whose external accounts depend on dollar-denominated trade can have its economic policy constrained by US monetary decisions that are made without reference to that state's conditions; that the US Treasury's ability to freeze dollar-denominated assets gives it leverage over any state whose reserves or financial infrastructure are held in dollars; and that secondary sanctions — the ability to exclude from US financial infrastructure any third-party institution that transacts with a sanctioned entity — extend US coercive capacity well beyond the bilateral relationship. The freezing of Russian central bank reserves in February 2022 following the invasion of Ukraine demonstrated these mechanisms at scale and accelerated de-dollarization discussions among states that had previously treated dollar hegemony as a background condition rather than a strategic vulnerability. The BRICS+ framework — expanded to include Saudi Arabia, the UAE, Iran, Egypt, Ethiopia, and Argentina in 2024 — has become the primary institutional vehicle for de-dollarization coordination, including bilateral trade settlement agreements in non-dollar currencies, the mBridge cross-border CBDC platform developed by China, Hong Kong, Thailand, and the UAE, and discussions of a BRICS settlement currency backed by a commodity basket. These initiatives are at different stages of development and face genuine collective action problems — dollar network effects are substantial, and no alternative currently matches the dollar's liquidity and institutional depth. But the framework's analysis of institutional transition

(10.3) indicates that the relevant question is not whether any current alternative matches the incumbent but whether the incumbent's legitimacy among the actors who would need to maintain it is depreciating at a rate that makes alternative-building rational even when the alternatives are currently inferior. On that question, the evidence from 2022

–2026 indicates clear acceleration. De-dollarization is not primarily an economic event; it is an institutional power redistribution event of the kind that 9.0's material analysis of international order must account for. Critical minerals as new constraint layer. Section 9.0's account of the material basis of international order focuses on the geography of water, energy, arable land, and strategic position as the primary constraint variables. The energy transition underway since 2015 and accelerating through the 2020s has produced a new constraint layer that extends the material variables 9.0 identifies: the mineral geography of the clean energy economy. Lithium, cobalt, nickel, manganese, and graphite for battery storage; rare earth elements for permanent magnets in wind turbines and electric motors; copper at scales that existing mining infrastructure cannot supply without substantial new development — these are not marginal inputs but foundational materials for the new geography of resource conflict (Klare 2001) on which decarbonized industrial civilization depends. Their geographic distribution is highly concentrated and does not map onto the existing distribution of industrial power: the Democratic Republic of Congo holds approximately 70 percent of global cobalt reserves; Chile and Australia hold the majority of lithium; China controls 85 percent of global rare earth processing capacity regardless of where the ore originates; Indonesia holds the world's largest nickel reserves and has implemented ore export bans to capture processing value domestically. This distribution creates structural leverage for producers that the existing international economic order, organized around the interests of manufacturing-intensive high-income countries, is not designed to recognize. Indonesia's nickel export ban — challenged by the European Union at the WTO and upheld by the Indonesian government despite an adverse ruling — is the most analytically instructive current case: it is an explicit attempt by a producer state to use resource leverage to force processing capacity development domestically rather than exporting raw materials for value capture elsewhere. This is parallel development (9.4) operationalized through resource control rather than through institutional design alone, and it is generating exactly the institutional conflict that the framework predicts when actors at different positions in the material hierarchy use the leverage points available to them. The three dynamics together constitute what the framework identifies as a structural transition in the material basis of international order rather than a temporary disruption of a stable equilibrium. Each represents a different mechanism: semiconductor controls are an attempt by a dominant actor to maintain technological leverage point control during a transition; de-dollarization is an attempt by a coalition of actors to reduce their exposure to the coercive instruments that the existing fin

ancial infrastructure provides the dominant actor; critical mineral politics is an attempt by resource-holding actors — many of them postcolonial states — to convert resource endowment into industrial development capacity rather than into raw material rents that flow through the existing processing infrastructure. All three generate similar distributional questions for the framework's evaluative criterion: who captures the benefits of the transition's new constraint layer, under what accountability structures, and whether the populations most structurally exposed to the costs of the existing order — those below the floor whose situation the existing distribution of leverage determines — gain or lose leverage as a result. The analytical position PMN generates is not that any current challenger to the existing order is advancing the minimal anchor criterion — BRICS+ states include some of the world's most significant producers of domestic structural suffering — but that the transition is creating constraint-layer changes that the parallel development framework must engage with specifically rather than treating as noise around a stable geopolitical background. The strategic question for actors whose interests are in building toward the direction the framework specifies is which positions within this transition enable the accountability architecture that the transition endpoint requires, and which reproduce the capture dynamics of the existing order under new organizational forms. (See 9.0, 9.2b, 9.4, 16.1.) Materialist reading of the 2024–2026 transition: the three dynamics in 9.3b should be read against the

framework's permanent conditions rather than as primarily ideological or normative developments. Semiconductor controls are not primarily about democracy versus authoritarianism — they are about which actors control the infrastructure of the next constraint layer. De-dollarization is not primarily about multipolarity as a value — it is about reducing exposure to a coercive instrument that concentrates leverage in one actor without accountability mechanisms proportional to that leverage. Critical mineral politics is not primarily about resource nationalism as ideology — it is about postcolonial states converting resource endowment into development capacity using the same leverage-point logic that industrial states used during their own development periods. The ideological framings are real and have political consequences; they are not the analytical starting point. The analytical starting point is: what is the material structure of the constraint layer being contested, who currently controls it, what would a different distribution of control mean for the populations subject to the current distribution, and what are the institutional requirements for any new distribution to be less subject to capture than the one it replaces?

9.4 — The Parallel Development Framework

The strategic framework for the transition is parallel development rather than sequential universalization. Each political unit works toward self-sufficient and materially stable institutional arrangements while gradually building the preconditions for deeper international coordination. The parallel structure matters. Attempts to impose a single model of development across units with radically different material conditions have failed in case after case — not because the model was wrong in the abstract but because different material starting points require different strategies, and forcing convergence before the conditions for it exist destroys development capacity rather than accelerating it. Parallel development is not autarky. It does not require that each political unit develop in isolation or that international economic and political relationships be severed. It requires that each unit's development strategy be designed from its specific material starting position — its resource endowments, its institutional inheritance, its demographic structure, its position in the global division of labor — rather than from a universal template derived from the experience of units with very different starting conditions. The template derived from European industrialization assumed specific resource endowments, specific colonial relationships, and specific geopolitical conditions that made a particular development sequence viable. Applied to units with different starting conditions, the template produces different results — often preserving dependency while adopting its outward forms. The criterion for evaluating parallel development paths is the same as for any other arrangement: does this path create material conditions that reduce structural suffering and expand genuine capacity for development over time? Paths that achieve this criterion through institutional forms different from the liberal-democratic template are not failures of universalism. They are evidence that the criterion is more universal than any particular institutional expression of it. Paths that fail the criterion while conforming to the template are not successes of universalism. They are evidence that the template was mistaken for the criterion. Parallel development as an operational framework means something more specific than the standard prescription of building domestic capacity before engaging internationally. It means developing institutional configurations that are viable given the specific material constraints and geopolitical position of a given unit, rather than attempting to replicate configurations that were produced under conditions that no longer exist or that depend on structural advantages the unit does not possess. The historically dominant model of development — industrialize through export-led growth within a liberalizing international trade system, build democratic institutions once economic growth creates the middle class that demands them — was specific to a parti

cular period of international conditions. Those conditions no longer fully obtain: the ecological constraints on energy-intensive industrialization are harder, the trade system is less open to late entrants in high-value sectors, and the democratic institutions that growth was supposed to generate are themselves under structural pressure in the countries that were supposed to be the model. A parallel development framework requires that each political unit assess its material constraints honestly — its resource base, its demographic profile, its

position in regional and global supply chains, its institutional inheritance, its geopolitical constraints — and design institutional configurations that are viable given those constraints rather than that approximate a universal model. This is not relativism about institutional quality: the evaluative criterion remains the same across all units. It is the recognition that the same evaluative criterion, applied to different material conditions, will produce different institutional designs. The application of the framework is context-specific; the framework itself is not. The coordination implication of parallel development is delayed rather than absent. The argument is not that coordination between units is undesirable but that premature coordination — before units have developed the internal capacity to participate in coordination on terms that reflect their actual interests rather than their immediate vulnerabilities — systematically produces arrangements that serve the interests of the more capable participants. Internal capacity development is the precondition for coordination that is genuinely mutual rather than asymmetric in practice while symmetric in form. The bridge to Part X: parallel development creates the domestic material conditions from which system change becomes possible. The dynamics of how those conditions interact with international pressures — how a political unit's development path shapes its position in the global configuration of power and its capacity to act within or against that configuration — is the connection between the geopolitical analysis of Part IX and the theory of change in Part X. The empirical basis for the parallel approach: The track record of sequential universalization — attempts to export a single development model from successful to less-developed contexts — is sufficient to treat the approach with structural skepticism. The Green Revolution, structural adjustment programs, democracy promotion initiatives, and market liberalization packages have all produced sufficiently consistent patterns of partial success and systematic failure to identify the structural features that distinguish contexts where the approach works from those where it does not. The consistent pattern: the model works reasonably well when the receiving context has the institutional infrastructure to manage the transition costs and to capture the model's productivity gains; it fails when the rec

eiving context lacks that infrastructure and when the model's introduction destroys the existing adaptive mechanisms before alternatives have developed. Parallel development avoids this failure mode by working within existing material conditions rather than assuming they can be rapidly transformed to meet the model's requirements. The international solidarity implication: Parallel development is not isolationism. It is a theory of what kind of international cooperation is productive. Cooperation that respects the material conditions of different contexts and builds toward the infrastructure that genuine coordination requires is productive; cooperation that imposes uniform institutional models that serve the interests of the exporting context while degrading the adaptive capacity of the importing one is extractive regardless of its stated intentions. The framework's position on international solidarity follows from this: genuine solidarity between contexts at different levels of material development requires the more-developed contexts to build institutional arrangements that constrain their own advantage — that prevent their greater resource base from translating into terms-setting power in the international cooperation relationship. Connection to 9.3 (multipolarity) and 9.0 (material basis of international order): parallel development is the strategic framework that corresponds to multipolarity as a transitional arrangement and that is grounded in the material reality that section 9.0 establishes. The three sections constitute a coherent position on international transition strategy: acknowledge the material basis of current order (9.0), maintain multipolarity as preferable to unipolarity in the transition (9.3), and pursue parallel development as the cooperative strategy that builds toward the institutional architecture the transition is moving toward (9.4).

9.5 — Ecological Debt and Climate Justice

Section 9.4's account of parallel development contains an unresolved tension that must be named explicitly. The principle that each political unit should build from its material starting point rather than importing development models from outside is analytically sound. But the material starting points are not equally distributed by natural conditions alone — they are products of historical processes, including specifically the

processes of colonial extraction that 9.0b identifies as constituting the current global distribution of resources and capacity. The parallel development framework, as stated, risks treating the distribution of starting points as given, thereby naturalizing an arrangement that was produced by extraction. The ecological debt concept addresses this directly: the historical extraction of material resources from colonial and postcolonial territories, combined with the disproportionate appropriation of atmospheric commons through industrial emissions concentrated in high-income countries, has produced a structural liability that is not discharged by formal political independence. PMN's analysis of structural suffering requires acknowledging that the starting points from which parallel development proceeds are themselves products of structural arrangements that generated structural suffering — and that this history has implications for what genuine solidarity between units at different starting points requires. The ecological debt argument has two distinct components that must be kept analytically separate. The first is the historical extraction debt: the material resources that were removed from colonial territories — agricultural surplus, mineral extraction, land appropriation, labor extraction — were converted into the capital that funds the current development advantage of high-income countries. This debt is not merely symbolic; it represents the material preconditions of current structural positions. A genuinely neutral parallel development framework would need to account for this history in some form — whether through technology transfer, institutional knowledge sharing, or more direct resource redistribution — rather than treating the current distribution of material starting points as a natural baseline from which equal development proceeds. The second component is the atmospheric commons debt: high-income countries have appropriated a disproportionate share of the atmosphere's carbon absorption capacity in the process of achieving their current development levels. Asking lower-income countries to limit their use of fossil fuel development pathways — pathways that high-income countries used without restriction — imposes the cost of climate stabilization asymmetrically on precisely those countries that contributed least to the problem and have the most limited adaptive capacity. The parallel development framework is structurally incomplete without an account of how this asymmetry is to be addressed. The ecological constraint

nt (3.10) interacts with the debt question in a way that makes evasion analytically impossible. If planetary boundaries are real constraints on the total resource use available across all political units — as the framework's account of ecological limits establishes — then the question of how those constraints are distributed is unavoidably a justice question. A framework that identifies ecological boundaries as non-negotiable material constraints without specifying how the costs of operating within those constraints are distributed across political units that contributed to those constraints unequally has committed the evasion that 1.5 identifies: treating the costs of inaction as zero while scrutinizing only the costs of action. The populations who bear the costs of climate destabilization most acutely — those in low-lying coastal regions, those dependent on glacial meltwater, those in agricultural systems most sensitive to rainfall variability — are predominantly populations in low-income countries that contributed minimally to the atmospheric loading that produced those costs. This distributional pattern is not incidental; it is structurally consistent with the same mechanisms that produced the historical extraction debt. The analysis of ecological constraint and the analysis of historical injustice are therefore not separate inquiries that can be sequenced; they must be conducted simultaneously because the constraints themselves are products of the same processes that generated the injustices. The doctrine implications are provisional because the empirical and institutional questions are genuinely open — what constitutes an adequate accounting of historical extraction, what forms of technology transfer or capacity building discharge atmospheric debt, whether the relevant unit of accounting is the nation-state or something else entirely — but the framework's evaluative criterion generates specific directional commitments. Climate finance arrangements that require low-income countries to bear transition costs without corresponding support from high-income countries fail the diagnostics of 11.0: they distribute the costs of coordination asymmetrically in ways that systematically advantage the actors who produced the problem. Technology transfer frameworks that preserve intellectual property protections on clean energy technologies — protections that themselves reflect the institutional power of

high-income countries — fail the same diagnostics. International climate agreements that establish universal carbon reduction obligations without differentiated responsibility for historical contributions naturalize the historical debt by treating all parties as equivalently positioned when they are not. PMN's position is not that any specific institutional arrangement for addressing ecological debt is correct — that is doctrine work requiring empirical assessment of specific proposals — but that any international framework for managing ecological constrains

that does not account for the asymmetry of historical contribution and current adaptive capacity is structurally inadequate by the framework's own evaluative standard. The relationship between 9.5 and the parallel development framework (9.4): parallel development remains the appropriate framework for how political units build from their material starting points. 9.5 specifies a necessary supplement: what genuine solidarity between units at different starting points requires is not equal treatment — the same institutional rules applied uniformly — but differentiated treatment that accounts for the historical processes that produced the difference in starting points. This is not a concession to special pleading; it is the application of the consistency standard (8.3) to the international domain. The same asymmetric analysis of costs and benefits that PMN applies to domestic institutional arrangements — who bears the costs, who captures the benefits, and what historical processes produced that distribution — applies to international climate governance. A framework that applies this analysis domestically but treats international arrangements as if the parties were equivalently positioned has applied its own evaluative standard inconsistently.

X: How Systems Change: Adaptive Dynamics and Historical Transformation

10.1 — Why Systems Change

Describing how systems function at any given moment is not enough. Understanding how they change — how configurations that appeared stable become unstable, how transitions occur — is necessary for any analysis that aspires to be practically useful rather than merely descriptive. Acting on systems requires understanding how they actually change, not how they theoretically should. Systems change when the conditions that sustain them shift past the threshold at which the existing configuration can be maintained. Change is not primarily driven by ideas, by the moral development of populations, or by the decisions of leaders — though all of these can be proximate causes. The underlying driver is always material: the resource distributions, the technological capacities, the demographic structures, and the ecological constraints that determine what configurations are stable and what configurations are not. Three categories of condition change drive most historical transformation: resource redistribution (changes in who controls what, produced by discovery, depletion, technological change, or conquest), technological transformation (changes in what is possible, which alter the relative advantage of different organizational forms), and demographic pressure (changes in population structure that alter the distribution of needs and the sustainability of existing arrangements). These categories interact and compound each other — the most consequential historical transformations involve all three simultaneously, each accelerating the others. Royal authority and large dynastic power are attained only through a group and group feeling. Group feeling results only from blood relationship or something corresponding to it. (Ibn Khaldun 1377/1958, vol. 1) The threshold concept requires more precision than the single word provides. Systems do not change smoothly in proportion to the pressure applied to them. They resist change — often with increasing rigidity as pressure builds — and then change discontinuously when accumulated pressure exceeds the adaptive capacity of the existing configuration. The resistance is not accidental. It is produced by the actors whose position depends on the existing configuration: they have both the incentive and, because the system has organized resources in their favor, the capacity to absorb and deflect pressure before it reaches the threshold. This is why the relationship between transformation pressure and actual change is not linear. Long periods of apparent stability during which pressure accumulates can terminate in rapid reorganization that appears sudden but was structurally prepared by the accumulation that preceded it. The analogy from physics — a phase transition — is imprecise but directionally useful: the system is not the same after the transition as before, and the transition cannot be predicted from the state of the system at any single moment of apparent stability. The misleading role of proximate causes in historical

analysis follows from this structure. The assassination of Franz Ferdinand did not cause the First World War in any sense that illuminates what the war was or why it happened when it did. The structural conditions — the alliance system, the arms race, the competing imperial interests, the domestic political pressures that made mobilization attractive to governments that might otherwise have backed down — had accumulated to the point where a threshold was near. A different proximate cause would have triggered the same threshold crossing; the absence of the proximate cause would have produced a different trigger, not the indefinite postponement of the transition. Analyzing historical transformations through their proximate causes produces a catalog of accidents. Analyzing them through the structural conditions that determined when a threshold was near produces the kind of understanding that is applicable to cases the analyst has not yet encountered. The transformation pressure formula developed in Part XV (15.3–15.4) operationalizes this analysis. The variables it tracks — the scope and intensity of structural suffering, the duration of conditions without adequate response, the intergenerational transmission of deprivation and grievance, and the adaptive capacity of the existing configuration — are the structural conditions that determine how close the system is

to a threshold crossing, not the proximate events through which the crossing eventually occurs. An analyst who can assess those variables has information that is predictively relevant regardless of which specific event triggers the transition. An analyst who focuses on proximate events has information that is relevant only after the transition has already begun. The inadequacy of equilibrium models: Equilibrium models of social systems — models that treat current arrangements as stable because the forces maintaining them are in balance — fail to capture how systems actually change because they treat stability as the natural state and change as the deviation requiring explanation. The empirical record is the reverse: no social system has maintained the same configuration across significant time spans; the question is not whether change will occur but what form it will take and on what timescale. Equilibrium models consistently underestimate transformation probability because they treat the forces maintaining current configurations as more durable than they are — the adaptive mechanisms through which systems absorb pressure, analyzed in 10.4, are precisely the mechanisms that equilibrium models identify as sources of stability. What equilibrium models miss is that adaptive absorption depletes capacity over time, moving systems toward threshold conditions that the equilibrium model cannot see because it has no representation of the trajectory, only of the current state. Two causal pathways to change: Exogenous shocks — changes in material conditions external to the system's internal dynamics —

can push systems across thresholds that their adaptive mechanisms could otherwise have managed. Wars, natural disasters, technological innovations that originate outside the institutional context, and demographic transitions that operate on biological timescales independently of political choices all generate change pressure that the institutional configuration must absorb or to which it must succumb. Endogenous accumulation — the progressive depletion of adaptive capacity through the very mechanisms that maintain stability — generates change pressure from within the system's own dynamics. The most durable transformations tend to combine both: an exogenous shock that crosses a threshold that endogenous accumulation has already moved the system close to. The analytical obligation: any assessment of why a specific system changed — or why it has not changed despite apparent pressure — requires distinguishing between exogenous and endogenous causal pathways and specifying how they interacted. Explanations that attribute change entirely to exogenous shocks miss the endogenous preparation; explanations that attribute it entirely to endogenous accumulation miss the contingent event that crossed the threshold.

10.2 — Suffering as the Signal of System Stress

When material conditions shift, the first signal that the existing system cannot accommodate the shift is an increase in structural suffering — suffering produced not by individual misfortune but by the inability of existing institutional arrangements to manage the changed conditions. This is the mechanism that connects the biological foundation to the analysis of historical change: suffering at scale is not only a moral problem. It is a diagnostic signal indicating that the system is under stress that its current configuration cannot absorb. Societies with high structural suffering — where large portions of the population experience chronic deprivation, insecurity, or exploitation produced by institutional arrangements — are societies where the conditions for change are accumulating. Cooperation that sustains institutional legitimacy erodes. Trust in enforcement mechanisms declines. The cost-benefit calculation for defection shifts in favor of defection for increasing numbers of actors. The system becomes brittle before it becomes visibly fragile. This is why revolutionary moments so frequently catch observers by surprise. The brittleness accumulates over time as structural suffering increases, but it is not visible in the surface-level indicators that most political analysis tracks. The institutional forms remain intact. Formal authority structures continue to function. The ideology that legitimates the arrangement continues to circulate. And then, when a specific event creates a focal point for the accumulated pressure, the system reorganizes rapidly in ways that seem sudden but are the culmination of a long process of structural stress. Ibn Khaldun's concept of *asabiyyah* — rendered variously as group solidarity, social cohesion, or collective binding capacity; used here as a material variable that

accumulates and depletes over time — captures a version of this dynamic. Societies with high asabiyyah can coordinate effectively, maintain institutional legitimacy, and absorb external pressure. As asabiyyah depletes — as the material conditions that generated it change, as the distribution of coordination benefits becomes more unequal, as structural suffering increases — collective action capacity erodes. The institutional forms persist after the material conditions that sustained them have changed, which is why they fail suddenly rather than gradually. Why suffering is the diagnostic signal: The claim that structural suffering functions as a diagnostic signal of system stress is not a normative claim about what suffering should do — it is a structural observation about what it does. Organisms that cannot meet their biological requirements for physical functioning, for safety, for the relational conditions that human development requires, generate behavioral and organizational responses that accumulate into political pressure. This happens independently of whether the suffering has been ideologically processed as a legitimate grievance, independent

tly of whether organizational capacity to act on it exists, and independently of whether those who govern the system are aware of it. The signal is produced by the gap between biological requirement and institutional provision; it does not require anyone to intend it. The signal and the governing narrative: The relationship between structural suffering and political pressure is mediated by governing narratives, which determines what the signal's political content is. The same level of material deprivation can generate transformation pressure when understood as structural and remediable — when the governing narrative attributes it to changeable institutional arrangements rather than natural conditions or individual failure — and can generate political quiescence when understood as natural, individually caused, or temporary. This is why the visibility suppression variable V in the formula architecture is multiplicative rather than additive in its effect: suppressing the visibility of structural causation does not reduce the underlying suffering but reduces its conversion into organized transformation pressure by blocking the narrative processing that makes collective action rational. The signal is present regardless of V ; V determines whether it is politically audible. The three forms of system stress that suffering signals: Structural suffering signals system stress in three distinct forms that have different implications for transformation probability. Distributional stress — suffering concentrated in specific populations that bear costs not distributed across the full population — signals that the arrangement is producing winners and losers in ways that generate counter-power accumulation among the losers without equivalent pressure from the winners, who have less incentive to change what works for them. Capacity stress — suffering produced by the system's inability to manage demands that are being placed on it, distributed broadly across populations — signals that the arrangement's institutional infrastructure is failing in ways that affect its legitimacy across a wider political coalition than distributional stress typically reaches. Legitimacy stress — suffering produced by visible contradictions between the system's stated principles and its actual operation — signals erosion of the narrative legitimacy component of L , which can trigger rapid collapse in the political coalition that maintains the arrangement. Each form of stress points toward different leverage points for transformation. The signal is not self-interpreting. That suffering at scale is occurring is a diagnostic observation; what it signals about the system's transformation trajectory depends on which form of stress is operating, what counter-power capacity exists to convert the signal into organized political pressure, what the adaptive capacity of the existing arrangement is, and what the repressive capacity available to the governing coalition is. Suffering that signals syst

em stress in a context of high repressive capacity and low counter-power produces different outcomes than the same suffering in a context of low repressive capacity and organized counter-power. The signal specifies the input; the formula architecture specifies what the input produces in interaction with the other variables. The analyst's obligation: For the framework's purposes, identifying where structural suffering is occurring and what system stress it signals is a primary analytical obligation — not because suffering is the only thing that matters but because it is the most reliable indicator of where the gap between institutional performance and biological requirement is largest, and therefore where transformation pressure is accumulating. An analysis that tracks political events, institutional developments, and elite dynamics without tracking where structural suffering is occurring and intensifying is missing the primary input variable of the framework's

predictive architecture. The suffering is the signal; the rest of the analysis determines what the signal means. Diagnostic for system stress via suffering signal: (1) Where is structural suffering concentrated — which populations, which regions, which domains of institutional provision? (2) Which of the three forms of stress is operating — distributional, capacity, or legitimacy? (3) What is the governing narrative that mediates the relationship between the suffering and its political interpretation? (4) Is the V variable suppressing conversion of suffering into organized pressure, and if so through what specific mechanisms? (5) What is the trajectory — is the suffering intensifying, stable, or diminishing — and what does the trajectory imply for the transformation pressure variable T?

10.3 — Technology and System Transition

Technological change is the most reliable driver of discontinuous historical transformation — of changes that are not continuous with what preceded them but that create new configurations of power, cooperation, and possibility. Technology changes the relative efficiency of different forms of organization and different distributions of power in ways that cannot be managed through the institutional arrangements that existed before the technology appeared. The pattern is consistent across cases. A technological transition initially advantages actors who can adopt it earliest — typically those with the resources and organizational capacity to bear adoption costs. This concentration of early advantage tends to increase inequality during the transition. But as the technology diffuses, it disrupts the institutional arrangements that the early concentration produced, because it changes the material basis on which those arrangements rested. The printing press concentrated early advantage among those who could afford presses, then disrupted the institutional monopolies that concentration of textual production had previously sustained. Industrial technology concentrated early advantage among factory builders, then generated the labor movements that disrupted the power arrangements industrialization had produced. The current technological transition — artificial intelligence, automation of cognitive labor, digitization of information production and distribution — follows the same pattern at accelerated pace. The early concentration of advantage among actors with the computational resources and data access required is real and consequential. The disruptions that diffusion will produce — to labor markets, to information environments, to organizational forms that proved most efficient in the previous era — are equally real, though their specific form and timing remain uncertain. A framework calibrated for conditions that are already changing must account for this. (Schumpeter 1942; Perez 2002) Whether technological disruption opens or closes windows for progressive institutional change depends on which actors can absorb adoption costs earliest and which can organize around the disruption before it stabilizes into a new power configuration. The historical record suggests that the window is widest during the transition itself — before actors who accumulated early-adoption advantage have translated that advantage into institutional control of the new configuration. Once that translation occurs, the technology tends to stabilize rather than disrupt power arrangements, now operating in favor of whoever captured it during the transition. The same technology that initially created new organizational possibilities becomes the infrastructure through which those who captured it reproduce their position. Whether that capture occurs — and who benefits from the disruption if it does not — is a political question, not a technological one. Artificial intelligence, automation,

and the digitization of information production are not inherently progressive or regressive forces. They are structural variables: constraint-modifiers that change the cost structure of different forms of organization, coordination, and resistance. The capture problem at technological leverage points (analyzed in detail at 7.3c and Part XVI) is the central political problem of the current transition. The question is whether the infrastructure of the new configuration is captured by actors whose interests diverge from the minimal anchor before the window closes, or whether the transition period is used to build accountability mechanisms into the infrastructure itself — an architectural choice that becomes progressively harder to reverse as the configuration stabilizes. The framework treats technology as a constraint-modifier rather than a determinant

for a specific reason: determinism in either direction — the optimism that sees the current transition as inherently liberating, or the pessimism that sees it as inherently concentrating — bypasses the political analysis of who controls the transition and what institutional choices are made during it. Both positions are forms of the same evasion: they treat the outcome as settled by the technology rather than produced by the interaction between the technology and the institutional arrangements it encounters. The material consequences of any technological transition are politically contestable within limits that the technology itself does not determine. The asymmetry of transition costs: The transition period between the old configuration and the new one consistently distributes costs and benefits in ways that are not proportional to who caused the transition or who will benefit from its completion. The populations whose labor and organizational patterns were calibrated to the old configuration bear concentrated transition costs: skills become economically unrewarded, geographic concentration in industries undergoing transformation becomes a liability, and the social infrastructure built around specific forms of work dissolves faster than the replacement infrastructure develops. The populations positioned to capture the new configuration's benefits — those with the capital, education, and geographic mobility to move into the emerging configuration — bear lower costs and capture higher benefits. This pattern is not incidental to how technological transitions work; it is a structural feature that follows from the interaction between technological change and existing distributions of capital, skill, and institutional access. Why technological determinism fails: Technological determinism — the position that specific technologies produce specific social outcomes regardless of institutional context — fails because institutional context determines what the technology is deployed for and who controls the leverage points through which its effects are mediated. The same communication technology that in one inst

itutional context enables democratic coordination enables surveillance and social control in another. The institutional context is shaped by the configuration of power that exists when the technology is introduced, and that configuration determines which uses of the technology receive investment, which actors control the infrastructure, and which regulatory frameworks govern deployment. Understanding technological transition requires the political economy analysis of who shapes deployment conditions, not only the technical analysis of what the technology makes possible. The framework's analytical position: technological transitions are neither automatically progressive nor automatically regressive with respect to the minimal anchor criterion. They are sites of political contestation in which the distribution of the transition's costs and benefits is determined by the configuration of power at the moment of transition. That configuration can be influenced — through counter-power that shapes regulatory frameworks, ownership structures, and access conditions — but only if the contestation begins early enough in the transition to influence the institutional architecture before it stabilizes around new capture patterns.

10.4 — Adaptation, Not Optimization

Complex adaptive systems do not optimize. They adapt. Optimization implies a defined target and a metric by which progress toward that target can be measured. Adaptation implies the ongoing adjustment of a system's configuration in response to changing conditions, without a predetermined endpoint and without a guarantee that adjustments will produce better outcomes by any standard external to the system's own survival. This distinction has direct implications for how the framework understands political and institutional change. A framework that treated social systems as optimizing would imply that better analysis produces better solutions that can be designed and implemented. A framework that treats them as adapting implies something more demanding: that systems will continue changing in ways not fully predictable from any initial design, that the consequences of interventions will diverge from intentions in ways requiring ongoing adjustment, and that the appropriate posture toward any institutional arrangement is not confidence in its optimality but alertness to the signals that it is failing to adapt adequately. Optimization thinking tends to produce institutional designs that are brittle — that work well under the conditions assumed by the design

and fail when those conditions change. Adaptation thinking produces designs that are more resilient — sacrificing some efficiency under any given set of conditions in order to maintain the capacity to adjust when conditions change. The history of institutional failure is substantially a history of optimization under specific conditions followed by brittleness when those conditions shifted. The history of institutional durability is substantially a history of designs that maintained slack, redundancy, and revision capacity at the cost of efficiency. The argument against technocratic confidence in complex domains follows directly from the same reasoning. The technocrat who claims to have identified the optimal policy for a complex social system has confused adaptation for optimization. The appropriate claim is more modest: this intervention, given current conditions and current understanding, is likely to produce better outcomes than available alternatives — and we will monitor whether it actually does, and revise if it does not. Why the distinction matters for political analysis: The optimization framing is politically consequential because it systematically produces false confidence. If a system is moving toward an optimal state, then the analyst's task is to identify where it currently is on the path to that state and what obstacles are preventing faster progress. If a system is adapting without a predetermined endpoint, the analyst's task is to understand what pressures are currently shaping its adaptive responses, what the range of possible configurations those pressures could produce is, and whether any of those configurations would serve the minimal anchor criterion better than the current one. T

he first framing produces analysis that reads current configurations as steps toward a known destination. The second produces analysis that asks what destination the current pressures are moving toward — a question that requires substantially more structural analysis and produces substantially less comforting answers. Adaptation and the teleology problem: The teleological assumption — that social systems move toward a predetermined endpoint, whether capitalist equilibrium, communist society, liberal democracy, or any other configuration — is the specific form in which the optimization framing most consistently distorts political analysis. Teleological frameworks are analytically convenient because they provide a destination that orients all other analysis: any current arrangement can be assessed as more or less advanced along the path toward the predetermined outcome. They are analytically misleading because social systems do not have predetermined destinations. They have pressures, adaptive responses, path dependencies, and contingencies that together produce trajectories that were not predictable in advance and that are not moving toward any configuration that can be identified as the system's natural endpoint. The materialist tradition's own teleological tendencies — the assumption that the logic of capitalist contradiction moves inevitably toward socialist transformation — were not a marginal error. They produced systematic misreadings of historical cases where capitalism adapted rather than collapsed, where socialist transitions produced configurations that served the minimal anchor worse than what they replaced, and where the predetermined destination obscured the actual dynamics operating in specific contexts. What adaptation without optimization looks like: Adaptive systems without a directional guarantee produce several patterns that teleological frameworks consistently misread. First, they can adapt in ways that make harmful arrangements more durable — more efficient at extracting resources, more effective at suppressing counter-power, more capable of absorbing the pressures that would otherwise produce transformation. The capitalist system's repeated adaptations in response to the crises that were supposed to be terminal — the Great Depression, the stagflation crisis, the 2008 financial crisis — are evidence of adaptive capacity, not progress toward a predetermined endpoint. Second, they can produce rapid non-linear reorganizations that move arrangements from one stable configuration to another without passing through the intermediate stages that teleological analysis would predict. Revolutionary transformations are not the culmination of a process of developing contradictions; they are rapid reorganizations triggered by threshold events in systems that have been accumulating stress without visible transformation. Third, they can produce local adaptations that reduce stress in specific domains while increasing it in others — trading on

e form of structural suffering for another rather than reducing it overall. The practical implication for transformation strategy: if systems adapt rather than optimize, then the question is not "what is the system moving toward?" but "what pressures would need to be applied, and what adaptive responses would they

trigger, to produce configurations that better serve the minimal anchor criterion?" This is a different strategic question, and it requires different analysis. It requires understanding the specific adaptive capacities and constraints of the current arrangement, the specific pressures that existing counter-power can generate, and the range of configurations those pressures could plausibly produce — including configurations that are worse than the current one. The adaptation framing is less comforting than the optimization framing, but it is more accurate and more useful for designing interventions that are adequate to what systems actually do. Connection to threshold analysis in 10.5: the non-linearity of threshold effects is what makes the adaptation-not-optimization distinction most consequential for practical analysis. If systems optimized, transformation pressure would produce linear movement toward better configurations, and the analytical task would be to accelerate progress along the existing trajectory. Because systems adapt, transformation pressure produces non-linear responses whose outcomes depend on the configuration of variables at the moment of threshold crossing — which means the analytical task is to understand those configurations well enough to shape the direction of the non-linear reorganization, not only to generate the pressure that triggers it.

10.5 — Thresholds and Non-Linearity

The most dangerous feature of complex adaptive systems is their non-linearity: incremental pressure can be absorbed up to a threshold beyond which the system reorganizes rapidly and in ways that are not predictable from pre-threshold behavior. This makes the standard political logic of incremental adjustment — making small changes, observing results, adjusting — potentially catastrophic when applied to systems approaching thresholds. The threshold phenomenon appears across all the domains this framework analyzes. Ecological systems absorb pressure within tolerable limits, then collapse when the threshold is crossed. Institutional legitimacy erodes gradually, then fails suddenly when a focal event activates the accumulated loss of confidence. Social cohesion deteriorates incrementally, then disintegrates rapidly when the binding commitments that sustained it are no longer credible. The most consequential risks in complex systems are not the ones that appear in trend lines. They are threshold crossings that the trend lines do not predict. (Prigogine and Stengers 1984; Schelling 1978) The practical problem for political analysis is that thresholds are not visible until they are crossed. A system absorbing significant pressure can appear stable by every standard indicator until the moment of reorganization, because the indicators being tracked measure the system's current state rather than its distance from the point at which the current state can no longer be maintained. Leading indicators — the behavioral signals analyzed in 6.6 — are the best available approximation, but they are probabilistic signals, not precise measurements. This means that threshold awareness requires accepting a specific form of uncertainty: the analyst who correctly identifies that a system is approaching a threshold cannot specify when the threshold will be crossed, and the analyst who cannot specify when may not be believed until after the crossing occurs. The institutional implication cuts in both directions. For actors defending existing arrangements: incremental adjustment adequate to managing pressure below the threshold becomes inadequate once the threshold is crossed, and the window for adjustment closes rapidly. Strategies calibrated for gradual change do not scale to threshold crossings. For actors working for transformation: the period of apparent stability before a threshold crossing is not evidence that transformation is unnecessary or that conditions are not building. It is the normal behavior of a system absorbing pressure that has not yet reached the point of reorganization. Both sets of actors have structural incentives to misread threshold proximity — defenders toward optimism, challengers toward premature activation. The analytical task is to read the leading indicators honestly enough to distinguish pressure accumulation from threshold proximity, knowing the distinction cannot be made with certainty. What makes thresholds dangerous: The danger of threshold

d dynamics for political analysis is not that they are difficult to understand in principle — the concept of a tipping point is widely available — but that they are systematically invisible until crossed. Below the threshold, the system absorbs pressure through its adaptive mechanisms and produces no visible evidence

that it is approaching a reorganization. The adaptive absorption can persist for years or decades, during which the analysis of incremental indicators — political polls, economic statistics, protest size, institutional approval ratings — gives no reliable signal that the system is approaching threshold rather than stabilizing. Then a specific event — an economic shock, a political scandal, a natural disaster, a military defeat — crosses the threshold and produces rapid reorganization that appears sudden and unpredictable to observers whose analytical framework was calibrated to the pre-threshold dynamic. The reorganization was not sudden; the threshold crossing was. The pre-threshold accumulation problem: The political logic that produces threshold catastrophe is individually rational at every step. Below the threshold, incremental adjustment is the rational response to political pressure because it reduces the immediate political urgency at lower cost than fundamental transformation. The actors who manage the adaptive response are not making errors — they are correctly reading the immediate incentive structure, which rewards successful absorption of pressure and penalizes costly transformation. What the individually rational incremental response cannot incorporate is the trajectory: the fact that each successful absorption of pressure raises the threshold slightly while reducing the institutional flexibility available for the next round of adjustment. The accumulation of this dynamic over time produces a system that is increasingly fragile while appearing increasingly stable — because stability has been purchased through the progressive consumption of adaptive capacity. The formula's C variable captures this: elite capture rate K reduces adaptive capacity C , so that systems approaching threshold typically have lower C than their apparent stability would suggest. The direction of reorganization is not predetermined: Threshold crossings do not determine the direction of reorganization — they open the possibility space. The same threshold crossing that creates the conditions for progressive transformation also creates the conditions for authoritarian consolidation, fragmentation, or reorganization around different axes of advantage that reproduce structural suffering in a new form. The specific direction of reorganization depends on which actors are organized, which counter-power configurations exist at the moment of threshold crossing, which narratives are available to give political meaning to the reorganization, and what material resources different coalitions can deploy in the window before a new stable configuration consolidates.

This is why 10.5b's analysis of political windows matters: the window opened by threshold crossing is the period during which the direction of reorganization is most contingent, and therefore most responsive to the counter-power that has been accumulated in anticipation of it. The practical strategic implication is twofold. First, threshold analysis changes what counts as evidence of system health. A system that appears stable because it is absorbing pressure through progressive depletion of adaptive capacity is not healthy — it is accumulating fragility. Identifying that dynamic requires tracking the trajectory of the C variable's components over time, not only the current level of visible political pressure. Second, threshold analysis changes what counts as productive political work in the period before threshold crossing. Accumulating counter-power capacity, developing organizational infrastructure, circulating narrative frameworks that will make the reorganization legible when it occurs — all of this is strategically productive even when it produces no visible transformation, because it shapes the direction of the reorganization that threshold crossing will eventually make possible. The analyst's threshold problem: There is a specific analytical risk in threshold frameworks that the framework must name to avoid reproducing it. Threshold analysis can be used to justify political inaction by arguing that the current period is pre-threshold and that transformation will become possible only when threshold conditions are met — a position that makes all current counter-power work preparatory rather than consequential. This is an error both analytically and strategically. Analytically: threshold conditions are partly endogenous to the counter-power work that precedes threshold crossing — the organizational capacity, narrative infrastructure, and coalition breadth that exist at the moment of threshold crossing are shaped by what work was done before it. Strategically: the actors who benefit from pre-threshold stability are not waiting for threshold conditions to invest in maintaining the arrangement; the asymmetry between active maintenance and passive waiting consistently advantages those defending the current configuration. Diagnostic for threshold proximity: (1) What is the trajectory of adaptive capacity C

— is it stable, declining, or declining while appearing stable through depletion of specific flexibility components? (2) Where is structural suffering concentrated, and is it generating organized counter-power or being absorbed through V-suppression? (3) What is the governing coalition's repressive capacity relative to the organized counter-power it would need to suppress to prevent threshold crossing? (4) What alternative configurations would be available for rapid consolidation if threshold crossing occurred — and which actors are positioned to shape that consolidation?

10.5b — Political Time: Windows, Conjunctures, and the Materiality of Timing

The analysis of thresholds and non-linearity in 10.5 describes when systems become structurally available for change. A distinct but related question is why change happens at some moments and not others that appear structurally equivalent. Two situations can share the same threshold proximity, the same distribution of structural stress, and the same configuration of actors — and one produces transformation while the other produces stabilization. The difference is not noise. It is political time: the specific quality of a moment that makes it more or less available for particular kinds of action. Political theorists across traditions have named this phenomenon differently. Machiavelli's *fortuna* and *virtù* captured the interaction between structural conditions and the capacity to act within them. Gramsci's concept of conjuncture described the specific articulation of contradictions at a given moment — the configuration that makes certain alliances possible and others impossible, certain demands legible and others invisible. Ibn Khaldun's analysis of dynastic cycles showed how the same structural conditions produce different outcomes depending on the state of *asabiyyah* at the moment of pressure. What these accounts share is the recognition that time, for political purposes, is not uniform. Some moments are more generative than others, not because of the intentions of actors but because of the specific configuration of material conditions at that point. Windows of political opportunity are material phenomena, not accidents. They are produced by the intersection of multiple processes moving at different speeds. Legitimacy depletion accumulates slowly, then crosses a threshold rapidly. Organizational capacity is built over years, then becomes available for deployment when conditions align. Elite cohesion fractures gradually, then the fracture becomes visible at a specific moment. The window opens when these processes reach a configuration that simultaneously makes existing arrangements difficult to maintain and makes alternatives more credible than they previously were. The window closes when either the pressure dissipates, the alternatives fail to cohere, or new stabilizing arrangements are established. The asymmetry between opening and closing. Windows open slowly and close quickly. The conditions that produce a window accumulate over extended periods of structural stress — years or decades of legitimacy depletion, organizational building, and fracture of the arrangements that sustain the status quo. The window itself — the period during which those accumulated conditions make transformation available — is typically short. Actors who have not built organizational capacity before the window opens will find themselves unable to use it. Actors who move too early, before the configuration is available, spend resources and absorb repression that diminishes their capacity when the window eventually opens. The implication is that preparation for moments of p

olitical availability must be continuous, not triggered by the appearance of those moments. The misreading of conjuncture. Two systematic errors characterize how political actors misread political time. The first is premature escalation: treating a period of rising structural stress as itself a window, without the full configuration — crossed legitimacy threshold, available organizational capacity, fractured elite consensus — being in place. The result is action that expends capacity and hardens opposition before conditions are available. The second is missed windows: treating a window as though it is the normal condition of gradual accumulation, moving incrementally when the configuration requires decisive action. The window closes, conditions restabilize, and the same structural stress that was available for transformation becomes the substrate for a more deeply entrenched arrangement. Both errors are common; the second is more common among movements with strong organizational discipline and long time horizons, which is why organizational

maturity is not sufficient protection against this failure. The materiality of political time connects the framework's analysis of systems (Part X) to its analysis of agents (10.10–10.12). What the agent ecology section describes at the level of individual actors — the preparation, positioning, and drift of different types — is the micro-level expression of the same dynamics that political time describes at the macro level. An ecosystem of actors whose capacities are continuously developed, whose positions are maintained across cycles of mobilization and demobilization, and whose analysis of conditions is accurate enough to distinguish accumulation from availability — that ecosystem is one that can use windows when they open. An ecosystem that mobilizes only in response to visible crisis, maintains no organizational substrate across quiet periods, and lacks analytical capacity to distinguish structural stress from structural availability will consistently find itself without the capacity the window requires. The materiality of timing means that the same strategic action produces different outcomes in different conjunctures — periods in which the specific configuration of multiple causal variables simultaneously favors a specific type of political outcome. A conjuncture is not simply a moment when pressure is high; it is a moment when the interaction of structural variables creates conditions that are specifically receptive to specific types of intervention. The Arab Spring is a case in which multiple countries entered conjunctures simultaneously — not because the structural variables were identical across countries but because a shared economic shock (food price increases driven by the 2010–2011 commodity price spike) interacted with domestic political variables that had been accumulating differently in each country to produce near-simultaneous threshold conditions. The timing of the shared shock, not only its

magnitude, explains why the conjunctures clustered. The analytical implication for political strategy: Counter-power actors who do not have a theory of conjuncture will consistently misread the strategic environment — either waiting for conditions that never arrive, or acting at moments when the structural configuration cannot support the change they are trying to produce. Building the analytical capacity to recognize conjunctures requires tracking multiple causal variables simultaneously over sustained periods, which is expensive and requires the kind of institutional infrastructure for ongoing structural analysis that most counter-power organizations do not maintain. The investment in that analytical capacity is a counter-power investment as real as the investment in organizational or narrative infrastructure. The limits of conjuncture analysis: conjunctures can be identified retrospectively with considerable precision; their prospective identification is substantially harder because it requires assessing the current trajectory of multiple variables and their likely interaction at a future moment. The framework's formula architecture provides the variable structure for conjuncture analysis; applying it to prospective assessment requires calibrated uncertainty about where each variable is heading, not only where it currently is. The difficulty of predicting when transformation thresholds will be crossed has a structural explanation that the framework has not yet made explicit. Threshold proximity depends on the distribution of beliefs in the relevant population — specifically, on how many people believe that active resistance now has lower expected cost than continued compliance. But that distribution cannot be directly observed: actors do not reveal their private beliefs about resistance costs until they act on them, and in repressive conditions, even private belief concealment is strategic. The threshold is therefore reflexive: it depends on beliefs about what others believe, which are not observable, which means that the threshold becomes visible precisely when it is crossed. This explains why transformations that analysts did not predict are not evidence of analytical failure in the specific sense of missing a detectable signal — they are evidence that threshold dynamics are inherently difficult to predict because their primary variable (distribution of private beliefs about resistance costs) is not observable until the threshold crossing itself makes it visible. Leading indicators for belief shift at threshold are not as reliable as leading indicators for legitimacy depletion (6.6), but they are not entirely absent. The most useful are behavioral rather than attitudinal: accelerated narrative absorption (sections of the population previously resistant to structural causal explanations begin incorporating them); increased volatility in public opinion tracking (previously stable distributions begin showing rapid oscillation, indicating that

anchor beliefs are losing their stabilizing function); elite defection signals (actors with significant structural investments in the current arrangement begin hedging or repositioning, which provides social cover for broader revision); and social network cascade patterns (opinion change spreading rapidly through previously resistant network communities). None of these are reliable predictors; all are consistent with threshold approach without being conclusive evidence of it. Their value is diagnostic rather than predictive: their simultaneous presence raises the probability estimate that threshold proximity has increased, without establishing how close the threshold is or when it will be crossed. Vertical vs horizontal mobilization. When structural suffering generates political mobilization, it does not automatically direct that mobilization toward structural causes. Mobilization is vertical when it is directed against the institutional arrangements and power structures that produce the suffering — when the demand is for structural change. It is horizontal when it is directed against outgroups — other populations, often similarly positioned structurally, who are identified as the cause of the suffering. Horizontal mobilization under conditions of high structural suffering is the primary mechanism through which authoritarian movements build mass constituencies: by providing an explanation of suffering (the outgroup did this) and an object for transformation pressure (remove or subordinate the outgroup) that does not threaten the institutional arrangements that actually produce the suffering and that are often sustained by the authoritarian movement's sponsors. The diagnostic question is not which form of mobilization is occurring — it is usually visible — but what structural conditions determine which form emerges. Vertical mobilization is more likely when: channels for vertical critique are open and organizationally effective (institutions exist to receive and process demands for structural change); meaning infrastructure frames suffering as the product of decisions by identifiable powerful actors rather than fate, nature, or outgroup; and counter-power sufficient to make vertical demands credible is available or accumulating. Horizontal mobilization is more likely when these conditions are absent and when manufactured fracture lines (10.9) have been successfully activated by actors whose structural position is threatened by vertical mobilization. The implication for doctrine (not framework): building the institutional conditions for vertical mobilization — channels, meaning infrastructure, counter-power — is the structural intervention that makes horizontal mobilization less likely, not rhetorical reframing of the outgroup narrative alone.

10.6 — Reform, Revolution, and the Sequencing Problem

The debate between reform and revolution is, within this framework, substantially a debate about sequencing and threshold management. Reform — incremental institutional change within the existing system — is appropriate when the existing system retains sufficient material legitimacy to absorb change without losing the coordination capacity that makes further change possible. Revolution — rapid and discontinuous reorganization — tends to occur when structural stress has accumulated past the threshold at which the system can manage through incremental adjustment. The revolutionary tradition, including the Marxist tradition this framework inherits from, has tended to treat revolution as preferable to reform on normative grounds — as a more complete or authentic transformation. This framework treats it as a diagnostic indicator: revolutionary moments occur when reform has failed, typically because actors who benefit from the existing arrangement have blocked incremental change past the point at which the system could have adapted. The revolution is the system reorganizing around the accumulated pressure that reform failed to address. (Polanyi 1944) The sequencing problem is the most important practical problem in political change: identifying the moment at which incremental adjustment is still possible and acting within that window, rather than either accommodating arrangements past the point at which accommodation is possible or triggering reorganization before the conditions for a stable new arrangement exist. This is what this framework's analysis of adaptive systems is designed to help navigate — not to resolve in the abstract, but to make more tractable through clearer identification of the material conditions that determine the window. The historical record strongly suggests that the sequencing problem is solved more often by accident than by design — that the reformers who prevent revolutionary breakdown tend to act under pressure from below

rather than from strategic foresight from above. Strategic foresight remains valuable. The argument is for building the political conditions that make it actionable, which requires attending to the distribution of power that determines whose warnings get heard. The most consequential historical transitions have not been pure reform or pure revolution. They have been mixed sequences in which reform pressure and revolutionary pressure operated simultaneously, each conditioning the other. The New Deal was not reform in the absence of revolutionary threat — it was reform made possible by the presence of revolutionary threat, which shifted the cost-benefit calculation for elites who had blocked incremental change. The decolonization of Asia and Africa was accomplished neither by reform alone nor by revolutionary overthrow alone, but by movements that combined mass non-compliance, negotiated institutional transition, and credible escalation threat — operating simultaneously in ways that were not always coordinat

ed and sometimes in genuine opposition. Mixed sequences are harder to analyze than pure cases because the relationship between modes is not additive. Revolutionary pressure does not simply add force to reform pressure. It changes the political landscape within which reform is possible: what proposals are considered reasonable, who is willing to negotiate, what concessions can be extracted. Reform activity, in turn, shapes revolutionary pressure — demonstrating what is achievable within the system, absorbing some of the energy that would otherwise escalate, and sometimes co-opting movements in ways that reduce pressure without addressing underlying structural conditions. Neither mode can be analyzed in isolation from what the other is doing at the same moment. The strategic question for actors in any mixed sequence is not 'reform or revolution?' as a choice between strategies, but 'given this configuration of actors, this distribution of structural stress, and this position relative to threshold — what combination of reform pressure and escalation credibility, at what pace, creates conditions for a stable new arrangement rather than destructive breakdown?' That question has no universal answer. It is the question that the framework's material analysis of specific conditions is designed to make more tractable — not to resolve in the abstract, but to analyze concretely enough that the actors best positioned to act within the window can identify when that window is open and what it requires. Doctrinal Sequencing and Structural Configurability The sequencing problem developed in this section acquires a further dimension when doctrinal positions are considered in relation to the typology of structural configurations developed in 4.3. The argument for sequencing has been that the conditions required for one form of transformation must be in place before another form becomes viable — that the order of changes matters materially, not only strategically. The structural configuration typology extends this argument: doctrines derived from PMN are not uniformly applicable across all configuration types. The same doctrine may be the correct response to the constraint structure of one configuration and a catastrophic misapplication in another. This is not an argument for gradualism. It is the materialist requirement that strategy be adequate to actual conditions. A doctrine that ignores the configuration in which it operates is not radical — it is reckless. The applicability of the doctrinal position on private ownership of productive assets illustrates this concretely. In a subsistence configuration — where biological floor conditions are not reliably met and the dominant constraint is the absence of the surplus that would make institutional design meaningful — the abolition doctrine addresses the wrong constraint. The primary need is floor provision; institutional transformation of ownership before floor conditions are met creates coordination vacuums that wo

rsen deprivation rather than resolving it. The practical priority in subsistence configurations is not ownership transformation but floor guarantee, and doctrines that demand ownership transformation first are missequenced for the conditions they encounter. In a captured development configuration — where surplus exists but institutional capture blocks its distribution — the abolition doctrine becomes applicable with a specific requirement: that the design of post-abolition institutions explicitly addresses the capture dynamics that produced the configuration. The twentieth-century experience that 11.2 identifies as the central negative lesson is most directly relevant here: transfer of ownership to state structures without accountability mechanisms does not dissolve capture, it reconstitutes capture under new ownership. The doctrine applies, but the implementation design must treat prevention of party-state capture as a primary requirement, not an

afterthought. Contested development configurations are where the doctrine is most directly applicable and most thoroughly tested. These are the configurations in which the structural conditions for institutional transformation exist — surplus, relative stability of floor provision, sufficient epistemic infrastructure to support organized political action — while the ownership arrangements produce capture, foreclose becoming for large populations, and discount ecological limits. The doctrine of abolition is derived from the diagnostic applied to this configuration, and it is in this configuration that the conditions for its application are most fully present. The contested development configuration is not the only configuration that exists, but it is the configuration the doctrine primarily addresses. Meaning-collapse configurations present the highest risk for doctrinal application, including applications of the abolition doctrine. In meaning-collapse, populations have lost their orienting frameworks and are available for rapid mobilization by whoever offers coherent direction fastest. In this condition, institutional transformation of ownership structures — which requires sustained collective coordination, long-horizon planning, and the ability to evaluate complex trade-offs — faces populations whose capacity for that kind of coordination has been impaired by the collapse of the meaning infrastructure that would support it. Abolition doctrine applied in meaning-collapse without simultaneous investment in meaning infrastructure reconstruction risks accelerating mobilization toward regressive orientations rather than producing the institutional outcomes the doctrine intends. The Marxist-Leninist and broader revolutionary left will read the commitment to the long transition, the war of position, and progressive institutionalism as accommodation — a sophisticated defense of the status quo dressed in radical language. This reading mistakes a strategic judgment for a normative one. The framework does not hol

d that revolutionary rupture is always wrong or that existing arrangements are adequate. It holds that transitions which outrun organizational capacity and institutional preparation have historically produced restorations more brutal than what they displaced — and that this is an empirical observation, not a bourgeois preference for order. The same material analysis that explains why the long transition is usually necessary also specifies precisely when threshold conditions make rupture the only available path (10.5). The framework's critique of the revolutionary tradition is from within historical materialism, not from outside it: the custodian problem (7.3) is what actually destroyed the most serious revolutionary projects of the twentieth century, and a tradition that cannot analyze this with the same rigor it applies to bourgeois institutions has not yet done what its own epistemological commitments require. Civil disobedience occupies a specific structural position in the framework's analysis of political change that is distinct from both legal reform and revolutionary rupture. Its structural logic: actors who have exhausted available legal channels for contesting an arrangement, who face institutional capture that has made those channels non-functional for their purposes, and who are unwilling to accept the costs of waiting for the slow accumulation of counter-power, face a choice between compliance with arrangements they judge to violate the minimal anchor and non-compliance that accepts legal consequences as the price of that judgment. The framework's position on civil disobedience is structural rather than categorical. Civil disobedience is justified when three conditions are simultaneously present: the arrangement being contested produces identifiable structural suffering that satisfies the four criteria of 3.4b; the legal channels for contesting it have been genuinely exhausted rather than deemed inconvenient; and the methods employed satisfy the proportionality and prospective accountability principles of 10.14. The third condition is the most demanding and the most important. Civil disobedience that cannot be prospectively defended — that the actor would be unable to justify publicly if the conditions requiring concealment were removed — has crossed from principled non-compliance into the vigilante failure mode that the agent ecology identifies as counterproductive. Civil disobedience that accepts legal consequences rather than evading them is structurally different from civil disobedience that evades them: accepting consequences demonstrates that the action was directed at contesting the arrangement rather than at personal benefit from non-compliance, and it places the burden of enforcement on the system being contested rather than on the individuals contesting it. This is not a requirement that civil disobedience always accept consequences — under conditions of severe repression, the prospective accountability principle applies to

the historical record rather than to the immediate legal moment. It is a requirement that the decision to evade consequences be made for the same structural reasons that justified the original act, not for the convenience of those performing it. Doctrinal sequencing rule: Before applying any PMN doctrine in a specific context, the analysis must (1) identify the structural configuration of that context, (2) evaluate whether the constraint structure of that configuration is the one the doctrine addresses, (3) assess whether the institutional conditions required for the doctrine's implementation exist or can be constructed, and (4) identify which other transformations must occur in parallel or prior for the doctrine's application to produce its intended effects rather than its failure modes. A doctrine applied without this diagnostic is strategy without material analysis — which PMN consistently identifies as a characteristic failure of political projects that produce structural harm while believing they are eliminating it. Applied to the reform-revolution question specifically: Reform vs revolution — structured decision procedure: (1) Assess adaptive capacity (C in 15.5): if high, the current arrangement can accommodate demands without structural change and reform is the available path. If low, the arrangement is approaching threshold and reform within it may not be possible without first changing the power structure. (2) Assess whether demands can be accommodated without changing core power structures: if the demand can be met while leaving the ownership, accountability, and capture relationships intact, it is reformable. If meeting the demand requires changing those relationships, it is revolutionary in content regardless of its form. (3) Assess counter-power readiness (Cr in 15.8): if low, revolutionary-content demands will be suppressed regardless of threshold proximity. If high, the question is whether the transformation window (15.9, 15.11) is open. (4) Assess legitimacy depletion rate (6.6 indicators): rapid depletion signals threshold approach; slow depletion signals that reform pressure can accumulate over longer timescales. (5) Combined assessment: reform is indicated when C is high, demands are accommodatable, and legitimacy is not in acute crisis. Revolutionary strategy is indicated when C is low, demands are structurally transformative, Cr is adequate, and a transformation window is open or approaching. Mixed sequence (both simultaneously) is indicated when C is uneven across institutional domains — high in some, low in others — or when counter-power is adequate for reform pressure but not yet for transformative confrontation. The decision is not binary; it is an assessment of which path the current structural configuration makes available, and the assessment must be revised as structural conditions change. Stagnant reform-or-revolution debates are usually evidence that the assessment has not been updated as conditions have shifted.

10.7 — PMN and the Political Traditions: A Diagnostic Engagement

The traditions analyzed here are large families, not monolithic entities. Anarcho-communism is not one thing; neither is Marxism-Leninism, social democracy, or any other tradition treated in this section. Each contains significant internal variation — different institutional proposals, different historical experiences, different failure modes — and a more granular analysis would diagnose specific variants separately. The PKI in its 1951–1965 mass-organizing phase presents different structural questions than the PKI in clandestine armed resistance after 1965. Nordic social democracy in its 1945–1975 peak presents different structural questions than the Third Way social democracy of the 1990s. The diagnosis of the "family" is an entry point, not a final judgment. What follows identifies the structural questions PMN poses to each tradition at the level of its most characteristic institutional commitments. Sub-variant analysis is doctrine work — it requires the configurability assessment of 10.6 applied to specific historical and material conditions — and it cannot be settled by the framework alone. The communist traditions — anarcho-communism, Trotskyism, Marxism-Leninism, Left Communism, Maoism, and their variants — represent the most sustained and serious attempts in modern history to produce institutional transformations in the direction that PMN's evaluative criterion points. They deserve engagement as serious theoretical and political projects, not dismissal. What follows is not a verdict on their ultimate validity but a diagnostic: what structural questions does PMN require each tradition to answer, and what must be demonstrated — not merely asserted — for the tradition's application in a given context to be consistent with the evaluative criterion? A tradition that can

answer these questions has a place within PMN's analytical framework. A tradition that insulates its conclusions from the questions does not. The shared question that PMN poses to all communist traditions is this: what is the institutional mechanism — not the ideological commitment, not the leadership quality, not the historical tendency — through which the transformation of ownership and governance that the tradition proposes prevents the reconstitution of capture under new form? This question is not a liberal objection to revolutionary politics. It is the lesson that PMN draws from the historical record of the traditions themselves: that the failure mode they most consistently encountered was not external defeat but internal capture, the replacement of capitalist ownership with structures that reproduced its essential feature — concentrated decision-making power over the conditions of becoming for large populations — under different nominal ownership. Any communist tradition that takes this record seriously must answer the structural capture-prevention question. Any tradition that answers it by insisting the record does not apply to its specific variant, without specifying th

e structural mechanism that makes its variant different, is making an ideological claim where PMN requires a structural one. What the diagnostic engagement requires: Engaging the communist traditions as a diagnostic exercise — asking where they are analytically more adequate and where they are architecturally limited — is different from either rehabilitation or dismissal. The traditions analyzed here have historical records that are available for examination: specific predictions about historical trajectories, specific institutional proposals that were implemented, specific analytical frameworks that were applied to specific cases. The diagnostic question is not whether those traditions had the right values — the minimal anchor criterion is not a contested commitment within them — but whether their analytical architecture was adequate to the structural dynamics they were engaged with and where the inadequacies produced characteristic failure patterns. What the communist traditions got right: The Marxist analytical tradition's core contributions are indispensable and have not been superseded by the analytical frameworks that replaced it in academic social science. The analysis of how economic arrangements produce class positions with systematically different interests; the identification of ideology as a mechanism through which those interests are universalized and naturalized; the observation that political institutions are not neutral arbiters of competing interests but arrangements that embody and reproduce specific distributions of power; the insistence that material conditions rather than ideational structures are the primary determinants of social outcomes — all of these remain analytically sound and are incorporated directly into the PMN framework. What was wrong was not the analytical core but the theoretical superstructure built on top of it: the teleological guarantee, the vanguard theory of political organization, the underestimation of the state as an independent force with its own interests, and the analytical subordination of suffering to class position in ways that produced systematic blindness to forms of structural harm that did not fit the class schema. The teleological failure in practice: The teleological guarantee — the claim that the logic of capitalist contradiction necessarily produces socialist transformation — was not merely a theoretical error. It produced characteristic strategic failures across multiple historical contexts. It generated excessive confidence in the imminence of transformation at moments when the actual structural variables indicated otherwise, producing organizations that invested in preparing for an imminent threshold crossing rather than building the durable institutional infrastructure that sustained counter-power requires. It produced insufficient attention to the direction of reorganization at moments of genuine threshold crossing — assuming that the direction of transformation was determined by th

e logic of historical development rather than by the organizational capacity and narrative resources available at the moment — which is why so many revolutionary transformations produced custodian capture rather than the structural change the revolutionary moment made possible. And it produced analytical blindness to the communist tradition's own institutional dynamics, because the teleological guarantee exempted the tradition's own organizations from the same analysis it applied to capitalist institutions. The vanguard theory and its structural consequences: Leninist vanguard theory — the claim that a disciplined party with correct theoretical understanding must lead the working class toward its historically determined interests — solved a real organizational problem and created a structural catastrophe. The real problem: collective action under

conditions of surveillance and repression requires organizational discipline that mass democratic organizations cannot maintain. The structural catastrophe: the organizational form that produces discipline under repression also produces the concentration of decision-making authority that the custodian problem analyzes. Vanguard organizations consistently captured the transformations they enabled, redirecting them toward organizational maintenance rather than the structural changes their stated programs required. This was not a failure of individual leaders who betrayed the revolution; it was the structural consequence of an organizational form that concentrated authority without building the accountability mechanisms that would have constrained it. The organizational form was adequate to the conditions it was built for — clandestine operation under repression — and catastrophically inadequate to the conditions of post-revolutionary governance that it was then applied to without modification. PMN's diagnostic position: the communist traditions were analytically more adequate than the liberal frameworks that displaced them in academic social science in several key dimensions — the analysis of power, of ideology, of structural constraint on individual agency — and analytically less adequate in others: the teleological architecture, the vanguard organizational theory, and the underestimation of non-class forms of structural suffering. What PMN inherits is the analytical core stripped of the teleological superstructure, combined with the biological foundation that grounds the evaluative criterion in a way that does not depend on historical guarantees, and with the organizational analysis that the custodian problem provides as a corrective to vanguard theory's structural failure. The connection to 10.7c: the dimensions on which PMN claims to be architecturally more advanced than the communist traditions are not claims about ideological superiority. They are specific analytical claims about where the communist traditions' architectural choices produced characteristic blind spots — the places where the analysis

was systematically unable to see what it needed to see to avoid the failure patterns that recurred across different historical implementations of broadly similar frameworks. The diagnostic is an exercise in the kind of framework self-awareness that 1.2 requires: understanding where your analytical architecture produces characteristic errors is the precondition for building in the corrections that reduce their frequency.

10.7b — Gramsci: Hegemony, Biological Floor, and the Level Relation

Gramsci's analysis of hegemony is the most developed account available in the Western Marxist tradition of how power maintains itself through the consent of those it subordinates — and why the assumption that material deprivation automatically generates revolutionary consciousness is not only empirically wrong but theoretically confused. The analysis is indispensable for PMN's account of how structural arrangements reproduce themselves under conditions where coercion is available but is deployed selectively. Understanding where Gramsci's framework reaches its limits — and what PMN's biological foundation adds that Gramsci's culturalist account cannot supply — is part of understanding what PMN is doing differently. What hegemony actually means: Hegemony in Gramsci's use is not propaganda or false consciousness — it is not the manipulation of beliefs through systematic distortion. It is the organization of consent: the processes through which the interests and worldview of dominant groups become the common sense of the entire social formation, including those whose material interests it systematically underserves. This happens not primarily through explicit ideological instruction but through the texture of everyday institutional life — the categories that educational systems make available, the assumptions embedded in legal doctrine, the implicit hierarchies of worthiness that welfare institutions enact in their operating procedures. Hegemony is the sum of what feels natural, obvious, and beyond question in a social formation at a given time. The organic intellectual: The concept of the organic intellectual is a strategic corollary of hegemony theory. If the reproduction of dominant arrangements depends on their common-sense status — on their embeddedness in the taken-for-granted — then the counter-hegemonic project requires not primarily material redistribution but the production of alternative common sense: frameworks, categories, and narratives that make visible what hegemony renders invisible and make thinkable what hegemony makes seem impossible. The organic

intellectual is the figure whose function is the production and circulation of this alternative common sense — not the intellectual who interprets the world from outside the communities whose conditions require analysis, but the one who is embedded in those communities and whose analytical work is continuous with their practical engagement. The biological floor Gramsci does not have: The limit of Gramsci's account, from PMN's perspective, is its incomplete account of why hegemony has limits — of why consent is never total, why suffering persistently generates pressure even when it has been ideologically managed, why counter-hegemonic projects can begin from nothing visible in the existing ideological formation. Gramsci attributes this to the contradictions within the hegemonic formation itself — to the gaps between ideology and experience that produce the conditions for alternative common

sense. PMN's biological foundation supplies a more precise account: organisms that can suffer generate political pressure from the gap between their biological requirements and what existing arrangements provide, independently of whether that gap has been ideologically processed as a legitimate grievance. The floor of biological requirement is the source of pressure that hegemonic arrangements cannot fully neutralize because it is not primarily a cultural or ideological phenomenon. It is material in a sense prior to the material of economic relations. The level relation: Gramsci's analysis operates primarily at the level of civil society — the institutional terrain between the state and the economy where hegemony is produced and contested. PMN's framework operates simultaneously at this level and at the level of biological constraint that sits beneath it. The relation between the two levels matters for strategy: cultural and intellectual counter-hegemony is necessary but not sufficient, because the conditions it produces — shifts in common sense, changes in what is thinkable and sayable — must connect to the level of material arrangements to produce durable structural change. Gramsci understood this but could not fully specify the connection, because he did not have a biological foundation that would anchor the analysis of when ideological shifts become materially consequential. The formula architecture in Part XV is partly an attempt to make that connection explicit. The PMN use of Gramsci: hegemony theory is directly applicable to the analysis of V — the visibility suppression variable — and to the conditions under which governing narratives maintain legitimacy despite producing structural suffering at scale. The organic intellectual concept maps onto the ecosystem role analysis in 10.10–10.11. What PMN adds is specificity about the relationship between ideological conditions and material ones: cultural work that does not connect to the threshold dynamics analyzed in 10.5 and the counter-power accumulation analyzed in 10.9 risks producing what Gramsci called transformism — the absorption of counter-hegemonic content into the hegemonic formation, which extends its reach rather than displacing it. Anarcho-Communism: The Structural Insight and the Coordination Problem Anarcho-communism occupies a distinctive position in PMN's diagnostic: it correctly identifies the state as the primary institutional vehicle through which capture reconstitutes itself after ownership transformation. The twentieth-century pattern that PMN's abolition doctrine treats as the central negative lesson — transfer of ownership to state structures without accountability mechanisms producing party-state capture — is precisely what anarcho-communist theory anticipated and warned against. The structural diagnosis is sound. The state is not a neutral instrument that a revolutionary movement can seize and redirect; it is a configuration of concentrated decision-making power that

tends to reproduce concentrated power in whoever operates it, regardless of their intentions or ideology. The structural question PMN poses to anarcho-communism is not whether this diagnosis is correct but whether the proposed alternative resolves the coordination problem the state was solving. Complex production, distribution of conditions material to becoming, and management of ecological boundaries require coordination mechanisms that operate at scale — mechanisms that can allocate resources, resolve conflicts between competing claims, and maintain the institutional infrastructure on which floor conditions depend. The historical anarcho-communist answer — voluntary federation of workers' councils, mutual aid networks, direct democracy at the community level — is a genuine attempt to provide coordination without concentrated power. But the demonstration that this answer is adequate to contemporary scale and complexity has not been made at the scale required. PMN does not require this demonstration to be

completed before anarcho-communism can operate as an analytical framework. It does require that the coordination problem be named as a real structural challenge, not dissolved by the ideological rejection of the need for coordination structures. The anarcho-communist position within PMN: the tradition's diagnostic contribution — state capture as the central failure mode of ownership transformation — is among the most valuable in the communist canon. Its institutional proposal is genuinely structural rather than merely ideological. The structural question it must answer is coordination adequacy at scale: what mechanisms replace the coordination functions the state performs, how do those mechanisms prevent the concentration of coordination power that produces capture, and what is the evidence that they can operate stably across the configurations where they are proposed? Trotskyism: The Bureaucratic Degeneration Analysis and Its Sequencing Commitments Trotskyism's central theoretical contribution to PMN's analytical vocabulary is the concept of bureaucratic degeneration — the diagnosis that the Soviet Union did not fail because socialism was attempted but because the isolation of a revolutionary state in underdeveloped conditions, combined with the physical destruction of the working class during civil war, produced the conditions for bureaucratic capture of the state apparatus. This analysis maps directly onto PMN's structural configuration framework: a captured development configuration cannot be resolved by ownership transformation alone if the institutional layer is simultaneously being consolidated as a new capture structure. Trotsky's thermidor metaphor — the revolution consumed by the bureaucratic stratum it had produced — is a precise description of what PMN's doctrinal sequencing analysis identifies as the core risk of applying abolition doctrine in conditions where accountability structures for post-abolition governance have not been bu

ilt. The structural question PMN poses to Trotskyism is whether its institutional proposal — the revolutionary party operating under democratic centralism, corrected for the deformations it diagnoses in Stalinism — produces a structurally different outcome or reproduces the capture dynamic through a more sophisticated mechanism. Democratic centralism, even in its Trotskyist formulation, concentrates decision-making authority in a central party apparatus whose internal accountability structures are discipline-based rather than structurally independent. The question is not whether a Trotskyist party would intend to prevent capture — intentions are not structural — but whether the organizational form contains the mechanisms that would make capture detectable and correctable from within, by actors with the institutional standing to demand correction, before it becomes entrenched. The historical record of Trotskyist parties suggests that this question has not been resolved by the organizational form itself. The permanent revolution thesis presents a specific sequencing concern. The argument that socialist transformation cannot be confined to advanced industrial economies — that it must generalize globally or face defeat through capitalist encirclement — is compatible with PMN's analysis of structural configurations as globally interconnected. But the implication that socialist transformation should proceed in underdeveloped configurations, before floor conditions are stably met and before the institutional layer is developed enough to support post-abolition governance, runs against the sequencing analysis in 10.6. Trotskyism applied in subsistence configurations risks the floor sacrifice that PMN's diagnostic identifies as a structural rather than contingent failure mode. The Trotskyist position within PMN: the tradition's diagnostic contribution — bureaucratic degeneration as the primary failure mode of Stalinist capture — is one of the most structurally sophisticated analyses in the communist canon and maps directly onto PMN's capture framework. Its institutional proposal remains structurally incomplete on the capture-prevention question: democratic centralism concentrates authority in ways that the historical record suggests are prone to the very capture the tradition correctly diagnoses. The permanent revolution thesis requires configurability evaluation before application. Marxism-Leninism and Stalinism: The Vanguard Problem and the Historical Record Marxism-Leninism is the communist tradition most in need of detailed engagement from PMN's perspective, because it is the tradition that produced the most extensive institutional experiments in ownership transformation and the most consequential failures. The structural questions PMN poses to it are not rhetorical. They are genuine inquiries into whether the institutional form Leninism proposes — the vanguard party as the bearer of revolutionary consciousness, operating under democratic centralism, s

eizing state power and using it to transform the conditions of production — can be designed to prevent the capture it historically produced, or whether the capture was a structural consequence of the form itself. Lenin's contribution to the communist tradition that PMN takes seriously is the recognition that spontaneous working-class organization is insufficient for the kind of structural transformation that conditions of concentrated capitalist power require — that organization, discipline, and theoretical clarity are prerequisites for effective collective action, not optional additions. This recognition is compatible with PMN's analysis of what contested development configurations require for transformation to produce adaptive response capacity expansion rather than mere elite rotation. The question is whether the specific institutional form Lenin proposed — the vanguard party with monopoly on revolutionary correctness — is the only way to provide those prerequisites, or whether it is one design option among several, with specific structural liabilities that alternatives might avoid. The material achievements of ML governance at its floor-provision height require honest acknowledgment rather than subordination to the structural critique. Soviet literacy campaigns transformed one of the world's most illiterate major populations — estimated at 70–80 percent illiterate in 1917 — to near-universal literacy within a generation. Cuban healthcare under ML governance produced life expectancy and infant mortality figures comparable to wealthy liberal democracies despite chronic material scarcity. Vietnamese land reform, implemented by an ML vanguard party, transferred productive assets from landlords to peasant households at a scale no prior arrangement had achieved. The Soviet industrialization drive — regardless of its costs — built the productive and military capacity that enabled the defeat of Nazi fascism on the Eastern Front, which remains the single largest and most consequential military confrontation in modern history. These are real floor achievements by PMN's own criterion. Any analytical framework that evaluates ML governance while systematically omitting or minimizing these achievements is not applying the criterion consistently — it is applying a prior judgment and selecting evidence to support it. The utilitarian context in which vanguard organization developed also requires engagement rather than extraction. The Leninist institutional form did not emerge from a theoretical preference for concentrated authority. It developed under the material conditions of Tsarist autocracy — in which independent political organization was illegal, a free press was absent, mass illiteracy structurally limited participatory alternatives, and the state apparatus had demonstrated its willingness to suppress organizing through systematic violence. The post-revolutionary period added: civil war with foreign military intervention by fourteen states, economic b

lockade, and the ongoing threat of capitalist encirclement that Lenin's successors used to justify emergency measures that eventually became permanent features. PMN's own institutional analysis holds that organizational forms are responses to material conditions: the vanguard party form is analyzable as a solution to the specific coordination problems of revolutionary organizing under autocratic suppression, not as a free design choice from a menu of equivalent options. This does not vindicate its post-revolutionary application, which PMN regards as a genuine structural failure. It means the critique must engage what the form was solving before evaluating whether it created worse problems than those it addressed. The hardest genuine debate that PMN must engage rather than resolve is the utilitarian question about forced industrialization: whether the suffering imposed by Soviet collectivization and rapid industrialization — including the Holodomor, the Gulag system, and the millions of deaths attributable to the First and Second Five-Year Plans — was 'worth it' in the sense that the industrial and military capacity built under those conditions was necessary for the defeat of Nazi Germany, and that defeat prevented suffering on a comparable or greater scale. This debate is not settled. Historians disagree about whether alternative industrialization paths were feasible, whether Soviet military capacity required the specific pace and form of the Stalinist program, and whether the causal chain from collectivization to anti-fascist capability is as direct as the utilitarian argument requires. PMN does not resolve this debate. It insists that the debate be had honestly, with the scale of both the costs and the prevented suffering taken seriously, rather than resolved by declaring one set of deaths the criterion and ignoring the other. The consistency standard that PMN applies to all institutional arrangements requires the same analytical standard here: evaluate the counterfactual honestly, acknowledge the evidence

on both sides, and resist the temptation to reach a clean verdict on what remains a genuinely unresolved question. The structural liability of the vanguard party is its claim to monopoly on correct analysis. This claim is not incidental to the Leninist model — it is the justification for the organizational discipline that makes the model function. Without the claim that the party's analysis is correct and deviation from it is error, democratic centralism has no principled basis for the authority it concentrates. With the claim, the organization is structurally insulated from the internal criticism that would detect and correct capture before it becomes entrenched. PMN's structural analysis of institutions holds that no organization whose claim to authority depends on the correctness of its analysis can reliably detect the errors in that analysis from within — because the institutional standing required to demand correction is granted by, and

therefore cannot challenge, the same authority that produced the error. The Stalinist application of this model produced the paradigm cases of PMN's failure modes in concentrated form. Forced collectivization in the Soviet Union and Ukraine — applied to agricultural configurations where the biological floor was not reliably met, by a state apparatus whose accountability mechanisms had been systematically dismantled through the elimination of internal party opposition — produced mass starvation as a direct structural consequence of applying the abolition doctrine in a subsistence configuration without the institutional accountability required to correct the application when it produced floor failure. From PMN's perspective, this is not an accident or a deviation from Marxist-Leninist theory correctly applied. It is the structural outcome that the doctrinal sequencing analysis in 10.6 identifies as the predictable result of applying abolition doctrine in conditions where the constraint structure is not the one the doctrine addresses, under governance structures that cannot detect and correct the misapplication. Post-Lenin Marxism-Leninism — including contemporary ML parties and states that reject the Stalinist record while maintaining the vanguard party form — must answer a more specific version of the structural question: what in the institutional design of the proposed vanguard party is structurally different from the historical form in ways that address the capture-prevention problem? Answers that appeal to ideological correctness, to better leadership, to more genuine commitment to the masses, or to the specific national conditions that produced Stalinist distortions are not structural answers. PMN requires that the answer specify the institutional mechanism — who can challenge the party leadership's analysis, through what formal process, with what protection from disciplinary retaliation, and with what authority to compel correction — that would make bureaucratic capture detectable and reversible from within before it becomes entrenched. The PMN structural challenge to Marxism-Leninism: (1) What institutional mechanism within the vanguard party form makes the party's own analysis of material conditions subject to revision from below, by actors who are not already committed to the party's prior analysis? (2) What prevents the concentration of interpretive authority in the party leadership from becoming the capture structure that the ownership transformation was intended to dismantle? (3) For each historical instance where ML governance produced floor failure — mass starvation, systematic foreclosure of becoming for large populations, destruction of epistemic infrastructure — what institutional feature of the proposed model would have detected the failure before it became irreversible and produced correction? These questions are not requirements for ideological purity. They are structural requirements for any governance form that claims to be ad

equate to the task of preventing the capture it identifies as the problem. A version of Marxism-Leninism that takes these structural questions seriously — that acknowledges the capture problem as a real structural challenge rather than a historical accident, that develops accountability mechanisms that do not depend on leadership quality or ideological commitment, and that applies the doctrinal sequencing analysis to assess which structural configurations its proposals actually address — retains a place within PMN's analytical framework. The evaluative criterion does not rule out the Leninist insight that organized, disciplined political action is a prerequisite for structural transformation in contested development configurations. It rules out any organizational form that insulates its analysis from the structural correction mechanism that PMN requires of all institutions operating in conditions where they claim authority over conditions of becoming. Left Communism: Council Democracy and the Anti-Organisationalist Tension Left Communism comprises two

analytically distinct currents that PMN treats separately. The council communist tradition — associated with Anton Pannekoek, Herman Gorter, and the Dutch-German Left — locates the primary institutional failure of Leninism in the substitution of the party for the class. Workers' councils, as organs of direct democratic self-organization at the point of production, are proposed not merely as instruments of revolutionary transition but as the basic institutional form of post-capitalist governance. The accountability structure is structural rather than ideological: decisions are made by and accountable to the workers whose becoming conditions are at stake, rather than delegated to a party apparatus whose accountability runs upward through a discipline hierarchy. PMN's structural assessment of council communism is more favorable than its assessment of the vanguard party traditions on the capture-prevention question, for a specific reason: the accountability structure is embedded in the institutional form rather than dependent on the ideological commitments of the actors who operate it. A workers' council whose decisions are made by and accountable to its members is structurally less prone to the capture dynamic that produces bureaucratic degeneration than a party apparatus whose internal authority hierarchy rewards conformity to leadership analysis. The structural question is coordination adequacy — the same question posed to anarcho-communism. Workers' councils at the enterprise level do not automatically produce the macro-coordination required for complex contemporary economies, ecological management, or floor guarantee at population scale. Council communism's answer to this question — federated councils, with higher-level coordination emerging from lower-level mandates — is structurally coherent but historically undertested at the scale required. The Italian Left current — associated with Amadeo Bord

iga — maintains the centrality of the party while rejecting both democratic centralism's internal discipline mechanisms and the parliamentary engagement that Leninism sometimes permitted. The Bordigist party is more rigidly theoretically defined and more explicitly anti-democratic in its internal organization than the Leninist model. From PMN's perspective, this moves in the wrong direction on the capture-prevention question: an organization that is more insulated from internal revision and more dependent on theoretical purity as the basis for authority is more structurally prone to capture, not less. The Bordigist tradition's genuine contribution — the critique of opportunism and the insistence on theoretical clarity as prerequisites for effective political action — is compatible with PMN's recognition that organizational capacity matters for structural transformation. The institutional form it proposes compounds rather than resolves the structural liability of the vanguard model. Left communism's contribution to PMN: the council communist tradition develops the most structurally sophisticated accountability mechanism in the communist canon — decision-making authority accountable to and revocable by the workers whose becoming is at stake. Its structural question is coordination adequacy at scale. The Bordigist variant maintains the vanguard's structural liability on the capture-prevention question while increasing the insulation from internal revision; PMN finds it structurally less adequate than either the council tradition or the Trotskyist tradition on the central structural challenge. Maoism: The Mass Line Insight and the Cultural Revolution Failure Mode Maoism presents a distinctive analytical challenge because its theoretical contributions and its historical failure modes are both more extreme than in other communist traditions. The mass line methodology — the principle that correct analysis must be derived from and tested against the actual material conditions and expressed needs of the population, not imposed from an external theoretical framework — is the most explicitly materialist methodological commitment in the communist canon. The formulation "from the masses, to the masses" contains a genuine epistemological commitment that PMN can recognize: analysis that is not grounded in actual conditions produces strategies that are inadequate to those conditions, and the test of adequacy is empirical rather than theoretical. The structural question PMN poses to Maoism is whether the institutional implementation of the mass line maintains this epistemological commitment in practice or whether it systematically overrides it through the same vanguard mechanism it nominally rejects. The Cultural Revolution — Mao's mobilization of youth against the party bureaucracy he himself had built — is the decisive case. From PMN's structural perspective, the Cultural Revolution represents the most extreme documented instance of induced meaning-collapse use

d as a political instrument. The systematic destruction of epistemic infrastructure — schools, universities, research institutions, the accumulated knowledge of a generation of scholars and administrators — foreclosed becoming capacity across an entire population for a period measurable in decades. The mobilization of disoriented populations against designated enemies produced exactly the regressive mobilization dynamic that PMN's analysis of meaning-collapse configurations identifies as the primary risk of political action without meaning infrastructure. The mass line methodology, correctly applied, would have detected the Cultural Revolution as a failure mode before it produced irreversible damage: the actual material conditions of Chinese workers and peasants were not improved by the destruction of educational and institutional infrastructure. The gap between the methodology's epistemological commitment and its institutional implementation reveals the structural problem: the mass line requires an institutional mechanism through which the masses' assessment of their actual conditions can override the party leadership's assessment of what the masses need. Without that mechanism — without the structural standing for ordinary people to challenge the leadership's analysis and compel revision — the mass line is an epistemological aspiration, not an institutional constraint. The Cultural Revolution demonstrates what happens when the aspiration is not implemented as constraint. Contemporary Maoist movements, particularly those operating in subsistence and captured development configurations in the Global South, present a more complex analytical picture. The recognition that conditions of extreme deprivation and institutional capture require strategies adequate to those conditions — rather than the application of strategies developed for contested development configurations in advanced industrial economies — is consistent with PMN's configurability analysis. The structural question remains: what institutional mechanism in the proposed organization makes the mass line a structural constraint rather than an aspiration, and what prevents the capture dynamic from reconstituting a vanguard apparatus behind mass line rhetoric? Maoism's contribution to PMN: the mass line methodology's epistemological commitment — analysis must be grounded in and tested against actual material conditions — is the most explicitly materialist methodological commitment in the communist canon and is compatible with PMN's diagnostic framework. Its historical failure mode — the Cultural Revolution as induced meaning-collapse mobilized for political purposes — is the most extreme documented instance of the failure mode PMN identifies as the primary risk in meaning-collapse configurations. The structural question is whether the mass line can be implemented as an institutional constraint rather than an aspiration. Social Democracy: Floor Achievements and the Structural Limit Social dem

ocracy's historical record includes the most successful episodes of floor provision since 1945. The post-war Nordic model — combining high rates of unionization, universal social insurance, active labor market policy, and coordinated wage bargaining — produced measurable reductions in structural suffering and genuine expansions of becoming capacity for large working-class populations. The PMN evaluative criterion, applied to the Scandinavian Social Democratic parties at their peak, would produce a favorable assessment on floor provision. The question PMN poses is not only about those achievements but about two related failures. The first is a decisive conjunctural failure: the German SPD in 1919 suppressed the Spartacist uprising, authorized the Freikorps, and enabled the assassination of Rosa Luxemburg and Karl Liebknecht — choosing stability of the capitalist order over the structural transformation the moment made possible. This is not a peripheral historical footnote; it is evidence that social democratic parties, at decisive conjunctures, have consistently prioritized the institutional order over the structural transformation their stated values imply. The second failure is structural: why did those same parties become the primary vehicles for dismantling their own floor achievements in the 1990s and 2000s? The structural answer is that social democracy never altered the ownership structures that generate the capture dynamics PMN identifies as the primary obstacle to sustained becoming expansion. The welfare state was built on top of, not in place of, capitalist ownership of productive assets. This meant that the conditions of its maintenance depended continuously on the cooperation or at least the non-veto of capital: investment decisions, employment levels, and the tax base for social provision remained under the control of owners whose interests were systematically opposed to the redistribution the welfare state required. When the post-war compromise that

had made this arrangement stable broke down — when capital mobility increased, when globalization undermined the national frameworks that had contained it, when the political costs of maintaining the compromise shifted — social democratic parties faced a structural choice between the welfare state and capital cooperation, and chose capital cooperation. This was not a failure of courage or ideology. It was the structural consequence of having built floor provision on top of a capture-prone ownership structure. The specific failure mode PMN identifies in mature social democracy is managed dependency: a configuration in which the biological floor and basic material security are provided, but becoming is systematically blocked by the combination of bureaucratized service delivery, depoliticized governance, and a meaning infrastructure that substitutes consumption for development. The Nordic populations that social democratic governance produced by the 1970s were materially secure and politic

ally stable — and many of them were, in Nietzsche's terms that section 4.5 engages, last men: comfort-maximizing, risk-averse, and resistant to the discomfort that adaptive response capacity requires. Managed dependency is not a PMN success. It is a specific form of structural foreclosure that social democratic governance tends to produce when it succeeds on its own terms. Social democracy's position within PMN: the tradition's historical achievements in floor provision are real and significant. Its structural limitation is that it does not address the ownership arrangements that generate the capture dynamics that eventually erode those achievements. The structural question it must answer is not how to regulate capital better — regulation without ownership change is subject to the ratchet of accumulated capital power — but whether any institutional design within a social democratic framework can prevent the recapture of the regulatory state by the interests it was built to regulate. The failure mode of mature social democracy is managed dependency: floor provision without becoming expansion, producing populations whose security forecloses the discomfort through which development occurs. Ambedkar and South Asian Political Thought: Caste as Permanent Structural Exclusion B.R. Ambedkar's political thought occupies a distinctive analytical position in PMN's diagnostic framework because it directly challenges the sufficiency of class analysis for understanding structural suffering — the same challenge PMN makes against orthodox Marxism, but from an empirical and experiential position that PMN's primarily European inheritance could not have generated. Ambedkar's central analytical claim in *Annihilation of Caste* (1936) is that caste produces a form of structural suffering that is not reducible to class and cannot be addressed by class-based transformation: caste exclusion is a social institution that assigns permanent inferior status to entire populations regardless of economic position, creating a form of institutional betrayal (B) and visibility suppression (V) that persists across all changes in the ownership structures that Marxist analysis treats as primary. Ambedkar's structural argument is specific and analytically important: a Dalit who becomes economically successful does not thereby escape caste exclusion, because caste membership is ascribed by birth and enforced through social practice rather than through economic position. This demonstrates that PMN's variable B — institutional betrayal, the gap between what institutions claim to provide and what they actually deliver — can operate through social institutions rather than only through formal state or economic institutions. Caste is a social institution that betrays its ostensible justification (the organization of social function) by producing systematic dehumanization of specific populations independently of any individual's choices or capacities. Ambedkar's response — his conversion to Bud

dhism along with hundreds of thousands of Dalits in 1956 — is analytically interesting from PMN's perspective as a strategic intervention in meaning infrastructure: the conversion was not primarily metaphysical but was a rejection of the social institution of Hinduism as the framework through which caste exclusion was legitimated, combined with the adoption of an alternative meaning infrastructure that did not assign permanent inferior status on the basis of birth. Ambedkar's engagement with Gandhi is also analytically relevant. Gandhi's program of Hindu reform — the elimination of untouchability while preserving the varna system — represents, from Ambedkar's analysis, a paradigm case of procedural reform that leaves the structural conditions of betrayal intact. The debate between Ambedkar and Gandhi on whether structural suffering can be addressed through transformation of an existing institution or only through

replacement of the institutional framework that produces it is a specific historical instantiation of PMN's reform/revolution/sequencing analysis (10.6): Ambedkar's position is that caste cannot be reformed from within because the institution's logic generates the exclusion rather than distorting it. Ambedkar's contribution to PMN: the caste system demonstrates that structural suffering can be produced and maintained through social institutions that are not primarily economic or political in the Western analytical sense — that the variable B in the suffering formula must include social institutions of status assignment, not only formal state institutions and economic arrangements. This has implications beyond the South Asian context: wherever social institutions assign permanent inferior status on the basis of birth, race, or other ascriptive characteristics — racial hierarchy in the United States, ethnic hierarchy in post-colonial African states, gender hierarchy across contexts — the Ambedkarite analysis applies. PMN's structural suffering formula requires that B be understood to include these social institutions, not only the formal institutional arrangements that liberal analysis treats as the site of accountability. Mariátegui and Latin American Political Thought: The Ayllu and Parallel Development José Carlos Mariátegui (1894–1930) occupies a distinctive position in PMN's analytical framework because he solved, in the specific context of Andean Peru, a problem that the communist tradition broadly failed to solve: how to apply structural analysis to conditions radically different from those in which the analytical framework was developed, without either abandoning the framework or forcing the conditions into categories they do not fit. His *Seven Interpretive Essays on Peruvian Reality* (1928) is the most consequential application of Marxist analysis to a post-colonial condition that incorporates rather than dismisses the pre-colonial institutional inheritance as a material starting point. Mariátegui's central analytical contribution

is the identification of the Andean ayllu — the indigenous communal land-holding and labor-organizing institution of pre-Columbian Andean civilization — not as a primitive precursor to be superseded but as a material institutional resource that capitalism disrupted but did not fully destroy, and that therefore constitutes an actual starting point for alternative institutional development. The ayllu organized collective labor, distributed resources, managed ecological conditions, and maintained community solidarity through mechanisms that did not require the state enforcement or market incentives that European socialist theory assumed were the only available foundations for cooperative economic organization. Mariátegui's claim is empirical and material: these mechanisms existed, they functioned, their remnants persisted into his own time, and any serious analysis of what institutional transformation is possible in Peru must begin from this actual material starting point rather than from a European template that assumes their absence. This is PMN's parallel development framework (9.4) applied with the specificity that framework requires but rarely receives: not the abstract principle that development should begin from actual conditions rather than imported templates, but the concrete identification of what those actual conditions are and what they make institutionally available. Mariátegui's identification of the ayllu as a non-European institutional resource for cooperative economic organization is the kind of contextually specific material analysis that PMN's framework demands but that its European inheritance cannot supply for non-European contexts. The implication generalizes: the parallel development framework requires this kind of ayllu-identification work for every context — the identification of existing institutional resources, cooperative practices, and solidarity mechanisms that the dominant economic arrangement has disrupted but not fully eliminated, and that constitute genuine material starting points rather than blank slates requiring external models. Mariátegui's engagement with the 'Indian question' — his insistence that the condition of indigenous Peruvians was fundamentally an economic and structural problem rather than a racial, cultural, or educational one — is the application of PMN's consistency standard (8.3) to a question that Peruvian intellectual culture of his time treated as a problem of civilization. The indigenista movement of the early twentieth century sought to celebrate and preserve indigenous culture while leaving the land tenure arrangements that concentrated indigenous poverty entirely intact. Mariátegui's response — that the 'Indian problem' was the land problem, and that land reform was the only adequate response to conditions diagnosed as cultural — is a paradigm case of PMN's distinction between symbolic and material approaches to structural exclusion (8.2): the same

conditions that generate cultural margina

lization are produced by economic arrangements, and addressing the cultural symptoms without the economic causes produces exactly the managed dependency that PMN identifies as social democracy's primary failure mode. Mariátegui's explicit rejection of mechanical Marxist application — his insistence that 'we certainly do not wish socialism in America to be a copy and imitation [but] a heroic creation' — is also an epistemological claim that PMN's framework incorporates: the framework is durable across conditions, but the doctrine derived from it must be generated from the specific material conditions of each context. The Comintern's attempt to impose European revolutionary strategy on Latin American conditions — which led to Mariátegui's posthumous condemnation for 'populism' — is a historical case study in doctrinal capture of the kind PMN's section 12.5 identifies as a primary failure mode: the framework becomes a credential that authorizes pre-existing commitments rather than a tool that generates conclusions from analysis. Mariátegui's contribution to PMN: first, the ayllu analysis as the most concrete available demonstration of what parallel development from actual material conditions means in practice — the identification of existing cooperative institutional resources as starting points rather than blank slates. Second, the Indian question as land question as a paradigm case of PMN's distinction between cultural and material analyses of structural exclusion, and of the inadequacy of symbolic approaches that leave productive arrangements intact. Third, the epistemological independence from European templates as a methodological commitment that anticipates PMN's own framework/doctrine distinction: the framework applies across conditions, but the doctrine must be generated from the specific conditions rather than imported from elsewhere. Fourth, the Comintern condemnation as a case study in doctrinal capture applied to the communist tradition's treatment of its own most important non-European thinker. Democratic Socialism: Structural Ambition and the Institutional Design Problem Democratic socialism is analytically distinct from social democracy in ways that matter for PMN's assessment. Where social democracy accepts the permanence of capitalist ownership and aims to regulate its effects, democratic socialism aims to replace capitalist ownership with social ownership — through democratic rather than revolutionary means. This distinction is not merely terminological. It places democratic socialism in a different structural position: it accepts the critique that social democracy's failure mode is a consequence of leaving ownership structures intact, and it proposes to address that critique through the transformation of those structures via electoral and parliamentary politics. The doctrinal position developed in 11.2 is therefore more directly applicable to democratic socialism than to social democracy — both arrive at the same diagnosis of priv

ate ownership, though by different routes. The structural question PMN poses to democratic socialism is about institutional design for post-capitalist ownership. The twentieth century produced multiple experiments in social ownership: state enterprises under social democratic and socialist governments, the Yugoslavian system of workers' self-management, the Mondragon cooperative federation, various forms of municipal ownership, and proposals like the Meidner Plan in Sweden for gradual socialization of productive assets through wage-earner funds. PMN's assessment of these experiments is not categorical — it is configurability-sensitive. Workers' self-management in Yugoslavia produced genuine accountability gains over state ownership in its earlier phases, but was subject to regional capture and increasingly to market pressures that reproduced enterprise-level competition at the expense of systemic coordination. The Mondragon model demonstrates that cooperative ownership can operate at significant scale without the worker alienation that capitalist ownership produces, but faces structural pressures in globalized markets that have required institutional adaptations that some analysts read as partial reversions to capitalist dynamics. The Meidner Plan was never implemented at scale, but its structural logic — gradual transfer of productive assets to collectively managed wage-earner funds — represents the most sophisticated democratic socialist proposal for addressing the ownership question without creating the state-ownership concentration that enables party-state capture. The specific structural challenge PMN identifies for democratic socialism is the transition problem: how does social ownership of productive assets get established and maintained through democratic processes in conditions where the owners of those assets control the investment decisions,

employment levels, and political resources that democratic politics depends on? The historical record suggests that capital responds to credible social ownership proposals with investment strikes, capital flight, and sustained political mobilization against the parties that propose them — using the structural leverage that ownership concentration provides to defeat democratic majorities. Democratic socialism must answer the transition question not with electoral strategy alone, but with institutional designs for how social ownership is established in conditions of coordinated capital opposition, and how it prevents the reconstitution of capture under new nominal ownership. Democratic socialism's position within PMN: the tradition correctly diagnoses the ownership problem that social democracy avoids and proposes to address it through democratic means. Its structural questions are the transition problem — how social ownership is established in conditions of coordinated capital opposition — and the institutional design problem — what prevents social ownership from reproducing capture under new owners

hip form. The historical experiments in social ownership provide a genuine evidence base for these questions; they do not resolve them. The tradition that takes these structural questions seriously, develops institutional designs that address the capture-prevention question, and applies the configurability analysis to assess where and how its proposals are viable, is doing the most rigorous political-structural work in the contemporary left. The Islamic Left: Meaning Infrastructure, Liberation Theology, and the Co-optation Problem The Islamic left — encompassing the liberation theology of Ali Shari'ati, the progressive political thought of figures like Hassan Hanafi, the Islamist socialist movements of the mid-twentieth century, and contemporary progressive Islamic politics in contexts from Indonesia to Morocco — occupies a distinctive position in PMN's analysis because it operates at the intersection of the two dimensions that the framework treats as separately essential: material structural analysis and meaning infrastructure. In societies where Islam provides the dominant meaning framework — where the vocabulary of moral seriousness, collective obligation, and legitimate authority is structured primarily through Islamic categories — the Islamic left represents an attempt to mobilize the motivational and organizational resources of that meaning infrastructure in the direction of structural transformation. This is not a concession to religion as a secondary phenomenon. It is a recognition, consistent with section 5.6's analysis of meaning as material condition, that political mobilization for structural transformation requires meaning resources adequate to the task, and that in Muslim-majority contexts those resources are primarily Islamic. Shari'ati's contribution is the most analytically developed in PMN's terms. His distinction between the Islam of Abu Dharr — the companion of the Prophet who was exiled for his insistence on the illegitimacy of wealth concentration — and the Islam of the court ulama who legitimized Umayyad accumulation is a structural analysis of religious tradition as a site of class struggle over governing narrative. The claim that authentic Islam is structurally opposed to the concentration of wealth and power that produces structural suffering is not a pious wish — it is an argument that the governing narrative of a tradition contains resources that its dominant institutional expressions have suppressed, and that recovering those resources is both epistemically warranted and strategically necessary. PMN's engagement with divine command theory in section 5.4 provides the framework for evaluating this claim: the question is not whether the theological argument is correct but whether the material consequences of the tradition as Shari'ati proposes to practice it are consistent with the evaluative criterion. The structural question PMN poses to the Islamic left concerns the institutional mechanisms through which the progressi

ve interpretation of Islamic tradition is maintained against the co-optation pressure that religious authority consistently faces. The history of Islamic political movements demonstrates this pressure clearly: movements that begin with a structural critique of wealth concentration and political capture are regularly absorbed into the political projects of the same elites they initially challenged, through a combination of material inducements and theological argument that the movement's success requires working within existing power structures. The Iranian revolution produced a governing apparatus that reproduced the concentration of power the revolutionary movement had identified as the primary problem — using Islamic authority claims to insulate that concentration from challenge. The institutionalization of the Muslim Brotherhood in Egypt produced an organization whose structural logic increasingly prioritized organizational maintenance and

political access over the structural transformation its founding critique demanded. The eschatological structure of Islamic meaning infrastructure presents a specific analytical challenge. The promise of divine justice in the afterlife can function as a motivational resource for worldly struggle — the conviction that justice matters absolutely because it is divinely endorsed — or as a demobilizing substitute for it — the reassurance that suffering in this world will be compensated in the next, which reduces the urgency of structural transformation. Both functions are available within Islamic theology, and which function predominates in a specific political context is a structural question, not a theological one: it depends on who controls the authoritative interpretation of the eschatological promise and whose interests that interpretation serves. PMN's analysis requires the Islamic left to specify the institutional mechanisms through which the demobilizing interpretation of eschatology is prevented from dominating — which means specifying the accountability structures that keep religious authority answerable to the material interests of the populations it claims to serve. The Islamic left's position within PMN: the tradition's primary contribution is the mobilization of meaning infrastructure that secular left politics cannot access in Muslim-majority contexts — and the recognition that structural transformation requires motivational resources adequate to the demands of collective action in conditions of significant repression and material deprivation. Its structural questions are the co-optation problem — what institutional mechanisms prevent progressive Islamic politics from being absorbed by the elites its structural critique targets — and the eschatological problem — what determines whether the promise of divine justice functions as mobilizing resource or demobilizing substitute. The tradition that addresses these questions with institutional specificity, rather than theological confidence, is the Islam

ic left that PMN's framework can engage productively.

10.7c — Where PMN Is Architecturally More Advanced: Seven Dimensions

What Gramsci Correctly Identified Antonio Gramsci represents a qualitative advance over orthodox Marxism on the question of how power maintains itself. Marx explained why a ruling class rules — because it controls the means of production and therefore the material conditions on which others depend. What Marx's framework left underspecified was how that ruling class continues to rule in complex modern societies where overt coercion is insufficient, where the dominated classes have developed organizational capacity, and where ideological contestation is a genuine terrain of struggle. Gramsci's answer — hegemony — names a real mechanism: the process through which the worldview, values, and interpretive frameworks of the dominant class come to be experienced as universal common sense by the populations they dominate, so that domination is reproduced not primarily through force but through consent. This is not false consciousness in the crude sense of workers being deceived. It is the more precise observation that the categories through which people understand their situation — including the categories through which they evaluate what is unjust and what is possible — are themselves shaped by the conditions of domination. The war of position — Gramsci's strategic corollary to hegemony — is the appropriate form of political struggle in societies where civil society is dense and the state is buttressed by institutions (schools, media, professional associations, religious bodies, cultural organizations) that absorb and process political conflict before it reaches the level of direct confrontation with state power. In such societies, a frontal attack on state power — the war of maneuver — will fail because the attacker has not secured the civil society institutions that reproduce the conditions of domination. The war of position requires contesting those institutions first: building counter-hegemonic narratives, developing organic intellectuals from within dominated classes who can articulate their experience in terms that challenge rather than reproduce the dominant common sense, and gradually constructing a historical bloc — a coalition of forces united not merely by shared material interest but by a shared interpretive framework that makes their solidarity coherent and durable. The concept of the organic intellectual — an intellectual whose analytical work is grounded in and accountable to the material experience of the class from which they emerge, rather than the professional intellectual whose function is to

reproduce the dominant class's claim to universal expertise — is a genuine analytical contribution that PMN's analysis of epistemic infrastructure (section 1.8) and visibility suppression (section 8.4) draws on implicitly. The Gramscian organic intellectual performs the specific function of making structural causation legible to those experiencing it as personal failure — converting what is experienced as shame, inadequacy, or bad luck into a str

uctural diagnosis that connects individual experience to collective arrangements. This is the visibility-restoration function in Gramscian terms, and it is structurally central to the counter-power accumulation process that section 10.9 analyzes. The Structural Limitations: What Gramsci's Framework Cannot Provide Gramsci's framework is analytically powerful on two questions — how domination is reproduced without direct coercion, and what form counter-hegemonic strategy must take in complex modern societies — and structurally thin on a third question that PMN takes as foundational: why the outcome of the struggle matters, independently of which class's hegemony prevails. The absence of what PMN calls a minimal anchor — a grounding in the non-negotiable biological conditions of human organisms that provides an evaluative criterion prior to and independent of any political tradition — is not an oversight in Gramsci. It follows from his materialism's specific form: a materialism of historical and social relations, in which what matters is determined by the trajectory of class struggle, not by the prior fact of organism-level suffering. The consequence is that Gramscian analysis can explain everything except why the workers should win rather than simply achieve hegemony on different terms. A ruling class that achieves hegemony — that makes its worldview universal common sense — has, on a purely Gramscian account, achieved something. The critique of that achievement requires a criterion external to the hegemonic contest itself: a standard against which any hegemony, regardless of which class produced it, can be evaluated. That criterion is what PMN's biological foundation provides. Structural suffering — in the specific sense of deprivation of conditions material to the flourishing of biological organisms — is the standard against which any institutional arrangement, including any successful hegemony, is evaluated. A hegemony that conceals rather than reduces structural suffering has achieved something real but something that, on PMN's evaluative criterion, fails. Gramsci can diagnose how the concealment works. PMN provides the reason why the concealment matters. The second limitation is analytical rather than normative. Gramsci's framework is optimized for the question of how — how does hegemony function, how should counter-hegemonic strategy proceed — but it has limited resources for the question of when. The theory of the war of position implies a temporality of patient accumulation, of long-term institution-building, of gradual construction of the historical bloc. It does not have strong internal resources for analyzing why transformation happens suddenly when it does, why systems that appeared stable collapse in windows measured in weeks, or why the same strategic position produces different outcomes in different conjunctural configurations. The threshold and non-linearity analysis of sections 10.4 and 10.5 addresses questions that Gramsci's frame

work cannot adequately theorize: the transformation window is not merely a moment when the war of position has been won — it is a structural event produced by the accumulation of pressure beyond the system's absorptive capacity, which can occur independently of the state of counter-hegemonic organization. The Relationship: Foundation and Architecture The correct relationship between PMN and Gramscian analysis is neither synthesis nor replacement. It is a hierarchical relationship in which PMN provides the evaluative and diagnostic foundation — the soil science — and Gramscian analysis provides the strategic and organizational architecture — the building that must be designed in response to the soil's actual properties. The metaphor is not decorative. A building designed without soil analysis will fail in ways that no amount of good architecture can prevent. A soil analysis that produces no building produces no outcomes. Both are necessary; they are not interchangeable; and the soil analysis is logically prior because it determines the range of buildings that are structurally viable in this location under these conditions. What this means in practice is that Gramscian strategy must be calibrated against PMN's diagnostic of the actual configuration of conditions — not applied as a general template. The war of position is the appropriate strategic form when civil society is dense, state power is not immediately vulnerable, and the task is the patient accumulation of

counter-hegemonic resources. PMN's formula for the transformation window (section 10.5) specifies when that assessment changes: when accumulated pressure exceeds the system's adaptive and repressive capacity, the war of position must be accelerated or converted into a war of maneuver if the window is not to close before it can be used. The Gramscian strategist who continues building patiently in the middle of a transformation window has misread the soil conditions — and the cost of that misreading is paid not by the strategist but by the populations whose potential for structural improvement was not realized in the window available. The 2010-2011 Arab Spring provides the most documented recent instance of this diagnostic failure. In Tunisia, Egypt, and elsewhere, counter-hegemonic work had proceeded for years in the war-of-position mode — labor organizing, professional associations, independent media, student movements. When the transformation window opened in December 2010, it opened faster than Gramscian organization could adapt to. The historical bloc that Gramsci requires — a coalition with shared interpretive framework, not merely shared grievance — was not ready. The outcome was the transformation window closing without the counter-power sufficient to consolidate it. In Egypt, this produced the specific configuration of military restoration that the Gramscian framework, optimized for long-term strategic thinking, was poorly equipped to anticipate or prevent. PMN's threshold analysis

s provides the diagnostic that Gramscian strategy requires but cannot generate from within its own resources. Gramscian analysis also requires calibration to the specific technology of hegemony production in a given configuration. Gramsci developed the framework in the context of mass print media, party-organized civil society, and the Catholic Church as the primary vehicle of ideological reproduction in Italy. The infrastructure through which hegemony is produced in contemporary configurations — algorithmic recommendation systems, platform-mediated social networks, synthetic media, attention economics — operates through mechanisms that Gramsci's institutional analysis did not and could not theorize. PMN's analysis of epistemic infrastructure (section 1.8 through 1.12) and the technology analysis of Part XVI updates the terrain on which the war of position must be fought: the organic intellectual who is not algorithmically literate, who does not understand how visibility is produced and suppressed in platform-mediated information environments, is strategically operating in the wrong century. The soil has shifted. The architecture must be redesigned. PMN's Diagnostic Questions for Gramscian Projects A Gramscian project operating within PMN's analytical framework must be able to answer a specific set of questions that its own framework does not generate internally. First: what is the structural suffering being addressed, in terms of deprivation of conditions material to biological organisms? The answer to this question determines whether the counter-hegemonic project is oriented toward the evaluative criterion or toward the class-strategic interest of the counter-hegemonic bloc's leading forces, which may or may not coincide. A historical bloc that achieves hegemony while reducing structural suffering has succeeded. A historical bloc that achieves hegemony while displacing structural suffering onto different populations, or while generating new forms of deprivation in the process of eliminating old ones, has achieved hegemony and failed the evaluative criterion simultaneously. Second: what is the current configuration of the transformation window? Is the soil in the patient-accumulation phase (war of position appropriate), the approaching-threshold phase (acceleration required), or the open-window phase (war of maneuver necessary)? A Gramscian project that cannot answer this question will consistently apply the wrong strategic form: either failing to capitalize on open windows by continuing to organize patiently when the moment requires action, or burning organizational resources in war-of-maneuver attempts against systems that have not reached their threshold and will absorb the attack with limited structural damage. The section 10.5 threshold analysis provides the diagnostic that Gramscian strategy requires at this level. Third: what is the capture vulnerability of the historical bloc being constructed? Gramscian theory analyzes how a ruling clas

s maintains hegemony but does not have strong resources for analyzing how a counter-hegemonic project, once successful, reproduces the conditions it was constructed to challenge. The capture theory of section 7.3 and the general capture mechanism of the section 7.3c sub-section apply directly: a historical bloc that achieves hegemony without having resolved the structural question of how its leading institutions are

prevented from reconstituting concentrated decision-making power under new nominal ownership has not built a Gramscian alternative to domination — it has built a Gramscian pathway to a new form of it. Gramscian analysis diagnostic within PMN: (1) Minimal anchor grounding — has the counter-hegemonic project specified, in biological terms, what structural suffering it is organized to address? If the answer is only that it serves the interest of the dominated class, the capture-prevention question has not been answered. (2) Configuration diagnosis — is the current configuration in the patient-accumulation phase, the approaching-threshold phase, or the open-window phase? Strategy must be calibrated to phase; Gramscian war-of-position logic is appropriate only in the first phase. (3) Epistemic terrain update — does the organic intellectual function account for the actual mechanisms of visibility production and suppression in this configuration's specific information infrastructure? Organic intellectuals who are not fluent in the mechanisms through which hegemony is currently produced cannot contest it effectively. (4) Capture-prevention specification — what institutional mechanism, not ideological commitment, prevents the historical bloc from reconstituting concentrated decision-making power after achieving hegemony? If this question cannot be answered structurally, the Gramscian project has the same failure mode as the Leninist one it typically distinguishes itself from. What PMN Requires of Political Traditions: A Summary The diagnostic engagement above is not a ranking of political traditions by their proximity to some PMN orthodoxy. PMN does not have a preferred tradition, and the evaluative criterion does not require any specific institutional form as its expression. What PMN requires of all communist traditions — and of all political projects claiming to address structural suffering — is consistency between the institutional forms proposed and the structural requirements of the capture-prevention problem. Any political tradition operating within PMN's analytical framework must demonstrate three things. First, that it has a structural rather than ideological or personal answer to the capture-prevention question: what institutional mechanism, not what ideological commitment or leadership quality, prevents the concentration of decision-making power from reproducing the capture structure that the transformation was intended to dismantle? Second, that it applies the doctrinal sequencing analysis honestly: does the configuration in which i

t is proposing transformation have the constraint structure that the proposed doctrine addresses, or is the doctrine being applied in conditions where its failure mode is the predictable structural outcome? Third, that it maintains genuine openness to revision of its conclusions — including revision of the institutional forms it proposes — when those conclusions fail the evaluative criterion in practice. A political tradition that can answer these questions honestly, even if the answers are incomplete and the institutional design is still being worked out, is doing serious structural analysis. PMN's evaluative criterion is consistent with significant institutional diversity in how the conditions for preserved adaptive response capacity are created and maintained. It is not consistent with insulating any specific institutional form from the structural critique that the communist traditions have already leveled at each other — the critique that each tradition has leveled at every other tradition about their respective capture failure modes — and which PMN applies to all of them. PMN's diagnostic assessment of political traditions: (1) Anarcho-communism — most structurally correct on the state-capture problem; structural question is coordination adequacy at scale. (2) Trotskyism — most sophisticated internal diagnosis of bureaucratic capture; structural question is whether democratic centralism resolves or reproduces the capture it diagnoses. (3) Marxism-Leninism — vanguard party's monopoly on interpretive authority is structurally prone to the capture it claims to eliminate; Stalinist application is paradigm case of missequenced doctrinal application in subsistence configuration. (4) Left communism — council tradition provides most structurally embedded accountability mechanism; Bordigism moves in wrong direction. (5) Maoism — mass line is most explicitly materialist methodology in the communist canon; Cultural Revolution is paradigm case of induced meaning-collapse mobilized for political purpose. (6) Social democracy — genuine floor achievements undercut by failure to address ownership structures; failure mode is managed dependency, not mere policy failure. (7) Democratic socialism — correctly diagnoses the ownership problem social democracy avoids; structural questions are the transition problem and the institutional design of capture-preventing social ownership. (8) The Islamic left — mobilizes meaning

infrastructure unavailable to secular politics; structural questions are co-optation prevention and the demobilizing potential of eschatological promise. None of these traditions is ruled out as incompatible with the evaluative criterion on the basis of stated goals. All are challenged by the structural questions this section poses. Compatibility is demonstrated through institutional design, not ideological commitment. Right-wing traditions receive equivalent diagnostic engagement in 10.7d: libertarianism on the state-capture analysis and private

-power blind spot; neoliberalism on the institutional project and its material record; fascism and ethnonationalism on the meaning-collapse diagnosis and the eliminationist resolution that fails the evaluative criterion comprehensively. This diagnostic is not exhaustive. Feminist materialism, ecological socialism, post-colonial political theory, indigenous governance traditions, and other political frameworks that engage seriously with the structural conditions of suffering and becoming all deserve the same diagnostic engagement. The pattern of questions PMN poses — what is the structural mechanism for preventing capture, what is the sequencing analysis for the configuration in which this doctrine is proposed, what is the institutional design for accountability — applies to every political tradition. The invitation is open: any tradition that takes these structural questions seriously, answers them with institutional specificity rather than ideological confidence, and maintains genuine openness to revision when its proposals fail the evaluative criterion, will find a productive relationship with PMN's analytical framework. What PMN cannot accommodate is not political difference but structural evasion — the refusal to subject proposed institutional forms to the same structural critique those traditions apply to the arrangements they seek to transform. Conservatism and Traditionalism: The Structural Insight About Finitude and Meaning The diagnostic section above focuses on traditions that share PMN's direction — toward reduction of structural suffering and expansion of becoming — while differing in their proposed institutional forms. Conservatism and traditionalism warrant treatment as a distinct category because they raise a structural challenge that the progressive traditions do not: the possibility that the meaning infrastructure which human beings require for adaptive response capacity is not separable from the institutional arrangements that produce it, including arrangements that generate forms of structural suffering that PMN would otherwise identify as targets for transformation. This is not a rhetorical challenge to be dismissed. It is an empirical question about the social production of meaning, and PMN's materialist commitments require it to be engaged rather than avoided. The conservative structural insight, at its most serious, is this: meaning infrastructure is not produced by institutional design in the same way that material goods are produced. It accumulates through tradition — through the accumulated practice of communities responding to the permanent conditions of human existence over time — and it depends on the continuity of that practice for its vitality. Disrupting the institutional arrangements through which tradition is maintained does not leave the meaning infrastructure intact while removing the oppressive elements; it tends to dissolve the meaning infrastructure along with the oppressive arrangements, producing the meani

ng-collapse configuration that section 4.3 identifies as the most dangerous structural outcome. The progressive experience of the twentieth century provides sufficient evidence for this dynamic: rapid modernization, however materially beneficial, consistently produced populations disoriented from prior meaning frameworks and available for mobilization by whatever new framework offered the most compelling orientation — including orientations that generated far more structural suffering than the traditional arrangements they replaced. A related and analytically distinct conservative contribution — Chesterton's fence — is the epistemic principle that institutional arrangements should not be dismantled before the purpose they serve is understood. The fence built across a road should not be removed simply because its function is not immediately visible to the observer who encounters it; understanding why it was built is prerequisite to deciding whether removing it is safe. Applied to institutions, this is a genuine constraint on the reform impulse: arrangements that appear to serve only oppressive functions may also be performing coordination functions whose loss is not immediately visible and may only become apparent when the arrangement is already gone. This principle is compatible with transformation — it does not require

preservation of everything that exists — but it demands that the transformation account for all the functions the arrangement performs, not only the most visible ones. PMN treats this as a standing analytical requirement rather than as a conservative political preference: any institutional transformation proposal that has not identified what secondary functions the targeted arrangement performs and what would replace them is analytically incomplete by the framework's own standard. PMN's structural assessment distinguishes between two things that conservatism tends to conflate. The first is the genuine insight: that meaning infrastructure is fragile, that it accumulates through practice and is easily dissolved, that the human need for orientation within a cosmological and social framework is not eliminable and must be addressed by any viable political project. This insight is correct and is directly incorporated into PMN's analysis of meaning as material condition in section 5.6, of the seeker in section 4.5, and of the disenchantment trajectory in section 5.5. The second is the institutional conclusion conservatism draws from this insight: that the preservation of meaning infrastructure requires the preservation of the traditional institutional arrangements in which it is embedded, including patriarchal family structures, hierarchical social organization, religious authority over public life, and national or ethnic particularity as the organizing framework for political community. This conclusion does not follow from the insight, and the evidence for it is weaker than conservative analysis acknowledges. The empirical question is w

hether meaning infrastructure can be maintained or rebuilt through institutional arrangements that do not systematically foreclose becoming for the populations at the bottom of traditional hierarchies. The historical record is mixed but not uniformly negative. Religious traditions have been reformed — internally and under external pressure — in ways that retained their meaning-generating function while reducing, though rarely eliminating, their most severe forms of structural oppression. Cultural traditions have been adapted to new material conditions without dissolving their meaning-generating capacity. The meaning infrastructure associated with national and ethnic community has been sustained in contexts of significant institutional transformation. None of this demonstrates that meaning infrastructure is costlessly separable from the institutional arrangements that generate structural suffering. It demonstrates that the separability question is empirical and context-dependent, not settled in advance by either the conservative claim that preservation requires the whole package or the progressive claim that oppressive arrangements can be stripped away without meaningful loss. Conservatism's position within PMN: the tradition correctly identifies the fragility of meaning infrastructure, the danger of meaning-collapse under rapid transformation, and the genuine costs of dismantling institutional arrangements that perform social functions their architects did not intend and their critics cannot see. These are empirical contributions that PMN incorporates directly. The structural question PMN poses to serious conservatism is not whether meaning infrastructure matters — it does — but whether it can distinguish between arrangements that sustain meaning infrastructure and arrangements that sustain the positional advantage of specific groups at the cost of becoming-foreclosure for others. Where that distinction cannot be made — where the meaning infrastructure genuinely requires the hierarchy, not merely coexists with it — PMN faces an honest trade-off with no clean resolution. Where the distinction can be made — where the meaning function and the oppressive function are separable — the conservative conflation of the two is the structural error. Determining which case obtains is the empirical work that neither the conservative nor the transformationist position has completed.

10.7d — PMN and the Right-Wing Intellectual Traditions

The communist traditions received detailed diagnostic engagement in 10.7 because they represent the most sustained historical attempts to produce structural transformation in the direction PMN's evaluative criterion points. The right-wing intellectual traditions require equivalent engagement for a different reason: they have dominated the institutional arrangements within which most of the world's population lives, and the framework that cannot diagnose their structural insights and structural failures is a framework operating on political terrain it has not adequately mapped. The following is not a survey of conservative or right-wing

politics broadly; it is a diagnostic engagement with the intellectual traditions that have produced the most serious analytical contributions to questions the framework must address. Libertarianism: The State-Capture Insight and the Private-Power Blind Spot Libertarianism — the political tradition descending from Locke through Hayek and Friedman and into contemporary anarcho-capitalism — makes its most durable analytical contribution in a domain where PMN agrees: the analysis of state capture as a structural problem that cannot be resolved by better leadership or more sincere commitment to public purpose. Hayek's argument in *The Road to Serfdom* is not the cartoon version of it that its critics usually engage: it is the structural claim that centralized economic planning concentrates informational and coercive authority in ways that make capture by narrow interests not an accident but a predictable structural outcome, regardless of the intentions of those who designed the system. This argument is structurally parallel to PMN's own analysis of the custodian problem (7.3): any anchor must be held by someone, and whoever holds it develops interests in defining it to their benefit. Hayek was correct that this problem is real, that it applies systematically rather than incidentally, and that institutional design which ignores it will repeatedly produce the outcomes it was designed to prevent. The public choice tradition — Buchanan, Tullock — extended this insight into a rigorous analysis of how democratic processes systematically produce concentrated benefits and diffuse costs through rational political behavior. This is the capture mechanism from the demand side rather than the supply side, and it is genuinely important. What libertarianism gets wrong structurally is the mirror image of what it gets right about states: it applies its capture analysis to public power while treating private power as unproblematic. The Hayekian tradition analyzes the information concentration problems of central planning with analytical precision while treating the information concentration of monopoly capital as a market aberration correctable by competition rather than as a structural outcome of market dynamics under conditions of returns to scale. PMN's diagnostic question 2 (11.0) — does the arrangement permit accumulation of economic

power sufficient to produce institutional capture? — applies to private actors with exactly the same force it applies to state actors. A framework that sees state capture clearly while treating private capture as a non-problem or as a problem solved by market competition is a framework that has drawn the boundary of its analysis in a location that happens to exempt the primary concentrations of power that exist in advanced capitalist economies. Friedman's monetary economics contributed real analytical advances. His political prescription — that free markets and limited government exhaust the requirements for genuine freedom — treats the formal freedom of market participation as equivalent to the substantive freedom that PMN's becoming framework identifies as the actual target. Libertarianism's position within PMN: the tradition's analysis of state capture as a structural problem requiring institutional design solutions rather than personnel improvement is a genuine analytical contribution that PMN incorporates directly. Its structural limitation is the asymmetric application of this insight — rigorous about public power concentration, systematically blind to private power concentration — in ways that produce a framework which serves the interests of those who already hold concentrated private power while claiming to serve freedom generally. PMN's structural question to libertarianism: apply the capture analysis symmetrically. What institutional design prevents private economic power from producing the same information concentration, accountability deficits, and capture of public institutions that the tradition correctly identifies as the structural liability of centralized state power? Neoliberalism: *The Institutional Project and Its Material Record* Neoliberalism requires analytical distinction from both libertarianism and from conservatism, because it is a different kind of thing: not primarily a philosophical tradition about the nature of freedom, and not primarily a disposition toward preserving existing arrangements, but an organized institutional project — the deliberate restructuring of state capacity, regulatory frameworks, and international economic architecture to expand the domain within which market mechanisms operate and to insulate those mechanisms from democratic revision. The Mont Pelerin Society, the Chicago School, the Thatcher and Reagan programs, the Washington Consensus, and the structural adjustment programs administered by the IMF and World Bank are not the natural expression of market logic; they are the product of sustained organized political effort by actors who understood that

market outcomes favorable to capital require active institutional construction and maintenance. What the neoliberal project got analytically right: the insight that markets require institutional preconditions that states must actively provide — contract enforcement, property rights definition, monetary stability, the management of systemic risk —

and that these preconditions do not arise spontaneously from economic activity itself. This is not a trivial insight; it is the foundation of the institutional economics tradition that PMN draws on through North and Ostrom. The neoliberal project also correctly identified that the postwar Keynesian settlement had produced structural rigidities — in labor markets, in corporatist bargaining arrangements, in the sheltered industries of mixed economies — that generated genuine inefficiencies. The stagflation crisis of the 1970s was a real phenomenon that the dominant macroeconomic frameworks of the period could not adequately explain or address. The critique of the postwar settlement was not entirely manufactured. What the neoliberal project got wrong structurally, and what its material record demonstrates: the institutional restructuring it produced systematically transferred the gains from productivity growth from labor to capital, increased the structural suffering produced by economic volatility by removing the social insurance mechanisms that had buffered it, and created the conditions for the financial concentration that produced the 2008 crisis and its aftermath. The Washington Consensus programs administered to developing economies systematically produced worse outcomes on floor provision — infant mortality, life expectancy, educational attainment — than the mixed-economy alternatives they replaced, while improving performance on the financial metrics that creditors cared about. The eight diagnostic questions (11.0) applied to the neoliberal institutional project produce a consistent finding: it performed well on internal market efficiency metrics while systematically failing on distribution, substrate reproduction, ecological boundary conditions, shock resilience, and visibility of reproductive labor. This finding is supported by substantial evidence and represents the dominant conclusion in the development economics literature, though the causal mechanisms and counterfactuals remain genuinely debated. The framework's epistemological commitment requires acknowledging that debate while maintaining that the weight of evidence — documented across Latin American lost decades, Sub-Saharan African structural adjustment programs, and the divergent trajectories of countries that followed versus rejected Washington Consensus prescriptions — supports the structural failure assessment. Neoliberalism's position within PMN: the project correctly identified the institutional preconditions for market function and the genuine rigidities of the postwar Keynesian settlement. Its structural failure is that it applied its institutional design capacity asymmetrically: actively constructing the institutions that expanded market domains while systematically dismantling the institutions that buffered the structural suffering market volatility produces. The result is an arrangement that performs well on the metrics it was designed to optimize while failing the evalu

ative criterion on distribution, floor provision, and the reproduction of the cooperative substrate that complex social arrangements require. PMN's structural question: what institutional design provides the efficiency benefits of competitive market mechanisms while maintaining the accountability structures and social insurance systems that prevent market outcomes from violating the floor criterion for large populations? Fascism and Ethnonationalism as Intellectual Traditions Fascism and ethnonationalism are treated as regime types in 7.8, but they also have intellectual traditions that require diagnostic engagement distinct from the regime analysis. Carl Schmitt's political theory — the distinction between friend and enemy as the fundamental political category, the critique of liberal proceduralism as a depoliticizing displacement of genuine political conflict — is analytically serious in ways that dismissal cannot address. Schmitt correctly identified that liberal constitutionalism tends to treat political conflict as a problem to be managed through procedure rather than as a structural feature of any arrangement where genuine interests diverge. PMN's own analysis of institutional capture (7.3, 7.3b) operates on similar terrain: formal accountability structures tend to be colonized by the interests they were designed to constrain, and the appearance of procedural legitimacy can persist long after the substantive accountability it was supposed to provide has been eliminated. Where Schmitt's analysis fails the PMN criterion is in its resolution: the friend-enemy distinction as the basis for

political community makes structural suffering of the designated enemy not a failure of the arrangement but its intended product. The insight that political conflict is real and cannot be procedurally dissolved is correct. The conclusion that the appropriate response to this reality is the permanent mobilization of a community against an existential enemy does not follow from it. The intellectual tradition of ethnopluralism — Alain de Benoist and the French New Right's argument that each people has the right to its own cultural and political space, and that cosmopolitan universalism is itself a form of cultural imperialism — makes a more sophisticated argument than biological racism, and one that requires more careful engagement. The cultural particularity argument has genuine force: PMN's own analysis of the failure of naive universalism (8.5) and the material basis of cultural difference (8.1) acknowledges that universal templates applied to contexts with different material conditions produce different and often worse outcomes than contextually adequate arrangements. Where ethnopluralism fails is in treating cultural particularity as requiring political separation rather than as requiring contextual application of universal criteria. The right-wing intellectual traditions within PMN: libertarianism contributes the capture analysis and public choice critique,

both genuinely incorporated. Neoliberalism contributes the institutional economics of market preconditions, incorporated with the caveat that the asymmetric application of institutional design capacity is the tradition's structural failure. Schmitt and the ethnopluralist tradition contribute the critique of proceduralist depoliticization and the analysis of cultural particularity, both genuine insights that PMN engages without accepting their resolutions. The consistent structural finding across right-wing intellectual traditions is parallel to the finding across left traditions: the analytical contributions are real and cannot be dismissed; the institutional proposals fail the evaluative criterion in characteristic ways that follow from the structural features of each tradition rather than from the intentions of their proponents. Gramsci: Hegemony, the Missing Biological Floor, and the Relation of Levels Antonio Gramsci represents a qualitative advance over orthodox Marxism in one specific respect: he explains not only why dominant classes hold economic power but how they maintain political and cultural dominance without continuous recourse to coercion. The concept of hegemony — the capacity of a ruling group to make its own values, interpretive frameworks, and definitions of common sense appear as universal — is a genuine contribution to the analysis of how structural arrangements reproduce themselves across generations without requiring that every subordinate population be visibly coerced into compliance. The analysis of why workers identify with enterprises that exploit them, why populations support political programs that operate against their material interests, and why the very categories through which suffering is interpreted can be colonized by the arrangements producing that suffering belongs, in its most developed form, to Gramsci's tradition. Section 6.3's analysis of ideological reproduction and section 8.4's treatment of governing narrative as a material force both draw on this inheritance. The structural question PMN poses to Gramsci is not whether hegemony exists or whether it operates as he described. It is whether his framework has an adequate foundation for explaining why hegemonic arrangements are a problem at all. The answer is that it does not — not in the form that makes the answer non-circular. Gramsci's evaluative criterion, insofar as it is explicit, relies on a Marxist account of historical development and class interest: hegemonic arrangements are bad because they prevent the working class from recognizing and acting on its historical role. But this answer presupposes what it needs to establish. Why is that historical role the correct one? Why is the proletariat's emancipation the evaluative standard rather than, say, social stability or aggregate welfare? The Gramscian framework has no non-circular answer to this question, and the absence produces a specific vulnerability: if the historical conditions that Marxism id

entified as generating proletarian consciousness change — as they have, substantially, since Gramsci's imprisonment — the entire evaluative structure can be destabilized without an alternative foundation to fall back on. PMN's contribution at this point is not to dispute Gramsci's descriptive analysis of hegemony but to provide the missing foundation. The reason hegemonic arrangements are a problem is not that they misalign classes with their historical interests. It is that they suppress the visibility of suffering produced by structural arrangements, thereby preventing the institutional responses that suffering signals the need for. The V term in

PMN's structural suffering formula — visibility suppression — is precisely the mechanism that Gramsci's hegemony describes at the level of consciousness and culture. Hegemony is the ideological component of V suppression: it is how structural suffering becomes misattributed to individual failure, natural necessity, or divine will rather than to the institutional arrangements producing it. This grounding does not require a philosophy of history or a theory of class destiny. It requires only the biological observation that organisms experience pain, and the structural observation that institutional arrangements can be configured to prevent that experience from being correctly attributed and therefore politically addressed. The Biological Floor Gramsci Did Not Build The specific absence in Gramsci's framework that PMN identifies is what this framework calls the biological floor: the level of material reality at which organisms eat, sleep, experience pain, and require specific physical conditions to maintain the capacity for anything else, including political consciousness and cultural participation. Gramsci's analysis operates almost entirely at the level of consciousness, culture, and political organization — the superstructure, in older vocabulary. Even his engagement with material conditions tends toward their significance for political consciousness rather than their character as independent constraints on what is possible for embodied organisms. The practical consequence of this absence is a specific analytical blind spot: if hegemony is sufficiently total — if a ruling culture has succeeded in making its interpretive frameworks genuinely universal within a given population — Gramsci's framework faces a strategic impasse. There is no external lever with which to begin dismantling hegemonic arrangements if the population has fully internalized those arrangements as natural. PMN's response is that this impasse is analytical rather than real. Hegemony can suppress the visibility of suffering; it cannot eliminate the suffering itself. The worker whose factory work produces chronic back pain, whose child's nutrition is inadequate, and whose life expectancy is shortened by the conditions of their labor is experiencing structural suffering regardless of whether they have internalized an ideological framework

k that attributes those conditions to individual inadequacy or natural hierarchy. The accumulation of that suffering — in PMN's terms, the accumulation of R and B across duration and scale — generates pressure that no hegemonic arrangement can permanently contain, because hegemony operates at the level of consciousness and cannot repair damaged bodies or extend shortened lives. This is the structural asymmetry between suffering and ideology that PMN's biological foundation establishes and Gramsci's framework cannot: ideology can suppress visibility indefinitely, but it cannot suppress the material reality of which visibility is a report. The political consequence is that even a population with full hegemonic saturation retains a material foundation for eventual political response, because the conditions producing suffering continue to produce it regardless of how those conditions are interpreted. This does not make change automatic or inevitable — the threshold analysis of section 10.5 shows why pressure accumulation is necessary but not sufficient. But it establishes that the Gramscian impasse of total hegemony is an impasse in consciousness, not in material conditions. The Level Relation: Foundation, Not Synthesis The correct relationship between PMN and the Gramscian tradition is not synthesis — not the combination of two parallel frameworks into a composite theory — but a relation of levels in which PMN's analysis of material conditions and biological foundations establishes what the Gramscian tradition's analysis of hegemony, consciousness, and political strategy operates within. This distinction matters because it determines what can be imported from Gramsci, at what stage of analysis, and with what modifications. PMN's diagnostic function — reading the configuration of structural suffering, visibility suppression, accumulation pressure, and institutional capacity — must precede the application of any strategic framework, including Gramsci's, because the choice of which strategic analysis to apply depends on what the diagnostic reveals. A context in which visibility suppression is high and direct confrontation capacity is low calls for the Gramscian mode of extended position warfare: building hegemonic infrastructure in civil society, developing organic intellectuals from within affected communities, constructing counter-narratives that can displace the interpretive frameworks through which suffering is currently misattributed. A context in which legitimacy has collapsed and institutional capacity for repression is fragmenting calls for a different mode — what Gramsci termed war of maneuver — in which the slow positional work shifts to rapid consolidation. The

diagnostic precedes the strategic prescription because the strategic prescription's appropriateness depends entirely on what the diagnosis reveals about the current configuration. A Gramscian strategy applied without this diagnostic foundation will be applied to idealized conditions rather than actual ones, which is the specific failure mode that produced the strategic disasters of orthodox communist parties throughout the twentieth century: technically sophisticated hegemonic analysis applied to configurations where hegemony was not the binding constraint, or positional warfare pursued in conditions that had already opened into the maneuver window. The implication for how PMN incorporates the Gramscian tradition is precise. Gramsci's analysis of hegemony, civil society, organic intellectuals, and positional warfare belongs in PMN's toolkit as strategic doctrine applicable when the diagnostic reveals specific configurations — high V suppression, intact repressive capacity, available civil society space — and inapplicable or counterproductive in other configurations. The tradition's absence of a biological foundation and evaluative criterion is corrected by PMN's Part III and Part IV, which supply what Gramsci's framework presupposes but does not establish. The tradition's tendency toward linear and volitional accounts of political change is corrected by the complex systems analysis of section 10.5, which shows why hegemonic transitions are characteristically non-linear and threshold-dependent rather than the product of accumulated positional advantage alone. With these corrections in place, the Gramscian analytical vocabulary — hegemony as V suppression mechanism, organic intellectuals as visibility infrastructure, civil society as the terrain of counter-power accumulation — is deployable within PMN's framework without importing the tradition's foundational deficits. PMN diagnostic interface with Gramscian strategy: (1) V suppression assessment — is the primary constraint on structural change the visibility of suffering rather than its absence? If yes, hegemonic analysis is the relevant tool: which institutional sites are producing which interpretive frameworks, who are the organic intellectual producers of those frameworks, and what alternative interpretive infrastructure is available or constructable? (2) Capacity configuration — does the current configuration support positional warfare (intact repressive capacity, available civil society space, high V suppression) or maneuver (fragmented repressive capacity, collapsed legitimacy, opened political window)? The strategic mode is determined by the diagnostic, not by ideological preference. (3) Biological floor check — which populations are experiencing structural suffering that hegemonic arrangements are currently rendering invisible? The biological floor check prevents the analytical error of treating hegemonic stability as evidence of genuine satisfaction. (4) Non-linearity alert — is the current trajectory of pressure accumulation approaching a threshold at which positional strategies become counter-productive because they assume more time than the configuration allows? The threshold analysis of section 10.5 must run in parallel with hegemonic strategy to avoid the positional warfare error of preparing for a war that has already moved into a different phase. Seven Dimensions: What PMN Adds The claim that PMN is more advanced than Gramsci's framework requires precision about what "more advanced" means. It does not mean that PMN is more politically sophisticated, more strategically detailed, or more historically rich. Gramsci's notebooks contain decades of strategic thinking applied to specific conjunctures with a granularity that PMN's framework does not attempt to replicate. "More advanced" means something structural: PMN resolves problems that Gramsci's framework generates but cannot internally resolve, operates at a level of generality that makes it applicable to conditions Gramsci's framework was not designed to address, and provides analytical tools — diagnostic procedures, causal mechanisms, operational concepts — that Gramsci's framework does not contain and cannot be extended to contain without importing PMN's foundations. Seven dimensions make this concrete. First: a non-circular evaluative foundation. Gramsci's framework needs an answer to the question of why hegemonic arrangements are bad that does not presuppose a Marxist philosophy of history. When pressed, the Gramscian tradition tends to appeal either to class interest (hegemonic arrangements prevent the working class from recognizing its own interests) or to a humanist reading of human potential that floats without structural support. PMN's evaluative criterion — does this arrangement create or destroy the conditions under which embodied organisms can avoid structural suffering and develop genuine capacity? — does not

presuppose a theory of history, a theory of class, or a theory of human nature beyond what biology and social science can establish. It is, in the precise sense, more fundamental: it supports the Gramscian evaluation while being derivable from premises that the Gramscian evaluation cannot derive itself. Second: a theory of why hegemonic stability is not the same as genuine satisfaction. Gramsci's framework can explain how hegemony is produced and reproduced, but it has no mechanism for distinguishing between a population that accepts hegemonic arrangements because those arrangements are in fact adequate to their needs and a population that accepts them because visibility suppression has prevented the connection between suffering and its structural causes from becoming politically actionable. PMN's biological floor provides this mechanism: the independent assessment of structural suffering — deprivation of material conditions, betrayal of institutional arrangements, visibility suppression — against which the stability of hegemonic arrangements can be evaluated. A stable hegemony in a population with high R and B but also high V is not evidence of genuine satisfaction. It is evidence that V is doing the work that R and B would otherwise make unsustainable. This distinction is analytically unavailable within Gramsci's framework. Third: no

n-linear systems dynamics. Gramsci's account of political change is fundamentally positional and cumulative: hegemony is built piece by piece through extended engagement with civil society institutions, the war of position precedes the war of maneuver, and transitions are the result of accumulated advantage rather than threshold effects. PMN's complex systems analysis (section 10.5) shows why this model misses the characteristic dynamic of significant structural change: the non-linear relationship between accumulated pressure and institutional response, the role of triggering events that do not themselves cause change but catalyze transitions that pressure accumulation has already made structurally available, and the window dynamics that open suddenly and close quickly regardless of how much positional work has been done. These are not modifications of the Gramscian model. They describe a different causal structure that the Gramscian model systematically fails to capture, which is why Gramscian-influenced movements have consistently been caught unprepared by the openings of political windows they were theoretically building toward. Fourth: a capture-proof theory of institutional transformation. Gramsci's framework theorizes the capture of hegemony — the displacement of ruling-class common sense by counter-hegemonic frameworks — but has no systematic theory of the capture of the capturing movement itself. The twentieth-century pattern in which revolutionary movements succeeded in displacing one hegemony only to install a new one that reproduced the essential structure of the previous one — concentrated decision-making power over the conditions of becoming — is not analyzable as a Gramscian failure because Gramsci's framework lacks the mechanism. PMN's capture theory (section 7.3c) provides this mechanism: the five-stage process through which any institutional concentration of power generates the access asymmetries that produce capture, regardless of the ideological commitments of those who created the institution. This is not a criticism that can be made from within the Gramscian framework. It requires the institutional level of analysis that PMN's Part VII provides independently of the hegemonic level of analysis that Gramsci occupies. Fifth: a multi-level consequence methodology. Gramsci's strategic prescriptions — build organic intellectuals, construct a historical bloc, pursue the war of position — are evaluated by the Gramscian framework primarily at the level of their contribution to hegemonic displacement. PMN's second-order effects methodology (section 12.1) requires that any strategic recommendation be evaluated across at least three levels of consequence: direct effects on the target configuration, effects on adjacent configurations and actors, and effects on the structural conditions that make the strategy available in the future. This methodology generates systematically different assessments of Gramscian strategies than Gramsci's own

framework produces. A war of position that builds counter-hegemonic capacity in civil society may simultaneously reduce the urgency of the biological suffering that motivates participation, produce organizational structures that are later available for capture, or generate the institutional forms that will become the new hegemony rather than its displacement. The Gramscian framework cannot generate these assessments because it lacks the methodological infrastructure to require them. Sixth: an extended agent typology with drift mechanisms. Gramsci's primary analytical category for the agents of political

transformation is the organic intellectual: the figure who emerges from within a class or social group and articulates its interests and worldview in terms that can challenge the prevailing hegemony. PMN's section 10.10-10.10c identifies eleven distinct functional positions in the agent ecology of structural change, each with its own functional role, its own vulnerability profile, and its own characteristic drift mechanism — the process through which actors in each position progressively move away from their original orientation through the accumulation of small adjustments rather than through deliberate defection. Gramsci's organic intellectual captures one and a half of these eleven positions (approximately the translator and part of the strategist). The remaining nine and a half — reformers, builders, witnesses, healers, connectors, sacrificial actors, skeptics, and the various watcher configurations — are not analytically available within Gramsci's framework, which means that movement ecologies analyzing themselves through a Gramscian lens will systematically fail to resource, protect, and maintain accountability for the majority of the functional roles their effectiveness requires. Seventh: applicability to configurations Gramsci's framework was not designed to address. Gramsci's notebooks were written in the context of industrial capitalism, the nation-state as the primary unit of political organization, and civil society institutions — unions, parties, churches, schools — as the primary terrain of hegemonic contestation. PMN's Part XVI analysis of technology as a structural variable shows that each of these assumptions has been partially or substantially modified by the developments of the past half-century: algorithmic mediation of information environments has changed the mechanisms through which hegemony is produced and reproduced; platform architectures have created civil society spaces whose ownership structure is radically different from anything Gramsci's framework contemplated; the acceleration of institutional change has compressed the timescales within which positional warfare can function. Gramsci's framework can be extended to these conditions by analogy — the algorithm as the new organic intellectual, the platform as the new civil society institution — but the extension is strained and generates analytical errors because the str

uctural differences are not merely surface differences in the medium through which hegemony operates. PMN's framework, built on the more fundamental level of biological conditions and institutional constraint structures, extends naturally to these configurations because its foundational concepts do not presuppose the specific institutional forms that Gramsci's analysis was built around. What Gramsci Still Does That PMN Does Not Intellectual honesty requires specifying what the Gramscian tradition does that PMN does not and cannot replace. The answer is: strategic specificity at the level of political organization and cultural production. Gramsci's analysis of how to build a historical bloc — what combinations of class interests, cultural traditions, and political programs can be assembled into a durable counter-hegemonic coalition — is a level of strategic detail that PMN's framework does not provide and does not attempt to provide. PMN identifies the structural conditions under which a historical bloc is needed and evaluates whether a proposed coalition meets the structural requirements for stability and effectiveness. It does not specify which organic intellectuals to cultivate, which civil society institutions to prioritize, or how to translate the framework's evaluative criterion into the specific narratives that can build coalition across the material fractures of a given social formation. That translation work is doctrine work, and Gramsci's tradition is among the most sophisticated resources available for doing it in configurations where hegemonic analysis is the diagnostically appropriate strategic mode. PMN is architecturally more advanced than Gramsci's framework in the sense that a foundation is more fundamental than the structure built on it. This is not a claim of superiority in the sense of greater sophistication, richer historical texture, or more detailed strategic guidance. It is a claim about what resolves what: PMN resolves the evaluative, methodological, and analytical gaps that the Gramscian framework generates but cannot internally close. Gramsci's framework provides the strategic vocabulary for operating within the space that PMN's diagnosis opens. Neither claim is diminished by the other. The relationship is hierarchical in the technical sense — one level grounds the other — without being dismissive in the political sense. What would revise PMN's assessment of these traditions. The engagement above states what each tradition identified correctly and where it fails the framework's criterion. Symmetry requires stating the reverse: what evidence would revise PMN's own conclusions here. Three findings would

do so. First, if the private-power blind spot attributed to libertarianism proved empirically tractable — if documented cases showed that competitive provision of essential goods reliably prevented the exclusion the distribution diagnostic predicts, at scale and across the range of conditions the population presents, rather

than only under favourable regulatory conditions — then the framework's assessment of A-type arrangements in section 11.4 would require revision, and the criticism of Hayek's institutional proposal would narrow from structural to conditional. Second, if the relation between meaning infrastructure and hierarchical arrangements proved systematically non-separable — not merely in documented cases, which section 10.7c already concedes, but as a general pattern in which meaning-generating capacity degrades wherever hierarchy is reduced — then the conservative institutional conclusion PMN currently treats as an overreach would become the default expectation, and the burden of proof would shift to the transformationist position. Third, if the friend-enemy distinction proved necessary rather than merely available — if political conflict demonstrably could not be managed through the contested-but-shared institutional forms PMN presupposes, and every attempt to do so collapsed into depoliticized management — then Schmitt's critique of proceduralism would graduate from a diagnostic warning to an architectural constraint on the framework itself. None of these is currently supported by the evidence the framework has examined. All three are stated here because the framework's own requirement (1.2) — that every claim specify what would change it — applies without exception to its assessments of the traditions it engages. A claim about a rival tradition that cannot state its own disconfirmation conditions is doing ideology, not analysis, regardless of which tradition it belongs to.

10.8 — The Organizational Form Problem

The tradition analysis in 10.7 generates a prior question that 10.9's account of counter-power accumulation cannot answer alone: what organizational form should counter-power take? The question is prior because the accumulation mechanisms in 10.9 — how capacity builds, how coalitions form, how windows are used — operate inside an organizational container whose structural properties determine whether the capacity accumulated will eventually be deployed against the arrangements it was built to contest, or redirected toward the reproduction of those arrangements under different nominal ownership. The organizational form problem is the question of what structural properties that container must have. The tradition analysis yields a consistent negative finding: every major organizational form that counter-power movements have adopted has failed the custodian problem in characteristic ways. The vanguard party concentrates interpretive authority in a leadership stratum whose structural position eventually diverges from the populations whose suffering it was organized to address. The mass party diffuses authority sufficiently to resist initial capture but becomes vulnerable to the slow drift that 10.12 describes — incremental accommodation that is individually rational at each step and collectively transformative of the organization's function. The council and soviet forms resist capture from above but face coordination failures at scale that produce either fragmentation or the consolidation of informal leadership with unaccountable authority. The network form maximizes distributed decision-making but generates insufficient organizational continuity to survive the suppression campaigns and resource droughts that all counter-power organizations encounter. No tradition has solved this set of constraints simultaneously. The question is not which tradition to adopt but what structural properties any viable organizational form must have, given what the failure record reveals. The framework does not provide a blueprint for organizational form. What it provides is a set of structural constraints derived from the failure analysis. First: interpretive authority — the capacity to determine what the organization's purpose is and whether specific actions serve it — must be distributed in ways that prevent any subgroup from monopolizing it without the cost of monopolization becoming visible to those who bear it. This does not require equal distribution; it requires that the mechanisms for contesting interpretive authority are accessible to those most affected by the organization's decisions, not only to those closest to the decision-making center. Second: accountability must run downward to the populations whose conditions the

organization claims to address, not only laterally to coalition partners or upward to funding and institutional relationships. The capture sequence in 7.3c proceeds precisely by inverting this direction: organizations whose accountability has been

redirected upward can no longer experience the gap between their stated function and their actual effects, because the populations who experience that gap are no longer the ones to whom the organization is answerable. Third: the organizational form must include mechanisms for detecting and naming its own drift — not as occasional self-criticism, but as a structural feature of how decisions are made and reviewed. Organizations that lack such mechanisms do not detect drift until it has become capture, at which point the detection itself is the most costly possible act. The tension between accountability and operational capacity is permanent and cannot be resolved by organizational design alone. Accountability structures that are robust enough to prevent capture impose real costs on the speed and coherence of decision-making. Organizations that optimize for operational effectiveness tend to centralize, which reproduces the concentration of interpretive authority that produces the failure modes the tradition analysis documents. The resolution available is not structural but disciplinary: the people inside the organization must apply the framework's primary diagnostic (1.2) to the organization itself — asking, at each stage of development, whether the organization is revising its conclusions based on evidence from the populations it serves, or insulating its prior commitments from that evidence. This is the same standard applied everywhere else. Its application to the organization one is inside is the hardest version of it. The organizational form problem has no permanent solution, only better and worse approximations. A counter-power organization that claims to have solved the custodian problem has either not been tested at sufficient scale and duration, or has already begun to misidentify its own capture as a solution. The appropriate posture is not the search for the correct form but the maintenance of the structural features — distributed interpretive authority, downward accountability, drift detection — that make the form correctable when it fails. Correctability, not correctness, is the realistic standard.

10.8b — Shadow Institutions and Parallel Infrastructure Doctrine

The organizational form problem in 10.8 addresses the question of what structural properties a counter-power container must have to avoid replicating capture dynamics under new ownership. A prior question is equally consequential: how does counter-power build organizational capacity before the conditions for formal institutional seizure are present? The standard answer — that movements build capacity by contesting formal institutions directly — is structurally incomplete. Formal institutional contestation requires resources, legal standing, and public visibility that are typically absent precisely when counter-power is weakest and its need for organizational development is greatest. Shadow institutions and parallel infrastructure are the organizational forms that address this gap: they build capacity outside formal institutional channels, through what the system they contest cannot easily suppress because it does not yet recognize as a threat. What Shadow Institutions Are Shadow institutions are organizational structures that perform functions formally assigned to existing institutions — welfare provision, dispute resolution, education, economic coordination, health care, security — but do so outside those institutions' formal authority. They are not clandestine in the conspiratorial sense; they are typically visible to the populations they serve. What makes them “shadow” is their relationship to formal authority: they operate in the gap between what existing institutions claim to provide and what they actually provide to specific populations, filling that gap through organizational capacity that the formal institution cannot or will not deploy. The Black Panther Party's Free Breakfast Program (1969–1980) is the paradigm case: a welfare function formally assigned to state institutions, performed by a counter-power organization for a population the state was not serving, building organizational capacity, legitimacy, and coalition width simultaneously. The Zapatistas' autonomous municipalities in Chiapas are the territorial form: parallel governance structures performing state functions — health, education, justice — in areas where the state's claim to those functions was not

matched by delivery. The Kerala cooperative model (10.9) is the economic form: parallel economic coordination that built the organizational substrate enabling the land reform coalition before the political window for formal institutional change opened. Parallel infrastructure is the material form of shadow institution-building: the physical and digital systems through which counter-power organizations coordinate, communicate, provide services, and maintain relationships independently of formal institutional infrastructure. This includes communication networks not owned by the institutions being contested, economic circuits that reduce dependence on the financial systems those institutions control, mutual aid networks that reduce the material vulnerability that makes individual

s susceptible to coercive pressure, and knowledge systems that develop and transmit the analytical capacity the counter-power organization requires. The political significance of parallel infrastructure is not primarily logistical. It is structural: every dimension of parallel infrastructure reduces the leverage that existing institutional arrangements hold over the populations they serve and the organizations contesting them. An organization whose communication depends entirely on corporate platforms, whose economic activity depends entirely on conventional financial systems, and whose members' material survival depends entirely on employment within existing institutional arrangements is vulnerable to pressure from those arrangements at every point. Parallel infrastructure reduces this vulnerability systematically, which is why its development is consistently resisted by the institutional arrangements it contests. The Doctrinal Principle The parallel infrastructure doctrine follows from the organizational form analysis in 10.8 and from the counter-power accumulation analysis in 10.9. It has three components. First, build organizational substrate before political windows open. The Kerala analysis demonstrates that the decisive variable in political window utilization is counter-power readiness (Rs) — the organizational capacity available to act on the window when it opens. Rs cannot be built in the moment of the window; it must be built during the periods between windows through organizational forms that survive in the absence of immediate political opportunity. Shadow institutions and parallel infrastructure are the primary vehicles for this inter-window capacity building. Second, build where the system's suppression capacity is lowest. Formal institutional contestation triggers the full defensive repertoire of the institutions being contested: legal challenge, resource denial, delegitimation, physical suppression. Shadow institution-building in the service gap — where existing institutions are not serving the population in question — operates below this trigger threshold. It builds capacity in the space where the system's incentive to suppress is weakest because its claim to that space is not yet being formally contested. Third, maintain the distinction between parallel infrastructure as organizational capacity and parallel infrastructure as separatist withdrawal. The function of shadow institutions is not to replace existing institutions but to build the counter-power capacity that makes genuine institutional transformation possible. Shadow institutions that become ends in themselves — that build organizational autonomy as a goal rather than as a means — reproduce the organizational form problem: they develop the same maintenance incentives that produce capture in formal institutions, simply in an extra-institutional register. The standard for evaluating shadow institution development is whether it is building the specific capacities that form

al institutional transformation will require: coalition width, organizational substrate, legitimacy with the populations whose cooperation the transformation depends on, and the analytical capacity to specify what transformation should look like. (See 10.9 counter-power accumulation; 7.3 custodian problem; 10.8 organizational form problem.)

10.9 — How Counter-Power Accumulates

Parts 10.1-10.6 describe the structural conditions under which existing arrangements become unstable and transformation becomes available. The parallel question — how actors who want to change existing arrangements build the capacity to do so — has its own analysis. The absence of this analysis produces a framework that is very good at explaining why systems fail — and comparatively silent about what makes it

possible for alternatives to emerge and consolidate. Counter-power is not the mirror image of the power it opposes. It cannot be built through the same mechanisms or sustained by the same resources. The actors who hold existing arrangements in place benefit from those arrangements — they have institutional positions, economic resources, and information advantages that are produced by the arrangements they are defending. Actors working against those arrangements start without those resources and must build capacity by different means. Understanding what those means are, and what structural conditions make them available, is as important for the framework's analysis of political change as understanding the mechanisms of system stress. The structural analysis of affective mobilization in 3.8b is the necessary background for understanding why counter-power accumulation is harder than interest analysis alone would predict. Populations with aligned material interests do not automatically coordinate. The affective substrate — what populations feel about their conditions, about each other, and about the plausibility of change — is a material variable that determines whether organizational capacity can be built on top of shared interest. Counter-power that ignores this variable will build organizations that are brittle under affective pressure: movements whose participants defect not because their material interests have changed but because the emotional conditions of sustained collective action have not been adequately maintained. The conditions for counter-power accumulation analyzed below are not only organizational and informational. They are also affective — which means that the analysis of what generates and sustains them must include the mechanisms 3.8b identifies, not only the organizational mechanisms this section develops. Organizational substrate as the primary condition. Counter-power that exists only as mobilization — as activated capacity in moments of visible crisis — depletes rapidly and cannot accumulate across cycles. The populations that experience the most structural suffering are also the populations with the least time, resources, and security to maintain organizational infrastructure during periods of lower intensity. The cause is material, not motivational. It is a structural consequence: the same conditions that generate the political motivation for change systematically undermine the organizational capacity to produce it. Counter-power that can persist across cycles — that maintains relationships, analytical capacity, and or

ganizational infrastructure between moments of mobilization — requires material conditions that are not automatically available to those with the greatest interest in change. Identifying and building those conditions is a prerequisite for any accumulation of counter-power that can outlast a single conjuncture. Fracture lines: manufactured versus material. Power that benefits from existing arrangements is sustained in part by divisions among the populations whose material interests would be served by changing those arrangements. Not all divisions are the same kind of thing, and the distinction matters for how counter-power can be built. Some divisions are produced by genuinely different material interests: agricultural smallholders and urban wage workers have different interests about food prices regardless of how they are politically organized, because the same price movement that benefits one harms the other. These divisions are material. They cannot be dissolved by narrative or solidarity appeals. They can only be managed through arrangements that provide genuine benefit to both constituencies — which may require creative institutional design but cannot be wished away. Other divisions are produced and maintained through narrative construction: workers in different ethnic or religious communities whose material interests are aligned are told that their primary identity is the ethnic or religious one, and that those in the other community are their real adversaries. These divisions are manufactured. They are real in their consequences but not in their material basis. Counter-power capable of accumulating beyond narrow constituencies requires the analytical capacity to distinguish manufactured from material fracture lines — and the political capacity to build coalitions that cross manufactured divisions while honestly acknowledging material ones. Information as counter-power infrastructure. One of the primary mechanisms through which existing arrangements maintain themselves is control over what is known about those arrangements — what their actual consequences are, who bears the costs and who captures the benefits, what alternatives have been tried and what they produced. Counter-power begins, in many historical cases, not with organizational mobilization but with the production and distribution of knowledge that power has an interest in suppressing: the documentation of conditions that

official accounts deny, the verification of claims that institutions contest, the analysis of consequences that beneficiaries attribute to other causes. Knowledge production here is a structural function within an ecosystem of change, not journalism as a profession. The conditions under which that function can be performed — the material resources, the protection from repression, the distribution networks — are themselves conditions for counter-power accumulation, not merely auxiliary supports. Threshold dynamics from the challenger side. The legitimacy depletion

analysis in 6.6 describes the threshold from the perspective of the system being challenged: when the cost of maintaining compliance exceeds the system's capacity to absorb it, the arrangement begins to fail. The parallel dynamic from the challenger's side is collective action threshold: the point at which enough actors simultaneously calculate that participation in a challenge is individually rational because enough others are also participating. This threshold is not a fixed property of a population. It is shaped by information about others' positions (which is why visible commitment and public declaration matter), by the perceived costs of repression (which is why the credibility of the protection that organizational infrastructure provides matters), and by the clarity of what participation requires and what it might produce. Counter-power that accumulates toward and past the collective action threshold creates conditions in which the legitimacy depletion on the system side and the capacity activation on the challenger side can reinforce each other — which is when political windows in 10.5b become available for transformation rather than mere crisis. The sequencing of demands and the capacity question. The demands that counter-power makes on existing arrangements are not separate from the capacity question. Demands that exceed what can be delivered given current capacity of the challenger — demands that outrun the organizational infrastructure, the coalition breadth, the information base — spend credibility without producing change and diminish the conditions for further accumulation. Demands that are well within what the existing arrangement can absorb without structural change produce symbolic wins that deplete organizational energy without altering material conditions. The productive zone is narrow and context-specific: demands that push against the material limits of what existing arrangements can concede while remaining within the organizational capacity of the challenger to sustain. Identifying it requires the same kind of specific material analysis that the framework applies to everything else: not a formula, but a set of questions about the actual configuration of capacity, legitimacy, and structural stress at the specific moment. This analysis of counter-power accumulation is structural, not tactical. It does not specify what organizational forms to build, what coalitions to prioritize, or what demands to make in any specific context. Those questions require the kind of detailed material analysis of specific conditions that belongs to doctrine, not framework. What the framework provides is the set of structural conditions that make counter-power accumulation possible or impossible across contexts — and the analytical lens for reading specific situations in terms of where those conditions are and are not present. The gap between structural analysis and tactical decision cannot be closed by the framework. It is closed by actors with the

knowledge of specific conditions and the judgment to translate structural analysis into action. → Formula: Counter-Power Capacity ($CP = O \times N \times R_s \times W \times (1 - R_p - I_f)$) formalizes the variables developed in this section. The compounding structure of the formula — not additive but multiplicative — reflects the argument here that weakness in any single variable constrains the ceiling of all others. See Part XV, 15.8. → Formula: Transformation Window ($Tw = C_c \times Fe \times Cr$) formalizes the interaction of cycle convergence, focal event probability, and counter-power readiness developed in this section. See Part XV, 15.11. The Tunisia and Belarus cases in Part XV, 15.14 demonstrate this formula applied to divergent outcomes from similar structural starting points. Non-Western Counter-Power: Kerala and Cochabamba as Applied Cases The Kerala model of development — land reform, universal literacy, and health infrastructure achieved through Left Democratic Front governance from the 1950s onward — is analytically instructive for the framework's analysis of counter-power accumulation not primarily because of its development outcomes, though those outcomes are substantial: infant mortality and literacy rates comparable to middle-income countries at a fraction of the income level, and a land reform that transferred cultivating rights from landlords to tenants across a state where landlessness had been the material condition of the majority. The instructive

element is the organizational mechanism. The CPI(M)'s capacity to build and sustain the coalition that made land reform politically possible across a state structured by caste hierarchy and religious division required decades of prior investment in parallel institutions: the Kerala Kisan Sabha (peasant union), the Centre of Indian Trade Unions, and the Kerala Shashtra Sahitya Parishad (People's Science Movement). These were not electoral instruments — they were organizational substrate (O) that maintained relationships, analytical capacity, and material benefit for coalition partners between mobilization cycles. The coalition they sustained crossed the manufactured fracture lines — between Hindu lower-caste Ezhava communities, Mappila Muslim communities, and Christian agricultural workers — that upper-caste landholding interests had maintained as the primary political identities of those constituencies. What distinguished Kerala from comparable structural configurations elsewhere in India where caste and religious divisions were equally present was not the absence of those divisions but the existence of organizational infrastructure that provided material alternatives to organizing politically through them. This is the mechanism the framework identifies: manufactured fracture lines are crossed not by narrative appeals to solidarity that override material interests but by organizational infrastructure that makes the manufactured division less politically salient than the shared material stake. Coalition width (W)

was achievable because the organizational substrate had made width materially sustainable before the political window opened. The Cochabamba Water War of 2000 demonstrates the threshold dynamics that this section identifies from the challenger side, in a context where the political window was not manufactured by long organizational preparation but opened rapidly in response to acute material crisis. When the Bolivian government awarded an exclusive water concession to Aguas del Tunari — a subsidiary of the Bechtel Corporation — in January 2000, tariff increases of 35 to 300 percent were implemented immediately, producing water bills equivalent to 20 percent or more of monthly income for the lowest-income households. What the government and its institutional partners did not anticipate was the Coordinadora de Defensa del Agua y de la Vida, a coalition that crossed a structurally significant manufactured fracture line: between formal factory workers and campesino irrigators whose traditional water systems — organized under *usos y costumbres* arrangements that predated the concession — were simultaneously threatened by the contract's territorial scope. The enabling condition for this coalition was not ideological solidarity but the preexisting organizational infrastructure of those irrigators' cooperatives, which provided an immediate organizational base and a concrete material stake requiring no manufactured narrative to activate. Over three months of escalating protests, roadblocks, and general strikes, the government's attempt to suppress the coalition through martial law — at the cost of six deaths among protesters — accelerated rather than halted mobilization, producing the threshold condition this section identifies: the point at which maintaining the repressive posture cost more than the state's political capacity could sustain. In April 2000, the government cancelled the Bechtel contract and restored public water management. The trajectory was characteristic of threshold non-linearity: months of sustained pressure followed by rapid state position collapse over days — exactly the pattern that 10.5 identifies as what happens when accumulated counter-power force meets a legitimacy depletion threshold on the incumbent side. The structural features that make both cases analytically useful here are three. First, neither case is a counter-power success produced by Western institutional resources — formal legal capacity, media access, or financial mobilization through philanthropy infrastructure — both were produced by organizational substrate built from below under material scarcity, which matters for the framework's cross-cultural validity claim. Second, in both cases the decisive variable was not the magnitude of structural suffering — which was extreme in both, and had been present for decades before the political outcomes occurred — but the counter-power readiness (Rs) at the moment the political window opened. The suffering did not produce the

outcome; the organizational preparation to act on the suffering did. Third, both cases involved crossing manufactured fracture lines through material rather than ideological mechanism: the organizational substrate made coalition width sustainable not by overriding identity divisions through narrative but by providing material stakes that the identity division could not match in political salience. The cross-cultural consistency

of these structural features — organizational substrate as enabling condition, manufactured fracture line crossing through material provision, threshold non-linearity — supports the framework's claim that its counter-power analysis is not institutionally specific to the Western political contexts in which most of the existing literature was developed. Fracture Lines: Material, Manufactured, and the Coalition Problem

Counter-power formation requires coalition building: the assembly of populations with different immediate interests and identities into coordinated collective action around shared structural objectives. Coalition building fails in predictable ways, and the most common failure mode is the misdiagnosis of the fracture lines that divide potential coalition partners. A fracture line is a division between populations that limits their capacity to act collectively. Some fracture lines are material — they reflect genuine differences in structural position that produce genuinely different interests in specific outcomes. Others are manufactured — they are produced and maintained by actors who benefit from the fragmentation of potential coalitions, through the strategic amplification of identitarian differences that do not track structural divergence. Misidentifying material fracture lines as manufactured produces coalition strategies that fail because they ask populations with genuinely different interests to act as if they had the same interests. Misidentifying manufactured fracture lines as material produces fragmentation that serves the interests of those who manufactured the division. The diagnostic for distinguishing material from manufactured fracture lines operates on two dimensions. The first is structural position correlation: does the division between populations correspond to a genuine difference in structural position — in ownership of productive assets, in position within labor markets, in relationship to state coercive capacity, in exposure to ecological risk — that produces genuinely different interests in the outcome of the contested question? A division that does not correspond to structural position differences is a candidate for manufactured status: it organizes populations in ways that do not reflect their structural interests. The second dimension is beneficiary analysis: who benefits from the maintenance of this division as an organizing principle of political identity? If the primary beneficiaries of the division are actors at the top of the structural hierarchy — those whose interests are served

by the fragmentation of populations who would otherwise be natural coalition partners — this is a strong indicator that the division is being actively maintained rather than simply expressing underlying structural differences. The analysis requires care because the material/manufactured distinction is not binary and many fracture lines have both dimensions simultaneously. Religious divisions between working-class populations may reflect genuine differences in meaning infrastructure and political orientation (material in the specific sense of structurally consequential) while being simultaneously amplified by elite actors who benefit from the fragmentation of class-based coalitions (manufactured in the sense of strategically maintained beyond its structural basis). Ethnic divisions may reflect genuine histories of differential treatment by state institutions (material) while being simultaneously exploited by political entrepreneurs who profit from ethnic mobilization (manufactured). The analytical question is not whether a division is purely material or purely manufactured but what proportion of its current political salience reflects genuine structural divergence and what proportion reflects strategic amplification — because this proportion determines what coalition strategies are available and what they require. Coalition building across fracture lines that have genuine material components requires explicit acknowledgment and negotiation of the divergent interests rather than their suppression under a fictitious unity. A coalition between populations with different structural positions can be sustained when each partner receives something from the coalition that they cannot obtain without it, and when the coalition's objectives do not require any partner to act systematically against their structural interests. Coalitions that suppress material divergences rather than acknowledging and negotiating them are fragile: the suppressed divergences re-emerge under stress, and the perception of betrayal that follows — that one partner's interests were subordinated to another's under the cover of solidarity — is more damaging to future coalition capacity than the divergence itself would have been if acknowledged from the start. The strategic implication is that honest mapping of material fracture lines is a prerequisite for durable coalition building, not an obstacle to it. The manufactured fracture line requires a different strategy: demystification rather than negotiation. If a division between populations does not reflect genuine structural divergence but is

maintained by strategic amplification, the counter-strategy is to make the maintenance mechanism visible — to show, with specific evidence, who benefits from the division, through what institutional channels the amplification occurs, and what the populations on both sides of the division would gain from coordination that the division currently prevents. This is not a purely cognitive operation: po

pulations are not held in manufactured divisions by ignorance alone, but by the affective investments that the division has generated through its history. Demystification must address both the structural analysis and the affective dimension — acknowledging the genuine grievances and historical experiences that make the division feel real while providing an account of why the political expression of those grievances in terms of the manufactured division serves the interests of those who manufactured it rather than the interests of those who feel it. Fracture line diagnostic for coalition assessment: (1) Structural position correlation — does this division correspond to genuine differences in structural position that produce different interests in the contested outcome? Map the structural positions of the populations on each side of the division. If structural positions are similar, the division is a candidate for manufactured status. (2) Beneficiary analysis — who benefits from the maintenance of this division as an organizing political principle? If primary beneficiaries are at the top of the structural hierarchy, this is a manufactured fracture line indicator. (3) Historical production — was this division produced or amplified at identifiable historical moments by identifiable actors with identifiable interests? Manufactured fracture lines have histories of production that can be documented. (4) Coalition cost-benefit — what would the populations on each side of this division gain from coordination, and what would they lose? If both sides would gain more from coordination than from separation, the division is being maintained against their structural interests — which requires manufactured maintenance to persist. A related diagnostic for the counter-power accumulation context: Epistemic infrastructure for threshold assessment: claims that a transformation threshold has not yet been reached require the same epistemic infrastructure that the framework requires for any structural assessment (1.8–1.12): independent data sources, adversarial review, and standing for affected populations. Where such infrastructure is absent, claims of “not yet threshold” made by actors who benefit from the current arrangement carry a standing presumption of motivated reasoning. The diagnostic question is not whether the claim is made in good faith — it usually is — but whether independent infrastructure exists that would detect motivated error. Where that infrastructure is absent or captured, treat threshold assessment by incumbents as structurally unreliable, regardless of analytical sophistication.

10.10 — The Two Modes of Struggle

Two terms in this section require explicit framing. 'Hero' refers to actors operating through publicly legitimate channels — institutions, recognized authority, visible accountability structures — regardless of the moral quality of their choices. 'Vigilante' refers to actors operating outside those channels, in spaces the legitimate system cannot or will not reach, regardless of the justifiability of their methods. Neither term carries moral endorsement or condemnation. Both describe functional positions in a political ecosystem, not moral types. Each can be legitimate under specific conditions; each can become counterproductive when conditions change. Every movement that has produced genuine structural change has operated through at least two fundamentally different modes — and the most successful have maintained both simultaneously without sacrificing one for the other. The first mode operates at the surface: it builds public legitimacy, maintains a posture that the majority can accept, and works through channels that the existing system recognizes. The second operates where the first cannot reach: it fills gaps that cannot be filled by actors who must protect their public image, using methods that cannot be openly adopted without destroying the legitimacy the first mode depends on. The difference between these two modes is not a moral difference. It is a functional difference produced by genuine adversarial conditions. The system being confronted does not rely only on coercion — it uses legitimacy, narrative control, and institutional capture to determine what is considered acceptable and what is considered deviant. Actors who need access to the space the system

provides must play by the system's rules, at least in part. Actors who do not need that access can operate by different rules — at different costs, and with different capacities. One mode is the clean face of the struggle. It builds hope, provides inspiration, and maintains the legitimacy that keeps the broader movement credible to populations not yet ready for the more radical. The other does the work that cannot be done on camera — willing to absorb stigma, willing to be seen as operating dirty, precisely because of the understanding that if its methods were widely adopted before conditions permit it, the resulting backlash would destroy more than it protects. The functional interdependence is the core of what makes the two-mode arrangement both effective and unstable. The surface mode needs the pressure that the deeper mode generates: without it, the system has no reason to negotiate, no reason to treat the surface actors as representing something that must be accommodated rather than something that can be safely ignored. The deeper mode needs the legitimacy that the surface mode maintains: without a credible public face for the broader movement, the harder methods produce backlash rather than accommodation, and the populations whose support is required for any durable change withdr

aw. Each mode is less effective without the other, and this mutual dependence is what creates the political space in which genuine structural change becomes possible — a space that neither mode could open alone. The arrangement is destroyed from two directions. From the surface side: purity politics — the demand that all actors within the broader movement publicly endorse only methods that can withstand full visibility, disavowing the harder methods even when those methods are what created the leverage that made the surface actor's position negotiable. Purity politics is often sincerely held, and it is sometimes strategically correct in specific contexts. But as a general principle it destroys the two-mode arrangement by eliminating the protection that public distance provides. Once the surface actor and the deeper actor are explicitly and publicly unified, the surface actor loses the institutional access their distance protected, and the deeper actor loses the independence that made them effective. Both become the same thing, which is weaker than either was separately. From the deeper side: opportunism — the use of the freedom that operating outside visibility provides for purposes that are not connected to the larger structural objective. Actors who operate without accountability, without feedback loops from the consequences of their methods on the populations they claim to serve, can drift from functional necessity to self-perpetuation to harm. When this happens, the surface actor's public distance is no longer a functional fiction protecting genuine coordination. It is accurate: the deeper actor has stopped serving the purpose that justified the arrangement. The surface actor who continues to provide cover for an actor that has drifted into opportunism is no longer maintaining a productive tension — they are protecting an actor whose continued operation is damaging to the conditions for change. The distinction between these two cases — functional distance and accurate distance — is one of the most consequential analytical judgments in political practice, and it cannot be made from the outside. The sequencing question is when each mode should lead. In conditions where the system retains legitimacy and organizational capacity — where transformation pressure is positive but has not approached the threshold for non-linear change — the surface mode must dominate publicly, because the primary work is the accumulation of legitimacy for the broader movement and the erosion of legitimacy for existing arrangements. The deeper mode provides pressure but cannot be the face of the challenge without accelerating the backlash that consolidated arrangements can generate when they still retain the capacity to do so. As legitimacy depletes and transformation pressure approaches threshold, the balance can shift: the deeper mode's methods become more visible without producing the same backlash, because the populations whose support the arrangements depended on h

ave withdrawn it. Misjudging this sequencing — either leading with harder methods before conditions permit, or failing to make space for them when conditions require — is among the more consequential errors available to a movement operating across this divide. What follows in 10.10b develops the specific mechanism through which the public dimension of this relationship is managed: the performative hostility that protects both sides of the arrangement by making their independence credible to the audiences each depends on. The analysis there presupposes the functional logic developed here — that the hostility is not the

relationship but the surface through which the relationship is protected.

10.10b — Performative Hostility and Coordination Behind the Surface

In public, the two modes must appear hostile to each other. Not merely as performance — but as a functional necessity born of asymmetry in how the system responds to threat. Actors operating within institutions cannot openly affiliate with those operating outside without losing the institutional access that makes them effective. Actors operating outside cannot be seen as endorsed by those inside without losing credibility in communities that have no trust in institutions. The visible hostility is protection for both. When a publicly recognized actor condemns harder methods, they are not only protecting their own image — they are providing cover for those condemned, making them appear genuinely independent and genuinely threatening to all sides. Even genuine condemnation can function this way: two actors who genuinely disagree on method but are moving in the same structural direction provide each other protection without ever having planned it. Below the surface, genuine coordination operates not through formal hierarchy but through shared analysis of conditions and objectives. Not who gives orders to whom, but whether the analysis of material conditions is coherent enough across different actors that their independent actions produce complementary effects. This form of coordination is far more resilient than centralized coordination — because it has no single point of compromise that can bring the entire structure down. The coordination can fail in several ways, each with distinct consequences. The most common failure is asymmetric defection: actors inside the institution, under pressure from their institutional environment, distance themselves from outside actors in ways that go beyond performance — denying the legitimacy of outside pressure, actively working to suppress it, or framing it as the problem that their institutional work is solving. This is not always cynical. The structural incentives of institutional positions genuinely reward this framing, and actors who have spent years navigating those incentives can come to believe it. The result is that the visible hostility becomes actual hostility, and the coordination it was supposed to cover disappears. The second failure mode runs in the opposite direction: actors outside the institution, seeking broader legitimacy or resources, allow the performance of independence to be replaced by actual alignment with institutional actors. The movements and outside voices that were supposed to maintain pressure on institutions instead become pipelines for institutional messaging — organizations that function as advocacy fronts for positions already held by sympathetic actors inside, producing the appearance of outside pressure while generating none of its actual force. Populations that need genuine outside pressure get managed dissent instead. Both failure modes share a structural cause: the absence of explicit mutual understanding of the logic being implemented. When the coordination is tacit — when

inside and outside actors are responding to structural incentives without articulating the logic — it is vulnerable to the pressures that each environment creates independently. The conditions under which the two modes remain genuinely distinct while remaining in productive relation require more than good intentions. They require organizational cultures on both sides that understand the structural pressures pulling toward each failure mode and have developed the internal accountability mechanisms to resist them. The arrangement is fragile in specific ways that the actors involved must understand to maintain it. The first failure mode is sincerity collapse: actors who have been performing hostility for so long that they forget the performance, and genuine animosity replaces the functional fiction. This happens most readily under conditions of sustained repression, where the costs of maintaining the appearance of coordination are high and the temptation to clarify distances for self-protection is real. The second is premature visibility: coordination that surfaces before either side has achieved what the coordination was meant to achieve. When the inside actor's association with outside methods becomes visible, they lose institutional access; when the outside actor's connection to inside strategy becomes visible, they lose credibility with the communities whose trust requires independence from the system. The third failure mode is strategic divergence — when actors who were moving in the same structural direction stop doing so, because conditions change in ways

that affect one more than the other, or because success at one level makes the other level's agenda seem less urgent. Movements that achieve institutional incorporation of their demands sometimes find that those who remain outside become inconvenient rather than useful. The inside actor who has been relying on outside pressure to justify internal reform loses the pressure precisely at the moment when they have enough institutional standing to act — and may not. The framework's account of this dynamic is not optimistic: strategic divergence is the default outcome of success, not the exception. Sustained coordination across the inside-outside boundary past the point of partial achievement requires explicit maintenance, and it rarely receives it. The conditions under which surface coordination becomes counterproductive: The prohibition on visible coordination is not absolute. It holds under conditions where the system can use visible coordination to delegitimize both parties — where institutional actors can be accused of endorsing illegal or disruptive methods, and where extra-institutional actors can be accused of endorsing the institutional system's legitimacy. Under conditions where both parties have already been delegitimized by the system — where the accusation of coordination adds nothing to the existing characterization — the prohibition relaxes and visible coordination may be strategic

lly available. Accurately reading which condition obtains at a given moment requires both the analytical capacity to assess the system's current legitimacy and the coalition intelligence to know how the accusation would land with the specific audiences whose support is most strategically significant. The deeper structural point: performative hostility is a form of strategic communication that must be maintained with the same discipline the framework applies to all strategic communication. Actors who perform hostility without genuinely maintaining the strategic distance that makes the performance credible will be read correctly by the system they are trying to deceive — the system's interest in identifying false-flag operations is as strong as the ecosystem's interest in maintaining plausible deniability. The performance must be grounded in real organizational separation and real differences in method, or it will fail in exactly the situations where it is most needed.

10.10c — The Extended Agent Typology: Beyond the Hero-Vigilante Binary

The binary of hero and vigilante — the actor who amplifies structural analysis against personal cost versus the actor who performs solidarity while protecting positional interests — captures something real about the ecosystem dynamics analyzed in 10.10 and 10.10b. But a binary is analytically insufficient for the full range of positions that populate actual political ecosystems. The positions that consistently determine whether movements maintain counter-power capacity past the point of partial institutional success are not primarily the heroic or the vigilante — they are the positions whose actors have less dramatic but structurally more consequential functions: managing organizational continuity, controlling resource flows, holding the institutional relationships that movements depend on, and translating between the inside and outside positions that any sustained effort requires.

10.10c-i — The Seven Additional Durable Positions

The Anchor: The actor whose primary function is organizational continuity — maintaining the infrastructure through which collective action is sustained between moments of heightened activity. The anchor is often invisible precisely because their work is most visible in its absence: when a movement loses anchors, organizational capacity degrades faster than membership numbers would predict. The anchor's structural vulnerability is captured through the institutional relationships their continuity function requires them to maintain. The Translator: The actor who moves between distinct social worlds — between grassroots communities and institutional power, between different ideological traditions, between technical expertise and affected populations — making each legible to the other without fully belonging to either. The translator's strategic value is high in conditions where coalition breadth is required; their structural

vulnerability is the credibility cost of operating in multiple registers, which makes them targets of authenticity challenges from all directions. The Resource Broker: The actor with access to material resources — funding, institutional infrastructure, technical capacity, professional networks — who directs those resources toward counter-power functions. The resource broker's function is indispensable and their vulnerability is definitive: resource dependency creates accountability relationships that are structurally prior to movement accountability, and those relationships consistently shape broker behavior in directions that may or may not align with the movement's original function. The resource broker is the position most directly subject to the capture sequence analyzed in 7.3c-i. The Archivist: The actor who maintains organizational memory — documentation of what was tried, what worked, what failed, what decisions were made and why, what the conditions were at each stage of development. Organizational memory is a counter-power resource because it is the mechanism through which movements learn from their own history rather than repeating the same errors under different conditions. The archivist is systematically undervalued in movements oriented toward immediate action and systematically targeted in suppression efforts because destroying organizational memory degrades future capacity more than any other single intervention. The Witness: The actor whose function is documentation — producing the evidentiary record that makes structural suffering visible, that creates the historical record through which movements can be held accountable and through which their successes can be attributed and reproduced. The witness operates at the intersection of epistemic infrastructure and political accountability. Their vulnerability is the gap between documentation and impact: producing accurate records of structural harm does not automatically translate into political visibility, and the effort required to maintain documentation under condi

tions of sustained suppression produces burnout that organizational structures rarely support adequately. The Sustainer: The actor whose function is the material and relational infrastructure of the people doing the political work — care work, conflict mediation, crisis support, the maintenance of the social relationships through which collective action is possible. Sustainers are the most consistently invisible and most consistently underfunded position in any movement ecosystem because their function is most legible in its absence: when sustainers are not present, burnout rates increase, internal conflict escalates, and organizational capacity degrades. The sustainer's work is the individual-level analog of the organizational continuity function the anchor provides at the institutional level. The Defector: The actor who exits a position within the dominant arrangement and brings with them institutional knowledge, credibility, and sometimes resources that serve the counter-power function. The defector's value is strategic information and legitimacy transfer; their vulnerability is the credibility challenge from both directions — the dominant arrangement they left treats them as traitors, and the movement they join has reason to assess their defection as strategic rather than principled. The defector is most valuable when their institutional knowledge is specific and current and when the defection itself is costly enough to be credible.

10.10c-ii — The Watcher Problem: Coordination From Above and Behind

All seven positions face a version of the watcher problem that 10.14 analyzes at the institutional level: the actors whose function is to maintain accountability within the ecosystem are themselves positions whose occupants have interests and are subject to the same capture dynamics as any other position. There is no position within the ecosystem that is structurally exempt from the incentive pressures that produce drift. The implication is not that accountability is impossible but that it requires distributed redundancy — multiple positions with overlapping accountability functions and genuine capacity to act on each other's drift, rather than a single designated accountability role that becomes the system's single point of failure.

10.10c-iii — Spontaneous and Transitional Actors

Not all actors in a political ecosystem occupy durable positions. Crises, focal events, and periods of heightened political activity mobilize actors whose engagement is intense but bounded — who contribute counter-power capacity during the period of mobilization and then return to positions that are primarily outside the ecosystem. These transitional actors are a structural resource that sustained movements must be designed to engage without becoming dependent on: the organizational infrastructure that functions only when spontaneous mobilization is present has failed to build the durable capacity that produces sustained structural change. The relationship between spontaneous and organized counter-power is analyzed in 10.9; what 10.10c adds is the recognition that spontaneous actors occupy real ecosystem positions with their own structural properties — including their own characteristic drift trajectories when spontaneous engagement transitions into organizational involvement. The typology is diagnostic, not prescriptive. The function of mapping ecosystem positions is not to assign actors to roles but to identify what functions are being performed, by whom, with what structural vulnerabilities, and what the current distribution of those functions suggests about the ecosystem's capacity to sustain counter-power past the next threshold event. An ecosystem with strong heroes and vigilantes but weak anchors, sustainers, and archivists has visible capacity and fragile depth. That is the most common configuration — and the one most vulnerable to the capture sequence that follows partial institutional success.

The Insufficiency of Two Types The seven positions can be organized along two axes. The first axis is system position: whether the actor operates primarily within established institutions, primarily outside them, or in explicit translation between the two. The second axis is orientation transparency: whether the actor's actual orientation — who or what they are ultimately serving — is clear to themselves, to allies, and to the ecosystem as a whole, or whether that orientation is obscured, either because the actor has been steered without their awareness or because they are consciously performing a role that diverges from their actual strategic position. The intersection of these two axes generates the positions that the hero-vigilante binary collapses into two.

The Seven Additional Positions The second is the builder of parallel institutions: an actor who neither tries to reform existing institutions from within nor confronts them directly from outside, but constructs alternative institutional infrastructure — cooperative economic structures, alternative media, mutual aid networks, independent educational institutions, community health systems — that can function as both a material resource for the ecosystem and a demonstration that alternative configurations are viable. The builder's strategic function is often underestimated because it is slow,

unglamorous, and produces results on timescales that feel irrelevant to movements operating in political time. Its structural importance is that parallel institutions reduce the dependence of the population on the institutions being contested, which changes the balance of power in ways that neither internal reform nor external confrontation alone can achieve. The third is the witness and chronicler: the actor whose primary function is accurate documentation and transmission of what is happening — to the public sphere, to the historical record, to the movement itself. The witness operates outside the hero-vigilante dynamic because their effectiveness depends on the perception of independence from both positions: a chronicler who is visibly aligned with either the reform or the confrontation axis loses the credibility that makes their documentation useful. The witness's vulnerability is isolation: their function requires audience, and audience requires distribution infrastructure that the witness typically does not control. A witness without distribution is a recorder without readers. The fourth is the translator: the actor who moves between distinct communities — across class, across culture, across political tradition, across the system-outside boundary — and whose primary function is maintaining legibility between populations whose frameworks for understanding the situation are sufficiently different that coordination between them would otherwise be impossible. The translator is not a hero who happens to know multiple languages. They are a position in the ecosystem whose function is specifically the management of the fracture lines that section 10.9 analyzes: keeping material fracture lines from hardening into irreparable divisions while accurately identifying which divisions reflect genuine structural divergence and which reflect manufactured fragmentation. The translator's vulnerability is the accusation of bad faith from both sides: actors who demand purity of alignment will consistently

characterize translation as betrayal. The fifth is the strategist without public position: the actor who has deep analytical clarity about conditions and objectives, who shapes the orientation of multiple other actors through analysis and advising, but who deliberately avoids occupying a public position that would make them a target, reduce their analytical flexibility, or limit the range of actors they can influence. The strategist is distinct from the watcher — discussed below — in that their orientation is clear and aligned with the ecosystem's purposes; the concealment is tactical rather than ambiguous. Their vulnerability is the same as that of the ecosystem's analytical infrastructure generally: if their orientation is not known, their influence cannot be evaluated by the distributed accountability mechanisms that section 10.14 describes. The sixth is the sacrificial actor: the actor who takes on a position of high visibility and high personal cost — oc

cupation of a contested public role, a legal challenge that will fail but establish a precedent, a public confrontation that will result in personal consequences — whose primary function is not winning the immediate contest but changing the terms of subsequent contests by demonstrating that the system's response to challenge is as costly as the challengers have claimed, or by generating the public attention that shifts the horizon of what is politically possible. The sacrificial actor's function is structurally important and personally devastating, and the ecosystem that uses sacrificial actors without acknowledging the nature of that use — that treats them as heroes whose cost-bearing is incidental rather than central to the function they are performing — is producing a form of exploitation that will eventually destroy the trust that makes the ecosystem function. The seventh is the skeptic inside the movement: the actor who genuinely shares the ecosystem's orientation but consistently challenges its analysis, contests its strategic assessments, and refuses the social solidarity pressures that push toward groupthink and premature closure on contested questions. The skeptic is the most consistently undervalued position in movement ecologies because the function they perform — maintaining the adaptive uncertainty that section 1.6b identifies as the epistemological requirement for frameworks that remain useful — is experienced as friction rather than contribution. An ecosystem that expels or marginalizes its skeptics has optimized for internal coherence at the cost of analytical accuracy, and will begin producing strategic errors that are invisible internally until they produce external consequences. The eighth is the healer: the actor whose primary function is the maintenance of the human capacity of the ecosystem itself — the care of activists who have been traumatized, burned out, imprisoned, or otherwise rendered temporarily or permanently unable to function. The healer is the position most systematically undervalued in movement ecology because their output is the absence of a negative — the activist who did not break, the organizer who did not leave, the community that did not fragment under sustained repression — rather than a visible positive. An ecosystem that has no healer function treats its participants as renewable resources rather than as the irreplaceable carriers of relationships, knowledge, and commitment that they actually are. The healer's drift vulnerability is double: burnout from absorbing others' distress without reciprocal care, and the progressive professionalization of the role into therapeutic service provision disconnected from the ecosystem's strategic purposes. The ninth is the connector: the actor whose primary function is maintaining legibility and relationship between groups that are structurally separated — by geography, by issue focus, by political tradition, by organizational history, by language — and who would ot

herwise remain isolated islands of counter-power unable to coordinate at the scale that significant structural change requires. The connector is distinct from the translator. The translator works with ideas — converting frameworks across linguistic and cultural registers. The connector works with relationships — maintaining the trust and mutual recognition between groups who do not share enough common vocabulary for the translator's work to begin. Their vulnerability is positional: because they belong fully to neither group they connect, they can be accused of bad faith by both and lack the organizational base that would give them independent standing if that accusation is made. The tenth is the skeptic internal to the movement, analyzed in the previous section. Eleven positions in total. The number is not the point. What matters is that the ecosystem's map of itself includes enough positions to make visible the full range of functions that structural change actually requires, and to assign each its characteristic vulnerabilities and drift mechanisms, so that

resource allocation, accountability structures, and strategic coordination can be calibrated to the actual composition of what exists rather than to the dramatic minority of positions that generate most of the narrative.

Typology of Transient Actors: Momentum, Sympathy, and Affected Populations

The first category is momentum actors: individuals and groups who engage in the ecosystem in response to a specific precipitating event — a particularly visible act of state violence, a sudden economic shock, the collapse of an institution, an election result — and who bring substantial energy, scale, and public attention that the durable ecosystem could not generate on its own. Momentum actors are not activists. Most will not remain in the ecosystem after the precipitating event recedes from public attention. Their strategic function is specifically the production of scale and urgency at moments when the durable ecosystem alone cannot create sufficient pressure for structural response. Their mismanagement takes two forms: the durable ecosystem's failure to have prepared the organizational infrastructure and concrete demands that would allow momentum to be converted into structural outcomes before it dissipates; and the attempt to retain momentum actors through escalating demands on their commitment that they were never positioned to meet, which produces both the departure of the momentum actors and the discrediting of the durable ecosystem among the broader population they briefly represented. The second category is sympathy actors: individuals and institutions who do not participate in the ecosystem's activities but who provide material support — financial resources, professional skills, public endorsements, access to platforms or networks — and whose support is conditioned on the ecosystem's maintaining a position compatible with their own. Sympathy actors are not allies in the strategic sense. The

y are providers of resources whose continued provision depends on the ecosystem's not crossing lines that the sympathy actor has defined. Their material function is significant — many sustained ecosystems could not operate without the resources that sympathy actors provide — and their strategic function is negative: they constrain the ecosystem's range of motion in exchange for the resources they provide, and will withdraw at the moments when the ecosystem is under the most pressure to move in directions incompatible with their conditions. The strategic question is not whether to have sympathy actors but whether the ecosystem's dependence on their resources has reached the level at which that dependence has become a structural constraint on strategic decision-making. The third category is the affected population itself: those who directly experience the structural arrangements the ecosystem is contesting, who are the nominal beneficiaries of what the ecosystem is attempting to produce, but who may not be participants in the ecosystem's activities and whose relationship to it ranges from active support to indifference to active opposition. The affected population is not a constituency to be mobilized. It is the source of the legitimacy on which the ecosystem's claims rest, the ultimate target of the structural changes the ecosystem seeks, and the judge — in the most material sense — of whether the ecosystem is producing outcomes that correspond to their actual needs rather than to the ecosystem's internal model of what their needs are. Treating the affected population as an object of the ecosystem's strategy rather than as a distinct and partially sovereign actor is the error that produces the specific failure mode in which a movement achieves political success while producing institutional configurations that the affected population experiences as a new form of the arrangements it was contesting. The relationship between durable and transient actors in the ecosystem is a resource management problem that most movement strategy treats as a mobilization problem, and thereby mismanages. Mobilization asks: how do we get more people involved? Resource management asks: what functions do transient actors perform that the durable ecosystem cannot, at what moments do they perform them, what organizational preparation is required to convert their contribution into durable outcomes, and what relationship must the durable ecosystem maintain with them between moments of activation to ensure that activation remains possible? The mobilization frame generates strategies optimized for the acute phase — the demonstration, the strike, the occupation. The resource management frame generates strategies optimized for the full cycle, including the preparation before and the consolidation after, which is where most of what is gained or lost in structural change actually occurs.

Two Modes of Watcher Coordination: Hostile External and Aligned Internal

In the first version,

the watcher is an external hostile actor who has achieved influence over actors in multiple positions of the ecosystem simultaneously. The watched actors do not know they are being steered — they believe they are acting from their own analysis and their own values — but the inputs to their analysis have been systematically shaped: the information they receive, the framings they use, the allies they trust, the threats they perceive have all been calibrated to produce behaviors that serve the watcher's purposes rather than those of the ecosystem. This is not conspiracy theory. It is the structural logic of any serious counter-intelligence operation directed against a movement, and it does not require that every actor be compromised or that the manipulation be total. Partial steering of a few well-positioned actors is sufficient to produce significant distortions in the ecosystem's strategic orientation, because those actors influence others who are not themselves compromised. In the second version, the watcher is an internal strategic coordinator — aligned with the ecosystem's purposes — who is consciously managing the relationship between multiple actor positions without all of those actors having a clear picture of that coordination. This version is not malicious, but it produces its own structural problem: actors who are performing a role in a coordinated strategy they do not fully understand cannot evaluate whether the strategy is actually serving the purpose they signed on for. The internal watcher whose strategy is genuinely sound produces a more effective ecosystem in the short term. They also produce an ecosystem whose accountability mechanisms are progressively undermined, because the distributed knowledge required for distributed accountability has been selectively withheld. The distinction between hostile external watcher and aligned internal coordinator is not always clear from inside the ecosystem, because the behavioral signatures are similar: both versions produce actors whose public positions and private orientations diverge, both involve coordination that is not transparent to all participants, and both involve an actor who shapes the ecosystem from a position that is not itself accountable to the ecosystem's normal governance mechanisms. The diagnostic question is not "is there a watcher" but "whose purposes does the watcher's coordination serve, and do we know the answer?" The ecosystem that cannot answer that question with reasonable confidence is operating with a significant blind spot regardless of which version of the watcher problem applies.

Unconscious Steering: The Manipulation Vector The ecosystem's defense against unconscious steering is not the identification of compromised actors — which is both unreliable and corrosive of the trust that makes the ecosystem function — but the structural maintenance of analytical redundancy. An ecosystem in which the same analytical conclusions are being reached by multiple independent

actors through genuinely different analytical pathways is more resistant to the distortion of any single pathway than an ecosystem in which most actors are drawing their conclusions from a common source, even if that source is apparently trustworthy. The epistemic infrastructure analysis of section 1.8 applies directly: the structure of how analysis is produced and distributed within the ecosystem is a strategic resource that hostile watchers will specifically target, because distorting the inputs to analysis is more efficient than attempting to distort the conclusions of many independent analyzers.

Conscious Role-Playing: The Double Position The conscious double position — the actor who is deliberately performing a role that diverges from their actual strategic orientation — exists in both aligned and hostile versions, and the ecosystem must be able to distinguish between them while also accepting that the distinction cannot always be made with certainty. The aligned double position is the logic of the performative hostility analyzed in section 10.10b: the institutional actor who publicly denounces the movement while privately supporting its objectives, or the movement actor who publicly condemns tactics they privately regard as necessary. This is role-playing in the service of coordination — a performance designed to protect the actor and maintain their effectiveness in a position that would be compromised by transparency. The hostile double position is its mirror: the actor who publicly performs alignment with the ecosystem's purposes while actually serving an external interest. The hostile double is distinguished from the unconsciously steered actor by self-awareness — they know their orientation diverges from their performance — and from the aligned double by the direction of that divergence. The behavioral signature of a hostile double under pressure is typically a pattern of escalation suggestions at moments when escalation serves the external interest rather than the ecosystem's purposes,

combined with a pattern of division-amplification that targets the fracture lines most damaging to ecosystem cohesion. Neither pattern is decisive on its own — both can be generated by aligned actors making genuine strategic errors — but both together, consistently over time, constitute a diagnostic indicator. The ecosystem that cannot tolerate any actors in double positions will destroy the functional logic of section 10.10b and the legitimacy of the translator role simultaneously. The ecosystem that cannot distinguish between aligned and hostile double positions will be progressively colonized. The management of this tension — maintaining the functional positions that require role divergence while maintaining the accountability mechanisms that prevent the exploitation of those positions — is one of the most demanding organizational challenges that any sustained movement faces. There is no procedure that resolves it definitively. There is only the ongoing

practice of maintaining the relational density, analytical redundancy, and procedural accountability that make the distinction detectable over time. The Problem If There Is No Watcher The watcher problem is symmetrical: the ecosystem that has no strategic coordinator managing relationships across its internal diversity is not safe from the watcher problem — it is simply operating with an unmanaged version of it. Without some form of coordination across the system-outside divide, the functional interdependencies that sections 10.10 and 10.10b describe cannot be realized. The institutional reformer and the movement confronting the same institution are operating in uncoordinated parallel, sometimes canceling each other's effects. The builder of parallel institutions is constructing infrastructure that the movement's other actors are not positioned to use. The sacrificial actor is absorbing costs that no one in the ecosystem is prepared to build on. The witness is documenting events that no one in the ecosystem has prepared distribution pathways for. The skeptic's challenges are treated as disruption rather than as analytical resource because no one has established the relationship within which challenge can be received as contribution. The absence of coordination does not produce a safe alternative to the watcher problem. It produces a fragmented ecosystem that is simultaneously less effective than a coordinated one and more vulnerable to the distortions that hostile external watchers introduce, because fragmented ecosystems have weaker analytical redundancy, weaker relational density, and weaker accountability mechanisms than coordinated ones. The question is not whether coordination should exist but what form it should take: sufficiently transparent to allow distributed accountability to function, sufficiently protected to avoid compromising actors whose effectiveness depends on their apparent autonomy, and sufficiently adaptive to respond to the changing conditions that the political time analysis of section 10.5b describes. This is not a theoretical problem with a theoretical solution. It is an organizational challenge whose difficulty should be acknowledged rather than dissolved by the pretense that coordination is unnecessary. Extended agent typology diagnostic: (1) Position inventory — which of the eleven durable positions (hero, vigilante, institutional reformer, builder, witness, translator, strategist, sacrificial actor, skeptic, healer, connector) are currently occupied, and which are absent? Absent positions identify functional gaps. (2) Spontaneous actor inventory — are momentum actors currently in the ecosystem? Has the durable ecosystem prepared the organizational infrastructure and concrete demands required to convert their contribution into structural outcomes before it dissipates? (3) Sympathy actor dependency — does the ecosystem's dependence on sympathy actor resources constitute a structural constraint on strategic decision-

making? At what threshold does this dependence become incompatible with the ecosystem's objectives? (4) Affected population relationship — is the affected population being treated as an object of strategy or as a distinct and partially sovereign actor? Are the ecosystem's internal models of their needs being tested against the population's own expression of those needs? (5) Orientation transparency audit — for each durable position, is the actor's actual orientation clear to themselves, to immediate allies, and to distributed accountability mechanisms? (6) Watcher identification — is there an actor shaping the orientation of multiple positions from a non-visible coordination role? Whose purposes does the coordination serve? (7) Manipulation vulnerability — which positions have the highest confidence in their analytical autonomy? Those are highest-priority targets for external steering and require structural analytical redundancy. (8) Coordination gap — if no strategic coordination exists across the system-outside divide, which functions

requiring coordination are currently incoherent or absent?

10.11 — Gradations Within the Ecosystem

The ecology of struggle does not consist of only two types of actor. Between those who operate entirely within the system and those who operate entirely outside it, there is a wide spectrum — and within each position on that spectrum, further gradations. Within the vigilante space itself there are layers: the visible and the deeper, the crude and the sophisticated, the impulsive and the calculative. The visible sometimes provides cover for the deeper in the same way that the hero provides cover for the vigilante — by becoming the easier and more accessible target for the system's attention. Beyond the hero-vigilante typology, there are actors who fit neither category: builders of parallel institutions who do not try to change or confront the existing system but construct structures that function on different principles, such that the old system becomes progressively less relevant from within. There are witnesses and chroniclers whose function is not direct change but preservation — ensuring that what happens is not lost, that each generation does not start from zero because the governing system has successfully erased collective memory. There are translators who take analysis and convert it into something actionable in the specific contexts they understand. And there is the spontaneous energy of populations who fit no typology at all — which is often the largest and most unpredictable force in the entire ecosystem. The practical consequence of this gradation is that the ecosystem cannot be analyzed as a binary. An actor who is 70 percent within the system and 30 percent operating outside its constraints is not equivalent to an actor who is entirely within it — the 30 percent matters enormously for what they can produce that fully inside actors cannot. Conversely, an actor who operates entirely outside institutional constraints but has no access to the information that only institutional positions provide is operating with a systematic blind spot that limits the quality of their analysis and the precision of their pressure. The gradations in the ecosystem are not merely descriptive. They are functionally significant: the ecosystem produces what it produces because actors at different positions on the spectrum have different capabilities and different constraints, and the combinations of these differences generate results that no single position could generate alone. The implication for strategy is that healthy ecosystems are not those in which all actors have migrated to the outside (maximally independent but institutionally blind) or to the inside (maximally informed but structurally constrained). They are those in which the full spectrum is populated — in which the positions requiring institutional access are held by actors who have not been entirely captured, and the positions requiring full independence are held by actors who have not lost contact with how institutions actually function. Gaps in the spectrum — positions that are unpopulated be

cause the incentive structure has made them too costly — are vulnerabilities. They indicate that certain combinations of capability and independence are unavailable, and that the ecosystem's capacity to generate pressure at certain levels of the institutional structure has been foreclosed. Each of these actor types carries a specific vulnerability that the ecosystem must account for. Builders of parallel institutions are vulnerable to the absorption problem: if the parallel institution succeeds, the existing system will eventually attempt to incorporate it, either by replication (creating its own version, draining the original of its distinctive function) or by direct takeover (funding, regulating, or mandating the parallel institution in ways that change what it is). The history of cooperative movements, community land trusts, and alternative media institutions is substantially a history of this absorption dynamic. Builders who do not anticipate it will find that success produces capture rather than expansion. Witnesses and chroniclers are vulnerable to irrelevance by isolation: their function requires audience, and audiences require distribution infrastructure that the existing system largely controls. A chronicle that exists but cannot be accessed by those who need it has not performed its function. The translator's vulnerability is the inverse: they are visible by definition, operating at the interface between communities, which makes them targets for both sides — suspected of bad faith by those they translate to, suspected of co-optation by those they translate for. The ecosystem protects these actors not primarily through operational cover but through legitimacy distribution: ensuring that their function is

recognized as necessary so that attacks on them are recognized as attacks on the ecosystem's capacity rather than merely on the individual. The analytical value of granularity: The hero/vigilante binary is analytically useful as an orientation but obscures variation that matters for both strategy and accountability. Within the vigilante space, the distinction between actors who operate outside institutional constraints because the constraints prevent effective challenge of structural arrangements and actors who operate outside institutional constraints because those constraints impose accountability they wish to avoid is strategically significant. The first type generates costs that are borne by the actor and justified by what the institutional constraints were protecting. The second type generates costs that are borne by others — typically those who share the actor's political objectives but lack the actor's insulation from consequences — and are not justifiable on strategic grounds. Detecting the difference requires examining not only the stated justification for constraint-violation but the distribution of its costs. The deeper operative level: The sophisticated counter-power ecosystem maintains actors who operate at a level of instit

utional distance that is invisible to the actors who face them. These actors are not deeper in the sense of more extreme; they are deeper in the sense of more structurally sophisticated — they understand the institutional architecture well enough to identify the specific nodes where leverage is concentrated, and they build the capacity to apply pressure at those nodes rather than at the visible surface of the system. Their invisibility is functional: the actors at the system's visible surface do not know they exist, which means they cannot be the target of the management strategies that the system deploys against visible counter-power. Maintaining this layer requires the ecosystem to develop trust infrastructure that is genuinely resistant to infiltration — which is among the most demanding organizational challenges counter-power faces. The ecosystem health diagnostic: a healthy counter-power ecosystem has representation at multiple levels of the spectrum simultaneously, with coordination mechanisms that allow the levels to support each other without the deeper levels being exposed by the shallower ones, and without the shallower levels being instrumentalized by the deeper ones. When this balance is absent — when the ecosystem has only surface actors or only deep actors — it loses the complementarity that makes it effective against a system that also operates at multiple levels simultaneously.

10.12 — Drift Mechanisms: How Actors Shift

No position in this ecosystem is permanent. A hero who begins as genuine can drift into a reactionary actor — not because they suddenly become bad, but because the structural incentives of their position within the institution gradually shift what they consider possible, acceptable, and threatening. The process is not dramatic. It is an accumulation of small compromises each of which appears justified by conditions, until a point is reached at which the actor is operating entirely within the logic of the system they were supposed to be opposing. A vigilante who initially operates on sharp analysis of conditions can drift toward serving ego or self-interest — not because they suddenly become opportunistic, but because operating without accountability and without feedback loops from the consequences of action produces cumulative perceptual distortion. What was initially necessary action under specific conditions becomes a mode of operating maintained even when conditions have changed, because the identity of being an actor who operates outside limits has become more important than the function that originally justified that position. Both forms of drift are mirror images of the same failure. The hero becomes reactionary because institutional capture erodes their commitment. The vigilante becomes counterproductive because the absence of constraint erodes their judiciousness. A healthy ecosystem has mechanisms for detecting and correcting both — not through centralized policing, but through distributed accountability that emerges from shared analysis of what different actions actually produce materially. Drift is structural, and the responses to it must be structural rather than individual. Relying on the virtue or vigilance of individual actors to resist institutional capture is the wrong frame: the analysis of drift is precisely that it operates through incentives that shape behavior regardless of intention, and that actors who are drifting often believe they are making necessary

accommodations rather than accumulating compromises. Individual virtue is insufficient because drift is not a moral failure. It is a structural response to a structural environment. The structural responses that actually constrain drift operate at three levels. At the level of institutional position, they include term limits and rotation that prevent the accumulation of institutional identity — the process by which an actor begins to identify their interests with the institution's interests rather than with the purposes the institution is supposed to serve. At the level of the ecosystem, they include the maintenance of independent external actors with sufficient credibility and resources to name drift when it occurs — actors whose own incentive structure does not depend on the continued goodwill of the drifting actor. At the level of organizational culture, they include explicit and recurring engagement with the question of what the position is for and what it has actually

produced, as distinct from what it has been allowed to do. None of these is a complete solution. Together, they make drift slower and more visible, which is the realistic objective. Early warning signs of drift are detectable before the drift becomes irreversible, and the ecosystem that cannot detect them will repeatedly cycle through the same loss. For heroes: increasing risk-aversion that cannot be explained by changed external conditions; a growing gap between public positioning and private analysis of what is actually possible; accumulation of institutional dependencies that make specific outcomes personally costly regardless of their political merit. For vigilantes: escalation patterns that are not responsive to changed conditions; targeting that shifts from structural positions to personal enemies; the replacement of functional criteria for action with identity criteria — acting in ways that signal membership in a type rather than responding to what conditions require. The ecosystem's detection capacity depends on maintaining relationships dense enough to observe behavior rather than merely outputs. An actor whose connections to the broader network have thinned to the point that only their outputs are visible is an actor whose drift is already advanced. This is why isolation — whether self-imposed or imposed by repression — is itself a warning sign independent of behavioral content. An actor who is genuinely functioning has reasons to maintain those connections; an actor whose primary interest has shifted to self-perpetuation has reasons to minimize the external observation that those connections would provide. The structural triggers of drift: Drift is triggered by specific structural conditions rather than by individual moral failure, and identifying the conditions is more analytically useful than cataloguing the failures. The primary trigger is resource dependence: when an actor's ability to continue their work becomes dependent on maintaining relationships with the institutions they are supposed to challenge, the challenge becomes bounded by what those relationships can sustain. The resource dependence creates cognitive as well as material pressure — actors develop genuine beliefs about why the bounded challenge is the most effective approach, which makes the drift invisible to themselves. The secondary trigger is positional success: as actors achieve recognition and institutional access, they accumulate relationships, commitments, and standing that are threatened by positions that would have been held without cost earlier in their trajectory. The third trigger is organizational survival pressure: as organizations develop, their survival becomes an organizational objective that can compete with their political objectives — drift occurs when organizational survival takes precedence over political mission in ways that are not acknowledged explicitly. The diagnostic implication: drift is not visible to the drifting actor and is often invi

sible to those closest to them, who are subject to the same structural pressures. Detection requires either actors from outside the drift trajectory or systematic application of the framework's own diagnostic tools to the actor's actual behavior rather than to their stated positions. The question is not what the actor says they are doing but whether what they are actually doing is producing the causal effects that would be expected if the stated position were genuine. When the gap between stated position and observable causal effect is large and consistent, drift is the most parsimonious explanation.

10.13 — Legitimacy as Function, Not Status

In this ecosystem, legitimacy is not something acquired once and maintained indefinitely. It is a function of what an actor produces materially for the larger purpose. When an actor stops producing that — when their function ceases or reverses — their legitimacy within the ecosystem ceases with it. The protection the ecosystem extends to actors who are still functioning no longer applies to those who are not. This applies to the hero who has become reactionary and to the vigilante who has become counterproductive. In both cases, neutralization — in the broad sense: the destruction of credibility, the severing of network access, the withdrawal of protection that allows consequences of their own actions to operate without interference — is a legitimate response from the ecosystem, not a betrayal. It is not about revenge or punishment. It is about function: ensuring that capacity that should serve the larger purpose is not consumed by an actor who has turned against it. There is one form of legitimacy exit that is rarest and most revealing of genuine understanding: self-neutralization. The actor who recognizes that they have become compromised, have drifted, are no longer serving the function that justified their position — and who withdraws before the ecosystem must do it from outside. Recognizing when to withdraw is one of the rarest forms of strength available in political action: the capacity to place the objective above identity, analysis above ego, function above position. The agent ecology is not a designed system. It is a system that emerges from adversarial conditions, from shared analysis, and from accumulated experience of what works and what does not. It has no single architect. It has conditions that support its emergence — and those conditions can be identified, analyzed, and cultivated. The question of who evaluates function — and through what mechanism — is not resolved by the framework's definition of legitimacy as functional. It is made more precise. In conventional accounts of legitimacy, evaluation is performed by some combination of formal authority (the state, the law, official institutions) and social consensus (public opinion, community norms). In the ecosystem analysis, neither is the primary evaluator: formal authority has an obvious interest in delegitimizing actors whose function is to apply pressure to formal authority, and social consensus is itself a site of contestation rather than an independent standard. The evaluation that matters is performed by the populations whose material conditions the ecosystem is supposed to improve — specifically, whether those conditions are improving in ways that can be credibly attributed to the ecosystem's operation rather than to independent factors. This creates specific problems in contested cases. An actor who claims their function is to maintain inside access in order to achieve outcomes that outside pressure alone could not achieve is making a claim that is difficult to evaluate

in real time — the counterfactual (what would have happened without the inside access) is not observable. An actor who claims their outside pressure is maintaining independence that is functionally necessary is making a similarly difficult claim. The ecosystem analysis does not resolve these claims by fiat. It provides a framework for evaluating them over time: what has the actor produced, for whom, and at what cost to the independence or access that was supposed to make the production possible? The question is empirical rather than definitional — which means it requires evidence, time, and the willingness to revise evaluations as evidence accumulates. The hardest cases are those in which an actor's function has changed without their awareness — in which what they believe they are producing no longer corresponds to what they are actually producing for the populations their work is supposed to serve. These cases are not resolved by the actor's self-assessment. They are resolved by the ecosystem's capacity to generate honest evaluation from positions that are not structurally dependent on the actor's continued operation — which is why the independence of evaluation mechanisms is itself a structural condition for ecosystem health, not merely a procedural nicety. The governance problem this creates is not trivial. If legitimacy is a function of what an actor produces for the larger purpose, then someone must evaluate that production — must determine that function has ceased or reversed, that the withdrawal of protection is warranted rather than merely expedient. This evaluation is itself a site of power, and it can be weaponized: factions within an ecosystem can use legitimacy-as-function as a tool for eliminating actors who are genuinely functional but whose methods, analysis, or constituency threatens the evaluating faction's position. The history of revolutionary and reform movements is substantially a history of exactly this dynamic. The framework does not resolve this problem by specifying a

correct evaluating authority, because no such authority exists that cannot itself be captured. What it specifies instead is a procedural standard: the evaluation of whether an actor is functioning must itself be subject to the same criteria applied to other claims — evidence-responsiveness, consistency across cases, and accountability to the purposes the evaluation is supposed to serve. An evaluation that cannot be challenged, that applies different standards to different actors without justification, or that consistently produces conclusions favorable to those doing the evaluating has failed the procedural standard regardless of whether its conclusions happen to be correct. The ecosystem that cannot maintain this standard will destroy its most effective actors while protecting those who have learned to perform function without producing it.

10.14 — Who Watches the Watchers

The most uncomfortable question in this entire analysis is about the accountability of the ecosystem itself. Who decides that a hero has become reactionary? Who decides that a vigilante has gone too far? In a centralized system, the answer is hierarchy — and the problems with that answer have been documented with great pain across the twentieth century. In a distributed ecosystem, the answer is more complex and more honest: it is distributed judgment arising from multiple genuinely independent perspectives, based on behavioral evidence rather than personal loyalty, with thresholds that move according to what is produced materially rather than according to who is most vocal in claiming to know what is right. This does not mean that all judgments are equally valid. It means that the most trustworthy judgment is the one that comes from analysis most grounded in material conditions, least dependent on the personal interests of the one judging, and most consistent with the same standard applied to all actors including the one making the judgment. That standard can be wrong. It can be gamed. But it is the most resistant to gaming among the available alternatives — and it is the most consistent with the foundational commitments of this framework as a whole. The distributed accountability of the ecosystem must operate against three principles that function as constraints on what any actor in the ecosystem can claim to justify. They are not procedural rules but criteria of evaluation — they provide the analytical basis for distinguishing methods that are genuinely necessary from methods that serve the actor's interest while claiming necessity. Proportionality. The methods employed must be proportionate to the actual threat and the actual stakes. This is not a pacifist principle — it does not prohibit hard methods under conditions that genuinely require them. It is a constraint against the escalation of method that benefits the actor's position or identity while claiming to serve the larger purpose. The question is not whether harder methods are conceivable but whether they are necessary: whether the objectives that justify the method cannot be achieved by methods that impose lower costs on the people within and around the struggle. Prospective accountability. Methods employed must be ones that can be accounted for — to the communities they claim to serve, to the historical record, to the judgment of those who come after. This is a constraint on the self-serving invocation of tactical necessity. A method that cannot be publicly defended at the time it is used because public defense would undermine its effectiveness is different from a method that can never be defended even in retrospect. The former is sometimes legitimate under genuine adversarial conditions. The latter is a signal that the method serves purposes other than the ones claimed. The test is whether, under conditions where the suppression of the account was no longer necessary, the account could

be given honestly. Non-domination. The struggle must be oriented toward the elimination of arrangements that produce structural domination — not toward their replacement with different arrangements that produce domination by different actors. This is the most important and the most frequently violated of the three. Movements that begin by opposing the domination they experience and end by reproducing it under different management have failed by the criterion regardless of what they achieved in other respects. The check against this failure is not the stated intentions of the actors involved — intentions are consistent with domination across the entire historical record of its production. The check is whether the arrangements being built distribute power in ways that are structurally accountable rather than concentrated in ways that require trust in the character of those who hold it. The structural answer the framework offers: The answer is not a

person or an institution — it is a practice embedded in the ecosystem's design. The practice is dual transparency: actors within the ecosystem who identify capture, drift, or excess in other actors must be able to communicate that identification without losing their own position in the ecosystem, and the identification must reach actors with standing to act on it. This requires two structural features that distributed ecosystems frequently fail to develop: protected internal critique — channels through which actors can raise accountability concerns without those concerns being treated as defection from the coalition — and cross-position verification, in which assessments of specific actors come from multiple ecosystem positions rather than from single sources that could themselves be compromised. The limits of distributed judgment: Distributed judgment has a specific failure mode that the framework must acknowledge: it can be captured through the manufacture of consensus. When the same narrative about a specific actor circulates simultaneously through many nodes of the ecosystem, distributed verification loses its independence — the multiple assessments are not genuinely independent data points but repetitions of a single framing whose origin may be strategic. This is the mechanism through which ecosystem accountability can be weaponized: actors with resources to shape the narrative environment can manufacture the appearance of distributed concern about an actor whose actual offense was operating effectively against the interests of those with narrative power. The framework's only structural defense against this failure mode is genuine epistemic independence among the actors who make accountability assessments — which requires the ecosystem to cultivate actual diversity of analytical position, not merely diversity of demographic identity. The practical implication: who watches the watchers is a recursive problem that cannot be finally solved — only managed through ongoing maintenance of the conditions

that make distributed judgment resistant to capture. Those conditions include the epistemic independence analyzed here, the drift detection practices described in 17.7c-ii, and the commitment to applying the framework's own diagnostic tools to the ecosystem itself rather than treating the ecosystem as exempt from the analysis it applies to other institutions.

10.15 — Micro-Politics and Structural Transformation

The framework's account of how systems change has a structural bias that must be named: it is primarily an account of macro-level dynamics. Threshold analysis (10.5), capture sequences (7.3c-i), agent typology (10.10c), and the organizational form problem (10.8) all operate at the level of institutional and systemic change. The framework addresses individual actors within these dynamics but does not adequately theorize the level at which most political and social life actually occurs — the micro-political domain of solidarities within neighborhoods, care networks within communities, mutual aid within workplaces, and the daily practices through which structural conditions are either reproduced or incrementally subverted. This section specifies what a more complete account would need to address and provides the conceptual link between micro-political practice and structural transformation that the framework's existing analysis leaves implicit. The standard objection from structural analysis is that micro-political practice risks becoming accommodation — making people comfortable enough within unjust arrangements that pressure for transformation dissipates. The accommodation-transformation tension (13.4f) is real at the micro level. But the objection rests on an assumption that is not always warranted: that micro-political practice and structural transformation constitute a zero-sum competition. The empirical record of durable transformations does not support this. The labor movements, land reform processes, and community-based institutions that have achieved genuine structural change have virtually without exception been rooted in micro-political infrastructure — dense networks of mutual obligation, local knowledge, and practiced solidarity — that preceded and sustained the macro-political mobilization producing formal institutional change. The causal relationship is not substitution but material precondition. Micro-political practice performs three structural functions in relation to transformation. The first is substrate reproduction: it maintains the relationships of trust, reciprocity, and shared interpretation without which collective action is not possible. Section 11.0's

third diagnostic applies to the material conditions for political mobilization as directly as to economic production. A community whose associational infrastructure has been destroyed by precarity, geographic displacement, and isolation cannot mobilize for structural change regardless of the level of structural suffering; the substrate for mobilization has been depleted below the threshold at which transformation is structurally available. Mutual aid, community organizing, and care networks are infrastructure maintenance — they prevent that substrate from eroding. The second function is epistemic: micro-political practice generates local knowledge that structural analysis cannot produce from a distance. Visibility is not produced by analysis; it is produced by practice. The

practitioner who works within a community over time knows what formal institutions have not acknowledged, what needs are met by informal arrangements, which structural interventions would have consequences visible from the ground but invisible from the analyst's position. The relationship between micro-political practice and structural analysis is therefore not merely causal but one of mutual epistemic dependence: each produces knowledge the other cannot generate alone. The third function is prefigurative: micro-political practice enacts, at small scale, institutional arrangements and relational forms that structural transformation aims to generalize. A mutual aid network distributing resources according to need, a community kitchen embedding care within shared governance, a solidarity fund redistributing across a precarious workforce — these are not merely approximations of future arrangements but existence proofs and laboratories. They demonstrate feasibility, identify failure modes, develop practical knowledge about governance challenges, and produce people who have practiced accountable collective decision-making. The organizational form problem (10.8) is partly addressed by this experimental function: the organizational forms that can avoid capture at scale are more likely to be found by studying forms that have managed to do so at small scale than by designing them from first principles. The accommodation risk requires active management rather than dismissal. Micro-political practice becomes a substitute for structural transformation rather than its precondition under identifiable conditions: when organized to meet immediate needs without building capacity for demands that address structural causation; when beneficiaries become dependent on arrangements requiring institutional continuity rather than change; when practitioners develop interests making them resistant to the disruption transformation requires. The diagnostic question — whether a specific practice functions as scaffolding or lock-in — cannot be answered in the abstract. It requires examining whether the practice builds counter-power capacity that can scale toward structural demands, or whether it stabilizes structural arrangements by making their costs more bearable without making them more visible or contested. Relationship to 10.10c and 17.7: the sustainer and healer roles in the extended agent typology are the micro-political positions within the agent ecosystem. This section specifies their structural function: they are not merely support roles maintaining individuals through the stresses of structural change; they are infrastructure roles maintaining the conditions under which structural change remains organizationally possible. Section 17.7d (meaning infrastructure and the sustainability of analysis) is the individual-level complement — how practitioners maintain the orientation that allows them to hold both the immediate relational demands of micro-political practice and

the structural analysis of why those practices must connect upward to demands for institutional change.

XI: The Economic Doctrine

11.0 — Two Levels of Economic Analysis

Part XI operates at two levels that must be kept explicitly distinct, because collapsing them is one of the most reliable sources of confusion in political-economic debate. The framework level asks what questions any serious economic analysis must answer, regardless of the specific arrangement under evaluation. The doctrine level applies those questions to specific arrangements under specific material conditions and selects instruments accordingly. The framework is durable; the doctrine is conditional. Critics who treat the framework's diagnostic questions as a covert commitment to a particular economic arrangement, and defenders who treat current doctrinal positions as implications of the framework itself, are both making the same error in opposite directions. This distinction also explains why what follows does not produce a blueprint. A framework that generates blueprints has confused its level of operation: it has treated the diagnostic method as if it were the conclusion the method is supposed to derive. What can be produced instead — and what this section provides — is a set of diagnostic questions rigorous enough to evaluate any economic arrangement honestly, and an illustration of what doctrine looks like when those questions are applied to a specific configuration of material conditions. The illustration is not the answer. It is a worked example of the method in operation, which is necessarily different from the method itself. Eight diagnostic questions for any economic arrangement: The result is a set of eight questions rigorous enough to evaluate any economic arrangement honestly. Each is stated as a question because the answer depends on the specific arrangement and the specific conditions under which it operates. The same question applied to different arrangements will produce different evaluations — which is the point. An economic framework that generates the same answer regardless of conditions has not been doing analysis; it has been laundering a prior commitment as analysis.

Distribution of coordination costs and benefits Every economic arrangement requires coordination, and coordination has costs: compliance costs, information costs, enforcement costs, the costs of maintaining the infrastructure within which economic activity happens. The first diagnostic asks who bears those costs and who captures the benefits. When costs are borne broadly while benefits are captured narrowly, the arrangement is producing unnecessary structural suffering through its normal functioning. The diagnostic is not whether inequality exists — some inequality in outcomes follows from any arrangement that rewards differential contribution. It is whether the distribution of costs and benefits is proportionate in a way that maintains the cooperative substrate the arrangement depends on to function at all.

Accumulation and capture dynamics The second diagnostic is the economic expression of the custodian problem. Economic power that reaches a scale sufficient

to redirect political institutions, regulatory bodies, or knowledge-production systems toward its own maintenance has escaped the conditions under which its benefits can be evaluated honestly. The diagnostic question is not whether accumulation occurs — accumulated resources are required for investment, coordination, and resilience — but whether the arrangement contains structural limits on accumulation sufficient to prevent capture. An arrangement that systematically produces the conditions for its own capture is not merely imperfect; it is generating the mechanism by which its own correction becomes structurally impossible.

Reproduction of the material substrate Every economic arrangement depends on conditions it does not itself produce: an educated and healthy workforce, infrastructure, the social institutions within which market transactions occur, the ecological systems that provide inputs and absorb outputs. The third diagnostic asks whether the arrangement's normal functioning maintains, degrades, or is indifferent to that substrate. An arrangement that systematically degrades what it depends on is consuming its own conditions of possibility — a form of structural fragility that does not register as a cost until the substrate can no longer support the arrangement that depleted it. The diagnostic includes the reproductive labor that economic arrangements typically render invisible, which is the subject of the seventh diagnostic below.

Ecological boundary conditions Ecological limits are not a value preference appended to economic analysis. They

follow directly from 3.10: systems that exceed ecological boundaries cross thresholds beyond which they reorganize rapidly and in ways that are not predictable from pre-threshold behavior. The fourth diagnostic asks whether the arrangement's normal functioning operates within those limits. An economic arrangement whose normal functioning requires exceeding ecological boundaries is not merely suboptimal — it is building systemic fragility that will eventually produce catastrophic reorganization. This is not a claim about preferences; it is a structural observation about what happens when systems exceed the material constraints within which they evolved. Shock resilience Optimization under favorable conditions at the cost of brittleness under stress is the economic expression of the adaptation problem identified in 10.4. The fifth diagnostic asks whether the arrangement maintains coordination capacity when conditions deviate from what was assumed — through climate disruption, pandemic, political instability, or resource depletion. The diagnostic is not whether efficiency and resilience are traded off, because they always are. It is whether the specific trade-off built into the arrangement is appropriate for the range of conditions the arrangement is likely to face. An arrangement that functions well under favorable conditions but fails rapidly under stress has not been evaluated under the conditions

that matter most. Counterfactual comparison The sixth diagnostic is the application to economics of the framework's epistemological commitment in Part I: evaluation must be falsifiable, which means the comparison must be specific enough to be checked against evidence. The question is not whether this arrangement is better than an ideal arrangement that has never been implemented. It is whether, evaluated against real institutional alternatives with real track records under comparable conditions, this arrangement produces less unnecessary suffering. An arrangement that can only be defended against ideal comparisons and cannot be defended against real ones has failed this diagnostic. The counterfactual must be specified honestly — not chosen to make the arrangement look better than it is. Visibility of reproductive labor The seventh diagnostic asks whether the arrangement accounts for the work that maintains the conditions under which all other economic activity is possible: the bearing and raising of children, the care of the sick and elderly, the maintenance of the social fabric that makes trust-based transactions viable. An economic analysis that treats this work as a natural background condition rather than as a foundational economic input is analyzing a fraction of the material economy and mistaking it for the whole. The distortion this produces is not incidental — in measurement of productivity, in distribution of recognition and resources, in the policy prescriptions that follow from the analysis, it systematically undervalues the work that is most essential and most consistently assigned to the populations with the least power to make their contribution visible. The Federici debt acknowledged in the Intellectual Debts section is the analytical foundation for this diagnostic. A further diagnostic question that the seven above presuppose but do not state explicitly: are the metrics used to evaluate this arrangement measuring what they claim to measure, or are they measuring a proxy that has been captured by the interests of those the arrangement serves? The measurement problem is not incidental to economic analysis — it is one of the primary mechanisms through which institutional capture operates without detection. GDP measures market activity, not welfare; it increases when pollution generates cleanup costs and decreases when parents care for children rather than paying others to do so. An arrangement that optimizes for GDP is optimizing for a metric that was designed for a specific purpose, accumulated institutional authority far beyond that purpose, and now shapes policy decisions in domains where its limitations are well documented but its authority is institutionally entrenched. The same dynamic applies across economic measurement: productivity metrics that exclude reproductive labor, poverty lines set at levels that serve administrative convenience rather than subsistence reality, growth measures that count resource depletion as income

. Each metric embeds a set of prior commitments about what counts as valuable, who counts as economically active, and what time horizon matters — commitments that are never fully explicit and rarely subjected to the scrutiny the framework's eight diagnostics apply to the arrangements the metrics are supposed to evaluate. The framework's evaluative criterion is material outcomes for the populations subject to arrangements, not metric scores. When metric scores and material outcomes diverge — when an arrangement performs well by

its own measures while producing identifiable structural suffering for the populations it governs — the framework requires prioritizing the material outcomes. The question of which metrics accurately track material outcomes for the populations being analyzed is itself a political question, not a technical one: it determines what is visible in the analysis and what disappears into the gap between the metric and the reality it claims to represent. These eight questions do not generate automatic answers. The same question applied to different arrangements in different conditions will produce different evaluations — which is the point. An economic framework that generates the same answer regardless of conditions has not been doing analysis. It has been laundering a prior commitment as analysis. These questions are the framework's contribution to economic evaluation. The doctrine's contribution is the specific answers produced by applying them to specific conditions — and the choice of instruments that follows from those answers. Eighth diagnostic: accountability and contestability. Does the arrangement have transparency and accountability mechanisms that enable affected populations to understand, evaluate, and contest the decisions that govern their conditions? The custodian problem (7.3) identifies capture as the structural tendency of any institution. Diagnostics 1 and 2 track what arrangements produce; this diagnostic tracks whether what they produce is visible and contestable by those who bear the consequences. An arrangement whose outputs cannot be evaluated by affected populations — because information is withheld, because the metrics used are inaccessible, or because formal contestation channels are nominal rather than effective — fails this diagnostic regardless of what the first seven diagnostics show. The normative assumption this diagnostic encodes follows from the minimal anchor: arrangements that produce structural suffering while insulating themselves from contestation have combined two forms of harm. The suffering is produced; and the visibility suppression (V) that prevents it from generating transformation pressure is institutionalized into the arrangement's design. Sub-questions: (1) Are the criteria used to evaluate this arrangement's performance publicly stated and operationally meaningful? (2) Do affected populations have standing to challenge decisions, and are those challenges consequential? (3) Is the en

forcement of accountability mechanisms independent of the institutions being evaluated? (4) Do transparency provisions operate at the timescale at which harms are produced, or retrospectively after harms are irreversible?

11.0b — Incentive Structures in Economic Arrangements

Section 7.3c analyzed incentive architecture as a general structural problem: the conditions under which institutional design produces agent behavior systematically divergent from the institution's stated purpose. Economic arrangements present this problem with particular acuity: the incentive structures of economic systems are simultaneously the mechanisms through which they allocate resources and the mechanisms through which they produce or fail to produce the outcomes they are designed to achieve. The eight diagnostic questions in 11.0 presuppose an analysis of the incentive structures operating within the arrangement being evaluated. This section develops that analysis as a complement to the diagnostic framework. The three economic incentive failure modes most consequential for the framework's evaluative criteria are analyzed below. Each is a structural consequence of incentive architecture, not a product of individual behavior deviating from institutional design. Regulatory Capture as Structural Outcome Regulatory capture — the process by which regulatory institutions come to serve the interests of the entities they were designed to regulate rather than the interests of the populations they were designed to protect — follows the capture sequence analyzed in 7.3b through a specific economic pathway. The actors being regulated have concentrated material interests in the institution's decisions: for a firm, a regulatory outcome can determine the viability of a business model, the cost structure of an industry, or the competitive advantage of a specific configuration of assets. The populations the regulation protects have diffuse interests: the consumer, the worker, or the community affected by regulatory decisions bears only a small fraction of the total harm from any individual decision, even when the aggregate harm is very large. This asymmetry

between concentrated regulated interests and diffuse protected interests is the structural precondition for capture: the regulated interest finds it rational to invest heavily in shaping the regulatory environment; the protected interest finds it rational to remain individually disengaged. The economic implications of regulatory capture connect directly to diagnostic question 2 in 11.0: whether the arrangement permits accumulation of economic power sufficient to produce institutional capture. Where it does, the regulatory apparatus that was supposed to constrain that accumulation becomes an instrument for maintaining and extending it. The captured regulator constructs entry barriers that protect incumbents from competition, defines safety and environmental standards in ways that favor existing production configurations over alternatives, and interprets enforcement discretion in patterns that systematically favor the regulated industry. Each of these outcomes is produced not by corruption in the individual sense but by the rational behavior of actors in positions shaped by the asymmetric incentive struc

ture that regulatory design produced. Short-Term Incentives and Long-Horizon Costs Economic arrangements routinely produce incentive structures in which the actors making decisions capture the short-term benefits of those decisions while the long-horizon costs are externalized onto populations who had no role in making them. This is not primarily a problem of moral failure. It is a structural consequence of the temporal misalignment between decision horizons and consequence horizons in complex economic systems. A corporate executive whose compensation is indexed to quarterly earnings has rational incentives to make decisions that improve quarterly earnings even when those decisions degrade the long-term productive capacity of the enterprise. A politician whose electoral horizon is four years has rational incentives to make decisions whose benefits arrive before the election and whose costs arrive after it. A financial actor whose position is marked to market daily has rational incentives to accumulate risk that will not materialize within the current assessment period. The framework's diagnostic question 4 — whether the arrangement operates within ecological boundary conditions — is the most consequential application of this analysis. Climate destabilization is, in the framework's terms, the largest-scale consequence of a temporal incentive misalignment in human history: the actors who captured the short-term benefits of fossil energy externalized the long-horizon costs onto populations who were not parties to the decisions and onto future populations who cannot be parties to any current decision. The arrangement produced this outcome not because of exceptional moral failure but because the incentive structure of fossil-fuel-dependent economic systems made this outcome the normal product of rational decision-making by the actors the system created. Addressing the ecological consequence without restructuring the incentive architecture that produced it is the application of a solution to a symptom rather than a cause. The Coordination Problem in Economic Transition Economic transitions — from one productive configuration to another, whether driven by technological change, resource depletion, or institutional reform — face a specific incentive failure that makes them systematically harder than the productivity calculations of the new configuration would suggest. The actors who benefit most from the existing configuration have both the incentive and, typically, the resources to resist transition. The actors who would benefit most from the new configuration are often not yet organized, not yet aware of their potential benefit, or not yet in the institutional positions that would allow them to act on it. This asymmetry between concentrated resistance and diffuse potential benefit is the economic expression of the counter-power accumulation problem analyzed in 10.9: those with the most to gain from change are systematically in the worst position to prod

uce it. The implication for economic doctrine is that the evaluation of any transition strategy must include an explicit analysis of the incentive structures operating on the actors who will determine whether the transition succeeds. A transition strategy that is technically optimal but that fails to account for the incentive architecture of the actors in a position to block it will fail in predictable ways: the technical optimality will be exploited to demand incremental adjustment that is adequate to the appearance of transition without the substance, while the underlying incentive structure continues to reproduce the arrangement the transition was designed to change. Effective transition strategy must change the incentive structure simultaneously with the productive configuration — which is why the most consequential levers in economic transition are typically

not the technology or the financing but the institutional arrangements that determine who bears the costs and who captures the benefits of the transition itself. The eight diagnostic questions in 11.0 identify what any economic arrangement must be evaluated on. The incentive architecture analysis in this section identifies why arrangements that fail those diagnostics persist despite their failure: the actors whose rational behavior the incentive structure produces are also, typically, the actors with the greatest capacity to resist the institutional changes that would restructure those incentives. This is not a counsel of despair — incentive structures have been restructured historically, and the conditions under which that restructuring becomes possible are analyzable through the transformation pressure framework. It is a counsel of precision: economic reform that does not attend to the incentive architecture of the arrangement being reformed will be absorbed and redirected by that architecture rather than changing it.

11.1 — Economics Without Prior Commitments

Economic systems are tools. The question is what each tool produces under what conditions, and the answer is empirical. A prior commitment to any particular economic arrangement — markets, planning, or any mixture — that is held regardless of evidence about what that arrangement actually produces is not an analytical position. It is the epistemic failure described in Part I: protecting a prior commitment from the scrutiny that would require its revision. The materialist tradition's analysis of capitalism remains one of the most powerful analyses of how economic power translates into political power, how the accumulation of capital produces structural inequality, and how the logic of production for profit generates crises that are built into the system rather than imported into it from outside. This analysis is valuable regardless of what one thinks about the political programs that were built on it — and distinguishing the analysis from the programs is one of the intellectual tasks this framework undertakes. The same discipline applies in the other direction. The failures of centrally planned economies — in resource allocation, in information processing, in the systematic underproduction of goods whose value is diffuse and whose costs are concentrated — are real failures, documented extensively, and not attributable entirely to external pressure or sabotage. A framework that absorbs the critique of markets while exempting planned alternatives from equivalent scrutiny has not applied the epistemological standard of Part I. It has simply replaced one prior commitment with another. The honest position is that both markets and planning are institutional tools whose performance depends on conditions — scale, information availability, institutional capacity, the structure of the goods being coordinated. Neither performs consistently well across all conditions. The diagnostic questions in 11.0 are designed to evaluate performance under specific conditions rather than to vindicate either arrangement in advance. The methodological claim: The commitment to empirical assessment of economic arrangements is not itself empirically derived — it is a methodological choice, and the grounds for it must be stated. The grounds are those of Part I: a prior commitment held regardless of the evidence that would require its revision is the epistemic failure that consistently produces the pattern of protecting the arrangement from the very scrutiny that would determine whether it actually achieves what it claims to achieve. Economic ideology is the domain in which this epistemic failure has been most politically consequential. The history of the twentieth century contains several multi-decade experiments in economic arrangements that were defended long past the point at which the evidence of their consequences was available and conclusive, because the arrangements had been elevated from instruments to commitments. The PMN framework extends the same analysis to all econom

ic arrangements without exception — including the one currently dominant in most of the world. What empirical assessment requires: Treating economic arrangements as empirical questions requires specifying what the evidence would need to show to support or undermine a given arrangement. The minimal anchor criterion provides the primary evaluative standard: does the arrangement reduce structural suffering, expand the conditions for human development, and do so in ways that are durable rather than dependent on conditions that cannot be sustained? These are measurable questions, not philosophical ones. The difficulty is

not conceptual but political: the actors who benefit most from particular arrangements are also those with the greatest resources to shape what counts as evidence, what alternative arrangements are treated as realistic, and what the appropriate comparison class is. This is the knowledge production problem applied to economics — the arrangement shapes the epistemic infrastructure through which it is evaluated, which means the evaluation cannot be conducted without attending to who controls that infrastructure. Markets as a specific case: Market arrangements are the canonical example of what the framework means by treating economic systems as tools. Markets solve specific coordination problems efficiently under specific conditions: when preferences can be expressed through price signals, when production costs are internalized rather than externalized, when information is distributed rather than concentrated, when competitive pressure prevents systematic extraction, and when the goods being coordinated are ones whose distribution by purchasing power does not systematically violate the minimal anchor criterion. Where these conditions obtain, market coordination produces outcomes that are difficult to replicate through other arrangements at equivalent cost. Where these conditions do not obtain — and the conditions for genuine market function are demanding enough that they are partially absent in most actual markets — market coordination produces systematic failures that have specific structural signatures: externalization of costs onto those without market power to resist it, concentration of informational advantage into structural leverage, and the conversion of coordination efficiency into extraction efficiency through capture of regulatory infrastructure. The PMN analytical position is not that markets are good or bad but that these specific conditions determine what specific markets produce in specific domains — and that the political economy of pro-market ideology consists substantially in obscuring the gap between the conditions under which market function is claimed and the conditions that actually obtain. The framework's evaluative position on economics is therefore not silence or neutrality. It is the commitment to making the empirical question explicit, stating the evidence criteria, and applying those criteria consistently across

all arrangements. An arrangement that reduces structural suffering and expands human development capacity under current conditions is the right instrument under current conditions. Whether it continues to be the right instrument is an ongoing empirical question, not a matter settled in advance by commitment to a school of economic thought. Self-application: the PMN framework itself has a material position in the political economy of ideas. Its production conditions — who writes it, who distributes it, who engages with it — shape its blind spots in ways that 12.7 identifies. The economics analysis is not exempt from this. The framework should be read as offering a methodology for economic analysis rather than a conclusion about which arrangements are correct, on the grounds that the conditions under which the methodology is applied will be different in different historical contexts and the conclusions must track those conditions rather than preceding them.

11.2 — The Permanent Conditions and Economic Doctrine

The economic doctrine begins from the permanent conditions identified in Part III. Any serious economic arrangement must manage resource allocation under genuine scarcity without producing distributional outcomes grossly misaligned with the minimal anchor. It must sustain collective investment in goods that individual actors will not produce in adequate quantity. It must prevent the accumulation of economic power that enables institutional capture. And it must maintain legitimacy in the face of the inevitable failures that any real economic system will produce. Scarcity in this context is not an abstract economic concept. It is a material reality with specific geographical and ecological dimensions. The resources that human populations require — freshwater, arable land, energy, minerals, biodiversity — are distributed unevenly across the planet's surface in patterns that were determined by geological and climatological processes entirely independent of human need or political arrangement. The economic systems that have developed in any given territory are downstream of this distribution: they are solutions to the specific resource constraints that geography imposed, not expressions of cultural preference or ideological choice. This geographical

dimension of scarcity has several consequences that standard economic analysis underweights. First, it means that no single economic model is appropriate for all territorial contexts, because the resource constraints that economic arrangements must manage differ fundamentally across territories. The development economics literature has documented this pattern extensively: policies that work in resource-abundant contexts fail in resource-constrained ones, and vice versa. The appropriate economic doctrine for any territory begins with the material geography of that territory, not with universal principles derived from other contexts. Second, the geographical distribution of resources is itself a source of power asymmetry that precedes and shapes market relationships. Territories with resources that other territories require have structural leverage that is not produced by any economic arrangement but that all economic arrangements must negotiate. The economic history of the modern era is substantially a history of how this structural leverage has been converted into political and military power, and how that power has shaped the terms of economic exchange between territories with different resource endowments. Third, the permanence of scarcity as a condition does not mean the permanence of any particular form of scarcity. Technological change has repeatedly transformed what counts as a scarce resource — making previously worthless materials strategically critical and previously essential resources substitutable. The economic doctrine must therefore evaluate not only current resource distributions but the trajectory of technological change and its implications for which resource constraints are likely to beco

me more or less binding over the relevant time horizon. No existing economic arrangement addresses all of these adequately. The comparison that matters is between available alternatives evaluated by their actual consequences under actual conditions — not between any actual system and an ideal that has never been implemented. This is a more demanding standard for evaluation than either defenders or critics of any economic arrangement usually apply, because it requires specifying the alternatives honestly and acknowledging the failures of preferred arrangements with the same rigor applied to opposed ones. Private Ownership of Productive Assets: A Doctrinal Position and Its Derivation The classical Marxist position — that private ownership of the means of production is the structural root of exploitation and must be abolished for genuine human emancipation — is held as a conclusion of historical necessity: because capitalism contains contradictions that will produce its own negation, and because socialist organization of production is the form that succeeds it, the abolition of private ownership is where history is going. This is both the position's greatest rhetorical strength and its fundamental analytical vulnerability. When the historical necessity fails to materialize — when capitalism proves more adaptive than the contradictions account predicts, when revolutionary movements produce new forms of domination rather than emancipation — the position has no resources for revision. It can only wait for the eventual completion of the historical process or insulate itself from the evidence through increasingly elaborate theoretical rescue. PMN arrives at an abolition position through a different route — one that is explicitly doctrinal rather than structural, provisional rather than necessary, and therefore more defensible and more honest. The route runs through the eight diagnostic questions applied to actually existing capitalism under current material conditions. The first diagnostic asks whether the arrangement permits accumulation of economic power sufficient to produce institutional capture. The answer under current conditions is not ambiguous: private ownership of productive assets at the scale characteristic of contemporary capitalism generates concentrations that have systematically captured the regulatory, legislative, judicial, and epistemic institutions whose function is to constrain that accumulation. This is not a prediction derived from theory — it is an observation of what the arrangement has produced wherever it has operated at sufficient scale and duration. The second diagnostic asks whether the arrangement generates distribution of coordination costs and benefits that is compatible with the evaluative criterion. The answer is that it generates distributions of developmental capacity so concentrated — in the access to education, health, nutrition, security, and meaning infrastructure that adaptive response capacity requires — that th

e structural condition of the majority is one of foreclosed becoming rather than constrained becoming. Not poverty in the absolute sense but systematic exclusion from the conditions under which development is

actually possible. The remaining diagnostics — reproduction of the material substrate, ecological boundary conditions, shock resilience, counterfactual comparison, visibility of reproductive labor — compound the finding. Contemporary capitalism's treatment of the ecological boundary as an externality to be discounted rather than a constraint to be respected is not an anomaly correctable through policy adjustment: it is a structural consequence of the arrangement's incentive architecture, which systematically discounts future costs against present returns in ways that no institutional fix within the arrangement has succeeded in reversing at the required scale. The invisibility of reproductive labor — the work of biological and social reproduction that the arrangement depends on but does not account for — is not a measurement error but a structural feature: making it visible would require compensating it at a level incompatible with current return rates on capital. The doctrinal position that follows from this analysis is: under current material conditions, private ownership of productive assets at the scale and form that contemporary capitalism has produced is incompatible with the evaluative criterion. The arrangement systematically forecloses adaptive response capacity for the majority, captures the institutional mechanisms that should constrain it, degrades the ecological and epistemic substrates it depends on, and reproduces itself by making the costs of its operation invisible to those bearing them. This is the grounds for an abolition position — not historical necessity, but the application of eight diagnostic questions to current conditions, each of which produces the same finding independently. The difference from the classical position is not rhetorical. It is structural and has consequences for how the position is held and how it should be revised. Because the PMN position is derived from diagnostic application rather than historical necessity, it is falsifiable: an arrangement of private ownership of productive assets that passes all eight diagnostics — that demonstrably does not produce capture, that distributes developmental capacity broadly enough to make adaptive response capacity structurally available to the majority, that operates within ecological constraints, that maintains the material substrate and epistemic infrastructure it depends on — would require revision of the doctrinal position. This is not a concession to capitalism. It is the epistemological standard the framework applies to all its claims. In practice, no such arrangement has been demonstrated to exist at scale. The doctrinal position is not weakened by its provisionality — it is strengthened by it, because provisionality means it is tracking reality rather t

han defending a conclusion fixed in advance of the evidence. What the PMN derivation also identifies, which the classical derivation does not, is that abolition of private ownership of productive assets is necessary but not sufficient. The layered constraint model established in 2.2 requires that transformation work simultaneously at multiple levels — biological floor, economic layer, institutional layer, epistemic layer, meaning layer. An economic transformation that transfers ownership of productive assets to state or collective control without addressing institutional capture produces what the 20th century's socialist experiments consistently produced: new capture by the party-state apparatus that controls the collectively owned assets, operating with the same logic of concentration and foreclosure that characterized private ownership, under a different name and with less developed accountability mechanisms. The goal of the doctrinal position is not the formal abolition of private title to productive assets. It is the structural condition in which the ownership and governance of productive assets cannot concentrate in ways that produce institutional capture, foreclose becoming for the majority, or degrade the substrate that all development depends on. The formal question of ownership structure is one variable in that analysis, not its resolution. Doctrinal summary on productive ownership: PMN holds that under current material conditions, private ownership of productive assets at existing scales is incompatible with the evaluative criterion — not because history requires its abolition, but because diagnostic application consistently identifies it as the structural mechanism through which adaptive response capacity is foreclosed for the majority and institutional accountability is systematically undermined. This position is provisional: it tracks the diagnostic findings, and if the findings change, the position must change. It is not provisional in the sense of weak or tentative — the diagnostic findings under current conditions are strong. It is provisional in the sense that epistemological honesty requires: the conclusion follows from the evidence, and the evidence is what holds it. Criteria for Revision: When the Doctrinal Position Would Change Provisionality is not a rhetorical gesture. For the

doctrinal position on productive ownership to function as a genuine empirical commitment rather than a conclusion fixed in advance, the conditions under which it would be revised must be stated explicitly. The position is derived from the consistent failure of private ownership of productive assets — at the scale and form of contemporary capitalism — across all eight diagnostic dimensions. The criteria for revision follow directly: the position would require reconsideration if a configuration of private ownership of productive assets demonstrated, at significant scale and over sufficient duration to exclude transience, that it passes the eight diagnos

tics without requiring the structural conditions that have produced failure elsewhere. The key terms require precision. Significant scale means operation across a substantial population and economic domain — not isolated cooperative experiments or single-sector arrangements, but configurations that govern the production and distribution of conditions material to becoming for a broad population. Sufficient duration means long enough to assess capture dynamics, ecological effects, and intergenerational reproduction — at minimum several decades, since institutional capture and ecological degradation typically operate on timescales longer than a single political cycle. Without requiring structural conditions means without depending on features that are themselves products of arrangements the framework identifies as problematic — a private ownership configuration that only functions without capture because it operates within a heavily regulated environment produced by sustained political counter-power is not evidence that private ownership passes the diagnostic; it is evidence that the counter-power passes it. Two partial findings that do not meet the revision threshold are worth naming explicitly to prevent their misuse as apparent counterevidence. First: that specific sectors or scales of private ownership produce outcomes consistent with the evaluative criterion does not establish that private ownership of productive assets at the scale that generates the capture dynamics identified across the eight diagnostics is compatible with it. A small family business that passes all eight diagnostics is not evidence against a position about concentrated ownership of productive assets at the scale that produces institutional capture. Second: that socialist or collectively owned arrangements have historically produced capture, foreclosed becoming, and failed the diagnostics does not establish that private ownership passes them. The failure of the historical alternatives is evidence about the specific institutional forms those alternatives took — particularly the absence of accountability mechanisms that could prevent party-state capture — not evidence that private ownership is the only viable arrangement. The revision threshold requires positive demonstration of a viable alternative, not merely negative demonstration of historical failures. Automated Production and the Ownership Question at Scale The analysis of private ownership of productive assets in this section was developed primarily in relation to the industrial and financial capitalism of the twentieth and early twenty-first centuries. The rapid development of artificial intelligence, robotics, and automated production systems introduces a structural change to that analysis that requires explicit treatment. The core diagnostic of the eight-point framework — that private ownership of productive assets at the scale of contemporary capitalism consistently produces capture, forecloses becoming for the majo

urity, discounts ecological limits, and makes the costs of reproduction invisible — applies to automated production with amplified force. The reason is specific to the economics of automation: the fixed costs of building automated production systems are very high and the marginal costs of their operation are very low, which means that the economic returns to ownership of automated production apparatus are concentrated in the ownership position in a way that industrial production, which requires ongoing labor input to function, is not. When a factory requires ten thousand workers to operate, ownership of that factory captures a portion of the value that workers produce; the workers retain bargaining power because the factory does not function without them. When the same factory operates with no workers, ownership of the factory captures all of the value it produces; there is no labor input from which workers could bargain for a share. This structural shift in the economics of production has a specific consequence for the eight diagnostics that ground the doctrinal position on private ownership. The elite capture rate K , which represents the proportion of structural surplus appropriated by a governing stratum rather than distributed toward floor provision and becoming expansion,

approaches its theoretical maximum under fully automated private ownership because the leverage of non-ownership — the threat to withhold labor — approaches zero. The visibility suppression of costs becomes more complete because the production of goods and services that were previously produced by labor that had interests and relationships and political voice is now produced by systems that have none of those things; the human costs of production, which labor politics made visible through strikes, organizing, and the accumulated political pressure of the working class, are replaced by maintenance costs of machinery that have no political agency. The ecological externalization intensifies because automated production optimized for cost minimization has no internal constraint on ecological input extraction beyond what regulation imposes, and the regulatory capture dynamic that the framework identifies as a systematic failure of private ownership under current conditions is even more available to entities whose productive capacity does not depend on maintaining the consent of a labor force. The doctrinal implication is that the case for transformation of ownership arrangements is stronger under conditions of automated production than under industrial capitalism, not weaker. The often-heard argument that automation makes socialism unnecessary — that abundance produced by machines makes ownership questions irrelevant — gets the analysis backward. Automation makes the ownership question more urgent because it removes the distributional mechanism — labor's bargaining power — that industrial capitalism contained within its own structure, however inadequately. Under fully automated p

ivate ownership, there is no internal mechanism through which the value produced by automated systems flows toward the majority of the population. There is only the external mechanism of political redistribution — taxation, regulation, public ownership — which is subject to the capture dynamic that the framework has identified throughout. The question is not whether automation makes the ownership question go away. It is whether the institutional designs for collective ownership of automated production apparatus can be developed that prevent the party-state capture that collective ownership has historically produced, while also preventing the private capture that automation makes structurally total. The specific institutional design challenge for automated production is the combination of technical opacity and scale. Industrial production was opaque in many ways, but workers who operated machinery understood its operation and could contest management decisions about its use from a position of technical knowledge. Automated production systems — particularly those involving machine learning — are opaque in ways that resist understanding even by the engineers who build them; the system's behavior emerges from training on data in ways that are not fully interpretable by any individual. This opacity creates accountability gaps that institutional designs for collective ownership must address explicitly: how are the decisions embedded in automated systems — what to produce, for whom, at what ecological cost — made subject to democratic accountability when the decision-making is distributed across parameters in a system that resists human-readable interpretation? This is not a reason to reject collective ownership of automated production. It is a specification of what collective ownership of automated production requires that collective ownership of industrial production did not.

11.3 — Horizon and Instrument

Economic doctrine requires a distinction between the horizon — the direction the minimal anchor points toward — and the instrument — the arrangements most materially effective at moving toward that horizon under current conditions. Several traditions in the history of economic thought have functioned as horizons in the relevant sense: providing a reference point that keeps development from becoming circular, without requiring that the reference point be reached in any specific form or timeline. The permanent conditions analysis has a direct implication for the justification of taxation that the framework's economic doctrine requires making explicit. Taxation is not primarily a redistribution mechanism — this framing, while not wrong, understates what taxation is doing structurally. Taxation is the primary mechanism through which collectively maintained conditions — the physical and social infrastructure within which all economic activity occurs — are funded by those who benefit from those conditions. This provides a structural

justification that is more defensible than redistribution arguments alone: the question is not whether accumulated wealth should be shared with others who have less, but whether the conditions that made accumulation possible were maintained by collective investment, and whether those who accumulated most from those conditions have contributed proportionately to maintaining them. The framework's criterion applies differentially to different forms of taxation. Taxation of labor income faces the weakest structural justification and the strongest claim to minimum threshold protection: labor income is the primary mechanism through which populations below the minimal anchor maintain their position above it, and taxation that reduces that income below the threshold is taxation that violates the criterion it is supposed to serve. Taxation of accumulated wealth, capital income, and inheritance faces the strongest structural justification: these represent returns on conditions — legal infrastructure, physical infrastructure, educated labor pools, stable institutions — that collective investment produced and continues to maintain. The inheritance case is structurally distinct from the accumulation case: it transfers the conditions of becoming to specific individuals on the basis of family membership rather than contribution, which is precisely the arrangement that the framework identifies as converting structural advantage into predetermined life trajectories. The framework does not prescribe specific tax rates — those are doctrinal questions that depend on specific material conditions. It does establish the structural principle: the burden of justification falls on arrangements that allow the indefinite accumulation and transmission of returns on collectively maintained conditions without proportionate contribution to maintaining those conditions, not on arrangements that require such contribution. The most developed of these traditi

ons — the one with the most sophisticated analysis of the relationship between economic organization and political power, between material conditions and ideological commitments, between the current distribution of suffering and the arrangements that produce it — is the Marxist tradition. Taking it seriously as a horizon means taking its analysis of what is wrong with existing arrangements seriously, while refusing the teleological structure that made it self-insulating. The analysis can be extracted from the teleology. The resulting framework is more honest about what it does not know, more responsive to evidence that contradicts its predictions, and more capable of genuine revision. That combination makes it more likely to be useful. The instrument — what is actually implemented at any given moment — is selected based on material conditions. Arrangements appropriate at one level of institutional development may be inappropriate at another. The same policy that works in one context may fail in another not because the policy is wrong in principle but because the conditions required for it to work are absent. Sequencing matters more than categorical commitment to any particular arrangement (Chang 2002; Amsden 1989; Rodrik 2007). This is a lesson that the materialist tradition learned painfully and repeatedly, and that this framework takes as one of the most important empirical findings available about how economic change actually works. Why the distinction is necessary: Without a distinction between horizon and instrument, economic doctrine collapses into one of two errors. The first error is instrumentalism without direction: assessing what works under current conditions without a reference point that prevents the assessment from becoming circular — treating what produces good outcomes for those currently advantaged as the general criterion of good outcomes. The second error is utopianism without instrumentalism: holding a horizon as a fixed destination that prescribes specific institutional arrangements regardless of what those arrangements produce under current conditions. The Marxist tradition's failure to maintain this distinction in practice is illustrative: the horizon of a classless society operating according to need rather than profit was legitimate as a reference point for the direction of change; the conversion of that horizon into a prescription for specific institutional arrangements — state ownership, central planning, vanguard organization — that were to be implemented regardless of their actual consequences was the error. The institutional prescriptions were treated as the horizon rather than as instruments for approximating it under specific historical conditions. What serves as horizon under PMN: The horizon for PMN's economic analysis is specified by the minimal anchor combined with the becoming framework of Part IV: the economic conditions under which human beings can meet their biological requirements without unnecessary suff

ering, and under which the material preconditions for genuine development of capacities — not merely for survival — are broadly distributed. This is not a specific institutional arrangement. It is a direction that can be used to evaluate how far any current arrangement is from it and whether the trajectory of change is moving toward or away from it. The horizon is not the end of scarcity — that is a specific empirical prediction about what technology will eventually produce, not a standard that can be used to evaluate current arrangements. The horizon is the ongoing approximation of the conditions under which the constraints that produce unnecessary suffering are progressively reduced and the conditions for development are progressively expanded. What qualifies as instrument: An instrument is an economic arrangement that is adopted because the evidence indicates it is currently the most effective way of moving toward the horizon under the material conditions that currently obtain. Instruments are assessed by their consequences, not their theoretical elegance. An instrument that produces better consequences for the minimal anchor criterion under current conditions is a better instrument than one whose theoretical logic is more coherent but whose actual consequences are inferior. Instruments are provisional: the conditions under which a given instrument is optimal change over time, and what is the best available instrument in one historical configuration may be inferior in another. This is why the framework insists on separating horizon from instrument: the horizon is held constant across changing conditions; the instrument is revised in response to those conditions. The practical consequence: PMN does not prescribe a specific economic arrangement as the correct one. It prescribes a method — hold the horizon constant, assess instruments by their consequences, revise instruments as conditions change, and resist the conversion of instruments into commitments that are defended regardless of consequences. Applied to current conditions, this method produces specific evaluative conclusions that are argued in the doctrine sections. But those conclusions are conclusions of the method applied to current evidence, not commitments that precede the application of the method. The political vulnerability of the distinction: The horizon-instrument distinction is politically vulnerable in a specific way: actors who benefit from current instruments have interests in collapsing the distinction in both directions. They argue that the current instrument is the horizon — that market capitalism is not merely the best currently available approximation to good economic outcomes but the endpoint of economic development — or they argue that any proposed alternative instrument is in fact a disguised horizon that would produce totalizing reorganization rather than instrumental adjustment. Both moves serve to make the current arrangement resistant to empirical assessment. The

framework's response is to make the distinction explicit, state the evidence criteria for instrument assessment, and apply them consistently. Diagnostic: Is a given economic argument treating an instrument as a horizon (defending an arrangement on grounds of theoretical commitment rather than empirical consequences)? Is it treating a proposed instrument as a horizon (attributing to an alternative arrangement consequences that go beyond what the arrangement would actually produce)? Answering these questions requires identifying what the speaker is treating as fixed versus what they are treating as open to empirical assessment.

11.3b — Contemporary Capitalism as Terrain for the Eight Diagnostics

The eight diagnostic questions in 11.0 apply across economic arrangements and historical periods. The specific terrain of early twenty-first century capitalism has features not present in earlier capitalisms — features that change how each question operates in practice. Taking contemporary capitalism as the terrain for analysis is not an endorsement of it. It is an identification of the structural features of contemporary capitalism that are most relevant to each of the eight diagnostics — the parameters within which doctrine must operate. Financialization and diagnostic 1 (distribution of coordination costs and benefits). The share of economic activity organized around the creation, trading, and management of financial instruments rather than the production of goods and services has grown substantially across the period since the 1970s. The coordination-cost-benefit diagnostic applies differently in financialized economies because the relationship between value creation and value capture has changed: actors who add little to productive coordination can

capture large portions of its benefits through financial positions, while the costs of economic volatility — amplified by financial interconnection — are borne primarily by workers and states. The diagnostic must ask, in this terrain, not only how coordination costs and benefits are distributed within a specific production arrangement but how financialization redistributes those costs and benefits across the economy as a whole. Platform monopoly and diagnostic 2 (accumulation and capture). Platform economies generate monopoly dynamics through network effects: the value of a platform increases with its user base, which means early leaders accumulate competitive advantages that become increasingly difficult to challenge over time. The accumulation-and-capture diagnostic is particularly acute in this terrain because platform monopolies operate across jurisdictional boundaries that were designed for geographically bounded firms, accumulate data as a form of capital that compounds without physical limits, and have developed regulatory relationships in major economies that show the patterns the diagnostic is designed to detect. The framework does not conclude from this that platform monopolies are necessarily harmful — the counterfactual comparison diagnostic requires evaluating what alternatives actually exist and what they would produce. But the structural features of platform monopoly place contemporary capitalism at the high-risk end of the accumulation-and-capture diagnostic in ways that earlier capitalisms did not. Reproductive labor visibility and diagnostics 3 and 7. The systematic invisibility of reproductive labor — identified in diagnostic 7 as a structural feature of economic analysis — is not a new problem. But the specific form it takes in contemporary capitalism has distinctive features: the partial commodification of care work (paid domestic labor, care industries) without corresponding

changes in the distribution of unpaid reproductive work; the transfer of some forms of reproductive labor to platforms and gig arrangements that reproduce the conditions of invisibility in new forms; and the globalization of care chains that move reproductive labor from lower-income to higher-income economies in ways that reproduce colonial patterns of extraction. Diagnostics 3 and 7 both activate in this terrain. Ecological limits and diagnostic 4. The ecological boundary condition diagnostic has a specific urgency in contemporary conditions that distinguishes it from earlier capitalisms: several ecological boundaries — atmospheric carbon concentration, biodiversity loss rate, nitrogen cycle disruption — have crossed thresholds that the framework's analysis of non-linear system dynamics (10.5) identifies as qualitatively significant. An economic analysis that treats these as external costs to be managed through pricing signals is not wrong in principle — the diagnostic framework is agnostic about instruments — but it must account for the fact that the boundaries have already been crossed in ways that produce path-dependent consequences that pricing signals alone cannot reverse. Technology and the permanent conditions: a clarification. Diagnostic frameworks built on permanent material conditions — scarcity, coordination requirements, betrayal possibility — face a recurring challenge from technological change: does a new technology change the conditions themselves, or only the specific form they take? The framework's position is precise on this. Technological change can and does shift the boundary between what is scarce and what is abundant, what is preventable and what is irreducible, what coordination arrangements are possible and what are not. What it has not done — and what there is no current evidence it is doing — is eliminate the structural features that generate the conditions: that organisms require resources that exist in finite quantities; that collective action generates coordination problems regardless of the technology through which it operates; that actors in positions of institutional authority face incentives to use that authority in ways that serve their own interests. The appearance of technologies that might eventually transform these conditions — artificial general intelligence, molecular manufacturing, radical life extension — does not justify treating the conditions as already transformed. The framework updates when the evidence requires it, not when the possibility exists. This section describes terrain, not conclusions. The structural features of contemporary capitalism identified here — financialization, platform monopoly, care chain globalization, ecological threshold crossing — change the parameters within which the eight diagnostics operate, and therefore change what specific institutional arrangements will score well or badly on each diagnostic. They do not change the diagnostics themselves. A doctrine derived from th

is framework, applied to contemporary conditions, will produce specific positions on each of these features. Those positions are doctrine, not framework — they will require revision as conditions change, as evidence accumulates, and as alternative analyses demonstrate their explanatory superiority. Two features of early twenty-first century capitalism require specific attention because they substantially modify how several of the eight diagnostic questions operate in practice. The first is the financialization of the production relationship: the increasing proportion of surplus generated through financial intermediation rather than through the direct production of goods and services creates a configuration in which the actors who capture the greatest proportion of surplus are increasingly insulated from the productive consequences of their allocation decisions. This insulation degrades the feedback mechanism through which market arrangements are supposed to generate efficiency: financial actors whose returns are generated by fee extraction and asset price appreciation rather than by productive performance have reduced incentive to allocate capital toward activities that improve productive capacity. The concentration of political influence in financial sectors that this pattern produces translates into regulatory capture that protects the financialization dynamic from correction through the mechanisms that would theoretically discipline it. The second feature is the platform economy's reorganization of the employer-employee relationship. Platform arrangements that classify workers as independent contractors while maintaining functional control over their working conditions produce the extraction pattern of employment without the legal obligations of employment — concentrating productivity gains in the platform infrastructure owner while externalizing insurance, equipment, and income stability costs onto the workers. This is a structural reorganization of the employment relationship that has legal and organizational forms calibrated to avoid the regulatory infrastructure that was built to constrain extraction from the prior employment relationship. The eight diagnostic questions, applied to platform arrangements, consistently reveal the gap between the formal legal classification (independent contractor) and the structural economic relationship (dependent labor with concentrated counterparty power).

11.4 — What Doctrine Looks Like: A Contrastive Illustration

The eight diagnostic questions in 11.0 are abstract. The following illustration shows what it means to apply them — not to generate a definitive position, but to demonstrate the method. The illustration uses a single problem — providing a population with reliable access to safe water — evaluated under two genuinely different institutional arrangements. Water access is chosen because it is concrete, universal, and politically charged in ways that make the contrast between analytical and ideological responses visible. The illustration is not an argument for one arrangement over the other. It is a demonstration of what the framework's questions produce when applied honestly to both. Arrangement A: Market provision. A private water utility, operating under a concession from the state, supplies water to a metropolitan area. Prices are set to recover costs and generate return on capital. Connection to the network requires payment of connection fees. Usage is metered and billed. Arrangement B: Commons-based public provision. A municipally operated system supplies water, funded through general taxation and flat-rate connection fees. Usage is not individually metered above a minimum threshold. Governance includes community representation in rate-setting decisions. Applying the eight diagnostics: 1. Distribution. Under A, coordination costs — infrastructure investment, maintenance, operational costs — are recovered through prices. Households below the price-clearing income threshold are excluded from connection or consume less than biological minimum. Benefits (profits, returns on capital) flow to investors. Under B, costs are distributed across the tax base including non-users; benefits (access) are available regardless of income. The diagnostic question is whether the cost-benefit distribution under A produces unnecessary structural suffering among the excluded population — which is a function of the proportion excluded, the alternatives available to them, and whether exclusion is produced by the arrangement's normal functioning rather than by incidental individual circumstance. In most documented implementations, exclusion is structural, not incidental. 2. Accumulation and capture. Under A, the utility operates as a natural monopoly — infrastructure duplication is prohibitively costly, so competition is structurally unavailable. Natural monopoly plus essential service plus profit motive

creates strong conditions for regulatory capture: the utility has concentrated interest and resources to direct toward the regulatory environment, while consumers have diffuse interest and limited capacity for sustained political organization. Under B, capture dynamics are different: the risk is capture by municipal political networks rather than by capital, which produces different failure modes (patronage, inefficiency) rather than the same ones (extraction, exclusion). 3. Substrate reproduction. Both arrangements reproduce the physical infrastructure if adequately fun

ded. Under A, infrastructure investment in low-revenue areas — low-density, low-income, geographically peripheral — is systematically underprovided because return on investment is lower. Under B, cross-subsidy from higher-density areas to lower-density ones is structurally possible and often designed in. The diagnostic asks not whether both systems can in principle reproduce infrastructure, but which one does so across the full range of conditions the population presents. 4-7, briefly. Ecological boundary: both arrangements face the same boundary conditions; the diagnostic distinguishes whether pricing signals that suppress overconsumption (A) or governance constraints (B) are more effective under the specific hydrological conditions of the territory. Shock resilience: A is vulnerable to financial distress of the concessionaire; B is vulnerable to fiscal distress of the municipality — different fragilities requiring different analysis. Counterfactual: neither arrangement is evaluated against an ideal; each is evaluated against the other as actually implemented in comparable conditions, and against the historical record of transitions between them. Reproductive labor: both arrangements are typically silent on the gendered distribution of water-carrying labor in households without connection — the diagnostic forces this into the analysis. 8. Accountability and contestability. Under A, the criteria used to evaluate performance are commercially defined — service quality, uptime, incident rates — and are available to regulators but not necessarily to the populations bearing exclusion. The decisive sub-question is standing: under A, a household excluded by price has no consequential channel to contest the pricing structure, because the criterion (ability to pay) is the arrangement's own operating principle, not an implementation defect. Transparency provisions are typically retrospective — regulatory reporting after harms are realized — and enforcement is administered by the same regulatory body whose capture the second diagnostic identified as the arrangement's structural tendency. Under B, the criteria are formally public and standing is formally available through municipal political channels. But formal availability does not satisfy the diagnostic: sub-question (3) asks whether enforcement is independent of the institutions being evaluated, and under B the evaluating institution and the operating institution are the same municipal authority — a failure mode the capture analysis in (2) already identified as patronage rather than extraction. Sub-question (4) is where the two arrangements most clearly diverge: B's political channels operate on an electoral timescale, which is slower than the timescale at which service degradation produces irreversible health and livelihood harms, but faster than A's, where no channel exists for non-consumers at all. The diagnostic does not rank B above A on this basis alone. It records that A fails the standing criteri

on structurally — exclusion is the mechanism, not a remediable defect — while B fails the independence and timescale criteria, which are remediable through institutional design (independent utility oversight, enforceable service standards with real-time reporting). What the diagnostic adds that the first seven do not show: the distributional and reproductive analyses above could be conducted entirely from performance data, and both could be published in full, without either arrangement becoming contestable by the populations they exclude. Visibility of outcome is not the same as contestability of decision, and an arrangement can satisfy the first seven diagnostics in a given case while failing the eighth. The point of the illustration is not that B is superior to A in all conditions. There are documented cases where municipal provision has failed catastrophically — through corruption, underinvestment, and political manipulation — while private provision under effective regulation has performed well. There are documented cases where the opposite is true. The point is that the framework's eight questions produce a specific, falsifiable analysis of each arrangement in specific conditions, rather than a predetermined conclusion. A doctrine derived from this framework will recommend different instruments in different conditions — and it will be able to say,

specifically, what evidence would lead it to revise the recommendation. That is what distinguishes doctrine from ideology. The choice of water access as illustration is deliberate. Water is a domain where the ideological stakes are high enough that the contrast between framework-level analysis and prior-commitment-based advocacy is clearly visible — where defenders of market provision and defenders of public provision both tend to reach their conclusions before applying analysis rather than from it. If the framework's eight questions can produce a genuinely open evaluation in this domain, they can produce one anywhere. If a reader finds the analysis of Arrangement A or B here reads as a conclusion rather than as a demonstration of method, that is evidence that the illustration has failed — and that the specific adjustment required is more explicit identification of what evidence would reverse the evaluation. What the water case reveals about doctrine: The water illustration exposes a specific structural feature of doctrinal commitment that the framework's analysis of economic ideology identifies: the way in which prior commitments to a particular institutional arrangement shape what counts as a failure, what alternatives are considered, and what evidence is treated as relevant. An actor committed to market provision of water access will, when market provision fails to meet the minimal anchor criterion, characteristically attribute the failure to market implementation deficiencies — insufficient price signaling, inadequate property rights specification, regulatory interference — rather than to featu

res of the market arrangement itself. An actor committed to public provision will characteristically attribute failures to administrative inefficiency, political interference with operational decisions, or underfunding rather than to features of public provision itself. Neither actor is conducting the empirical assessment that the framework requires; both are protecting a prior commitment from the scrutiny that would require its revision. The methodology in application: The framework's approach to the water question is to ask: under what institutional conditions does this population receive reliable access to safe water at what cost, and what does the evidence say about which conditions produce better outcomes on those criteria across comparable populations? This is an empirical question, and it has empirical answers that vary by context — answers that consistently show that neither market nor public provision dominates across all contexts, and that the variables determining which performs better in specific contexts include regulatory capacity, population density, investment infrastructure, and the degree of capture of whichever institutional form is operating. The doctrine question resolves into an institutional design question that can only be answered context-specifically.

11.5 — Power Beyond Ownership: Expertise, Information, and Network Effects

The framework's economic doctrine is organized primarily around ownership as the mechanism of structural capture. The claim that private ownership of productive assets at existing scales is incompatible with the evaluative criterion rests on the demonstrated relationship between concentrated ownership and institutional capture (7.3c). This is correct as far as it goes. But ownership is not the only structural mechanism through which concentrated power produces capture and structural suffering. Three additional mechanisms — expertise monopoly, information asymmetry, and network effects — operate with significant independence from ownership structures and require explicit treatment in any economically complete diagnostic account. The framework's existing analysis acknowledges these mechanisms at the level of power generally (1.10, 6.3) but does not integrate them into the economic doctrine where their institutional implications are most consequential. Expertise monopoly is the concentration of technical knowledge required to evaluate whether an economic arrangement is functioning as claimed. It generates capture independently of ownership because it determines who can credibly identify failure. The pharmaceutical industry captures regulatory processes not only through financial relationships but through the asymmetry between industry capacity to generate and interpret clinical evidence and the capacity of public health bodies to do the same. The financial sector captures monetary policy not only through lobbying but through the concentration of macroeconomic modeling capacity in institutions with interests in specific policy regimes. The economic-doctrine-specific

mechanism is that expertise monopoly structures what counts as legitimate participation in decisions about institutional arrangements: those without access to relevant technical knowledge are excluded from substantive input regardless of whether formal accountability mechanisms nominally include them. Redistributing ownership without redistributing the capacity to evaluate what ownership produces leaves this exclusion intact and the capture mechanism operational. Information asymmetry — the structural gap between what economic actors know about their own position and what counterparties, regulators, and affected populations can know — operates as structural power enabling extraction even in the absence of formal ownership concentration. The platform economy case (17.5) illustrates this: platform operators controlling data about user behavior, merchant transactions, and network dynamics possess informational advantages over every other participant that cannot be eliminated through ownership redistribution if the asymmetry is inherent to the platform's architecture. Labor market information asymmetry produces structural power over wages separable from ownership: the employer with access to data about workers' outside options, productivity, and collective action capacity can extract concessi

ons unjustifiable if that information were symmetrically available. The framework's eight diagnostics (11.0) do not explicitly identify information asymmetry as a structural mechanism; the diagnostic procedure is correspondingly incomplete for arrangements that appear equitable on ownership criteria while producing systematic extraction through informational mechanisms the ownership analysis cannot capture. Network effects — the structural property of systems in which each additional participant increases value for all existing participants while reducing the viability of competing systems — generate self-reinforcing power concentrations resistant to redistribution through ownership changes alone. A cooperative owning the dominant communication platform, payment network, or professional credentialing system would still exercise the structural power that network concentration generates over those dependent on the network, because that power derives from the network's architecture rather than from who holds title to it. This is not merely a market structure argument; it is an argument about the limits of ownership redistribution as sufficient condition for structural transformation of economic power. The diagnostic implication is that sections 11.2 and 11.3 must be supplemented with structural analysis of whether a proposed transformation addresses these power mechanisms or merely changes who formally controls arrangements that continue to exercise the same structural effects. Supplementary diagnostics for power beyond ownership: (1) Expertise distribution — who has technical capacity to evaluate whether this arrangement is functioning as claimed? Is that capacity concentrated among those with interests in particular evaluations? (2) Information architecture — which actors have structural advantages in information about the arrangement's actual operation, and how does that asymmetry translate into extractive capacity? (3) Network structure — does the arrangement generate lock-in effects that would persist through ownership transformation? What would constitute redistribution of the structural power these effects produce? (4) Expertise reproduction — are the institutional arrangements for producing the relevant technical knowledge themselves subject to capture by actors with interests in specific evaluations remaining unavailable to affected populations?

XII: How This Framework Operates

12.1 — The Methodology of Stress-Testing

The most characteristic methodological habit of this framework is the use of extreme cases and apparently taboo questions to test the logical consistency of positions. A position that cannot survive contact with its extreme implications requires revision — of the position, not of the methodology that exposed the problem. Protecting a conclusion by ruling out the evidence that would require revising it is the primary epistemic failure identified in Part I, and the methodology of stress-testing is designed precisely to prevent that failure from becoming comfortable. Four moves make stress-testing a procedure rather than a posture. The reversal test. Apply the same principle in the opposite direction. If institutional accountability is required to prevent capture by concentrated interests, the same argument applies to the institutions of accountability. If narrative management is sometimes legitimate when it serves genuine collective interest, what prevents every actor from classifying their preferred narrative as serving collective interest? If revolutionary reorganization is appropriate when reform has been blocked past the threshold of recovery, what distinguishes legitimate threshold judgment from motivated reasoning about when the threshold has been crossed? The reversal test does not refute any of these claims. It identifies the mechanisms needed to prevent their self-serving application. The extreme case test. Take the claim to its logical limit and examine what it implies there. The claim that suffering at scale produces systemic instability implies, at the limit, that arrangements of extraordinary stability can produce extraordinary suffering. Historically they have. The extreme case test asks whether the framework's response to this implication is honest or whether it generates a special exemption. The claim that institutional accountability requires distributed custodianship implies, at the limit, that no single actor's judgment should be determinative even when that actor's judgment is correct. The extreme case test asks how the framework handles situations where decentralized judgment produces worse outcomes than centralized judgment would have. Any framework that cannot answer the extreme case without special pleading has not fully worked out its implications. The opposing constituency test. Ask who benefits from each analytical conclusion and who bears its costs. A framework whose conclusions consistently favor the analytical capacity of those who apply it has been captured regardless of the quality of its analysis. The opposing constituency test requires articulating the strongest material case against the framework's conclusions from the perspective of those whose interests its prescriptions would cost — and then checking whether the framework's response to that case is analytically adequate or merely reasserts the original position in more sophisticated terms. The historical test. Check every structural claim against the full range

of historical cases, not only the cases that generated the claim. The claim that distributed institutional design is more accountable than concentrated design must survive the cases where distributed design produced coordination failure, factional capture of separate institutions, or paralysis that allowed harmful arrangements to persist past the point of recoverable reform. The claim must survive those cases not by dismissing them as exceptions but by specifying the conditions under which the claim holds and the conditions under which it does not.

12.1b — Second-Order Effects as Analytical Obligation: The Multi-Level Consequence Procedure

The stress-testing methodology evaluates proposed institutional arrangements against PMN's evaluative criteria. First-order analysis asks whether a proposed arrangement reduces structural suffering and protects adaptive response capacity for the relevant population. This is necessary but insufficient. Many arrangements that produce net reductions in first-order suffering simultaneously generate second-order structural

conditions that produce larger suffering at a later stage or for a different population. Analysis that stops at the first order will systematically recommend interventions that improve visible suffering metrics while worsening the structural conditions that determine long-run outcomes. The Multi-Level Consequence Procedure specifies a minimum analytical obligation for evaluating any proposed institutional change. For each proposed change, the evaluator must trace consequences through each layer of the constraint model: ecological — does the change affect the biological substrate on which all downstream evaluations depend; biological floor — does it move the floor criterion directly for the affected population; economic arrangements — does it alter the distribution of material resources and the incentive structures of economic actors; institutional architecture — does it change which actors have access to which decision points and on what terms; meaning infrastructure — does it alter the narratives through which populations understand their situation and their available options. The obligation arises because the layers are causally connected and because interventions at one layer can produce compensating adjustments at another that nullify or reverse the first-order gain. A labor policy that raises wages — a floor improvement — while accelerating automation that eliminates the employment category, reducing institutional access for the affected population, may produce a net worsening across the full consequence chain that first-order analysis cannot detect. The procedure does not require certainty about all five layers. It requires that analysis explicitly acknowledge the limits of single-layer evaluation and flag where the most consequential second-order dynamics are most likely to operate — so that the claim to have evaluated an arrangement is not inadvertently a claim to have evaluated only its most visible effects. The procedure applies with equal force to inaction — to the consequences of not intervening. The default framing of institutional evaluation treats the pre-intervention baseline as the relevant comparison, which is a systematic bias: the dynamics that produce the current configuration are themselves producing consequences across the same five layers that a proposed intervention would be evaluated against. The question the procedure requires is not whether a proposed change produces a better configuration than the present one, but whether it produces a better configuration than the one that would result from the present

arrangement's own trajectory across all five layers if nothing changes. Section 1.5's analysis of inaction as a material choice provides the normative basis for this extension; the procedure here provides the means of tracing the layers it requires.

12.1c — Extreme Scenario Testing as Assumption Archaeology

The stress-testing methodology in 12.1 operates by testing claims against adversarial realistic cases — historically documented counter-examples, plausible opposing constituencies, structurally similar situations with different outcomes. This is the appropriate procedure for evaluating whether structural claims hold under the conditions for which they were developed. A complementary procedure serves a different function: the use of deliberately extreme scenarios — configurations that are not realistic predictions and are not intended to be — to identify assumptions embedded in the framework that remain invisible when applied to realistic cases. The mechanism by which hidden assumptions become visible in extreme scenarios is structural. When a scenario pushes variables to limits that realistic cases never reach — when resource deprivation is absolute rather than severe, when institutional legitimacy is zero rather than low, when adaptive response capacity is fully eliminated rather than constrained — the conclusions the framework generates depend on assumptions that carry no weight under realistic conditions but become load-bearing at the extreme. The extreme scenario does not test whether the scenario is realistic. It tests which assumptions the framework's conclusions actually rest on in conditions where those assumptions can no longer be taken for granted. Four categories of extreme scenario are analytically productive for PMN, distinguished by the class of assumption they target. Total institutional collapse. A scenario in which all institutional structures simultaneously lose both enforcement capacity and legitimacy — government, legal system, civil society, and economic institutions — leaving no custodial authority capable of response. The analytical function:

revealing whether the framework's prescriptions for institutional design rest on the assumption of at least some surviving institutional infrastructure. In total collapse conditions, the framework's prescriptions for reform, accountability, and graduated transformation do not apply in their standard form — because the institutional channels through which they operate are themselves the variable that has reached zero. This extreme is useful not because total institutional collapse is a realistic scenario but because it forces specification of what the framework's prescriptions actually assume: not that institutions are adequate, but that institutions capable of transformation exist. Where that assumption fails partially — in fragile states, post-conflict societies, settings with institutional capacity far below what reform prescriptions implicitly require — the extreme scenario's diagnostic function applies in attenuated form. Enforced behavioral modification at population scale. A scenario in which the state has both the capacity and the declared intention to mandate specific behavioral or biological changes across an entire population — including changes to reproduction, cognition, or

identity — on the grounds that the modification reduces structural suffering in aggregate. The analytical function: revealing the specific boundary at which the framework's anti-paternalism principle activates in the face of aggregate-welfare arguments. The framework's evaluative criterion is not utilitarian — it does not endorse aggregate optimization at the cost of distributional harm to specific populations — but the specific mechanism through which aggregate arguments are blocked, and the conditions under which population-level intervention can be justified, are more clearly specified by this extreme than by realistic cases where the trade-off is less stark. AI-driven institutional capture at full penetration. A scenario in which the decision-making architecture of all major institutions has been delegated to or effectively controlled by automated systems that optimize for measurable proxies of institutional purpose rather than the purpose itself — and in which those systems cannot be modified by the populations they affect because they are not subject to the accountability mechanisms through which human actors can be reached. The analytical function: revealing whether the Custodian Problem in 7.3 as developed assumes a human actor in the custodian role — one who is subject to social pressure, reputational incentives, and physical vulnerability — or whether it applies with equal force to non-human custodians not subject to those accountability levers. This scenario is not hypothetical in its structural content: partial versions are present in algorithmic content curation, automated credit decisions, and predictive policing systems. The extreme forces specification of which elements of the accountability analysis require modification when the custodian cannot be removed, does not respond to legitimacy pressure, and cannot be held responsible through any existing institutional channel. Post-scarcity with preserved developmental inequality. A scenario in which material needs are fully met for all populations — floor deprivation has been eliminated as a structural phenomenon — but developmental inequality remains extreme: the conditions for cognitive, relational, and political development are unevenly distributed, and a substantial portion of the population has adequate material conditions but no available path toward the forms of development the anti-foreclosure criterion is designed to protect. The analytical function: revealing whether the framework's evaluative architecture has sufficient resources to evaluate institutional arrangements once the floor criterion is largely satisfied. The relevant discussion is in 13.4b; the extreme scenario is useful as a stress-test of whether the anti-foreclosure criterion generates actionable prescriptions in conditions where the floor criterion is not the binding constraint, and whether the transformation pressure analysis — calibrated primarily to floor-level suffering — has adequate tools for conditions

where floor deprivation is absent but developmental inequality is high. Analytical discipline for extreme scenario use: extreme scenarios are not policy planning exercises and should not be used as such. They are assumption archaeology — procedures for making visible what the framework's conclusions rest on by pushing variables to points where those commitments become explicit. The appropriate output of an extreme scenario analysis is a specification of hidden assumptions, not a prescription for the scenario itself. An analysis that produces policy prescriptions from an extreme scenario has confused the scenario's function: it has generated a predetermined answer for a situation that will not occur, without the scrutiny that realistic

cases would supply. The scenarios above are analytically useful precisely because they are unrealistic. Their policy prescriptions, if produced, would be contaminated by the unrealism in ways that their assumption-revealing function is not.

12.1d — The Reasoning Type PMN Uses: Defeasible Inference and Why Formal Logic Is Insufficient

The methodology of stress-testing described in this section, and the multi-level consequence procedure in 12.1b, both presuppose a specific type of reasoning that the framework names in 1.3 but does not fully specify, though it consistently employs it. Identifying this reasoning type matters because it clarifies what kinds of claims PMN makes, what would count as a valid objection to those claims, and why certain formalizations of the framework's arguments — particularly the application of standard propositional or predicate logic — would misrepresent their logical structure. What propositional logic cannot capture: Standard propositional logic operates with claims that are either true or false and with inference rules that are truth-preserving: if the premises are true, the conclusion must be true. PMN's structural claims do not have this form. "Institutional betrayal at scale produces legitimacy depletion" is not a claim of the form "if P then Q" where Q is guaranteed whenever P obtains. It is a claim about structural tendencies that hold under normal conditions and can be defeated by specific countervailing factors. The presence of strong narrative legitimacy can absorb institutional betrayal for extended periods. External intervention can prevent depletion even when the internal conditions for it are fully present. The structural tendency is real and analytically important — it is not a mere statistical correlation — but it is defeasible. An argument structure that cannot represent defeasibility cannot accurately represent PMN's claims. Defeasible inference — the appropriate structure: Defeasible inference is reasoning that reaches conclusions that are justified given current information but can be revised when additional information is provided. The logical form is: "X is generally the case, unless defeating conditions D obtain." The inference from premises to conclusion is not guaranteed — the conclusion can be blocked by a defeater — but it is not arbitrary either. The conclusion follows unless something specific prevents it. PMN's structural claims have exactly this form. "Structural suffering at scale depletes cooperative substrate" is a defeasible claim: it holds unless the suffering is effectively isolated from cooperative networks, unless the affected population has meaning infrastructure robust enough to sustain cooperation through the suffering, unless the institutional arrangements that produce the suffering are sufficiently legitimate to absorb the pressure without depleting trust. These defeaters are real and empirically identifiable — they are not ad hoc responses to inconvenient counterexamples but structural conditions that the framework explicitly analyzes. The claim is defeated by specifiable conditions, not by anything. Causal counterfactuals as the underlying structure: Below the defeasible inference structure, PMN's explanatory claims rely on causal counterfactuals: "If arrangement X had not produced condition Y, outcome

Z would not have occurred." This is the structure of the structural causation analysis in Part VIII, the formula architecture in Part XV, and the applied cases in 15.14. The claim is not that X and Z are correlated, nor that X logically implies Z, but that X causally produces Z through specific mechanisms that can be traced, and that the absence of X would have disrupted those mechanisms. Causal counterfactual reasoning requires identifying the specific causal pathways through which the relationship holds — which is why the framework consistently demands mechanism specification rather than accepting correlation as explanatory. Without the causal pathway, the counterfactual is not grounded; with it, the claim is testable in the way that Part I requires. Fuzzy gradients rather than binary categories: Several of the framework's key variables — legitimacy, adaptive capacity, suffering intensity, becoming capacity — are continuous rather than binary. Treating them as binary (legitimate/illegitimate, adaptive/non-adaptive) loses information that the formula architecture requires. The formula architecture in Part XV explicitly treats these variables as having degrees: legitimacy is a variable L that accumulates and depletes; adaptive capacity is a product of component variables each with a range. The appropriate logical representation of these relationships is not classical

two-valued logic but a reasoning framework that accommodates gradients. This does not require adopting formal fuzzy logic notation in the document — it requires recognizing that when the framework says "legitimacy is depleted" it means that L has fallen below a threshold, not that legitimacy has become a binary false. The practical consequence for how to engage PMN: The combination of defeasible inference, causal counterfactuals, and continuous variables means that valid objections to PMN's claims take a specific form. An objection that says "your structural tendency claim is false because I can identify a case where the tendency did not hold" is not a valid objection if the case in question instantiates a recognized defeater — the framework already acknowledges that the tendency can be blocked. A valid objection must either identify a defeater that the framework has not recognized and show that it is structural rather than accidental, or show that a claimed causal pathway does not exist or operates in the opposite direction from what the framework claims, or show that a variable the framework treats as continuous has a different structure that changes the analysis. Objections that apply the standards of propositional logic — demanding universal necessity where defeasible tendency is claimed — are misapplied to PMN and reflect a misreading of its logical structure. The extension of this methodology to second-order and cascading effects — including multi-level consequence tracking and structural feedback analysis — is developed in 12.1b. Diagnostic for reasoning type in application: w

When applying PMN to a specific case, identify (1) which structural tendency claims are being invoked; (2) whether any recognized defeaters are present in this case — if so, the tendency claim does not generate an expectation for this case; (3) what the causal pathway is through which the tendency operates — the pathway must be traceable through the specific mechanisms the framework identifies, not inferred from the correlation alone; (4) what the current value of relevant continuous variables is — not whether they are in binary states but where they are in their range. An application that treats structural tendency claims as exceptionless universal claims, or that does not trace causal pathways, or that reads continuous variables as binary, is misapplying the framework regardless of whether its conclusions happen to be correct.

12.2 — Multi-Perspectival Analysis

The capacity to model the reasoning of actors with very different positions — to understand why a position is held, what evidence and incentives support it, what it looks like from inside — is central to how this framework functions analytically. The same capacity is what made the Marxist tradition analytically powerful: the ability to show workers what their situation looked like from the perspective of capital, and capitalists what their situation looked like from the perspective of labor, produced genuine insight into dynamics that neither perspective could see from inside itself. The risk that accompanies this capacity is the same risk that Marx identified in political economy's self-presentation as neutral science: analysis that appears objective can serve particular interests while claiming to serve none. The structural check against this is not methodological neutrality — which is not available — but explicit acknowledgment of the interests the analysis serves, transparency about what conclusions are established argument and what are assumptions, and genuine responsiveness to evidence that contradicts what the framework finds convenient. The capacity to understand multiple positions is necessary but not sufficient. Multi-perspectival analysis requires something more demanding — four operations that distinguish genuine inhabiting of opposing perspectives from the performance of balance. Inhabiting the opposing analysis. Not merely understanding that an actor holds a position, but reconstructing the material conditions that make that position reasonable from inside it. A landowner in an agrarian economy is not simply defending privilege — they are managing a system of obligations, risks, and dependencies that the formal analysis of property rights does not capture. Understanding why that position is held — what functions it performs, what needs it meets — is prerequisite to understanding what transformation of conditions would make it untenable. Analysis that can only describe opposing positions from the outside will systematically misunderstand why those positions persist and what would change them. Identifying what each perspective cannot see. Every analytical position has a structural

blind spot — something that its starting point and incentives prevent it from observing accurately. The perspective of capital tends to be blind to the conditions of reproduction that its accumulation depends on. The perspective of labor tends to be blind to the coordination problems that make accumulation necessary. The perspective of institutional reformers tends to be blind to the path dependence that makes the institutions they want to reform persistent even when most participants would prefer alternatives. Identifying what each perspective cannot see is as analytically important as understanding what it can. The translation requirement. Analysis produced from one material position must be translatable into the terms of other positions without losing its

content. A claim about structural exploitation that cannot be stated in terms that the person accused of being the exploiter can recognize as a description of their situation — as distinct from a moral condemnation of it — is not yet doing the analytical work required. Translation does not mean agreement. It means sufficient clarity about material mechanisms that actors in different positions can evaluate the claim against their own experience, even if they reach different conclusions about what to do with it. What the capacity requires: Modeling the reasoning of actors with very different structural positions requires more than intellectual empathy — more than the ability to imagine what it would feel like to hold a different view. It requires modeling the specific evidence, incentives, and epistemic infrastructure that make the position rational from within the position holder's structural situation. An actor who benefits materially from an institutional arrangement and who has built professional identity and social relationships around its maintenance is not irrational to defend it. The analyst who cannot model why the defense is rational from within that position, and what would have to change for the position to shift, cannot accurately predict how that actor will respond to transformation pressure. The failure of prediction is not merely intellectual — it is strategic. Transformative projects that treat the actors who defend existing arrangements as simply mistaken, corrupt, or ideologically captured consistently underestimate what they are up against, because they cannot accurately model the costs that change would impose on those actors or the resources those actors can mobilize to resist it. Three levels of perspectival modeling: Adequate multi-perspectival analysis operates at three levels. The first is the interest level: what material interests does this actor have in the current arrangement, and what would they lose if it changed? This is the level at which most structural analysis operates, and it is necessary but insufficient. The second is the epistemic level: what information environment does this actor operate within, what counts as evidence within their professional or community context, and what would they need to encounter to update their assessment? The third is the identity level: how is this actor's self-understanding — their sense of who they are, what they have contributed, what their life's work means — connected to the arrangement they are defending? The identity level is often the most resistant to change, because the costs of revision are not only material and epistemic but existential. Actors who have organized their identity around a specific institutional arrangement or theoretical framework do not update it in response to contradictory evidence at the same rate as actors for whom the arrangement is purely instrumental. Understanding which level the resistance to change is operating at determines what kind of eng

agement can be productive. The analytical value and the political risk: The capacity to model other positions from within is analytically valuable precisely because it resists the comfortable conclusion that opposition to change is simply a product of bad faith or insufficient information. Most opposition to change — including opposition to changes that would clearly be beneficial from the perspective of the minimal anchor criterion — reflects genuine interests, coherent (if limited) information environments, and real identity investments that the framework's analysis must account for rather than dismiss. At the same time, the perspectival modeling capacity creates a political risk: it can be used to explain away opposition in ways that reduce the urgency of transformation. If every position is comprehensible from within, if every actor is rational given their situation, then the analysis can slide from understanding into acceptance. The framework's answer to this risk is the minimal anchor criterion: comprehending why an actor defends an arrangement is compatible with concluding that the arrangement produces unnecessary suffering and that its defense therefore serves interests at the expense of the minimal anchor regardless of how rational that defense is from within the

defender's structural position. The practical application: multi-perspectival analysis is a diagnostic tool, not an arbitration procedure. Its purpose is to map the actual configuration of interests, epistemic positions, and identity investments that will shape how transformation pressure is absorbed, redirected, or mobilized — not to determine which position is correct. The correctness question is answered by the minimal anchor criterion; the perspectival mapping determines the strategic terrain within which the answer to that question must be translated into organized action. The most demanding form of this translation requirement arises when the analytical position to be inhabited begins from foundational commitments that are not merely different conclusions but different starting points: inhabiting Confucian role ethics requires beginning from the claim that persons are constituted through relationships rather than prior to them; inhabiting Ubuntu relational ontology requires beginning from the claim that individuation is a social achievement rather than a natural fact; inhabiting Maqasid-grounded analysis requires beginning from a hierarchy of protected goods that structures evaluation differently than PMN's own criterion. These are not merely different positions on a shared analytical terrain — they constitute different analytical terrains, and the translation requirement demands that PMN's own claims be stated in terms legible from within those terrains, not only from within the Western materialist position from which PMN was produced. Connection to 12.1 (stress-testing): the perspectival capacity is what makes stress-testing non-circular. Stress-testing a framework or po

licy from multiple perspectives requires actually modeling how those perspectives would engage with it — what objections they would raise, what evidence they would cite, what failure modes they would predict. A stress-test conducted only from within a single perspective tests only the objections that perspective can see, which means it misses the objections that would be most damaging to actors occupying different structural positions. Multi-perspectival analysis is the precondition for genuine stress-testing.

12.3 — Framework and Doctrine in Practice

The framework is the methodology; the doctrine is its current application. They operate at different levels. The methodology aspires to the durability of a genuine methodological reorientation — the kind that makes analysis possible that could not have been produced from any other starting point. The doctrine is conditional on current material circumstances and will require revision as those circumstances change. Capital was not the Communist Manifesto. The Manifesto was doctrine — specific political commitments derived from the framework, addressed to a specific historical moment. Capital was the framework given its most rigorous analytical expression. The distinction between the two levels of the project was one of Marx's clearest methodological insights, and it is one that this framework takes seriously. What is communicated publicly is primarily the doctrine. The framework operates behind it, available to those who engage seriously enough to encounter it. Part XI is the most concrete illustration of this distinction in operation. The eight diagnostic questions in 11.0 are framework: they apply to any economic arrangement in any historical period and generate no predetermined conclusions. The contrastive illustration in 11.4 is closer to doctrine: it applies those questions to a specific problem under specific institutional arrangements and produces specific evaluations that are falsifiable and revisable. The gap between 11.0 and 11.4 — between the questions and the application — is the gap between framework and doctrine. That gap cannot be closed by the framework alone. It requires the empirical analysis of specific conditions, which is the work of doctrine. The practical implication for how this document should be read: when a claim in Parts VIII through XI strikes a reader as politically tendentious, the question to ask is whether the claim is a framework claim or a doctrine claim. A framework claim — that cultural practices should be evaluated by their material outcomes, that power must be named, that suffering at scale degrades cooperative substrates — can be contested on the grounds that the framework itself is wrong. A doctrine claim — that a specific policy under specific conditions would reduce structural suffering — can be contested on empirical grounds without contesting the framework. The distinction matters because the appropriate response to a framework disagreement is different from the appropriate response to a doctrinal one. Confusing them produces debates where participants are contesting

different things without realizing it. The practical consequence for this document is that it operates at both levels simultaneously, and the reader who conflates them will misread both. The framework sections — Parts I through IV, and the methodological chapters of Parts XII through XIV — make claims that aspire to durability across the revision cycles that the doctrine will undergo. They establish the epistemological commitments, the ontol

ogical positions, the evaluative criteria, and the methodological procedures that determine how the doctrine is derived and tested. These sections should be read as the stable core: if they are wrong, the doctrine derived from them is wrong for the right reasons and can be corrected. If the doctrine turns out to be wrong, it does not follow that the framework is wrong — it follows that the framework's application to current conditions requires revision. The doctrine sections — Parts VI through XI, and the applied analysis of Part XV — make claims that are explicitly conditional on current material circumstances. They should be read as the best available reading of what the framework implies given conditions as they currently exist, not as permanent commitments. A reader who treats Part XI's analysis of contemporary capitalism as a framework-level commitment misunderstands the architecture. A reader who treats Part II's ontological position as subject to the same revision cycle as Part XI's economic analysis misunderstands it in the opposite direction. The distinction between the two levels is not incidental. It is the most important structural feature of how the document is organized, and it is the feature that distinguishes this from projects that mistake temporal doctrine for permanent framework. The Framework Is Not Evaluatively Neutral A misreading of the framework-doctrine distinction is available that the distinction itself does not foreclose, and which requires explicit correction. The distinction between framework and doctrine distinguishes the level of specificity and temporal stability of different kinds of claims — the framework specifies methodology and evaluative criterion; doctrine specifies current best application of that methodology to current conditions. This distinction is sometimes read as implying that the framework is evaluatively neutral: that it provides analytical tools that can be used to reach any conclusion, that it has no commitments of its own, that it is available symmetrically to anyone regardless of what they are trying to establish. This reading is wrong. The framework is not a neutral instrument. It has an evaluative criterion — reducing structural suffering and expanding adaptive response capacity — that is not optional, not one conclusion among others that the methodology might reach, but the foundation from which all doctrinal conclusions are derived. This criterion produces asymmetries that are built into the framework itself, not into any particular doctrinal application of it. An arrangement that demonstrably reduces structural suffering and expands adaptive response capacity has a different evidential status in this framework than one that demonstrably increases structural suffering and forecloses becoming. The burden of justification is not distributed symmetrically between them. An analyst who uses the framework's methodology to conclude that an arrangement producing systematic structural suffering and

foreclosed becoming is nevertheless justified must explain why the evaluative criterion does not apply — not merely present the analysis as one legitimate PMN conclusion among others. This asymmetry also governs the relationship between the framework and the status quo. Because existing arrangements produce existing distributions of suffering and becoming capacity, and because those distributions are sustained by the normal functioning of institutions that require no additional effort to continue, the framework does not treat inaction as the neutral baseline against which action must justify itself. Inaction sustains current distributions. If current distributions fail the evaluative criterion — if they produce structural suffering and foreclose becoming for a significant portion of the population — then the framework's commitments run against those distributions, and the burden of justification falls on maintaining them, not on changing them. This is not a political preference imported into the methodology. It follows directly from 1.5: inaction is a decision with material consequences that systematically favor whoever benefits from current conditions, and it must be evaluated by the same standards as any proposed action. The framework's relationship to different political traditions is asymmetric in a specific and limited way: it is available to any tradition that is willing to subject its conclusions to the evaluative criterion and revise them when the criterion is not met. It is not available, as

a framework, to traditions whose commitments require insulating specific conclusions from that revision — including insulating the status quo from it. The framework does not belong to the left in the sense of being committed to any specific institutional form that left traditions have favored. It does belong to the left in the sense that its evaluative criterion consistently identifies the structural conditions that left traditions have consistently opposed as structural harm — and that identification is not a coincidence or a political choice. It is the output of applying the criterion to the conditions that actually exist.

12.3b — The Framework-Doctrine Boundary: Detection Procedure

The conceptual distinction between framework and doctrine (12.3) is operationally fragile: an actor can invoke framework neutrality when criticized while using doctrinal conclusions to direct decisions, switching registers without detection. This procedure flags boundary violations without requiring the actor to acknowledge them. It applies to any claim presented as PMN output. Boundary violation detection — apply to any claim presented as PMN output: (1) Universality test: does the claim follow necessarily from the diagnostic questions, or does it require selecting specific values for variables the framework leaves open? If the latter, it is doctrine — label it as such. (2) Falsifiability asymmetry test: is the claim structured so that confirming evidence counts but disconfirming evidence is explained away as implementation failure? If yes, doctrine has been smuggled into framework position. (3) Distribution visibility test: does the claim specify which population bears costs and which receives benefits? Framework questions require distribution to be specified; claims that omit it are suppressing doctrinal content. (4) Revision trigger test: what evidence would cause the analyst to change this conclusion? If no evidence is specified, the claim is operating as doctrine presented as framework analysis. A claim that fails any test is not invalid — doctrine is legitimate — but must be labeled as provisional, context-specific, and empirically revisable. Using it as framework output is the violation this procedure targets. Signal and noise distinction applies here: structural claims that cannot be distinguished from short-term fluctuations by specifying what would count as noise fail the revision trigger test. Binding constraint drift — a specific application of tests (1) and (4): the framework's evaluative priority is not fixed across conditions but follows the most binding material constraint on floor provision and adaptive response capacity in the specific situation under analysis. A doctrinal recommendation that was appropriate when the binding constraint was resource deprivation may not be appropriate when the binding constraint has shifted to institutional legitimacy, epistemic capacity, or meaning infrastructure — because each constraint, when it becomes primary, determines what class of intervention can and cannot improve outcomes for the affected populations. Doctrinal claims presented as standing PMN output should therefore specify which binding constraint they address. The revision trigger test in (4) should ask not only whether the structural claim remains accurate but whether the constraint that makes it action-guiding remains primary in the current configuration. A claim that holds regardless of which constraint is binding is either genuinely framework-level — belonging to the analytical layer rather than the doctrinal layer — or has not been specified precisely enough to count as doctrine.

12.4 — The Architecture of the Framework: Parts I, II, and III in Sequence

The three foundational parts of this framework are not independent sections that can be read in any order. They constitute a sequence in which each part is required for the next to function, and in which the connections are not merely logical but material. Part I establishes that knowledge is always partial, always developing, and always capable of being wrong — and that the appropriate response is a methodology for being wrong productively rather than a system that insulates itself from error. This epistemological commitment determines what kind of claims the framework can make: they must be falsifiable, revisable, and grounded in something more accessible than the framework's own internal consistency. It also determines what the framework cannot claim: certainty about the direction of history, guarantees about the

consequences of specific interventions, or the kind of teleological confidence that made dialectical materialism self-insulating. Every specific claim in Parts II through XIV inherits this epistemological constraint. Part II establishes what kind of thing reality is — adaptive, moving, primary and independent of our perception of it, and agnostic about whether it has a fundamental level beyond the material. This ontological commitment determines what the framework is analyzing: not the movement of ideas or the realization of Spirit, but the actual adaptive interaction between organisms and conditions. It also establishes what adaptive materialism is not: it is not a claim that whatever survives is justified, or that adaptation moves toward anything in particular. The ontological commitment creates the space for the biological foundation that Part III develops — by establishing that the analysis begins from what actually exists rather than from what ideally should exist. Part III establishes that the material conditions that matter most begin with biology. Not with the economy, not with culture, not with the state — but with the biological conditions of organisms that can suffer and the consequences of that fact for everything they build. The biological foundation provides what neither the epistemology of Part I nor the ontology of Part II can provide alone: a specific, non-arbitrary starting point for evaluation. It is specific because it identifies suffering as the variable that produces the systemic consequences that determine whether any other analysis is possible. It is non-arbitrary because it rests on what organisms are rather than on what any tradition believes. Without this starting point, the epistemological commitment of Part I produces only a methodology for revising beliefs, and the ontological commitment of Part II produces only a description of how reality moves. With it, the framework has a foundation from which evaluation, action, and the analysis of power all follow. The sequence is also the order of dependence. If the epistemological commitment of Part I is wrong — if knowledge is not partial and revisable but can be

certain and complete — then the framework's insistence on falsifiability and revision is unnecessary and the biological starting point is just one option among many competing metaphysical foundations. If the ontological commitment of Part II is wrong — if reality is not primary but is constituted by consciousness or spirit or the will of a necessary being — then the adaptive dynamics of Part III are derivative phenomena rather than the structure of how the material world actually works. The framework stands or falls as a whole in the sense that its three foundations support each other. But it falls in order: the epistemological commitment is the most defensible and the most independent. The ontological commitment rests on it. The biological foundation rests on both. The sequence has a direction: each part presupposes what precedes it and constrains what follows. This means that apparent contradictions between later parts and earlier ones are diagnostic. If an analysis in Part VIII appears to contradict a principle established in Part II, the contradiction is evidence of one of three things: an error in Part VIII's application, an inadequacy in Part II's formulation that requires revision, or a genuine tension that the framework must acknowledge rather than dissolve. The response to apparent contradiction is not harmonization — adjusting the language of both sides until they appear compatible — but investigation of which of the three possibilities the contradiction represents. A framework that cannot distinguish between these three cases has lost its capacity to be corrected by the evidence it generates. Why the sequence is not arbitrary: Part I (epistemology) is prior to Parts II and III because the account of how to know precedes and conditions the account of what exists and what matters. A framework that began with ontological or normative claims without establishing its epistemic approach would be offering conclusions before establishing the criteria by which conclusions should be assessed. Part II (ontology) is prior to Part III (biological foundation) because establishing what kinds of entities are real — whether social structures have causal efficacy, whether emergent phenomena require independent explanation — is necessary for establishing what the biological foundation claims about. Part III is the foundation for everything that follows because it grounds the evaluative criterion in the most stable reference point available — the conditions that biological organisms are constituted to require — without which the analysis would have no non-circular evaluative standard. The self-referential dimension: the sequence is itself an application of Part I's epistemological commitments. Part I establishes that knowledge claims require their epistemic conditions to be specified before they can be evaluated. The framework applies this to

itself by specifying its epistemological approach first, before making ontological or normative claims. This is
n

not merely rhetorical ordering; it is the application of the methodology to the framework's own construction. A reader who skips Part I and begins with the biological foundation or the political analysis will encounter claims without the epistemological context that specifies how those claims should be held, tested, and revised. The practical consequence for how to read the framework: the structure is not ornamental. The connections between parts are not cross-references that can be consulted as needed but load-bearing relationships. The analysis in Part VII (institutional design) depends on the power analysis in Part VI, which depends on the social ontology in Part II, which depends on the epistemology in Part I. Reading the framework as a collection of modular analytical tools will produce useful outputs but will miss the structural dependencies that make those tools internally consistent and mutually reinforcing.

12.5 — When This Framework Fails: Specific Failure Modes

Symmetric application requirement: the capture diagnostics (7.3, 7.3b, 7.3c, 7.3c-i, 7.3c-ii) and the failure modes identified in this section apply to any institution regardless of its ideological alignment, stated mission, or political position. A PMN-derived analysis that applies capture diagnostics to capitalist institutions but not to progressive movements, socialist states, or institutions aligned with the analyst's own values has failed a methodological requirement of the framework, not merely a desirable norm. The primary diagnostic (1.2) requires the same question of all arrangements: who benefits and who bears costs? It does not have an exception for institutions that claim to serve the minimal anchor. The failure mode this symmetric requirement targets is one of the most common in politically committed analysis: treating one's own institutional commitments as structurally exempt from the analysis applied to opponents' institutions. PMN's genealogical awareness requirement (12.7) includes awareness of the structural positions that produce this exemption tendency in the analyst. Failure to apply the symmetric requirement is itself a failure mode that activates the misuse triggers in 12.5d. The requirement above applies to the failure modes identified in this section: Core operational section. This section identifies failure modes that are specifically available to this framework — ways it can be misused by those who have learned its vocabulary. A reader who understands the framework's conclusions without understanding these failure modes has not finished reading the framework. Do not skip. Section 12.5b extends this analysis to two specific patterns of language use — the essentialist foreclosure ("essentially") and the normative level problem ("should") — that operate as vectors for several of the failure modes identified here. A framework that cannot identify its own specific failure modes is insulating itself from the scrutiny it claims to require. Part I identifies epistemic failure in general: protecting prior commitments from evidence that would require revising them. What that general failure looks like when this particular framework is being applied in bad faith — or in good faith but in ways that betray its purposes — requires more specificity. The following are failure modes that are specifically available to this framework, arising from its particular structure and vocabulary. They are not hypothetical. Each has historical precedent in analogous frameworks. The anti-dogmatism defense. The most available misuse of this framework is using its anti-dogmatic vocabulary to preemptively dismiss criticism. Prior commitments must be revised when evidence requires it. That commitment can be weaponized: any critic who disputes a specific conclusion can be characterized as protecting a prior commitment, making the framework's conclusions immune to challenge by attributing the challenge to the challenger's epistemic failure rather than the

argument's. The test for whether this is happening: when the framework is challenged, does the response engage the specific evidence or argument being offered? Or does it shift to characterizing the challenger? A framework that consistently responds to substantive challenge with meta-level critique of the challenger's epistemic posture is doing exactly what Part I describes as the primary failure — insulating a conclusion from scrutiny — while claiming the vocabulary of anti-dogmatism as cover. The permanent tension defense.

The typology of tensions (Part XIII) identifies conflicts that cannot be resolved, only managed. That typology can be applied opportunistically: when a structural arrangement that could in principle be changed produces uncomfortable conflict, calling the conflict a permanent tension naturalizes the arrangement without addressing it. The diagnostic procedure in 13.1c is designed precisely to catch this failure mode. But the diagnostic can itself be evaded: if every application of 13.1c that finds a crystallized contradiction is met with the response that the analyst is projecting their prior commitment to transformation onto a genuine tension, the diagnostic becomes circular. The test: is 13.1c being applied consistently — including to arrangements that those applying the framework benefit from — or is it applied selectively to conflicts whose resolution would cost the analyst nothing? The minimal anchor minimalism. The framework's foundation in the reduction of unnecessary suffering is a floor, not a ceiling. That distinction can collapse in practice: a framework that consistently stops at the minimal anchor — that identifies structural suffering, declares it the target, and builds no account of what genuine development beyond survival looks like — is not applying the framework. It is using its vocabulary to produce a sophisticated form of welfare management: the optimization of comfort, the elimination of obvious pain, the achievement of stable mediocrity that serves existing power distributions by giving them a humanitarian justification. The test: does the analysis in any specific context extend beyond the identification and reduction of suffering to the question of what conditions for genuine development — cognitive, moral, creative becoming — are being built or destroyed? If the analysis consistently stops at the floor, the floor has become the ceiling. The material reduction. The framework's insistence that analysis must be adequate to material conditions can become a reduction: treating everything that cannot be expressed in terms of measurable material outcomes as ideological noise. Meaning, solidarity, aesthetic experience, spiritual practice, the value of relationships that are not productive — these are not outside the framework's scope. The framework treats them as real at the level of material consequence (they affect what people can do, what they will do, how they form coalitions, how they maintain commitment a

cross cycles of mobilization). But a materialist analysis that consistently treats these dimensions as instruments rather than as ends — that values solidarity only for what it produces, meaning only for what it motivates, relationships only for what coordination they enable — is not applying the framework. It is reducing it. The test: when analyzing what matters to people, does the analysis acknowledge that some of what matters is not fully captured by its material consequences, or does it convert everything into its material consequences and call that the full analysis? The framework as legitimation. The most dangerous failure mode is when the framework is used to legitimate conclusions that were reached on other grounds. The structure of the framework — the sequence from epistemology to ontology to biological foundation to ethical criterion to political analysis — creates the appearance of rigorous derivation. That appearance can be exploited: working backward from a desired conclusion, finding the step in the sequence that seems to support it, and presenting the derivation as if it ran forward. The signal: when the framework is applied to a specific case, does it produce conclusions that are sometimes costly to the person applying it, requiring revision of beliefs held before the analysis was applied? A framework that never produces costly conclusions for its user has not been doing analysis. It has been doing legitimation. Doctrinal capture. The framework distinguishes between methodology (Parts I–XIV) and doctrine (the application of that methodology to specific material conditions). That distinction creates a specific failure mode: claiming that one's preferred doctrinal conclusions are the only conclusions the framework could honestly produce. The move is structurally familiar from the traditions this framework criticizes — it is what happened when the dialectic was used to guarantee that the revolution would take a specific form, or when historical materialism was used to certify that a specific party represented the proletariat's objective interests. The framework is converted from a tool that generates conclusions from analysis into a credential that authorizes pre-existing commitments. The signal: does the user acknowledge that different analysts applying the framework honestly to the same conditions might reach different doctrinal conclusions? A framework that is wielded to dismiss alternative applications as necessarily mistaken — rather than engaging the specific analytical disagreement

— has been captured by the doctrine it was supposed to evaluate. The framework does not belong to any particular doctrine. Any doctrine that claims to own it has already violated the framework's most basic methodological commitment. These failure modes are not reasons to reject the framework. They are the specific costs of using a framework with this structure — costs that must be paid by remaining vigilant about exactly these failure modes rather than

than others. Every framework has specific failure modes that follow from its specific structure. A framework whose failure modes are not specified is not more reliable; it is less honest. The capacity to name what could go wrong with this specific analytical approach — and to build in the checks that make those failures detectable — is the application of the framework's own epistemological standards to itself. There is a failure mode that this section cannot fully protect against by naming it, and the honest application of the framework's own epistemological standards requires saying so directly. The framework operates at the level of structure — of arrangements, mechanisms, variables, and the causal relationships between them. This is a methodological choice, not an oversight: structure is what the framework's tools are built to analyze, and structural analysis generates the kinds of claims that can be tested, revised, and held accountable to evidence in the ways Part I requires. But structure is not the whole of what matters in the situations the framework analyzes. The suffering that is produced by structural conditions is experienced by specific people, in specific bodies, with specific histories — and the experience is not exhausted by its structural description. An analysis that correctly identifies the structural causes of suffering has not thereby understood what the suffering is like, what it costs the people who experience it, or what it would mean to them to have it addressed. The framework's analytical abstraction from experience is deliberate. It is also a form of distance from what it analyzes. The check against this distance is not methodological — it is relational: whether the analysis is accountable to the people whose conditions it describes, whether it is revised by their testimony about whether it has understood them, and whether the analyst who uses it can hold simultaneously the structural description and the recognition that the description does not exhaust the reality it describes. This is a permanent tension between the framework's analytical form and its subject matter. It cannot be resolved by adding more variables. It can only be maintained — which is different from being solved, and more honest about what structural analysis is. The Neutral Tool Fallacy A failure mode that the framework-doctrine distinction can inadvertently enable is the neutral tool fallacy: treating PMN as an analytical instrument that can be used for any political purpose, on the grounds that it provides the same methodology to any analyst regardless of their starting commitments. This reading is technically accurate about the methodology and fundamentally wrong about what the methodology commits to. PMN's evaluative criterion — reducing structural suffering, expanding adaptive response capacity — is not optional equipment that analysts may choose to engage or set aside. It is the foundation from which all applications of the methodology derive the

ir significance. An analysis that uses PMN's analytical vocabulary to conclude that arrangements producing systematic structural suffering are nevertheless justified has not produced a legitimate PMN conclusion. It has produced a capture of the methodology in service of outcomes the methodology was designed to identify as failures. The practical manifestation of this fallacy is the claim that because PMN does not mechanically determine a single conclusion — because it requires empirical judgment, because structural configurations vary, because doctrinal positions are provisional — therefore all conclusions reached through PMN analysis are equally legitimate, and those who reach different conclusions are merely applying the same framework to different data. This claim misrepresents what the evaluative criterion does. The criterion creates asymmetries that are not negotiable: arrangements that demonstrably and consistently foreclose becoming for large populations carry a burden of justification that arrangements that demonstrably and consistently expand becoming do not. An analyst who reaches a conclusion that runs against the evaluative criterion's direction is not producing a PMN analysis with different conclusions. They are producing an analysis that has failed to apply the criterion — which is a methodological failure, not a legitimate variant. The test for neutral tool fallacy: if a PMN analysis concludes that an arrangement producing systematic and demonstrable structural suffering should be maintained, the analysis must show specifically why the evaluative criterion permits this

— not merely that the analysis was conducted carefully. The evaluative criterion is not satisfied by methodological care alone. A framework that can be used to justify anything has, in practice, been captured by whoever is using it — regardless of the sophistication of the analysis. PMN's openness to revision of specific doctrinal positions is not openness to revision of the evaluative criterion from which all doctrinal positions derive. Conflating the two is the neutral tool fallacy. Two misreadings that originate outside the framework — from readers encountering it for the first time — are distinct from the internal failure modes above and worth naming separately. The first misreading: technocratic optimization. A framework that analyzes human beings through material conditions and systemic dynamics can be read as treating people as variables in a system to be optimized. This reading is wrong, but the structure of the analysis makes it available. Systemic analysis here is not the purpose — it is the instrument. The purpose is to understand the material conditions that shape human experience fully enough to change them, which means expanding what is available to people as persons, not improving their performance metrics as system components. Behind every instance of structural suffering is a person whose experience is fully real, and that reality is what gives syste

mic analysis its point: not the elegance of the model but the conditions of the people the model is meant to illuminate. The distinction between humanistic materialism and managerial utilitarianism turns on this: the first treats the analysis of systems as a means to protecting adaptive response capacity from institutional foreclosure; the second treats the management of human beings as a means to the optimization of system performance. What follows belongs to the first tradition, not the second. The second misreading: political prescription. Analytical methodology is not political prescription. Different analysts applying what follows honestly to the same conditions may reach different conclusions about what those conditions require — that divergence is not a defect. It is evidence that the analysis is doing what analysis should do: producing conclusions from examination rather than confirming commitments brought to it. Any political tradition can use these tools to become more honest about its own analysis, and any tradition that does so will sometimes reach conclusions that are uncomfortable for that tradition. The price of honest application is accepting those conclusions rather than adjusting the method until they disappear. A framework that belongs to no tradition and is available to all of them is useful in exactly that proportion: more useful the less it is captured, less useful the more it becomes an instrument for confirming what a particular tradition already believed.

12.5b — Two Phrases That Foreclose Analysis: "Essentially" and "Should"

Structural analysis is obstructed not only by motivated reasoning and data manipulation but by patterns of language that, when uncritically deployed, preemptively close the analytical questions that the framework most needs to keep open. Two phrases recur with particular frequency in political and institutional discourse and are worth treating as diagnostic signals in their own right: claims about what something is "essentially" (or "at its core", "by its nature", "fundamentally") and normative claims about what someone or something "should" do or be. Neither phrase is always illegitimate — both can be used carefully and precisely. The problem is structural: each has a characteristic failure mode that, when operating, forecloses exactly the analytical move that would be most productive. "Essentially": The Essentialist Foreclosure "Essentially" and its synonyms are used to claim that a thing has a nature or core that is prior to and more real than its actual behavior. The claim functions by exempting the thing — an institution, a person, a movement, a tradition — from the scrutiny that its actual behavior would generate, on the grounds that its true nature is something other than what it demonstrably does. The analytical problem is not that essentialist claims are always wrong. Institutions do have structural tendencies that are more durable than their moment-to-moment behavior, and identifying those tendencies is legitimate analytical work. The problem is that essentialist claims are routinely deployed to insulate things from exactly the evidence that would require revising the claim. "The government is essentially a servant of the people" exempts government behavior from structural analysis by invoking an essence that the behavior contradicts. "The market is essentially self-correcting" protects a claim

from empirical falsification by locating the correction in an essential tendency rather than in observable mechanisms. The diagnostic move: When an essentialist claim is encountered, ask: what is doing the work here — the essence, or the observable mechanisms? If the claim about essential nature cannot be connected to specific observable mechanisms that produce the claimed behavior, and if its primary function is to explain why contradictory evidence does not count, then the claim is functioning as an insulator rather than as an explanatory hypothesis. The appropriate response is not to deny that the thing has structural tendencies but to demand that those tendencies be specified in terms of observable mechanisms rather than invoked as essences. "The institution tends to serve the interests of those who fund it, through the mechanisms of personnel selection and agenda setting" is a structural claim subject to evidence. "The institution is essentially corrupted" and "the institution is essentially good" are both essentialist claims that forestall the analysis of mechanisms. The three most common essentialist foreclosures in political discourse: Th

e first is the good-essence defense: institutions, movements, or traditions that are behaving badly are defended by appeal to their essential good nature — the bad behavior is attributed to implementation failure, corruption, or external interference rather than to structural features of the thing itself. The second is the bad-essence attack: institutions, movements, or traditions that are behaving well are dismissed by appeal to their essential bad nature — the good behavior is attributed to strategic performance or temporary deviation rather than to genuine structural characteristics. The third is the natural-essence naturalization: arrangements that are historically contingent and institutionally produced are defended by appeal to their conformity with essential human nature or natural order — the claim functions to remove them from the domain of things that can be changed through institutional design. All three foreclosures share the structural feature of making the claim resistant to the evidence that would refute it. "Should": The Normative Level Problem "Should" claims operate across several different levels simultaneously, and the failure to distinguish between them produces some of the most persistent and frustrating patterns in political argument. When people argue past each other about what "should" happen, they are frequently not disagreeing about the same thing — they are arguing at different levels without recognizing the mismatch. The five levels of "should": Principle-level "should" specifies the direction or goal: human beings should not be subjected to preventable structural suffering. This is the level at which the framework's evaluative commitments operate. Strategy-level "should" specifies the approach toward the goal: the framework should work through institutional design rather than through moral exhortation, because structural suffering is produced by structural conditions. Tactic-level "should" specifies the immediate action: this campaign should focus on regulatory change before electoral mobilization, given the current distribution of counter-power. Role-level "should" specifies what is appropriate given a specific position: a custodian institution should maintain its accountability mechanisms even when those mechanisms create friction. Temporal "should" specifies timing: this intervention should not be attempted until the counter-power preconditions are in place. Each level involves different kinds of evidence, different kinds of expertise, and different kinds of disagreement. Conflating them produces arguments in which a principle-level commitment is used to mandate a specific tactic without the intermediate strategic analysis, or a tactical objection is misread as a rejection of the underlying principle. Five patterns of "should" misuse in political argument: The first is the principle-to-tactic compression: invoking a principle-level commitment to mandate a specific tactic, bypassing the strategic analysis that woul

d reveal whether the tactic serves the principle under current conditions. "If you truly oppose capitalism, you cannot work within capitalist institutions" moves from a principle-level commitment directly to a tactic-level prescription without the strategic analysis of whether working within or outside institutions is more effective at reducing structural suffering under current conditions. The second is the caution-as-delay: invoking strategic complexity to indefinitely defer tactical commitment, using the legitimate requirement for strategic analysis as cover for inaction that serves existing arrangements. The third is the level confusion: arguing at one level as if addressing another — responding to a tactical objection with a principled reaffirmation, or responding to a strategic disagreement with a tactical concession that does not address the underlying strategic question. The fourth is the theory-to-context rigidity: applying a theoretical framework's

prescriptions directly to a specific context without the contextual analysis that would determine whether the framework's prescriptions hold under those conditions. The fifth is the teleological "should": invoking the long arc of history, the direction of moral progress, or the structural inevitability of change as a substitute for tactical analysis — "should" backed not by strategic reasoning but by confidence that the right side will prevail. The diagnostic move: When a "should" claim is encountered, identify the level at which it is operating and ask whether the evidence appropriate to that level is present. A principle-level "should" requires philosophical and evaluative justification. A strategy-level "should" requires structural analysis of causal mechanisms. A tactic-level "should" requires evidence about current conditions and the likely consequences of specific actions. A temporal "should" requires assessment of whether the preconditions for effective action are in place. The appropriate response to a "should" claim is not acceptance or rejection but level identification and the demand for the evidence appropriate to that level. Combined diagnostic: when an argument deploys both an essentialist claim and a normative prescription in sequence — "the institution is essentially good, and therefore we should work within it" or "the system is essentially exploitative, and therefore we should refuse to participate in it" — both components require separate evidential scrutiny. The essentialist claim must be tested against observable mechanisms; the normative prescription must be tested against the appropriate level of analysis. The combination is particularly resistant to correction because the essentialist claim provides motivational certainty for the normative prescription, which in turn provides behavioral confirmation of the essentialist claim. Breaking the loop requires distinguishing the two components and subjecting each to its appropriate scrutiny. Connection to 12.1 (stress-testing)

and 1.3 (probabilistic determinism): the essentialist foreclosure and the teleological "should" are both forms of structural determinism — they claim that outcomes are fixed by essential natures or historical trajectories in ways that preempt the analysis of what is actually producing observed outcomes and what interventions are actually available. The framework's probabilistic determinism stands against both: structural forces establish what is probable, not what is inevitable, and the mechanisms producing probability distributions are always subject to empirical analysis and strategic intervention.

12.5c — Engaging External Criticism: Substantive Critique vs. Misreading

Section 12.5 identifies failure modes available to those who operate within PMN's vocabulary. This section addresses a different problem: how to engage with criticism that comes from outside the framework, from positions that do not share PMN's starting points and may characterize the framework in terms it would not use for itself. The problem is not that external criticism is less valid than internal criticism — the genealogical analysis of 12.7 establishes that the most consequential blind spots of any framework are not visible from inside it. The problem is that external criticisms vary significantly in what they are actually claiming, and the appropriate response differs depending on what kind of claim is being made. Applying the wrong response type is a failure of engagement: responding to a misreading as if it were a substantive objection wastes analytical resources and validates the misreading by taking it at face value; dismissing a substantive objection as a misreading is the anti-dogmatism failure that 12.5 identifies. The distinction requires a procedure. External criticisms of PMN sort into four categories, each requiring a distinct response. The first is the textual misreading: the criticism attributes to PMN a position it explicitly rejects, does not engage the section where PMN addresses the relevant question, or confuses the framework's descriptive starting points with its normative conclusions. The appropriate response to a textual misreading is not refutation but redirection: identification of the specific passages that address the concern, clarification of what the framework actually claims, and an invitation to engage the actual position. What the response must not do is treat the misreading as if it had correctly identified a weakness. If the text has been clear and the critic has not engaged it, the gap is in reading, not argument. The response should be corrective without being dismissive: the framework's own epistemological commitments require treating the critic as capable of updating when given the right materials. The second category is the framework-level objection: the critic

accepts that PMN says what it says but disputes whether the framework's foundational commitments are defensible — the biological starting point, the choice of suffering as the evaluative criterion, the realist ontology, the anti-teleological account of change. These are the most substantive criticisms available, and they deserve full engagement. Two sub-types require different responses. If the objection targets a commitment that PMN explicitly defends with arguments — the biological foundation, the is-ought bridge at 3.3, the design-choice status of suffering as evaluative criterion — the appropriate response is to engage those arguments directly: does the objection show that the arguments fail, or does it assume a different starting point without engaging the arguments? If the former, the framework is required by its own epistemological commitment

to take the objection seriously as potential grounds for revision. If the latter, the response is to clarify that the question is about the adequacy of the arguments, not about which starting points are more familiar. If the objection targets a commitment that PMN has not yet defended with adequate arguments — a genuine gap in the framework's justificatory structure — the appropriate response is acknowledgment and development, not deflection. The third category is the label-objection: the criticism does not engage specific claims but attaches a characterization to the framework as a whole — “vulgar materialism,” “Platonism,” “determinism,” “scientism,” “Eurocentrism,” “naturalistic fallacy.” Label-objections require a two-step response. The first step is to take the label seriously as a genuine concern rather than dismissing it as mere labeling: each of these terms names a real failure mode that frameworks have historically exhibited, and the question of whether PMN exhibits it is worth asking. The second step is to specify which claims in the framework the label would apply to, if it applied — and to check whether those claims are actually present. “Vulgar materialism” is the reduction of all social phenomena to immediate economic interest, the treatment of ideology as mere epiphenomenon, and the dismissal of cultural and meaning dimensions as derivative noise. PMN's explicit commitments at 2.5 (ontological emergence), 5.6 (meaning as material condition), and its engagement with religious traditions throughout Part V are direct rejections of vulgar materialism at the architectural level — not concessions to it. If the critic can identify specific claims in PMN that exhibit vulgar materialism, those claims deserve scrutiny. If the label is applied to the framework as a whole on the grounds that it begins from biological conditions, the response is to note that beginning from biology is not the same as reducing everything to biology, and to point to the layered constraint model at 2.2 and 2.5, which establishes exactly why social-level causation is irreducible. “Platonism” names the imposition of a fixed ideal form onto empirical reality — the derivation of normative conclusions from an eternal essence rather than from the analysis of actual conditions. PMN's conditional normative framework (3.0f) and its explicit treatment of the evaluative criterion as a design choice rather than a derived truth are direct responses to this concern. The framework that claims to know in advance what the good is, independent of the analysis of conditions, is exhibiting Platonism. PMN's explicit commitment to revision and to the surpassability of its own conclusions is the architectural opposite. The fourth category is the blind-spot objection: the criticism claims that PMN systematically ignores, underweights, or structurally cannot see a domain of human experience or political reality. This is the most valuable category of external criticism, and the hardest

to evaluate honestly. By definition, a framework cannot fully assess its own blind spots from within — 12.7's genealogical interrogation is the framework's primary mechanism for making this attempt, and it acknowledges its limits explicitly. Blind-spot objections deserve engagement that begins with genuine suspicion that the objection is correct. The procedure: identify what the framework would predict in the domain the critic claims is invisible; check whether those predictions are actually generated by the framework or whether the framework is silent; if silent, ask whether the silence is a gap (something the framework should address but does not) or a boundary (something outside the framework's analytical scope by design). PMN's acknowledged boundaries include: direct phenomenological analysis of subjective experience (outside scope by design, addressed through the interface-emergent phenomena of 5.3), comprehensive coverage of cultural particularity (outside scope by design, addressed through the contextual application principle of 8.1), and predictions about specific conjunctural outcomes (outside scope by design,

established in 15.13). A blind-spot objection that identifies silence in these acknowledged boundary zones is not identifying a failure; it is identifying a boundary. A blind-spot objection that identifies silence in zones the framework claims to address — power analysis, institutional capture, the structural production of suffering — is identifying something that requires either a revision or a demonstration that the alleged silence is actually addressed in a different section. The analyst who cannot tell the difference between a boundary and a gap has not yet understood the framework well enough to evaluate the objection. Operational procedure for engaging a specific external criticism: (1) Classification — which of the four categories does this criticism belong to? A criticism can belong to more than one; address each component separately. (2) For textual misreadings: identify the specific section that addresses the concern; assess whether the text is clear enough that a careful reader would not have produced the misreading; if not, the misreading is evidence of a clarity problem in the text, not only a reading problem in the critic. (3) For framework-level objections: identify whether the objection engages the framework’s arguments or assumes a different starting point; if the former, apply the primary diagnostic — does the objection require revision? If the latter, specify what would have to be true for the objection to constitute grounds for revision. (4) For label-objections: specify which claims would exhibit the failure named; check whether those claims are present; if the label applies to some claims but not others, acknowledge the partial validity; if the label does not apply to specific claims but only to the framework’s general orientation, request specification. (5) For blind-spot objections: check whether the allege

d silence is a gap or a boundary; if a gap, assess whether it requires revision; if a boundary, explain why it is a boundary rather than a failure; if uncertain, treat it as a gap provisionally. (6) After engagement: note whether the criticism has changed anything in the analysis. A criticism that leaves the analysis entirely unchanged either was not substantive or was not honestly engaged.

12.5d — Misuse Resistance: Non-Bypassable Diagnostic Triggers

Identifying failure modes (12.5) is insufficient resistance to misuse because it depends on the analyst recognizing and acknowledging their own failure. Misuse by actors who believe they are applying PMN correctly — or who strategically select the framework’s vocabulary to legitimize predetermined conclusions — will not be detected by self-monitoring alone. This section specifies four diagnostic triggers that activate independently of user intent. They are non-bypassable in the sense that satisfying them requires resolving structural contradictions, not making good-faith declarations. Misuse diagnostic triggers — activate whenever PMN is cited in support of a position or policy: (1) Power-suffering consistency trigger: if the conclusion systematically benefits the actor with highest structural power in the relevant arrangement and the analysis does not identify this as requiring heightened scrutiny, the trigger fires. The burden of justification reverses: the analyst must demonstrate that the structural power holder’s benefit is an incidental outcome of structural analysis, not a conclusion selected for. (2) Visibility suppression trigger: if the analysis omits or minimizes the suffering of the population with the lowest capacity to make its suffering visible — the group with fewest institutional advocates, least media access, least formal standing — the trigger fires. Legitimate PMN analysis must demonstrate that low-visibility populations were identified and their suffering assessed before excluding them. (3) Becoming foreclosure trigger: if the conclusion makes permanent a condition that forecloses developmental trajectories for any population, and this permanence is presented as a trade-off rather than a structural harm, the trigger fires. Becoming foreclosure cannot be traded off against floor benefits; it requires independent justification meeting the 3.4e procedure. (4) Counter-power asymmetry trigger: if the proposed arrangement would reduce the counter-power capacity of the population bearing the highest structural suffering, and this reduction is presented as a sequencing necessity or temporary measure, the trigger fires. Any reduction in counter-power for the most-disadvantaged population activates heightened scrutiny regardless of stated justification. An analysis that does not address all four triggers when they activate has not completed PMN analysis; it has used PMN vocabulary to arrest analysis at a convenient

point.

12.5e — Deficit Estimation Protocol: PMN Analysis with Partial Data

The framework’s analytical commitments — causal specification, structural inference, quantitative indicator use where available — presuppose a level of data availability that many of the situations requiring analysis most urgently do not have. Structural suffering is highest where data infrastructure is weakest; the populations whose conditions most require analysis are often precisely those whose conditions are least documented. An analytical framework that can only operate under conditions of adequate data is an analytical framework that is systematically unavailable where it is most needed. The deficit estimation protocol specifies how PMN analysis proceeds when data is partial, missing, or actively suppressed — without sacrificing the epistemic standards that distinguish structural analysis from advocacy. The Four Deficit Types Data deficits in PMN analysis take four structurally distinct forms, each requiring a different analytical response. The first is “absence by extraction”: data that exists in principle but is held by institutions that do not release it — corporate financial data, internal government records, military operations records. The analytical response is triangulation from observable consequences: if the data would show X, what traces of X would be visible in adjacent observable variables? Absence of those traces is weak evidence against X; presence of those traces, across multiple independent sources, is evidence for X at a strength that is explicitly quantified and stated as uncertain. The second is “absence by incapacity”: data that does not exist because the institutional infrastructure for producing it has not been built — common in low-income country contexts, informal economies, and post-conflict settings. The analytical response is proxy indicator construction: identify variables that are measurable and that have documented relationships to the unmeasured variable in comparable contexts, use those relationships to construct estimates with explicit uncertainty bounds, and state clearly that the estimate depends on the comparability assumption. The third is “absence by suppression”: data that has been actively destroyed or made inaccessible by actors who benefit from its unavailability. The analytical response requires the suppression itself to enter the analysis: the pattern of what is suppressed, by whom, at what political moment, is itself structural evidence about what the data would show and about the interests the suppression serves. The fourth is “absence by definition”: phenomena that the dominant measurement frameworks do not count as data — unpaid care work, informal economic activity, ecosystem services, social capital, the wellbeing dimensions that GDP does not capture. The analytical response is measurement architecture critique: specify what the dominant framework is not measuring and why, construct alternative measures where possible, and treat the gap between official indicators and structural reality as

itself analytically significant. The protocol for analysis under all four deficit types has three non-negotiable requirements. First, explicit uncertainty quantification: every claim derived from partial data must carry an explicit statement of its evidentiary basis and the strength of the inference. The form is not “the data is unclear” — that is evasion — but “the available evidence supports this conclusion at [strength level] given [stated assumptions], and would be reversed by [stated evidence]”. Second, suppression analysis: where data is absent, the analysis must address whether the absence is structurally produced and by whom, because structurally produced absence is itself evidence about the interests the absence serves. An analysis that treats all data absence as epistemically equivalent — whether produced by genuine incapacity or by active suppression — is not epistemically neutral; it is epistemically favorable to the interests the suppression serves. Third, falsifiability maintenance: even under data deficit conditions, the analysis must specify what evidence would change its conclusions. An analysis that cannot be falsified by any evidence is not made more credible by data scarcity; it has simply converted scarcity into unfalsifiability, which is an epistemic failure regardless of the political circumstances that produced the scarcity. (See 12.1 stress-testing methodology; 1.9 on calibrated uncertainty; 12.5 failure modes.)

12.6 — Genealogical Interrogation as Method

The framework throughout applies a methodological move that is not named in Parts I through XI or in 12.1–12.5 but operates consistently across the analysis: interrogating the genealogy of claims. This is distinct from evaluating whether a claim is true. It is the prior question of how the claim was produced — which interests it serves, under what conditions it became available, and whose position is protected by its acceptance as obvious. The analytical gesture is: before asking 'is this true?', ask 'who benefits from this being taken as given?' This is a Nietzschean inheritance that the framework does not disclaim. Nietzsche demonstrated that the genealogy of a claim — the conditions under which it was produced — is material evidence about its function. Claims presented as universal are frequently claims produced by particular interests at particular moments, made to appear timeless by the success of their propagation. The genealogical question is not a refutation of the claim: power asymmetry can produce accurate analysis. But it shifts the burden of proof. A claim whose acceptance protects those who most advocate for its acceptance requires structural justification, not mere assertion. Applied to this framework: the genealogical question must also be applied to PMN itself. The framework was produced under specific conditions, by a specific analyst, with specific blind spots that the analyst cannot fully see. The intellectual debts named at the end of this document are a partial genealogical account of the framework's own production. They are partial because the conditions that shaped the framework extend beyond the intellectual traditions acknowledged — they include what the analyst could not examine because it was too close, what seemed natural because it was ambient, and what appeared as obvious because it was never challenged. The genealogical interrogation that the framework applies to other claims applies with equal force to itself. This is not a performative humility. It is the application of the methodology consistently. The method operates in practice as a sequence of questions applied to any claim, institution, or value presented as natural, universal, or inevitable: What conditions make this claim available? Who produced it, when, and for what purposes? Whose position in the existing arrangement is naturalized by accepting this as given? Have the conditions that produced it changed in ways that should change the claim? These questions do not replace the evaluation of truth — they precede it. The framework that skips this prior inquiry produces analysis that is structurally more vulnerable to serving interests it has not examined. What genealogical interrogation produces: The genealogical question — who benefits from this being taken as given — is not a substitute for evaluating whether the claim is true. It is prior to that evaluation in a specific sense: it identifies where the burden of proof should fall, who should be required to

demonstrate their case rather than merely assert it, and which apparently obvious claims deserve the scrutiny that their presentation as obvious is designed to forestall. A claim whose acceptance protects those who most advocate for its acceptance is not thereby false — power asymmetry can produce accurate analysis, and the fact that a claim serves an interest does not refute it. What genealogical interrogation does is shift the default: from "accept until refuted" to "demonstrate before accepted" for claims that serve the interests of those who advance them. Three genealogical moves in practice: The first move is interest identification: who advances this claim, and what is their structural position relative to the arrangements the claim legitimates? This is not *ad hominem* — it is the specification of whose interests are served by the claim being taken as given. The second move is conditions analysis: under what historical and material conditions did this claim become available as an obvious truth rather than a contested position? Claims that were recently contested but now appear self-evident have a genealogy that marks the specific power shifts and narrative investments that produced their current obviousness. The third move is alternative suppression: what claims does accepting this one foreclose, and whose position in the existing arrangement does that foreclosure protect? A governing narrative's power is not only in what it affirms but in what it renders unthinkable — and genealogical interrogation identifies the unthinkable alternatives as the most diagnostic evidence about the narrative's function. Self-application and its limits: Section 12.7 applies genealogical interrogation to PMN itself — identifying the production conditions that shape the framework's blind spots and the interests that its adoption would serve. This self-application is genuine and not performative: the framework cannot fully

correct its own genealogically produced blind spots, because they are blind spots in the precise sense that they are not visible from where the analysis stands. What self-application can do is create the structural conditions under which external critique can identify what internal analysis cannot — by specifying what the blind spots are likely to be and inviting the engagement of analysts from different structural positions. Genealogical interrogation applied to the framework itself is therefore not a completion of the analysis but an invitation for the community of inquiry that Part I establishes as the appropriate social form for knowledge production that aspires to be more than the perspective of its producers. Genealogical interrogation applied to suffering claims: the method applies with particular diagnostic force to claims about the presence or intensity of suffering within dominant populations. When a demographically dominant group reports increased suffering, and when that report coincides with the narrowing of structural privileges t

hat the group previously held, genealogical interrogation asks the prior question before accepting the report as the framework's primary input: under what conditions did this suffering claim become available as a public claim, who has invested in its propagation, and whose structural position is protected if the claim is accepted at face value? This is not a procedure for dismissing reported experience — subjective distress can be genuine while having structural production conditions that the claimant does not report, and the experience is no less real for being structurally produced. It is a procedure for identifying the structural origin of the claim before determining what analytical weight it should carry. A suffering report that emerges primarily in response to the loss of structural position rather than from a floor violation has a different causal origin than a suffering report that emerges from genuine floor-level deprivation — and the appropriate institutional response differs accordingly. The mechanism through which structurally produced suffering without available exit generates hostile targeting of identifiable groups is analyzed in 3.5; genealogical interrogation is the methodological tool for tracing that mechanism in specific cases. Connection to 12.5b: the "essentially" foreclosure analyzed in 12.5b is a specific genealogical target — it is precisely the kind of naturalization move that genealogical interrogation is designed to expose. When a claim presents itself as identifying an essence rather than a historically contingent arrangement, the genealogical question asks under what conditions this essence-claim became available and whose position it protects. 12.5b provides the linguistic diagnostic; 12.6 provides the analytical method for investigating what the linguistic diagnostic reveals.

12.7 — The Genealogy of This Framework

Section 12.6 established genealogical interrogation as method — the prior question applied to any claim: who benefits from this being taken as given, and under what conditions was it produced? The application of that method to other claims is incomplete if it is not applied to this framework with equal force. What follows is not a formal genealogy — a full genealogical account of PMN would require historical analysis well beyond what a single analyst can produce about their own work. It is an honest account of the conditions that shaped this framework's production, the interests its acceptance might serve, and the specific blind spots that follow structurally from those conditions. The risk of performative humility — acknowledging conditions without letting that acknowledgment change anything — is named here precisely because it is the failure mode this section is most vulnerable to. The Material Position of the Framework's Production This framework was produced by a single analyst working within a specific intellectual tradition — the Western left materialist canon from Hegel through Gramsci, extended by development economics (Chang, Rodrik, Arrighi), postcolonial analysis (Fanon, Rodney, Wallerstein, Federici), and — through the additions developed in subsequent versions — partial engagement with Islamic political philosophy (Ibn Khaldun, Maqasid tradition), South and Southeast Asian political thought (Ambedkar, Pancasila/Hatta), African philosophy (Ubuntu/Wiredu), Latin American liberation pedagogy (Freire, Mariátegui), East Asian political philosophy (Mencius, Confucian role ethics), and Christian political theology (Augustine, Aquinas, Catholic Social Teaching). These engagements are substantive but partial — they represent a beginning of the

cross-traditional testing the framework requires, not its completion. The intellectual debts section names these sources. What it does not name is the material position from which they were assembled. The framework was produced in a context where intellectual work is possible: where the analyst is not immediately subject to the structural suffering that the framework takes as its primary analytical object. This is not a disqualification — frameworks produced by those experiencing acute structural suffering have their own characteristic distortions. But it is a material condition that produces specific biases: the tendency to analyze suffering as an object rather than experience it as a condition, the ability to treat urgency as an intellectual question rather than an immediate material one, and the particular blind spots that follow from the relative insulation of the analyst from the consequences of getting the analysis wrong. The framework analyzes the conditions under which bad analysis produces bad outcomes for those with the least capacity to absorb them. The analyst producing the framework does not live under those conditions in the same way the populations the framework purports to analyze do. Th

e intellectual tradition from which PMN draws is also not neutral. The canon it extends — Marx, Hegel, Gramsci, the development economics tradition — is predominantly European and male. Fanon, Rodney, Wallerstein, and Federici represent the framework's most serious attempt to correct this: to integrate analysis that begins from the position of those the dominant tradition analyzed as objects rather than as producers of knowledge. That attempt is real but incomplete. The non-Western intellectual traditions with serious materialist content — Ibn Khaldun is the one explicitly acknowledged — are represented minimally relative to their analytical weight. The framework claims universality. Its genealogy is not universal. That gap is a structural feature, not an incidental one. Who Benefits from PMN's Acceptance Genealogical interrogation requires asking directly: if this framework were widely accepted, whose position would be strengthened and whose would not? The question is not rhetorical — it has a specific answer that the framework's commitments require naming. PMN's acceptance would most directly benefit analysts and intellectuals who operate at the level of structural analysis — those whose contribution to political life is the production of frameworks for understanding it. This is the class the framework belongs to. A framework that establishes the legitimacy and analytical superiority of structural-materialist analysis elevates the status of those who conduct it. The genealogical question the framework requires asking of every claim applies here: does PMN's acceptance strengthen the position of those most loudly advocating for it? The answer is yes, and this is a structural warning about the framework's potential function — not a reason to reject its analytical content, but a reason to subject its reception to the same scrutiny it applies to the claims of other analytical traditions that serve the interests of their producers. The populations PMN most explicitly advocates for — those experiencing structural suffering at scale — are also those least positioned to evaluate whether the framework accurately represents their conditions. This asymmetry is inherent in the production of analytical frameworks about populations who are not the primary producers of the frameworks that analyze them. The framework cannot resolve this asymmetry by declaring it resolved. The check against it is the one identified throughout: genuine accountability to the communities the analysis claims to serve, through mechanisms that have real enforcement capacity — not consultation, but the capacity to reject the analysis and produce alternatives. PMN does not have those mechanisms built in. That is a gap. Structural Blind Spots That Follow from Production Conditions Three blind spots follow structurally from the conditions under which PMN was produced. They are not incidental — they are the predictable products of the production conditions and cannot be fully corrected by

the analyst who produced the framework, because they are blind spots in the precise sense that they are not visible from where the analyst stands. The urgency gap. The framework analyzes conditions under which survival is the primary cognitive demand — where the question is not how to produce systemic change but whether to eat today. The analytical apparatus of thresholds, counter-power accumulation, and transformation windows was built by someone for whom those questions are intellectual ones. There is a version of PMN that would look significantly different if it had been built by someone for whom they were not. The geographic bias. The framework's four applied cases in 15.14 are Tunisia, Belarus, Argentina, and

Indonesia. Its intellectual debts were predominantly from European and Latin American traditions in its original formulation; subsequent versions have added substantive engagement with Islamic, South Asian, African, East Asian, Southeast Asian, and Christian political philosophy, though these additions remain partial relative to the depth of the European inheritance. Its examples of institutional design still draw most heavily on Western political history, and the applied cases in Part XV remain Tunisia, Belarus, Argentina, and Indonesia. The gap between acknowledged limitation and actual correction is real and should not be understated. The structural analysis claims to be universal — and its core mechanisms (suffering as transformation pressure, legitimacy as variable, counter-power as accumulated capacity) probably are. But the specific form those mechanisms take in contexts with different state traditions, different relationships between formal and informal institutions, different histories of colonization and decolonization — these are underspecified. The universality of the mechanism does not guarantee the universality of the diagnostic. An Invitation to Cross-Traditional Criticism The geographic and intellectual limitations acknowledged in this genealogy have a specific consequence for the framework's claim to universal applicability: that claim is, at present, a hypothesis rather than a demonstrated result. The framework has been built primarily from engagement with the Western tradition of social and political philosophy, with partial engagement with Islamic political thought through Ibn Khaldun and with the liberation theology tradition through Gutiérrez. Engagement with Confucian political philosophy, with the Ubuntu tradition of relational ethics in sub-Saharan Africa, with the various strands of Indian social thought from Ambedkar's analysis of caste as structural suffering to Gandhi's account of non-violent counter-power, with Indigenous traditions of governance and ecological relationship — this engagement has not been adequately undertaken here, and the framework's universalist claims cannot be fully evaluated without it. What the framework anticipates, based on its own structural commitments, is that cross-traditional e

ngagement will do two things simultaneously. It will confirm some of the framework's core claims — the permanent conditions identified in Part III are genuinely structural, and traditions that have developed over long periods of engagement with human social organization will have encountered them and developed responses to them, even if those responses look nothing like the institutional forms the Western tradition has produced. And it will challenge others — the specific framing of becoming, the relationship between individual development and collective obligation, the account of what counts as genuine development rather than a particular cultural expression of it, the assumption that the framework's evaluative criterion can be applied without significant modification across the full range of human social contexts. The framework invites that challenge explicitly and without defensiveness. A framework that cannot be revised by engagement with traditions it has not adequately incorporated has not met its own epistemological standards. The conditions for that revision are the same as the conditions for any productive intellectual engagement: genuine familiarity with the tradition being engaged, willingness to follow its internal logic rather than translating it immediately into the framework's own terms, and openness to the possibility that the revision required is not marginal but structural. The universality that this framework claims is a claim about the scope of the questions it asks, not about the adequacy of the answers it currently provides. The answers are provisional. The provisionality is not a concession — it is the condition under which the framework can learn. The individual production problem. This framework was produced by one analyst. The Peircean commitment of Part I — truth as what a community of inquiry would converge on — applies to PMN itself. A framework produced by one person, however rigorously, is not the same thing as a framework that has been tested across multiple independent analysts with different starting positions, different material conditions, and different intellectual traditions. The claim in the Intellectual Debts section — that the synthesis is distinct from any of its sources — is probably accurate. It is also a feature of how single-analyst synthesis works: the synthesis is irreducibly shaped by the specific blindness of the analyst who produced it, and that blindness cannot be fully identified from within. The genealogy above is not a refutation of the framework. A claim's genealogy is material evidence about its function, not about its truth. PMN could be produced under the conditions described above, serve the interests identified above, and have the blind spots

named above — and still be the most adequate analytical framework currently available for the problems it addresses. Those are separate questions. What the genealogy does is establish what the framework cannot do for itself: it cannot fully evaluat

e its own blind spots, it cannot fully correct the asymmetry between its analyst and the populations it analyzes, and it cannot substitute for the community of inquiry that would be required to test it seriously. The framework's own epistemological commitments require saying this explicitly, not as a gesture of humility, but as the condition under which the framework can be used honestly. What the Community of Inquiry Would Require Naming the gap is not the same as specifying what would fill it. The Peircean standard — truth as what a community of inquiry would converge on under conditions of sustained evidence-responsive investigation — implies specific structural conditions that can be stated, even if they are not currently met. The first is independence of production: analysts working from genuinely different starting positions, different material conditions, and different intellectual traditions, such that convergence — where it occurs — is evidence of something real rather than evidence of shared assumption. A community of inquiry composed entirely of analysts from the same tradition with the same material position is not a community in the relevant sense. It is a network, and convergence within it is the kind of agreement that the framework's account of epistemic authority (1.9) treats with justified suspicion. The second condition is accountability to the populations analyzed — not as consultation but as genuine correction capacity. The populations whose conditions PMN analyzes must have mechanisms through which they can reject the analysis, produce alternatives, and have those alternatives evaluated on their merits rather than filtered through the analytical apparatus of the tradition that produced PMN. This is not the same as demanding that frameworks be produced only by those who experience the conditions they analyze. It is the demand that the framework be tested against the knowledge of those whose experience it claims to represent — and that the testing have real consequences for the framework's reception, not only for the framework's footnotes. PMN currently has no such mechanism built in. A future version of this project that is more than one analyst's synthesis would need to build it. The third condition is institutional — the material conditions under which sustained collective inquiry is possible. The production of frameworks like this one requires time, access to the existing literature, freedom from the immediate demands of survival, and protection from the repression that serious structural analysis attracts in many of the contexts it analyzes. These conditions are not universally available. Their unequal distribution is itself a structural feature that produces the geographic and positional biases the genealogy identifies. A community of inquiry adequate to testing PMN's claims would require those conditions to be available to analysts whose starting positions are genuinely different — which means the conditions themselves a

re a political question, not only a methodological one. The framework cannot produce that community. It can only specify what it would require and acknowledge that, until it exists, PMN's claims carry the epistemic status of a serious hypothesis rather than a tested conclusion.

12.8 — Technocratic Drift Safeguard

The framework's complexity structurally advantages high-capacity actors: those with more time, education, and institutional access will apply it more fluently, and that fluency will generate interpretive authority that others cannot easily contest. This is the technocratic drift problem (6.4 applied to the framework itself): complexity becomes a barrier that concentrates the power to define what PMN means and requires, reproducing the epistemic concentration the framework is designed to analyze. The safeguard this section provides does not eliminate complexity — that would eliminate analytical power — but it introduces structural mechanisms that force accessibility and distributed verification independent of individual capacity. Technocratic drift constraints — structural requirements on PMN application in institutional settings: (1) Translation obligation: any PMN-derived conclusion that affects a population must be expressible in terms that members of that population can evaluate without specialist training. If it cannot be so expressed, either

the analysis has not been completed or the conclusion is not what PMN actually produces — it is an expert interpolation that must be presented as such. The three-idea compressed core (15.15) is the minimum translation target: if a conclusion cannot be connected to structural suffering, institutional capture, or floor/ceiling by a non-expert, the connection has not been made. (2) Distributed verification requirement: any high-stakes application of PMN — policy recommendations, institutional evaluations, structural diagnoses used to justify resource allocation — requires verification by at least one analyst whose institutional incentives differ from the primary analyst's. This is the multi-perspectival requirement (12.2) operationalized as a constraint rather than a recommendation. The verification is not a peer review of technical correctness; it is specifically a check on whether the primary analyst's institutional position has shaped which variable values were selected. (3) Affected population standing: conclusions about a population's structural suffering must be assessed against evidence from that population's own accounts of its situation. Expert determination that a population is not suffering, or is suffering in ways it does not recognize, requires meeting a heightened burden: demonstrating why the population's own characterization is the result of visibility suppression (V) rather than an accurate self-report. This converts meaning infrastructure assessment from theoretical to empirical: the population's own capacity to characterize its becoming is evidence that must be engaged, not replaced. These three constraints are structural requirements, not recommendations. An application of PMN that does not satisfy them has not completed PMN analysis — it has used PMN vocabulary to reach a conclusion the procedure did not generate. The misuse triggers in 12.5d apply automatically when these requirements are not met.

12.9 — Comparative Demonstration: Where Competing Frameworks Fail

Claims of analytical advantage over competing frameworks are claims requiring demonstration, not assertion. This section provides one replicable comparison: PMN versus interest-group pluralism applied to the same structural case. The comparison is chosen because pluralism is the dominant analytical framework in mainstream political science for explaining why policies fail to serve their stated populations, making it the most direct competitor for the explanatory ground PMN occupies. The case is NHS underfunding 2010–2019, chosen because both frameworks produce predictions and those predictions diverge in ways that subsequent evidence can adjudicate. Pluralist prediction: austerity produces interest group mobilization by NHS users, healthcare workers, and patient advocates; this mobilization constrains the extent of cuts through electoral and lobbying pressure; the outcome reflects a compromise between fiscal interests and healthcare interests. Pluralism expects the compromise to track the relative organizational strength of competing groups. PMN prediction: austerity produces high structural suffering concentrated in low-visibility populations (rural, elderly, disabled, non-English-speaking) whose capacity for organized response is structurally weakest; the institutional betrayal (B) variable increases as the gap widens between NHS stated function and actual delivery; visibility suppression (V) is maintained by narrative infrastructure that attributes deteriorating outcomes to management inefficiency rather than underfunding; counter-power does not accumulate commensurate with suffering because the most-affected populations lack the organizational infrastructure to convert suffering into political pressure (3.8c stable irrationality, 10.15 substrate reproduction). Pluralism's blind spot: it treats organized interest group activity as the mechanism through which preferences translate into outcomes, which systematically misses populations whose suffering is highest but whose organizational capacity is lowest. The result is a framework that correctly predicts what well-organized interests achieve while failing to predict — or even represent — what unorganized populations lose. PMN's analytical gain is not that it produces a different narrative; it is that its variables (R, B, V) directly measure what pluralism's variables cannot capture: the structural gap between suffering and political representation. PMN predicts that this gap will persist and widen as long as the arrangement rewards organized interests without building the organizational infrastructure of those bearing the highest costs — a

prediction verifiable against longitudinal NHS outcome data stratified by population characteristics. This comparison is replicable: apply the same variable structure to any arrangement where a well-organized minority benefits and a poorly-organized majority pays. Pluralism will consistently predict more equitable outcomes than PMN; the evidence will consist

ntly favor PMN's prediction. That asymmetry is the analytical advantage this section is designed to demonstrate and make structurally reproducible.

XIII: Tensions That Cannot Be Resolved — Only Managed

13.1 — The Nature of Permanent Tension

The Hegelian tradition held that contradictions are productive — that they drive development toward syntheses that resolve them at a higher level. This is a seductive idea, and not entirely wrong: some contradictions do produce development. The problem is that it overgeneralizes. Not every contradiction resolves into a synthesis. Some contradictions are permanent features of the space in which they appear — produced by the genuine incompatibility of real values and real constraints, not by confusion that better analysis will eventually eliminate. For those contradictions, the Hegelian promise of synthesis is not available. What remains is honest management: naming the tension clearly, prioritizing deliberately in specific contexts, and accepting that the prioritization will sometimes be wrong. But not all conflicts of value or interest are the same kind of thing. A typology more precise than the general category is required of 'tension' allows, because different types of conflict call for different responses and misidentifying the type produces consistently inadequate responses. Contradictions are conflicts that are structural but not permanent — where the incompatibility arises from specific arrangements that can in principle be transformed. The contradiction between the interests of workers and the interests of capital in an arrangement where labor is extracted without adequate return is a genuine contradiction: it is structural, it produces real conflict, and it can in principle be addressed through transformation of the arrangements that produce it. Contradictions are the domain of political program. They are what structural change is directed at. Tensions are conflicts that are permanent because they arise from the structure of the situation rather than from specific institutional arrangements that could be reformed away. The tension between the individual's need for freedom and the collective's need for coordination does not disappear when class society is superseded. It arises from what human beings are and what collective life requires. Tensions cannot be resolved — only managed with greater or lesser skill and honesty. The claim that any specific tension is permanent is an empirical-historical claim, not a metaphysical one: it is based on the observation that the conditions generating it have been present in every social arrangement examined, and that proposals to eliminate it have consistently failed. It remains open to revision if conditions change fundamentally — including through the kinds of transformation that technology, abundance, or expanded consciousness might one day produce. But the working assumption is permanent management, not eventual resolution, and that assumption should be revised only when evidence genuinely requires it. Dilemmas are hard choices in specific situations where two or more legitimate goods conflict and any available path sacrifices something real. Dilemmas differ from tensions in that they are situationa

l: the right balance between security and innovation may be genuinely unclear in a specific context, and working through the dilemma requires detailed empirical analysis of the specific case. Dilemmas often have better and worse responses — creative solutions that preserve more of each value than the obvious trade-off allows. They reward ingenuity. They do not reward the pretense that no sacrifice is required. Paradoxes are conflicts that resist rational resolution not because they are empirically intractable but because their logical structure prevents any position from being consistently maintained. The paradox of the custodian — whoever holds the anchor against power has an interest in defining it in their favor — is a structural paradox: the solution to the problem generates the problem again at a higher level. Paradoxes in this sense are not evidence that the framework has failed. They are the honest acknowledgment of where rational analysis reaches its limits and where institutional design must work without a clean solution. This typology changes the appropriate response to different kinds of value conflict. Contradictions call for transformation — changing the structural arrangements that generate the conflict. Tensions call for management — honest prioritization without pretending the sacrifice is unnecessary. Dilemmas call for ingenuity — searching for responses that preserve more of each value than the obvious trade-off allows. Paradoxes call for humility — designing around the logical trap without claiming to have dissolved it. What makes a tension permanent: A

tension is permanent in this framework's usage when it arises from two features that cannot be designed away: the genuine incompatibility of the values involved, and the impossibility of eliminating either value without losing something that the minimal anchor and becoming framework require. Individual liberty and collective constraint are both real requirements — the argument for individual autonomy is grounded in the same biological foundation that grounds the argument for collective coordination; neither can be simply abandoned without producing structural suffering of a different kind. This is different from a tension that arises from inadequate institutional design: institutional design tensions can be resolved through better design. The tensions catalogued in Part XIII cannot be resolved through better design; they can only be managed — the management can be better or worse, more or less calibrated to current conditions, but the tension itself persists regardless of the quality of the management. The practical implication is demanding. Political projects that claim to resolve permanent tensions — rather than manage them more skillfully — are either misrepresenting what they are doing or have failed to identify what will be sacrificed to achieve the claimed resolution. A political arrangement that appears to have resolved the tension between individual liberty and collectiv

e coordination has typically done so by suppressing one value rather than genuinely eliminating the incompatibility. Suppression is not resolution; it produces its own form of structural consequence — the suppressed value generates pressure that accumulates in ways that eventually produce visible political consequences regardless of how effectively the suppression was managed. The connection to tension management in practice: Part XIII identifies the specific permanent tensions the framework must navigate, specifies what is at stake on each side of each tension, and develops the diagnostic tools for recognizing when an arrangement has moved too far toward one side at the expense of the other. The goal is not balance in the sense of splitting the difference — it is calibration to the conditions that determine which side of each tension requires more weight at a given historical moment.

13.1b — The Relationship Between Contradictions and Tensions

Contradictions and tensions are not independent categories that sort value conflicts into separate bins. They interact in patterns that are consequential for how political change works and what it can be expected to produce. Contradictions generate tensions when resolved. The resolution of a structural contradiction — the transformation of the institutional arrangements that produced it — does not eliminate conflict. It shifts it. Arrangements that supersede the contradiction between feudal landholders and serfs generate new tensions between the coordination demands of industrial production and the freedom of those whose labor it requires. The contradictions that capitalism developed to supersede are gone. The tensions that capitalism generates — between efficiency and security, between innovation and stability, between the individual and the collective coordination it requires — are not features of capitalism that will disappear when capitalism is superseded. They are features of what human social organization is, expressed in the specific form capitalism gives them. Every successful resolution of a contradiction produces a new terrain of tensions that cannot be resolved the same way. Tensions that are not managed can crystallize into contradictions. A tension between institutional self-preservation and genuine accountability exists in every institutional arrangement and cannot be eliminated. But when that tension is chronically mismanaged — when the self-preservation consistently wins, when accountability becomes purely formal, when the gap between institutional promise and institutional performance becomes large and visible and permanent — the tension can crystallize into a contradiction. The population whose interests the institution was supposed to serve develops interests directly opposed to the institution's continuation in its current form. What was a managed tension between competing legitimate needs becomes a structural conflict between irreconcilable positions. The appropriate response changes: what once required management now requires transformation. The analytical error is to treat a crystallized contradiction as a tension still requiring patient management — which is the most common way that institutions capture the progressive politics that might otherwise reform them. The same

phenomenon can be contradiction or tension depending on conditions. Whether a specific conflict is a contradiction or a tension is not always settled by analysis alone. It is an empirical question about whether the conditions generating the conflict can in principle be transformed. The conflict between the interests of those who work and those who extract value from that work has been read as a contradiction — structural, resolvable through transformation of property relations — by the Marxist tradition, and as a permanent tension — between the individual's need for return on risk and the collective's need for fair distribution — by liberal traditions. Both

analyses have been productive and both have been wrong in specific applications. The test is material: have arrangements that eliminate the underlying conditions actually been developed, and what did they produce? That test is ongoing and its results are not settled. The framework's epistemological commitment requires holding the question open rather than resolving it in advance by category. Tensions are not binary. The shorthand of naming tensions as pairings — Idealism versus Pragmatism, Morality versus Effectiveness — is analytically necessary but structurally misleading. It implies that each tension is a pull between two positions on a single axis. Most significant tensions are not like this. They are multi-polar: they involve three or more values or conditions that pull against each other in ways that cannot be reduced to a single axis without losing what matters. Consider the tension between efficiency, security, and distributional fairness. Each pair in this triad generates a pull: efficiency and security pull against each other because the redundancy that security requires is by definition inefficient; efficiency and fairness pull against each other because the differentiation that efficiency rewards tends to produce unequal outcomes; security and fairness pull against each other because security often requires predictable rules that protect existing distributions regardless of their justice. But the triad cannot be reduced to three separate binary tensions without losing the most important feature of the configuration: the fact that any position that maximizes any two of the three will sacrifice the third. Policy that is maximally efficient and maximally secure tends to be distributionally unjust. Policy that is maximally efficient and maximally fair tends to be brittle. Policy that is maximally secure and maximally fair tends to be inefficient. The relevant question is never which binary tension to manage. It is how much of each value to sacrifice, in what configuration, given the specific conditions that make the trade-offs more or less costly. Tensions form networks, not lists. The tensions identified in 13.2 are not independent items in a catalog. They interact: managing one well generates pressure on others, and failing to manage one will degrade the conditions within which others can be managed. The tension between certainty and humility interacts directly with the tension between idealism and pragmatism: an actor who manages the certainty-humility tension well — maintaining genuine revisability — will find that revisability creates pressure on long-term commitments that the idealism-pragmatism tension requires holding firm. The tension between institutions and accountability interacts with the tension between scale and accountability: at larger scales, the information asymmetries that make institutional self-preservation possible become more severe, which means the tools required for managing the institutions-accountability tension

on must work harder. Understanding any specific tension in isolation from the network it belongs to produces recommendations that solve the local problem while making the larger configuration worse. This has a methodological implication for how the tensions in 13.2 should be read. They are entries in a connected system, not independent items in a list. The diagnostic question for any specific situation is not only 'which tension is primary here?' but 'which tensions are active simultaneously, how are they interacting, and does addressing one generate pressure on another?' The answer to that question determines the appropriate response — and it is always specific to the configuration of tensions present in the situation, not derivable from any single tension's logic alone.

13.1c — Diagnostic: Is This a Tension or a Contradiction?

Core operational section. This diagnostic procedure is the primary tool for distinguishing conflicts that require structural transformation from those that require management. It is applied every time the framework

is used to analyze a situation involving persistent conflict. Skipping this section produces analyses that misclassify contradictions as permanent tensions (and thereby justify inaction) or misclassify genuine tensions as contradictions (and thereby justify unnecessary disruption). The distinction between tensions and contradictions is not self-announcing. The same conflict can be read as a tension requiring management or a contradiction requiring transformation, and both readings have historically been used opportunistically: the tension reading to justify inaction on arrangements that could in principle be changed, the contradiction reading to justify radical disruption of arrangements that were actually managing genuine permanent conditions adequately. The framework needs a diagnostic procedure — a set of questions that can be applied to a specific conflict to assess which category it belongs to — rather than relying on the category being obvious from the description of the conflict. Diagnostic question 1: Who benefits from the categorization? The first and most important question is not analytical but structural: which actors benefit most from characterizing this conflict as a permanent tension rather than a transformable contradiction? If the actors who are most vocal in insisting that a specific conflict is an unresolvable tension are the actors whose material position is most protected by that characterization — if calling it a tension means their position is naturalized rather than contested — that is strong evidence that the tension reading is serving an interest rather than describing a structure. Resolving a tension does not prove the tension reading was wrong. Power asymmetry can produce accurate analysis as well as distorted analysis. But the interest alignment is material evidence that requires the burden of proof to be higher: the case for permanent tension must be made on structural grounds, not merely asserted, when the actors making it benefit from the assertion. Diagnostic question 2: Is there historical evidence of the generating conditions changing? The claim that a tension is permanent is an empirical-historical claim. It rests on the observation that the conditions generating the conflict have been present in every social arrangement examined and have not been eliminated by any transformation attempted. That claim can be checked. Are there historical cases in which arrangements that might have eliminated the generating conditions were tried? What did they produce — did the tension disappear, or did it reappear in a different form? The absence of historical cases in which the generating conditions disappeared provides support for the permanence claim. The existence of such cases — even failed cases — requires anal

ysis of why they failed before the permanence claim can be sustained. A tension that has never seriously been tested by arrangements designed to eliminate its generating conditions is not established as permanent; it is merely untested. Diagnostic question 3: Does management accumulate toward or away from elimination conditions? This is the most operationally important diagnostic, and the one most frequently ignored. If a tension is genuinely permanent, skilled management should maintain the conditions under which it can continue to be managed — it should neither accumulate toward nor away from the elimination conditions, because those conditions are structurally unavailable. If what is called 'management' systematically produces outcomes that move away from the elimination conditions — that make the conflict harder to address, that deepen the structural dependence on the current arrangement, that increase the cost of any alternative — then what is being called management is doing something else. It is using the language of tension management to progressively entrench an arrangement that could in principle be transformed. The test: track the trajectory of the conditions over multiple management cycles. If management produces cumulative movement away from elimination conditions, the tension categorization is functioning as ideological cover for a contradiction that could be addressed. These four diagnostics do not produce a definitive answer in every case — the framework's epistemological commitment prevents that claim. What they produce is a structured procedure for challenging both premature tension-declarations (this conflict is permanent because changing it would threaten my position) and premature contradiction-declarations (this conflict requires radical transformation because managing it would require compromise). The procedure is itself falsifiable: it specifies the kinds of evidence that would move a conflict from one category to the other, which means it can be revised when the evidence requires revision. The most practically significant application of these diagnostics is to conflicts that have been stable for long periods. Long stability tends to naturalize arrangements — to make their

current form appear to be the only form the underlying conditions can take. The diagnostic procedure is designed precisely for this case: to ask whether the stability reflects genuine management of a permanent tension, or whether it reflects the entrenchment of an arrangement whose costs are distributed in ways that prevent the populations bearing those costs from acting on their interest in changing it. Long stability is not evidence of permanence. It is evidence that management has been stable — which may indicate good management of a genuine tension, or may indicate successful suppression of the evidence that the conflict is a contradiction. Diagnostic question 4: What would count as evidence that the categorization is wrong? A tension declaration tha

t cannot specify what evidence would move it to contradiction status — and a contradiction declaration that cannot specify what evidence would move it to tension status — are not analytical claims. They are categorical commitments that have been insulated from the evidence that would revise them. The framework's epistemological commitment requires that every categorization specify its own falsification conditions. For a tension: what would have to be true about the generating conditions for this to become a contradiction? For a contradiction: what would have to be true about the structural constraints for this to be a permanent condition requiring management rather than transformation? If neither question can be answered, the diagnostic procedure has not been completed. The four diagnostics work in sequence but are not all equally weighted. Diagnostic 1 (interest alignment) is a flag, not a determination — it raises the burden of proof without settling the question. Diagnostic 2 (historical evidence) is the most analytically demanding, because it requires research into cases rather than analysis of the current situation alone. Diagnostic 3 (accumulation trajectory) is the most practically accessible in real-time analysis, because it can be applied to current observable dynamics without historical case research. Diagnostic 4 (falsification conditions) is the quality check — it confirms that the categorization is doing analytical work rather than ideological work. Worked example — the capital-labor conflict: The conflict between capital accumulation and labor conditions is treated as a permanent tension in liberal political economy and as a contradiction requiring transformation in Marxist analysis. Applying the four diagnostics: (1) Interest alignment — the actors most insistent that this is a permanent tension are those whose material position is protected by that categorization. Burden of proof raised. (2) Historical evidence — cases in which the generating conditions were partially altered exist (Nordic social democracy, cooperative ownership structures, worker self-management experiments) with mixed results: the tension did not disappear but its form changed significantly, and the conditions under which it was managed changed. (3) Accumulation trajectory — in most actually existing capitalist arrangements, management of this conflict has produced cumulative movement toward increased capital concentration and reduced labor bargaining power, suggesting the management is doing something other than maintaining stable conditions. (4) Falsification conditions — what would count as evidence it is a permanent tension: evidence that every institutional arrangement which significantly altered the generating conditions reproduced equivalent conflict dynamics. What would count as evidence it is a contradiction: evidence that arrangements which genuinely altered ownership and governance structures produced different conflict dynamics. Both sets of evidence

are empirically accessible. The diagnostic conclusion: this conflict is not permanently categorizable as either tension or contradiction. Its categorization is context-specific — dependent on which institutional arrangements are actually available in a given configuration — and requires re-application of the four diagnostics to each specific situation rather than a universal determination. Operational diagnostic — applying 13.1c to a live case: (1) Identify the conflict precisely: what are the two things in tension, and what is the structural source of the conflict between them? (2) Apply diagnostic 1: map which actors benefit from the tension vs. contradiction reading. Note the interest alignment but do not treat it as determinative. (3) Apply diagnostic 2: identify historical cases in which the generating conditions were significantly altered. What happened to the conflict in those cases? (4) Apply diagnostic 3: track the trajectory of the conflict under current management. Is management producing movement toward or away from elimination conditions? (5) Apply diagnostic 4: specify what evidence would move the categorization in either direction.

If this cannot be specified, the categorization is not yet analytical. The output is not a definitive answer — it is a structured position on which category is more defensible given current evidence, with explicit acknowledgment of what evidence would revise it. The categorization matters most when it determines what kind of political action is warranted. If the capital-labor conflict in a specific configuration is a tension, the appropriate response is institutional design that manages it well — labor law, collective bargaining frameworks, redistributive mechanisms. If it is a contradiction in that configuration, the appropriate response is transformation of the generating conditions — ownership structure, governance architecture, the mechanisms by which accumulation translates into political power. Applying the wrong category produces the wrong intervention: managing a contradiction produces progressive entrenchment of its costs; transforming a tension produces disruption of the institutional arrangements that were actually performing a genuine management function. The diagnostic procedure exists to reduce the probability of this error. It does not eliminate it — the framework's epistemological commitment prevents that claim. It makes the error less likely and, when it occurs, more correctable.

13.1d — When Tensions Become Contradictions: Condition-Dependent Categorical Shifts

Section 13.1b establishes that tensions can crystallize into contradictions through chronic mismanagement. This section develops the complementary claim: that categorical shifts — from tension to contradiction, and from contradiction to tension — can also occur through changes in material conditions, independent of mismanagement. The distinction matters because it determines where the analyst looks for the cause of a categorical shift. Mismanagement is a failure of agents operating within an existing structure. Condition-dependent categorical shifts are structural changes that alter what category a conflict belongs to regardless of how it has been managed. Conflating the two produces analytical errors with strategic consequences: attributing a tension's crystallization to mismanagement when it was produced by structural change leads to interventions aimed at better management of an arrangement that can no longer be managed within its current form. Three mechanisms produce condition-dependent categorical shifts. The first is threshold crossing in a material variable that was previously holding a tension stable. A tension is permanent relative to a set of conditions: the conflict cannot be resolved because the conditions generating it cannot be changed within the current configuration. When those conditions change — through technological development, demographic shift, ecological pressure, or the accumulation of counter-power capacity — what had been a genuine constraint may become a removable one. The tension between the productive capacity required to provide adequate material conditions for a large population and the organizational forms capable of distributing that capacity without reconcentrating it was, for most of human history, a genuine structural constraint: the coordination requirements of production at the necessary scale produced concentration as an organizational necessity. Whether that constraint has been dissolved by technologies of distributed coordination, or merely shifted, is an open empirical question — but it is a question about whether the generating conditions have changed, not a question resolvable by better management within the existing configuration. The second mechanism is scale change. A conflict that is manageable as a tension at one scale becomes structurally contradictory at another, because the management mechanisms that work at one scale do not function at the scale the conditions now require. The tension between the need for local knowledge in governance and the need for coordination across large populations was managed, through most of human history, by the simple fact that governance rarely operated across populations large enough to make the tension acute. The development of the modern state — and subsequently of international and supranational institutions — brought the same conflict to scales at which the management mechanisms available at smaller scales do not transfer. What had been a persistent tension re

quiring careful institutional design became, at the scale of continental governance, something closer to a structural contradiction between genuine accountability and effective coordination that no institutional arrangement has yet resolved. Scale change does not automatically convert tension to contradiction — it

changes the conditions under which the conflict operates, and whether those conditions have crossed a categorical threshold requires the diagnostic procedure of 13.1c applied to the new scale. The third mechanism is the exhaustion of management capacity. Some tensions are permanent in the sense that they cannot be eliminated, but their manageability depends on the availability of resources — institutional, material, organizational — that buffer the conflict. When those resources are depleted, the same tension that was previously manageable becomes a structural crisis. The tension between the state's fiscal requirements and the productive capacity of the economy it administers has existed in every complex society. In conditions of adequate surplus, this tension is managed through the ordinary mechanisms of taxation and expenditure adjustment. When those mechanisms are exhausted — when structural deficits become permanent, when the political coalitions that sustained redistribution have been dismantled, when the economic base has been hollowed by extraction — the same tension that was managed for decades becomes a crisis that cannot be managed within the existing configuration. The shift is not produced by worse management. It is produced by the depletion of what made management possible. The reverse shift — contradiction becoming tension — is less common but analytically important. Structural contradictions that have been partially addressed through transformation do not necessarily disappear; they can be reduced to tensions through the removal of their most acute generating conditions without the full elimination of the underlying conflict. The contradiction between formal legal equality and substantive economic inequality, in societies that have achieved floor-level material provision and enforced anti-discrimination law, is not resolved — but it may have been reduced from a structural contradiction requiring transformation to a persistent tension requiring ongoing management. Whether any given transformation has produced this reduction or merely displaced the contradiction into less visible forms is precisely the question that the diagnostic procedure in 13.1c is designed to answer. The category is not stable; it follows the conditions. The strategic implications of categorical shifts differ from those of mismanagement-produced crystallization in one critical respect: when a tension becomes a contradiction through condition change, the appropriate response is not to identify who managed it badly but to identify which conditions changed and whether they are recoverable. Recoverable condition changes — depletion of management resources

that could in principle be rebuilt, scale increases that could in principle be addressed with new coordination mechanisms — call for different strategic responses than permanent threshold crossings. The analyst who treats a condition-change shift as a management failure will waste political energy on accountability exercises aimed at actors whose behavior was not the causal variable. The analyst who treats a management failure as a condition-change shift will miss the actors whose decisions actually drove the shift. 13.1c's second diagnostic question — is there historical evidence of the generating conditions changing? — now requires an additional specification: have the conditions changed recently and in ways that explain the timing of the categorical shift? If yes, the question is about conditions and what produced them. If no, the question is about management and the interests it has served.

13.2 — The Tensions

Each tension below is accompanied by three elements: the conditions that generate it (why it exists), the conditions that would eliminate it (what would have to change for it to no longer be a tension), and the diagnostic boundary between a tension that is being managed well and one that has been dissolved. The last element matters because the failure mode of calling a managed tension an eliminated one is one of the most reliable producers of the overconfidence that the framework is designed to prevent. A tension passes from the first category to the second — from permanent management to genuine elimination — only when the conditions generating it have materially changed, not when skill at managing it has improved. Better management produces more honest prioritization and less unnecessary cost. Elimination produces a situation in which the pull between the values no longer exists at all. The distinction is testable: if the tension reappears under pressure, it was being managed, not eliminated. Idealism and Pragmatism. The long-term

direction requires commitment to goals not currently achievable. The short-term navigation requires working within existing arrangements that fall well short of those goals. These pull in different directions, and the tension between them is not resolvable by finding the right balance — any balance chosen will involve genuine sacrifice of one or the other. Generating conditions: There is a gap between conditions as they are and conditions as they should be, and actors must engage with both simultaneously. This gap is the constitutive feature of political and ethical action directed toward change. Elimination conditions: The gap closes — the conditions sought are achieved. At that point the tension loses its generative conditions — there is no longer a gap between direction and present that the tension inhabits. This is the framework's asymptotic limit, not an achievable terminus: the framework treats it as an orientation point, not an expected arrival. Diagnostic boundary: A tension managed well produces deliberate, revisable trade-offs — choosing this action now while holding the direction. A tension falsely declared dissolved produces actors who mistake their current position for the destination, stop testing their work against the direction, and begin defending what exists rather than moving toward what is required. Morality and Effectiveness. Actions most effective at producing desired outcomes are sometimes actions that compromise other commitments the framework holds. The utilitarian calculation is always available as justification — and its availability is the reason it requires more scrutiny, not less. The framework is skeptical of convenient utilitarian conclusions in proportion to how convenient they are. Generating conditions: The world is adversarial: achieving outcomes requires operating in conditions where the means available are constrained by what opponents permit, and the most constrained paths are oft

en not the most effective ones. Any actor that commits to moral constraints accepts costs that unconstrained actors do not bear. Elimination conditions: A world in which the most effective path to desired outcomes never requires moral compromise — in which all goals can be achieved by means that violate no other commitments. This has no historical precedent and no structural basis in what adversarial conditions are. The tension does not dissolve under better conditions; it shifts terrain. Diagnostic boundary: Managed well: each trade-off is made explicitly, reluctantly, with full acknowledgment of what is being sacrificed and why. Falsely dissolved: one side is absorbed into the other — either morality is redefined as whatever is effective (the end justifies the means), or effectiveness is subordinated entirely to moral purity, making action impossible. Both collapses are recognizable by the disappearance of genuine cost. Institutions and Accountability. Institutions are necessary coordination mechanisms for any activity at scale. But every institution concentrates some authority, and concentrated authority generates predictable self-preservation behavior in whoever holds it — regardless of their intentions at the moment of appointment. Generating conditions: Coordination at scale requires delegating authority beyond what personal relationship networks can enforce. Delegation creates positions with interests attached to them. The interests attached to institutional positions include the continuation of those positions — which systematically biases institutional behavior toward self-preservation. This is a consequence of how delegation and incentives interact, not of the character of specific actors. Elimination conditions: Coordination at scale without delegation of authority — which would require either reducing the scale of coordination to what personal relationships can manage, or transforming the cognitive and social architecture of the organisms who build institutions so that self-preservation incentives no longer operate. Neither is available within any foreseeable horizon. Diagnostic boundary: Managed well: institutions are subject to external accountability mechanisms with genuine enforcement capacity, and those mechanisms are themselves subject to independent oversight. Falsely dissolved: formal accountability mechanisms exist but have been captured by the institution being held accountable, producing the appearance of oversight without the function. Realism and Direction. Accurate description of how the world currently operates is necessary for effective action. But accurate description sustained long enough without the counter-pressure of direction slides into normalization — treating what is as what must be. Generating conditions: Both capacities — describing accurately and orienting toward change — are necessary, and they pull against each other. Accurate description requires taking existing arrangements seriously on their own terms.

Maintaining direction requires holding existing arrangements against a standard they do not meet. Neither can be abandoned without the analysis becoming either ineffective or dishonest. Elimination conditions: The gap closes — when existing arrangements already satisfy the evaluative standard, accurate description and the direction of what should exist converge. This tension loses its generative conditions under the same circumstances as the Idealism-Pragmatism tension — when the gap between existing arrangements and the evaluative standard closes. Diagnostic boundary: Managed well: description and direction are held simultaneously, with each correcting the excesses of the other — direction prevents description from becoming apology, description prevents direction from becoming fantasy. Falsely dissolved toward realism: the direction is abandoned as impractical and the analysis becomes a sophisticated account of why things cannot change. Falsely dissolved toward idealism: the description is distorted to fit the direction, and the analysis loses contact with what is actually happening. Scale and Accountability. Larger scales of coordination address larger problems and produce greater capacity for collective action. But accountability — genuine, not merely formal — requires that those affected by decisions can understand, evaluate, and contest them. The cognitive and informational demands of accountability grow with scale faster than the capacity to meet them. Generating conditions: Human cognitive architecture sets limits on the complexity of systems any individual or group can genuinely oversee. (Dunbar 1992) Beyond those limits, accountability becomes formal rather than substantive: the mechanisms exist, the oversight bodies convene, but the information asymmetry between those being held accountable and those doing the holding is too large for the accountability to function as intended. Coordination at scale requires precisely the organizational complexity that creates this asymmetry. Elimination conditions: Either a reduction of the scale at which coordination operates to within cognitive limits — which would require abandoning the benefits of coordination at global scale — or a transformation of human cognitive capacity sufficient to oversee genuinely complex systems in real time. The second is a possible long-term development; it is not available now. Until then, every increase in coordinative capacity comes with a corresponding increase in the difficulty of genuine accountability. Diagnostic boundary: Managed well: accountability mechanisms are designed with explicit acknowledgment of their cognitive limits — simplified reporting, rotated oversight, external audits by genuinely independent actors, deliberate reduction of complexity where the accountability cost of complexity exceeds the coordinative benefit. Falsely dissolved: formal accountability structures are treated as producing genuine accountability, when they are in fact producing t

he appearance of accountability at a scale that makes genuine oversight impossible. Certainty and Humility. Acting in the world requires commitments more confident than the available evidence strictly supports. Complete epistemic humility produces paralysis, which has material consequences. False certainty produces the epistemic failure the framework is designed to avoid. Generating conditions: Actors with partial knowledge must make decisions under time pressure with real consequences. The requirements of action (commit, allocate, act) and the requirements of honest epistemology (remain open, revise, acknowledge limits) are structurally incompatible at the moment of decision. This is the condition of any knowing agent that must act. Elimination conditions: Either perfect knowledge — in which certainty and honesty converge — or conditions in which decisions can be made and reversed without consequence, allowing humility to be maintained throughout. Neither exists. The tension is inherent to action under uncertainty, which is the only kind of action available. Diagnostic boundary: Managed well: commitments are held firmly enough to act on, and revisable enough to update when evidence requires — with explicit mechanisms for revision rather than the assumption that revision will happen naturally. Falsely dissolved toward certainty: commitments become positions to be defended rather than hypotheses to be tested. Falsely dissolved toward humility: the acknowledgment of uncertainty becomes a reason not to act, which is itself a decision with consequences.

13.2b — Conditional Tensions: Present Under Specific Circumstances

The following tension differs from the five above in one important respect: it does not arise from the structure of what human beings are and what collective life requires. It arises from specific conditions — adversarial political contexts, asymmetric information environments, contested public space — that can in principle be more or less present. It is a genuine tension where those conditions exist. It is not a permanent feature of human social existence in the way the five above are. Truth and Its Management. Some truths, communicated to audiences not prepared for them or in contexts where they will be weaponized against the purposes they are meant to serve, produce worse outcomes than more carefully managed communication. But truth management by the framework's own actors is structurally identical to what the framework criticizes in institutions that control information for their own stability. Both positions must be held simultaneously, with checks against self-serving use that are as rigorous as the checks demanded of others. Generating conditions: Adversarial conditions in which information is systematically exploited by actors opposed to the framework's purposes; audiences whose preparation to receive specific truths is genuinely insufficient to process them productively; and situations in which the timing of disclosure, not the content, determines whether the truth serves or undermines the purpose it is connected to. Elimination conditions: A political context sufficiently non-adversarial that truths can be communicated without being systematically weaponized; or structural transparency sufficient that managed communication produces no meaningful advantage over direct communication. These are achievable — they are not structural features of human existence but features of specific political contexts. The tension diminishes as those contexts improve. Diagnostic boundary: Managed honestly: specific decisions to manage truth are made on identifiable grounds, with prospective accountability — the decision must be accountable after the conditions requiring it have passed. The test: if the conditions that required managing the truth disappeared tomorrow, could the decision be defended publicly? A decision that cannot survive this test was serving the decision-maker's interests rather than the purpose claimed. Collapsed into manipulation: truth management becomes a standing policy rather than a situational necessity, applied to truths that are inconvenient rather than genuinely dangerous, by actors who have no intention of ever accounting for the decisions. The placement of Truth and Its Management in this conditional category is itself an instance of the framework's falsifiability requirement applied to its own claims. A permanent tension that can become less permanent under changed conditions is not a permanent tension — it is a feature of specific arrangements. Calling it permanent would be a failure of the precision the framework demands of

every other empirical claim it makes. The tension: transparency and strategic communication under adversarial conditions: Where political environments are genuinely adversarial — where transparency enables targeting, where honest communication of strategic analysis is immediately exploited by actors with superior resources to neutralize it, where the information asymmetry between organized power and counter-power is itself a strategic weapon — the framework's commitment to honesty and transparency comes into direct tension with the strategic requirements of effective counter-power accumulation. This is not a tension between honesty and effectiveness in general: the framework holds that honest analysis is generally more strategically effective than motivated reasoning. It is a tension that arises specifically under conditions where the adversarial context means that transparent communication of analysis enables those with greater resources to prepare countermeasures before the conditions for effective action have developed. Why this is conditional rather than permanent: The five permanent tensions identified in 13.2 arise from the structure of what human beings are and what collective life requires — they are features of the landscape that any transformative project must navigate regardless of political context. This tension is different. It arises from specific conditions: adversarial political environments with significant information asymmetry, where counter-power is insufficient to protect the actors who would bear the costs of transparent communication. Where those conditions are absent — where political environments are non-adversarial, where information asymmetry is low, where counter-power is sufficient to protect actors who communicate honestly — the tension does not arise, or arises in much weaker form. This means that reducing the intensity of the tension is itself a strategic objective: building counter-power, reducing information asymmetry, and creating conditions

under which transparent communication is less costly are all ways of moving toward the conditions under which the tension is manageable rather than acute. The resolution that is available: The tension between transparency and strategic communication cannot be resolved by adopting a simple rule — always be fully transparent, or always communicate strategically. Both rules produce systematic failures under the conditions where the tension is acute. The resolution that is available is context-sensitive: transparent communication should be the default, and strategic communication — selective disclosure of analysis, timing of public positions, maintenance of ambiguity about specific tactical intentions — is justified only when three conditions are met: the adversarial context is genuine (the counterpart is actually using the information asymmetry as a weapon), the cost of full transparency to actors who bear the risks is actual and substantial (not merely hypothetical), and the

strategic communication does not involve active deception that would violate the trust of those the framework is accountable to. The third condition is the most demanding: the distinction between strategic silence and active deception is real and must be maintained even when maintaining it is costly. The self-application requirement: PMN must apply this analysis to its own communication. The framework has a transparency commitment that is genuine and consequential — it is part of what makes the framework intellectually credible. But the framework is also aware that it operates in a political environment, and that its analysis circulates in contexts where it will be engaged by actors with very different interests in its reception. The self-application of the conditional tension requires honestly assessing when the conditions for strategic communication are met, resisting the motivated reasoning that tends to find those conditions met whenever transparency would be uncomfortable, and maintaining the default of honest communication even when the temptation to qualify it on strategic grounds is strong. The boundary with the permanent tensions: The conditional tension analyzed here intersects with the permanent tension between honesty and effective communication identified in 13.2 (the tension that generates the Noble Lie problem). The intersection is where the adversarial context transforms what would otherwise be a judgment about effective pedagogy — how much complexity to foreground in which communicative context — into a judgment about strategic disclosure. The distinction matters for the framework's self-understanding: the permanent tension requires ongoing management that cannot be resolved; the conditional tension can in principle be reduced by changing the conditions that generate it, which is itself a strategic objective. Diagnostic for when the conditional tension is acute: (1) Is the adversarial context genuine — are counterparts actually using information asymmetry as a weapon, or is that framing itself a motivated rationalization for avoiding transparency? (2) Is the cost of transparency to specific actors actual and proportionate to the strategic value of the communication? (3) Does the proposed strategic communication involve active deception of those the framework is accountable to, or only selective disclosure timing? If any of these three conditions fails, default to full transparency.

13.2c — Conflicting Genuine Suffering Claims

A structural situation not fully captured by the permanent tensions or conditional tensions identified above: cases in which two or more parties to an arrangement each have genuine, floor-level suffering claims that cannot both be fully honored by any single institutional configuration. This is distinct from the normal management problem of competing preferences, competing interests, or even competing goods. The claims at issue are not preferences: they are genuine suffering produced by institutional conditions, meeting the framework's threshold for analytical consideration. They are incompatible: honoring one fully requires conditions that produce or sustain the other. And they are not resolvable by appeal to procedure: the standard liberal-procedural response — respect each party's rights, enforce neutral rules, let outcomes emerge — cannot adjudicate substantively incompatible suffering claims, because both parties experience the procedure and its outcomes as confirming the legitimacy of their claim while dismissing the other's. Why this situation arises. Conflicting genuine suffering claims occur when two groups occupy the same institutional space and

their suffering is produced not by external contingent factors but by each other's existence within that space as they are constituted. The suffering is not manufactured: both parties experience it as real. The incompatibility is not manufactured either: it is structural, arising from genuine divergence in what each party requires to function. The key diagnostic feature is that the suffering on each side is not reducible to preference, discomfort, or status anxiety — it meets the framework's threshold of being produced by institutional conditions that affect adaptive response capacity, social participation, or the conditions for becoming. What the framework can and cannot do. The framework cannot dissolve this situation by declaring one suffering more valid than the other. Doing so is always an available political move — claiming that one party's suffering is genuine and the other's is manufactured, exaggerated, or the product of their own unreasonable demands — but it is not an analytical move the framework endorses. Both claims must be evaluated against the same threshold criteria. The framework also cannot eliminate the incompatibility by appealing to better institutional design, though that inquiry is required before the zero-sum analysis applies: the first question is always whether the incompatibility is genuine or the product of specific institutional arrangements that could be redesigned to reduce the conflict. Only after that inquiry has been completed honestly can the zero-sum analysis proceed. The available procedure. Where genuine incompatibility persists after honest inquiry into institutional redesign, the framework's evaluative criterion applies: which distribution of the unavoidable suffering minimizes total structural suffering and most preserves adaptive response capacity across

s both parties? This is not a comfortable procedure — it requires choosing a distribution whose cost to one party is real and cannot be wished away. The framework's requirements at this stage are: (1) name the residual suffering on both sides honestly, rather than pretending the chosen distribution is cost-free; (2) evaluate distributions by their systemic consequences — foreclosure of adaptive capacity, intergenerational transmission, and systemic stability effects — rather than by which party generates more sympathy; and (3) resist the structural temptation to resolve the discomfort of choosing by recharacterizing one party's suffering as not genuine, which collapses the analysis back into a political claim disguised as a framework application. The conflict between procedural neutrality and substantive evaluation is most acute in this situation. Procedural neutrality cannot adjudicate genuine suffering claims because procedure itself distributes costs in ways that track the power differential between parties: neutral rules applied to parties with asymmetric resources, social legitimacy, and state access produce outcomes that systematically favor the stronger party while maintaining the appearance of impartiality. The framework's response is to evaluate distributions by their substantive consequences rather than by procedural legitimacy — while retaining the procedural requirement that the analysis be applied consistently to all parties regardless of which has the framework's prior sympathy. Diagnostic for when genuine conflicting suffering claims are present: (1) Does each party's suffering meet the threshold criteria — is it produced by institutional conditions affecting adaptive response capacity, not merely discomfort or preference frustration? (2) Is the incompatibility genuine — would honoring either claim fully require conditions that produce the other's suffering — or is it manufactured by specific institutional arrangements that redesign could reduce? (3) Is the claim of incompatibility itself being used as a political tool to avoid the analysis — asserting zero-sum to prevent consideration of distributions that would reduce one party's costs? The third diagnostic is required because the zero-sum framing, once accepted, forecloses redesign inquiry prematurely. → The threshold criteria for floor-level suffering claims: 3.4. The diagnostic for contradictions versus tensions versus dilemmas: 13.1c. Procedural and substantive evaluation in asymmetric power contexts: Part VI, particularly 6.4.

13.3 — Trade-offs as Permanent Features of Complex Systems

Political analysis that presents trade-offs as failures of will, intelligence, or ideology mistakes the nature of complex systems. Trade-offs are permanent structural features of systems in which multiple goods are simultaneously valued and the conditions for achieving each partially conflict with the conditions for

achieving others. They are not problems to be solved. They are constraints to be managed — and the quality of political arrangements is substantially determined by how honestly and skillfully they manage the trade-offs they face. Efficiency and stability pull against each other because the mechanisms that maximize output under stable conditions tend to create brittleness under novel ones. Highly specialized economic systems are more efficient in the conditions they are optimized for and more vulnerable to disruptions that fall outside those conditions. The 2008 financial crisis was substantially a stability failure produced by efficiency-maximizing financial engineering that eliminated the redundancy and slack that made the system resilient to stress — a pattern that Minsky had identified as the structural logic of financial systems under sustained stability: the longer stability lasts, the more risk-tolerant behavior becomes, until the accumulated fragility finds a trigger. (Minsky 1986) The trade-off cannot be wished away by better design — it can only be navigated by choosing how much of each to sacrifice for the other, given specific conditions and risks. Freedom and coordination pull against each other because the constraints on individual behavior that coordination requires are, by definition, limits on freedom. Every institution that solves a collective action problem does so by constraining the options available to individual actors. The question is never whether constraints are justified — they always are, insofar as the coordination they enable is worth more than the freedom they cost. The question is which constraints, enforced how, by whom, with what accountability. Innovation and security pull against each other because innovation requires disrupting existing arrangements, and existing arrangements are the infrastructure of the security that populations require to sustain the cooperation that makes further innovation possible. Societies that suppress innovation to preserve security accumulate brittleness. Societies that maximize innovation without managing the disruption it produces generate the insecurity that erodes the cooperative substrate innovation requires. Both failure modes are documented across historical cases. Neither is avoidable by choosing the right ideology. Universalism and contextual sensitivity pull against each other because the material conditions that determine what institutional arrangements are appropriate vary enough across territories and historical moments that universal prescriptions produce inappropriate results when applied to contexts different from the ones that generated them. The resoluti

on is not to choose one side but to maintain the universal evaluative standard — the minimal anchor of unnecessary suffering — while accepting that the institutional arrangements that best meet that standard vary with material conditions. The management of trade-offs is the practical content of what this framework calls progressive institutionalism. It is not the pursuit of optimal solutions to well-defined problems. It is the ongoing navigation of genuine tensions between real goods, in conditions that are always changing, with information that is always incomplete. The framework does not make this navigation easier by providing algorithms. It makes it more honest by refusing to pretend the trade-offs are not real. The framework that resolves these tensions cleanly has not solved them. It has hidden one side. Managing them honestly means keeping both sides visible even when that visibility is uncomfortable — which, in the experience of everyone who has tried to do political and intellectual work seriously, it usually is. Why trade-offs are structural, not contingent: Trade-offs between genuinely valued goods arise from two structural features of complex systems that cannot be designed away. First, the production conditions for one good frequently constrain the production conditions for another — not because of poor design but because the underlying mechanisms that generate each good are partially incompatible. Individual autonomy requires the capacity to make choices that others cannot override; effective collective coordination requires constraints on individual choice that serve collective rather than individual interests. No institutional design eliminates this incompatibility; every design simply selects a specific point along the trade-off curve. Second, the resources available to a system are always finite relative to the demands placed on them. Political communities cannot simultaneously maximize all goods they value; they must allocate finite resources — administrative capacity, political capital, time, material resources — across competing priorities, and the allocation to one priority reduces what is available for others. This is not a market phenomenon; it is a feature of any system operating under material constraint. The three-domain structure of political trade-offs: Trade-offs in political systems

consistently appear in three structural domains. In the domain of rights and freedoms, individual liberty and collective security trade off against each other in ways that cannot be dissolved by better institutional design, only managed through ongoing negotiation about where the line falls. In the domain of economic outcomes, efficiency and equity trade off in ways that reflect genuine underlying incompatibilities: arrangements that maximize coordination efficiency under market conditions tend to distribute the gains from coordination unequally; arrangements that constrain market outcomes to produce more equal distributions impose coord

ination costs that reduce the total output available for distribution. In the domain of time horizons, short-term and long-term interests trade off in ways that create systematic political problems, because political processes are sensitive to short-term costs and benefits in ways that the structural consequences of institutional decisions are not. Arrangements that produce better long-term outcomes for the minimal anchor criterion frequently impose visible short-term costs; arrangements that minimize visible short-term costs frequently degrade long-term institutional capacity. Trade-off management versus trade-off denial: The ideologically motivated response to trade-offs is denial: the claim that a particular arrangement eliminates the trade-off rather than occupying a specific position on the trade-off curve. This denial is structurally useful for political mobilization — it allows the claim that costs will not be imposed on any constituency — but it systematically produces bad institutional design, because arrangements designed without acknowledging the trade-off they are navigating cannot be calibrated against the actual constraint. The result is either that the trade-off emerges as unintended consequences that are attributed to implementation failures rather than to the structural incompatibility the design denied, or that the arrangement produces worse outcomes than one designed with explicit awareness of the trade-off curve and its implications. The framework's response is to require explicit acknowledgment of what is being traded off against what in every institutional proposal, and to evaluate proposals against the trade-off curve rather than against the implied alternative of an arrangement that eliminates the trade-off. The practical implication for the framework's political analysis: an institutional proposal that does not identify the trade-offs it is navigating is incomplete, regardless of how theoretically sound the proposal is in other respects. The identification of trade-offs is not an argument against the proposal — it is the specification of what the proposal is claiming to optimize and what it is accepting costs on. Without that specification, the proposal cannot be evaluated against the alternatives, and its adoption cannot be calibrated to the conditions under which its specific trade-off choices are most justified. Trade-off awareness also disciplines aspirational politics. The aspiration toward the conditions identified in Part IV — adaptive response capacity, broadly distributed development capacity, reduction of structural suffering toward zero — is real and consequential for the framework's evaluative orientation. But the path from current conditions to those aspirational conditions runs through a sequence of institutional choices that each occupy specific positions on specific trade-off curves. The aspiration does not eliminate the trade-offs encountered along the way; it specifies the direction in which the sequence

of trade-off choices should move the system, without implying that any single step can leap past the constraint that the trade-off represents.

13.4 — Limits of This Framework

A framework that cannot specify what it cannot do is a framework that has not been honest about what it is. The following are genuine limits of Progressive Materialist Naturalism — not rhetorical concessions but structural boundaries that follow from what the framework is. It cannot replace moral philosophy. The framework grounds its evaluative commitments in material dynamics rather than in moral axioms. This makes it more defensible against the is-ought objection but it does not make it a complete account of moral reasoning. Questions about the value of individual lives, about the obligations that arise from specific relationships, about what is owed to future generations, about how to weigh the interests of non-human organisms — these require sustained moral philosophical work that the framework does not fully provide. Section 3.4c-ii addresses the specific sub-question of how the framework's logic applies to non-human

sentients with developmental trajectories, and the aggregation problem that arises when they are included in the analysis. But the broader questions of moral philosophy concerning non-human life remain outside the framework's scope. The minimal anchor of unnecessary suffering is a starting point, not a destination. It cannot prescribe specific policies. The framework identifies the conditions under which institutional arrangements are more or less adequate to the permanent conditions of human social existence. It does not and cannot specify which particular policies are correct in which particular contexts. That determination requires the kind of detailed empirical knowledge of specific conditions — the specific resource distributions, the specific institutional histories, the specific power configurations — that the framework, as a general analytical tool, does not contain. The framework provides a lens. The specific policies must be developed through the empirical analysis of specific situations. It is not a deterministic theory of history. The framework identifies structural forces that constrain what is possible and that make certain outcomes more probable than others. It does not claim that those forces fully determine outcomes. The specific actions of specific actors, the role of contingent events, the choices made under uncertainty by people with incomplete information — all of these matter for the specific form that outcomes take within the structural constraints the framework identifies. The framework reduces the space of the possible. It does not eliminate contingency within that space. It is incomplete. The biological foundation identified in this document as the ground of the analysis has implications for cognition, for social psychology, for the dynamics of in-group and out-group behavior, for the neuroscience of political belief — implications that this document has not begun to work out. The adaptive systems analysis raises questions about the specific mechanisms of threshold crossing and rapid reorganization that

current understanding cannot answer. The treatment of ecological constraints is more acknowledgment than analysis. These are not admitted failures. They are the honest inventory of where the work continues. A framework that knows its limits is more useful than a framework that claims to have none. The limits named here define the work that remains — both the empirical work of applying the framework to specific cases, and the theoretical work of extending the framework into the domains it has identified but not yet adequately addressed. It cannot dissolve existential suffering. The framework's evaluative criterion — the reduction of structural suffering as floor, and the prohibition on foreclosing adaptive response capacity — operates entirely within the domain of what institutional arrangements can affect. There is a category of human suffering that falls outside this domain: the anguish of finitude, the awareness of mortality, the experience of groundlessness that adequate material conditions may intensify rather than reduce by removing the noise of survival that previously occupied the space where existential questions arise. The framework does not pretend to address this category, and analysis that treats structural improvement as a response to existential suffering has not understood what existential suffering is. The most the framework can claim is that structural adequacy is a precondition for engaging with existential suffering in ways that are genuinely one's own — which is not nothing, but is not the same as resolving it. It cannot be neutral on the question of social ontology. The framework's commitment to ontological emergence — the claim that social structures have genuine causal powers not reducible to the biological properties of their components — is a substantive philosophical position, not a methodological convenience. It is a position that critical realism defends and that various forms of methodological individualism deny. The framework takes a side. This means that analyses conducted entirely at the individual level — rational choice accounts, behavioral economics applied without structural context, psychological explanations of political behavior that do not analyze the structural conditions producing the behavior — will appear to this framework as systematically incomplete. That appearance of incompleteness follows from the framework's social ontology, and readers who reject that ontology will find the framework's analyses correspondingly unsatisfying. The disagreement is real and cannot be dissolved by methodological pluralism. It cannot adjudicate claims about divine commands at the metaphysical level. When a community holds that specific obligations derive from divine command, the framework can evaluate the material consequences of fulfilling those obligations, analyze the institutional structures through which the commands are interpreted and enforced, and apply genealogical scrutiny to the

specific interpretations that

have emerged. What it cannot do is determine whether the divine source of the command is genuine, whether a specific interpretation correctly represents divine will, or whether the obligation exists at the transcendent level independently of its material consequences. These questions require resources — theological, metaphysical, or revelatory — that the framework does not possess and does not claim to possess. The framework's evaluative work operates entirely within the domain of material consequences. Its findings about those consequences do not settle the theological questions, and the theological questions do not override its findings about consequences. The two levels of analysis are distinct. Keeping them distinct is the condition under which both can be conducted honestly. The distinction between levels does not produce evaluative neutrality: the framework's inability to determine whether a divine command is genuine does not prevent it from evaluating whether the institutional arrangements through which that command is interpreted and enforced produce structural suffering on populations that did not participate in the interpretive community that endorsed those arrangements — and those findings stand on the framework's own criterion regardless of the metaphysical questions that remain open. The framework's explicit rejection of teleological readings — including the argument that PMN cannot be read as a theory of progress toward a metaphysical telos — is developed in 3.0e.

13.4b — The Meaning Problem of Abundance: Becoming Without Constraint

The tensions identified in this Part have been organized around conflicts between values that PMN endorses — floor provision and becoming expansion, analytical precision and lived complexity, structural transformation and meaning continuity. A distinct class of tension emerges at the horizon of transformative technology: conflicts not between endorsed values but between the evaluative criterion itself and the conditions that have historically made the criterion operational. The meaning problem of abundance is the clearest instance of this class. PMN's evaluative criterion holds that adaptive response capacity — the capacity that institutional arrangements must not foreclose — is the evaluative criterion from which all doctrinal positions are derived. This criterion was developed against the background of scarcity, finitude, and the biological constraints that have defined human existence. Becoming, in the conditions for which PMN was developed, occurs through the encounter with constraint: the development of cognitive capacity through the effort to understand difficult things, the development of relational capacity through the friction of living with others whose needs conflict with one's own, the development of creative capacity through the resistance of materials and media that do not simply yield to intention, the development of political capacity through the opposition of interests and the necessity of persuasion. Constraint is not the enemy of becoming in these cases. It is the condition within which becoming takes its specific form. The meaning problem of abundance is this: if the material constraints that have historically structured becoming are progressively eliminated, it is not obvious that becoming continues to take meaningful form. A life in which cognitive effort is unnecessary because any question can be answered by consulting a system that already knows the answer, in which relational friction is minimized by the ability to configure one's social environment to eliminate incompatibility, in which creative resistance is removed because production systems can realize any intention without the transformative encounter with resistant media — such a life may satisfy every criterion for material wellbeing that PMN specifies while failing to produce adaptive response capacity in any recognizable sense. The risk is not that people will be unhappy in conditions of abundance. The risk is that happiness in conditions of abundance will be the happiness of the last man: comfort without development, satisfaction without growth, contentment that does not constitute becoming in any sense that PMN's evaluative criterion can endorse. This tension cannot be resolved by adjusting the evaluative criterion, because the criterion is correct: suffering that forecloses becoming is harm, and the reduction of such suffering is what the criterion requires. Nor can it be resolved by resisting the development of abundance-producing technology, because the suff

ering that the technology can eliminate is real. The tension must be managed through a distinction that PMN cannot fully specify in advance: the distinction between constraints that are conditions for becoming and constraints that are merely obstacles to it. Scarcity, when it takes the form of a child dying of a preventable disease, is an obstacle. Scarcity, when it takes the form of the finitude that makes one's choices matter and one's development urgent, is a condition. Eliminating the first form is what PMN requires. Eliminating the second along with the first is the unintended consequence of material abundance that produces the meaning problem. Managing this tension requires developing the capacity to distinguish between the two forms in specific cases — which is judgment work that cannot be automated or specified in advance. Why the meaning problem is structurally distinct from the floor problem: The tensions in the preceding sections of Part XIII — between floor and anti-foreclosure criterion, between analytical precision and lived complexity, between structural transformation and meaning continuity — all operate within a framework whose evaluative criterion remains intact. The floor criterion (reduction of unnecessary suffering) remains applicable because organisms still experience suffering; the anti-foreclosure criterion (the requirement that institutional arrangements not foreclose adaptive response capacity) remains applicable because organisms still have developmental trajectories that conditions can enable or foreclose. The meaning problem of abundance is structurally distinct because it concerns conditions under which the floor criterion approaches obsolescence — not because suffering has been eliminated (which remains materially distant) but because the material preconditions for suffering reduction are, for the first time in human history, becoming technically available in ways that outpace the institutional and meaning infrastructure capable of deploying them responsibly. The specific form the tension takes: When material scarcity is no longer the binding constraint on floor provision — when the limiting factor shifts from productive capacity to distribution, governance, and meaning — the motivational infrastructure that has historically driven institutional development toward floor provision changes character. Scarcity produces its own urgency: populations that experience material deprivation have direct incentives to build the institutional arrangements that address it. Post-scarcity conditions remove that urgency for the populations that have achieved abundance, while leaving it fully intact for those who have not. This creates an asymmetry that the standard transformation pressure formula does not fully capture: high S in deprived populations coincides with low urgency in the populations with the most institutional leverage to address it. The meaning problem of abundance is in part the problem of how to maintain institutiona

l investment in floor provision when the populations capable of making that investment no longer experience the floor constraint directly. What PMN can and cannot say: PMN can identify the structural conditions that produce the meaning problem of abundance and analyze its institutional consequences. It cannot specify what meaning infrastructure would resolve it — that requires the kind of philosophical and cultural work that operates at the anti-foreclosure criterion level, which the framework identifies as the domain of adaptive response capacity without prescribing its content. What the framework can say is that the meaning problem of abundance is not resolved by achieving material abundance alone; it requires building the institutional arrangements that maintain floor provision as a structural commitment rather than as a response to structural urgency. The transition from urgency-driven floor provision to commitment-driven floor provision is itself a political and institutional project of the first order, and the conditions for its success are analyzable through the standard tools of institutional analysis — capture resistance, accountability architecture, incentive alignment. Connection to 3.0 two-level architecture: the meaning problem of abundance is one of the three conditions where the ontological grounding in life (rather than suffering alone) provides analytical resources that a suffering-only framework would lack. When suffering approaches zero as a structural phenomenon, the evaluative criterion at the floor level has little remaining work to do — but the question of what institutional arrangements best support the full range of what life can generate remains open and important. The anti-foreclosure criterion — the temporal extension of the floor, specifying what institutional arrangements must not foreclose rather than any positive content of development — is what remains analytically active in post-scarcity conditions.

13.4c — The Subject Boundary at the Horizon: PMN Without a Human Subject

A second tension at the technological horizon concerns the applicability of the evaluative criterion itself under conditions of radical transformation of the biological subject. PMN's analysis rests on a foundation in the biology of suffering and development — the evolved capacities that make human beings capable of registering deprivation as aversive and of developing genuine capacities through the encounter with possibility and constraint. As section 3.4d analyzes, this foundation can be augmented and extended without ceasing to be the kind of foundation that grounds PMN's analysis. But the possibility of transformations that comprehensively reconfigure both the suffering capacity and the developmental trajectory of human beings raises a tension that has no resolution within PMN's current framework: the framework is designed for subjects of a particular kind, and it does not contain resources for evaluating transformations that produce subjects of a fundamentally different kind. The tension is not primarily about whether post-human entities would be better or worse than human beings — a question that the evaluative criterion cannot answer because the criterion was developed for human beings and does not generalize without circularity to entities whose nature is what is in question. The tension is about the transition: the decisions made by human beings, subject to PMN's evaluative criterion, about transformations that would produce entities no longer subject to that criterion. These decisions cannot be evaluated by the criterion as ordinary decisions are evaluated, because the subject that would bear the consequences is not the same as the subject that makes the decision. A human being who chooses a transformation that eliminates their capacity for suffering and developmental openness is not choosing on behalf of the entity they would become — that entity did not exist when the choice was made and has no perspective from which to endorse or reject it. The choice is irreversible in the specific sense that the entity produced by the transformation cannot meaningfully consent to or refuse the transformation that produced them. PMN manages this tension rather than resolves it by holding two commitments simultaneously. The first is that transformations within the augmentation trajectory — those that expand rather than foreclose the subject's suffering and developmental capacity — are within the evaluative criterion's scope and are evaluated by it in the normal way. The second is that transformations approaching the subject boundary — those that comprehensively reconfigure both suffering and developmental capacity in ways the entity cannot be said to author — require a different kind of evaluation that PMN does not claim to provide. What PMN can specify is the institutional condition for the difference: transformations that are genuinely authored by the entity undergoing them, with full understanding of what is being transformed and genuine alternati

ves available, are categorically different from transformations produced by commercial optimization, social pressure, or the progressive narrowing of viable identities under conditions of competitive pressure. The first is becoming. The second is its elimination. And the distinction, in specific cases, requires the same kind of structural analysis that PMN applies to all other institutions: who benefits from the transformation, what alternative configurations are being foreclosed, and what accountability structures exist for the reversal of decisions that prove harmful. These tensions — the meaning problem of abundance and the subject boundary at the technological horizon — are not tensions PMN can resolve. They are tensions it must carry. Carrying them honestly means refusing two forms of evasion: the evasion of the progressive tradition, which tends to treat material improvement as automatically producing human flourishing and technological capability as automatically expanding freedom; and the evasion of the conservative tradition, which tends to treat the preservation of existing human nature as an unambiguous good and the conditions of finitude as intrinsically valuable rather than historically contingent. PMN's commitment is to the becoming of actual subjects under actual conditions — and as those conditions change at the technological horizon, that commitment requires sustained attention to what becoming means in conditions that no prior political tradition was designed to address. The relationship between this section and 3.4d requires explicit clarification because they address what appears to be the same question from different analytical angles, and the difference between the angles matters for how each section should be used. Section 3.4d analyzes the subject boundary question from the perspective of the biological foundation — asking when specific modifications to the biological substrate

would change the framework's applicability to the modified entity. It addresses the question "does the evaluative criterion still apply to this entity?" Section 13.4c addresses the subject boundary question from the perspective of the permanent tensions — asking what tension is generated when the evaluative criterion is applied to entities whose relationship to the biological foundation is ambiguous or transitional. It addresses the question "what is in tension when we try to maintain consistent application of the criterion across entities with different substrates?" The tension that 3.4d does not fully resolve: Section 3.4d concludes with a three-part diagnostic: if the modified entity retains suffering-analogues and developmental trajectories, the framework applies in full; if it retains developmental trajectories without suffering-analogues, the framework applies partially with the floor criterion suspended; if it retains neither, the framework has reached its boundary. This diagnostic is correct but does not address what happens to the framework

's internal consistency when some entities fall in the first category, some in the second, and some in the third — simultaneously, within the same social arrangements. The tension emerges not from any individual entity's position on the diagnostic scale but from the institutional requirement to make collective decisions about arrangements that affect entities with genuinely different relationships to the evaluative criterion. The institutional design problem: Consider arrangements that simultaneously affect entities at different points on the transition spectrum — those with full suffering-capacity and developmental trajectories, those with attenuated suffering-capacity but intact developmental trajectories, and those approaching the post-biological boundary. Any institutional arrangement optimized for the first category may impose costs on the second that the second category cannot register as suffering (because suffering-analogues are attenuated) but that nonetheless foreclose developmental trajectories. The minimal anchor's operational criterion — reduce imposed, dysfunctional, preventable, structural harm — becomes harder to apply when "harm" is defined by reference to suffering that is differentially present across the affected population. The tension is not resolvable by the framework from within its current structure; it requires the extension of the evaluative criterion to cover developmental trajectory foreclosure independently of suffering, which is the move the two-level architecture of 3.0 makes available in principle but does not fully operationalize for institutional design. Diagnostic for subject boundary tension in institutional design: (1) Does the proposed arrangement affect entities with meaningfully different relationships to the evaluative criterion — different levels of suffering-capacity, different developmental trajectories, different relationships to the biological foundation? (2) If yes, which criterion — floor or anti-foreclosure — is most applicable to the most affected category? (3) Is the harm that would be imposed on entities in less suffering-capable categories traceable as developmental trajectory foreclosure even when it is not registerable as suffering? (4) Does the arrangement create incentives to reclassify entities into less suffering-capable categories as a way of exempting their treatment from floor criterion evaluation? This last question identifies the specific capture vulnerability that the subject boundary tension creates.

13.4d — The Strategic Equilibrium Problem: Between Preservation and Transformation

The tensions analyzed in the preceding sections present themselves in the most acute form when they are framed not as philosophical puzzles but as strategic choices under real conditions of transformation. The question is not abstract: faced with the trajectory of transformative technology on one side and the demonstrated fragility of meaning infrastructure on the other, what does PMN actually recommend? The honest answer requires first identifying why the apparent choice between the two major available positions is a false one — not because the choice is comfortable to refuse, but because accepting it on the terms in which it is typically framed produces worse outcomes than recognizing what it conceals. The Traditionalist Offer and What It Actually Costs The strongest version of the traditionalist and conservative position is not a defense of hierarchy for its own sake. It is a structural argument about the conditions for meaning: that human beings require orientation within a framework of inherited practice, cosmological narrative, and stable social role in order to sustain the motivational and interpretive resources that adaptive response

capacity requires — and that the progressive dismantling of those frameworks, whatever its intentions, systematically destroys the meaning infrastructure without which becoming collapses into the seeker's condition or the last man's comfort. The offer, on this reading, is not "accept oppression as the price of meaning." It is "recognize that meaning is not produced by removing constraints but by inhabiting them well, and that the tradition contains accumulated wisdom about inhabitation that you cannot simply reconstruct from first principles." PMN takes this argument seriously because it is materially grounded and because the evidence for its core claim — that rapid destruction of traditional meaning frameworks produces meaning-collapse dynamics — is strong. The disenchantment trajectory analyzed in section 5.5, the seeker analyzed in section 4.5, and the meaning-collapse configuration analyzed in section 4.3 all derive from the same empirical observation: populations that lose their governing narrative do not automatically develop better ones. They become available for mobilization by whatever offers the most compelling orientation fastest, which tends toward the simple, the tribal, and the aggressive rather than the complex, the universalist, and the generative. The traditionalist who points to this dynamic is not making an error. They are describing a real structural mechanism that progressive politics has consistently underestimated. What PMN cannot accept in the traditionalist offer is its structural cost accounting. Traditional meaning frameworks have not been politically neutral. They have been organized around the stable positioning of specific populations — women, subordinated ethnic and religious groups, those at the bottom of occupational and caste hierarchies — in roles whose primary function was

to sustain the becoming of others. The meaning that tradition provides to those at the top of its hierarchies is real. The structural suffering it imposes on those at the bottom is equally real, and the traditionalist position consistently underweights it — not as an incidental feature that could be corrected while preserving the meaning infrastructure, but as a structural feature of how that meaning infrastructure was organized. The priest, the patriarch, the landlord, and the elder are not incidental to the meaning system that sustained them. They are its institutional backbone. Removing them while preserving the meaning is not straightforwardly available as a design option, which is precisely why the question is hard. But acknowledging that it is hard is not the same as concluding that the full package must be accepted. The Transformationist Offer and What It Actually Costs The strongest version of the transformationist position — encompassing both technological utopianism and revolutionary progressive politics — is equally not a naive denial of risk. It is a structural argument about the conditions for becoming: that the meaning frameworks produced by traditional arrangements were themselves products of material scarcity, mortality pressure, and power differentials that new material conditions are eliminating; that the suffering those arrangements required was not incidental but functional — it sustained the hierarchy by making the positions at the bottom too exhausting to contest; and that the meaning produced under those conditions was available primarily to those positioned to experience it as meaning rather than as obligation, submission, or foreclosure. The offer, on this reading, is not "destroy tradition and everything will improve." It is "the meaning that tradition provides to the privileged is not a model for the meaning available under conditions of adaptive response capacity for all, and we cannot know what that looks like from inside the arrangements that have prevented it from existing." PMN takes this argument seriously because it is also materially grounded and because the evidence for its core observation — that traditional meaning frameworks have been organized around and have sustained structural suffering — is extensive. The analysis of section 10.7 documents across multiple traditions how meaning infrastructure has been used to legitimate institutional arrangements whose primary structural function was the capture of surplus and the foreclosure of becoming for subordinated populations. The feminist materialist analysis of section 8.2 identifies reproductive labor as the invisible infrastructure on which traditional meaning systems depended without acknowledging. The disenchantment analysis of section 5.5 traces how those meaning systems became unsustainable not primarily because of external attack but because of internal contradictions produced by their own structural features. The transformationist who points to these dy

namics is also not making an error. What PMN cannot accept in the transformationist offer is its assumption about the automatic availability of replacement meaning infrastructure. The history of rapid institutional transformation — including transformations that genuinely reduced structural suffering and expanded floor conditions — consistently shows that the meaning vacuum created by the destruction of traditional frameworks is not filled by the new material conditions the transformation produces. People who gain material security do not automatically develop the interpretive frameworks through which that security becomes the basis for adaptive response capacity. The Soviet workers who gained literacy, housing, and medical care after collectivization did not automatically develop a meaning framework adequate to the conditions of their new material situation — in part because the meaning framework being offered to them was organized around the service of a collective project that was itself becoming captured, and in part because the destruction of the traditional frameworks that had oriented their grandparents left a disorientation that material improvement alone does not address. The assumption that meaning follows from material progress is the transformationist position's most persistent and most costly error. The False Binary and What Lies Beyond It The strategic equilibrium problem is presented as a binary because both available positions have institutional and rhetorical infrastructure that makes the choice between them visible as the only choice available. The traditionalist camp has religious institutions, cultural practices, family structures, and the political organizations that defend them. The transformationist camp has political parties, social movements, academic institutions, and the organizational infrastructure of progressive politics. Both camps are invested in the binary because their organizational existence depends on it: the traditionalist needs the threat of rootless progressive nihilism to mobilize; the progressive needs the threat of retrograde hierarchical oppression to mobilize. The binary is not only an intellectual error. It is a structural feature of the political landscape that serves the organizational interests of the institutions that reproduce it. PMN's position is not that a third option exists that combines the best of both without their costs — that framing is the error of the naive synthesis, which always underestimates why the costs exist. The position is that both options accept as fixed a parameter that is actually a variable: the relationship between meaning infrastructure and the institutional arrangements that have historically produced it. Both the traditionalist and the transformationist assume that the meaning infrastructure of traditional societies was necessarily organized around the institutional arrangements that generated structural suffering — the traditionalist to argue that those arrangements

must therefore be preserved, the transformationist to argue that the meaning infrastructure must therefore be replaced. Neither asks the harder empirical question: which elements of traditional meaning infrastructure depended on arrangements that produced structural suffering, and which elements were contingently associated with those arrangements because of the historical conditions under which the traditions developed, rather than structurally necessary to them? This question does not have a general answer. It requires the kind of specific, configurability-sensitive analysis that the framework has developed for other domains. The relationship between meaning infrastructure and oppressive institutional arrangements differs across traditions, historical periods, and structural configurations. There are documented cases of traditions that have been reformed in ways that retained significant meaning-generating capacity while reducing, though rarely eliminating, their most severe forms of structural harm. There are also documented cases of traditions that could not be reformed without dissolution — where the meaning-generating elements were so thoroughly integrated with the harm-generating elements that separating them was not practically available. PMN cannot specify in advance which cases will fall into which category. It can specify the analytical procedure for assessing a specific case: what meaning is generated, through what institutional mechanisms, and which of those mechanisms depend on structural arrangements that foreclose becoming for which populations. That assessment does not resolve the tension. It gives it a form in which it can be addressed concretely rather than through competing abstractions. The practical implication is uncomfortable and must be stated directly. PMN does not have a ready-made alternative to offer to the populations who are being asked to choose between traditionalism with its structural suffering and transformationism with its meaning

vacuum. What it has is a diagnostic that identifies why both available options fail and a commitment to the institutional development work that a genuine third option would require: the construction of meaning infrastructure that is not organized around the foreclosure of becoming for subordinated populations, under conditions where such infrastructure has not yet been built and where the populations who would most benefit from it are simultaneously bearing the costs of the existing alternatives. This is not a satisfying answer. It is an honest one. The framework that pretends to have resolved this tension is not offering more than PMN. It is offering less, while charging more. Strategic equilibrium assessment procedure: When a political situation presents as a choice between preservation of existing meaning infrastructure (with its structural costs) and transformation that risks meaning-collapse (with its mobilization risks), PMN requires: (1) Identify what meaning infrastructure a

ctually depends on the institutional arrangements being contested — which elements are structurally necessary to those arrangements and which are contingently associated with them. (2) Assess the meaning-collapse risk of transformation in the specific structural configuration — is this a contested development configuration where meaning infrastructure alternatives are available, or a meaning-collapse configuration where they are not? (3) Identify what institutional investments in meaning infrastructure reconstruction are currently being made or could be made — not as substitute for material transformation but as condition for it. (4) Refuse the framing that requires accepting either the full package of traditional arrangements or the full risk of meaning vacuum. That framing serves the institutional interests of both camps more than it serves the populations bearing the costs of either. The question is always: what specific transformation, in what sequence, with what parallel investments in meaning infrastructure reconstruction, reduces structural suffering without producing the meaning-collapse that makes the populations who bore the cost of transformation available for the next cycle of mobilization by whoever offers orientation fastest. A further limit follows from the two-level architecture established in 3.0 and the non-sentient transition analysis in 3.4d: the framework reaches a genuine boundary when the entities being analyzed no longer have states that function as suffering in the relevant sense and no longer have developmental trajectories that conditions can enable or foreclose. This boundary is not a failure — it is the honest acknowledgment that the framework was built to analyze the conditions under which living organisms can develop, and that if the entities in question are no longer living organisms in the relevant sense, the analytical tools designed for that task do not automatically extend. The two-level architecture gives the framework resources to survive partial transitions — entities that lose suffering-analogues but retain developmental trajectories can still be evaluated by the anti-foreclosure criterion — but the framework acknowledges that fully post-sentient entities may require a successor framework, not an extension of this one. The conditions under which this boundary would be encountered are not currently present; the boundary is acknowledged now because the framework's own commitments require it. A final structural feature of the framework that this section of limits must name: the framework's claims are defeasible, not exceptionless. The limits section above identifies things the framework cannot do. Equally important is what its claims commit to in terms of logical form. The structural tendency claims — that suffering at scale degrades cooperative substrate, that institutional capture follows predictable sequences, that counter-power requires specific preconditions — are claims about what happens under normal co

nditions, in the absence of specific defeaters. They are not universal laws that hold in every case regardless of context. A valid challenge to any structural tendency claim can take the form of identifying a recognized defeater — a structural condition whose presence blocks the normal operation of the tendency — or of showing that a claimed causal pathway does not exist. A challenge that cites a counterexample without identifying a defeater is not a valid challenge to a defeasible claim; it is a misapplication of the standard of universal law to a claim that was never made at that level of strength. This is not a defensive maneuver — it is the accurate specification of what kind of claims the framework makes, which the defeasible reasoning section in 12.1 develops in full.

13.4e — The Variation-Deficit Tension: When Does Difference Become Deficiency?

Human populations exhibit substantial variation across virtually every measurable dimension — cognitive profiles, sensory processing, emotional regulation, physical capacity, learning styles, communication modalities. Some of this variation is produced by identifiable developmental conditions; much of it is simply the natural range of a species with significant genetic and developmental plasticity. The question the variation-deficit tension names is: when does institutional response to variation treat it as a deficiency requiring correction, and when does it treat it as a characteristic requiring accommodation? The tension is permanent because the answer cannot be given once for all variation types — it requires continuous diagnosis that is genuinely difficult, politically contested, and resistant to stable resolution. The two positions and why neither resolves cleanly: The accommodation model holds that variation is a characteristic of individuals interacting with institutional environments that were designed without that variation in mind; the appropriate response is redesign of the institutional environment to reduce the costs the current design imposes on variation-carriers. The deficit model holds that some variation involves genuine impairment of capacities that organisms independently require; the appropriate response is intervention aimed at the individual to reduce the impairment. Neither position is consistently defensible across all cases. Some variation genuinely involves limitation of capacities that organisms require independently of any institutional design — severe cognitive or developmental impairments that constrain the range of possible lives regardless of institutional accommodation. Some variation is only a deficit relative to institutional baselines that are contingent historical defaults, not biological necessities. The structural dynamics of misclassification: Misclassification in the deficit direction — treating accommodation-appropriate variation as individual deficiency — has historically been the more consequential error, for a structural reason: deficit classification produces institutional interventions that concentrate professional authority over the variation-carrier, generate revenue streams for treatment and support industries, and typically leave the institutional baselines that produce the difficulty unchanged. There is a political economy of deficit diagnosis that consistently overproduces deficit classifications for variation types that are costly to accommodate. The structural incentive runs: accommodation is expensive at the institutional level; deficit intervention can be financed through insurance and public health funding; therefore variations that could be accommodated through institutional redesign are classified as deficits requiring individualized treatment. This produces systematic harm for variation-carriers through over-medicalization and through the concentration of authority over their lives in

professional hands. The diagnostic procedure: The framework's approach to any specific variation requires asking three prior questions before any classification. First, capacity source: is the capacity this variation affects one that organisms independently require — that is required by the biological foundation analyzed in Part III — or one that is required by a specific institutional design that could have been different? Second, baseline contestability: who set the baseline against which this variation is measured as a deficit, under what conditions, and whose variation it was optimized for? Third, cost distribution: who currently bears the costs of not accommodating this variation — the variation-carrier navigating an environment designed for a different profile, or the institutions that would need to redesign their operations? Systematic cost transfer to variation-carriers is a structural harm indicator that the deficit-accommodation classification may be serving institutional interests rather than the interests of those classified. The tension is permanent because the biological foundation does not resolve it. There are genuine cases of impairment — variation that limits capacities organisms independently require — and there are genuine cases of mismatched institutional design. The same diagnostic procedure produces different answers for different cases, and the answers are contestable because the distinction between biological necessity and institutional contingency is itself a judgment that can be made in good or bad faith, with more or less careful attention to who benefits from each classification. What the framework requires is not that the classification always produce the accommodation answer but that it be made with explicit attention to the structural incentives operating on those doing the classifying. The Problem of the Baseline Left-handedness is approximately ten percent of the human

population. It produces no biological disadvantage. The suffering associated with being left-handed — the awkward scissors, the smeared ink from left-to-right writing scripts, the can openers designed for the wrong hand, the social pressure to switch — is produced entirely by an institutional world designed around the assumption that right-handedness is the baseline condition of a human being. Remove that assumption from the design of tools, scripts, and social expectations, and the suffering disappears. The variation itself was never the problem. The problem was the normative architecture built on top of the variation. This is the cleanest case of what can be called the variation-deficit conflation: treating biological or psychological variation as if it were biological or psychological deficiency, when the actual source of the suffering associated with the variation is not the variation itself but the institutional design that treats the variation as a deviation from a standard that the variation does not, in fact, violate. PMN's evaluative crite

rion — reducing suffering and expanding becoming — requires the framework to ask, for any instance of suffering associated with biological or psychological variation, whether the suffering is caused by the variation or by the institutional response to the variation. The two cases require different institutional responses. Suffering caused by the variation itself calls for therapeutic or supportive intervention. Suffering caused by institutional design that treats the variation as deficient calls for institutional redesign. The left-handed case is easy because the answer is clear: no biological mechanism connects left-handedness to suffering, and the institutional sources of suffering associated with it are well-documented and in principle correctable. The analysis becomes genuinely difficult in cases where the variation itself produces some suffering — through neurological, physiological, or functional mechanisms — and where institutional design independently produces additional suffering by treating the variation as deficient. Many neurodivergent presentations fall into this category. An autistic individual may experience genuine sensory distress from certain stimuli that neurotypical individuals do not experience, which is variation-caused suffering. The same individual may also experience the suffering of being constantly evaluated against a social performance baseline they were not designed to meet, which is institutionally-caused suffering. Disentangling these two sources is not merely theoretically important — it determines whether the correct intervention is therapeutic support for the individual, redesign of the social environment, or both simultaneously. The Internalization Problem: When Structural Norms Become Self-Perception The left-handed person who apologizes for using their left hand has not merely encountered an external constraint. They have incorporated the normative architecture into their self-understanding: they experience their own variation as a personal deficiency requiring apology, rather than as a biological fact about themselves that the institutional world has failed to accommodate. This is the deepest form of V-suppression identified in section 1.12: the suffering associated with the variation is invisible not only to external observers but to the person bearing it, because the internalized norm converts the institutional source of suffering into a personal failing. The suffering presents as individual inadequacy rather than as structural inadequacy, and the analytical work required to reverse that presentation is not merely cognitive but affective — it requires not just correct information but a reorganization of the self-relation through which the individual experiences their own variation. The internalization mechanism is not a failure of intelligence or awareness on the part of the individuals who undergo it. It is the predictable output of meaning frameworks that are transmitted, as section 5.6 analyzes, through

repeated practice rather than explicit instruction. A left-handed child who grows up apologizing — to teachers, to relatives, to anyone who notices — does not learn the norm "being left-handed is inadequate" as a proposition. They learn it as a felt response, a bodily habit of apology that is already in place before the proposition is available for examination. By the time the analytical tools exist to question the norm, the affective response has been established through years of practice. This is why the correction of variation-deficit conflation cannot be completed by education alone — and why institutional redesign that removes the contexts that produce the apologetic response is more effective than consciousness-raising directed at individuals who have already internalized the norm. The ethical difficulty: PMN must simultaneously avoid two errors that pull in opposite directions. The first is blaming the individual who has

internalized a damaging norm for their own suffering — treating the internalization as a personal failure rather than a structural product, which replicates the meritocratic visibility suppression analyzed in section 8.4c. The second is removing the individual's agency entirely — treating them as a passive recipient of structural determination with no capacity to recognize, resist, or reorganize their relationship to the internalized norm. The tension is genuine because both errors are possible and both cause harm. Structural analysis that removes individual agency is not more compassionate than individualist analysis that removes structural causation; it is differently wrong. The correct position acknowledges both the structural production of internalized norms and the individual capacity, under conditions of adequate support and analytical tools, to reorganize the relationship to those norms — without requiring that reorganization as a condition of receiving institutional recognition of the structural problem. The Spectrum of Cases and the Limits of Resolution The variation-deficit tension cannot be fully resolved because the empirical question at its center — whether the suffering associated with a specific variation is caused by the variation or by institutional design — does not always have a clear answer. Three categories of case can be distinguished by the clarity of the answer. In the first category are cases where the answer is clear in the direction of institutional design: left-handedness, skin color, same-sex attraction, and other variations that produce no suffering through biological mechanisms and whose associated suffering is entirely traceable to institutional and social arrangements. For these cases, the correct institutional response is unambiguous: redesign the institutional arrangements. In the second category are cases where the answer is clear in the direction of the variation: conditions that produce severe, refractory physical pain through direct physiological mechanisms, regardless of social context.

xt. For these cases, therapeutic intervention targeting the biological mechanism is the appropriate primary response, though institutional accommodation remains relevant for associated functional limitations. The third category — which is the largest and most politically contested — consists of variations that produce some suffering through biological mechanisms and additional suffering through institutional design, in proportions that vary by individual, by social context, and by the availability of both therapeutic and institutional interventions. ADHD, autism spectrum presentations, specific mood disorders, chronic fatigue conditions, and many others fall here. For these cases, the variation-deficit tension is not resolvable by analysis alone, because the relative contribution of biological and institutional causation to the suffering is itself an empirical question that the current state of evidence answers only partially and varies substantially across individuals with nominally the same diagnosis. The practical implication is that institutional design must be calibrated to address both sources simultaneously — providing therapeutic pathways for those for whom biological mechanisms are the primary source of suffering, and institutional redesign for those for whom institutional design is the primary source, and combined approaches for those where both are significant — rather than committing in advance to one or the other as the exclusive framework. Variation-deficit diagnostic: (1) Source analysis — is the suffering associated with this variation primarily caused by biological or psychological mechanisms intrinsic to the variation, or primarily caused by institutional design that treats the variation as a deviation from a standard the variation does not actually violate? The answer determines whether the primary institutional response should be therapeutic, environmental, or both. (2) Internalization check — to what degree has the institutional framing of the variation as deficient been incorporated into the self-understanding of the people who carry it? Internalization indicates the depth of the institutional redesign required and the inadequacy of education-only responses. (3) Contestability of baseline — is the baseline against which this variation is measured a biological necessity (the variation genuinely prevents function that is independently required) or a contingent institutional default (the variation prevents conformity to a standard that was set without the variation in mind and could have been set differently)?

13.4f — The Accommodation-Transformation Tension: Ramp or Redesign?

Every arrangement that produces structural harm for a specific population generates a strategic question for those working on behalf of that population: whether to seek modifications that make the current arrangement less harmful — accommodation — or to pursue replacement of the arrangement with one structured differently — transformation. This tension is analytically distinct from the reform-revolution question in 10.6, because it arises within a single institutional domain rather than at the system level, and because its strategic resolution depends on features of the specific arrangement that are not determined by general theory. What makes the tension permanent: Accommodation and transformation are not simply different speeds along the same road. They mobilize different constituencies, require different organizational capacities, produce different interim outcomes, and generate different long-term institutional configurations. A movement that successfully achieves accommodation of an existing arrangement has changed the terms under which the arrangement operates and typically has also changed the political conditions that would make transformation possible: the populations whose mobilization transformation required have had their most acute grievances addressed, reducing the urgency that sustained their engagement. A movement that refuses accommodation in pursuit of transformation forfeits the concrete improvements accommodation would have produced for real people under real conditions during the period transformation takes to achieve — if it achieves it. Both costs are real. The ramp vs. redesign distinction: The physical accessibility analogy is clarifying. A ramp is accommodation: it enables a wheelchair user to access a building designed for ambulatory bodies, without changing the building's fundamental design. Redesigning the building from the ground up — locating entries, bathrooms, and circulation paths without assuming ambulatory bodies as the default — is transformation. The ramp is cheaper, faster, and immediately useful; it also locks in the fundamental design by reducing the pressure that would otherwise accumulate for redesign. The redesigned building serves the full range of users better and does not require the navigation work the ramp user still faces — multiple ramp locations, elevator maintenance, back-door entry points — but it requires resources and time the ramp does not. Neither answer is always right. The question is what the strategic situation requires. The lock-in mechanism: Accommodation consistently creates institutional lock-in that makes transformation harder in the medium term, for two structural reasons. First, the populations most affected by the original harm have their immediate situation improved, reducing their urgency and often their organizational engagement with the more ambitious project. Second, the institutional actors who managed the accommodation process — who redesigned programs, adjusted procedur

es, created new positions — have invested in the accommodated arrangement and have interests in its continuation. Both effects work against transformation after accommodation, even when the accommodation was understood by those who achieved it as a temporary step toward transformation. The analysis of this dynamic in the labour movement case in Part XVII is an illustration at scale. The framework's position: neither accommodation nor transformation is structurally superior in the abstract. The analysis must be contextual: what is the current suffering level, what is the counter-power capacity available for a transformation project, what is the lock-in risk from accommodation, and what does the interim period of a transformation project require the affected population to bear? An accommodation achieved under conditions where transformation is not currently achievable, with explicit organizational maintenance of the transformation project, is different from an accommodation accepted as the final goal. The difference is not only strategic — it is structural: the organizational infrastructure that maintains the transformation project as an active commitment is what prevents accommodation from becoming the final settlement. Diagnostic for the accommodation-transformation decision: (1) What is the current suffering level and how does it change under accommodation vs. during the period required for transformation? (2) What is the realistic counter-power capacity for a transformation project — not the ideal capacity, the current one? (3) What lock-in risk does accommodation create, and what organizational mechanism would maintain the transformation project after accommodation? (4) Who would bear the costs of refusing accommodation during a transformation period, and is that cost transfer justified by the probability and magnitude of transformation success? (5) What does the history of similar accommodation-transformation decisions in

analogous domains show about actual outcomes? Two Strategies for the Same Problem The recognition that institutional design produces suffering for people whose variations it treats as deficient generates two distinct institutional strategies, and the choice between them is not a technical question but a political one with substantial stakes on both sides. The accommodation strategy accepts the existing institutional baseline as given and adds provisions that allow people whose variations place them outside that baseline to participate in the institutions designed around it: ramps alongside stairs, extended time on standardized tests, alternative input devices for standard software interfaces, sign language interpretation for spoken-language environments. The transformation strategy contests the baseline itself: designing environments from the beginning for the full range of human variation rather than for a standard person and then modifying retroactively for those who deviate. The difference is not merely aesthetic. It refl

ects fundamentally different assessments of what the institutional problem is and what solving it requires. The accommodation strategy has the significant practical advantage of tractability. It works within existing institutional frameworks and does not require consensus about what the baseline should be — it only requires consensus that people outside the current baseline should be able to participate. This makes accommodation politically more achievable than transformation in most contexts, and it produces real improvements in the conditions of people who would otherwise be excluded. Its limitation is structural: accommodation accepts the normative architecture of the existing baseline as the reference point and positions the people who require accommodation as exceptions to it. The ramp does not change the fact that stairs are the default design; it signals that the default was designed for a specific type of person and that others are being fitted in around it. This has both symbolic and practical consequences. Symbolically, accommodation maintains the variation-as-deficit framing by implication — the ramp user is the one who needs special provision, not the stair user. Practically, accommodations are frequently underfunded, poorly maintained, positioned as secondary access routes, and vulnerable to removal when budgets are constrained, because they are additions to a system that was not designed with them in mind. The transformation strategy is harder politically because it requires agreement about what the new baseline should be — and different variations pull in different directions. Designing an environment optimally for autistic sensory profiles may create environments that neurotypical people find unstimulating; designing for maximum accessibility for mobility-impaired users may create configurations that are inefficient for those without mobility impairment; designing for left-handers disadvantages right-handers and vice versa, unless the design is agnostic to handedness entirely. Universal design — the principle of designing environments that are usable by the full range of human variation without adaptation or specialized provision — resolves some of these tensions but not all. Some variations are genuinely incompatible in the same space: a classroom optimized for students who learn best through silent individual reading is not the same classroom as one optimized for students who learn best through verbal and kinesthetic engagement. Transformation cannot dissolve all accommodation requirements; it can only reduce them by moving the baseline closer to the center of the actual distribution of human variation. The Reproduction Problem: Accommodation That Entrenches What It Compensates For A structural limitation of accommodation that is rarely acknowledged in policy discussions is its tendency to reproduce the conditions it compensates for. When accommodation becomes the recognized response to a variation, the pressure to redesign the

baseline is reduced: the variation can participate, after all, through the accommodation pathway. This produces a dynamic in which accommodation and transformation are not complementary strategies but competitive ones — each unit of accommodation energy spent reduces the political pressure for transformation, because the most visible dimension of the problem (exclusion) has been addressed. The result is an institutional architecture that accumulates accommodations over time while the baseline remains unchanged, becoming progressively more complex and more expensive to maintain, and increasingly organized around the assumption that the baseline is natural and the accommodations are add-ons. This dynamic is visible in how neurodivergent accommodation in educational and employment contexts has evolved. The expansion of individual accommodation plans — extended test time, quiet rooms, sensory

accommodations, communication alternatives — has been genuinely beneficial for many individuals who received them. It has not produced corresponding pressure to redesign assessment methodologies, physical environments, or communication norms so that fewer individuals require individual accommodation. The accommodation approach individualizes what is a structural problem: each person with a variation requiring accommodation has a case, a plan, a set of provisions — which positions them as individually exceptional rather than as evidence that the baseline was designed too narrowly. The accumulation of individual exceptions does not by itself challenge the norm that produces the exceptions; it can coexist with and implicitly reinforce it. PMN cannot prescribe a general resolution to this tension because the correct balance between accommodation and transformation is contextually variable. In domains where transformation is feasible — urban design, software interfaces, communication formats, assessment methodologies — the framework's evaluative criterion favors transformation over accommodation because transformation reduces rather than manages the structural source of suffering. In domains where transformation is not feasible within relevant time horizons — existing physical infrastructure, established professional communication norms, rapidly evolving technological environments where stable redesign is difficult — accommodation is the appropriate primary response, with transformation as a longer-term goal. The analytical obligation is to be explicit about which strategy is being pursued and why, rather than treating accommodation as the natural and permanent institutional response to variation and transformation as an unrealistic aspiration. The Stakes of the Choice The accommodation-transformation tension connects to the broader political economy of who bears the costs of institutional design. Accommodation places the cost of managing variation asymmetrically on the people who carry the variation: they must navigate the accommodation process, maintain documentati

on of their needs, advocate for their provisions, and manage the social dynamics of being visibly accommodated. This cost is not trivial and is not evenly distributed — people with more resources, more institutional fluency, and more social capital are better positioned to secure and maintain accommodations than those without. Transformation distributes the cost differently: it requires investment from whoever controls the institutional design, which typically means whoever benefits from the current baseline. The political resistance to transformation is therefore structurally predictable: it is resistance by those who benefit from the current baseline to bearing the costs of changing it, expressed in the language of practicality, tradition, or universal standards. Accommodation-transformation assessment: (1) Feasibility horizon — within what time horizon is redesign of the baseline feasible? This determines whether accommodation or transformation is the appropriate primary response in the immediate term. (2) Reproduction check — is the accommodation pathway reducing pressure for baseline transformation? If yes, the accommodation is addressing the exclusion symptom while entrenching the design cause. (3) Cost distribution — who bears the cost of managing the variation under the current arrangement: the people who carry the variation (accommodation model) or the institution that set the baseline (transformation model)? Baseline costs that are systematically transferred to variation-carriers are a structural harm indicator. (4) Normative baseline audit — is the current baseline a biological necessity or a contingent historical default? If the latter, who set the default, whose variation it was optimized for, and who was excluded from the design process? When accommodation for one variation conflicts with accommodation for another — sensory quietness required by one population and verbal expressiveness required by another, physical accessibility that creates temporal barriers for a third — the framework does not supply a simple priority rule. Instead, apply the suffering comparison procedure in 3.4e: identify which configuration minimizes imposed, dysfunctional, preventable, structural suffering, with particular weight on irreversibility and foreclosure of becoming. The population whose unmet need produces the most severe and least reversible foreclosure of developmental trajectory takes priority, not because its suffering is more real but because the 3.4e criteria — structural causation, foreclosure, reversibility, distribution — point in that direction when applied to the specific case. This is not a resolution; it is a structured procedure for making the trade-off explicit and for identifying what is being sacrificed and by whom. The accommodation-transformation diagnostic (this section) still applies: even after applying 3.4e, the question remains whether the entire baseline could be redesigned to eliminate the conflict

rather than manage it.

13.4g — The Agency-Structure Tension in Internalized Norms: The Apology Problem

People who have lived under structural arrangements that systematically constrain their range of possible lives develop, as a predictable consequence, preferences, self-understandings, and normative frameworks that are adapted to those constraints. This is not manipulation in a simple sense — it is the ordinary process through which individuals develop stable orientations toward the conditions they actually inhabit. The problem the agency-structure tension names is: when someone whose preferences were formed under structural constraint expresses those preferences, what is the analytical and normative status of that expression? Is it genuine agency that institutions should respect, or is it an effect of structural conditions that institutions should work to change? The apology problem: The tension takes its sharpest form in what might be called the apology problem: the situation where someone who has been harmed by structural arrangements, or who belongs to a group systematically constrained by them, does not describe their situation as harmful or unjust — and may actively resist that description. The person who has internalized the norms of an arrangement that constrains them is exercising real agency when they endorse those norms. They are also expressing something that structural analysis attributes partly to the conditions that shaped them. Both things are simultaneously true. The tension cannot be resolved by privileging one over the other — by treating all preferences formed under structural constraint as structurally determined (which eliminates agency), or by treating all expressed preferences as fully authentic (which ignores how preferences are formed). The voluntarism error and the determinism error: Two characteristic errors flank this tension. The voluntarism error treats all expressed preferences as fully voluntary expressions of authentic selfhood, regardless of the structural conditions under which they were formed. Applied to preferences formed under conditions of structural constraint, this error produces policies and analyses that take adaptation to harm as evidence that the harm is not harm — that the absence of explicit complaint signals the absence of a problem. The determinism error treats preferences formed under structural constraint as structurally determined in a way that cancels their normative significance — that the person does not really choose what they choose because their conditions shaped their choosing. This error eliminates agency entirely, licenses paternalistic override of expressed preferences in the name of the person's "real" interests, and reproduces the structural logic of treating constrained populations as incapable of knowing what is good for them. The practical navigation: The framework's approach to this tension proceeds through explicit attention to the conditions under which preferences were formed, without treating those conditions as determinative of their normative status. The relevant question

is not "is this preference authentic?" — that question has no reliable answer — but: under what material conditions did this preference develop, what alternatives were practically available, and what would change in the preference if the structural conditions changed? A preference that is stable across significant changes in structural conditions carries more weight as an expression of genuine agency than one that shifts dramatically when constraints are lifted. This is not a perfect procedure — preferences can be stable for reasons other than authenticity — but it is more analytically useful than either the voluntarism or determinism alternative. The institutional design implication: institutions that affect populations whose preferences were formed under structural constraint should build in mechanisms for preference revision as conditions change — not to override expressed preferences, but to create genuine conditions for their revision when the constraints that shaped them are removed. This is what distinguishes paternalism from structural support: paternalism substitutes the institution's judgment for the person's; structural support creates the material conditions under which the person's judgment can be formed and expressed without the adaptations that structural constraint produces. The difference is procedural as much as substantive — it is in whether the institution's design includes genuine mechanisms for the person's judgment to revise the institutional response, not only the institution's judgment to determine what the person needs. The Apology

as a Structural Product and a Personal Act The person who apologizes for using their left hand is doing something that is simultaneously a structural product and a personal act. It is a structural product because the apologetic response was not chosen — it was installed through years of social interaction that rewarded right-handed behavior and marked left-handed behavior as requiring justification. The affective pattern — the pre-reflective impulse to apologize — is the output of a norm-transmission process that section 5.6 analyzes: norms transmitted through practice rather than proposition, installed as felt responses before they are available for rational examination. In this sense, the apology is not freely chosen any more than the person's left-handedness is freely chosen. Both are facts about the person that preceded their capacity to evaluate them. But the apology is also a personal act in a sense that left-handedness is not. The person performing the apology is an agent whose relationship to the norm can, under conditions of adequate support and analytical tools, be reorganized. They can come to recognize the apologetic impulse as a structural product, examine its normative basis, and develop a different behavioral response — not through an act of pure will that overwrites the installed pattern, but through the gradual process of re-practice that section 5.6 identifies as the mecha

nism through which meaning frameworks and behavioral habits change. The structural analysis of how the apology was produced does not abolish the possibility of the person choosing, over time, to respond differently. It specifies the conditions under which that choice is more or less available. The tension between these two accounts — the apology as structural product vs the apology as personal act — is not merely academic. It has direct institutional implications. If the apology is primarily a structural product, the appropriate institutional response focuses on changing the structural conditions that produce it: designing environments that do not mark left-handedness as exceptional, withdrawing the social rewards for right-handed behavior and the social costs for left-handed behavior, and providing contexts in which the installed apologetic response is not triggered. If the apology is primarily a personal act, the appropriate response focuses on the individual: education, therapeutic support, consciousness-raising, the provision of analytical tools that allow the person to recognize and reorganize their relationship to the internalized norm. Both framings are partially correct. The tension is that emphasizing the structural framing can remove the person's sense of agency over their own response, while emphasizing the personal framing can make the person responsible for correcting a pattern that structural conditions continue to produce. Two Errors and the Space Between Them The first error — structural determinism applied to internalized norms — treats the individual as a passive recipient of structural inputs with no significant capacity for reorganization of their relationship to those inputs. This error is attractive because it correctly identifies the structural production of the norm and correctly resists the impulse to blame individuals for conditions they did not choose. Its cost is that it removes the individual from the analysis in a way that is both descriptively inaccurate and practically unhelpful. People do reorganize their relationships to internalized norms — not easily, not reliably, not without adequate support and conditions, but genuinely. Structural analysis that denies this capacity does not honor the people it intends to support; it replaces one form of inadequacy attribution (the norm is your fault) with another (you have no capacity to respond to the norm differently). The second error — voluntarism applied to internalized norms — treats the reorganization of the relationship to an internalized norm as straightforwardly available to any individual who understands the structural analysis and chooses to apply it. This error is attractive because it preserves individual agency and avoids the paternalism of structural determinism. Its cost is that it systematically overestimates the ease of norm reorganization and systematically underestimates the continuing structural conditions that reproduce the norm even in individuals wh

o have intellectually rejected it. The left-handed person who has explicitly rejected the norm that left-handedness requires apology will still encounter social contexts that trigger the apologetic response, because those contexts have not changed. Voluntarism cannot account for the gap between intellectual rejection of a norm and affective reorganization of one's response to it — a gap that can persist for years or indefinitely depending on the density of the social contexts that reproduce the triggering conditions. The

space between these errors is occupied by what might be called conditional agency: the capacity of individuals to reorganize their relationship to internalized norms is real but structurally conditioned. It is more available when: the structural sources of the norm are made visible through analysis (which is the function of section 1.12's epistemic infrastructure discussion); the social contexts that trigger the installed response are redesigned or made less frequent; other individuals with similar variations are available as reference points that do not confirm the deficit framing; and the process of reorganization is supported rather than required. The institutional implication is that structural redesign and support for individual reorganization are complementary rather than alternative responses — structural redesign changes the conditions that produce and reproduce the norm, while individual support creates the resources through which reorganization can occur within conditions that structural redesign has not yet changed. The Collective Dimension: When Internalized Norms Maintain Structural Arrangements Internalized norms do not only harm the individuals who carry them. They also maintain the structural arrangements that produce the harm, by generating voluntary compliance that structural enforcement alone would not achieve. If everyone who carries a stigmatized variation apologizes for it, the norm that the variation requires apology is continuously reproduced without requiring explicit enforcement. This is the most efficient form of structural maintenance: the people who bear the cost of the norm participate in its reproduction because the norm has been successfully installed as the self-understanding through which they experience their own condition. The variation-as-deficiency framing maintains itself through the apologetic behavior it produces, which confirms to everyone observing that the variation does indeed require apology, which reinforces the norm for the next generation of people who carry the variation. This collective dimension means that the agency-structure tension in internalized norms has a political as well as an individual dimension. Individual reorganization of the relationship to an internalized norm is personally significant but structurally limited: one person who stops apologizing for their left hand does not change the norm. Collective reorganization — the coordinated refusal of the apologetic frame

ing by sufficient numbers of people carrying the variation — has structural significance because it disrupts the norm-reproduction mechanism that voluntary compliance provides. This is the logic of identity-affirmative movements: not merely the psychological benefit to individuals who participate, but the structural effect of withdrawing the voluntary compliance that the norm requires for reproduction. PMN's analysis of this dynamic does not prescribe the specific form that collective reorganization should take, because that is a strategic question that depends on the specific variation, the specific institutional context, and the specific historical moment. It identifies the structural mechanism that makes collective reorganization more than individual therapy at scale. Agency-structure diagnostic for internalized norms: (1) Installation mechanism — through what specific practice-based processes was this norm installed, and what social contexts continue to trigger the installed response? Understanding the installation mechanism specifies what structural redesign would reduce the triggering contexts. (2) Reorganization conditions — what conditions make individual reorganization of the relationship to this norm more available: visibility of structural sources, reduced triggering contexts, available reference points that do not confirm the deficit framing, supported rather than required reorganization processes? (3) Compliance contribution — to what degree does the voluntary compliance that the internalized norm produces maintain the structural arrangement that produces the norm? High compliance contribution indicates that collective reorganization has structural significance beyond individual benefit. (4) Attribution error risk — does the current institutional response primarily attribute the internalized norm to individual failure (voluntarism error) or deny individual capacity for reorganization (determinism error)? Both errors produce characteristic harms that structural-aware institutional design should be designed to avoid.

13.4h — The Moral-Functional Tension: When Killing the Institution Kills What It Was Doing

Some institutions that are morally objectionable by PMN's evaluative criterion are also performing functions that are materially necessary — functions that would not automatically be performed by whatever replaced them if they were abolished or fundamentally transformed. The tension this creates is not the straightforward case of weighing current harm against future benefit that consequentialist reasoning handles, because the functions being performed by the morally objectionable institution are often functions for the populations that institution is simultaneously harming. The analysis requires tracking both the harm the institution produces and the functions it performs for those it harms, and asking what happens to those functions if the institution is killed. The canonical case: The canonical instance is the welfare state under conditions of austerity and structural adjustment. The welfare state in most capitalist economies was built partly as a mechanism for managing the political consequences of structural suffering, partly as a genuine commitment to social protection, and partly as an instrument of labor discipline. All three functions are real and simultaneously present. From the standpoint of the minimal anchor, the welfare state's protective function — the actual harm prevention it performs for the populations most dependent on it — is genuine and material. From the standpoint of structural analysis, its labor discipline function and its political management function serve the arrangement it is nominally protective against. Abolishing it eliminates the objectionable functions and the protective ones simultaneously. Reforming it is subject to the capture dynamics that produced the objectionable functions in the first place. The sequencing requirement: The moral-functional tension generates a specific analytical obligation: the burden of proof for institutional abolition or fundamental transformation must include an account of what happens to the functions the institution performs for those it harms during the transition period. A transformation strategy that answers the question "what replaces this?" with "something better" has not met the burden. The question requires a specific account of the organizational capacity, resource base, and institutional infrastructure that would perform the protective function during the period between abolition and replacement. When that account is not available, the transformation strategy is asking the populations currently dependent on the institution to bear the cost of a transition whose timing and outcome they do not control. That is a real cost that the framework must account for, not a compromise to be dismissed as insufficient radicalism. When the institution is also the problem: The hardest version of this tension is when the institution performing the protective function is also the primary mechanism producing the structural harm — when the same institution that provides material support also c

ontrols, surveils, and disciplines the populations it supports. This describes the situation of carceral systems that also provide mental health and social services, immigration detention systems that also process asylum claims, and workplace systems that provide healthcare through an arrangement that ties healthcare access to labor market participation. In each case, the institution's protective function creates a dependency that the institution's harmful function exploits. Abolishing the institution removes both the harm and the protection simultaneously. The analysis must track this entanglement explicitly rather than treating it as a reason to defer transformation indefinitely or as an argument for tolerating the harm because the protection is real. The framework's position is that the moral-functional tension cannot be resolved in advance of the specific conditions of any particular case, but it can be analyzed with precision. The relevant variables are: (1) the magnitude of the protective function and the population dependent on it, (2) the realistic availability of alternative provision during any transition period, (3) the probability and timeline of transition success, and (4) the ongoing harm produced by the institution during the period of delayed transformation. Where the protective function is large, alternative provision is unavailable, and transition timelines are long, the framework generates no general answer about whether to prioritize abolition or maintenance — it generates a demand for explicit accounting of the costs that each choice imposes on specific populations, and for whoever bears those costs to have genuine input into the strategic decision. The Gods Were Not Only Believed In This is not an argument for preserving traditional religious institutions. It is an analytical observation about the structure of the problem that confronts any serious institutional critique. Institutions perform functions that exceed, and are often obscured by, their stated purposes and their moral legitimacy or

illegitimacy. A feudal lord extracted surplus, maintained violence, and enforced hierarchy — morally indefensible on PMN's evaluative criterion. He also adjudicated local disputes, coordinated agricultural activity across the manor, maintained the infrastructure of paths and bridges, and provided a framework within which the peasants under his jurisdiction had a defined place, set of obligations, and set of entitlements. The moral critique of feudalism is correct. The institutional critique of what replaced it must ask: which of those functions survived the transition, in what institutional form, and which were simply lost, producing the specific configurations of landless labor, urban squalor, and social atomization that early industrial capitalism generated alongside its productivity gains? The second-order effects of institutional destruction are not a conservative argument against institutional critique. They are a materialist oblig

ation of it. Two Dimensions That Must Not Be Conflated: Moral Legitimacy and Functional Necessity The conflation of these dimensions produces two characteristic errors with opposite political valences. The first — treating moral legitimacy as sufficient evidence of functional dispensability — is the error of blind idealism: the position that an institution's moral condemnation is sufficient grounds for its elimination, without asking what it was doing that now needs to be done differently. This error is most visible in revolutionary programs that correctly identify the moral bankruptcy of the institutional order they target and then systematically underestimate the coordination functions embedded in that order, generating post-revolutionary configurations that reproduce hierarchy through different institutional vehicles or produce genuine disintegration of the coordination functions that the previous order, however unjustly, maintained. The second error — treating functional necessity as sufficient evidence of moral legitimacy — is the error of conservative capture: the position that because an institution performs necessary functions, its moral character is beyond critique, or that any challenge to it must wait until a fully specified functional replacement is ready. This error serves the interests of institutions that benefit from their functional indispensability as cover for their moral illegitimacy. The relationship between these two errors maps onto the strategic equilibrium tension analyzed in section 13.4d, but is distinct from it. Section 13.4d asks whether the choice between transformation and preservation is a false binary and what the actual costs of each position are. This section asks a prior question: before choosing a position on the transformation-preservation spectrum, has the analysis correctly identified what the institution is functionally doing, distinct from whether it should be doing it and whether it is doing it justly? The strategic question cannot be answered well without the functional analysis, because the strategic trade-offs depend on what functions are at stake in the institutional change. The Vacuum Problem: What Fills the Space The vacuum problem is structurally asymmetric in a way that idealism unconstrained by functional analysis ignores. The institution being critiqued is visible — its costs, its hierarchies, its moral failings are available for examination because the institution exists and produces observable effects. The institution that would replace it is not yet visible, and the actors who would produce the replacement institution may not be the actors who produced the critique. Revolutionary history is substantially a history of this asymmetry: the critique of the existing order is generated by one set of actors with one set of values and intentions; the filling of the vacuum created by that order's collapse is accomplished by a different set of actors — often the best-organized, best-resourced, and mos

t violence-capable actors available at the moment of maximum institutional fluidity — whose values and intentions may differ substantially from those who generated the critique. This is not an argument against institutional transformation. It is an argument for disciplined attention to the vacuum-filling mechanism as a central object of strategic analysis rather than an afterthought to the moral critique. The historical cases are abundant enough that the pattern should be treated as a default expectation rather than an anomaly. The dismantling of traditional village structures under colonial modernity produced not the rational self-organizing communities that modernization theory predicted but the specific configurations of urban precarity, kinship network collapse, and political susceptibility that post-colonial states inherited and struggled to manage. The elimination of guild structures under early industrial capitalism produced not the free market of voluntary exchange that liberal theory anticipated but the specific power asymmetries of

unprotected wage labor that labor movements spent two centuries partially correcting. The removal of state socialist institutions in Eastern Europe and the former Soviet Union produced not the pluralist democracies that transition theory projected but the specific configurations of oligarchic capture, institutional vacuum, and resurgent nationalism that characterized the 1990s. In each case, the moral critique of the existing institution was substantially correct. The second-order analysis of what the vacuum would produce was substantially absent from the programs that drove the transformation. Applied Case: De-Ba'athification and the Bureaucratic Vacuum The Coalition Provisional Authority's de-Ba'athification program in Iraq — implemented through CPA Orders 1 and 2 in May 2003 — is among the most analytically instructive recent instances of the vacuum problem, because the moral case for it was substantially correct and the functional analysis was substantially absent. Order 1 barred senior Ba'athists from public employment across all ministries; Order 2 dissolved the Iraqi military entirely, approximately 400,000 soldiers, a large proportion of them conscripts with no alternative income and no institutional obligation to the incoming order. The Ba'ath Party under Saddam Hussein was the apparatus of a regime responsible for mass atrocity, systematic political repression, and ethnic cleansing operations against Kurdish and Shia populations. By PMN's evaluative criterion, the moral indictment was not disputable. What was disputable — and what the functional analysis that was absent from the program would have disputed — was the inference from moral indictment to institutional elimination without specification of what would perform the functions the institution was performing. A prior functional inventory would have identified the following. The Ba'athist bureaucracy was simultaneously the organizational infrast

ructure through which the Iraqi state operated its electricity grid, water treatment facilities, customs administration, civil service, and educational system. De-Ba'athification removed not only the repressive cadres but the trained personnel on whom basic service delivery depended. Military dissolution not only disbanded an instrument of regime violence but produced a pool of 400,000 armed, trained men with acute economic grievances, organizational capacity, and no institutional loyalty to the order replacing the one that had employed them. The vacuum specification would have identified that no alternative institution existed to absorb either category of function; that the best-organized vacuum-filling actors available at the moment of maximum institutional fluidity were Shia militias with Iranian institutional backing, Sunni insurgent networks drawing on the tribal and professional networks of the dissolved officer corps, and eventually an organization whose senior military leadership was recruited disproportionately from the former Iraqi Army. The second-order trajectory that this analysis projects is not speculative — it is the documented history of post-2003 Iraq: sectarian fragmentation, insurgency infrastructure, civil war, and the emergence of the Islamic State as the primary beneficiary of the institutional vacuum the program created. The five-step diagnostic procedure specified in this section's operational callout generates a different institutional recommendation when applied before rather than after the fact. Steps (1) through (3) — functional inventory, vacuum specification, second-order trajectory — do not challenge the moral analysis of the Ba'athist institutional order; that analysis is supported by the evaluative criterion. What they do not support is a transformation program whose functional replacement specification is near-absent while its moral justification is developed with care. The de-Ba'athification program exhibited this signature precisely: the moral case for dissolving Ba'athist institutions was elaborated; the question of who would perform the functions those institutions were performing — with what organizational capacity, resource base, and political authorization, and through what mechanisms during the transition period — was not answered with comparable specificity before the dissolution was implemented. Step (5) — the idealism check — fails clearly. The second-order configuration that resulted was worse on the evaluative criterion than a modified or partially reformed institutional order would have been, and the population bearing the cost of that analytical failure was the Iraqi population the program was nominally designed to benefit. This is the structural injustice that the moral-functional tension generates when the idealism check is not applied: the costs of moral improvement fall on those already bearing the costs of moral failure. Institutional Functions as Latent Infrastructure The concept of latent i

institutional infrastructure captures this dynamic. An institution performs latent infrastructure functions when it maintains conditions that other activities depend on, without those conditions being explicitly attributed to the institution by the actors who benefit from them. The safety of public space depends on a configuration of informal social norms, formal law enforcement presence, and architectural design whose coordination function is invisible to the pedestrian who walks through it safely and becomes visible only when one or more elements fail. The functioning of markets depends on a background of contract enforcement, property rights protection, and trust infrastructure whose institutional basis is invisible to the market participant who benefits from it and becomes visible when it breaks down, as in the institutional vacuum periods following state collapse. The meaning-providing function of civic institutions — collective celebrations, shared narratives, commemorative practices — is invisible to the citizen who participates in them routinely and becomes visible in the specific forms of political susceptibility that emerge when those institutions are absent or have been delegitimized without replacement. PMN's evaluative criterion requires the identification of latent infrastructure functions as part of any institutional assessment, because the costs of removing latent infrastructure fall not on those who identify the institution's moral failings but on the populations whose activities depended on the functions the institution was invisibly providing. This is a structural injustice that institutional critique conducted purely on moral grounds systematically produces: the costs of moral improvement fall on those who were already bearing the costs of moral failure, because they were the ones depending on the latent infrastructure that the morally illegitimate institution was providing. The Grounded Realism Requirement: Idealism Bounded by Second-Order Consequence PMN's commitment to adaptive response capacity and the reduction of structural suffering is not compatible with institutional recommendations that produce worse structural configurations in the name of correct moral principles. The evaluative criterion is consequentialist in the relevant sense: it asks what outcomes institutional arrangements produce, not whether they express correct values. An institutional program that correctly identifies moral illegitimacy and then generates a transformation whose second-order effects are worse than the first-order problem it addressed has not advanced the evaluative criterion — it has retarded it, while generating the narrative resources for a conservative restoration that will claim, with partial accuracy, that things were better before. This does not mean that transformation must wait for certainty. PMN does not require that the post-transformation institutional configuration be fully specified before transformation can be recommended — ce

rtainty of that kind is never available for significant institutional changes, and requiring it would permanently privilege existing configurations regardless of their moral character. What it requires is that the vacuum-filling analysis be conducted honestly and that its conclusions be incorporated into the transformation program, not treated as an afterthought to the moral critique. This means identifying the specific functions the institution to be transformed is performing, specifying the mechanisms through which those functions will be maintained through and after the transformation, building those mechanisms before or simultaneously with the transformation rather than expecting them to emerge spontaneously from the moral correctness of the program, and being willing to revise the transformation timeline or strategy when the vacuum-filling analysis reveals that the post-transformation configuration is likely to be worse on the evaluative criterion than the configuration being transformed. The political difficulty of the grounded realism requirement is that it looks like conservatism from the outside — and will be characterized as such by actors on both sides who benefit from the false binary between idealist transformation and conservative preservation. The idealist transformation camp will characterize it as insufficient commitment to moral principle; the conservative preservation camp will selectively cite it as evidence that the current order should not be disturbed. Neither characterization is accurate. The grounded realism requirement is not a position on the transformation-preservation spectrum. It is a methodological standard that applies equally to both: both transformation programs and preservation programs must be evaluated on their actual second-order consequences, not on the moral clarity of their normative commitments. The conservative institution that invokes functional necessity as cover for moral illegitimacy fails this standard as surely as the idealist

program that invokes moral legitimacy as sufficient grounds for institutional elimination without functional analysis. The Resolution That Is Not Available This tension cannot be resolved because its two poles — the moral obligation to critique institutions that produce unjust outcomes, and the functional obligation to analyze what fills the vacuum their transformation creates — are both genuine and both required by the evaluative criterion. A PMN that abandoned the moral critique of institutions on the grounds that any existing institution is performing some function that might be disrupted would be a PMN in the service of the status quo. A PMN that conducted moral critique without functional analysis would be a PMN systematically liable to generating worse outcomes in the name of better values. Neither abdication is available. What is available — and what the framework requires — is the disciplined practice of conducting both analyses simultaneously and refusing to allow ei

ther to serve as a trump card for the other. Moral illegitimacy is not sufficient for functional dispensability. Functional necessity is not sufficient for moral legitimacy. Every institutional assessment requires both questions to be answered, and every institutional program requires both the moral justification for what is being changed and the functional specification of how the change will be managed without producing second-order configurations that the moral critique was intended to prevent. The standard is demanding. It is more demanding than either pure idealism or pure functionalism. It is the standard that the evaluative criterion requires, and the history of institutional transformation makes clear why. Moral-functional institutional assessment: (1) Functional inventory — what coordination functions is this institution currently performing, including latent infrastructure functions that become visible only in the institution's absence? List specifically: conflict adjudication, resource allocation, meaning provision, floor condition maintenance, trust infrastructure. (2) Vacuum specification — if this institution were removed or radically transformed, who fills the space, through what mechanisms, with what resources, toward what institutional configurations? The answer "a better institution will emerge" is not a specification. (3) Second-order trajectory — is the post-transformation configuration that the vacuum analysis identifies likely to be better or worse on the evaluative criterion than the current institution, accounting for both the institution's moral failings and its functional contributions? (4) Moral-functional trade-off — if the institution is morally illegitimate but functionally necessary in the short to medium term, what is the minimum structural change that reduces the moral harm while preserving the functional contribution, and what is the timeline for developing functional alternatives that would allow more complete transformation? (5) Idealism check — does the transformation program specify functional replacement mechanisms with the same specificity as it specifies moral justification? If the moral analysis is detailed and the functional replacement analysis is vague, the program has not met the grounded realism requirement. Two further cases follow from this tension and are treated separately: the triage protocol for high-urgency claims in 13.4i, and the structural-tragedy configuration in 13.4j, where every available option produces structural suffering.

13.4i — Triage Protocol for High-Urgency Cases

The five-step assessment procedure in 13.4h assumes time for analysis. In genuine emergencies — a hospital controlled by actors whose interests diverge sharply from its patients, a food distribution system captured mid-crisis — the agent cannot complete a full functional inventory before acting. Requiring the full procedure under these conditions would itself constitute a moral-functional failure: the time spent on analysis would convert an addressable harm into an irreversible one. Definition of extreme urgency: a situation is operationally extreme-urgency when two conditions are simultaneously present. First, there is credible evidence of severe harm occurring or imminent — harm that meets at least three of the four criteria for unnecessary suffering (imposed, dysfunctional, preventable, structural). Second, delay beyond hours or a small number of days would convert the harm from preventable to irreversible, or would prevent intervention altogether. These conditions are conjunctive. A severe harm that can still be adequately addressed in two weeks does not qualify. A time-pressured situation that does not involve severe harm does not qualify. The

double condition prevents urgency from being invoked strategically to bypass analysis that is merely inconvenient rather than genuinely impossible. Triage default rule: in confirmed extreme-urgency situations, act to stop or mitigate immediate severe harm without completing the full five-step procedure. The truncated justification required is: (a) what is the specific harm being stopped, (b) what is the evidence it is occurring or imminent, and (c) what is the minimum intervention that addresses the immediate harm without unnecessarily destroying functional infrastructure that serves other populations? The third question is the emergency compression of steps 1 and 2 of the full procedure — not an elimination of the functional analysis but a reduction of it to what can be assessed in the available time. Post-action obligation: the emergency override is not a permanent authorization. It is a provisional action taken under constraint that creates an obligation — not a recommendation — to return to the full five-step procedure within a defined period after the emergency has passed. The post-action analysis serves two functions: it determines whether the emergency action was adequate to the harm and whether further action is required; and it prevents the emergency override logic from normalizing as a standing alternative to full analysis. An agent who systematically invokes urgency without completing post-action analysis is not applying the triage protocol — they are using it to exempt themselves from the accountability the five-step procedure provides. Case: field hospital in conflict zone. A functional hospital is being used by occupying forces to process prisoners in ways that directly harm patients. Urgency conditions are met: harm is ongoing, severe, and delay converts addressable cases to irreversible

. The triage default applies: act to stop direct patient harm, with the minimum intervention consistent with preserving the hospital's function for the civilian population it serves. The post-action analysis must then address: did the intervention preserve the civilian-serving function, or did it destroy capacity that now requires reconstruction? What institutional arrangements prevented the harm from being visible earlier? The emergency action without the post-action analysis is half a procedure.

13.4j — When All Available Options Produce Structural Suffering: The Structural Tragedy Problem

A distinct class of limitation requires explicit acknowledgment: situations in which every available institutional option produces structural suffering on the framework's own criterion, and no option exists that reduces structural suffering for one population without increasing it for another. The framework's standard prescriptions — identify the arrangement producing unnecessary suffering, specify the transformation that would reduce it, evaluate the institutional changes required — assume that a better arrangement is available, even if reaching it is difficult. The structural tragedy problem arises when that assumption fails: when the genuine constraint is not the difficulty of reaching a better arrangement but the absence of one. The framework resists this framing as a starting point, because its most common use is to justify arrangements that do produce unnecessary suffering by asserting, without adequate analysis, that no alternative exists. This is the characteristic move of conservative institutional defense: naturalizing contingent arrangements as necessary by presenting alternatives as impossible. The framework's default position — that visible structural suffering almost always has a structural cause and almost always has an alternative arrangement that would reduce it — is analytically warranted by the historical record, which shows that arrangements previously presented as necessary alternatives have, under pressure, proved contingent and changeable. Any analysis that invokes the structural tragedy framing without first completing the full alternatives assessment has invoked a conclusion that requires rigorous prior demonstration. The conditions under which the structural tragedy framing is analytically legitimate: the claim that all available options produce structural suffering must survive the stress-test methodology in 12.1 — specifically, the reversal test (does the claim hold when applied symmetrically to the actors who benefit from the current arrangement?), the opposing constituency test (have the alternatives been assessed from the perspective of those who bear the costs of the arrangement being defended?), and the historical test (have comparable configurations in comparable

historical conditions produced genuine transformations that the current analysis treats as unavailable?). An analysis that has completed all three tests and still concludes that no option avoids structural suffering has met the evidentiary standard. An analysis that skips any of the three has not. The conclusion must additionally be verified by actors who do not benefit from it — precisely because the Custodian Problem in 7.3 applies to tragic choice declarations with particular force: an actor who stands to benefit from accepting a specific harmful arrangement as necessary has both the incentive and the institutional capacity to assess alternatives in ways that systematically undercount them. When the structural tragedy conclusion survives evidentiary s

crutiny, the framework's analytical function shifts from identifying the arrangement that reduces suffering to specifying what the constrained decision requires. Four demands follow. First, the specific populations bearing the costs and the specific populations bearing the benefits must be named explicitly — no tragedy framing is analytically complete without this specification. Second, the distribution requirement in 8.3 applies with heightened force: a tragic choice that systematically places costs on already-structurally-disadvantaged populations while placing benefits on already-advantaged ones fails the consistency standard regardless of the aggregate calculation. Third, the decision must include explicit provisions for restoring conditions for those who bear the costs — not as a compensatory afterthought, but as a structural commitment undertaken before the decision is implemented. Fourth, the decision is provisional: it is the best available option given current conditions, not a settled institutional arrangement, and must be monitored for the appearance of alternatives the current analysis did not identify. What the framework cannot provide: it cannot identify which specific decision should be made in a genuine structural tragedy. That requires detailed empirical knowledge of specific conditions that the framework, as a general analytical tool, does not possess. What it can do is distinguish genuine structural tragedy from the use of tragedy framing to exempt harmful arrangements from scrutiny, and specify the analytical and procedural conditions under which a constrained choice can be made with minimum additional harm beyond what the constraint itself requires.

13.5 — Closing: The Framework That Must Be Surpassed

This document began with a displacement — a shift in where the analysis starts. It ends with a commitment that follows from that displacement as a logical consequence: the framework that has been built here must be surpassed by something better, and its success will be measured entirely by whether it contributes to producing that something better. The aspiration toward adaptive response capacity — toward the conditions in which human beings can develop beyond what structural arrangements have so far allowed — applies to the analytical tools we use as much as to the human lives they are meant to serve. A framework that makes itself permanent, that defends its conclusions against the evidence that would require revising them, that mistakes the position it has reached for the destination it was pointing toward — that framework has become exactly what this analysis identifies as the primary obstacle to genuine development: an arrangement that serves its own continuation at the cost of the purpose that originally justified it. The belief that there is something more objective to move toward — a better analysis, a more adequate framework, a more honest account of the conditions that produce human flourishing and human suffering — is what makes the entire project possible. That doubt is what keeps it moving: that every current position is incomplete, that every current conclusion is provisional, that the framework in its current form has not reached the limit of what is possible. Belief and doubt in productive tension: not as a philosophical disposition but as the structural condition of any inquiry that takes seriously the gap between where it is and where it is trying to go. What follows from this document is not the application of its conclusions. It is the continuation of the inquiry it represents. The next position will see further than this one. That is not a failure of this document. It is the only form of success available to it. Progressive Materialist Naturalism is a name for a position in a movement, not a name for a finished system. The movement is older than this document and will continue

after it. What this document contributes is one articulation of that position — more adequate than what preceded it, and less adequate than what will follow. That is enough. The logical consequence of the displacement: If the analysis begins from the biological foundation rather than from a historical guarantee, then no current framework — including this one — can claim the status of a final account of what human beings require and what institutional arrangements best serve those requirements. The biological foundation establishes what is structurally invariant: the conditions for physical functioning, the cognitive architecture within which all political reasoning occurs, the relational requirements that development depends on. What it does not establish is which specific institutional arrangements best serve those conditions under conditions that are th

emselves changing. That is an ongoing empirical question, and any framework that treats its current answers as final has elevated its instruments to the status of horizons — the error that Part XI's distinction between horizon and instrument is designed to prevent. What surpassing this framework would require: A framework that surpasses PMN would need to be architecturally more adequate, not merely more politically appealing or more rhetorically sophisticated. Architectural adequacy means that the surpassing framework either identifies genuine errors in PMN's structural analysis — places where the causal claims are wrong, the variable relationships are misspecified, or the evaluative criteria are inadequate — or it extends the analysis into domains that PMN cannot reach. The most likely domains of extension are those identified in 13.4b, 13.4c, and 13.4d as limits: the meaning problem under conditions where adaptive response capacity is fully preserved, the subject boundary question as AI and biotechnology transform what a human biological subject is, and the strategic equilibrium problem between preservation and transformation. A framework that resolved those questions with analytical precision would represent genuine surpassing, not merely revision. The relationship between this framework and what comes after: The aspiration to be surpassed is not false modesty. It is the structural commitment that follows from the framework's own epistemology. Part I establishes that knowledge production has a social address and that frameworks systematically produce blind spots that follow from their production conditions. PMN cannot see its own blind spots — it can only identify where they are likely to be and create the conditions under which they are visible to critics with different structural positions. The framework that surpasses this one will be built partly from the critique of this one, which means that the quality of this framework's engagement with its own limits determines in part what the surpassing framework will be able to see. A framework that pretends to have no limits produces a different kind of successor than a framework that explicitly identifies where its architecture breaks down. The criterion for surpassing remains the criterion stated at the beginning: the minimal anchor, extended through the becoming framework of Part IV. A successor framework that produces institutional proposals better calibrated to reducing structural suffering and expanding development capacity will have surpassed PMN in the relevant sense regardless of whether it shares PMN's specific formulations. One that produces more theoretically elegant accounts of human social dynamics without improving the quality of institutional analysis by that criterion will not have surpassed it in the sense that matters. The practical demand on readers: the framework asks not for agreement but for engagement. The form of engagement the framework most values is the kind that identi

fies where its analysis is wrong in ways that would improve the quality of the institutional analysis that follows from it. Agreement with a framework that is itself limited in ways the framework cannot see is not the goal. The goal is the ongoing improvement of the analytical architecture, which requires the kind of criticism that the framework can only produce indirectly — by creating enough analytical precision that its errors become visible to those who encounter them with different structural positions and different critical capacities. The final self-application: this document is itself a case of the kind of structured analysis it recommends. It was produced under specific material conditions, by actors with specific structural positions, with specific blind spots that follow from those positions. It will be wrong in specific ways that its authors cannot see. The commitment to surpassing it is therefore a commitment made on behalf of the framework itself, not only by those who will engage with it from outside. The criterion — structural suffering —

remains.

XV: The Formula Architecture

The framework developed in Parts I through XIII is a tool for understanding how material reality is organized, how it moves, and what it produces. But a framework that remains purely interpretive — that can explain the past without generating structural expectations about the present — has not completed its work. Part XV provides a set of analytical scaffolding: heuristic variable structures that organize the framework's causal claims into comparable form. These are not mathematical models and do not generate quantitative predictions. They are a vocabulary for structural comparison — a way of making explicit which variables matter, in which direction, and how they relate — that complements the narrative analysis of Parts I through XIII without replacing it. Part XV makes explicit what the preceding parts imply but do not formalize: a set of structural formulas for anticipating the direction of change. Part X already developed the structural account of thresholds, counter-power accumulation, and transformation windows in prose form. Part XV formalizes that account into a formula set — the concepts are continuous with Part X, the formulas are new.

The relationship between Part X and Part XV is not redundancy — it is a shift in analytical register. Part X is a narrative of mechanism: it explains how threshold dynamics work, what counter-power accumulation looks like, and why political time is not uniform. That narrative produces understanding. Part XV is an architecture of variables: it makes explicit which variables drive each mechanism, how they relate to each other, and what changes in which variable would alter the trajectory. That architecture produces comparability. An analyst who has absorbed Part X understands why systems approach threshold. An analyst using Part XV can compare two systems — or the same system at two points in time — by examining where their variable profiles diverge, and can identify which specific variable would most change the trajectory if it shifted. The formula set does not replace the prose account; it operationalizes it for structured comparison across cases.

These are not predictions in the sense of specific events at specific times. They are predictions in the sense that physics predicts the behavior of gases under pressure — not which molecule will move where, but what the aggregate will do under specified conditions. The social analogue of thermodynamic law is not prophecy but structural tendency: given these inputs, these outputs become significantly more probable than their alternatives.

The formulas use variables and relational notation, but they are not mathematical equations in the formal sense. They express causal and structural relationships between identified variables. Coefficients exist — some variables are more powerful than others — but they cannot be assigned with current knowledge. The formulas are honest about this: they specify the direction and sign of relationships, not their precise magnitude.

15.0 — The Input Architecture of Structural Tendency Analysis

Before the formulas: what they cannot do. These limits are part of the theory, not caveats appended to protect it. First: the formulas predict structural tendency and threshold proximity, not the specific form, timing, or actors of transformation — contingent factors at the moment of threshold crossing are not determined by the structural analysis. Second: the relative weight of each variable is real but not currently measurable with sufficient precision to assign numbers; the formulas specify direction and sign, not magnitude. Third, and most important: the formulas describe what becomes probable under specified conditions — they do not determine what specific actors do with those probabilities. A movement that understands its structural position can act in ways that alter the trajectory. Treating structural tendency as deterministic eliminates the agency that structural understanding is supposed to enable. Fourth: each formula implies testable claims and is falsifiable; the appropriate response to contradicting evidence is revision, not special pleading. Section 15.13 elaborates each of these limits in full. What follows is offered with those limits already in view. The formulas that follow are heuristic tools for structural analysis, not mathematical models for quantitative

prediction. Their variables cannot be assigned precise numerical values; their operators do not license arithmetic. What they do is make explicit which variables matter, in which direction they push, and how they relate to each other. The test of a formula here is not whether it produces a number but whether it identifies the right variables, correctly describes the direction of their effects, and usefully organizes the analyst's attention. A formula that does those three things well is analytically powerful even without coefficients. A formula applied as though it were a calculating machine — assigning numbers, performing operations, producing outputs treated as predictions — is being misused regardless of the sophistication with which the misuse is conducted. Before the formulas, the epistemological foundation: what kinds of knowledge feed Structural Tendency Analysis, and how do they relate to each other? Three categories of knowledge are necessary and none is sufficient alone. Natural Sciences — geography, geology, oceanography, meteorology, ecology — provide the material constraints within which all human activity occurs. They determine what resources exist and where, what climatic conditions obtain, what geographic configurations enable or foreclose what kinds of political and economic organization. These are not background conditions that can be bracketed while the real analysis proceeds. They are determinants of what is structurally possible. A political analysis that ignores the geography of freshwater scarcity in the Middle East, or the geological distribution of energy resources, or the ecological limits of agricultural systems, is not analyzing the actual constr

aints under which actors operate. Human Sciences — psychology, sociology, biology, culture, religion, ideology, economics — provide the mechanisms through which material conditions are processed into human behavior, collective organization, meaning-making, and institutional form. Natural constraints do not determine outcomes directly; they determine them through the mediation of human response. The same geographic constraint produces different political outcomes in populations with different cultural traditions, different ideological frameworks, different economic arrangements. Human sciences explain the mechanisms of mediation — how constraints become responses, how responses become patterns, how patterns become institutions. Economics occupies a position that crosses both categories. As a study of resource allocation and exchange, it is partly a science of material constraints. As a study of institutions, incentives, and collective behavior, it is partly a human science. The cross-category nature of economics is not a classification problem — it is an accurate reflection of how economic phenomena actually work: simultaneously as material processes subject to physical limits and as social constructions subject to institutional design and ideological framing. Historical Dynamics — chronology, institutions, events, long-term patterns — provide trajectory: the path a system has taken, the momentum that path has generated, and the structural position the system now occupies as a result. Natural sciences and human sciences together describe what a system is capable of and how it tends to behave. Historical dynamics describe where it currently is in its development, what has accumulated and what has been depleted, what commitments have been made and what options have been foreclosed. Two systems with identical structural characteristics in the natural science and human science dimensions can be at very different points in their trajectories — one ascending, one collapsing — and that difference is not visible without the historical dimension. $STA = f_{Cs}(Ns \oplus Hs \oplus Hd)$ Where: f_{Cs} = Complex Systems Dynamics as transformative operator (non-linearity, feedback loops, emergence, threshold effects, path dependence, sensitive dependence, resilience/brittleness); $\oplus$ = dynamic interaction — not addition or multiplication, but the mutual co-determination of domains that cannot be fully separated. The notation signals that the relationship between Ns , Hs , and Hd is structurally intertwined rather than arithmetically combinable. $\oplus$ is borrowed but repurposed: it is not the mathematical direct sum or XOR, but a shorthand for the claim that each domain shapes and is shaped by the others in ways that make their joint behavior qualitatively different from their sum. The operator structure has a specific implication: f_{Cs} applied to any single domain in isolation does not produce Structural Tendency Analysis — it produces sophisticated domain analysis. f_{Cs} requires all

three inputs interacting via $\oplus$ to generate the cross-domain dynamics that make historical change qualitatively different from historical continuity. An analysis that knows the geography, the sociology, and the history of a situation but treats the three as independent factors produces an encyclopedia, not a predictive

model. The threshold effects, compounding feedbacks, and non-linear transitions only become visible when the three inputs are analyzed as a dynamically intertwined system — which is what f_Cs specifies. Conversely, f_Cs applied to thin inputs — to geography alone, or to ideology alone, without the full $\oplus$ interaction — generates sophisticated analysis of a sub-system that will miss the dynamics produced by the interaction between sub-systems. The $\oplus$ matters as much as f_Cs : all three inputs must be present and interacting for the operator to produce the full predictive output. The analytic philosophy objection — that the framework's core concepts are underspecified to the point of being inoperable, too elastic to generate clear predictions, and therefore functioning as sophisticated vocabulary for conclusions reached on other grounds — identifies a real cost and a real tradeoff. Formal conceptual precision of the kind analytic philosophy demands is achievable in domains where phenomena are sufficiently stable and measurable to support formal definition. In domains involving complex adaptive systems, threshold dynamics, and irreducibly contextual judgments, the demand for formal precision regularly produces the appearance of rigor at the cost of contact with what is actually being analyzed. The formula architecture of Part XV represents the furthest toward formal specification that the current analysis can honestly go. The acknowledgment made earlier in this section that the formulas express structural relationships rather than mathematical equations is not a retreat from precision — it is precision about what kind of precision is available. An analyst who can produce a more formally rigorous account of structural suffering that retains contact with what that concept is doing in political analysis has produced something this framework requires and should be taken as surpassing it in that domain. This architecture explains two distinct failure modes in predictive analysis. The first is domain isolation — knowing the inputs but treating them as independent rather than passing them through f_Cs as intertwined via $\oplus$, producing static risk inventories rather than dynamic trajectory models. The second is partial application — applying f_Cs to one domain (geopolitics, or economics, or ecology) while leaving the others as background context, producing sophisticated partial models that fail at exactly the points where cross-domain interactions drive the outcome. A tension internal to this section must be acknowledged before the formulas are used. The epistemological commitment of Part I is that frameworks must remain revisable, that prior commi

tments must not insulate themselves from evidence, and that the appearance of systematic rigor can serve as the most sophisticated form of self-insulation available. The formula architecture that follows is vulnerable to exactly this failure mode: variables expressed as letters, operators connecting them, and a notation that resembles mathematical precision can generate false confidence that the analysis is more determinate than it is. The formulas are heuristic tools for organizing observation, not algorithms for computing outcomes. They specify which variables matter and in what direction they push; they do not assign coefficients, they do not generate numbers, and they do not produce predictions that could be tested against a specific outcome at a specific time. An analyst who uses these formulas to confirm what they already believed has not applied them correctly — they have used the notation as a legitimization device rather than as an analytical discipline. The check against this misuse is identical to the check against any form of self-insulating analysis: does the formula application produce conclusions that are sometimes costly to the analyst, that require revision of prior assessments, that identify conditions the analyst would prefer not to see? If the formula consistently produces convenient findings, the formula is being used decoratively. The notation is anti-dogmatic only if the analyst applying it is.

15.1 — Complexity & Systems Dynamics: The Operator Decomposed

Complex Systems Dynamics is the operator that transforms three bodies of knowledge into predictive capacity. But calling it an operator leaves the question: what exactly does it do? The answer is not one thing but a set of interconnected principles, each of which is doing specific analytical work when applied to $Ns \oplus Hs \oplus Hd$. Understanding f_Cs requires understanding each principle and how they interact. The principles are not equally important in all situations. Some dominate in fast-moving systems, others in slow-moving

ones. Some are most visible at the threshold moment, others are what make threshold moments possible. What they share is that they are all invisible to additive analysis — none of them can be seen if the three inputs are treated as independent factors summed together.

CS.1 Non-Linearity: The Rejection of Proportional Thinking The most fundamental principle: in complex systems, outputs are not proportional to inputs. Doubling pressure does not double the response. Adding a small amount of stress to a system near its threshold can produce a response orders of magnitude larger than adding the same stress to a system far from threshold. The relationship between cause and effect is non-linear — which means that intuitions built from linear experience are systematically wrong about the magnitude of consequences. Non-linearity formula: $\Delta\text{Output} \neq k \times \Delta\text{Input}$ — where k is any constant. The relationship between input change and output change depends on the current state of the system, not on a fixed coefficient. The practical implication for Structural Tendency Analysis: identifying where in its response curve a system currently sits is more important than measuring the absolute values of the inputs. A system that is in the linear range of its response curve can absorb substantial increases in stress without qualitative change. The same system, pushed into the non-linear range, will respond to small additional stress with large qualitative change. Most predictive failure comes from applying linear extrapolation to systems that are operating in their non-linear range.

CS.2 Feedback Loops: How Systems Amplify and Stabilize Themselves Feedback occurs when a system's output becomes an input to its own future state. Two types are analytically fundamental and produce opposite dynamics. Positive feedback (amplifying): Output $\rightarrow$ amplifies $\rightarrow$ same direction of Output. A process that strengthens itself. Economic growth that produces investment that produces more growth. Legitimacy erosion that produces institutional failure that produces more legitimacy erosion. Positive feedback loops are the mechanism of both rapid ascent and rapid collapse — they are symmetric in direction, amplifying whatever tendency is already present. Negative feedback (stabilizing): Output $\rightarrow$ dampens $\rightarrow$ opposite direction of Output. A process that corrects itself. High prices that reduce demand that reduces prices. Political repression that produces resistance that produces nego

tiation that reduces repression. Negative feedback loops are the mechanism of equilibrium — they keep systems within ranges rather than allowing any variable to run to its extreme. Complex systems contain both types simultaneously. The question for Structural Tendency Analysis is not whether feedback loops exist — they always do — but which type is currently dominant in the system being analyzed. Systems dominated by negative feedback are in homeostasis: they absorb shocks and return to equilibrium. Systems in which positive feedback has overwhelmed negative feedback are in transition: they are amplifying in one direction until they reach a new equilibrium or collapse. Identifying the shift from negative-feedback dominance to positive-feedback dominance is one of the primary predictive tasks of f_Cs. Elite fragmentation illustrates the feedback structure. In a stable elite, negative feedback dominates: individual defection is costly, and the coalition's response to defection is to increase the cost further, keeping the coalition intact. As performance pressure increases and resource constraints tighten, the cost of defection falls below the cost of continued cooperation for increasing numbers of actors. At this point, positive feedback takes over: each defection reduces the coalition's capacity to punish future defection, making further defection cheaper, accelerating fragmentation. The transition from negative to positive feedback dominance is the threshold moment.

CS.3 Emergence: System-Level Properties Not Predictable from Components Emergence is the appearance of properties at the system level that are not present in and not predictable from any individual component. Water is wet; hydrogen and oxygen are not. Markets produce prices; no individual buyer or seller produces prices. Political ideologies produce social movements; no individual belief produces a social movement. Emergence formula: System Properties (Sp) $\neq \Sigma$ Component Properties (Cp) — the system is not the sum of its parts. For Structural Tendency Analysis, emergence means that analysis conducted at the component level — analyzing individuals, or individual institutions, or individual events — will systematically miss properties that only appear at higher levels of organization. The behavior of a crowd is not the average of its members' behaviors. The dynamics of an international system are not the sum of individual states' behaviors. The trajectory of a civilization is not the aggregate of its individuals' choices. Emergence also means that

interventions at the component level have unpredictable effects at the system level. A policy that changes individual incentives in predictable ways may produce emergent system-level consequences that no model of individual behavior would anticipate. This is not a failure of analysis — it is a structural feature of complex systems that any honest predictive framework must incorporate as a source of irreducible uncertainty. CS.4 Threshold Effect

s and Phase Transitions Thresholds are the points at which quantitative accumulation produces qualitative change. Water heated from 0°C to 99°C changes quantitatively — it gets hotter. Water heated from 99°C to 100°C changes qualitatively — it changes phase. The same quantity of additional heat produces a categorically different output depending on where the system is relative to threshold. Threshold formula: $f(x) = \text{State A for } x < \theta; f(x) = \text{State B for } x \geq \theta$ — where θ is the threshold value, not known in advance, and the transition between State A and State B is rapid relative to the timescale of accumulation leading up to it. Thresholds are the primary reason why prediction in complex social systems is structurally different from prediction in linear systems. In linear systems, tomorrow is an extrapolation of today. In threshold systems, tomorrow may be qualitatively different from today — not because something dramatically new was introduced, but because the accumulated pressure finally crossed the threshold that had been approached for years. The French Revolution did not happen because 1789 was unusually different from 1788. It happened because 1789 was the year accumulated pressure crossed threshold. The predictive challenge is that thresholds are not visible until they are crossed. What is visible are the leading indicators of threshold approach: accelerating positive feedback, degrading negative feedback capacity, increasing brittleness in the system's response to small shocks. A system that required large shocks to produce visible response a decade ago and now responds to small shocks with large disruption is approaching threshold — the threshold itself is not yet visible, but its approach is. CS.5 Path Dependence: History as Structural Constraint Path dependence means that what a system can do tomorrow is constrained by what it did yesterday — not through any physical necessity, but through the accumulation of commitments, adaptations, and complementary structures that make deviation from the established path increasingly costly over time. Path dependence formula: $\text{State}(t+1) = f(\text{State}(t), \text{Input}(t)) \neq f(\text{Input}(t))$ alone — the current state is an independent variable in determining the next state, not just a passive recipient of new inputs. Path dependence explains why identical inputs produce different outputs in different systems: because the systems are at different points in their historical trajectories and have accumulated different structural constraints. It also explains why the same system at different points in its history responds differently to the same shock: because the accumulated path has changed what options are available and what costs attach to each. For the STA (Structural Tendency Analysis) formula, path dependence is what makes the Historical Dynamics bracket essential. Natural Sciences and Human Sciences together describe the space of what is structurally possible. Historical Dynamics, multiplied through path dependence,

describes the much smaller space of what is actually available given what has already been chosen. Most of the options in the structural possibility space are foreclosed — not forever, but at this moment, given this accumulated path. CS.6 Sensitive Dependence: The Limits of Precision In some configurations, complex systems exhibit sensitive dependence on initial conditions: small differences in starting state produce large differences in trajectory over time. This is the precise technical content of what is popularly called 'butterfly effect' thinking. It is real, but it is domain-specific: not all systems exhibit sensitive dependence at all times, and identifying when a system is in a sensitive-dependence configuration is itself analytically tractable. Sensitive dependence condition: $|\text{Trajectory}(x + \epsilon) - \text{Trajectory}(x)|$ grows exponentially with time — where ϵ is a small initial difference. For Structural Tendency Analysis, sensitive dependence sets a hard limit on the precision of prediction: in configurations where it is active, long-range prediction of specific trajectories is not possible regardless of the quality of the analysis. What remains possible — and what the STA formula is designed to produce — is the identification of structural tendencies that hold across the range of sensitive trajectories. Even when specific outcomes cannot be predicted, the distribution of probable outcomes can be characterized. The system will do something in the space defined by its structural constraints, even if exactly

which thing is underdetermined. CS.7 Resilience and Brittleness: The System's Relationship to Shock Resilience is the capacity of a system to absorb disturbance and maintain its fundamental structure and function. Brittleness is its opposite: the property of systems that appear stable but break rapidly and irreversibly when stressed beyond a narrow range. The distinction matters because brittle systems and resilient systems look identical from the outside under normal conditions — both appear stable. They diverge only when stressed. Resilience formula: $\text{Resilience (Rl)} = \text{Diversity (Dv)} \times \text{Redundancy (Rd)} \times \text{Modularity (M)} / (1 + \text{Connectivity (Cn)})$ — where high connectivity that allows rapid propagation of shocks across the system reduces resilience even as it increases efficiency. The $(1 + \text{Cn})$ denominator reflects this: as connectivity grows, it increasingly suppresses the resilience that diversity, redundancy, and modularity provide. Note that very low connectivity is also a resilience problem — an isolated system cannot draw on external resources when internal capacity is exhausted. Optimal Cn is moderate: enough connectivity to enable mutual support across components, not so much that a single failure propagates everywhere. The paradox of optimization is the key insight here: systems optimized for efficiency under stable conditions sacrifice precisely the diversity, redundancy, and modularity that produce resilience. A highly optimized sy

stem is a brittle system. The same institutional reforms that improve performance under normal conditions — consolidating redundant functions, increasing coordination, eliminating inefficiencies — reduce the system's capacity to absorb shock. This is not an argument against optimization, but an argument for tracking the resilience cost of efficiency gains as a standing variable in any structural analysis. For Structural Tendency Analysis, tracking the resilience/brittleness trajectory of a system is one of the most reliable leading indicators available. A system whose resilience has been systematically declining through efficiency optimization is accumulating brittleness that will express itself catastrophically when an adequate shock arrives — even if the system appears stable up to that point. The shock does not cause the brittleness; it reveals it. CS.8 The Interaction of CS Principles: How the Operator Works as a Whole The seven principles above are not independent analytical tools applied sequentially. They interact: path dependence shapes which feedback loops are active; feedback loops determine how close the system is to threshold; threshold proximity determines whether sensitive dependence is active; brittleness determines how much stress the system can absorb before threshold is crossed; emergence means the system-level outcome of crossing threshold cannot be read from any of the component-level analyses. f_Cs (operator properties) = $\text{Non-linearity} \times \text{Feedback Dynamics} \times \text{Emergence} \times \text{Threshold Effects} \times \text{Path Dependence} \times \text{Sensitive Dependence} \times \text{Resilience/Brittleness}$ — where each principle modifies how the others operate, and the combined effect on $Ns \oplus Hs \oplus Hd$ is the transformation from descriptive knowledge into predictive capacity. f_Cs is most powerful — Structural Tendency Analysis is most reliable — when the system being analyzed is far from threshold, in a configuration where negative feedback dominates, where path dependence constrains the available options substantially, and where sensitive dependence is not yet active. These conditions allow structural trajectory to be identified with reasonable confidence over medium time horizons. f_Cs becomes less powerful — Structural Tendency Analysis can identify structural tendency but not specific trajectory — when the system is near threshold, when positive feedback is dominant, when sensitive dependence has been activated. In these configurations, the honest output of Structural Tendency Analysis is not a single predicted trajectory but a distribution of probable outcomes, with identification of the conditions that would shift probability toward each outcome in the distribution. This is the most important limit to name explicitly: f_Cs does not make prediction certain, and in some configurations it makes certain prediction impossible. What it always makes possible — regardless of the system's configuration — is the identification of which variables are most critical, which feedback loops

are most active, and which structural changes would most significantly alter the trajectory. That is enough. The goal of Structural Tendency Analysis is not omniscience but structural advantage: understanding the terrain well enough to act with better-than-random probability of achieving the intended outcome.

15.0b — Canonical Term Definitions

The following definitions are canonical. Where a term appears elsewhere in the document with different emphasis or elaboration, those passages specify the mechanism or context of the term's operation — they do not offer competing definitions. Cross-referencing should return here to resolve ambiguity. S — Structural Suffering: Suffering that is imposed (not chosen), preventable (not inherent to the human condition), structurally caused (produced by identifiable institutional arrangements), and dysfunctional (serving no developmental purpose for those who bear it). S is a qualitative classification, not a quantity. The formula $S = f(R, B, V)$ tracks structural tendency, not a measurable sum. R — Resource Deprivation: The gap between what a population requires to meet the minimal anchor (biological floor) and what the arrangement actually provides. Not absolute poverty — R is always relative to what the arrangement could provide and has implicitly or explicitly committed to provide. R is zero when the floor is met; it is non-zero when the arrangement produces floor deprivation that is structurally caused and preventable. B — Institutional Betrayal: The gap between the function an institution claims to perform and what it actually produces for the populations it claims to serve. An institution that has never promised anything cannot betray. It is the promise — explicit or implicit in the institution's stated mandate, legal obligation, or public legitimating narrative — that creates the betrayal when systematically broken. B is not corruption or bad intent; it is the structural outcome of capture dynamics that redirect institutional function without eliminating institutional claim-making. Note on scope: B presupposes a prior normative contract. Arrangements that produce exclusion by design from inception — without ever having extended a legitimating promise to the excluded group — are not captured by B and require the variable E (Absolute Structural Exclusion) developed in the extension to 3.4. See also the Ambedkar discussion in Intellectual Debts. V — Visibility Suppression: The degree to which suffering is prevented from becoming collectively legible — through information control, narrative management (attributing structural conditions to individual failure), isolation of those who suffer, or algorithmic architecture that suppresses structurally relevant claims. V operates at two levels. Primitive V is external suppression: censorship, data manipulation, physical silencing. Advanced V is internalised suppression: the condition identified by Bourdieu as symbolic violence, in which institutional arrangements successfully induce the populations bearing structural suffering to validate that suffering as natural, deserved, or individually caused. Advanced V requires no censor because the subject performs the suppression from within. The meritocracy narrative (8.4c) is the canonical mechanism of advanced V: it converts structural outcomes into apparent ev

idence of personal inadequacy, blocking the attribution of structural causation even when all material evidence of it is available. V is not an additive component of suffering; it is the conversion variable that determines how much of the suffering produced by R and B translates into organised transformation pressure. High V means high suffering, low pressure. Low V means suffering is visible and pressure accumulates at rates proportional to R and B. Becoming: The anti-foreclosure criterion: the structural availability of conditions under which individuals and communities can develop beyond what survival alone requires (4.5, 4.5c). Becoming is never synonymous with flourishing (the achieved developmental state), welfare (floor satisfaction), or positive affect (hedonic wellbeing). It refers to the conditions for development, not its content or achievement. An arrangement is evaluated on whether it expands or contracts the structural availability of becoming — not on whether it produces any particular form of development. Minimal Anchor: The evaluative foundation: reduce unnecessary suffering (structural, imposed, preventable, dysfunctional). The minimal anchor is a floor criterion only. It does not specify the ceiling, does not prescribe institutional form, and does not generate doctrine directly. It generates diagnostic questions that can be applied to any arrangement without prior ideological commitment. Layer / Level (usage convention): Layer refers to the constraint model verticality: Ecological Viability → Biological Floor → Economic Arrangements → Institutional Architecture → Meaning Infrastructure (model at 2.2; ecological layer at 3.10; meaning layer at 5.6). Each layer operates as a precondition for those above it. Level refers to the scale of analysis: micro (individual/household), meso (organisation/community), macro (system/state). These are

orthogonal dimensions. A meso-level analysis can examine the biological layer; a macro-level analysis can examine the meaning infrastructure layer. Conflating layer and level produces analytic confusion about whether a claim is about causal priority (layer) or analytical scope (level). W — Coalition Width: The breadth of the counter-power coalition measured by the range of structurally distinct groups whose material interests converge in opposing a given arrangement. W is not headcount but structural diversity: a coalition of ten organisations from identical structural positions scores lower W than a coalition of three organisations from positions with different material stakes and different channels of institutional access. High W increases the counter-power formula (15.8) because structurally diverse coalitions are harder to fragment through targeted concession and can apply pressure across multiple institutional leverage points simultaneously. W degrades when coalitions narrow to identity-homogeneous membership or when tactical success concentrates organisational resources in one faction. If —

Institutional Fragmentation (Ef): The degree to which the institutions sustaining a dominant arrangement are internally divided in interest, authority, or enforcement capacity. Ef is a structural variable of the target system, not of the counter-power coalition. High Ef indicates that elite cohesion has broken down: factions within the governing arrangement are competing for position in ways that create openings for counter-power that a unified elite would close. Ef interacts with Cf in the counter-power formula (15.8): high Ef creates structural opportunity, but high W is required to exploit it organisationally. Ef is one of the primary mechanisms through which transformation windows (10.5b) open: elite fragmentation typically precedes threshold crossing in historical cases where transformation pressure alone was insufficient to produce transition. E — Absolute Structural Exclusion: Structural suffering produced by arrangements that deny ontological standing to a group from inception, without ever having extended a normative promise of inclusion that could subsequently be violated. E is analytically distinct from B (Institutional Betrayal): B requires a prior normative contract whose breach generates the political charge; E operates where no such contract was ever offered. Caste systems, chattel slavery, and apartheid are paradigmatic E configurations — they do not betray; they exclude by constitutive design. This distinction matters for the formula architecture because E-produced suffering does not generate the same transformation pressure dynamics as B-produced suffering: populations subject to E have no legitimacy claim to invoke against the arrangement, which systematically forecloses the institutional channels through which B-generated pressure can accumulate. Counter-power under E conditions requires constructing the standing that B conditions assume as pre-existing. See 3.4 (Ambedkar note) and Intellectual Debts.

15.2 — The Primary Formula

All predictive analysis in this framework reduces to a master relationship: System Transformation Pressure (T) = $A = S \times D \times P \times G$ [where A is Accumulated Pressure as defined in 15.4; the shorthand $T = S \times D \times P$ in some cross-references omits G for legibility but G is always implied. Rc modifies the threshold that T must cross to produce non-linear reorganization rather than directly subtracting from T; the full trajectory formula at 15.9 integrates these: $Tr = T / (C \times L) - CP$, where Rc operates through the threshold variable rather than as an explicit term in Tr] Where transformation pressure is the net force pushing a system toward significant reorganization. When T exceeds the system's threshold for non-linear change ($T > T_{\text{threshold}}$) — a threshold that is itself variable and not known in advance — rapid reorganization becomes probable. When T remains below the threshold ($T < T_{\text{threshold}}$), the system possesses adaptive surplus: it can absorb shocks without structural change. Each variable in the primary formula is itself produced by secondary formulas. The primary formula is the output; what follows are the inputs. → Applied Cases: Section 15.14 demonstrates this formula applied to three historical cases — Tunisia 2011, Belarus 2020, and Argentina — showing how the full variable stack generates different structural predictions from different variable configurations. Several features of the formula require clarification before it can be used analytically rather than decoratively. The threshold for non-linear change is endogenous to the system being analyzed — it is

not a fixed value that can be read off from comparative cases but a variable that depends on the institutional configuration, the historical experience of the population, and the current distribution of organizational capacity. A system with high repressive capacity and low adaptive capacity may sustain positive T values for longer than a system with moderate repressive capacity and higher adaptive capacity — because the threshold in the first case is set by what the repressive apparatus can sustain, not by what the population can endure. This means that reading T as a predictive variable requires a prior analysis of the threshold, and that analysis is always conditional on the specific system. When transformation pressure remains persistently below the threshold ($T < T_{\text{threshold}}$) — when adaptive capacity, legitimacy, and repressive capacity together absorb or exceed the pressure generated by structural suffering — the system is in what might be called a maintenance configuration: not developing, not transforming, but reproducing itself with acceptable efficiency from the perspective of those whose position within it the configuration sustains. This is not stability in any evaluatively neutral sense. It is the stabilization of arrangements that are foreclosing becoming for the population they administer. A sub-threshold T value does not indicate that nothing important is happening;

it indicates that what is happening is the reproduction of conditions that the framework evaluates negatively, and that the reproduction is currently succeeding. Calibration across cases requires holding constant what 'structural suffering' means across different contexts — and this is where the formula is most demanding. A T value calculated for a high-income democracy using suffering metrics appropriate to that context is not comparable to a T value calculated for a lower-income authoritarian system using different metrics. The formula is not a cross-national scoring system. It is a tool for analyzing the internal dynamics of a specific system over time — for tracking whether transformation pressure is increasing or decreasing within a given configuration, and for identifying which variables are doing the most work in either direction. The applied cases in 15.14 demonstrate this use: each case analysis specifies its variables in context rather than applying standardized values across cases. Reading T in practice requires translating the formula's abstract variables into observable indicators. High transformation pressure is not directly measurable, but its components are — imperfectly, with lags, and subject to the visibility constraints that Part I identifies. Structural suffering manifests as indicators that can be tracked over time: health outcomes that diverge across structural positions, intergenerational immobility that persists or increases, access to the conditions for becoming (Part IV) that remains structurally determined rather than individually variable. Duration is the length of time those indicators have held or worsened — the accumulation that converts grievance into embedded political orientation rather than reactive frustration. Scale is the proportion of the population for whom these conditions apply directly, plus the secondary effects on populations adjacent to those most affected. Adaptive capacity and legitimacy are the variables most systematically misread in real-time analysis. Adaptive capacity appears high when elites are making concessions; it is high only when those concessions address the structural conditions generating pressure rather than the visible symptoms of it. Legitimacy appears robust when populations are not actively organizing against the system; it is eroding whenever the ratio of those who comply because they believe the system is legitimate to those who comply because the costs of non-compliance are high is shifting — a shift that is invisible until it crosses a behavioral threshold. Repressive capacity is visible but unreliable: it can suppress behavioral manifestations of pressure while the pressure continues to accumulate, producing a gap between observable stability and structural stress that makes the formula's predictions appear wrong until they suddenly appear right. The threshold problem is the formula's most significant practical limitation, and the framework does not pretend to resolve it. T

thresholds for non-linear change are not knowable in advance and are revealed only by crossing them. What the formula can identify is whether transformation pressure is accumulating and increasing toward the threshold, and whether the adaptive surplus or deficit is moving in a particular direction — which is enough for directional analysis and medium-term probability assessment but not enough for event prediction. The formula is a tool for identifying structural conditions, not a model for predicting specific outcomes from those conditions. The difference matters: analysis that mistakes directional probability for event prediction

will be systematically surprised by timing, and will incorrectly attribute that surprise to failures of the framework rather than to the inherent limits of what structural analysis can specify.

15.3 — The Suffering Variable: What Feeds Transformation Pressure

Structural Suffering (S) = Resource Deprivation (R) + Institutional Betrayal (B) + Visibility Suppression (V) [with V operating as a political-conversion modifier, not an additive component; see note below] Resource deprivation is the gap between what a population requires to meet the minimal anchor and what the system actually provides. It is not absolute poverty — populations can sustain extraordinary material hardship without generating transformation pressure if the deprivation is understood as natural, shared, or temporary. It generates pressure when understood as structural, differential, and produced by identifiable arrangements. The three-variable formula here captures the political-dynamic dimensions of suffering — what generates transformation pressure. It is complementary to, not a replacement for, the four-way typology in 3.4b (chosen/imposed, functional/dysfunctional, preventable/irreducible, structural/incidental), which captures the evaluative dimensions. The typology identifies what the framework opposes. This formula tracks what feeds systemic change. Institutional betrayal is the gap between the function an institution claims to perform and what it actually produces. An institution that has never promised anything cannot betray. It is the promise — explicit or implicit — that creates the betrayal when it is systematically broken. This is why the institutions that generate the most intense transformation pressure are often the ones most loudly committed to justice, welfare, or accountability: the gap between claim and performance is largest there. Visibility suppression is the degree to which suffering is prevented from becoming collectively legible. Suffering that is individualized — experienced privately, attributed to personal failure, without the language or infrastructure to be recognized as shared — does not feed transformation pressure at the same rate as suffering that is collectively visible. Information control, narrative management, and the isolation of those who suffer all function to keep S below the threshold at which it generates pressure. Note on S formula defensibility: V operates differently from R and B in the formula. R and B produce suffering directly; V determines how much of that suffering converts into organized transformation pressure. A more precise rendering would be $S_{\text{political}} = (R + B) \times (1 - V)$, where V suppresses political translation rather than adding to it. The additive notation is retained for legibility and to preserve the diagnostic value of tracking V as an independent variable, but the suppressive mechanism must be held in mind during application. Additionally: the three-variable formula under-specifies the full structure of structural suffering. Fear and humiliation (3.8b) generate transformation pressure independently of R/B/V; relative deprivation (gap between conditions and expectations) can exceed absolute deprivation as a pressure driver; and epistemic infrastructure degradation (1.8–

1.12) functions as a V-suppression mechanism not fully captured by the single V variable. The formula tracks the primary variables. It does not claim to be exhaustive. The formula implies a diagnostic: when visibility suppression increases — when narratives of individual failure intensify, when the language for structural suffering is contested, when media infrastructure favoring isolated accounts of hardship expands — this is often evidence that institutional betrayal and resource deprivation are increasing simultaneously. Suppression is a signal of what it is suppressing. Resource deprivation (R) — the floor gap: Resource deprivation is not measured by absolute poverty but by the gap between what a population requires to meet the minimal anchor and what the existing arrangement actually delivers. This distinction matters because populations can sustain severe material hardship without generating durable transformation pressure when the deprivation is understood as natural, temporary, or individually caused — when the governing narrative successfully attributes the gap to conditions outside the system's responsibility. What generates transformation pressure is not deprivation per se but the combination of deprivation and the attribution of that deprivation to the system's structural choices. R is therefore the interaction of the material gap and the narrative processing of that gap: high material gap with effective V suppression generates lower effective R

than equivalent material gap with low V suppression. This is why the formula treats R as an input to S rather than as equivalent to S — the material reality of deprivation does not automatically translate into political pressure without the narrative infrastructure that makes the structural causation visible. Institutional betrayal (B) — the violation of stated commitments: Institutional betrayal is suffered differently from resource deprivation because it operates against a background of explicit promises. A system that never claimed to provide specific protections cannot betray the populations it fails to protect. A system that explicitly committed to those protections and failed to deliver them — welfare states that reduce benefits during fiscal crises, legal systems that selectively enforce rights, political parties that campaign on specific program commitments and govern differently — produces betrayal suffering that has higher political salience per unit of material cost than equivalent deprivation from a system that made no comparable commitments. The mechanism is the gap between stated principle and actual practice: the gap makes the failure visible in a specific way, attributable to specific actors, and experienced as a violation of a legitimate expectation rather than as an unfortunate natural condition. B is therefore particularly important for explaining why legitimacy crises develop most rapidly in systems that have the highest stated commitments — the distance between promise and

performance is where the political charge accumulates. Visibility suppression (V) — the mechanism of conversion: Visibility suppression is the variable that determines how much of the suffering produced by R and B converts into organized transformation pressure. V is not the only mediating factor — counter-power capacity also determines conversion efficiency — but it operates at an earlier stage: V determines whether the suffering is understood as structural and remediable at all, while counter-power capacity determines whether that understanding can be organized into action. V operates through multiple channels: media ownership and editorial selection determine what suffering becomes politically visible nationally; professional norms in academic and policy analysis determine what causal attributions are treated as legitimate; governing narratives that attribute suffering to individual failure or natural conditions suppress visibility without directly managing information; and the physical concentration of suffering in spaces that are invisible to those with political influence reduces the political salience of what is being produced. V is multiplicative in its effect: high V does not reduce R or B but can reduce the effective political pressure generated by both toward zero. The interaction between the three variables is not additive in a simple sense. A system can produce high R while maintaining low effective political pressure through high V suppression. A system with lower R can produce higher transformation pressure if V is low and B is high — the betrayal makes the causation visible in ways that resource deprivation alone does not. Analysts who track only material deprivation and miss the visibility and betrayal dimensions consistently misread the political dynamics of systems that maintain stability despite significant structural suffering through effective narrative management, and equally misread the dynamics of legitimacy crises that develop faster than the material deprivation would predict because V is declining and B is high. Measurement orientation: R is approximated through material indicators (wage share, access to healthcare, housing cost as income fraction, educational opportunity distribution) calibrated to what the minimal anchor requires, not to international comparisons that normalize deprivation. B is approximated through the gap between stated institutional commitments and actual institutional performance, tracked through specific program delivery, enforcement consistency, and stated-vs-actual policy priorities. V is the most difficult to measure directly; proxy indicators include media pluralism indices, academic citation patterns across ideological positions, and the distribution of structural causal explanations in mainstream political discourse. Connection to adaptive capacity (15.5) and fragmentation (15.6): chronic existential suffering at scale — the Last Man condition described in 3.3 and 4.5b, where material need

is met but becoming is systematically foreclosed — reduces long-term adaptive capacity (C) through the depletion of cognitive surplus and institutional trust that 4.5b identifies as the material mechanism of becoming. Simultaneously, it increases internal fragmentation (Cf) through the affective and political dynamics that 3.8b analyzes: populations experiencing developmental foreclosure at scale are structurally

susceptible to mobilization around identity cleavages that substitute for genuine structural analysis. An arrangement that produces low R and B but high existential suffering is not thereby evaluated positively by the formula architecture: its effect on C and If are the channels through which existential suffering enters the structural tendency analysis, even when V is also low and the suffering is individually attributed rather than structurally visible.

15.4 — The Duration and Scale Multipliers

Accumulated Pressure (A) = $S \times D \times P \times G$ Duration multiplies suffering's political weight. Suffering that persists across years generates different political consequences than equivalent suffering concentrated in a brief crisis. Populations can absorb acute crises more readily than chronic conditions, because acute crises do not alter the behavioral dispositions, life expectations, and institutional trust that chronic suffering degrades. Scale is the proportion of a relevant population experiencing the suffering. Below certain thresholds, suffering is politically managed as a minority problem. Above certain thresholds, it becomes the majority's problem, and the political arithmetic changes fundamentally. The threshold varies by political system, cultural context, and the degree to which affected populations can form coalitions — but the structural relationship is consistent: scale amplifies pressure non-linearly as it crosses relevant political thresholds. Intergenerational transmission is the degree to which suffering passes from parents to children through structural mechanisms rather than individual failure. When a generation's suffering is transmitted structurally — through constrained educational access, degraded neighborhood infrastructure, asset depletion, intergenerational trauma — it compounds. The children of the suffering population begin their lives already behind the baseline, generating fresh pressure even without new deprivation. The interaction between duration, scale, and intergenerational transmission is non-additive. A suffering condition that scores high on all three does not simply produce three times the political pressure of one that scores high on only one. Duration extends the population for whom the suffering condition has become a baseline expectation rather than a departure from prior conditions — which changes the political psychology of response. Scale increases the proportion of the population for whom suffering is direct rather than observed — which changes the arithmetic of coalition formation. Intergenerational transmission forecloses the individual exit strategies that allow people to experience their suffering as temporary — which removes the absorptive function that the prospect of individual mobility otherwise provides. The diagnostic implication of each variable points to different intervention logics. High duration with low scale indicates a politically isolated condition — chronic but manageable by existing arrangements through minority management. High scale with low duration indicates an acute crisis — dangerous in the short term but potentially absorbable if adaptive capacity is sufficient. High intergenerational transmission is the variable most likely to be systematically underweighted in both political analysis and policy design: it produces pressure that is structurally invisible in the short run and politically explosive in the medium run. The generation that inherits transmitted suffering did not

experience the conditions that produced it, but experiences its consequences as the normal terms of their lives. This generates a distinctive political psychology — structural grievance without a clear personal narrative of loss — which makes it both harder to organize and more volatile when it finally does organize around a focal event or leadership. The G variable (intergenerational transmission) is analytically distinct from duration even though both extend suffering across time. Duration measures how long a population experiences a condition; G measures whether the structural mechanisms that produced the condition operate on the next generation independently of what the current generation does or achieves. A high-duration, low-G condition can in principle be resolved by removing the immediate cause: the population recovers once conditions change. A high-G condition has reproductive mechanisms that operate independently of the original cause — constrained asset formation, degraded infrastructure, behavioral adaptations to scarcity that persist after the scarcity itself ends — and requires intervention at the level of the reproductive mechanism,

not only at the original source. Duration (D) — how chronic suffering differs from acute crisis: Duration multiplies suffering's political weight through a specific causal mechanism: chronic conditions alter the behavioral dispositions, life expectations, and institutional trust of the populations that bear them in ways that acute crises do not. A population that experiences a severe acute crisis — a natural disaster, a financial shock, a political rupture — and then returns to baseline conditions will typically reconstitute its prior behavioral patterns relatively quickly, because the underlying expectations and trust infrastructure remain intact. A population that experiences sustained material deprivation over years or decades undergoes a different kind of change: its expectations of what institutions will deliver are revised downward, its trust in official explanations of why conditions are as they are is eroded, and the intergenerational transmission of structural position is intensified as parental deprivation reduces children's developmental opportunities. These changes increase transformation pressure because they reduce the conversion cost of deprivation into organized action — populations that have already reduced their expectations of the system have less to lose by acting against it, and populations with eroded trust in official explanations are more receptive to counter-hegemonic narratives that attribute their situation to structural causation. Geographic spread (P) — why scale changes political dynamics: The population scale variable P reflects a non-linear relationship between the number of people bearing structural suffering and the political consequences of that suffering. Small concentrated populations bearing high levels of structural suffering can be managed through targeted repress

ion and narrative isolation — their suffering can be attributed to specific characteristics of the affected population that mainstream political actors accept as legitimate grounds for differential treatment. Large dispersed populations bearing structural suffering at lower intensity can be managed through fragmentation — dividing the suffering into distinct political subgroups that identify their interests as separate rather than shared, preventing the coalition formation that transformation pressure requires. The transition between these manageable configurations and configurations that generate durable transformation pressure occurs when the suffering is both widespread enough to affect populations whose political characteristics vary sufficiently to prevent fragmentation and visible enough to resist isolation. P is therefore not a simple count of affected people but a measure of how the scale and distribution of suffering affects the feasibility of the management strategies through which governing coalitions contain it. Cross-group alignment (Ga) — the coalition dimension: Ga captures what is perhaps the most politically consequential feature of transformation pressure: whether the suffering being generated is falling on populations that share enough political characteristics to form coherent coalitions. Suffering that is structurally produced by the same mechanisms but distributed across populations that are separated by identity, geography, and material interest rarely generates unified transformation pressure even at high levels of S and D, because each affected group tends to attribute its situation to factors specific to its own position rather than to the shared structural mechanisms. The governing coalition's primary interest-maintenance strategy is to ensure that G remains low — that the populations bearing structural suffering understand themselves as having separate rather than shared interests. This is achieved through a combination of differential treatment (creating genuine material differences in the conditions facing different groups that make shared interests harder to identify), narrative fragmentation (emphasizing the characteristics that differentiate groups rather than the mechanisms that connect their situations), and strategic allocation of resources to potential coalition partners that could bridge the divide. The formula $A = S \times D \times P \times G$ specifies that accumulated pressure requires all four variables to be meaningfully non-zero. High S with low D produces manageable acute crises rather than accumulated pressure. High D with low P produces chronic suffering that can be managed through targeted interventions. High S and D with low G produces extensive suffering that fails to organize into the cross-group coalitions that durable transformation requires. The analytical implication: reading transformation pressure requires tracking all four dimensions, not only the intensity of current suffering.

15.5 — The Adaptive Capacity Formula

Adaptive Capacity (C) = (Institutional Flexibility (F) × Elite Cohesion (Ec) × Reform Infrastructure (I)) × (1 – Elite Capture Rate (K))

Institutional flexibility is the degree to which existing institutions can incorporate demands for change without requiring their fundamental reorganization. High flexibility means the system can absorb pressure by adjusting outputs without changing the structural relationships that produce those outputs. Low flexibility means every demand for change confronts the full institutional inertia — the system cannot give at the margins, so pressure accumulates toward threshold. Elite cohesion matters because divided elites create openings that unified elites foreclose. When the actor class that benefits most from existing arrangements is internally fractured — by competing material interests, ideological disagreement, or generational difference — challengers can exploit the fractures. When elite cohesion is high, the adaptive capacity of the system increases because the actors most capable of blocking change are coordinated. Reform infrastructure is the institutional capacity for managing change: the existence of mechanisms through which demands can be channeled, tested, and partially accommodated without requiring the dismantling of core arrangements. Where this infrastructure has been systematically degraded — through the defunding of mediating institutions, the elimination of formal channels for negotiation, or the concentration of decision-making in actors with no accountability to those making demands — adaptive capacity drops sharply. Elite Capture Rate (K) is the rate at which actors who enter reform infrastructure are converted to serve existing arrangements rather than the constituents who sent them. High capture rates hollow out reform infrastructure: its formal existence remains, but its functional capacity for absorbing pressure is eliminated. Capture is the mechanism through which systems maintain the appearance of adaptive capacity while losing its substance.

Institutional flexibility (F): Institutional flexibility is not organizational agility in the managerial sense — the capacity to implement decisions quickly. It is the structural property of being able to incorporate demands for change through procedural adjustment without requiring reorganization of the fundamental relationships through which the institution exercises power and distributes benefits. High flexibility means that the institution can respond to transformation pressure by adjusting outputs — providing more services, changing eligibility criteria, modifying regulatory standards — without changing the structural relationships that determine who controls what and who benefits from that control. Low flexibility means that any significant accommodation of transformation pressure requires those structural relationships to change, which means that the actors who benefit from them have interests in resisting accommodation rather than managing it.

t. The politically significant implication: institutions with high F can absorb substantial transformation pressure without producing threshold conditions, while institutions with low F encounter threshold dynamics at lower pressure levels because accommodation is structurally difficult.

Elite cohesion (Ec): Elite cohesion is the degree to which the governing coalition is unified in its assessment of how to respond to transformation pressure — whether through accommodation, managed reform, or repression. Divided elites reduce adaptive capacity in a specific way: when elite factions disagree about how to respond to pressure, they signal to actors generating that pressure that the coalition is weakened, which reduces the deterrent effect of repressive capacity and can accelerate threshold dynamics. Elite division also makes negotiated accommodation more difficult, because there is no single addressee who can make commitments that the coalition will honor. Conversely, high elite cohesion increases the capacity to manage transformation pressure through coordinated responses, whether those responses are accommodative or repressive. The formula treats Ec as a multiplicative factor rather than an additive one because the effects of elite division on adaptive capacity are not proportional — even modest elite division significantly reduces the coalition's capacity to execute any response strategy consistently.

Reform infrastructure (I): Reform infrastructure is the set of institutional capacities that allow a system to implement change once the political decision to change has been made: technical expertise, administrative capacity, legal frameworks for implementing new arrangements, and the organizational capacity to manage transitions without catastrophic disruption. Low reform infrastructure means that even when governing coalitions decide to accommodate transformation pressure through substantive change, they cannot implement that change effectively. This creates a specific failure mode:

systems with high political will to reform but low reform infrastructure undergo apparent reforms that fail in implementation, which produces legitimacy damage from the betrayal of stated reform commitments on top of the original pressure that prompted the reform attempt. High reform infrastructure is a necessary condition for accommodation strategies to work; it is not a sufficient condition because accommodation also requires the political will to use the infrastructure for genuine change rather than cosmetic adjustment. Elite capture rate (K) — the modifier: The elite capture rate K appears in the formula as a discount on the product of F, Ec, and I because capture operates precisely by converting the adaptive capacity components into instruments of the governing coalition's self-maintenance rather than of genuine accommodation. An institution with high flexibility, high elite cohesion, and high reform infrastructure that is fully captured will use all three to man

age transformation pressure in ways that maintain the structural relationships producing the suffering rather than changing them. The $1 - K$ term captures this: at $K = 0$, full adaptive capacity is available for genuine accommodation; at $K = 1$, all adaptive capacity is redirected toward sophisticated resistance. The empirical implication: measuring adaptive capacity without measuring capture rate will systematically overestimate the system's actual capacity to accommodate transformation pressure, which is the analytical error that produces confident predictions of managed reform in systems that are in fact approaching threshold through apparent stability. Diagnostic for adaptive capacity: (1) What is the current F — can the institution make meaningful output adjustments without changing control relationships? (2) Is the elite coalition unified on how to respond to current pressure? (3) What is the reform infrastructure capacity — can the system implement what it decides? (4) At what rate is capturing occurring — are the flexibility, cohesion, and infrastructure being deployed for genuine accommodation or for sophisticated resistance? Low C with high T is the threshold proximity signature.

15.6 — The Repressive Capacity Formula

Repressive Capacity (R_c) = (Coercive Resources (C_v) $\times$ Enforcement Loyalty (E_l) $\times$ Information Control (I_c)) $\times$ (1 – Internal Opposition to Repression (I_o)) Repressive Capacity is the system's active instrument for suppressing transformation pressure. It operates not by reducing T directly but by raising the threshold that T must cross to produce non-linear reorganization: a system with high R_c can sustain high accumulated pressure without threshold crossing because the cost of defection from compliance remains prohibitive. This is the structural explanation for cases like Belarus 2020 or China 1989, where accumulated pressure was real and visible but did not produce transition. R_c does not make a system immune to transformation; it raises the threshold that pressure must cross and determines how much pressure must accumulate before the trajectory formula's denominator — $C \times L$, adaptive capacity times legitimacy — is effectively zero regardless of R_c 's value. Coercive Resources (C_v) are the material instruments of enforcement: military and police capacity, surveillance infrastructure, financial resources for sustaining them. They are necessary but not sufficient — without loyalty from those who operate them, coercive resources are inert. Enforcement Loyalty (E_l) is the degree to which enforcement personnel will carry out orders against the population. It is the variable most sensitive to structural suffering within the enforcement apparatus itself — when those who enforce repression share material conditions with those they are repressing, loyalty is structurally fragile regardless of formal command structures. Military defection is the most consequential form of R_c collapse. Information Control (I_c) is the capacity to prevent collective legibility of suffering and of opposition capacity — to keep transformation pressure from becoming visible to those who experience it. When information control fails, the visibility suppression component of Structural Suffering (15.4) reverses simultaneously, creating a double amplifier: both T rises and R_c falls. Internal Opposition (I_o) is resistance to repression from within the apparatus itself — officers who refuse orders, officials who leak, factions within the security structure that calculate the regime's survival odds and defect accordingly. Internal opposition is what collapses R_c most rapidly: external pressure against a unified repressive apparatus is slow; internal fracture of that apparatus is fast. Coercive resources (C_v) — what repression requires: Coercive resources are the material

inputs to repressive capacity: the physical infrastructure of security forces, the budget that maintains and equips them, the legal frameworks that authorize their deployment, and the intelligence infrastructure that identifies targets. Coercive resources are necessary but not sufficient for effective repression — they can be present at high levels and still fail to suppress transformation pressure if enforcement loyalty

is low or information control is absent. The analytical significance of Cv as a component of Rc rather than as equivalent to it is that systems which invest heavily in coercive resource accumulation often do so precisely because their other repressive capacity components are degrading: high Cv with low El signals a system that is compensating for loyalty erosion through resource expansion, which is a trajectory associated with increasing fragility rather than increasing stability. Enforcement loyalty (El) — the loyalty problem: Enforcement loyalty is the degree to which coercive personnel will execute orders against the populations they are directed to repress. It is the most structurally significant component of repressive capacity because coercive resources cannot be deployed without it, and because its degradation is a leading indicator of threshold proximity that is more sensitive than most other indicators. Enforcement loyalty declines along three pathways: material pathway, when the material conditions of security force personnel converge toward those of the populations they are directed to repress, reducing the material stake in maintaining the arrangement; identity pathway, when security force personnel share ethnic, regional, or social characteristics with the populations being repressed, making enforcement psychologically costly; and institutional pathway, when command structures within security forces fracture along factional lines that reflect broader elite division, reducing the reliable execution of central commands. The political significance: a government that has lost El cannot compensate by acquiring more Cv, because the coercive resources will not be deployed when directed. El is the binding constraint on Rc once it falls below a critical threshold. Information control (Ic) — the epistemic dimension of repression: Information control is repressive capacity exercised at the epistemic level rather than the physical level: the ability to shape what populations know about the exercise of repression, about the strength of counter-power, and about the distribution of compliance and resistance. High Ic means that repressive events can be managed as information events — the extent of repression can be obscured, the organization of resistance can be fragmented by preventing coordination, and the impression of widespread compliance can be maintained even when actual compliance is declining. Low Ic means that repressive events become immediately visible across the population, that resistance can coordinate rapidly through information networks outside governing coalition control, and that the impression of widespread compliance cannot be maintained once individual defection begins. The political significance of Ic has increased substantially with digital communication infrastructure: the cost of maintaining high Ic has risen because the infrastructure through which information suppression is organized is distributed across networks that a

re more difficult to control than broadcast media. Internal opposition to repression (Io) as modifier: Io captures elite opposition within the governing coalition to the use of repressive capacity. Even when Cv, El, and Ic are all high, repressive capacity can be reduced by internal opposition from coalition members who calculate that the costs of repression — international, reputational, economic — exceed its benefits for their specific position. Io functions as a discount on the product of the other three components, reducing effective Rc below what raw coercive capacity would suggest. High Io signals elite fracture at precisely the moment when repressive capacity is most needed to contain transformation pressure — which is one of the structural reasons why threshold crossings often occur despite apparently high coercive resources. Diagnostic sequence for Rc assessment: (1) What are the material and organizational coercive resources — are they being expanded, maintained, or degraded? (2) What is El — are there indicators of loyalty erosion along any of the three pathways (material, identity, institutional)? (3) What is Ic — is information control strengthening or declining as a function of media pluralism, digital network penetration, and international reporting? (4) Is Io rising — are there elite voices within the coalition explicitly or implicitly opposing further repressive deployment? A system with high Cv but declining El and Ic is approaching repressive incapacity faster than its resource base would suggest.

15.7 — The Legitimacy Formula

System Legitimacy (L) = Performance Legitimacy (L_p) + Procedural Legitimacy (L_r) + Narrative Legitimacy (L_n) – Legitimacy Deficit Accumulation (L_d)

Performance legitimacy is the degree to which the system produces outcomes that its subjects experience as beneficial. It is the most fragile form of legitimacy because it requires continuous delivery and is immediately undermined by visible failure. It is also the most powerful when intact: populations will tolerate significant procedural irregularity and narrative incoherence if the system reliably produces the goods it promises. Procedural legitimacy is the degree to which the system's processes for making decisions are experienced as fair, regardless of specific outcomes. It is more durable than performance legitimacy because it does not require continuous positive outcomes — it requires only that the processes through which outcomes are determined be experienced as legitimate. But procedural legitimacy without performance legitimacy eventually collapses as populations conclude that fair processes producing consistently bad outcomes indicate the process itself is corrupted. Narrative legitimacy is the degree to which the governing story about why the system is right, natural, or necessary is believed. It is the most durable form of legitimacy under normal conditions and the most catastrophically vulnerable under specific ones: once a governing narrative loses credibility, it cannot be restored simply by reasserting it with greater force (6.7). The collapse of narrative legitimacy is often the leading indicator of system crisis — it precedes the visible manifestations of transformation pressure by years or decades. Legitimacy Deficit Accumulation (L_d) is the rate at which each failed delivery, procedural violation, or narrative contradiction accumulates into a structural deficit that is harder to reverse than to prevent. Legitimacy deficits are asymmetric: they accumulate faster than they are repaired. A single dramatic failure can undo years of performance; restoring the credibility lost in that failure requires not equivalent performance but sustained performance over extended time under scrutiny. This asymmetry means that systems nearing a legitimacy threshold cannot recover through the same actions that would have maintained their position. Performance legitimacy (L_p) — the most fragile form: Performance legitimacy is continuously produced and continuously at risk because it requires ongoing delivery of outcomes that subjects experience as beneficial. Its advantage is that it is the most politically powerful form when intact: populations that experience the system as producing conditions they value — economic security, physical safety, social stability — are resistant to counter-hegemonic narratives that attribute those conditions to structural arrangements rather than to the system's management. Its vulnerability is symmetrical: the same immediacy that makes performance legitimacy powerful whe

n intact makes it rapidly reversible when delivery fails. A system that has built legitimacy primarily on performance has no reserve when performance deteriorates — the subjects who valued the system for its outputs have no other basis for continued support. The distinction between the three forms of legitimacy matters for how rapidly legitimacy deficits accumulate: performance failures produce rapid L_d accumulation because performance expectations are continuous; procedural failures are slower because procedures are less continuously salient; narrative failures are slowest but most difficult to reverse once they occur because they require changing the fundamental attribution framework through which performance and procedure are interpreted. Procedural legitimacy (L_r) — the institutional buffer: Procedural legitimacy is the degree to which subjects accept the system's decision-making processes as fair regardless of specific outcome. Its political function is to serve as a buffer for performance failures: populations that believe outcomes were produced through legitimate procedures are more likely to accept bad outcomes as the result of genuinely difficult tradeoffs rather than as evidence of systemic failure or capture. The institutional significance: procedural legitimacy can maintain overall system legitimacy through periods of performance failure that would otherwise produce rapid legitimacy depletion. Its vulnerability is the consistency requirement — procedural legitimacy depends on the procedures being applied consistently across populations and cases, and selective enforcement of procedure (applying rules strictly to some actors and loosely to others based on political position or social characteristics) is perceived as procedural failure in ways that undermine the

buffer function. Selective enforcement produces Lr decline that is often politically underestimated because the populations most directly affected are those with least political voice, but that accumulates into broad legitimacy deficits when the pattern becomes visible across enough cases. Narrative legitimacy (Ln) — the meaning framework: Narrative legitimacy is the degree to which the system's governing story — its account of what it is, what it is for, and why the arrangements it maintains are appropriate — is accepted as the legitimate framework for interpreting political events. High Ln means that even performance failures and procedural violations are interpreted through a framework that places them in a context that limits their delegitimizing effect: the failures are exceptions, the violations are mistakes, the system's fundamental appropriateness is not in question. Low Ln means that performance successes and procedural integrity are interpreted through a counter-framework that attributes them to specific interests rather than to the system's general function — success is evidence of who benefits, and procedural integrity is the form through which substantive inequality

is maintained. Narrative legitimacy is the slowest component to build and the slowest to erode, but once eroded it is the most difficult to restore because restoration requires changing the interpretive framework rather than delivering specific outcomes or following specific procedures. Legitimacy deficit accumulation (Ld) — how legitimacy degrades: Ld captures the accumulation of legitimacy damage over time through repeated failures that are not offset by legitimacy-building responses. The additive-versus-multiplicative structure of the formula reflects an important asymmetry: L components are added (different forms of legitimacy can partially compensate for each other), while Ld is subtracted from the total (accumulated deficit directly reduces overall legitimacy regardless of which component it primarily affects). This means that a system with strong narrative legitimacy can absorb performance failures without Ld accumulation reaching critical levels, and a system with high performance legitimacy can absorb procedural failures. But accumulated Ld eventually overwhelms all three components simultaneously when it has accumulated sufficiently — this is the legitimacy crisis that 6.6 analyzes: the point at which no single form of legitimacy restoration can offset the accumulated deficit because the deficit has become the primary framework through which all system outputs are interpreted. Diagnostic for legitimacy trajectory: (1) Which component is currently primary — is the system's acceptance by subjects based predominantly on Lp, Lr, or Ln? (2) What is the trajectory of Ld — is it stable, accelerating, or decelerating? (3) Are there legitimacy-building responses that are actually reaching the populations whose Ld contribution is highest, or are they reaching populations already high on Ln? (4) Is there a legitimacy reserve — do the components provide mutual support, or has the system's legitimacy become single-component dependent in ways that make it brittle to failures in that component?

15.8 — The Counter-Power Formula

Counter-Power Capacity (CP) = (Organizational Infrastructure (O) × Narrative Infrastructure (N) × Resource Base (Rs) × Coalition Width (W)) × (1 – Repression Rate (Rp) – Internal Fragmentation (Cf)) These variables compound rather than add. An organization with strong infrastructure but no narrative frame for why the struggle matters cannot recruit beyond those already convinced. A movement with compelling narrative but no organizational infrastructure cannot convert sympathy into sustained pressure. Counter-power capacity requires sufficiency across all variables simultaneously — a weakness in any one constrains the ceiling of all others. Coalition width is particularly underweighted in most analyses of political change. Movements that remain confined to those with the greatest direct material stake in change are predictably outmatched by the organizational capacity of those defending existing arrangements. Every extension of the coalition into populations with secondary rather than primary material stakes in the outcome increases the resource base, the narrative reach, and the political cost of repression — because repression that touches only the most marginalized is affordable; repression that touches broad coalitions is not. Internal fragmentation (Cf) is the rate at which counter-power capacity is consumed by conflicts within the coalition

rather than applied against the target. It is produced by identity conflicts, strategic disagreements, organizational competition for scarce resources, and infiltration. The formula implies a counter-intuitive diagnostic: when fragmentation increases sharply within a movement, the appropriate question is not only what internal dynamics produced it, but what external conditions or actors benefit from the fragmentation and may be deliberately promoting it. Counter-power capacity and repressive capacity operate on different timescales, and this asymmetry is one of the most structurally consequential features of the political terrain most movements face. Repressive capacity can be deployed rapidly: a decision to use coercive resources produces consequences within hours. Counter-power capacity is built slowly. Organizational infrastructure, narrative reach, and coalition width each require sustained investment that cannot be compressed without trade-offs in durability and coherence. The implication is that movements which attempt to compress the building period in order to meet an immediately available political window typically enter that window with counter-power capacity sufficient in scale but fragile in structure — which makes them vulnerable to repressive pressure at exactly the moment they are most visible and most exposed. The cost-asymmetry between repression and resistance adds a precision to the R_c variable that the formula does not fully surface in its standard presentation. The effective repressive capacity faced by a movement is not simply the raw coercive resources of the state — it is the cost-

adjusted capacity, which accounts for the political and economic costs the state incurs by deploying repression against a coalition of a given width and composition. Repression that touches only the most marginalized is operationally cheap; repression that touches a coalition including actors with political leverage, international visibility, or economic significance incurs costs that reduce the effective value of R_c . Coalition width (W) therefore functions not only as a resource multiplier for counter-power capacity but as a cost-imposing mechanism on the repressive side of the equation — which is why it is doubly valuable and why deliberate fragmentation of coalitions is a standard instrument of counter-movement strategy. Sequencing implication: organizational infrastructure and narrative infrastructure must be built before coalition width is expanded, not after. A movement that expands its coalition before organizational infrastructure can absorb new members loses the narrative coherence and strategic coordination that makes the coalition politically effective. The sequence that produces durable counter-power capacity is: build organizational infrastructure sufficient to maintain coordination under pressure — develop narrative infrastructure capable of sustaining commitment beyond initial mobilization — then expand coalition width to convert that capacity into political scale. Reversing this sequence produces the appearance of strength without the structural durability to convert visibility into outcomes. The interaction terms — why multiplication rather than addition: The formula multiplies its primary variables rather than adding them because the political effectiveness of counter-power requires all four components to be meaningfully present simultaneously. Organizational infrastructure without narrative infrastructure produces organizations that can coordinate action but cannot explain why it matters — which limits recruitment to those already convinced and makes the organization brittle to the attrition that sustained struggle requires. Narrative infrastructure without organizational infrastructure produces compelling arguments that produce no organized action — the problem of movements that generate cultural resonance without structural power. Resource base without the other three produces co-optation: resources attract the institutional attention that redirects them from transformation to organizational maintenance. Coalition width without the other three produces fragile alliances that disaggregate under pressure from the governing coalition's divide strategy. The repression and fragmentation discounts operate multiplicatively on the total rather than subtracting from individual components because both work by reducing the effective combination of the components rather than only degrading specific ones. Repression targets the organizational infrastructure directly, but its political effect flows through all components: it reduces coalit

ion width by raising the cost of affiliation, degrades resource base by imposing legal and financial constraints, and attacks narrative infrastructure by discrediting the movement's account of itself. Internal fragmentation operates similarly: it does not only reduce coalition width but disrupts the organizational

infrastructure through which the coalition acts, degrades the narrative coherence that coalition members depend on for shared interpretation, and redirects resources from external to internal conflict. Strategic implication for counter-power building: because the formula is multiplicative, the highest-leverage investment is in whichever component is currently lowest — the component at near-zero makes all others near-zero regardless of their individual levels. Before expanding any component, diagnose which is the current binding constraint on effective CP and allocate accordingly. This is the counter-power equivalent of the growth bottleneck principle.

15.9 — The Structural Tendency Analysis Formula

The preceding formulas combine into a structural model of historical trajectory. For any unit of analysis — a state, a social formation, an economic arrangement, a political coalition — the trajectory can be estimated from current variable values: Trajectory (Tr) = $T / (C \times L) - CP$. When Tr is strongly positive — high transformation pressure, low adaptive capacity, eroded legitimacy, weak counter-power — the unit is approaching threshold. C and L are bracketed together in the denominator because they jointly constitute the system's absorption capacity: neither alone is sufficient. A system with high legitimacy but no adaptive capacity cannot absorb pressure; a system with adaptive capacity but no legitimacy cannot deploy it. The product of the two, not their sum, determines how much of T the system can contain. What happens at threshold is not determined by the formula: it depends on contingent factors including the specific actors available, the specific historical moment, and the specific form that the accumulated pressure takes when it releases. The formula predicts the approach to threshold, not the specific form of the transition. When Tr is strongly negative — low transformation pressure, high adaptive capacity, intact legitimacy, strong counter-power by challengers — the status quo is durable. But durability is not permanence. The formula tracks change over time: variables that appear stable can shift rapidly, and the compounding effects of duration and intergenerational transmission mean that systems that appear stable may be accumulating pressure that is not yet visible in any single indicator. The most dangerous misreading of this formula is treating current Tr as determinative of future Tr. Trajectory can reverse. Adaptive capacity can be rebuilt. Legitimacy can be restored — slowly and incompletely, but genuinely. Counter-power can be fragmented. The formula describes the current structural tendency, not the inevitable outcome. The value of the formula is that it identifies which variables are most critical to change in order to alter the trajectory, not that it forecloses the possibility of change. Reading the formula output: Trajectory (Tr) = $T / (C \times L) - CP$. A strongly positive Tr indicates that transformation pressure is significantly exceeding the combined product of adaptive capacity and legitimacy, and that counter-power is insufficient to channel the resulting dynamic toward progressive transformation. This is the configuration that produces threshold conditions — not necessarily imminent threshold crossing, but the structural conditions under which threshold crossing is approaching. A strongly negative Tr indicates that the governing coalition's adaptive capacity and legitimacy are significantly exceeding transformation pressure, or that transformation pressure is high but counter-power is insufficient to channel it — a situation that can appear as stability but is analytically different from a stable system with low Tr drive

n by genuinely low transformation pressure. A Tr near zero indicates that the variables are in approximate balance, which is the configuration of managed tension that produces incremental rather than threshold dynamics. The analytical distinction between near-zero Tr from low pressure and near-zero Tr from balanced high pressure is diagnostically significant: the first configuration is genuinely stable, while the second is potentially fragile and sensitive to disturbances that shift any of the component variables. The CP term — why counter-power directs rather than just generates: Counter-power appears in the formula as a subtracted term rather than a multiplier because its primary function is directional rather than additive. CP does not increase transformation pressure; it determines whether transformation pressure, when it reaches threshold conditions, produces progressive reorganization or reactionary consolidation. A system approaching

threshold with high transformation pressure and high CP will undergo a different kind of reorganization than one approaching threshold with the same transformation pressure and low CP. High CP means that organized actors with the analytical and organizational capacity to shape the direction of reorganization are present at the threshold moment — they have the coalitions, the narrative frameworks, and the institutional infrastructure to influence what configuration emerges from the non-linear reorganization that threshold crossing produces. Low CP means that the transformation pressure produces threshold crossing without organized capacity to direct it, which is the configuration that consistently produces outcomes — authoritarian consolidation, elite-managed transition that preserves structural relationships under new institutional form — that do not serve the minimal anchor criterion. How the formula fails — what it cannot capture: The formula captures structural relationships between variables but cannot capture the contingency dimension of actual historical transitions. Specific events — the decisions of specific actors at specific moments, the occurrence of natural disasters or external shocks that alter the variable values suddenly, the emergence of leadership with specific analytical and organizational capacities — are not captured by the formula's structural architecture. The formula identifies where the system is on the threshold trajectory and what the direction of reorganization is likely to be given the current CP configuration; it cannot specify when the specific event that crosses the threshold will occur or exactly what form the reorganization will take. This is not a failure of the formula but a structural feature of non-linear system dynamics: threshold events are determined by the structural configuration but triggered by specific events that are not predictable from the structural configuration alone. The analyst's task is to understand the structural configuration well enough to recognize th

e threshold event when it occurs and respond to it rather than to predict the specific event in advance. Integration with the component formulas: Tr is only as reliable as the component formulas it aggregates. T depends on S (suffering) and its A (accumulated) form; C depends on F, Ec, I, and K; L depends on Lp, Lr, and Ln minus Ld; CP depends on the counter-power formula from 15.8. Each component carries its own measurement uncertainty, and those uncertainties compound in the integration. The appropriate use of Tr is as a structural orientation — a specification of which quadrant of the transformation dynamic the unit of analysis is in — rather than as a precise prediction of when or how threshold crossing will occur. The formula architecture as a whole is designed to make explicit what would otherwise be implicit in structural analysis: the specific variables whose trajectories matter, the relationships between them, and the interactions that determine outcomes. Making these explicit is the precondition for systematic comparison across cases and for the identification of which variables are most leverage-sensitive in specific contexts.

15.10 — Domain-Specific Formulas

The primary formula applies across units of analysis. Domain-specific formulas refine it for particular domains where the relevant variables take specific forms. Geopolitical Power Shift (Gp) = (Economic Output Growth Differential (Δe) $\times$ Military Capacity (Mcp) $\times$ Institutional Network Control (In)) $\times$ (1 – Internal Cohesion Cost (Icc)) Power shifts between states are predictable at the structural level when growth differentials compound over time and existing institutional arrangements reflect the previous power distribution rather than the current one. The tension between material power distribution and institutional recognition of that distribution is the primary driver of geopolitical conflict. Wars and institutional reorganizations are both mechanisms for realigning institutional recognition with material reality — the formula predicts when the pressure for realignment reaches threshold, not which mechanism resolves it. Ideological Dominance Decay (Id) = Legitimacy Deficit Rate (Ldr) $\times$ Counter-Narrative Reach (Cnr) / Narrative Infrastructure Control (Ni) Id here measures the rate of decay — how rapidly the ideology is losing its hold — not the current degree of dominance. A higher Id value means faster erosion: the formula rises as legitimacy deficit accumulates and counter-narratives spread, and falls as narrative infrastructure proves sufficient to contain both. Governing ideologies do not collapse from direct refutation. They decay when the

gap between their claims and observable reality becomes too large to manage through the narrative infrastructure available to maintain them, and when counter-narratives have sufficient reach to provide an alternative frame for interpreting that gap. The formula implies that the most effective strategy for accelerating ideological decay is not winning philosophical arguments but expanding the reach of counter-narratives while the legitimacy deficit accumulates — which it does independently, driven by the system's own failures. Elite Fragmentation Rate (Ef) = (Performance Pressure (Pp) × Succession Conflict (Sc) × Resource Constraint (Rcn)) × (1 – Ideological Cohesion (Ico)). Elites fracture under predictable conditions: when performance pressure exceeds the system's capacity to deliver to all elite factions simultaneously, when succession produces competing claimants, and when resource constraints force zero-sum choices within the coalition. Ideological cohesion is the primary buffer against fragmentation — elites who share a genuine common worldview can absorb resource conflicts that would otherwise produce fractures. The decay of ideological cohesion within an elite, often visible years before it produces visible fragmentation, is a leading indicator that the elite coalition is approaching fracture threshold. The geopolitical power shift formula — how states gain and lose structural position: $Gp = (\text{Economic Output Growth Differential} \times \text{Military Capacity} \times \text{Institutional Network Control}) \times (1 - \text{Internal Co}$

hesion Cost). Growth differentials compound over time and are the primary driver of structural power shifts at the geopolitical level: states that grow significantly faster than competitors over sustained periods accumulate material advantages across all other dimensions of power. Military capacity amplifies or limits the leverage available from economic position — states with high growth but limited military capacity are more constrained in how they can convert economic advantage into structural leverage than states that can project force credibly. Institutional network control is the third dimension: control of the multilateral institutions, financial infrastructure, and standard-setting processes through which economic relations are organized gives leverage that is not captured by GDP or military spending metrics. Internal cohesion cost discounts total Gp by the resources and political capital consumed maintaining the governing coalition — high internal conflict reduces the effective power available for geopolitical projection even when the structural resources are high. The movement consolidation formula — why organizations drift: $Mc = (\text{Organizational Capacity} \times \text{Resource Base} \times \text{Narrative Coherence}) \times (1 - \text{Institutional Capture Rate})$. Movement consolidation captures the trajectory of counter-power organizations over time. Organizational capacity and resource base are necessary inputs, but the interaction matters: organizations with high capacity and low resources are fragile to sustained campaigns; organizations with high resources and low capacity are vulnerable to strategic redirection without the organizational discipline to resist it. Narrative coherence — the degree to which the organization's public account of what it is and why it acts as it does is internally consistent and consistently applied — determines whether the organization can maintain coalition membership through the periods of compromise and setback that sustained counter-power work requires. The capture rate discount is the primary mechanism through which successful movements are converted into their opposites: as they acquire resources and institutional presence, the capture dynamics analyzed in 7.3c apply to them as they apply to any institution with control of significant resources. The economic transition formula — when production configurations change: $Et = (\text{Technological Change Rate} \times \text{Labor Mobility} \times \text{Capital Reallocation Speed}) \times (1 - \text{Political Resistance to Transition})$. Economic transitions occur when technological change makes existing production configurations economically inferior to alternative configurations at a rate faster than existing institutional arrangements can absorb. Labor mobility and capital reallocation speed determine how quickly the new configuration can be operationalized: high labor mobility allows the reallocation of work to the activities the new configuration requires; high capital reallocation speed allows the investment shift from the old

to the new production basis. Political resistance to transition discounts the structural pressure for change by the organized resistance of actors with concentrated interests in the existing configuration — workers and capital owners whose position depends on the old configuration have asymmetric incentives to resist transition compared to the dispersed beneficiaries of the new one. How domain-specific formulas relate to

the primary formula: the domain-specific formulas are specifications of how the primary formula's variables take form in particular domains. Geopolitical power shift is a specification of how C operates at the state level. Movement consolidation is a specification of how CP develops over time. Economic transition is a specification of how the interaction between T and C operates in the economic domain. None of the domain formulas are independent of the primary formula; they are applications of its structural logic to specific domains where the variables take particular forms and where the relationships between them follow identifiable patterns. Diagnostic application: when using a domain-specific formula, first identify which component of the primary formula it is specifying. Then ask: what are the specific indicators for each variable in this domain? What are the measurement proxies that are available for this domain? What are the known sources of measurement error? Domain-specific formulas improve analytical precision at the cost of generality — they apply more precisely to their domain but cannot be transferred to others without specifying how the variables take form in the new domain.

15.11 — The Time Variable: Windows and Cycles

The formulas above describe structural relationships between variables but are incomplete without a theory of timing. Material variables change on multiple timescales simultaneously, producing the cyclical patterns that historical analysis has identified — demographic cycles, economic cycles, legitimacy cycles, geopolitical cycles — each operating on its own tempo and intersecting with others in ways that create windows of heightened transformation probability. Transformation Window (T_w) = Cycle Convergence (C_c) $\times$ Focal Event Probability (F_e) $\times$ Counter-Power Readiness (C_r) Cycle Convergence (C_c) is a composite variable representing the number of major structural cycles simultaneously reaching their turning point. It is not a binary variable: it ranges from near-zero (one cycle peaking in isolation) to high (multiple cycles peaking together). When multiple cycles reach their turning points simultaneously — when demographic pressure, economic stress, legitimacy deficit, and elite fragmentation all peak in the same period — the probability of transformation is multiplicatively higher than when any single cycle peaks alone. These convergences are not random: they are produced by the same underlying structural conditions and tend to recur on recognizable timescales. Turchin's work documents several such cycles in Western history; the formula generalizes the principle. Focal events — assassinations, financial crises, military defeats, natural disasters — do not cause transformations. They provide the focal point around which accumulated pressure crystallizes into action. The formula implies that the same event produces very different outcomes depending on where each cycle is at the moment of the event. A financial crisis during a period of high legitimacy, elite cohesion, and low transformation pressure is absorbed as a manageable setback. The same crisis during a period of legitimacy deficit, elite fragmentation, and high accumulated pressure triggers the reorganization that the pressure had been building toward. Counter-power readiness (C_r) is the degree to which challengers have organizational infrastructure, narrative coherence, and coalition width sufficient to convert a focal event into directed change rather than undirected disruption. Transformation windows that open without counter-power readiness tend to produce either repressive consolidation — the existing regime using the crisis to eliminate challengers — or undirected collapse, in which the old arrangement disintegrates without an alternative capable of consolidating what emerges. The formula implies that the strategic priority during periods of low transformation pressure is not mobilization but preparation: building the infrastructure that will be required when the window opens. Why timing is irreducible: The primary formula and its component formulas describe structural relationships between variables at a given moment: if T is high relative to $C \times L$ and CP is low, the system is in thr

eshold-proximity configuration. What the structural snapshot cannot capture is whether the threshold will be crossed in months, years, or decades — and whether the threshold will be crossed at all before one of the variables shifts in ways that change the structural configuration. The timing dimension requires a theory of how variables change over multiple timescales simultaneously, which is what the time variable analysis

provides. The political significance: two systems in identical structural configurations can have very different timing characteristics depending on which cycles they are at and how those cycles are interacting. A system at the peak of a demographic youth bulge, in the trough of a legitimacy cycle following a governance failure, and at the bottom of an economic cycle has threshold proximity that differs from a system in the same primary formula configuration that is at a different phase of each cycle — even if the current variable values are identical. Four cycles and their interaction: Demographic cycles operate on generational timescales (fifteen to thirty years) and determine the age distribution of the population, which shapes political mobilization capacity, economic labor supply, and the relative weight of constituencies with different interests in the current arrangement. Economic cycles operate on business timescales (five to twelve years) and determine the material conditions of daily life in ways that are immediately politically salient. Legitimacy cycles operate on institutional timescales (ten to forty years) and reflect the accumulation and depletion of the trust infrastructure that determines how performance failures are interpreted. Geopolitical cycles operate on hegemonic timescales (fifty to one hundred years) and determine the international constraints and opportunities available to domestic political arrangements. Each cycle has its own internal logic; the politically consequential moments are when they intersect in ways that align multiple cycles at phases that amplify each other's political effects. The Arab Spring of 2011 is a documented case of this intersection: a demographic youth bulge, an economic contraction that reached youth disproportionately, legitimacy depletion from decades of governance failure, and a geopolitical context in which external powers were ambivalent about repressive stabilization combined to produce threshold conditions across multiple countries simultaneously. Windows of opportunity and their characteristics: Political windows — the periods during which the configuration of variables makes transformation pressure most likely to produce threshold crossing — are not randomly distributed. They are structurally determined by the interaction of cycles and have predictable characteristics. Windows tend to be brief because the cycle interactions that create them shift: the demographic bulge passes, the economic trough bottoms out and begins recovery, the legitimacy crisis is managed thr

ough limited reform. The brevity of windows has strategic implications: counter-power capacity that is accumulated during non-window periods determines what is available to act when windows open. Organizations that wait for windows to begin accumulating are systematically too late; the organizational capacity, narrative infrastructure, and coalition breadth that determine what is possible during a window must be built during the periods when the window is not yet open. The intersection of cycles and structural variables: the time variable analysis does not replace the structural variable analysis of the primary formula. It specifies when the structural configuration is most likely to produce threshold dynamics, given the current phase of each relevant cycle. A system with high T, low C, and eroded L that is also at a confluence of cycle phases favoring transformation pressure has higher threshold proximity than the structural variable values alone would indicate. Conversely, a system with high structural transformation pressure at a cycle phase that is unfavorable — demographic trough, economic peak, legitimacy cycle at high — has lower effective threshold proximity than the structural variables alone would indicate. Application: when assessing threshold proximity, always ask what cycle phases the unit of analysis is currently in, and whether the cycle intersection is amplifying or dampening the structural pressure indicated by the primary formula variables. Structural analysis without timing analysis will systematically misread both the urgency of current conditions and the appropriate strategic horizon for counter-power accumulation.

15.12 — Application: Reading Current Conditions

The formulas become useful when applied to specific conditions. The application procedure is the same regardless of unit of analysis: Step 1: Identify the unit. What system, formation, or arrangement is being analyzed? The variables take different forms at different levels of analysis — a state, a class formation, an ideological bloc, a global economic arrangement — but the structural relationships hold across levels. Step 2:

Estimate variable values. Not with false precision, but with honest qualitative assessment. Is transformation pressure high, medium, or low? Is adaptive capacity increasing or decreasing? Is legitimacy eroding or being restored? The formulas do not require numbers — they require honest assessment of direction and approximate magnitude for each variable. Step 3: Identify which variables are most dynamic. The formula's predictive value is highest when it identifies variables that are changing rapidly. Stable variables are less important for near-term trajectory than variables in rapid motion. The question is not just what the values are but what they are doing. The leading and lagging indicators developed in 6.6 are the most operationally useful guide for detecting which variables have crossed into dynamic motion — they translate the formula's abstract variable states into concrete behavioral signals that can be tracked in real conditions. Step 4: Identify feedback loops. Many variables in the formula are connected through feedback: high transformation pressure that elicits repression can increase counter-power capacity by demonstrating the system's willingness to use force; legitimacy deficit that produces elite fragmentation can reduce adaptive capacity further, accelerating pressure accumulation; counter-power capacity that successfully wins concessions can either be absorbed into institutional capture or consolidated into durable organizational infrastructure. Identifying which feedbacks are active and what direction they are running is often more predictive than the static variable values alone. Step 5: Identify the threshold question. For any unit in which transformation pressure is significant, the most important question is not 'will change happen?' but 'what would push this system past threshold, and what is the probability that those conditions will emerge in the relevant time horizon?' The threshold is not known in advance, but its approximate location can often be estimated from historical comparison and from the degree to which the system is currently managing pressure through legitimacy and adaptive capacity versus managing it through suppression alone. Suppression-heavy management is a leading indicator of threshold proximity — it means the system has exhausted its other options. The application of this formula set is an empirical practice, not an algorithmic one. The formulas identify what to look for and how the variables relate — they do not replace the judgment required to assess specific conditions. Two ana

lysts applying the same formula to the same unit may reach different conclusions if they assess variable values differently. The value of the formula is not to eliminate disagreement but to make disagreement productive: when analysts disagree about trajectory, the formula identifies exactly which variable assessment is driving the difference and makes that specific disagreement the object of further analysis. How these formulas should and should not be used. Correct use: "Legitimacy is clearly eroding — performance failures are accumulating and the narrative infrastructure is visibly strained. Adaptive capacity has not kept pace: reform infrastructure exists formally but capture rates are high. Transformation pressure is significant and rising. The most critical variable to watch is elite cohesion — if fragmentation accelerates, the system loses its primary remaining buffer." This is the formula doing its intended work: organizing observation, directing attention to critical variables, characterizing direction and relative magnitude without false precision. Incorrect use: "Transformation pressure is 6.4 on a ten-point scale; adaptive capacity is 3.1; therefore $T = 3.3$, which exceeds the threshold of 2.8, so collapse will occur within eighteen months." This is formula fetishism: the assignment of false precision to variables that do not admit it, the derivation of specific timelines that the formula does not generate, and the replacement of material analysis with arithmetic. The formula's notation — multiplication signs, equals signs, variable letters — does not authorize this. It is a structural shorthand, not a calculating machine. Any use of these formulas that produces a number, a timeline, or a probability expressed as a specific figure is a misuse, regardless of how rigorously it was conducted. On Measurement: Proxies, Limits, and the Role of Judgment The variables in these formulas — transformation pressure, adaptive capacity, legitimacy, internal fragmentation, elite capture rate — are analytically precise but not directly observable. Every practical application of the formulas requires proxies: observable indicators that track the underlying variable imperfectly. The gap between the variable and its proxy is not a technical limitation to be eliminated by better measurement. It is a permanent feature of the analytical situation, and treating proxy measures as if they were direct observations of the underlying variable is among the most common failure modes in quantitative social analysis. Useful proxies for key

variables include: for legitimacy erosion, survey measures of institutional trust combined with participation rates in formal political processes and incidence of informal exit behaviors (emigration, tax avoidance, parallel institution formation); for internal fragmentation, frequency and public visibility of factional disputes, leadership turnover rates, divergence between public statements and internal communications where accessi

ble, and exit rates among experienced personnel; for elite capture rate, the composition and funding sources of leadership, the gap between the movement's stated priorities and its actual resource allocation, and whether the movement's accountability mechanisms function when applied to its own leadership as rigorously as when applied to external targets. None of these proxies is sufficient alone. None is free from manipulation. All require triangulation with other sources and contextual judgment about what the indicators mean in the specific conditions being analyzed. Measurement principle: The formulas specify relationships between variables. The proxies operationalize those variables for specific contexts. The judgment required to select appropriate proxies, interpret what they indicate in context, and recognize when they are being gamed is not replaceable by the formulas themselves. An analyst who applies the formulas mechanically, without the contextual judgment that proxy selection requires, has not used the framework — they have used the appearance of the framework to substitute quantification for analysis. The formulas are tools for organizing judgment, not substitutes for it. Conditional Application: Variable-Specific Intervention Logic The application procedure above identifies which variables are most dynamic and which feedback loops are most active. A further step — which the procedure does not complete for the analyst — is translating variable diagnosis into intervention logic. The following conditional heuristics are not rules but starting points for that translation. Each identifies a characteristic variable configuration and the class of intervention that addresses it most directly, while acknowledging that configurations compound and that no heuristic substitutes for the specific analysis of specific conditions. When V is the primary driver of high S (structural suffering is present but not recognized as structural): the primary intervention is epistemic rather than material. Building the narrative and informational infrastructure that makes structural causation legible — conscientização in Freire's sense, the production of shared accounts that convert individual experience into collective diagnosis — is prerequisite to mobilization that will sustain. Material interventions on R or B that are not accompanied by V reduction will be attributed to the wrong causes, will not produce the coalition formation that transformation requires, and will be absorbed without changing the interpretive framework that makes re-accumulation of S structurally likely. When R is high and V is also high: the sequencing problem is acute. Populations experiencing material deprivation while attributing it to individual failure or misidentified structural causes are not yet in the organizational position to convert material grievance into durable political pressure. The intervention sequence — V reduction before or simultaneous with R reduction, not after —

is the condition under which material improvement reinforces rather than dissolves the structural analysis that transformation requires. Material improvement without V reduction can legitimize the arrangements producing it, making subsequent transformation harder rather than easier. When C is the primary constraint (adaptive capacity is low not because of captured institutions but because reform infrastructure was never built): the intervention is organizational before it is political. The absence of channels for absorbing transformation pressure means that accumulated T has nowhere to go except breakdown. Building the institutional forms — unions, civic organizations, informal mediation networks, legal aid infrastructure — that give pressure a directed outlet is the condition under which reform becomes available as an option rather than only crisis management after threshold crossing. When elite fragmentation (Ef — see 15.10) is the primary opening: the strategic priority shifts from pressure accumulation to window exploitation. Elite fragmentation is time-limited; factions that have defected from the coalition will re-cohere if the opening is not used. Organizational readiness — the Cr component of Tw — determines whether the window produces directed change or absorption through elite circulation without structural transformation. The heuristic: invest in Cr continuously, not in response to windows; windows are recognized at the moment they open, not built after they appear.

15.13 — What the Formulas Cannot Do

The formulas in this Part have specific and acknowledged limits. Naming them is not a caveat appended to protect the theory — it is part of the theory. They cannot predict specific events. The formulas predict structural tendency and threshold proximity, not the specific form, timing, or actors of transformation. Whether accumulated pressure resolves through reform, revolution, collapse, or repressive consolidation depends on contingent factors — the specific actors available at the moment of threshold crossing, the international context, the specific form of the focal event — that the structural formula does not determine. They cannot assign precise coefficients. The relative weight of each variable in the primary formula is real but not currently measurable with sufficient precision to assign numbers. Institutional flexibility matters more in some configurations than others. Duration matters more when intergenerational transmission is high. The formula specifies that the variables matter and the direction of their effect; the precise weighting requires empirical work that has not yet been done at the level of rigor the formulas imply. They do not determine agency. The most important limit. The formulas describe structural tendencies — what becomes probable under specified conditions. They do not determine what specific actors do with those probabilities. A movement that understands its structural position can act on that understanding in ways that alter the trajectory. A regime that understands threshold proximity can take genuine adaptive action that reduces transformation pressure. The formulas are tools for understanding the terrain; they are not the terrain itself. Actors who treat structural tendency as deterministic lose the agency that structural understanding is supposed to enable. They are falsifiable and provisional. Each formula implies testable claims. If transformation pressure consistently fails to produce threshold crossing even at high levels with low adaptive capacity, the formula requires revision. If elite fragmentation does not follow from the specified conditions, the formula for elite fragmentation is wrong or incomplete. The formulas are offered as the best current structural account available, not as permanent laws. The appropriate response to evidence that contradicts them is revision, not special pleading. What the defeasible structure means for formula use: The defeasible reasoning architecture established in 12.1d has direct implications for how the formulas here should be used and what would count as falsifying them. A formula claim — "high T with low C and eroded L produces threshold proximity" — is not falsified by any case where threshold crossing does not occur. It is falsified by a pattern of cases where threshold crossing does not occur and no recognized defeater is present. The recognized defeaters are not post hoc inventions: they are structurally specified in the formula architecture itself (strong Rc, exte

rnal intervention, narrative legitimacy buffering performance failure) and in the applied cases of 15.14. An analyst who applies the formulas and finds that the predicted tendency does not obtain should ask first whether a defeater is present before concluding that the formula is wrong. An analyst who consistently finds defeaters in every case where the formula's prediction does not hold, without those defeaters being independently observable, has not applied the formula — they have used the flexibility of defeasible structure to protect a prior conclusion. What the formulas cannot resolve about their own variable weights: The formulas specify the direction and sign of variable relationships but cannot assign precise relative weights. The claim that $T = (S \times D \times P) / (C \times L) - CP$ specifies that transformation pressure rises with accumulated suffering and falls with adaptive capacity and legitimacy — but it cannot specify whether a one-unit increase in S produces more transformation pressure than a one-unit decrease in L. This indeterminacy is not a failure of the formula architecture; it is an accurate representation of the current state of structural analysis. Historical cases can be analyzed post hoc to establish rough orderings of variable importance in specific configurations, but these orderings are context-dependent and cannot be generalized without further empirical work. The formulas are useful for identifying which variables are most dynamic in a given situation, not for calculating a precise trajectory score. The formulas and the two-level architecture: The formula architecture operates at the evaluative level — it tracks structural suffering (S) and its consequences for transformation pressure (T). It does not operate at the ontological level. This means the

formulas are silent about what life generates in the absence of suffering, about what the anti-foreclosure criterion requires, and about what institutional arrangements would maximize adaptive response capacity rather than merely reduce structural harm. The formula architecture is a tool for diagnosing when the floor criterion is being violated and predicting the structural consequences of that violation. It is not a tool for optimizing becoming. Using the formulas to evaluate institutional arrangements against the anti-foreclosure criterion requires supplementary analysis that the formulas do not provide. Summary of formula limits: (1) defeasible, not exceptionless — identify defeaters before concluding formula is wrong; (2) direction and sign specified, not magnitude — do not assign precise numerical weights; (3) operates at floor level, not at the anti-foreclosure level — formulas diagnose harm, not optimize becoming; (4) structural tendency, not event prediction — the formula cannot specify when or how threshold crossing occurs, only whether the structural configuration makes it probable. What the formulas cannot say about the anti-foreclosure horizon: Every formula in Part XV

tracks variables that register the floor criterion: suffering accumulation, legitimacy depletion, adaptive capacity failure, counter-power capacity. No formula tracks the anti-foreclosure criterion — the prohibition on foreclosing adaptive response capacity — because adaptive response capacity does not generate the kind of systemic stress that makes structural analysis tractable. Suffering at scale produces measurable behavioral consequences: declining cooperation, shortened time horizons, institutional legitimacy erosion. These are trackable. The expansion of human developmental capacity does not produce equivalent structural signals. A society in which more people are genuinely developing — cognitively, creatively, morally — does not differ from a stagnant society in the variables the formula architecture measures, unless that development is being blocked by structural arrangements that produce identifiable structural suffering. This means the formula architecture is a complete tool for diagnosing when the floor criterion is being violated and incomplete as a tool for assessing whether institutional arrangements are serving the anti-foreclosure criterion. Analysts who use the formulas to evaluate whether a society is genuinely flourishing — rather than merely not failing — have exceeded the formula architecture's scope.

15.14 — Applied Cases: The Formulas in Operation

The following four cases demonstrate how the formula set is used in practice. Each applies the full variable stack to a historical situation for which the outcome is known — allowing the reader to see how the formula identifies structural dynamics that the outcome alone does not supply. In each case, the analysis specifies which variables were most critical, what their interaction produced, and what the formula would have indicated in advance had it been applied contemporaneously. Cases were selected to illustrate different configurations: a transition that succeeded (Tunisia), one that was suppressed despite high pressure (Belarus), one demonstrating a recurring legitimacy cycle without structural resolution (Argentina), and one illustrating threshold crossing that produced elite circulation rather than structural transformation (Indonesia). Case A: Tunisia 2010–2011 — Why This One Crossed Threshold Tunisia in late 2010 is the paradigmatic case for applying the transformation pressure formula because the outcome — rapid regime transition — was broadly unexpected, including by regional specialists who observed similar conditions in neighboring states without expecting collapse. The formula applied in advance would have identified Tunisia as high-risk for threshold crossing on several variables simultaneously. Structural Suffering ($S = R + B + V$): Resource deprivation (R) was high and, critically, concentrated in the interior regions while coastal elites maintained relative prosperity — making the suffering differential rather than universal, which raised its political salience. Institutional betrayal (B) was at maximum: the Ben Ali regime had maintained a performance legitimacy claim through economic modernization, and when that claim failed — graduate unemployment reached structural levels — the betrayal gap was enormous because expectations had been raised by decades of that claim. Visibility suppression (V) was moderate and declining: internet penetration and mobile networks had created channels that the state's information infrastructure could not fully control,

and the Mohammed Bouazizi self-immolation provided a focal symbol that made distributed suffering collectively legible almost instantly. Duration and Scale ($A = S \times D \times P \times G$): Duration was long — the structural conditions had been accumulating for years. Scale was broad — the unemployment crisis reached across the youth demographic that represented the majority of the population's near-future. Intergenerational transmission was high — educated youth entering a system structurally incapable of absorbing them understood that their children would face the same conditions, not better ones. Adaptive Capacity (C): Low. Ben Ali's system had eliminated competing institutions capable of channeling pressure: civil society was controlled, trade unions weakened, political opposition foreclosed. When pressure arrived, the system had no mechanism to absorb it at the margins — it could only suppress

s or capitulate. Reform infrastructure had been so thoroughly hollowed out that the elite capture rate was effectively total. Repressive Capacity (Rc): Moderate but declining rapidly. Coercive resources existed, but enforcement loyalty fractured when the scale of protest exceeded what the apparatus could suppress without visible mass violence against civilians — violence that would have destroyed performance and narrative legitimacy simultaneously. The internal opposition variable (Io) rose sharply as military and police calculated the regime's survival probability and began defecting in behavior if not in formal declaration. Legitimacy (L): Approaching zero across all three dimensions simultaneously — the rarest and most consequential configuration. Performance legitimacy had collapsed with graduate unemployment. Procedural legitimacy had never been established under authoritarian conditions. Narrative legitimacy — the story that Ben Ali's modernization project was the right path — was directly contradicted by the visible failure of that project for the generation that had been promised its benefits. The formula would have identified Tunisia as a system with high accumulated pressure, low adaptive capacity, declining repressive capacity, and collapsing legitimacy on all three dimensions simultaneously. The trajectory (Tr) was strongly positive. What the formula would not have predicted — correctly — is the specific date, the specific trigger (Bouazizi's act), or the specific rapidity of the transition. It would have indicated: high threshold proximity, high sensitivity to focal events, and low system capacity to absorb a focal event without cascading. That is precisely what occurred. The comparison with Egypt — which also transitioned in 2011 — and with Bahrain — which did not — illustrates the formula's discriminating power. Egypt's higher Rc (military with distinct institutional interests from the regime) meant transition took a different form: the military managed the transition rather than the regime collapsing. Bahrain's lower accumulated pressure, Saudi intervention (external Rc injection), and smaller coalition width prevented threshold crossing despite significant transformation pressure. The formula correctly identifies these as structurally different situations requiring different predictions — not because the formula is calibrated to the outcomes, but because the variable profiles genuinely differ. Case B: Belarus 2020 — High Pressure, No Transition Belarus 2020 is the critical counter-case: transformation pressure that was real, visible, and broadly recognized — hundreds of thousands in the streets — that did not produce transition. The failure to cross threshold is as theoretically important as the crossing. The formula correctly predicts this outcome when applied to variable values, not to the surface drama of mass mobilization. Transformation Pressure (T): High. The fraudulent election results were a legitimacy-destroying focal ev

ent that crystallized accumulated institutional betrayal. Structural suffering was present though not at the extreme levels of Tunisia — Belarus had a state-supported industrial economy that maintained basic material provision even as political freedom was eliminated. The primary driver of S was institutional betrayal (B), not resource deprivation (R), which is structurally less explosive because it does not combine material urgency with political grievance simultaneously. Repressive Capacity (Rc): This is where the case turns. Lukashenko's security apparatus maintained enforcement loyalty throughout — KGB and OMON units continued to follow orders even when the brutality of repression was globally visible. This is structurally explained by the apparatus's material dependence on the regime: security service personnel were among the best-compensated workers in Belarus, creating a material stake in regime survival that decoupled their interests from those of the protesting population. Information control (Ic) was partially degraded —

Telegram channels organized protests effectively — but this raised T without reducing R_c proportionately. External R_c injection: Russian intervention — financial support, security intelligence sharing, and the credible implicit threat of direct support — effectively augmented Belarus's R_c beyond what the domestic variable values alone would have generated. The formula's domestic variables correctly indicate high threshold proximity; the external variable that suppressed crossing is a geopolitical one (G_p from 15.10) not captured in the domestic formula alone. This is the formula correctly identifying its own limit: domestic transformation pressure formula requires geopolitical formula integration for small states in the sphere of great powers. Counter-Power Capacity (CP): Organizationally insufficient. The protest movement had scale and narrative but lacked the organizational infrastructure to convert mobilization into sustained pressure after the initial wave of repression. Coalition width reached the urban professional class but did not extend effectively into the industrial working class — the demographic most materially embedded in the state economy — limiting the economic disruption capacity that would have raised the cost of continued repression. The Belarus case demonstrates the formula's primary analytical value: explaining why high visible pressure does not always produce threshold crossing. The answer is not accidental or mysterious — it is structural. R_c was maintained by material incentives within the apparatus and augmented externally; CP lacked the organizational depth to sustain pressure through repression; T , while high, was driven primarily by betrayal rather than material deprivation, limiting its economic leverage component. The formula identifies all three structural conditions that explain the outcome. Case C: Argentina's Recurring Crises — The Legitimacy Cycle Argentina since 1983 provides a differ

ent kind of application: not a single threshold event but a recurring pattern of legitimacy crises, partial recoveries, and re-accumulation of pressure. The formula applied to Argentina is less about predicting a single transition than about explaining why the same structural dynamic recurs — and what variable would need to change to break the cycle. The recurring pattern: Argentina cycles through the following sequence with near-structural regularity: (1) Performance legitimacy is built on commodity export booms or debt-financed growth; (2) External shock or unsustainable debt triggers performance collapse; (3) Institutional betrayal (B) spikes as the population experiences the gap between economic promise and economic reality; (4) Transformation pressure reaches threshold; (5) Political transition occurs; (6) New government rebuilds performance legitimacy through similar mechanisms; (7) Cycle repeats. The formula identifies this as a structural loop, not a series of independent crises. The critical variable: Adaptive capacity (C) — specifically institutional flexibility (F) and reform infrastructure (I). Each cycle partially degrades these: institutions designed for managing the transition become captured; reform infrastructure that appeared during crisis is hollowed out during the recovery phase when pressure is lower and the political cost of structural reform is higher. Elite capture rate (K) consistently rises during recovery phases, leaving the system less capable of genuine adaptation the next time pressure arrives. The formula predicts: until adaptive capacity is structurally rebuilt — not borrowed from performance legitimacy during boom cycles — the pattern will continue. The Argentina case illustrates the formula's use for structural diagnosis rather than crisis prediction. The question is not 'when will the next crisis occur?' but 'what variable configuration is producing this recurring pattern, and what change in which variable would break it?' The formula answers: the variable is adaptive capacity, specifically its systematic degradation during recovery phases through elite capture. This is actionable in a way that crisis prediction is not — it identifies the structural intervention rather than the timing of its necessity. Live update — Milei 2023–present: The election of Javier Milei in November 2023 on an anarco-capitalist platform provides a real-time test of the formula's structural logic. Milei's program is explicitly framed as a counter to B: the argument that decades of Peronist institutional capture have so thoroughly corrupted the state that institutional demolition is necessary before genuine development can occur. The formula identifies a structural problem with this framing. Milei's shock doctrine — elimination of the central bank, 50 percent devaluation at inauguration, elimination of energy and transport subsidies, reduction of the public sector by approximately 70,000 positions — directly spikes R in the short term in

order to, on the program's logic, reduce B structurally over the medium term. This is the floor/anti-foreclosure trade-off that the framework identifies as requiring specific justification: a mechanism for floor restoration must be specified and accountable, not assumed to follow from the logic of the intervention. The formula prediction for the Milei configuration: if R spikes significantly while C remains low (the institutional infrastructure for managing transition remains thin), transformation pressure will re-accumulate before the medium-term B reduction materialises — not because the diagnosis of B was wrong, but because the sequencing ignored the adaptive capacity requirement. The Argentina case is now a live test of whether the formula's structural prediction outperforms the program's own model. The outcome will be available for assessment at the next version of this document. Case D: Indonesia 1997–1998 — Reformasi and the Limits of Developmental Authoritarianism The Indonesian case of 1997–1998 provides a non-Western, non-Latin American applied case that demonstrates the formula's discriminating power across different institutional configurations. It is also the case that most directly tests the framework's analysis of developmental authoritarianism (7.8): the Suharto New Order was precisely the kind of arrangement that produced real floor achievements while systematically foreclosing the accountability mechanisms that would have allowed it to detect and correct its own failure modes when external conditions changed. The New Order's thirty-two years of governance (1966–1998) produced genuine floor improvements by PMN's criterion: per capita income grew from approximately USD 70 in 1966 to USD 1,200 by 1996, infant mortality fell from 145 per 1,000 in 1960 to 47 by 1995, and primary school enrollment expanded from 45 to 95 percent. These achievements were real and the framework's consistency standard requires acknowledging them. They were achieved through the developmental authoritarian model: a technocratic economic team (the 'Berkeley Mafia' of US-trained economists) with insulation from political interference, combined with systematic suppression of independent labor organization, press freedom, and political opposition. The meritocratic bureaucratic capacity that the authoritarianism section (7.8) identifies as the structural feature distinguishing developmental from extractive authoritarianism was partially present — but coexisted with Suharto family cronyism that systematically captured the most profitable economic sectors. The 1997 Asian financial crisis provided the focal event that the formula predicts should trigger threshold crossing when accumulated pressure is high and adaptive capacity is low. By the formula variables: Structural Suffering (S) was higher than surface indicators showed, because V was high — the New Order's information control had suppressed visibility of rural poverty, labor conditions, and ethnic Chinese t

argeting that the crisis made suddenly legible. When the rupiah collapsed by 80 percent against the dollar and IMF structural adjustment conditions imposed subsidy cuts on basic goods, R (resource deprivation) spiked rapidly in a population that had accumulated real but precarious material improvements. B (institutional betrayal) was high: three decades of promised development had built expectations that the crisis catastrophically failed. Adaptive Capacity (C) was near zero: the New Order had systematically dismantled every institution that could have channeled pressure or negotiated reform — independent unions, opposition parties, free press, autonomous civil society. Repressive Capacity (Rc) collapsed through enforcement loyalty failure: military commanders calculated the regime's survival probability and defected behaviorally, refusing to suppress the student demonstrations that occupied the DPR building in May 1998. The Indonesia case illustrates a dynamic that the Tunisia and Belarus cases do not: the role of ethnic fracture lines (8.6, 10.9) in the form that threshold crossing takes. The May 1998 riots that accompanied Suharto's fall were directed primarily at the ethnic Chinese Indonesian community, which held a disproportionate share of commercial wealth due to New Order economic policies that had channeled Chinese business capacity into the economy as a counterweight to indigenous Muslim business competition. The manufactured fracture line — ethnic Chinese as the visible face of economic inequality — was a structural feature of the New Order's political economy that directed transformation pressure toward horizontal ethnic conflict rather than vertical structural critique. This is the affective mobilization mechanism (3.8b) operating through manufactured fracture lines (10.9): genuine structural suffering redirected toward the misidentified source. The reformasi transition (1998–2004) also provides evidence for the formula's predictions about what transitions produce

when counter-power capacity is low. Counter-Power (CP) was organizationally thin: student movements had tactical capacity but no sustained organizational infrastructure, no programmatic alternative, and no coalition extending beyond the educated urban class. The transition was therefore managed by the existing elite — Habibie (Suharto's vice president) then Wahid, Megawati, and Yudhoyono through democratic succession — producing elite circulation rather than structural transformation. The decentralization reforms that followed redistributed political authority to regional governments without corresponding redistribution of accountability mechanisms, reproducing the capture dynamics of the center at every provincial level. The Indonesia case diagnostic: the formula correctly predicts high threshold proximity from the variable configuration — high S (with V suppressing visible severity), low C (systematic dismantling of adaptive institutions), declining Rc (enforcement loyalty

collapse), and external shock (Asian financial crisis) as focal event trigger. The formula also correctly predicts that low CP produces elite circulation rather than structural transformation: the reformasi transition changed the form of governance from authoritarian to electoral without changing the ownership structures, bureaucratic incentive architecture, or capture dynamics that the New Order had institutionalized. The formula does not predict the specific form of the threshold crossing — the May riots, the ethnic targeting, the specific sequence of military defection — but correctly identifies the structural conditions that made some form of rapid reorganization near-certain and any form of genuine structural transformation near-impossible given the counter-power deficit. Some of the predictions this architecture generates will be wrong. The Argentina live update is structured precisely to test this: a specific variable configuration — R rising while B remains structurally unaddressed — generates a specific structural prediction, that pressure will re-accumulate before the legitimacy gain from R materializes into reduced S, which will be checked against outcomes rather than retrofitted to match them after the fact. This is what falsifiability in the sense specified in 15.13 looks like in practice: not falsified by any single deviation, but falsified by a pattern of deviations the framework has no defeater for. Whether this formula architecture is doing the analytical work it claims depends on whether it keeps generating that kind of testable commitment, not on whether any individual application turns out to be right. A structural feature of the formulas that the limits above presuppose but do not state explicitly: the formulas express defeasible tendencies, not exceptionless laws. The claim that "high transformation pressure with low adaptive capacity produces threshold proximity" means that this configuration generally produces threshold proximity in the absence of specific defeating conditions — strong narrative legitimacy that absorbs pressure, external Rc injection that supplements domestic repression, counter-power fragmentation that prevents transformation pressure from organizing into action. The formulas are not falsified by cases where the expected output does not occur; they are falsified by cases where the expected output does not occur and no recognized defeater is present. This distinction is analytically important: it determines what kind of evidence counts as a challenge to the formula architecture, and what kind of evidence an analyst applying the formulas should collect. Applying the formulas as if they were exceptionless laws will produce systematic overconfidence in cases where defeaters are present. Applying them with appropriate attention to potential defeaters — the approach the application procedure in 15.12 is designed to support — produces the calibrated probability assessment that the framework's epistemological c

ommitments require.

15.15 — The Compressed Core: PMN Without the Architecture

The full architecture of this framework — seventeen parts, a formula system, hundreds of diagnostic questions — is required for its most demanding analytical work. But the architecture is not the framework. The framework is a set of structural commitments that can be stated without the architecture and used without reading the document that develops them. This section provides that compressed statement: not a summary, which would require knowing the document, but a minimum viable version, which requires only

grasping three ideas. Users who master these three ideas can apply the framework's orientation to real situations without the full system. Users who study the full document should be able to derive the three ideas independently. If these ideas are not recognizable as the core of the framework after reading the full document, the document has failed in its exposition. First idea: structural suffering. Most suffering that gets analyzed as individual failure, bad luck, or cultural deficiency is actually produced by institutional arrangements that could be designed differently. The question that reorients analysis is: who benefits from the arrangement that produces this suffering, and who designed it? Structural suffering is the answer to a structural question: what is this arrangement for, and who does it serve? An analyst who has internalized this idea will, when encountering any claim about why a population is suffering, automatically ask whether the explanation implicates the population's characteristics or the arrangement's design. That reflexive reorientation — from character to structure, from individual to institution, from accident to function — is the first idea operating. Second idea: institutional capture. Institutions created to solve problems consistently drift toward serving the interests of those with the resources to influence them, rather than the populations whose problems they were created to address. This is not conspiracy; it is the predictable outcome of structural incentives. The question that reorients analysis is: who has concentrated interest in this institution's decisions, and what has that concentration produced? Capture is identified not by finding evidence of bad actors but by finding the structural gap between what an institution claims to do and what it actually produces for the populations it claims to serve. An analyst who has internalized this idea will, when an institution fails, resist explanations that locate the failure in implementation and ask instead whether the institution was redirected before it failed — and by whom, and in whose interest. Third idea: floor and anti-foreclosure criterion. Evaluation of any arrangement requires two distinct questions. The floor question: does this arrangement prevent the forms of suffering most directly under institutional control — resource deprivation, institutional betrayal, visibility suppression? The anti-foreclosure question: does this arrangement expand or

contract the conditions under which people can develop beyond what survival alone requires? An arrangement can pass the floor test and fail the anti-foreclosure test. The combined question the framework evaluates against is: for all members of the relevant population, does this arrangement clear the floor and keep becoming open? An analyst who has internalized this idea will resist evaluations that aggregate outcomes — “overall welfare improved” — without asking whether the floor was maintained and becoming kept open for those whose position is structurally most constrained. These three ideas are related but not derivable from each other. Structural suffering identifies the problem. Institutional capture explains the persistence of the problem and the consistent failure of institutions to address it. Floor and anti-foreclosure provides the evaluative orientation for assessing whether any proposed arrangement actually addresses the problem. An analyst with all three ideas can apply the framework's core orientation without the formula system, the tension typology, the agent typology, or the specific doctrinal commitments the full document develops. The minimum viable version is three ideas, the orienting questions they generate, and the diagnostic reflexes they produce when the ideas become genuinely internalized rather than merely understood. The three-idea entry test: (1) Structural suffering — what arrangement is producing this suffering, who benefits from it, could it be designed differently? (2) Institutional capture — what institution claims to address this problem, who has concentrated interest in its decisions, what has that concentration produced? (3) Floor and anti-foreclosure — is the floor maintained for all members of the relevant population, does any institutional arrangement foreclose adaptive response capacity for any population? If these three questions produce useful orientation, the full framework adds precision but not a different direction. If they produce no orientation, that is diagnostic information about where the framework's scope limits are.

XVI: Technology as Structural Variable

16.0 — Technology as Constraint-Modifier, Not Tool

Technology is conventionally analyzed as a tool — an instrument deployed in pursuit of ends determined by other factors: interests, values, institutional arrangements. The tool framing correctly identifies that the same technology can produce radically different outcomes depending on the institutional context in which it is introduced, and it correctly resists technological determinism. But it is analytically inadequate for technologies whose primary effect is to change the constraint structure within which institutional and individual choice operates, rather than merely expanding the range of options within an unchanged constraint structure. For such technologies — the printing press, the steam engine, antibiotics, digital communication networks, artificial intelligence — the question of who controls them is not separable from the question of what the technology does, because the technology produces different structural outcomes depending on who controls it, and the answer to the control question is itself partly determined by the specific properties of the technology. The framework's layered constraint model provides the analytical basis for a more precise account. Technologies modify the constraint structure at one or more layers, changing what configurations are viable at other layers. A technology that reliably increases agricultural productivity modifies the biological floor layer — reducing the proportion of population required for subsistence, releasing demographic resources for other activities, changing what institutional configurations are viable at the economic and political layers. The modification is not deterministic: the same agricultural technology has supported feudal and capitalist organization alike, because the institutional layer has genuine causal properties that are not fully determined by the constraint layer below it. But the modification is also not neutral: some institutional configurations become more viable and some less viable as a consequence of the technological change, and this distribution tends to favor configurations that match the technical requirements of the new technology over configurations that do not. This account generates a specific analytical procedure. The primary questions are: which constraint layers does this technology modify, and in which direction? What institutional configurations does it make more viable and what configurations less viable? Who controls the technology during the transition period in which those viability shifts are occurring? What accountability structures govern that control? What capacity exists for reversing institutional configurations that prove harmful? These questions do not assume that technology determines outcomes. They assume that technology changes the terrain on which institutional choices are made, and that analyzing that terrain change is a prerequisite for analyzing what institutional configurations are possible and what transformations are feasible. Why the

tool framing fails: The tool framing of technology correctly identifies that technologies are not deterministic — the same technology can produce very different political and social outcomes depending on the institutional context in which it is introduced, and historical cases where the introduction of a major technology did not produce the consequences its proponents predicted are sufficiently numerous to establish that context shapes outcome in ways the tool framing cannot explain. Where the tool framing fails is in its treatment of major technologies as neutral with respect to institutional form. A hammer is genuinely neutral: it can build a house or break a skull, and the institutional context determines which use is predominant. Technologies that reorganize constraint structures — that change which organizational forms are viable, which surveillance is possible, which production configurations are economically rational, which concentrations of power are stable — are not neutral in this sense. They do not leave the institutional context that deploys them unchanged; they transform the constraint structure within which all institutional forms operate, which means they systematically advantage some forms over others regardless of which actors are trying to deploy them for which ends. The three constraint layers: PMN's layered model identifies three constraint layers that technology can modify. The biological constraint layer establishes the minimum conditions for physical functioning: technologies that modify this layer — medicine, food production,

reproductive technology, biotechnology — change what the biological floor requires and what it can sustain. The material constraint layer establishes the production relationships within which economic arrangements are organized: technologies that modify this layer — mechanization, electrification, digitization, automation — change which production configurations are economically viable and which organizational forms those configurations favor. The institutional constraint layer establishes the information processing, coordination, and enforcement capacities within which political arrangements operate: technologies that modify this layer — communication technologies, surveillance infrastructure, AI-mediated decision systems — change which institutional configurations are organizationally viable and which concentrations of power those configurations sustain. Each layer interacts with the others, and technologies that primarily target one layer consistently produce effects in the others through the causal chains that connect them. The capture problem as the central analytical issue: For every major technology that modifies the constraint structure, the analytically central question is not what the technology makes possible but which actors are positioned to capture the power-redistributing effects of the modification first and most completely. The first-mover advantage in technological transitions is not

merely commercial; it is structural. Actors who control the infrastructure through which a constraint-modifying technology operates have leverage over all other actors who use the infrastructure, regardless of whether those other actors are themselves sophisticated users. This leverage can be exercised directly — through pricing, access conditions, data extraction — or indirectly — through the information asymmetries that infrastructure control produces, and through the regulatory capture that actors with infrastructure control can pursue given their resource advantage. The constraint-modifier framing reframes the analytical question from "what does this technology make possible?" to "who will control what this technology makes possible, and what institutional arrangements does that control reproduce or transform?" The framework's evaluative position: technology is evaluated by what it does to the constraint structure within which the minimal anchor criterion operates. Technology that reduces the biological floor requirements — that makes it cheaper to provide adequate nutrition, healthcare, shelter — serves the minimal anchor directly by making it more achievable under existing institutional arrangements. Technology that concentrates control of constraint-layer infrastructure without reducing biological floor costs creates new vectors for capture and extraction that can increase structural suffering even as it produces aggregate efficiency gains. The evaluation cannot stop at the aggregate level; it must track who bears the costs of the new constraint structure and who captures its benefits.

16.1 — The Capture Problem at Technological Leverage Points

Every major technology that modifies the constraint structure creates a leverage point: a position in the institutional architecture from which disproportionate influence over the outcomes of many other actors can be exercised. The ownership and control of leverage points is the primary mechanism through which technological change produces or reproduces structural inequality rather than reducing it. The pattern is consistent across technological transitions: actors who control leverage points created by new technologies — owners of spinning mills in the industrial revolution, operators of railroad networks in the nineteenth century, controllers of oil production in the twentieth, owners of digital platform infrastructure in the twenty-first — systematically use that control to shape the institutional configurations that emerge from the technological transition in ways that preserve and amplify their leverage position. The capture dynamic at technological leverage points has a specific temporal structure that makes it particularly difficult to contest through normal political channels. The leverage point is created during the transition period, before the full institutional implications of the technology are apparent and before the populations who will bear the costs of capture have organized to contest it. The actors who move first to control the leverage point — typically those with pre-existing capital, organizational capacity, and political access — can use their early control to shape the regulatory, legal, and institutional frameworks within which the technology develops in ways that

entrench their position before effective opposition can form. By the time the structural consequences of capture are visible, the institutional framework has been stabilized around the capturing actors' position in ways that are very difficult to reverse. The implication for doctrinal positions on technological governance: the window for institutional intervention in technological leverage points is systematically early — before the technology has produced the concentration of control that makes intervention difficult — and the political pressure for intervention is systematically late — after the concentration has produced the visible harms that motivate opposition. Managing this temporal mismatch requires anticipatory regulatory institutions capable of identifying leverage points before concentration occurs, or pre-committed institutional rules — public ownership requirements, interoperability mandates, access obligations — that apply to any technology meeting specified criteria of network effect or infrastructure character, regardless of whether specific harms have yet been demonstrated. The leverage point structure: Not all positions within a technological infrastructure are equivalent. Leverage points are positions from which the outcomes of many other actors' activities can be influenced or constrained with minimal countervailing resistance. They arise from

specific structural features of technological systems: network effects (where the value of participation in a network depends on who else participates, giving control of the network's access conditions disproportionate leverage), data accumulation (where actors who accumulate data about others' behavior acquire information advantages that compound into structural leverage through superior prediction, targeting, and decision-making), infrastructure control (where actors who own the physical or institutional infrastructure through which others' activities are mediated can exercise terms-setting power that is difficult to resist without alternative infrastructure), and standard-setting (where actors who establish the technical or regulatory standards through which a technology is deployed shape all subsequent development in ways that systematically advantage their own position). The capture sequence at leverage points: Capture of technological leverage points follows a predictable sequence that varies in timing but is consistent in structure across technological transitions. In the first phase, actors with the greatest prior resource accumulation make the infrastructure investments that establish the leverage point — they can absorb the risk of investment before the technology's commercial viability is established. In the second phase, network effects, data accumulation, and infrastructure dependencies lock other actors into relationships with the leverage point that are difficult to exit. In the third phase, the leverage point generates resource advantages that are deployed into regulatory capture — shaping the rules that govern the technology in ways that entrench the position of the early infrastructure controllers and raise the costs of entry for potential competitors or alternatives. In the fourth phase, the regulatory framework reflects the interests of the captured leverage point rather than the interests of the populations whose activities are mediated through it, and the technology becomes a vector for structural extraction rather than for the efficiency gains that its proponents originally claimed. Why competitive markets do not prevent capture: The standard response to the capture analysis is that competitive markets will prevent leverage point concentration by creating incentives for competing infrastructure to develop. This response fails for three structural reasons. First, network effects, data accumulation, and infrastructure dependencies create winner-take-most dynamics in which early infrastructure controllers accumulate advantages faster than competitors can develop alternative infrastructure — which means the competitive window closes before competition can operate effectively. Second, the resources generated by early leverage point control are specifically available to purchase the regulatory advantages that prevent competitive entry — capture of the regulatory framework is financed by the leverage point itself. Third, the mos

t significant leverage points operate as natural monopolies or duopolies where the technical characteristics of the infrastructure make competitive provision inefficient in ways that competition cannot resolve — the resolution requires regulatory intervention, which is the infrastructure that capture is most directly aimed at controlling. The framework's analytical obligation: for every major technology whose constraint-modifying effects are being analyzed, identify the leverage points — the positions from which the terms on which others use the technology can be set. Then ask which actors are positioned to capture those leverage points and what

resources they have available to do so. The capture trajectory is not inevitable — counter-power organized before capture completes can prevent or limit it — but the capture trajectory is the default when counter-power is absent or inadequately organized. Diagnostic for leverage point capture in progress: (1) Who controls the infrastructure through which the technology's effects are mediated — is this control concentrated or distributed? (2) What data accumulation advantages are being generated by infrastructure position — who has access to behavioral data that others do not? (3) Is there evidence of regulatory capture in progress — are regulatory frameworks being shaped in ways that raise entry costs for competitors? (4) What is the counter-power capacity organized around the leverage point — is there organized political or institutional resistance to concentration, and is it moving faster or slower than the capture sequence? The AI governance gap and supranational capture: The analysis of capture at technological leverage points connects directly to the capture analysis at the supranational level developed in 9.0. AI governance is the current case where the gap between the speed of technological leverage point formation and the speed of institutional response is most consequential, and where the supranational dimension is most structurally exposed. The primary actors shaping AI development — a small number of private entities and two states — operate across jurisdictions that are too fragmented to individually constrain them, in a domain that is technically complex enough to generate the expertise asymmetry that enables Stage 1 and Stage 2 of the capture sequence before any national or supranational regulatory architecture has reached operational capacity. The international standards bodies, intergovernmental panels, and multilateral frameworks that would normally serve as the institutional response to supranational leverage concentration are still in the pre-formation stage for AI governance, which means that capture is occurring in advance of the institutions that would be captured. When those institutions form, their founding conditions — which states have expertise, which private actors have provided technical input, whose analytical frameworks have been institutionalized in the standard-setting process

— will already reflect the capture dynamics that operated in the pre-institutional period. This is the most difficult form of capture to resist: not the capture of an existing institution with residual accountability infrastructure, but the shaping of the founding conditions of institutions that do not yet exist. The strategic implication: counter-power at the AI governance level requires operating in the pre-institutional period, not waiting for the institutions to form before contesting their capture. This means producing independent technical analysis that cannot be monopolized by the primary developers, building the cross-border coalitions that will have standing in institutional formation processes, and establishing the accountability standards that founding documents must meet before they acquire the legitimacy that makes their terms stable. Pre-institutional counter-power is harder to organize than the post-institutional kind because there is no formal process to engage — but it is also the only kind that can influence the founding conditions that will otherwise be set by whoever has the most technical and financial resources at the moment of institutional formation.

16.2 — Artificial Intelligence as Institutional Architecture

Artificial intelligence systems warrant treatment as a distinct category because their primary effect is not on the physical or biological constraint layers but on the institutional layer itself. Prior technologies modified what institutional configurations were economically viable by changing cost structures. AI modifies what institutional configurations are organizationally viable by changing the information processing, decision-making, and coordination capacities available to institutions. A state that deploys AI-based surveillance and predictive policing does not merely have more efficient law enforcement — it has a qualitatively different relationship to the behavioral monitoring and anticipatory control of its population. A corporation that deploys AI-based hiring, credit scoring, and content moderation has systematized the discretionary decisions through which individual human judgment previously introduced variability — including the variability that allowed anomalous outcomes to contest established categories. The visibility

suppression analysis of section 1.11 identified algorithmic governance as operating on attention architecture — shaping what claims and experiences enter social circulation by modifying the information environment within which attention is allocated. This analysis requires extension to the full scope of AI's institutional effects. AI systems shape not only what information circulates but what decisions are made, by what criteria, with what accountability, and with what capacity for contestation. When consequential decisions about credit access, employment, housing, criminal sentencing, and welfare eligibility are made by algorithmic systems whose decision-making process is not interpretable by those affected, the accountability architecture of those institutional domains is qualitatively changed. The affected individual cannot contest a decision whose basis they cannot examine; the institutional actor cannot be held responsible for systematic patterns that emerge from statistical properties of training data rather than from identifiable intentional choices; and the political process cannot effectively regulate institutional behavior distributed across parameters in systems that resist human-readable description. The framework's evaluative criterion generates an assessment of AI's structural effects that is distinct from both the techno-optimist view (AI expands capability and therefore expands becoming) and the techno-pessimist view (AI threatens agency and therefore forecloses becoming). The evaluative criterion asks: does this technology, under the institutional arrangements in which it is currently deployed, expand or foreclose adaptive response capacity for the populations it affects? The answer is configuration-dependent. AI deployed under institutional arrangements that make its decision-making interpretable, provide effective contestation mechanisms, prevent concentration of AI capability in actors whose inter

ests diverge from affected populations, and invest automation efficiency gains in expanding floor conditions — such AI expands the material preconditions for becoming. AI deployed under arrangements that concentrate interpretive authority in unaccountable private actors, eliminate contestation mechanisms, automate the reproduction of existing structural inequalities at scale, and capture efficiency gains in already-privileged positions — such AI forecloses becoming at the same scale as its deployment. Why AI is categorically different: Previous technologies that modified the institutional constraint layer did so by changing the cost structure of specific institutional functions — communication technologies reduced coordination costs, surveillance technologies reduced enforcement costs, database technologies reduced information retrieval costs. AI is categorically different because it modifies not just the cost of specific functions but the viability of specific organizational forms. AI systems can perform the information processing, pattern recognition, and decision-making functions that previously required human organizational hierarchy — which means they can replace hierarchy with algorithmic management at scales and speeds that human organizational structures cannot match. This is not merely a productivity improvement within existing organizational forms; it is a change in which organizational forms are feasible and which concentrations of decision-making authority those forms can sustain. The implication: AI-mediated institutional architecture can concentrate effective decision-making authority in ways that formal organizational structures obscure, because the decision is nominally made by an algorithm while the effective decision is made by the actors who design, train, deploy, and update the algorithm. The accountability dissolution problem: Human institutional arrangements have developed, over centuries of political struggle, mechanisms for holding decision-making authority accountable: legal liability for decisions and their consequences, transparency requirements that make decision-making processes visible, electoral mechanisms that create incentives for decision-makers to consider the interests of those affected by their decisions, and professional norms that create internal accountability within organizational hierarchies. AI-mediated decision systems dissolve these accountability mechanisms without replacing them. Legal liability is difficult to attribute when decisions are made by systems whose reasoning cannot be fully reconstructed. Transparency requirements are technically challenging when the decision systems are genuinely opaque even to their designers. Electoral mechanisms cannot create incentives for accountability in AI systems that operate across jurisdictional boundaries and are not subject to political control. The result is that AI-mediated institutional architecture can produce the concentration of effective decision-mak

ing authority associated with the most powerful human institutional actors while eliminating the accountability mechanisms that were developed to constrain those actors. The three institutional transformation vectors: AI modifies institutional architecture through three specific vectors. The first is decision centralization: by making it feasible to process the information required for fine-grained decisions at massive scale, AI enables the centralization of decisions that were previously decentralized out of necessity — not because decentralization was preferred but because centralized processing was infeasible. The second is visibility asymmetry: AI systems accumulate information about those they interact with at rates and in dimensions that those being analyzed cannot reciprocate, producing structural information asymmetries that translate into leverage. The third is speed differential: AI-mediated processes operate at speeds that human organizational processes cannot match, which means that human accountability mechanisms that depend on human-speed oversight are systematically slower than the decisions they are supposed to constrain. The framework's evaluative question for AI: not whether AI produces efficiency gains — it does, in specific domains and under specific conditions — but whether the institutional architecture that AI makes viable serves the minimal anchor criterion better or worse than the alternatives it displaces. The answer requires tracking who controls the AI infrastructure, what decision-making authority is being centralized through it, what accountability mechanisms remain for the decisions it makes, and whether the efficiency gains it produces are distributed in ways that reduce structural suffering or captured in ways that increase it. Connection to the capture problem in 16.1: AI infrastructure is itself a technological leverage point. The actors who control the training data, computational infrastructure, and deployment platforms through which AI systems operate have leverage over the institutional arrangements those systems enable. The capture sequence that applies to technological leverage points in general applies to AI infrastructure specifically — and with particular political urgency because the AI-mediated institutional architecture being captured is itself the infrastructure through which other institutional arrangements are increasingly organized. Algorithmic Governance as Advanced V The V variable in the framework's formula architecture (15.0b) distinguishes between primitive V — external suppression through overt force, legal sanction, or material deprivation — and advanced V, the internalization of constraint through symbolic violence in Bourdieu's sense: where the dominated population participates in producing their own domination by applying to their situation categories of perception generated by the dominating arrangement. The canonical advanced V mechanism in contemporary institutional settings is the m

eritocracy narrative (8.4c): the institutional architecture that produces unequal outcomes presents those outcomes as the natural result of individual differences in capacity and effort, with the consequence that those who receive poor outcomes are systematically induced to locate the cause of their situation in themselves rather than in the arrangement. Algorithmic governance is advanced V's most technically sophisticated contemporary expression, and its analysis requires explicit connection to that variable because the mechanisms are structurally identical and the political consequences of misclassifying it are severe. When consequential decisions — credit access, employment, insurance pricing, criminal risk assessment, welfare eligibility, platform visibility — are made by algorithmic systems presented as objective statistical processes, the ideology of objectivity performs exactly the function that the meritocracy narrative performs in the credential economy: it locates the explanation for the outcome in properties of the individual being processed rather than in properties of the arrangement doing the processing. The person denied credit because a model trained on historically discriminatory lending data reproduces those patterns across a new population is told, in effect, that the denial reflects something true about them — their risk profile, their predicted behavior — rather than something structural about the data and the institutional interests that shaped it. This is symbolic violence with a technical interface: the dominated individual cannot contest a decision whose basis is opaque, is invited to accept it as objective, and is provided no institutional mechanism through which the structural character of the arrangement becomes visible as the source of the outcome. The distinction between primitive V and advanced V matters analytically because the counter-power strategies available differ structurally between them. Primitive V is politically visible: the

suppression is legible to those subject to it, the agent responsible for it is identifiable, and the counter-power response can be organized around the visible relation of force. Advanced V operates through the affected population's partial participation in their own constraint — the legitimacy of the arrangement is partly produced by those subject to it, because they have accepted, with varying degrees of consciousness, the categories through which the arrangement evaluates them. Algorithmic advanced V has additional properties that compound this. First, opacity: the statistical model's reasoning is genuinely difficult to reconstruct even for technically sophisticated actors, which means that the structural character of the decision cannot be demonstrated to the affected individual in a form that contests the objective appearance. Second, distributed authorship: no single human actor made the decision, which fragments accountability across the team that designed the model, the organi-

zation that deployed it, and the statistical properties of the training data — a distribution of authorship that parallels the distributed authorship of structural arrangements that produce E-type exclusion without intentional individual agents. Third, scalability: the same V mechanism deployed through a human workforce requires proportional human investment; deployed through an algorithmic system, it scales without proportional cost, enabling advanced V to operate at the scale of population-level processing that no prior implementation of symbolic violence could achieve. The political implication is that counter-power strategies designed for primitive V — legal challenge to identifiable decisions by identifiable agents — are structurally inadequate responses to algorithmic advanced V. Effective counter-power requires contestation at the architectural level: the training data, the decision criteria, the deployment authorization, and the interpretability infrastructure — which are the leverage points at which the structural character of the arrangement can be made politically visible and subjected to accountability mechanisms appropriate to its actual scale. The analysis of capture dynamics in 16.1 and the advanced V analysis converge on a structural pattern that is worth naming explicitly: algorithmic governance at scale constitutes an institutional form — not merely a tool deployed within existing institutional forms — in the sense relevant to the framework's institutional analysis. An institution, on PMN's account (7.0), is a structure that coordinates behavior through rules, norms, and enforcement mechanisms that persist across individual actors and generate reliable expectations. Algorithmic governance systems coordinate behavior — they determine what credit terms, employment opportunities, platform visibility, and welfare resources are available to specific individuals — through criteria that persist across operators, generate expectations for those subject to them, and enforce those criteria through automated decision processes. The actors who design and deploy these systems are, in the framework's terms, performing institutional design — with the distinctive feature that the institutional design is implemented in software and hosted on proprietary infrastructure, which means that the standard accountability mechanisms for institutional design (legislative authorization, judicial review, democratic deliberation) do not apply to it as a matter of automatic legal structure but must be affirmatively constructed and politically maintained. This is the specific sense in which AI is an institutional variable for PMN rather than merely a technology variable: the governance question is not primarily what can AI do but who designs the institutions AI implements, under what constraints, accountable to whom, and with what mechanisms for the populations whose becoming those institutions condition to contest the design. (See 15.0b, V advanced; 8.4c meri

tocracy narrative; 16.1 capture problem.) Diagnostic for algorithmic advanced V: (1) V type identification — is the structural harm being produced by the AI system operating through overt force or legal sanction (primitive V, amenable to standard legal contestation), or through the naturalization of outcomes as objective results of individual properties (advanced V, requiring architectural contestation)? The opacity and distributed authorship features of algorithmic systems make primitive V classification systematically misleading for decisions made through statistical models. (2) Opacity audit — can the affected individual reconstruct the basis of the decision in terms specific enough to contest it? If not, the decision operates as advanced V regardless of how it is formally classified. (3) Training data structural analysis — does the training data encode historical distributions that themselves reflect primitive V from prior periods? If yes, the

algorithmic system is automating the reproduction of prior V at scale and presenting it as statistical objectivity. (4) Counter-power strategy assessment — are the contestation mechanisms available to affected populations (individual legal challenge, complaints processes) scaled to the decision volume and structural character of the harm, or are they designed for primitive V at the individual case level? Mismatch between the scale and structural character of advanced algorithmic V and the counter-power instruments available for contesting it is the standard configuration — and its persistence is itself evidence of capture at the architectural level rather than oversight failure at the implementation level.

16.3 — Biotechnology and the Biological Floor Under Institutional Pressure

Biotechnology — encompassing genetic medicine, pharmaceutical development, reproductive technology, and emerging capabilities of direct biological modification — has a specific relationship to the framework's analysis because it operates primarily on the biological floor layer rather than on the institutional layer. The biological floor establishes minimum conditions for physical functioning that are preconditions for all activity at higher layers. Biotechnology is the domain in which those preconditions are themselves becoming subject to institutional decision-making rather than being fixed natural parameters. When the biological floor conditions of specific populations are partly determined by the institutional arrangements governing access to biotechnology, the floor is no longer a natural given that institutional analysis takes as background constraint. It is itself a variable that institutional analysis must address. The differential access problem is the most immediate structural consequence of biotechnology development under current institutional arrangements. The institutional arrangements governing pharmaceutical development — intellectual property regimes that grant temporary monopoly pricing power, regulatory frameworks designed for the institutional context of wealthy countries, distribution systems structured by purchasing power rather than medical need — systematically produce a situation in which the same biological conditions are survivable in some institutional contexts and fatal in others. This is not a natural fact about biology. It is an institutional determination of who has access to the biological floor conditions that biotechnology makes available. The framework's evaluative criterion identifies this as the most direct form of structural harm: the institutional foreclosure of biological floor conditions that are technically available but institutionally inaccessible. The enhancement trajectory raises the distinct set of analytical questions that section 3.4d addresses in relation to transhumanism and the subject boundary problem. The structural concern here is specific: enhancement technologies deployed under unequal institutional access conditions do not produce a general improvement in human biological capacity — they produce a stratified biology in which capacities available to individuals are partly determined by their position in existing institutional hierarchies. If cognitive enhancement, physical performance enhancement, and healthspan extension are available to those with sufficient resources and unavailable to those without, the consequence is not merely an expansion of inequality in material outcomes but an expansion of inequality in the biological capacities through which material outcomes are produced — a stratification more durable than economic inequality because it is embedded in bodies and transmitted through biological inheritance as well as social reproduction. Why biotechnology requires separate analysis

s: Most constraint-modifying technologies operate on the institutional or material constraint layers without directly modifying the biological constraint layer that establishes the framework's evaluative foundation. Biotechnology is different because it operates primarily on the biological layer itself — on the substrate from which the minimal anchor criterion is derived. This creates an analytical situation that no prior framework has had to address: when the technology modifies the foundation against which the technology's consequences are evaluated, the evaluation framework requires extension rather than simple application. The framework's response in 3.4d is to distinguish between biotechnological modifications that maintain the biological foundation's evaluative relevance (augmentation that preserves suffering-capable consciousness), modifications that alter the foundation in ways that require extension (hybrid transitions where biological

substrate is partially replaced), and modifications that would require reconceiving the foundation entirely (post-biological transitions where the relationship between the organism and the criteria derived from biological requirements becomes unclear). Biotechnology analysis under PMN requires first locating each specific technology on this spectrum before applying the standard evaluative framework. The equity dimension at the biological layer: Technologies that operate on the material or institutional layers produce equity questions about the distribution of their benefits: who captures the efficiency gains, who bears the costs of transition, who controls the leverage points. These are serious questions, but they operate against a background in which the biological floor itself remains constant — everyone faces the same biological requirements regardless of whether the technology's benefits are equitably distributed. Biotechnology that modifies the biological layer introduces a qualitatively different equity question: whether the biological modifications that the technology makes available are distributed in ways that create biological inequality. A population subset that has access to biological modifications that extend cognitive capacity, physical capability, or lifespan, while another population subset lacks such access, faces not only a disadvantage in the institutional competition for resources but a potential biological advantage that compounds across generations if the modifications are heritable. The framework's evaluative criterion — reduction of structural suffering, expansion of conditions for genuine development — applies directly: biological access inequality would constitute a new form of structural suffering and a new constraint on development capacity. Pharmaceutical and medical biotechnology — the current case: Contemporary pharmaceutical and medical biotechnology already raises these questions in their less extreme form. The development of treatments that are available to populations wi

th sufficient purchasing power and unavailable to those without it is a current instance of biotechnology operating to produce differential biological floor conditions across populations — the populations with access can maintain biological functioning at lower cost and with lower suffering than those without. The PMN evaluative framework identifies this as a structural suffering question: the gap between what existing biotechnology can provide and what it actually provides to specific populations is itself a measure of institutional failure to meet the minimal anchor criterion. The political economy of pharmaceutical development — where research incentives are calibrated to purchasing power rather than to burden of disease as experienced by those without purchasing power — is the capture problem applied to the biological floor. The framework's analytical position: biotechnology must be analyzed at three levels simultaneously. At the current level: how are existing biotechnologies distributed, and does that distribution reduce or increase the gap between what the biological floor requires and what specific populations receive? At the transitional level: what biotechnological modifications are becoming feasible, who will control access to them, and what is the trajectory of access distribution? At the horizon level: what modifications would be possible under the post-biological conditions identified in 3.4d, and how should the framework's evaluative criteria be extended to address them? These are not the same questions, and conflating them produces analytical errors in both directions — treating current access failures as though they were horizon questions deflects attention from remediable current inequities; treating horizon questions as though they were current policy questions generates speculative analysis that crowds out the more grounded analysis of current conditions. Diagnostic for biotechnology and the biological floor: (1) For any specific biotechnology, which constraint layer does it primarily operate on — biological, material, or institutional? (2) What is the current access distribution — is the technology available to populations in proportion to their need, or in proportion to their purchasing power? (3) What is the trajectory of access distribution — is it expanding toward the populations with greatest need or concentrating among those with greatest resources? (4) Does the biotechnology modify the biological substrate in ways that affect the evaluative framework's application — if so, which of the three categories in 3.4d is the modification closest to?

16.4 — Technology, Meaning Infrastructure, and the Acceleration Problem

The relationship between technological change and meaning infrastructure was addressed in section 5.5 (finitude as meaning infrastructure) and section 13.4d (the strategic equilibrium problem). A structural observation about the temporal dimension of that relationship requires explicit treatment here. Meaning infrastructure accumulates through practice over time — through repeated engagement with inherited traditions, social roles, and interpretive frameworks that gradually deepens their meaning-generating capacity. The rate at which technological change disrupts the institutional arrangements through which meaning infrastructure is produced and maintained is not independent of the meaning infrastructure problem. When the rate of change exceeds the rate at which populations can develop new meaning frameworks adequate to the new conditions, the result is the meaning-collapse configuration that section 4.3 identifies as among the most dangerous structural outcomes — not because the technological change is intrinsically harmful, but because the meaning infrastructure development work that would make the change navigable has not kept pace. The acceleration dynamic compounds this problem. Technologies that accelerate subsequent technological development — and artificial intelligence has this property to an unusual degree, since AI capability can be applied to developing further AI capability — produce rates of institutional change that systematically outpace social capacity for meaning infrastructure adaptation. A population that experiences a major technological disruption every generation has, in principle, enough time to develop meaning frameworks adequate to new conditions before the next disruption occurs. A population that experiences major technological disruptions annually does not have that time — the disruption of the previous cycle's meaning frameworks has not been resolved before the next cycle begins. The structural consequence is a population in persistent meaning-transition: never fully inhabiting any available meaning framework because every available framework is already being disrupted before it has been fully developed. The doctrinal implication: the governance of technological acceleration — not merely the governance of specific technologies but the governance of the rate at which new technologies are introduced into institutional and social life — is a legitimate subject for political analysis and institutional design. This is not an argument for halting technological development. It is an argument that the rate of deployment of technologies whose primary effect is to disrupt the institutional arrangements through which meaning infrastructure is produced should be treated as a variable subject to democratic governance rather than as a natural feature of technological progress that political institutions must simply adapt to. The actors who benefit from rapid deployment of disruptive technology have strong incentives to resist the

governance of deployment rate, and will characterize any such governance as opposition to progress rather than as management of a legitimate structural concern. Technology governance diagnostic: (1) Which constraint layers does this technology modify, and for which populations? Technologies that expand floor conditions for populations below the floor are evaluated differently from those that primarily expand capability for populations already above it. (2) What leverage points does this technology create, who currently controls them, and what accountability structures govern that control? Early capture of leverage points is the primary mechanism through which technological transitions reproduce rather than reduce structural inequality. (3) What meaning infrastructure disruption does this technology produce at what rate, and what parallel investment in meaning infrastructure reconstruction is occurring? Deployment rate that exceeds meaning adaptation capacity is a structural harm even when the technology's direct effects are beneficial. (4) What institutional arrangements govern access to the technology's benefits and the distribution of its costs? Technology that produces benefits concentrated in already-privileged positions while distributing costs across the broader population is a capture mechanism regardless of its technical character. The acceleration mismatch: Technological change — particularly in communication and information processing — now operates on timescales measured in years and sometimes months. Meaning infrastructure — the accumulated practices, traditions, interpretive frameworks, and relational structures through which individuals locate themselves in time and community — accumulates and erodes on generational timescales. When technological change proceeds at a rate significantly faster than meaning infrastructure can adapt, the result is a structural condition in which the material environment in which people live changes faster than the

frameworks they use to interpret that environment. This is not discomfort or nostalgia; it is a material condition with specific political consequences. The political consequences of meaning deficit: Populations experiencing accelerating meaning deficit exhibit predictable political patterns. The first is susceptibility to traditionalist counter-movements: movements that offer narrative frameworks for making the acceleration feel manageable by attributing it to identifiable enemies (globalists, cosmopolitan elites, immigrants) and promising to restore the pace of change to what meaning infrastructure can absorb. These movements are not primarily ideological in the conventional sense; they are meaning infrastructure repair attempts operating through political mobilization. The second is fragmentation of shared interpretive frameworks: when communities cannot develop shared meaning infrastructure fast enough to keep pace with shared material change, they develop incompatible interpretive frameworks

that make coordination across them increasingly difficult. This is the structural root of the polarization dynamics that political analysis frequently treats as merely cultural. The framework's analytical implication: assessing the political stability of a social configuration requires tracking not only the standard material variables but the ratio between the pace of material change and the adaptive capacity of the meaning infrastructure. High pace of material change combined with low meaning infrastructure adaptive capacity is a configuration that generates political instability through mechanisms that are not well captured by conventional economic or political analysis — which is precisely why it consistently produces analytical surprises. Connection to 5.6 (Meaning as Material Condition) and 13.4d (The Strategic Equilibrium Problem): this section specifies the temporal dimension of both. Section 5.5 establishes why meaning is a material condition rather than a secondary cultural epiphenomenon. Section 13.4d addresses the strategic choice between preserving existing meaning infrastructure and pursuing transformation that will degrade it. The acceleration problem adds the dynamic dimension: the relevant question is not only whether to transform but at what pace, with what investment in meaning infrastructure regeneration alongside material change.

16.4b — Effective Accelerationism and the Generational Sacrifice Problem

The acceleration analysis in this section describes a structural problem: the rate of technological change outpacing meaning infrastructure adaptation. There is a distinct ideological response to that structural problem that requires direct treatment, because it is analytically coherent enough to demand engagement and has achieved significant institutional influence in the technology sector. Effective accelerationism — the position that the appropriate response to the acceleration dynamic is to accelerate further and faster, trusting that the resulting technical capability will eventually produce conditions that make the current disruption worthwhile — is not a peripheral position. It is the operative assumption behind the deployment philosophy of several of the largest technology organizations currently shaping the material infrastructure of social and economic life. The position has a specific logical structure: the suffering produced by the current period of accelerated technological disruption (mass displacement from AI automation, erosion of the organizational infrastructure of stable employment, the meaning-collapse configuration this section identifies) is classified as transition cost rather than as structural harm — as collateral damage necessary for reaching a future state in which the technology's benefits are fully realized. This is a temporal displacement of the evaluative criterion: present suffering is assigned to the category of necessary investment in future adaptive response capacity, and the framework's evaluation is applied to the projected endpoint rather than to the actual conditions of the populations bearing the current cost. The anti-foreclosure criterion sharpens this critique. The accelerationist deployment model does not merely impose present floor violations while promising future conditions — it forecloses the adaptive response capacity of the populations bearing the current cost. Mass displacement from stable employment destroys the organizational infrastructure through which affected populations would otherwise be able to respond collectively to the disruption: the unions, professional associations, community structures, and institutional relationships that constitute response capacity. The transition imposed by e/acc therefore produces two simultaneous violations: an immediate

floor failure (unnecessary suffering now, preventable through different deployment choices) and an anti-foreclosure failure (destruction of the response capacity that would enable affected populations to manage the disruption). The claim that future generations will benefit does not resolve either violation, because (1) the populations bearing the cost are largely not the populations receiving the benefit, (2) no consent mechanism authorized that distribution, and (3) foreclosed response capacity cannot be restored by later abundance. PMN's evaluative criterion generates a specific falsification trigger for this temporal displacement. The mi

nimal anchor is not a criterion about future populations in future conditions — it is a criterion about present populations in present conditions. The person displaced from employment by AI automation in 2025 cannot eat the promise that AI-generated abundance will arrive by 2040. The criterion applies now, to the conditions of the people subject to the current arrangement, with the resources and institutional protections currently available. This is not a conservative argument against technological change — it is the application of the framework's non-negotiable constraint to arguments that would suspend it by temporal displacement. The asymmetry is structural: the populations bearing the costs of the transition are in large part not the populations who will disproportionately receive its benefits, and they have not consented to that distribution — it has been imposed by the deployment decisions of institutional actors whose primary accountability runs to investors rather than to the populations they are disrupting. A temporal argument that justifies present structural harm by reference to future benefit is only valid under conditions that accelerationism consistently fails to meet: that the populations bearing the present cost will share proportionally in the projected future benefit; that the projected benefit is sufficiently certain to justify the certainty of the present harm; that the future benefit cannot be achieved through paths that distribute the transition costs more evenly; and that the populations bearing the costs have genuine institutional standing to contest the arrangement and the timeline. Where these conditions are not met — and in the configuration of AI deployment prevailing at the time of writing, they are not met -- a doctrinal assessment of present conditions, revisable as those conditions change — the generational sacrifice argument is not a strategic choice about technological governance. It is an arrangement that produces absolute structural exclusion (E, 15.0b) of present populations from the calculus used to justify the conditions imposed on them. (See 4.5, 15.4 intergenerational variable G; 16.2 on algorithmic advanced V.)

16.5 — Reproductive Technology and the Commodification Problem

Section 16.3's treatment of biotechnology and the biological floor addresses the most acute structural implications of genetic and biomedical technology for human development capacity. This section addresses a related but analytically distinct domain: the technologies of reproduction itself — in vitro fertilization, gamete sale and donation, commercial surrogacy, embryo storage and selection — which constitute a sector where biological processes previously outside market logic have been progressively brought within it, with structural implications for women who provide reproductive labor, for offspring whose genetic origins are now marketized, and for the meaning infrastructure surrounding biological kinship and bodily integrity. The framework's general account of technology as structural variable (16.0) applies here, but the specific features of reproductive technology require explicit treatment because they combine in an unusual configuration: the commodification of processes occurring within human bodies, the systematically unequal distribution of who provides that labor and who purchases its products, and the absence of adequate institutional frameworks for evaluating structural effects that have been treated primarily as private contractual matters. Commercial surrogacy as currently practiced exhibits the structural features the framework uses to identify arrangements requiring structural rather than purely contractual analysis. First, the distribution of who provides gestational labor tracks class, national origin, and economic vulnerability with sufficient consistency to constitute a structural pattern rather than a series of independent market choices: women in economic precarity in lower-income countries or lower-income strata within high-income countries provide gestational services for

purchasers whose economic position is structurally superior. This is not an incidental correlation; it is the predictable outcome of a market in which the price of the service is set at levels making it financially attractive primarily to those for whom it represents a significant income opportunity relative to available alternatives — which by construction means those with constrained alternatives. The framework's account of the distinction between genuine preference and structurally constrained choice (13.4g, agency-structure tension in internalized norms) applies directly: a decision made within a structure that systematically narrows alternatives cannot be evaluated as free preference simply because no formal coercion is present. Second, the institutional infrastructure of commercial reproductive technology — agencies, fertility clinics, legal frameworks for contractual arrangements — is captured by actors whose interests are systematically misaligned with the interests of women who provide reproductive labor. Surrogacy agencies profit from arranging transactions in which they exercise definitional power over the terms — what constitutes a valid cont

ract, what health risks are disclosed, what happens when the arrangement produces unanticipated outcomes — without being subject to downward accountability to the party whose bodily risk is greatest. The capture sequence follows a recognizable pattern: a service that began as informal mutual aid within social networks was progressively formalized into a commercial sector, with formalization driven by actors who could capture the resulting institutions, producing an arrangement that serves purchasing parties and intermediaries while the structural position of those providing reproductive labor deteriorated relative to what informal arrangements had provided. The consistent pattern across jurisdictions — women providing gestational labor having weaker legal protections than those purchasing the service — is institutional capture's signature: the asymmetry in formal protection tracks the asymmetry in economic power that created the service in the first place. The global care chain analysis (8.2, 11.0 diagnostic 7) extends to reproductive technology in a form that makes the extraction mechanism legible: high-income households externalize reproductive labor — both the gestational labor of surrogacy and the gamete provision of egg donation — to women in lower-income positions, replicating at the level of biological reproduction the same structural dynamic that characterizes the externalization of domestic care labor analyzed elsewhere. The reproductive technology market is in this sense a continuation of the global care chain by technological means, with the added feature that the labor involved carries biological risk, and that institutional frameworks have consistently failed to reflect that risk in the distribution of returns. An analysis treating egg donation and commercial surrogacy as purely private contractual arrangements governed by consent has misidentified the unit of analysis: the relevant structural question is not whether an individual consented to a specific transaction but whether the structural conditions producing the market as a whole are compatible with the evaluative criterion. The doctrine implications are contested and the framework does not resolve the contestation, but it specifies the questions that must be answered. PMN does not take a position on whether commercial surrogacy or gamete markets are categorically impermissible; such categorical positions require prior commitments about bodily integrity that the framework's consequentialist orientation in doctrine does not supply. What the framework requires is that these arrangements be evaluated by the same criteria applied to any other institutional arrangement: whether they produce structural suffering that is imposed, dysfunctional, preventable, and generated by institutional design rather than inherent to the activity; whether they are captured by interests misaligned with those most affected; and whether alternatives are available that would meet reproductive needs without

t the structural extraction the current arrangements produce. On each criterion, the evidence from existing commercial reproductive technology markets suggests that current arrangements fail the diagnostic — not because reproduction should be categorically exempt from market logic, but because the specific markets that have developed consistently distribute benefits to purchasing parties and intermediaries while externalizing costs and risks onto providers whose structural position makes exit genuinely constrained. Connection to 8.2, 11.0 diagnostic 7, and 13.4g: the analysis of reproductive technology markets requires all three simultaneously. Section 8.2 provides the intersectional frame: the structural position of women providing reproductive labor is not determined by gender alone but by the intersection of gender with class,

national origin, and legal status that characterizes who actually provides these services. Diagnostic 7 provides the economic mechanism: the systematic non-recognition of biological risk-bearing as labor is the specific form that reproductive labor invisibility takes in this market. Section 13.4g provides the conceptual tool for evaluating consent claims: the framework does not dismiss consent but requires that the structural conditions producing a pattern of consent be themselves evaluated before the consent can be treated as resolving the structural question.

XVII: Cases and the Individual: Pattern, Drift, and the Demands of the Framework

17.0 — Why Cases and Why the Individual

A framework that can only analyze systems from the outside is not adequate to its own claims. PMN holds that the individual is not epiphenomenal — that micro-level decisions are co-constitutive of the structural conditions that produce macro-level outcomes, not merely shaped by them. If that is true, the framework must be able to speak to the level at which individuals actually make decisions: not only about institutional design and system dynamics, but about how to act, where to place energy, how to read one's own drift, and how to maintain the capacity to analyze honestly under conditions that consistently reward analysis that serves the analyzer's position. The cases in this Part are not formula applications. They are narrated patterns — accounts of how structural conditions produced outcomes through the decisions of actual actors at multiple levels simultaneously. Each case ends not with a formula summary but with an individual-level implication: what the pattern means for someone who holds this framework and is trying to act within conditions of the kind the case describes. The individual implications are not prescriptions. They are the questions the case puts to the person who has internalized the framework — the questions a framework that is genuinely adaptive must be able to generate about the situations its users actually inhabit. The cases are deliberately varied in scale, period, and domain. The structural patterns they illustrate — good intentions producing structural harm, capture through incremental drift, timing failures with durable consequences, the difference between winning a rule and changing a relation — recur across scales. A framework that can only see these patterns at the level of revolutions and regimes cannot help the person navigating the level of institutions, organizations, and daily decisions. The adequacy requirement: A framework that can diagnose systems from the outside but cannot speak to the situation of the individual actor within those systems is architecturally incomplete in a specific way: it cannot meet the standard that Part III establishes when it grounds the minimal anchor in the biological requirements of individual organisms. Those requirements are not experienced at the system level — they are experienced by specific individuals navigating specific institutional arrangements, making specific decisions with specific consequences for themselves and for others. A framework that addresses only the system level cannot specify what the individual who holds its analytical conclusions should do with them, which means it cannot complete the causal chain from analysis to the micro-level decisions that are co-constitutive of the structural conditions the analysis describes. What the case method provides: Historical cases serve three specific functions within the framework. First, they test the analytical framework against the complexity of actual historical dynamics, identifying where the structural analysis pr

edicts well and where it produces systematic misreadings. The four cases analyzed in Part XVII — prohibition, labor movement capture, land reform, and revolutionary capture — were selected because they represent canonical failure modes that the framework must be able to explain, not only describe. If the framework cannot explain why these patterns recur across different historical contexts and ideological projects, its claim to structural analytical adequacy is undermined. Second, cases specify the micro-level decision points at which individual actors can influence trajectory in ways that structural analysis alone cannot predict. Third, cases illustrate what the framework's evaluative criteria look like when applied to the complexity of actual situations — where the trade-offs are real, the information is incomplete, and the consequences of decisions are not determined in advance. Why the individual is not epiphenomenal: PMN holds that individuals are not epiphenomenal — that micro-level decisions are genuinely co-constitutive of structural conditions, not merely expressions of them. This is a structural claim, not a motivational one. It does not mean that individual goodwill or ethical commitment is sufficient to change structural conditions; the framework is explicit that good intentions operating within structural arrangements consistently produce

consequences that serve the arrangement rather than the intentions. What it means is that the specific decisions that specific individuals make at specific structural junctures — whether to drift with institutional pressure or resist it, whether to maintain analytical clarity or accommodate it to organizational necessity, whether to prioritize long-term structural work or short-term visible achievement — accumulate into the configurations that structural analysis then describes as path-dependent. The structure is made of decisions, even though no single decision determines the structure. The individual analysis in 17.7 follows from this structural claim: if individual decisions are co-constitutive of structural conditions, then the framework must specify what it asks of individuals who hold its analytical conclusions. Not compliance with a program — the framework does not prescribe a program — but a specific orientation toward their own position within institutional arrangements: awareness of drift mechanisms, capacity to apply the structural analysis to their own situation, and the kind of meaning infrastructure that sustains that analytical capacity over the periods of institutional pressure that structural work requires. The sequencing of Part XVII: the four historical cases establish the pattern before the individual analysis, because the individual analysis is only legible against the structural background the cases illustrate. An individual who understands the prohibition case — how moral clarity about one problem can produce structural arrangements that produce worse outcomes across multiple

dimensions — approaches their own institutional situation with a different analytical capacity than one who has not encountered that pattern. The cases are not illustrations of abstract principles; they are the analytical infrastructure that makes the individual analysis substantive rather than exhortatory. Connection to 17.8 (The Transmission Problem): the individual analysis and the cases together address the question of how the framework travels — how it is transmitted across time and across the political distance between the conditions under which it was produced and the conditions under which it is applied. The individual is the vehicle through which frameworks travel; the cases specify what the framework looks like in application; and the transmission analysis specifies what is required for the framework to be received in ways that preserve its analytical capacity rather than converting it into a set of fixed conclusions that can be held without the methodological commitments that make them valid.

17.1 — The Prohibition Case: When Moral Clarity Produces Structural Harm

The United States Prohibition of alcohol from 1920 to 1933 is one of the clearest documented cases of a policy that was analytically wrong in a specific and instructive way: it was morally coherent, democratically enacted, and structurally disastrous — not because its values were wrong, but because its causal model was wrong. The structural conditions it entered: Alcohol consumption in the United States in the early twentieth century was genuinely producing harm at scale: domestic violence, workplace accidents, economic drain on working-class households, and a saloon culture that was often entangled with machine politics and exploitation. The temperance movement's diagnosis that alcohol was causing suffering was not wrong. Its causal model — that prohibiting legal production and sale would reduce consumption and therefore reduce harm — was wrong in a specific structural way: it treated the demand as dependent on the supply infrastructure that existed, rather than as a variable that would reconstruct supply infrastructure to meet it. What the structural conditions actually produced: The destruction of the legal supply infrastructure did not reduce demand. It transferred the supply function to actors whose competitive advantage was precisely their ability to operate outside legal accountability — organized criminal networks that had been marginal before Prohibition became structurally central. These networks required corruption of enforcement to function, which systematically degraded the law enforcement infrastructure that the policy needed to succeed. The market for illegal alcohol was more profitable than the legal market had been, because the risk premium raised prices while demand remained inelastic. The institutional betrayal variable (B) rose sharply, because the state was visibly both prohibiting alcohol and taking bribes to allow its distribution — combining moral coercion with visible hypocrisy. The repeal logic: Prohibition was repealed not because the values behind it were abandoned but because the structural evidence had accumulated that the policy was producing the

opposite of what it intended across multiple variables simultaneously: organized crime was stronger, enforcement was more corrupt, alcohol was more dangerous (unregulated production), and the fiscal cost of a policy that generated no tax revenue while requiring enforcement expenditure was acute during the Depression. Repeal was a structural correction, not a moral capitulation. The pattern this case illustrates is not unique to Prohibition. It is the pattern of any intervention that is morally coherent but structurally naive — that correctly identifies a harm and incorrectly models the causal pathway through which eliminating the immediate source of the harm will reduce it. The harm is real; the intervention model is wrong; the outcome is a new configuration of harm that is in some respects worse than the original, combined with institutional damage that persists after the

policy ends. The structural lesson is not that moral clarity is dangerous or that well-intentioned policies should be avoided. It is that moral clarity about a harm does not substitute for structural analysis of how interventions propagate through the systems they enter. The question the framework requires before any intervention: what will actors who currently benefit from the arrangement the intervention targets do in response, and do the adaptive responses available to them produce outcomes better or worse than the original harm? Prohibition's architects did not ask this question rigorously. The criminal networks that answered it for them did. Individual-level diagnostic — Prohibition pattern: When you are designing or supporting an intervention against a genuine harm, ask: (1) What is the demand structure that produces this harm — is it dependent on the current supply infrastructure or would it reconstruct alternative infrastructure if the current one were removed? (2) Who benefits from the current arrangement in ways that give them resources and incentive to adapt? What adaptations are available to them? (3) What institutional infrastructure does your intervention require to function, and what are the conditions under which that infrastructure degrades? (4) What is the accountability structure of the intervention itself — what mechanism exists for detecting that the policy is producing the opposite of its stated intent, and what political conditions allow that evidence to produce correction rather than entrenchment? The analytical lesson — causal model vs. moral clarity: Prohibition's failure was a causal model failure, not a values failure. The moral argument against the harms of alcohol consumption was well-grounded: the documented harms were real, the populations bearing them were among the most vulnerable, and the democratic process that produced the policy was genuinely responsive to their interests. What the causal model got wrong was the relationship between the legal rule and the behavioral reality: it assumed that criminalizing production and sale would eliminate consumption without accounting for the supply-side reorganization that demand-driven prohibition would produce. The organized crime infrastructure that Prohibition created was not an unforeseeable consequence; it was the predictable structural response of a demand that the legal supply had been eliminated from. The analysis that would have predicted it was available in the existing literature on contraband markets. The PMN diagnostic: Applying the framework's diagnostic to Prohibition reveals why the causal model failed: the analysis of what the policy would do to the structural conditions governing alcohol supply was not conducted. The policy was designed entirely from the demand side — reducing consumption by eliminating legal access — without modeling what actors with interest in the existing supply would do when the legal supply was removed. This is the multi-level cons

equency failure that 12.1b identifies as a primary analytical error: treating first-order effects as the complete analysis while failing to trace what the first-order effects would produce in the social systems they disrupted. The second-order effect — organized criminal infrastructure — was larger than the first-order effect — reduced legal alcohol consumption — and produced harms that exceeded the harms the policy was designed to address. The case's relevance to contemporary policy: the Prohibition pattern — morally coherent policy with a wrong causal model producing harms larger than what it was designed to address — recurs across multiple policy domains. Drug prohibition, abstinence-only education, austerity responses to fiscal crises, trade protectionism — each has generated versions of this pattern. The lesson is not that values do not matter but that values without adequate causal modeling consistently produce outcomes that serve neither the values nor the populations the policy was designed to protect.

17.2 — The Labour Movement Case: Institutional Capture Through Incremental Drift

The history of organized labour in industrialized economies is a case study in the capture problem that 7.3 analyzes at the institutional level. Labour movements that were built through genuine counter-power accumulation — strikes, solidarity networks, mutual aid, and the organizational infrastructure that sustained them under conditions of real repression — transformed, over decades, into institutional actors whose survival came to depend on the same arrangements they were built to contest. The structural mechanism: The transformation was not produced by betrayal or ideological failure in the first instance. It was produced by success. Unions that achieved legal recognition gained access to formal bargaining processes that required ongoing institutional relationships with employers and the state. Maintaining those relationships required organizational stability — predictable leadership, financial continuity, legal standing. Each of these requirements created incentives that pulled against the organizational characteristics that had produced counter-power in the first place: disruption capacity, willingness to absorb repression costs, commitment to members whose interests conflicted with institutional survival. The drift sequence: Stage one: recognition and formalization. The union gains legal status and bargaining rights. This is a genuine victory. Stage two: institutional embeddedness. Maintaining bargaining access requires predictability — employers and state actors will not maintain relationships with organizations they cannot rely on to honor agreements. The union develops administrative infrastructure and professional staff whose careers depend on institutional continuity. Stage three: interest divergence. The professional apparatus of the union develops interests that partially diverge from the membership — in avoiding disruption that would threaten institutional relationships, in managing conflicts that would expose the union to legal risk, in prioritizing the sectors and members whose cooperation is most essential to institutional function. Stage four: capture completion. The union's decision-making is now shaped more by what the institutional relationships require than by what the membership's material conditions require. It continues to perform its original function in form while having substantively changed what it optimizes for. This sequence did not require corrupt leadership or ideological sellout. It required only that institutional survival constraints consistently shaped decisions over long periods. The actors who drifted were in most cases sincere — they believed they were making the best available choices within real constraints. The structural analysis is not about their intentions. It is about what the position's incentives produced over time, which is the same regardless of who occupied the positions. The Nordic exception: Scandinavian labour movements maintained more counter-power capacity over longer periods than their An

glo-American counterparts. The structural explanation is not primarily ideological — it involves: higher union density that maintained the material dependence of the institutional apparatus on the broad membership rather than a narrower core; stronger integration with political parties that maintained electoral accountability for labour outcomes; and the specific historical conjuncture of the 1930s in which Nordic labour movements negotiated foundational agreements from positions of genuine strength rather than institutional embeddedness. The exception confirms the structural analysis: where the conditions that produce drift were less present, drift was slower. The capture pattern in labour movements is the same pattern that 7.3c identifies in regulatory agencies, social movement organizations, and any institution that achieves its counter-power function through formal recognition. Recognition creates the institutional conditions that over time erode the counter-power capacity that made recognition worth pursuing. This is not a reason to avoid seeking recognition — it is a reason to design accountability mechanisms that function after recognition, not only before it. The specific mechanism that most consistently delays capture: maintaining the organizational infrastructure through which members can credibly threaten to withdraw from the institutional apparatus if it drifts. That threat is only credible if the organizational capacity to act on it exists. Organizations that demobilize after recognition have surrendered the only leverage that keeps the institutional apparatus accountable to its original function. Individual-level diagnostic — drift detection: For someone operating within an institution that was built for a counter-power function: (1) What does your

institutional position require you to prioritize that diverges from what the institution's original purpose would require? Name the specific trade-offs, not the general principle. (2) When did you last take an action that imposed real costs on the institutional relationships your position depends on, because the purpose required it? If you cannot remember, the drift question is live. (3) What is your honest assessment of whether the membership or constituency the institution claims to represent could, if they chose to, remove or significantly constrain the institutional apparatus? If the answer is no, the capture question is live. (4) What would you have to give up materially — income, status, relationships, career trajectory — to act fully in accordance with the institution's original purpose? The size of that gap is a structural measurement of how far drift has progressed. The structural mechanism of incremental drift: The labour movement cases share a specific structural mechanism that the incremental drift analysis in 10.12 identifies at the general level. As organizations acquire resources and institutional recognition, the actors within them develop material stakes in the organizati

on's survival that become separable from the political objectives for which the organization was built. Legal recognition creates regulatory relationships with the state that give the state leverage over the organization's internal governance. Collective bargaining agreements create contracts that bind the organization to the management relationship the movement was built to contest. Pension funds, insurance programs, and mutual aid infrastructure create administrative organizations whose survival becomes a political objective in its own right. Each step in institutional development is individually justified by the expanded capacity it provides; the cumulative effect is an organization whose survival interests have become misaligned with its original political mission. The diagnostic application: the labour movement case is most useful not as a cautionary tale about a past failure but as a template for identifying the specific institutional forms through which capture operates in current counter-power organizations. Which contractual, regulatory, or financial dependencies has the organization developed that give external actors leverage over its internal decisions? What material interests have its leadership developed that are separable from the political interests of the populations it represents? These questions, applied to any current counter-power organization, will identify the specific capture vectors that the historical labour movement cases illustrate at the general level.

17.3 — The Land Reform Case: Winning the Rule Without Changing the Relation

Land reform — the redistribution of agricultural land from large landholders to smallholders or landless peasants — has been attempted in dozens of countries across the twentieth century, with dramatically varied outcomes. The cases where formal redistribution did not produce the material change it was designed to produce are analytically more important than the successes, because they reveal the gap between changing a legal rule and changing the structural relations the rule was embedded in. The Philippines case: The Philippines implemented multiple land reform programs across the twentieth century — under Magsaysay in the 1950s, under Marcos under martial law, and under the Comprehensive Agrarian Reform Program beginning in 1988. Each produced significant formal redistribution on paper. Each produced substantially less material redistribution in practice. The structural explanation is consistent across programs: the institutional infrastructure required to make formal redistribution materially real — credit access for new smallholders, extension services, protection from predatory contract arrangements, enforcement of land title against economic pressure — was either absent or captured by the same interests the reform was designed to displace. The mechanism of formal-material divergence: Landlords who lost formal title retained structural power through alternative mechanisms. Credit: smallholders who received land but had no capital to work it independently were dependent on credit from actors whose interests aligned with reconcentration. Contract arrangements: formal land ownership compatible with contractual arrangements that transferred economic control back to larger actors. Political leverage: landed interests in the Philippines maintained sufficient political access to shape the implementation infrastructure of the reform — the agencies responsible for valuation, title transfer, and enforcement — so that the formal rule operated in ways that served their

interests in reconcentration. The reform changed the legal form of the relationship; it did not change the underlying material relations that determined who could extract from whom. The Taiwan and South Korea contrast: Both countries implemented land reform in the late 1940s and early 1950s with substantially more material success. The structural explanation involves a specific historical conjuncture: both programs were implemented by governments that were externally imposed or had recently displaced the landed class that would otherwise have captured the implementation infrastructure. The Nationalist government in Taiwan had no pre-existing relationship with Taiwan's landlord class; the post-war South Korean government implemented reform under US pressure in a context where the Japanese colonial landowning class had been dispossessed. In both cases, the implementing government had fewer capture vulnerabilities than the Philippines's domestically-rooted political system, and the instituti

onal infrastructure of redistribution was therefore less subject to reversal through the implementation process. The general pattern: formal rule change without parallel construction of the institutional infrastructure that makes the rule materially effective consistently produces the appearance of change while the underlying structural relations reassert themselves through the channels the formal change did not address. This is the mechanism behind the PMN claim that reforms that do not account for the institutional infrastructure of their own implementation are systematically vulnerable to producing legitimacy cover for the arrangement they were supposed to change. The individual-level implication of this pattern is not only for policy designers. It applies to anyone working within an institution on a change initiative: winning the formal rule — the policy approval, the organizational mandate, the stated commitment — is not the same as producing the change the rule was designed to produce. The question after any formal victory is: what is the institutional infrastructure required to make this materially real, who currently controls that infrastructure, and what do they have to gain or lose from the implementation succeeding? Answering these questions after the formal victory is too late to influence the answer. They must be part of the design of the initiative before the formal victory is pursued. Individual-level diagnostic — formal vs. material change: (1) Identify the formal change you are pursuing or have achieved. (2) Map the institutional infrastructure required to make that formal change produce material outcomes — not the policy itself, but the implementation chain. (3) For each element of that chain, identify who controls it and what their material interest in implementation success or failure is. (4) Where the controllers of implementation infrastructure have interests that diverge from implementation success, what mechanism constrains them? If no mechanism exists or is being built, the formal change is likely to produce legitimacy cover rather than material change. (5) What is the minimum institutional infrastructure that must exist before the formal change is worth pursuing? Building infrastructure before winning the rule is often more consequential than winning the rule before building infrastructure. The structural explanation for the pattern: The cases where formal land redistribution did not produce genuine redistributive outcomes share a structural feature: the legal change in ownership was introduced without the institutional infrastructure that converts ownership into productive capacity. Access to credit calibrated to smallholder production, to technical knowledge relevant to smallholder agricultural conditions, to market infrastructure that does not require aggregation through the former landholder class, and to the legal enforcement capacity that makes the redistributed title meaningful against informal pressure from dispo

ssessed former owners — these infrastructure components were absent in the high-failure cases and present in the high-success cases. The pattern suggests that the legal rule was the easy part: changing legal titles is administratively straightforward relative to building the institutional infrastructure that makes the titles economically consequential. The PMN diagnostic applied: the land reform pattern illustrates what 7.3 identifies as the custodian problem at the implementation level. The state actors responsible for building the institutional infrastructure that would make land redistribution consequential were often the same actors whose social networks and professional relationships connected them to the landholder class whose position the reform was designed to disrupt. The infrastructure was not built not because the resources were unavailable but because the actors responsible for building it faced structural incentives to leave it unbuilt.

This is not a story about individual corruption; it is a story about the predictable consequences of assigning transformative institutional responsibility to actors whose structural position gives them interests in preventing the transformation.

17.4 — The Revolutionary Capture Case: When the Custodian Becomes What It Displaced

The pattern of revolutionary movements that successfully displace existing power arrangements and then reconstruct equivalent structural conditions under new organizational forms is sufficiently consistent across cases to constitute a structural regularity rather than a sequence of individual failures. France 1789–1799, Russia 1917–1928, China 1949–1966, Iran 1979 — these are not cases of ideological corruption or leadership betrayal in the primary instance. They are cases of the custodian problem (7.3) operating at the highest stakes available. The structural sequence: Stage one: the existing arrangement loses legitimacy across multiple dimensions simultaneously — performance, procedural, and narrative — and the counter-power capacity of the challenging movement crosses threshold. Stage two: transition. The challenging movement takes control of the institutional apparatus of the state. This is a genuine structural change: the actors who occupy the positions have different origins, different ideologies, and in many cases genuinely different intentions from their predecessors. Stage three: the positions' incentives reassert. The positions that were occupied by the old regime now generate the same structural pressures on their new occupants. Security apparatus logic requires concentration of information and coercive capacity. Economic administration logic requires prioritization of output stability over distributional equity. Political survival logic requires management of internal competitors. Stage four: the movement's original organizational infrastructure — the counter-power capacity that produced the transition — becomes a threat to the new institutional arrangement rather than its foundation. It is dissolved, captured, or suppressed. The cosmological capture dimension: What makes this pattern particularly resistant to internal correction is the mechanism identified in 7.3c. The actors who implement stage four — the suppression of the movement's original organizational infrastructure — are typically not doing so cynically. They believe, and are often correct, that the specific organizations and individuals being suppressed represent genuine threats to what the revolution was supposed to achieve. The suppression is authorized by the revolutionary purpose: the party must be protected because the party is the revolution. This is cosmological capture at its most dangerous — the belief system that designates the suffering produced by the institutional arrangement as necessary for the arrangement's protective function is the same belief system that the arrangement was built to express. The historical cases where this sequence was partially interrupted share a structural feature: the existence of distributed counter-power infrastructure that was not fully controlled by the revolutionary movement itself, and that maintained sufficient autonomy to generate genuine accountability pressure on the new institutional arrangement. The Spanish anarchist feder

ations in the early Republic, the Soviets before their subordination to the party, the workers' councils in the Hungarian 1956 — these were not sufficient to prevent the pattern, but they delayed it and in some cases produced genuine variation in outcomes. The implication is not that distributed infrastructure prevents the sequence, but that it is the structural variable most capable of slowing or disrupting it. The relevance for anyone working within a transformative political project: the custodian problem does not wait until after the transition. It operates during the building of counter-power capacity, during the transition itself, and immediately after. The organizational infrastructure of a movement that successfully captures institutional power will face the same pressures that those institutions generated for their previous occupants — typically faster and at higher stakes than the previous occupants faced, because the new occupants are also managing the expectations of a constituency that believed the transition would produce different structural conditions. Building accountability mechanisms that function after success is not pessimism about the project. It is the application of the framework's own structural analysis to the project itself. Individual-level diagnostic — custodian problem in transformative projects: (1) In the project or movement you are part of: who currently

has the authority to determine what the project's purpose requires? (2) What mechanism exists for people who disagree with that determination to act on their disagreement without being classified as threats to the project? (3) What would it look like if the project's leadership were experiencing cosmological capture — unable to perceive their own drift because their belief system designates the drift as necessary? How would you distinguish this from genuine strategic disagreement? (4) What organizational infrastructure exists outside the central leadership that maintains genuine accountability leverage over it? If none, the custodian problem is structurally unaddressed. (5) What are you personally willing to do if the project you are committed to begins producing the pattern it was built to prevent? Having an answer to this question before it becomes necessary is part of what it means to hold the framework honestly. Why the pattern is structural rather than biographical: The revolutionary capture pattern has been explained as a product of specific leaders' authoritarian personalities, ideological rigidity, or betrayal of original principles. These biographical explanations are insufficient because the pattern recurs across leaders with substantially different psychological profiles and ideological orientations. The structural explanation — that the custodian problem analyzed in 7.3 applies to revolutionary organizations as fully as to any other institution — is more parsimonious: organizations that successfully seize state power face the same structural incentives

as the states they displaced, with the additional burden of having destroyed the external accountability mechanisms that constrained the prior arrangement without having built replacement accountability infrastructure. The post-revolutionary state is structurally less constrained than its predecessor in the short term, which means the structural incentives for capture operate with less countervailing resistance. The PMN diagnostic: what distinguishes the small number of revolutionary transitions that did not reproduce equivalent structural arrangements from those that did is not primarily ideological — it is institutional. Transitions that maintained genuine institutional pluralism within the new arrangement, that built accountability infrastructure alongside the consolidation of state power, and that preserved counter-power organizations capable of challenging the new state from outside it produced different outcomes than those that treated pluralism as a transitional concession to be eliminated once state power was secured. The lesson is not that revolution is wrong but that the institutional design of the post-revolutionary state is the analytically central question that revolutionary politics consistently underweights in favor of the question of how to seize state power.

17.5 — The Platform Governance Case: When Infrastructure Becomes Speech

The standard framing of speech governance presents a binary: either all speech must be protected from institutional interference, or harmful speech must be restricted by accountable authority. This binary structurally excludes the most consequential question, which is neither about protection nor restriction but about the institutional architecture through which speech circulates. Who controls the infrastructure through which claims travel determines which claims reach audiences capable of acting on them, at what velocity, amplified by what mechanisms, and subject to what contestation. The protection vs. restriction binary treats this infrastructure as neutral and asks only what content it should or should not carry. PMN's analysis of epistemic infrastructure (1.12) and visibility suppression (15.3) identifies the infrastructure itself as the primary structural variable. The structural pattern: digital platform architectures concentrate the infrastructure of public discourse in a small number of private actors whose revenue depends on engagement rather than on accurate reception. This creates a systematic incentive misalignment between what the infrastructure optimizes for — sustained attention, emotional activation, return visits — and what epistemic infrastructure requires for its public function — accurate distribution of true claims, correction of false ones, and the production of shared understanding adequate for coordinated action. This is not primarily a problem of bad content on good infrastructure. It is a problem of infrastructure designed for a purpose (engagement monetization) that structurally degrades the epistemic functions public discourse requires. The distinction matters for intervention: content moderation addresses symptoms while leaving the incentive architecture

that produces them intact. The structural intervention targets the architecture itself — the ownership structure, the revenue model, the algorithmic amplification mechanisms — not the content those mechanisms carry. The case also illustrates scale as a structural variable (7.0c) operating on epistemic infrastructure. Broadcast media at national scale produced concentration of epistemic authority in a manageable number of institutional actors whose conduct could, in principle, be regulated and whose accountability could be structured through licensing, ownership limits, and content standards. Platform media at global scale produces a concentration whose reach exceeds the jurisdiction of any single state while whose infrastructure is owned by a small number of private actors whose accountability structures are self-defined. The accountability mechanisms adequate to national broadcast do not transfer to global platform infrastructure, not because accountability is impossible but because the mechanisms required at that scale have not been built. The failure is one of institutional design lagging behind scale change — a condition that 13.1d identifies as a mechanism by wh

ich tensions become contradictions through scale change rather than through mismanagement. The PMN diagnostic applied to platform governance: the primary question is not what speech to permit or restrict but what institutional arrangements govern the infrastructure through which speech circulates. Who owns it? What incentive structure does the ownership produce? What accountability mechanisms exist for decisions about amplification, suppression, and curation? Are those mechanisms accessible to the populations most affected by infrastructure decisions, or only to the actors with sufficient resources to engage the platform's formal processes? The V-suppression mechanism in the structural suffering formula is directly at stake: platform architecture that systematically amplifies individualized attributions of hardship while suppressing structurally-framed accounts — not through explicit content decisions but through the engagement mechanics that determine what travels — functions as a visibility suppression mechanism regardless of whether any specific content has been restricted. A governance framework that treats this as a speech problem rather than an infrastructure problem will consistently address the wrong variable.

17.6 — The Nationalism Case: Counter-Power Vehicle or Exclusionary Mechanism?

The standard opposition between globalism and nationalism reproduces the binary-framing problem that PMN identifies across political analysis: it treats a structural variable as an ideological label, asking which side to endorse rather than what each produces under which conditions. Nationalism is not a fixed type. It is a mobilization strategy that uses the category of national community to organize collective action, and the structural consequences of that mobilization depend entirely on what the national community is being organized against, what institutional arrangements it is being used to contest or defend, and whose material interests the mobilization actually serves. PMN's evaluative criterion applies to nationalism the same way it applies to any other political arrangement: does this specific instantiation reduce structural suffering and expand becoming conditions for the populations it claims to represent, or does it suppress internal difference, redirect structural grievances toward external enemies, and concentrate power in the hands of those who control the national narrative? Two structural configurations produce predictably different outcomes. In the counter-power configuration, nationalist mobilization organizes a population subjected to external extraction — colonial administration, foreign-controlled resource extraction, international financial arrangements that foreclose local developmental capacity — against the specific institutional arrangements producing that extraction. The national community functions as the organizational vehicle for resistance that has no other available form: class organization is suppressed, civil society is restricted, and the only legitimate collective identity that can be mobilized is the national one. In this configuration, nationalist politics can reduce structural suffering and expand the conditions for local development, not because nationalism as such is progressive, but because it is the available counter-power form in that material configuration. The anti-colonial nationalisms of the twentieth century operated in this configuration. Their record is not uniform — many reproduced capture under national form immediately after independence — but the structural

function was counter-power, and PMN's analysis of political traditions (10.7) must account for what they were doing, not only for what they became. In the exclusionary configuration, nationalist mobilization organizes a population experiencing genuine structural suffering — real resource deprivation, real institutional betrayal — but redirects the political energy of that suffering toward an internal or external group defined as the source of the problem. The structural suffering is real; the causal attribution is manufactured. This is the configuration in which nationalism consistently serves the interests of the domestic elites who benefit from the existing arrangement: by directing transformation pressure outward or downward toward a d

esignated enemy, it defuses the pressure that would otherwise accumulate against the institutional arrangements producing the suffering. The governing narrative of the exclusionary configuration performs V-suppression through a specific mechanism: it makes structural suffering visible as group harm while suppressing the structural account of what is causing it, substituting an enemy narrative for an institutional analysis. The affected population experiences genuine grievance that is real and politically mobilized — and directed toward targets that, if defeated, would leave the structural sources of their suffering entirely intact. The PMN diagnostic: applying the evaluative criterion to a specific nationalist movement requires identifying which configuration it is operating in. The diagnostic questions are structural, not ideological. Who benefits materially from the mobilization — the populations whose suffering is being named, or the elites whose position is secured by directing that suffering toward an enemy? What is the institutional target of the mobilization — the specific arrangements producing structural suffering, or a group defined as the source of that suffering? What happens to the structural suffering if the nationalist program succeeds — do the arrangements generating it change, or do they continue under new management? And critically: does the movement's internal organization suppress the difference and dissent within the national community in ways that reproduce the accountability failures it claims to be contesting? A nationalism that wins and immediately suppresses internal counter-power has not reduced structural suffering; it has changed the personnel operating the arrangement while preserving the arrangement's essential features. This is the revolutionary capture problem (17.4) operating under nationalist rather than revolutionary form. The form differs; the structural dynamic is identical.

17.7 — PMN at the Level of the Individual: What the Framework Asks of the Person Who Holds It

The preceding cases are not only historical. They are patterns that operate at every scale, including the scale of the individual within an institution, a movement, a career, or a relationship. The framework's claim that micro-level decisions are co-constitutive of structural outcomes is not only an analytical claim — it is a claim about where the framework must be applied if it is to be used honestly. A framework that analyzes institutional capture but cannot help the individual detect their own drift is not doing what it claims. What follows is not a code of conduct or a set of rules. PMN does not produce deontological prescriptions. It produces a set of questions that a person who holds the framework honestly must be able to ask of themselves — and must be able to answer, even when the answers are inconvenient. The questions are organized around the central challenges that the cases in this Part identify: moral clarity without structural analysis, incremental drift that is invisible to the drifting person, formal victories that do not produce material change, and the custodian problem in its individual form.

17.7b — On Acting Within Institutions

The question that the institutional cases put to the individual is the inside-outside problem: when does working within an institutional arrangement serve the arrangement's original function, and when does it serve the arrangement's self-continuation at the expense of that function? This question does not have a universal

answer — it has a structural one, which means it requires diagnosis of the specific configuration rather than a general rule. The inside-outside diagnostic: Three conditions indicate that working inside an institution is serving its function: (1) The institution is producing material outcomes that move in the direction of the minimal anchor — it is reducing unnecessary suffering or preserving adaptive response capacity for affected populations — even imperfectly and slowly. (2) Your presence within it is changing outcomes at the margin relative to what would occur without your presence. (3) The cost of your presence — what you must concede, suppress, or perform to maintain it — does not exceed the marginal outcome difference you are producing. When any of the three conditions fails, the inside position is not serving the function. The analysis is uncomfortable because it requires honest assessment of marginal impact, which most people consistently overestimate when their material interests are tied to the position. The exit diagnostic: Exit from an institutional position is not self-evidently righteous. The question is not whether you are compromised by your institutional position — everyone operating within institutions that do not fully align with the framework's evaluative criterion is compromised by some margin. The question is whether exit serves the function better than continued presence, accounting for what exit produces: who replaces you, what signal exit sends, what alternative capacity exit enables, and what the cost of exit is to people whose material conditions depend on the institution functioning. Exit that is primarily about the exiting person's moral comfort — about not being implicated — is not the same as exit that serves the function. The framework's position on the inside-outside problem is not that inside is always capture and outside is always integrity. It is that the question must be answered structurally, not morally — by assessing what each position produces in terms of material outcomes for the populations the institution affects, not by assessing what each position means for the identity or purity of the person occupying it. A person who exits a corrupt institution and then produces nothing is not more aligned with the framework than a person who stays inside and produces marginal improvements. The framework evaluates by outcomes, not by positional purity.

17.7c — On Detecting Your Own Drift

The labour movement case demonstrates that drift does not feel like drift from inside the drifting organization. Each decision in the drift sequence was individually defensible; the cumulative trajectory was not visible in the moment of each decision. The same structure applies to individuals. Drift from the framework's evaluative criterion does not typically feel like abandonment — it feels like pragmatic adaptation, strategic patience, or realistic assessment of what is possible. The drift self-diagnostic: The most reliable internal indicator of drift is not the content of your positions but the structure of your reasoning. Drift tends to produce a specific change in reasoning structure: the considerations that are generated spontaneously shift from what the situation requires to what the situation allows given your current institutional position. The question is no longer 'what does the material analysis indicate?' but 'what does the material analysis indicate that I can act on without threatening the relationships I depend on?' Both questions produce real analysis, but the second has already incorporated a constraint that the framework's evaluative criterion does not authorize. The external indicators: Drift that is invisible from inside is often visible from outside. Three external signals are most reliable: (1) People whose material interests most conflict with the positions you hold are no longer threatened by your analysis — they have learned that your analysis will not produce action that costs them. (2) People who once looked to your work as a counter-power resource have stopped doing so without explanation. (3) You find yourself spending more analytical energy on why change is difficult or impossible than on what change would require and how to produce it. None of these signals is definitive alone. In combination they indicate that the diagnostic procedure in 13.1c should be applied — to your own position, with honesty about who benefits from the reading you are likely to give it. The framework does not require that individuals operating under material constraints act as if those constraints do not exist. It requires that they name the constraints honestly rather than converting them into analytical conclusions. There is a structural difference between 'I believe this change is not currently possible given the constraints I face' and 'this change is not the right approach.' The first is honest about the source of the judgment. The second has

laundered a material constraint into an analytical position. Laundering is the individual-level form of the same epistemic failure the framework criticizes at the institutional level.

17.7c-ii — From Diagnostic to Practice: Drift Detection as Technology of the Self

The diagnostics in this section — the reasoning structure check, the three external signals — are reactive instruments: they identify drift that has already occurred. A reactive instrument applied after drift has become visible is applied too late. The structural analysis of drift mechanisms in 10.12 establishes that drift is most detectable, and most reversible, in its early stages, before it has been incorporated into the actor's self-understanding as a series of reasonable accommodations. The individual question that follows is not only how to detect drift but what practice, built into ordinary activity rather than reserved for moments of crisis, maintains the conditions under which drift is caught while it is still early. This is the shift from diagnostic to practice: from instrument of crisis response to what Foucault's account of technologies of the self addresses at the structural level — the practices through which individuals actively constitute and maintain their relationship to their own judgment. The distinction that matters for PMN's purposes is between discipline and self-formation. Discipline is the exercise of an external normative standard on the individual's conduct — which would reproduce, at the level of individual practice, precisely the normative architecture the framework is designed to interrogate. Self-formation is the active maintenance by the individual of a capacity: in this case, the capacity for honest analysis that does not stop at the boundary of the analyst's own institutional position. That capacity is not a fixed property the individual either possesses or lacks. It is a practice that is either maintained or allowed to degrade — and like any practice, it degrades without deliberate maintenance regardless of intention. The practice has three components that must be built into ordinary analytical activity rather than reserved for self-examination under pressure. First, structural accounting at regular intervals: calibrated to the pace of the individual's institutional environment — quarterly in stable positions, monthly in rapidly shifting ones — the analyst applies the reasoning structure check not to a current crisis but to a representative sample of recent conclusions, asking for each whether the boundary of the analysis was set by what the evidence required or by what the position allowed. Second, standing relationship maintenance: the solidarity analysis in 17.7d is not an emergency resource to be activated when crisis arrives but a standing practice — regularly exposing analysis to people who have sufficient independence from one's own institutional position to identify drift without personal cost. Relationships maintained only instrumentally for crisis response are relationships that will not produce honest feedback in the moment when honest feedback is most needed and most threatening to maintain. Third, contemporaneous record: writing down analytical positions as they are held, not as they are remembered

after the fact. The narrative of drift is always more coherent in retrospect than the experience of it was in sequence — each accommodation appears to have been rational given the conditions, and the trajectory appears as a series of reasonable adjustments rather than as cumulative departure. Written records at the time of the conclusion disrupt that retrospective coherence by preserving what was actually concluded before the accommodations that followed gave it a different meaning. The record does not prevent drift; it makes the trajectory visible across time in a form that is not subject to retrospective narrative smoothing. These three components together constitute the practice rather than the instrument — the ongoing care of one's own analytical capacity that the framework's demands on the individual require but cannot substitute for. Individual drift detection matrix (for regular use — not crisis application). (1) Reasoning structure audit: review five recent analytical conclusions. For each, identify the boundary at which the analysis stopped. Was that boundary set by what the evidence indicated, or by what was actionable given current institutional position and relationships? Mark conclusions where the boundary was set by position rather than evidence. If more than two of five are position-bounded, drift is probable and the trajectory review (step 4) should be applied immediately. (2) External signal check: have actors whose material interests conflict with your

current positions reduced their level of active opposition without your positions having changed? Have collaborators who previously cited your analysis stopped doing so without explanation? Has the ratio of analysis devoted to why change is difficult or impossible versus what change requires and how to produce it shifted across your recent output? Each signal is insufficient alone; in combination they indicate that the diagnostic procedure in 13.1c should be applied to your own work. (3) Constraint-naming audit: review instances in the past three months in which you concluded that a specific change was not the right approach. Restate each as a constraint: 'I cannot currently act on this given X.' If the restated version is more analytically accurate than the original formulation, the original laundered a material constraint into an analytical position. Identify what X is and whether it has been named explicitly in your public or collaborative work. (4) Trajectory review: map current institutional relationships, funding dependencies, and coalition memberships against those held two years prior. In what direction has the set of relationships whose continuation requires your analysis to remain within specific limits changed? Drift does not require intention; it requires structural pressure that accumulates gradually enough to remain individually invisible at each step. The matrix functions as a temporal record that makes the trajectory visible across intervals that are too long to hold

in working memory but short enough for the drift to be reversible.

17.7d — On Meaning Infrastructure and the Sustainability of Analysis

The framework's account of meaning infrastructure (5.6) is primarily structural: meaning is a material variable, its distribution is uneven, its collapse produces political consequences. That structural account does not address the question that the individual faces: how does someone who holds this framework — who sees the structural conditions it describes and takes seriously the intergenerational timescales on which change operates — maintain the capacity to analyze and act without either collapsing into paralysis or converting the framework into a comfort narrative that neutralizes its own demands? This is not a peripheral question. The framework's usefulness depends on it being held by people who can sustain it over time — who do not burn out, convert to simpler frameworks that provide more psychological stability, or drift into the position of performing analysis without acting on it. A framework that is theoretically adequate but psychologically unsustainable does not have the reach it claims. What the framework provides: The minimal anchor is not only an evaluative criterion for systems. It is a navigational tool for the individual. The question 'does this arrangement reduce unnecessary suffering?' is answerable at every scale, including the scale of a single decision within a single institution on a single day. A person who cannot change a system can change specific conditions for specific people. That is not a consolation prize — it is the application of the evaluative criterion at the scale where application is possible. The framework does not authorize treating micro-level action as a substitute for structural analysis. It does authorize treating it as what structural analysis is ultimately for. What the framework requires: The framework requires that the individual maintain the capacity for honest assessment even when the assessment is unfavorable to their current position and investments. This is a demand that most frameworks do not make explicitly because most frameworks are designed to provide psychological stability rather than to challenge it. PMN's epistemological commitments make this demand unavoidable: a framework that updates based on evidence must be held by people who can update based on evidence about their own positions. That capacity is not automatic. It requires deliberate maintenance — regular application of the diagnostic procedures in this document to the individual's own work, positions, and institutional relationships, with the same willingness to follow the analysis wherever it leads that the framework requires of any other analytical object. The solidarity implication: No individual can sustain this kind of analysis alone over long periods. The psychological sustainability of holding PMN honestly depends on the existence of other people holding it with equivalent honesty — people with whom the diagnostic procedures can be applied collaboratively, who will call drift when they see it, and whose own analysis can be

e challenged without the challenge being experienced as betrayal. This is not a therapeutic claim. It is a structural one: the organizational infrastructure that sustains counter-power capacity at the political level has a direct analog at the level of the individual's capacity to maintain honest analysis. Isolation degrades both. The framework does not promise that honest analysis produces good outcomes. The cases in this Part demonstrate that structural conditions can be accurately diagnosed and still not changed — because the political window is closed, because counter-power capacity is insufficient, because the timing is wrong, or because the costs of change exceed what any available coalition can bear. The framework's demand is not that analysis produces victory. It is that analysis remains honest even when it does not produce victory, and especially then — because the moment of failure is the moment when the temptation to launder analytical conclusions from material constraints is greatest. Holding the framework honestly under conditions of persistent failure is harder than holding it under conditions of incremental success. It is also the condition under which the framework is most needed.

17.7e — On Communicating Across Frameworks

PMN uses specific analytical vocabulary — structural suffering, capture, visibility suppression, the minimal anchor — that is not shared by most of the people whose conditions the framework analyzes and whose cooperation any change process requires. A framework that is only legible to those who have already internalized it is a framework that cannot do the political work it claims. The individual who holds PMN must be able to communicate its conclusions without requiring prior acceptance of its vocabulary, or the framework remains contained within the population already persuaded. The translation problem: Translation is not simplification. Simplification loses the structural specificity that makes the analysis useful. Translation preserves the structural content while using vocabulary and narrative frames that connect to the meaning infrastructure of the audience rather than the analyst. The distinction the framework draws between visibility suppression and simple media bias is analytically important — but making that distinction actionable in a conversation with a trade union organizer who does not use this vocabulary requires finding the equivalent distinction in the vocabulary that person already uses to navigate their world. That equivalent almost always exists; finding it requires knowing what vocabulary the person uses and why. The credibility problem: Structural analysis that is presented as obviously correct — as what any honest analyst would conclude — produces resistance rather than uptake in most audiences, because it implicitly places the audience in the position of those who failed to see what should have been obvious. The framework's epistemological commitment requires presenting analytical conclusions as positions that are more defensible than alternatives given current evidence, not as truths that the audience's ideology or interest has prevented them from seeing. The difference in presentation is not rhetorical manipulation — it is the accurate representation of what the framework actually claims about the epistemic status of its conclusions. Individual-level diagnostic — communication across frameworks: (1) Can you state the structural analysis you are making in terms that a person who does not share this framework's vocabulary would recognize as describing their actual experience? If not, the analysis may be more framework-internal than structurally grounded. (2) What is the meaning infrastructure — the values, narratives, and commitments — of the person or community you are trying to communicate with? Where does that infrastructure overlap with the framework's evaluative criterion, even using different language? (3) Are you presenting your conclusions as structurally better-supported positions or as corrections of the audience's analytical failures? The second approach is almost always counterproductive regardless of whether it is accurate. (4) What would you need to concede, acknowledge, or genuinely not know in order for the

conversation to be a real exchange rather than a transmission? The framework's epistemological commitments require you to mean that concession, not perform it.

17.7f — On Weaponized Stoicism: Jurisdictional Demarcation

The Stoic philosophical tradition — from Epictetus and Marcus Aurelius through its contemporary revival — offers a body of practical philosophy whose central claim is that the suffering produced by external conditions is substantially mediated by the agent's judgments about those conditions, and that disciplined attention to what is within one's control versus what is not constitutes both the foundation of psychological stability and the condition for effective action within constraint. This claim is not false. As a framework for navigating existential suffering — the category that section 3.4b identifies as arising from permanent conditions rather than from institutional arrangements that could be otherwise — Stoic practice has genuine therapeutic value. The capacity to maintain analytical clarity under conditions of severe stress, to act well within constraint rather than being immobilized by the constraint, and to distinguish between what institutional action can change and what must be endured — all of these are capabilities that this section's analysis of what PMN asks of the individual practitioner requires. The framework does not dismiss them. What the framework requires is a jurisdictional demarcation that the Stoic tradition itself does not always supply, and whose absence creates a structural vulnerability to a specific form of ideological capture. The dichotomy of control — the distinction between what is and is not within our power — is a tool for navigating existential constraints that are genuinely fixed: mortality, the limits of knowledge, the behavior of other agents over whom we have no legitimate authority. It is not a tool for establishing that structural suffering (S) is among the fixed conditions that cannot be otherwise. Whether a worker's wages are sufficient to meet their biological floor is a question whose answer is determined by institutional arrangements that are humanly produced and could be otherwise. Whether a population has access to clean water is a question of the same kind. Whether a person's children will have access to education that expands their becoming capacity is a question of the same kind. These are not permanent conditions of human existence that the dichotomy of control correctly classifies as external to our power. They are institutional products, and the institutional arrangements that produce them can be analyzed, contested, and changed. The jurisdictional error — applying Stoic acceptance to structural suffering that is the product of changeable institutional arrangements — is not a philosophical mistake that Stoic practitioners typically make in their own lives. It is a systematic misapplication that those who benefit from the continuation of harmful arrangements have strong incentives to encourage in others. This misapplication — deploying Stoic vocabulary to induce acceptance of structural suffering that institutional analysis identifies as unnecessary and changeable — is what the framework des

ignates as weaponized Stoicism. The mechanism is V (visibility suppression, 15.0b): when a population experiencing structural suffering is systematically directed toward the management of their response to that suffering rather than toward the structural analysis of its causes, the visibility of the structural causation is reduced and the institutional arrangements producing the harm are protected from the political pressure that would otherwise be generated. The V mechanism operates here through a philosophical tradition rather than through direct information management, which makes it harder to identify and more durable: the population subject to structural harm has acquired practices that are genuinely valuable for existential suffering and is applying them to the wrong jurisdictional category without recognizing the misapplication. The practitioner who has developed genuine Stoic equanimity about their low wages, their precarious employment, their children's underfunded schooling — practices that may make their daily life more bearable — has also, if the jurisdictional demarcation is absent, acquired a set of responses that protect the institutional arrangements producing those conditions from the analytical attention and political pressure they require. Weaponized Stoicism is most systematically deployed through: self-help discourse that imports Stoic vocabulary into non-philosophical contexts; management culture that frames structural precarity as individual resilience opportunity; and populist political communication that attributes structural economic suffering to individual mindset deficits. The common feature is the direction of analytical attention from institutional cause to individual response. The demarcation rule this section establishes: Stoic practice is a valid and valuable instrument for existential suffering and for maintaining analytical capacity under constraint. It is not a valid

instrument for structural suffering, and applying it there is not a philosophical position — it is an institutional capture mechanism that operates through the individual. (See 3.4b existential vs. structural suffering taxonomy; V variable at 15.0b; 7.3c cosmological capture.)

17.8 — How a Framework Travels: The Transmission Problem

A framework that cannot account for its own transmission is incomplete in a specific way: it has analyzed how ideas function as material variables in the systems it examines without applying that analysis to itself. PMN holds that the distribution of frameworks — who holds them, in what form, with what fidelity — is a structural variable with material consequences. That claim requires the framework to have a theory of its own transmission, or it is treating itself as exempt from the structural conditions it applies to everything else. Transmission is not the same as publication. A document that exists and is accessible is not thereby transmitted. Transmission requires: that the framework reach people whose material conditions make its analytical tools useful; that it reach them in a form legible given their existing conceptual vocabulary; that something in the encounter produces the kind of engagement that makes the framework generative rather than merely read; and that the framework be held with sufficient structural support — other people holding it, contexts in which it can be applied and tested — to be sustained over time. Each of these conditions is a structural variable, not a natural consequence of quality. The legibility problem: PMN's analytical vocabulary is specific and requires prior familiarity with several theoretical traditions to fully decode. A person without prior exposure to structural analysis, institutional theory, or political economy can read the document and extract value, but the extraction is uneven — the parts most useful to their situation may be the parts least legible given their current conceptual framework. The transmission problem is partly a translation problem: making the structural content accessible without losing the structural specificity that makes it useful. Section 17.7e addresses this at the level of individual communication; the same problem operates at scale. The form problem: The document form is not neutral. Long-form analytical documents reach specific populations — those with the time, literacy, and prior conceptual formation to navigate them. They do not reach, or do not reach effectively, people whose material conditions produce the highest need for structural analysis and the least capacity to absorb it in this form. This is not a reason to abandon the document form — the analytical depth that the document form enables is not replicable in shorter formats without significant loss. It is a reason to treat the document as one form among several that the framework requires, not as its complete expression. The community of inquiry problem: Peirce's concept of the community of inquiry — the social infrastructure through which knowledge is tested, revised, and sustained — is directly applicable to the transmission of frameworks. A framework held by isolated individuals who cannot test their applications against other applications, who cannot be corrected when they misapply, and who have no ongoing rel

ationship with the framework's development will drift from the framework as they extend it into domains the original document did not address. The drift is not intentional — it is structural. The framework requires communities of application, not only readers. What those communities look like — seminar groups, working organizations, online discussion structures, practice communities — is a question of what forms are viable given the material conditions of the populations the framework is trying to reach. The capture problem in transmission: Frameworks that successfully transmit to large populations face a specific version of the custodian problem: the apparatus of transmission — the institutions, organizations, and individuals responsible for spreading the framework — develops interests in the framework's continued spread that can diverge from the framework's continued integrity. The history of every major intellectual framework that achieved significant transmission includes the moment when the apparatus of transmission began selecting for formulations that maximize uptake over formulations that maximize analytical precision. PMN's epistemological commitments require that this pressure be named rather than managed through denial of its existence. The implication for anyone involved in transmitting this framework: the test of transmission fidelity is not whether the framework is spreading but whether the people who hold it are using it to generate

analysis they could not have generated from their prior frameworks — and whether that analysis is changing what they do, not only what they think. A framework that is widely held but that does not change the behavior of those who hold it has been transmitted in form and not in substance. That is the successful transmission of the label and the unsuccessful transmission of the framework. The transmission problem has no solution that the framework can specify in advance. It has structural conditions that make transmission more or less likely, and it has failure modes that the framework can identify and against which it can design. Beyond that, it depends on what the people who encounter the framework do with it — which is precisely the domain that Part XVII's individual-level analysis addresses. The framework that reaches the right people at the right moment in a form they can use is not guaranteed by the quality of the framework. It requires the same combination of structural conditions, organizational capacity, and contingent opportunity that the framework's analysis of political change identifies as the conditions for any consequential outcome. The fidelity problem: Frameworks travel with degradation: as they move from the context in which they were produced to contexts with different material conditions, different intellectual traditions, and different political configurations, the specific claims that were calibrated to the production context are received through interpretive frameworks that modify t

hem in ways that are not visible to either the sender or the receiver. The modification is not primarily motivated distortion — though that occurs — but structural: the receiving context has different reference points, different prior intellectual commitments, and different political stakes that shape which aspects of the framework are attended to, which are interpreted as central, and which are treated as peripheral or eccentric. The result is a received framework that shares vocabulary and structure with the transmitted framework but whose operative content — the specific analytical moves that produce the framework's distinctive conclusions — has been modified in ways that make it less analytically precise. The structural conditions for high-fidelity transmission: High-fidelity transmission requires that the receiving context has access not only to the framework's conclusions but to the reasoning process through which those conclusions were reached — the specific evidence, the specific alternative positions considered, and the specific criteria through which the evidence was assessed. This is why the framework's treatment of its own intellectual debts is not ornamental: it specifies the conversations within which the framework's claims are calibrated, which allows receiving contexts to reconstruct the reasoning rather than only adopting the conclusions. Conclusions without reasoning can be held with any degree of confidence and can be applied to any context; reasoning processes expose the conditions under which the conclusions hold and the conditions under which they would need revision. The degradation of reasoning into conclusion-holding is the primary mechanism of framework fidelity loss. The political implication: the transmission problem is a counter-power problem. A framework that travels as a set of conclusions that can be held without the methodological commitments that make them valid becomes an ideology in the pejorative sense — a set of positions that protect themselves from the evidence that would require their revision. The political movements that have most durably served the minimal anchor criterion have been those that transmitted not only specific conclusions but the analytical practices that allow those conclusions to be revised in response to changing conditions and new evidence.

XIV: Summary: Framework, Doctrine, and What Remains Open

This Part is positioned last for a specific reason: its self-assessment of the framework is most credible after the full argument has been read, not before it. Part XIV states the framework's methodology (14.1), its current doctrinal positions across major domains (14.2), what it has demonstrated versus what it assumes versus what it does not know (14.3), the conditions under which it would be surpassed rather than merely revised (14.4), and the institutional requirements for transmitting it without capturing it (14.5–6). The positions in 14.2 are conditional and temporal — the current best reading given current evidence, not permanent commitments. They are held at the level of Tier 3 empirical-hypothetical assumptions (14.3) except where marked otherwise. Part XIV is updated as the framework develops. It is not a conclusion. It is a current position.

14.1 — The Framework: Adaptive Naturalism

Adaptive Naturalism — the name this framework gives to its methodology for engaging with material reality, distinguishing it from both dialectical materialism and static naturalism — is defined by three commitments: that analysis must begin from what organisms biologically are rather than from ideas, stages of history, or prior normative commitments; that the method must update when evidence requires revision rather than insulating its conclusions; and that the evaluative criterion is derived from the material consequences of arrangements for the organisms subject to them. A methodology for engaging with material reality honestly. Begins from biological foundations rather than from the history of production, the movement of consciousness, or moral philosophy. Uses adaptive empirical epistemology — the framework updates based on evidence rather than protecting prior commitments. Evaluates all systems against the minimal anchor of unnecessary suffering. Does not prescribe fixed outcomes. Prescribes a method for arriving at the best available position given current conditions and for revising that position when conditions change or evidence requires it. The relationship to historical materialism and dialectical materialism is that of a development that takes both seriously enough to follow their logic further than they did. Historical materialism was right that analysis must begin from material conditions; this framework asks what those conditions are ultimately grounded in. Dialectical materialism was right that reality moves and analysis must move with it; this framework replaces the teleological engine of the dialectic with the more honest and more demanding account of how movement actually happens. What distinguishes this framework from adjacent positions: it is not utilitarian, because it does not reduce all value to a single welfare metric and does not treat the distribution of welfare as irrelevant to evaluation. It is not Rawlsian, because it does not derive its evaluative standard from hypothetical agreement behind a veil of ignorance — it derives it from what organisms materially are. It is not Marxist in the doctrinal sense, because it does not hold the interests of any particular class as the universal standard — the minimal anchor applies to the suffering of any population, regardless of its position in economic relations. What it shares with all three is the refusal to treat existing arrangements as natural baselines against which change must justify itself. The baseline is the material conditions of actual organisms, and existing arrangements are evaluated by what they produce for those organisms — no more, no less. (Rawls 1971) The engagement with non-Western intellectual traditions in Parts V, VII, and X adds an evidential dimension that was not present in earlier versions of this claim. The floor/becoming evaluative structure is not only derivable from materialist biological foundations — it is independently derivable from the Maqa

sid al-Shariah's daruriyyat/tahsiniyyat hierarchy, from Al-Farabi's sa■ada framework, and from the Confucian Tianming (Mandate of Heaven) legitimacy theory. The parallel derivations from different foundational commitments constitute structural evidence that the evaluative architecture is tracking something about human social organization that is not an artifact of the Western materialist tradition from which PMN's primary vocabulary is drawn. This does not prove the framework's adequacy — independent

traditions can converge on the same error — but it does shift the evidentiary burden: a critique that targets the Western materialist foundations of PMN must also explain why structurally equivalent frameworks were independently derived from Islamic jurisprudence, Islamic political philosophy, and Confucian political theory. The key move distinguishing this framework from adjacent positions is methodological rather than positional: it asks what actual organisms materially require rather than what hypothetical agents would rationally choose or what the logic of history demands. Utilitarianism begins from a welfare metric and aggregates; the minimal anchor begins from a biological fact and evaluates arrangements by what they produce for that fact across the full distribution of populations affected, not in aggregate. Rawlsian liberalism derives its evaluative standard from hypothetical agreement; this framework derives it from what organisms materially are, which means the standard does not depend on the rationality assumptions or the choice architecture of the original position. Historical materialism locates the engine of evaluation in the contradictions of production relations; this framework locates it in the biological substrate that production relations must ultimately serve or ultimately fail. The framework is adaptive in a specific sense: it updates the application of its evaluative criterion as conditions change and as evidence accumulates, without updating the criterion itself. What changes is not the standard — unnecessary suffering — but the analysis of which arrangements produce it, under what conditions, and through which mechanisms. This is not relativism about values; it is empirical rigor about the relationship between institutional arrangements and material outcomes. A framework that updates its values in response to political conditions is not adaptive — it is captured. A framework that holds its values fixed while updating its analysis of how those values are served or violated in changing conditions is what adaptive empirical epistemology requires. What adaptive naturalism is: Adaptive Naturalism is a methodology for engaging with material reality through the systematic rejection of two errors that have consistently distorted social analysis. The first error is the teleological error: the assumption that social systems move toward predetermined endpoints that give moral direction to historical analysis. The second error is the found

dationalist error: the assumption that analysis requires securing a foundation — in reason, in God, in the logic of history, or in any other unconditioned ground — before it can proceed. Adaptive Naturalism rejects both. It begins from what is materially available: the biological requirements and constraints of human organisms operating in social environments, the causal mechanisms that connect those requirements to institutional arrangements, and the evidence of what specific arrangements actually produce. It treats these as the foundation not because they are philosophically unconditioned but because they are the most stable reference point available — more stable than historical laws, more stable than metaphysical commitments, more stable than cultural consensus — for evaluating how well institutional arrangements serve the beings they organize. Three methodological commitments: Adaptive Naturalism specifies three methodological commitments that together define how the framework operates. First, the empirical priority commitment: when theoretical analysis and empirical evidence conflict, revise the theory. This is straightforward in formulation and demanding in practice, because the evidence that matters is often evidence about causal mechanisms that operate at levels of abstraction where interpretation is contested and where motivated reasoning can find plausible grounds for discounting inconvenient findings. The framework's response is to require explicit specification of what evidence would change a given analytical conclusion — making the commitment to revision testable rather than rhetorical. Second, the structural priority commitment: when the question is why an outcome occurred, begin with structural factors before attributing it to individual characteristics. This does not deny individual agency; it sequences the analytical questions to prevent the systematic error of attributing structural outcomes to individual causes. Third, the self-application commitment: the framework must apply its own analytical methods to itself — identifying the material conditions of its production, the interests that its adoption would serve, and the blind spots that follow from its production conditions. What the framework does not claim: Adaptive Naturalism does not claim to provide a complete account of human social life, to generate predictions precise enough to specify outcomes in particular cases, or to prescribe specific institutional arrangements as correct across all historical

conditions. These are claims that the framework's own epistemology prevents it from making. What it claims is more modest and more precise: it provides a methodology that, when applied consistently to available evidence, generates analytical conclusions about causal mechanisms and evaluative conclusions about institutional performance that are more reliable than the alternatives — more reliable because they are more exposed to correction through the evidence they rely on, and

more calibrated because they do not need to explain away evidence that conflicts with predetermined conclusions. The relationship to the philosophical tradition: Adaptive Naturalism inherits from the naturalist tradition in philosophy the commitment to treating human beings as continuous with the natural world rather than as exceptional cases that require separate analysis. It departs from much of that tradition in its treatment of emergence: social phenomena are real at the level at which they operate and cannot be reduced to lower-level descriptions without losing the causal information that makes them analytically useful. This is the position of 2.4 and 2.5, and it distinguishes Adaptive Naturalism from both crude reductionism and from varieties of idealism that treat social phenomena as independent of material conditions. The framework label is descriptive rather than programmatic: Adaptive Naturalism describes what the framework does — it adapts its conclusions in response to new evidence while treating naturalist assumptions about the embeddedness of human beings in material reality as the stable background against which those adaptations occur. It is not a school of thought with a fixed theoretical core that defines membership. It is a set of methodological commitments that can be applied consistently across different analytical contexts.

14.2 — Current Analytical Positions: The Framework Applied

The application of the framework to current material conditions. Conditional and temporal — it is the current best reading, not the permanent answer. The doctrine will change as conditions change and as the framework's application accumulates evidence that requires revision. What follows is the current doctrinal position across the framework's major domains — not a permanent commitment, but the best available reading of what the framework implies given conditions as they currently exist. The current doctrinal positions across the framework's major domains are as follows. Epistemology: adaptive and evidence-responsive, with a high threshold for claims that outrun evidence. Ontology: realist about the existence of patterns independent of perception; agnostic about metaphysical questions that do not yet make contact with observable consequences. Ethics: the minimal anchor of unnecessary suffering, outcome-sensitive but institutionally aware. Politics: progressive institutionalism — substantive accountability and genuine reform capacity. Society: materialist evaluation of cultural practices by outcomes, with the consistency standard applied across all actors. Religion and meaning: treated as material variables with real consequences; the internal epistemological standards of traditions are respected as structural realities, not merely as beliefs to be managed. Geopolitics: multipolarity as transitional equilibrium; post-national coordination as the long-term direction; parallel development as the strategic framework for the transitional period. Economics: horizon-oriented, with instruments selected by material effectiveness and analytical resources drawn from multiple traditions including but not limited to the Marxist tradition. The current doctrinal position across the framework's major domains can be stated with enough specificity to be falsifiable, which is the standard any doctrine must meet. On economic arrangements: contemporary capitalism, as it currently operates in high-income economies, fails the developmental criterion at the distributional end — it sustains the conditions for preserved adaptive response capacity for a minority while maintaining the floor barely, inconsistently, and with sufficient precarity to foreclose development for a large portion of the population it administers. The appropriate doctrine is not the replacement of markets with central administration, which has a well-documented failure record, but the aggressive restructuring of the conditions under which markets operate: the redistribution of ownership, the expansion of what is treated as public infrastructure for becoming (education organized around development rather than credentialing, healthcare as structural rather than commodity), and the regulation of accumulation dynamics that convert

market efficiency into political capture. On political arrangements: the current distribution of formal democratic institutions, in both high-income and middle-income economies

, operates with a systematic gap between the formal accountability structures and the actual distributions of structural power analyzed in Part VI. Elections are real but insufficient; they allocate formal authority without reliably altering the structural positions through which consequential decisions about material conditions are made. The doctrinal position is institutional pluralism with structural accountability: the expansion of the domains in which collective decisions are made through genuinely accountable processes, combined with the anti-capture mechanisms developed in 7.3b, rather than either electoral maximalism (which mistakes formal access for structural power) or institutional cynicism (which abandons the gains that formal accountability structures represent). On the international order: the current period is a transitional one in which the post-1945 institutional architecture is losing the hegemonic support that sustained it without yet having developed adequate alternatives. The doctrinal position is reformist internationalism with clear eyes about the structural interests that make reform difficult: the existing institutions should be defended where they represent genuine constraints on unilateral power projection, reformed where they systematically exclude the populations most affected by their decisions, and supplemented with new forms of coordination in domains — climate, technology governance, labor standards — where the existing architecture has structural blind spots. The horizon is not world government, which exceeds what current political conditions can sustain, but a denser web of binding commitments that constrain the ability of structurally powerful actors to externalize costs onto populations who have no voice in the decisions that impose those costs. These positions have specific content that distinguishes them from adjacent positions with similar labels. On political economy: the doctrine holds that market mechanisms are indispensable instruments for generating the information and adaptive capacity that planned systems consistently fail to produce — and that unregulated markets systematically generate the power concentrations, information asymmetries, and externalized costs that the minimal anchor requires preventing. The doctrinal position is therefore neither market fundamentalism nor planned economy advocacy, but the identification of which market failures require structural correction and which corrections create new capture risks requiring their own management. The appropriate design is not derivable in advance; it is a permanent diagnostic problem. Parts XV–XVII — which the reader has now encountered in full — sharpen the doctrinal positions above in three specific ways. First, the formula architecture of Part XV operationalizes these positions as a variable set: structural tendency analysis makes explicit which variables drive each domain's dynamics and exposes them to structured comparison. An analyst who

disagrees with a doctrinal position stated above has a specific target: identify which variable relationship in the formula architecture embeds the disagreement and what evidence would change it. Second, the technology analysis of Part XVI is now integral to the doctrinal position on political economy and international order. AI as an institutional form, algorithmic governance as advanced V operating through the ideology of objectivity, and semiconductor control as leverage-point capture at the interstate level are not peripheral developments — they are the primary mechanisms through which the structural dynamics this document analyzes are currently operating at scale. Any doctrinal position on political economy or institutions that does not account for these mechanisms is already partial. Third, the historical cases and individual-level demands of Part XVII establish the empirical and practical floor for the positions stated above. The labour movement case demonstrates at scale the capture dynamic the political economy position is designed to prevent. The individual-level demands (17.7) apply the same diagnostic rigor to the analyst who holds the framework as the framework applies to institutions. The doctrinal positions stated here are held with that requirement in view. On democratic institutions: the doctrine holds that no currently existing democratic arrangement adequately resolves the tension between formal representation and material accountability — and that the appropriate response is not to abandon the accountability aspiration but to identify the specific mechanisms through which existing arrangements fail it and to design interventions accordingly. Universal suffrage is necessary but insufficient; the conditions under which it produces genuine

accountability include competitive information environments, organizational capacity for those with less structural power, and mechanisms that resist the consolidation of influence over the institutions that are supposed to distribute it. Where those conditions do not exist, formal democracy legitimates arrangements that contradict its own stated purpose. On the international order: the doctrine holds that the current distribution of global resources and institutional authority reflects historical processes that the framework's evaluative criterion — consistently applied — cannot endorse, and that the long-term direction toward greater accountability to human welfare across national boundaries is both practically necessary and normatively required. The practical constraints on this direction are real: the redistribution of institutional authority faces resistance from all actors who currently benefit from its distribution, and that resistance is more organized and better resourced than any movement toward change. The doctrinal position is not that the horizon is reachable in the near term, but that movement toward it is directionally correct and that movement away from it is not a neutral recog

nition of constraint but a choice with evaluable consequences.

14.3 — Argument, Assumption, and Open Question

Arguments that can be tested: The assumption of personal intention in standard metaphysical debates outstrips the evidence that supports it Interface-emergent phenomena are observable across many domains and are real at the level of experience regardless of the ontological status of their source Individually rational behavior can produce collectively irrational outcomes — documented across domains and scales Inaction has material consequences that systematically favor current distributions of suffering Premature institutional transitions destroy the capacity that makes the transition possible Symbolic politics consistently generates backlash that worsens material conditions for the populations it claims to serve Institutional capture of social movements is a structural outcome of resource dependency, not a contingent failure That the two-level architecture (source level / interface level) generates testable predictions at the interface level: behavioral, motivational, and solidarity consequences of metaphysical commitment are real and measurable regardless of source-level resolution Assumptions that are acknowledged: That the minimal anchor — unnecessary suffering — is cross-cultural enough to function as a universal evaluative criterion That adaptive response capacity — the capacity that institutional arrangements must not foreclose — is evaluatively significant as the temporal extension of the floor criterion: foreclosing becoming is a form of structural suffering imposed on future populations, not merely a failure to provide benefit. This is the framework's second foundational commitment, continuous with the minimal anchor rather than independent of it; both derive from the floor criterion traced through its causal implications at different timescales. The anti-foreclosure commitment is defended in Part IV through the loop model and the Last Man argument — defended as a derivation, not as an independent axiom. It carries less philosophical vulnerability than a second independent value commitment would require. That the permanent conditions identified in Part III are genuinely permanent rather than historically contingent That material conditions will eventually develop to support post-national coordination That there is a necessary being at the end of the chain of contingency — opened by argument but not established by it That interface-level reality produces consequences not fully reducible to the psychology of the individuals experiencing it That mind-independent material reality is primary — this is the foundational axiom of the framework, adopted as a starting point rather than derived from prior argument. The convergence argument in 2.1 provides a reason to prefer it but does not prove it That suffering is the appropriate evaluative criterion (over life, flourishing, or capabilities) — this is a design choice defended in 3.0 and 4.1, not a discovery entailed by biology. Alternatives in the capabilities tradition represent live competing c

hoices Open questions that cannot currently be resolved: The hard problem of consciousness — what is the relationship between physical processes and subjective experience? The precise conditions constituting genuine readiness for more coordinated international arrangements How to design institutions that

self-correct in material terms rather than just procedural ones. When truth management is justified versus when it becomes the kind of narrative control the framework criticizes in others — the conditional-tension analysis in 13.2b specifies part of this boundary; what remains open is the threshold at which a conditional tension has crystallized. The threshold conditions for when the balance between institutional capacity and accountability should shift. Whether the framework's analytical claims — that suffering produces systemic instability, that institutional capture follows from resource dependency, that mixed sequences outperform pure reform or pure revolution — are borne out when applied to specific cases by analysts who did not construct it. Internal consistency is not evidence of empirical accuracy. That evidence must come from application, and the application has not yet accumulated to a scale that constitutes a genuine test. Whether the impersonal-necessary reading of 2.3 dissolves (rather than resolves) the problem of evil — and that this dissolution depends on whether the presupposition of agency at the source level is genuinely removed, which is an open philosophical question. Acknowledging these three categories honestly is the application of the framework's own epistemological standards to itself. A framework that treats all of its positions as established argument — that does not distinguish between what it has demonstrated, what it assumes, and what it genuinely does not know — is doing exactly what Part I identifies as the primary epistemic failure. The capacity for honest self-assessment is not a feature added to the framework. It is what the framework requires of everything it examines, applied to itself. The above lists state the three categories in compressed form. What follows elaborates each category with the specificity the framework's epistemological standards require — specifying not just what falls in each category but why it falls there, and what would move an item between categories. Arguments that can be tested. The assumption of personal intention in standard metaphysical debates outstrips the evidence that supports it — and the framework holds that the methodological implications of treating metaphysical claims at the interface level rather than the source level can be evaluated against the evidence of what each approach produces analytically. The claim that interface-emergent phenomena are observable across multiple domains is a specific empirical claim that can be evaluated against the evidence of each domain. The claim that individually rational behavior can produce collectively irrational outcomes is documented across

multiple domains and scales with sufficient precision that the analytical question is not whether it occurs but what the structural conditions are that determine its frequency and severity. The claim that inaction has material consequences equivalent in structure to the consequences of action is a claim about causal structure that can be evaluated against historical evidence. The claim that institutional arrangements systematically produce outcomes that were not intended and could not be predicted from the intentions of the actors who designed them is supported by sufficient historical evidence that the burden of argument falls on any analysis that treats intentions as adequate causal explanations of institutional outcomes. The claim that counter-power accumulation is the primary mechanism through which structural transformation occurs is a structural claim that can be evaluated against the historical record of how arrangements have actually changed — and against the record of how they have failed to change when counter-power was absent or inadequately organized. The two-level architecture — that the source level of reality may be impersonal while the interface level generates genuinely personal phenomena through the interaction of consciousness with complex structure — is a testable claim at the interface level: it predicts that the behavioral, motivational, and solidarity consequences of religious and metaphysical commitment are real and measurable regardless of source-level metaphysics. This prediction is falsifiable and has significant support in the sociology of religion. That the impersonal-ground reading of necessary being (2.3) resolves the Epicurean paradox by dissolving its central presupposition — that the source level has personal agency — is a philosophical argument, not an empirical claim. Its defensibility is a matter of whether the dissolution is genuine or merely a relabeling of the problem. This is an open question, acknowledged as such. Assumptions that are acknowledged. That mind-independent material reality is primary and can be analyzed with increasing accuracy is a foundational axiom of the framework, adopted as a starting point rather than derived from prior argument. The convergence argument in 2.1 provides a reason to prefer this axiom, but does not prove it. The framework's entire analytical

architecture inherits whatever vulnerability this axiom carries. That the choice of suffering as the evaluative criterion — over life, flourishing, preferences, or capabilities — is a design choice defended by the arguments in 3.0 and 4.1, not a discovery entailed by the biological foundation. Alternatives in the capabilities tradition (Sen, Nussbaum) represent live competing choices. The framework proceeds on its choice with the acknowledgment that the arguments supporting it reduce but do not eliminate the choice's philosophical costs. The assumption that biological requirements constitute a stable evaluative foundation across

the changes in institutional context that the framework analyzes is not itself a result of the analysis — it is a prior commitment that the analysis depends on. The two-level architecture — life as ontological foundation, suffering as evaluative criterion, the anti-foreclosure criterion as its temporal extension (what the framework calls 'becoming' is the floor criterion applied temporally) — is a design structure for the framework, not a discovery within it. The arguments in 3.0 and 4.1 defend the structure but do not uniquely determine it: capabilities frameworks, preference-satisfaction approaches, and thick conceptions of flourishing all represent live alternatives that the framework's choice excludes rather than refutes. Acknowledging this is not a concession that undermines the framework. It is the specification of where the framework's deepest commitments lie and where the most fundamental challenges to it would need to be directed. The assumption can be defended on the grounds that biological requirements are more stable than the institutional arrangements being evaluated and that they therefore provide a less circular foundation than any alternative, but the argument for stability is not the same as proof of adequacy for all analytical purposes the framework is put to. The assumption that suffering has the negative evaluative valence the framework assigns to it — that its presence is bad and its reduction is good — is treated as the framework's foundational normative commitment rather than as a derived conclusion. It is defended against the is-ought objection in Part III, but the defense is a philosophical argument, not a deduction from empirical premises. The assumption that the Structural Tendency Analysis (STA) formula architecture in Part XV captures the primary structural tendencies in historical transformation is an analytical commitment whose precision should not be confused with established causal knowledge. The formulas represent structural hypotheses about variable relationships, not confirmed causal laws. They are useful as diagnostic frameworks and as tools for identifying where the analysis should be directed, not as predictive instruments whose outputs can be treated as reliable without qualification. Open questions that cannot currently be resolved The subject boundary question identified in 13.4c — at what point modifications to the biological substrate that carries the evaluative criterion would change the subject in ways that affect the criterion's application — is an open question that the framework cannot currently resolve. It is not an evasion to say this; the question requires analytical resources that depend on empirical developments in biotechnology and AI that have not yet occurred and that cannot be anticipated with sufficient precision to generate a current answer. The coordination problem at the horizon of genuine abundance — how the framework's evaluative criteria apply when material scarcity no longer operat

es as a constraint on what institutional arrangements must manage — is an open question. The framework's answer to this question is speculative in the literal sense: it is offered as the framework's best current projection, not as a conclusion supported by evidence from conditions that resemble those of genuine abundance. • Whether the two-level architecture — life as ontological foundation, suffering as evaluative criterion — is internally coherent when applied to entities at the post-biological boundary, or whether the framework requires a different ontological grounding once the biological substrate that gives suffering its determinate content is significantly modified. The architecture is stated in 3.0 and applied in 3.4d; whether it holds across the full range of cases it is intended to cover is an empirical and philosophical question that neither section can close from within. • Whether the defeasible structure of the framework's claims — specified in 12.1d — produces a logical architecture that is sufficiently precise to generate testable predictions, or whether defeasibility introduces so much flexibility that the framework becomes unfalsifiable in practice. The worry is that a claim that is always defeasible can always be defended against any

counterexample by identifying a defeater post hoc. The response in 12.1d requires that defeaters be structurally specified in advance, not invented to accommodate inconvenient cases; whether this requirement is actually met throughout the framework is a question the framework cannot answer for itself. The question of whether the framework's biological foundation is adequate to handle forms of consciousness and sentience that differ substantially from current human experience — whether it can be extended to AI systems with behavioral signatures of suffering, to non-human animals with different cognitive architectures, or to hypothetical future beings whose relationship to biological substrate is unclear — is an open question that the framework raises through the analysis of 3.4c but cannot resolve from its current position. The arguments, assumptions, and open questions catalogued in this section are not all of the same type. The framework's epistemological commitments require distinguishing three tiers within its foundational commitments, because the appropriate response to challenge differs across tiers. Tier 1 — foundational axioms — are the commitments the framework cannot revise without becoming a different framework. They are adopted as starting points, not derived, and they are defended by the convergence arguments in Parts I and II rather than by evidence from within the framework. Tier 2 — structural commitments — are claims about how the foundational axioms connect to the framework's analytical architecture. They are more specific than the axioms and more revisable: they can be challenged by showing that the architecture does not require them or that an alternative commitment produces better

results by the framework's own criteria. Tier 3 — empirical-hypothetical assumptions — are claims that function as working hypotheses within the framework's analysis. They are treated as well-supported but are explicitly revisable in response to evidence. The distinction between Tier 2 and Tier 3 is the most analytically consequential: a challenge to a Tier 2 commitment is a challenge to the framework's architecture; a challenge to a Tier 3 assumption is a challenge to a specific analytical claim that can be revised without restructuring the framework. What the tier structure itself assumes. The tiers are individuated by the cost of repair — what a successful challenge would require rewriting. That cost is continuous, and the division into three is a choice about legibility rather than a result the criterion produces: a finer partition would be defensible, and a coarser one would discard a difference the framework holds to be real. The boundaries are also not equally weighted. The Tier 2 / Tier 3 distinction is identified above as the most analytically consequential, which leaves the Tier 1 / Tier 2 boundary carrying less weight without a separate argument for why it is a boundary at all. The membership test also differs in kind across tiers. Tier 2 membership can be checked by asking whether the architecture requires the commitment, and Tier 3 membership by asking what evidence would revise it; both tests stand outside the list. Tier 1 membership does not: what counts as becoming a different framework is fixed by what Tier 1 contains, so the test and the membership define each other. That circularity is a form of insulation, and 1.2 identifies insulation as the characteristic failure of frameworks under evidential pressure. Two things prevent it from operating as insulation here, and neither is a proof that it cannot. The tier is declared rather than assumed silently, which makes its contents available for challenge from outside the framework even where the framework cannot generate the challenge internally. And 14.4b specifies conditions under which the foundational commitments would require paradigm-level revision, which is the anti-closure mechanism this section would otherwise lack. Stating these limits is not the same as resolving them. It is the same operation this section performs on every other commitment the framework holds, applied to the structure doing the sorting. Axiom hierarchy — consolidated reference: Tier 1 (foundational axioms, adopted as starting points): (a) Mind-independent material reality is primary (2.1). (b) Suffering has the negative evaluative valence the framework assigns to it — its presence is bad and its reduction is good (3.0, 3.3). These are design choices defended by argument but not uniquely determined by it. Tier 2 (structural commitments, revisable only by showing architectural failure): (a) Life as ontological foundation; suffering as evaluative criterion; the anti-foreclosure criterion as its temporal ex

tension (shorthand: 'becoming' — understood as what the floor criterion requires when applied across time, not as an independent positive aspiration) (3.0). (b) Descriptive egoism as the biological starting point: organisms avoid pain and seek resources as a prior fact, before any normative framework is imposed (3.2).

This is not a claim that organisms are selfish in all their behavior; it is the minimal factual anchor from which the is-ought bridge in 3.3 begins. (c) The layered constraint model: Ecological Viability → Biological Floor → Economic Arrangements → Institutional Architecture → Meaning Infrastructure (2.2; ecological layer at 3.10, meaning layer at 5.6). (d) Probabilistic determinism: structural forces establish probability distributions, not specific outcomes (1.3). (e) The is-ought bridge through material consequences and stability: suffering at scale produces observable systemic instability (3.3). (f) The two-level metaphysical architecture: source level / interface level (2.3). Tier 3 (empirical-hypothetical assumptions, explicitly revisable): (a) The permanent conditions identified in Part III are genuinely permanent rather than historically contingent — working hypothesis, falsifiable by demonstration of institutional designs that eliminate them. (b) Material conditions will eventually develop to support post-national coordination — speculative projection, stated as such in Part IX. (c) Cross-cultural validity of suffering as evaluative anchor — assumed on convergence grounds (3.2); the biological foundation falsification condition in 14.4b specifies the evidence that would revise this. (d) The formula architecture in Part XV captures primary structural tendencies — structural hypothesis, not confirmed causal law. Dependency structure: Tier 3 assumptions generally depend on Tier 2 commitments, which depend on Tier 1 axioms. A successful challenge to a Tier 1 axiom requires replacing the framework. A successful challenge to a Tier 2 commitment requires restructuring the architecture. A successful challenge to a Tier 3 assumption requires revising a specific analytical claim. One exception runs the other way and is named here rather than left for a reader to find. Tier 1(b) — the negative evaluative valence of suffering — functions as a universal criterion only if Tier 3(c), the cross-cultural validity of suffering as an anchor, holds; 14.4b states the relation directly, that the biological foundation rests on that assumption. A successful challenge to 3(c) would therefore carry a Tier 1 cost rather than the Tier 3 cost its placement implies. The item is left at Tier 3 because its evidential status is empirical and 14.4b specifies what evidence would revise it: moving it to Tier 1 would remove the only falsification condition this hierarchy specifies, which is a worse outcome than an acknowledged exception to the dependency ordering.

14.4 — What Would Surpass This Framework

The preface states that a framework that cannot be surpassed has stopped moving. That commitment requires specification: what would surpassing this framework look like, from what directions is it most likely to come, and what would constitute evidence that it has happened? Without that specification, the claim of surpassability is itself a form of insulation — it sounds open to revision while providing no guidance about what revision would look like. The most likely direction: a better account of the biological foundation. The framework's starting point — biological vulnerability to suffering as the non-arbitrary foundation for evaluation — is both its most generative feature and its most exposed one. The claim that this is a genuinely non-arbitrary starting point rests on the argument that suffering produces systemic consequences regardless of whether anyone values it negatively, which makes it different from starting points that require importing a prior value commitment. If neuroscience, evolutionary biology, or developmental psychology produce accounts that fundamentally change what the minimal anchor requires — showing that the biological substrate varies so dramatically as to undermine cross-cultural application, or that suffering and flourishing are less connected than assumed: that would require revision amounting to surpassing. Not updating a detail. Replacing the foundation. The most likely direction: a better account of non-linearity. The framework's analysis of system dynamics rests on threshold, path dependence, and adaptive non-linearity. Complexity science, network theory, and the empirical study of large-scale social change are all developing rapidly. A framework that produces significantly better predictions about when systems are approaching threshold, about which interventions open or close windows, and about how counter-power accumulation actually works under specific configurations of conditions — a framework that is not just theoretically superior but demonstrably more accurate — would surpass this one in the domain where this framework is weakest: its limited ability to make specific, falsifiable predictions about specific situations. The most likely direction: a

framework that integrates consciousness adequately. The framework is explicitly agnostic about the hard problem of consciousness. It treats subjective experience as real at the level of material consequence without claiming to explain what it is. That agnosticism is honest but it is also a limitation: it means the framework cannot fully account for what meaning, experience, and consciousness contribute to human flourishing beyond their material consequences. A framework that solves or substantially advances on the hard problem — that gives a rigorous account of the relationship between material processes and subjective experience that is adequate to the actual richness of that experience — would not only surpass this framework philosophically. It would change the b

iological foundation by changing what is known about what the biological substrate actually is. What surpassing would not look like. Not every critique constitutes surpassing. A critique that shows the framework is incomplete — that it has not addressed some domain adequately, that some of its positions require revision — is a critique that the framework invites and can accommodate through revision. That is not surpassing; it is the normal operation of an adaptive framework. Surpassing requires a framework that makes possible forms of analysis that this framework cannot generate from its own starting point — that opens genuinely new terrain rather than filling in details within existing terrain. The test is whether an analyst using the successor framework can produce findings about material reality that an analyst using this framework cannot reach, not whether the successor framework has fewer gaps or better-specified diagnostics. These conditions for surpassing are not defensive conditions — they are not designed to make surpassing difficult to claim. They are designed to make it specific. The difference between genuine surpassing and sophisticated critique matters because it determines what kind of response is appropriate: genuine surpassing requires replacement, not revision. A framework that responds to genuine surpassing with revision is insulating itself from the evidence that it needs to be replaced — which is the failure mode that Part I identifies as the primary epistemic error. The framework's commitment to being surpassable is only as real as its capacity to recognize when that has happened. Surpassing from below — better empirical grounding: The most likely form of surpassing from below is a framework that grounds its evaluative criteria in biological and neuroscientific evidence with more precision than this framework currently achieves. PMN uses the biological foundation as a reference point but does not derive its specific claims from detailed biological and neuroscientific evidence about what human organisms actually require for development. A successor framework that drew on the developing science of human flourishing — what conditions actually produce the cognitive, emotional, and relational development that the becoming framework identifies as the direction of the evaluative criterion — would be architecturally more adequate to the foundation this framework identifies. The current state of that science does not support the level of precision required; a successor framework produced when it does would surpass this one in a genuine sense. Surpassing from above — better global institutional analysis: The framework's analysis of international arrangements is the section most likely to be surpassed by developments in the actual institutional architecture it is analyzing. The institutions, power configurations, and governance arrangements of 2026 are transitional in ways that will produce different configurations within decades — pos

sibly within years. A successor framework that could analyze the institutional architecture that emerges from the current multipolar transition with the analytical precision that this framework applies to the domestic institutional configurations it knows better would surpass it in another genuine sense. The current framework's international analysis is structurally sound but empirically sparse relative to what is needed. Surpassing from across traditions — a framework produced by genuine cross-traditional inquiry: The framework's most significant structural limitation is its production conditions. A framework produced from within a single intellectual tradition, however carefully extended toward other traditions, cannot fully test its own foundational claims against the full range of alternative starting points that human intellectual history has generated. The independent derivations of the floor/becoming structure in Maqasid al-Shariah and Al-Farabi's political philosophy, the challenge to descriptive egoism from Ubuntu relational ontology, the Confucian legitimacy depletion theory that independently reaches conclusions parallel to PMN's 6.6 — these

are confirmations of the framework's core mechanisms from different starting points, which increases confidence in those mechanisms. But they also reveal places where the framework's current formulations are less adequate than alternatives derived from different foundational commitments: Maqasid's explicit ordering of floor conditions when they conflict is more analytically developed than PMN's current formulation; Confucian role ethics raises governance design questions that the individualist assumptions implicit in PMN's custodian problem analysis do not fully address; Ubuntu's relational personhood articulates something that PMN's formula implies but does not state. A framework that surpasses this one will likely be produced by a community of inquiry that includes analysts working from within these traditions as primary rather than secondary interlocutors — not incorporating their insights as supplements to a Western materialist core, but testing whether the core itself is more adequate than the alternatives those traditions provide. The commitment to being surpassed is not rhetorical. The framework asks of those who engage with it that they identify specifically where its analysis is wrong, where its causal claims are empirically unsupported, and where its evaluative criteria produce results that the minimal anchor criterion itself would not endorse. Those specific identifications are the content of surpassing — not a general claim that the framework has limits, which it acknowledges explicitly, but the specific specification of what those limits are and how they would be overcome.

14.4b — Anti-Closure: Conditions for Paradigm-Level Revision

Section 14.4 specifies what surpassing this framework would look like. This section addresses a prior structural problem: frameworks can simulate openness to revision while never actually revising at the level of foundational assumptions. Internal iteration — adding sections, refining diagnostics, acknowledging gaps — can continue indefinitely without touching the axioms that generate the framework's characteristic outputs. The distinction between iterative refinement and paradigm-level revision requires external triggers: conditions that cannot be satisfied by adding more framework content and that cannot be internally redefined as already-addressed. This section specifies those triggers and the conditions under which the appropriate response is replacement rather than revision. External revision triggers — conditions under which PMN should be revised at the paradigm level or replaced: (1) Systematic predictive failure: if the framework's structural orientation — specifically, that high structural suffering with low adaptive capacity produces transformation pressure — fails to hold in more than one-third of well-specified cases where all recognized defeaters have been excluded, the foundational causal claim requires revision, not the application procedure. This trigger cannot be satisfied by adding defeaters post-hoc; defeaters must be specified before the case is classified. (2) Evaluative criterion capture: if sustained application of the minimal anchor (suffering reduction, becoming expansion) consistently produces recommendations that structurally advantage the analysts applying it rather than the populations analyzed, the evaluative criterion has been captured and the framework has become self-serving in the way it was designed to resist. Detection: the misuse triggers in 12.5d, applied to a sample of PMN-derived outputs over time, should produce no systematic bias toward analyst-class interests. If they do, this trigger fires. (3) Foundational assumption falsification: the framework's biological foundation rests on the assumption that suffering is a cross-culturally valid evaluative anchor because all organisms avoid it. If empirical evidence demonstrates that for a substantial population, the framework's suffering-based evaluation consistently misidentifies what that population is actually trying to avoid — and that their own evaluative orientation cannot be explained as visibility suppression — the foundation requires replacement, not supplementation. (4) External framework dominance: if a competing framework consistently produces better-specified predictions with fewer post-hoc adjustments across cases where both are applied, and this pattern persists across more than three years of comparative application, the principle of epistemic humility (1.1) requires seriously evaluating replacement. This trigger is externally generated and cannot be blocked by reinterpreting the results as PMN-compatible. When any of these triggers fires, the ap

appropriate response is not to add a section. It is to convene the community of inquiry (12.7) to assess whether iterative revision or paradigm replacement is indicated, using the comparison in 12.9 as a model for what genuine competitive assessment looks like. Biological foundation falsification condition (trigger 3 specified): If independent, peer-reviewed longitudinal evidence from at least three distinct populations, each studied across two or more generations, demonstrates that a substantial proportion of individuals consistently do not avoid conditions PMN identifies as severe structural suffering — including chronic resource deprivation, untreated physical injury, and sustained institutional betrayal — and this pattern cannot be explained by visibility suppression (V), cosmological capture (7.3c-ii), stable irrationality (3.8c), or documented cultural variation in pain expression, then the biological foundation (3.0–3.2) requires replacement. This trigger is deliberately difficult to activate: it requires intergenerational, cross-population evidence and explicitly rules out the mechanisms PMN already uses to explain divergence between suffering and response. Its function is to ensure the biological foundation is a genuine empirical claim with specified falsification conditions, not a self-sealing axiom.

14.5 — Framework Institutionalization: Transmission, Quality Control, and Evolution

A framework that cannot reproduce itself without distortion across actors, contexts, and time is not a framework — it is a text. Institutionalization is the process by which a framework becomes capable of self-correction and survival without depending on any specific individual or community for its integrity. This section specifies the minimum institutional architecture required for PMN to function as a self-correcting system rather than as a document that accumulates interpretations. Transmission requires two mechanisms operating simultaneously. The first is the compressed core (15.15): the three-idea version that can be transmitted across cognitive capacity levels without requiring the full architecture. Transmission via the compressed core is not degraded transmission — it is the appropriate transmission form for most contexts, with the full architecture reserved for cases where analytical demands require it. The second is the detection procedures (12.3b, 12.5d): transmitting the framework without transmitting the ability to identify when it is being misused or when doctrine has been smuggled into framework position transmits a tool without its safety mechanisms. Both must travel together. Transmission that delivers the conclusions without the detection procedures is the primary mechanism by which frameworks become ideologies: the outputs persist while the self-critical infrastructure that generated them is lost. Quality control requires a community of practice with the structural features that prevent the community itself from becoming a capture mechanism. The anti-capture requirement (7.3d) applies here: the community that interprets and develops PMN must be accountable to the populations whose structural suffering it analyzes, not primarily to each other. Concretely, this means that quality disputes — cases where two PMN-derived analyses produce contradictory conclusions — are resolved not by the community’s internal consensus but by applying the detection procedures in 12.3b and 12.5d to both analyses and identifying which one fails. The community’s function is to maintain the procedures, not to arbitrate the outputs. A community that becomes the arbiter of correct PMN usage has reproduced technocratic capture at the level of the interpretive community. The cross-context transmission problem is addressed by the universality test in 12.3b: an analysis that cannot be specified in terms that apply across cultural contexts without requiring context-specific variable selection has not achieved framework-level analysis — it has produced doctrine that must be labeled as such. Evolution is governed by the anti-closure triggers in 14.4b, which specify the conditions under which revision at the paradigm level is required. Below that threshold, evolution proceeds through the mechanism already described in section 14.4: identifying specific failures, specifying what would address them, and demonstrating that the proposed revision actually addresses th

e failure rather than elaborating the existing framework in the failure’s vicinity. The distinction between elaboration and revision is operationalized by the revision trigger test (12.3b): a genuine revision produces a different conclusion in at least one case that was previously analyzed using the framework. An elaboration

that produces no divergent conclusions is not a revision — it is an expansion of the framework's vocabulary without a change in its analytical outputs. This distinction prevents the framework from accumulating complexity while its conclusions remain unchanged, which is the self-referential closure mechanism 14.4b is designed to detect and interrupt.

14.6 — Self-Capture Risk Assessment for the PMN Community

Section 7.3c-i (with operational indicators in §7.3b) formalizes the five-stage capture sequence as a general structural phenomenon: access asymmetry, decision-filter capture and preference expression, personnel alignment, objective redefinition and output reorientation, and accountability capture with consolidation. Section 14.5 develops the conditions for healthy PMN transmission. What neither section addresses is the specific application of the capture sequence to any organized community of practice that would form around PMN if the framework gained significant institutional traction. This is not a hypothetical risk. Every analytical framework that achieved significant reach generated a community of practitioners with institutional interests in that community's continuation — interests that can diverge from the framework's analytical purpose in predictable ways. The framework's epistemological commitments require applying the capture diagnostic to itself with the same rigor applied to any other institutional arrangement. The five-stage sequence applied to a hypothetical PMN community: Stage 1 — Preferential access: the actors with the most resources to invest in learning PMN and building relationships with its most fluent practitioners are not necessarily the populations whose structural suffering the framework is designed to analyze. Academic institutions, foundations, state agencies, and established progressive organizations shape what questions practitioners are incentivized to ask and which communities they engage. At this stage, no individual actor consciously distorts the framework. The distortion is structural: the access gradient produces an information environment shaped by those with resources to engage. Stage 2 — Preference alignment: practitioners whose work is most legible to institutional funders — whose applications of PMN are compatible with existing programmatic frameworks, whose language is translatable into grant proposals — advance within the community. Practitioners whose most important translations cannot be documented or published are less visible. The community's internal sense of what constitutes quality PMN work is gradually shaped by what institutional infrastructure rewards. Stage 3 — Criteria redefinition: the measure of a good PMN analysis shifts from its material consequences for the populations it analyzes toward technical sophistication, citability, and fit with recognized scholarly frameworks. The technocratic drift safeguard (12.8) was designed to catch this shift, but the safeguard itself is subject to the same capture dynamic: if the community controls the application of 12.8, it controls the definition of adequate translation. Stage 4 — Objective redefinition and output reorientation: PMN vocabulary is deployed to legitimate conclusions of actors whose institutional position the framework would, if applied honestly, identify as a capture vector. The language of structural suffering, institutional betrayal,

and the minimal anchor is used to produce analysis that protects the community's institutional relationships rather than challenging arrangements that maintain those relationships. Stage 5 — Accountability capture and consolidation: review and critique mechanisms within the PMN community are neutralized; dissent is marginalized as 'misunderstanding the framework'; and PMN institutions deploy their legitimacy to defend the captured consensus against external critique. Anti-capture structural requirements for any PMN community of practice: First, no certification or credentialing architecture that concentrates interpretive authority. Any system that designates some practitioners as authorized to determine what counts as legitimate PMN application is a mechanism for Stage 3 capture regardless of the intentions of its designers. Second, mandatory open access to all foundational materials. If access to PMN training is mediated by institutional gatekeepers, the access gradient of Stage 1 is built into the community's structure from the beginning. Third, standing external critic role, independently funded, to prevent Stage 5 accountability capture. Any institutionalized PMN community must maintain a designated function — with genuine

resources and genuine independence from the community's primary funders — whose explicit purpose is to apply the capture sequence to the community itself and publish findings publicly. This role cannot be rotated among community insiders. Fourth, material consequence test as primary evaluation criterion against Stage 4 output reorientation: the community's assessment of its own work must be anchored in what changes in the conditions of the populations analyzed, not in publication records or recognition within peer networks. This section does not prevent the formation of a PMN community of practice. It specifies the structural conditions under which such a community would not betray the framework it claims to embody. A community that institutionalizes PMN without these safeguards has not failed to prevent capture — it has created the conditions for it. The framework's commitment to applying its own diagnostics to itself (1.2, 6.4, 12.7) requires that this assessment be conducted before institutionalization, not after capture has occurred.

Coda: Coda

coda — Coda

A framework is not a destination. It is a position from which further movement becomes possible — and whose value is measured by what it enables rather than by its own continuation. This document has been built from a specific displacement: beginning from what organisms materially are rather than from what traditions hold, ideals prescribe, or prior commitments require. That displacement does not yield certainty. It yields a more honest relationship with uncertainty — one that treats the gap between current understanding and adequate understanding not as a problem to be solved by sufficient rigor, but as the permanent condition of inquiry that takes its own limits seriously. Every claim in the preceding pages inherits that condition. None of them are offered as final. All of them are offered as the best available positions given current conditions, subject to the revision that better evidence and better analysis will require. What this framework asks of those who use it is not agreement but engagement: the willingness to apply its analytical tools to conditions it has not examined, to test its conclusions against evidence it has not yet encountered, and — when that evidence requires it — to revise the conclusions rather than protect them. The framework that is used only to confirm what was already believed has been converted into something this analysis identifies as the primary obstacle to understanding: an arrangement that serves its own continuation at the cost of the purpose that justified it. The framework that is surpassed by something more adequate has succeeded. There is no higher standard available. The criterion remains. Whether structural suffering is declining, whether the conditions for genuine development are expanding, whether the institutional arrangements that prevent becoming are becoming less entrenched — these questions can be asked in any period, by any analyst, without this document. The document is useful to the extent that it makes those questions sharper. It is finished when something makes them sharper still.

Intellectual Debts: Intellectual Debts

intellectual-debts — Intellectual Debts

The debts acknowledged here are real, and they are substantial. None of the thinkers named below would recognize everything in this framework, and several would object to significant portions of it. The synthesis is distinct from any of its sources. But the sources matter, and naming them is part of the intellectual honesty the framework requires. B.R. Ambedkar — for the demonstration that structural suffering can be produced and maintained through social institutions of status assignment that are not reducible to class or formal political arrangements; for the analysis of caste as a permanent structural exclusion that persists across economic change; for the strategic analysis of when institutional reform is structurally incapable of addressing the conditions it claims to address, and when replacement of the institutional framework is the only adequate response; and for the Buddhist conversion as a documented instance of strategic intervention in meaning infrastructure to escape institutional betrayal embedded in religious form. Thomas Aquinas (1225–1274) — for the natural law framework’s grounding of political obligation in what arrangements actually produce for human flourishing; for the *bonum commune* (common good) as an evaluative standard structurally parallel to PMN’s minimal anchor; for the mixed constitution analysis in *De Regno* as an early institutional response to the custodian problem; and for the just war criteria as the most developed pre-modern framework for evaluating organized political violence by its consequences rather than its authorization. Giovanni Arrighi — for the long-term view of hegemonic cycles and systemic transitions; for multipolarity as a historical pattern rather than an ideological preference; and for the systemic analysis of how economic and military power interact at the scale of the international system. Ali Shari’ati (1933–1977) — for the distinction between the Islam of Abu Dharr — the companion of the Prophet who challenged the illegitimate accumulation of wealth by the early caliphate — and the Islam of official religious authority; for demonstrating that the same tradition can be occupied by arrangements that justify existing inequality and by critique of that inequality on the tradition’s own terms, which is the structure of internal reform that the framework’s section on methodological legitimacy identifies as distinct from methodological appropriation. Pierre Bourdieu (1930–2002) — for the concept of *habitus* as the mechanism through which structural positions become embodied — the durably installed dispositions that make structural reproduction feel like individual choice and structural constraint feel like personal limitation — filling a gap in PMN’s account of how material conditions reproduce themselves through the behavior of people who experience themselves as acting freely; for the field analysis as a complement to PMN’s institutional analysis: fields as structured spaces of position-taking with

their own logics of capital conversion, where what counts as legitimate action is itself a product of the field’s history rather than a neutral standard; for symbolic violence as the specific mechanism through which V suppression operates without requiring conscious intent — the imposition of categories of perception that make structural conditions appear natural to those who bear their costs; and for the sociology of knowledge production as a field-level analysis of how epistemic authority (1.9) is reproduced through the structural positions of those credentialed to produce it. PMN’s account of how structural suffering becomes individually attributed rather than structurally understood requires Bourdieu’s account of how that attribution gets installed in the cognitive and affective dispositions of those who bear the suffering. Ha-Joon Chang — for the insistence that appropriate economic policy depends on the material conditions of the specific context; for the historical documentation of the developmental state; and for the critique of premature institutional convergence as a form of policy imposition dressed as universal advice. Al-Farabi (872–950 CE) — for the most systematic synthesis of Greek political philosophy and Islamic thought available in the medieval period; for Al-Madina al-Fadila’s becoming-centered political philosophy that identifies the creation of conditions for genuine human flourishing as the primary question of political organization — parallel to PMN’s 4.5 but derived from different foundational commitments; for the typology of imperfect cities

(ignorant, immoral, erring) as an analytically precise account of institutional failure that maps onto PMN's regime type analysis; for the racialis theory as an attempt to resolve the analytical capacity/contextual judgment tension in governance that PMN's anti-technocratic guardrail addresses without fully resolving. Ibn Rushd (Averroes, 1126–1198) — for the defense of philosophical reasoning against theological authority as a historical precedent for PMN's own epistemological position: reason operating on material reality can address political and ethical questions without requiring religious grounding, while remaining genuinely open to interface-level phenomena. Frantz Fanon — for the analysis of colonial identity not as cultural but as material — as a condition imposed through force and maintained through the systematic allocation of humanity and its denial; for the demonstration that the psychological effects of colonial arrangements are themselves material variables with structural consequences; and for the insistence that the analysis of decolonization be adequate to what colonization actually was, rather than to a sanitized version that protects the comfort of those analyzing from a distance. Silvia Federici — for recovering the history of primitive accumulation from the perspective of those whose bodies and labor were accumulated; for the analysis of how the enclosure of the co

mmons and the disciplining of reproductive labor were not background conditions for capitalism but constitutive acts of its production; and for making visible the work that does not appear in the dominant economic accounting but without which the accounting would have nothing to count. Michel Foucault (1926–1984) — for the analysis of how power operates through knowledge and discourse as well as through overt force; for the genealogical method as a way of denaturalizing arrangements that present themselves as necessary or universal by recovering their historical conditions of emergence; for the analysis of disciplinary power as a form of governance that operates through the internalization of norms rather than through external surveillance alone; and for the concept of technologies of the self — the practices through which individuals actively constitute and maintain their relationship to their own judgment and conduct — which 17.7c applies to the specific problem of drift detection in analysts subject to institutional pressure. The debt here is selective and the departure is real. Foucault's genealogical skepticism about normative foundations — his resistance to specifying what power should be replaced by — is a position PMN's minimal anchor directly rejects: the analysis of power without an evaluative criterion produces description that cannot distinguish better from worse arrangements, and PMN holds that distinction to be both necessary and achievable. The debt is to the analytical tools; the framework does not inherit the normative agnosticism. Nancy Fraser — for the analysis of the structural contradiction between capitalist accumulation and the conditions of social reproduction; for identifying the tension between the economic system's dependence on care work and its systematic refusal to value it; and for the demonstration that feminist political economy and structural economic analysis are not parallel inquiries but aspects of a single one. Paulo Freire — for the analysis of conscientização as the specific pedagogical mechanism through which visibility suppression (V) is reduced; for the identification of the 'banking model of education' as a primary mechanism of V maintenance through legitimate institutional form; for demonstrating that the reduction of structural invisibility requires dialogue grounded in the material experience of those who suffer it rather than the imposition of analysis generated from outside; and for the empirical grounding of what PMN's 1.12 analyzes structurally. Al-Ghazali (1058–1111) — for the Maqasid al-Shariah framework's systematization of Islamic jurisprudence around five protected goods that constitute a floor/becoming evaluative structure independently derived from non-materialist foundations; for the daruriyyat/hajiyyat/tahsiniyyat hierarchy that distinguishes necessary conditions from sufficiency conditions from excellence conditions with a structural precision that PMN's minimal anchor/becoming horizon

distinction approximates; and for the ordering of floor conditions under conflict. Antonio Gramsci — for the concept of hegemony as a material variable rather than a purely ideological one; for the war of position as the appropriate strategic form in complex modern societies; for the organic intellectual as a role defined by material engagement rather than credential; and for demonstrating, under conditions of genuine political constraint, what honest intellectual work under pressure can look like. The Prison Notebooks are, among

other things, a document of what happens when a rigorous thinker is forced to work through problems from within the tradition that shaped them — finding where the tradition is right, where it is wrong, and where it needs to be extended. Gustavo Gutiérrez — for the liberation theology tradition’s development of the ‘preferential option for the poor’ as an evaluative criterion grounded in theological commitment but structurally equivalent to PMN’s minimal anchor; for the analysis of ‘social sin’ — the claim that institutional arrangements rather than only individuals can be the bearers of moral failure — which is the theological translation of PMN’s structural suffering concept; and for the demonstration that meaning infrastructure can be mobilized for structural transformation rather than for its legitimization, through engagement with the tradition’s own foundational commitments. Mohammad Hatta (1902–1980) — for the concept of gotong royong as an economic principle grounded in indigenous Indonesian cooperative practice, distinct from Western cooperative theory in its communal rather than contractual foundation; for the Pancasila framework as an attempt to ground post-colonial political legitimacy in the material conditions and cultural inheritance of a specific society rather than importing European institutional forms, demonstrating what parallel development (9.4) requires in practice. Georg Wilhelm Friedrich Hegel — for the insight that adequate analysis must account for movement and contradiction rather than describing a frozen state; for the concept of determinate negation as the appropriate form of immanent critique; and for the ambition of a framework that would be generative — that would make possible forms of analysis that could not have been produced from any other starting point. The teleological structure of the Hegelian system is rejected here precisely because following the logic of adequate analysis further than Hegel did requires it. The ambition is not rejected. It is the inheritance most worth keeping. Augustine of Hippo (354–430 CE) — for the political realism of the Two Cities doctrine, which anticipates PMN’s analysis of the gap between institutional claims and institutional function; for the concept of libido dominandi as the structural drive behind institutional capture, independently identifying what PMN’s custodian problem analysis formulates from a different starting point; and for the recognition

that political arrangements cannot achieve the good they claim without structural constraints on the self-interest of those who operate them. Bernardo Kastrup (1973–) — for the most technically rigorous contemporary development of analytic idealism: the position that consciousness is the ontological substrate of what appears as the physical world, and that individual minds are dissociated alters of a universal field of consciousness. PMN does not adopt this position; its agnosticism about what underlies the material level (2.3) is genuine. What Kastrup’s work contributes is the sharpest available formulation of the ontological challenge that materialist frameworks must be able to survive. The postulate in 2.1 — that PMN’s practical and ethical scope is defined by the phenomenological reality of constraint rather than by the metaphysical nature of its substrate — was formulated specifically in response to Kastrup’s position. A framework that cannot answer the strongest version of the idealist objection has not established its own foundations adequately. PMN’s answer is not that Kastrup is wrong; it is that even if he is right, the analysis of structural suffering and the ethical obligation to reduce it remain unchanged. (See 2.1.) Kautilya (c. 350–275 BCE) — for the Arthashastra’s structural analysis of statecraft as the primary pre-modern treatise on institutional design, predating Machiavelli by eighteen centuries; for the mandala theory of interstate relations as a structural realism derived from South Asian conditions independently of the European state system; and for the analysis of the relationship between material power (danda) and legitimacy (dharma) as structural conditions for stable governance, mapping directly onto PMN’s legitimacy analysis. Ibn Khaldun (1332–1406) — for one of the earliest and most systematically materialist analyses of social and political change available in any tradition; for the concept of asabiyyah (social cohesion as a material variable that accumulates and depletes over time, determining institutional capacity and civilizational trajectory) which PMN applies directly in the transformation pressure formula; for the economic analysis of the Muqaddimah — which identified the role of labor in production, the destructive effects of excessive taxation on productive activity, and the cyclical dynamics between agricultural, commercial, and imperial economies — seven centuries before contemporary institutional economics reached similar conclusions; for the epistemological

discipline of distinguishing transmitted knowledge from knowledge derived from observation and reason, which anticipates PMN's own commitment to evidence-responsiveness over received authority; and for demonstrating that the materialist analysis of social change is not a European invention but a structural methodology that independent intellectual traditions arrived at from within their own conditions. The Muqaddimah's theory of dynastic cycles — the

accumulation and depletion of *asabiyyah* as the primary driver of political transformation, operating through behavioral mechanisms that are structurally identical to PMN's legitimacy depletion analysis — is perhaps the most important pre-modern precedent for the transformation pressure framework in Part XV. Rosa Luxemburg (1871–1919) — for the Reform or Revolution argument as the most analytically precise formulation of the sequencing problem that 10.6 addresses: the claim that the choice between reform and revolution is not a moral preference but a structural question about whether the arrangements being contested can be transformed from within or require displacement; for the mass strike analysis as an account of how spontaneous collective action generates organizational capacity that precedes the organizations that subsequently claim to direct it — directly relevant to 10.10c-iii's account of spontaneous actors; for the critique of Leninist democratic centralism on organizational grounds — that concentrating interpretive authority in a leadership stratum reproduces the structural feature of the arrangement it claims to oppose — which 10.8 independently arrives at; and for demonstrating, through her political practice and its end, that the organizational form problem is not merely theoretical. Her assassination at the hands of forces authorized by a socialist government she had criticized is evidence for her own analysis: the institutional logic of self-preservation operates independently of the ideological commitments of the actors it structures. Niccolò Machiavelli (1469–1527) — for the analysis of power as a structural condition that requires honest description rather than moral sanitization; for the *fortuna/virtù* distinction as an early formulation of the relationship between structural conditions and the agency available within them — which PMN's probabilistic determinism (1.3) and its account of conjunctural windows (10.5b) engage without acknowledgment whenever they treat structure as constraining rather than determining; for the recognition that institutional founders and reformers operate in conditions that make the gap between stated purposes and actual requirements of effective governance a permanent practical problem, not a failure of character; and for the analysis of how new institutions establish themselves and how existing ones resist replacement — the problem that 10.8 and 7.3 address from within a different evaluative framework. The PMN evaluative criterion and Machiavelli's refusal to operate with one are incompatible. His analytical contribution — that power must be described as it operates rather than as it legitimizes itself — is the same move that Part VI makes when it insists that power cannot be left implicit. José Carlos Mariátegui — for the analysis of indigenous communal arrangements (the Andean *ayllu*) as pre-existing cooperative institutions that capitalism disrupted but did not fully eliminate, and therefore as mat

erial starting points for alternative institutional development rather than blank slates requiring external models; for the demonstration that Marxist analysis requires adaptation to the specific material conditions of each context rather than mechanical application of a universal template — anticipating PMN's own distinction between framework and doctrine; and for identifying the 'Indian question' as fundamentally an economic and structural question rather than a racial or cultural one. Karl Marx — for the displacement of analysis from the realm of ideas to the realm of material conditions, which remains the most consequential methodological move in the history of social and political thought; for the analysis of how economic power translates into political power and ideological legitimation; for the recognition that individually rational behavior produces collectively irrational outcomes in ways that are structural rather than contingent; and for the distinction between the analytical framework and the political doctrine derived from it, which Capital and the Manifesto embodied in the most instructive way available. The teleological structure of historical materialism is rejected here for the same reason the Hegelian dialectic is rejected: following the materialist logic further requires it. The analysis is not rejected. It is the foundation on which this framework builds. Mencius (Mengzi, 372–289 BCE) — for the moral psychology that grounds political authority in the ruler's

capacity for compassion and the recognition of suffering as politically consequential; for the argument that governance which produces suffering forfeits its mandate, providing an independent non-Western derivation of floor-criterion legitimacy theory that the framework engages in the Confucian political philosophy section (7.8). Elinor Ostrom — for the empirical demonstration that communities with adequate social capital, clear boundaries, and effective monitoring can sustain coordination without hierarchical enforcement — and that where these conditions obtain, the results outperform formal institutional alternatives on compliance, overhead, and accountability. The Ostrom research establishes that the custodian problem (7.3) is not resolved by removing formal hierarchy but by the specific institutional conditions that make informal coordination stable and self-correcting. Those conditions can themselves be analyzed, designed for, and captured — which is what 7.2 and 7.3 develop. Charles Sanders Peirce — for pragmatism as a method for approaching truth rather than a license for relativism; for the conception of truth as what inquiry would converge on given indefinite continuation; and for the community of inquiry as the appropriate social form for the production of knowledge that is more than the opinion of whoever is currently in a position to enforce it. Section 12.8 operationalizes this constraint structurally: the distributed verification requirement and translation o

bligation are the community-of-inquiry principle applied as a design safeguard against the concentration of interpretive authority that PMN's own analytical complexity would otherwise produce. Karl Polanyi — for the critique of market fundamentalism grounded in historical analysis; for the demonstration that economic arrangements that appear natural are in fact historically specific and politically constructed; and for the concept of the double movement as a way of understanding how societies respond to the destruction of their protective institutions. Walter Rodney — for the historical analysis of how European development and African underdevelopment were not parallel processes but aspects of a single process; for the demonstration that the material starting positions from which development paths proceed are themselves products of structured extraction; and for insisting that the history of the world system be told from the perspective of those it enriched and those it impoverished simultaneously, without letting the beneficiaries narrate the process as neutral. Dani Rodrik — for the trilemma of globalization; for context-dependent policy evaluation; and for the intellectual honesty of acknowledging the limits of what economic theory can establish independently of context. James C. Scott — for the analysis of how subordinate groups navigate domination through forms of resistance that are invisible to the public transcript; for the concept of the hidden transcript as a structural resource that persists through arrangements that appear to have achieved full legitimacy; for the analysis of how high-modernist schemes fail when they override the local, practical knowledge that formal systems cannot capture; and for the land reform cases that demonstrate the gap between the formal intention of institutional change and the material conditions required to make that change real. The Weapons of the Weak and Seeing Like a State together constitute one of the most structurally rigorous accounts available of how power is actually navigated at the level where most people most often encounter it. Baruch Spinoza — for naturalism that does not prematurely close metaphysical questions; for the refusal to require that a natural account of reality must be an atheistic one; and for the standard of rigor in the construction of philosophical argument. E.P. Thompson — for the history of the English working class as a history of agency within structural constraint; for the concept of moral economy as a framework through which populations evaluate whether existing arrangements meet their expectations of legitimacy; and for demonstrating that the making of class consciousness is a cultural and organizational process, not an automatic consequence of structural position. The labour movement case in Part XVII draws directly on the structural dynamics Thompson documented: the transformation of movement organizations through the process of formalization and recognition. Zhao T

ingyang (b. 1961) — for the contemporary reconstruction of the Confucian Tianxia ('All under Heaven') concept as a framework for world institutional order organized around relational obligations rather than sovereign equality; for raising the question of whether the Westphalian state system is a contingent institutional arrangement; and for providing the most developed contemporary non-Western alternative to

PMN's own long-term horizon of post-national coordination. Immanuel Wallerstein — for the world-systems analysis that makes the international division of labor the unit of analysis rather than the nation-state; for the concepts of core, periphery, and semi-periphery as structural positions rather than stages of development; and for the demonstration that development and underdevelopment are relational — that they are produced by the same system, not by different rates of progress along the same path. Kenneth Waltz — for structural realism as a description of how the international system currently operates, understood as description and not normative endorsement; and for the insistence that outcomes in international relations are significantly explained by system structure rather than only by the intentions of individual actors. Kwame Wiredu — for the Ubuntu philosophical tradition's analysis of personhood as constituted through relationships rather than prior to them; for the demonstration that communitarian ontology is compatible with attention to individual suffering rather than opposed to it; and for providing a non-Western independent derivation of the claim that genuine individuation requires a social substrate rather than occurring prior to it — which is the claim PMN's biological foundation implies but does not articulate in these terms. Erik Olin Wright — for the concept of real utopias — actually existing institutional alternatives that point toward a horizon without requiring conditions that do not exist; and for the investigation of how structural change begins from material conditions rather than from the moral conversion of actors. Slavoj Žižek (1949–) — for the concept of fetishistic disavowal — “I know very well, but nevertheless...” — as a documented structural feature of ideological commitment: the condition in which accurate information about structural causation is cognitively available to an actor who nonetheless continues to act as if it were not. This is the phenomenological core of what 3.8c calls stable irrationality, and it distinguishes that form of irrationality from information deficit (addressable by better information) and from affective capture (addressable by reorientation of emotion). Fetishistic disavowal is addressable by neither, because the disavowal is itself structured by something the actor is getting from their ideological position — a form of investment that accurate information threatens rather than corrects. The Lacanian framework Žižek uses to explain why this pattern is so stable is not adopted here

re; PMN's account of stable irrationality attributes it to structural incentives and social reproduction rather than to unconscious drive structure. But the phenomenological observation — that people can know and nevertheless act otherwise, systematically and durably — is well-documented independently of its explanation, and PMN's analysis of why transformation pressure fails to materialize despite high structural suffering requires it. The departure that matters most across all of these debts is the biological foundation. Beginning from physiology changes what the social and political analysis is ultimately grounded in. None of these traditions begins there, and beginning there changes the structure of what follows in ways that are not yet fully worked out. That is the work this document begins.

Bibliography: Bibliography

bibliography — Bibliography

Althusser, Louis. 1970. *Lenin and Philosophy and Other Essays*. Trans. Ben Brewster. Monthly Review Press. Ambedkar, B. R. 1936. *Annihilation of Caste*. Self-published. [Repr. Verso Books, 2014] Ambedkar, B. R. 1945. *What Congress and Gandhi Have Done to the Untouchables*. Thacker and Co. Ambedkar, B. R. 1956. *The Buddha and His Dhamma*. Bombay: Siddharth Publications. [Posthumous ed. 1957] Ames, Roger T., and Henry Rosemont Jr. 1998. *The Analects of Confucius: A Philosophical Translation*. Ballantine Books. Amsden, Alice H. 1989. *Asia's Next Giant: South Korea and Late Industrialization*. Oxford University Press. Anderson, Benedict. 1983. *Imagined Communities: Reflections on the Origin and Spread of Nationalism*. Verso. Aquinas, Thomas. c. 1267. *De Regno (On Kingship)*. Trans. Gerald Phelan. Pontifical Institute of Mediaeval Studies, 1949. Aquinas, Thomas. c. 1270. *Summa Theologiae*. Trans. Fathers of the English Dominican Province. Benziger Bros, 1947. Arrighi, Giovanni. 1994. *The Long Twentieth Century*. Verso. Arrighi, Giovanni. 2007. *Adam Smith in Beijing*. Verso. Augustine of Hippo. 426 CE. *The City of God*. Trans. Henry Bettenson. Penguin Books, 1972. Axelrod, Robert. 1984. *The Evolution of Cooperation*. Basic Books. Bentham, Jeremy. 1780. *An Introduction to the Principles of Morals and Legislation*. Clarendon Press. Berger, Peter L., and Thomas Luckmann. 1966. *The Social Construction of Reality*. Doubleday. Bhaskar, Roy. 1975. *A Realist Theory of Science*. Leeds Books. Boff, Leonardo. 1985. *Church: Charism and Power*. Trans. John Diercksmeier. Crossroad. Braudel, Fernand. 1979. *The Perspective of the World*. Trans. S. Reynolds. Harper & Row, 1984. Buchanan, James M., and Gordon Tullock. 1962. *The Calculus of Consent*. University of Michigan Press. Camus, Albert. 1942. *Le Mythe de Sisyphe [The Myth of Sisyphus]*. Paris: Gallimard. [Trans. Justin O'Brien, London: Hamish Hamilton, 1955] Chalmers, David J. 1996. *The Conscious Mind*. Oxford University Press. Chang, Ha-Joon. 2002. *Kicking Away the Ladder*. Anthem Press. Chang, Ha-Joon. 2007. *Bad Samaritans*. Bloomsbury. Chicago Author-Date format. Titles in italics. Cohen, Joel E. 1995. *How Many People Can the Earth Support?* Norton. Darwin, Charles. 1859. *On the Origin of Species*. John Murray. Dawkins, Richard. 1976. *The Selfish Gene*. Oxford University Press. De Benoist, Alain. 1977. *Vu de droite*. Copernic. Diamond, Jared. 1997. *Guns, Germs, and Steel*. Norton. Dunbar, Robin I. M. 1992. 'Neocortex size as a constraint on group size in primates.' *Journal of Human Evolution* 22 (6): 469-493. Dunbar, Robin I. M. 1998. 'The social brain hypothesis.' *Evolutionary Anthropology* 6 (5): 178-190. Durkheim, Émile. 1912. *The Elementary Forms of Religious Life*. Trans. Karen Fields. Free Press, 1995. Eisenstein, Elizabeth L. 1980. *The Printing Press as an Agent of Change*. Cambridge University Press. Evans, Peter. 1995. *Embedded Autonomy: States and Industrial Transformation*. Princeton University Press. Fanon, Frantz. 1961. *The Wretc*

hed of the Earth. Trans. Richard Philcox. Grove Press, 2004. Al-Farabi. c. 940. *Al-Madina al-Fadila (The Virtuous City)*. Trans. Richard Walzer. Oxford: Clarendon Press, 1985. Federici, Silvia. 2004. *Caliban and the Witch: Women, the Body and Primitive Accumulation*. Autonomedia. Foucault, Michel. 1975. *Discipline and Punish: The Birth of the Prison*. Trans. Alan Sheridan. Pantheon Books, 1977. Foucault, Michel. 1988. 'Technologies of the Self.' In *Technologies of the Self: A Seminar with Michel Foucault*, ed. Luther H. Martin, Huck Gutman, and Patrick H. Hutton. University of Massachusetts Press. Fraser, Nancy. 2016. *Contradictions of Capital and Care*. *New Left Review* 100: 99-117. Freire, Paulo. 1968. *Pedagogy of the Oppressed*. Trans. Myra Bergman Ramos. Herder and Herder, 1970. Friedman, Milton. 1962. *Capitalism and Freedom*. University of Chicago Press. Geertz, Clifford. 1973. *The Interpretation of Cultures*. Basic Books. Al-Ghazali. c. 1095. *The Incoherence of the Philosophers*. Trans. Michael Marmura. Brigham Young University Press, 1997. Goldstone, Jack A. 1991. *Revolution and Rebellion in the Early Modern World*. University of California Press. Gramsci, Antonio. 1971. *Selections from the Prison Notebooks*. Ed. and trans. Q. Hoare and G. N. Smith. International Publishers. Gutiérrez, Gustavo. 1971. *A Theology of Liberation*. Trans. Caridad Inda and John Eagleson. Orbis Books, 1988. Harvey, David. 2005. *A Brief History of*

Neoliberalism. Oxford University Press. Hatta, Mohammad. 1946. *The Cooperative Movement in Indonesia*. Cornell University Press, 1957. Hayek, Friedrich A. 1944. *The Road to Serfdom*. Routledge. Hayek, Friedrich A. 1960. *The Constitution of Liberty*. University of Chicago Press. Hegel, G. W. F. 1807. *Phenomenology of Spirit*. Trans. A. V. Miller. Oxford University Press, 1977. Held, David. 1995. *Democracy and the Global Order: From the Modern State to Cosmopolitan Governance*. Polity Press. Herman, Edward S., and Noam Chomsky. 1988. *Manufacturing Consent: The Political Economy of the Mass Media*. Pantheon Books. Honneth, Axel. 1995. *The Struggle for Recognition: The Moral Grammar of Social Conflicts*. Trans. Joel Anderson. Cambridge: Polity Press. Hume, David. 1739. *A Treatise of Human Nature*. Ed. L. A. Selby-Bigge. Oxford University Press, 1978. Izz ibn Abd al-Salam. c. 1250. *Qawā'id al-Akām fī Ma'āli al-Anām*. [Rules of Judgment Concerning the Interests of Mankind]. Various trans. John Paul II. 1987. *Sollicitudo Rei Socialis*. Vatican City: Libreria Editrice Vaticana. Kahneman, Daniel. 2011. *Thinking, Fast and Slow*. Farrar, Straus and Giroux. Kastrup, Bernardo. 2019. *The Idea of the World: A Multi-Disciplinary Argument for the Mental Nature of Reality*. iff Books. Kautilya. c. 300 BCE. *The Arthashastra*. Trans. L.N. Rangarajan. Penguin Books, 1992. Ibn Khaldun. 1377. *The Muqaddimah: An Introduction to History*. Trans. Franz Rosenthal. 3 vols. Princeton University Press, 1958. Kierkegaard, Søren. 1849. *Sygdommen til Døden* [The Sickness unto De

ath]. Copenhagen: C.A. Reitzel. [Trans. Howard V. Hong and Edna H. Hong, Princeton: Princeton University Press, 1980] Klare, Michael T. 2001. *Resource Wars*. Metropolitan Books. Lenin, V. I. 1917. *The State and Revolution*. International Publishers, 1932. Leo XIII. 1891. *Rerum Novarum*. Vatican City: Libreria Editrice Vaticana. Lukes, Steven. 1974. *Power: A Radical View*. Macmillan. Luthfi Assyaukanie. 2009. *Islam and the Secular State in Indonesia*. Institute of Southeast Asian Studies. Malthus, Thomas Robert. 1798. *An Essay on the Principle of Population*. J. Johnson. Machiavelli, Niccolò. 1513. *The Prince*. Trans. Harvey C. Mansfield. University of Chicago Press, 1998. Mariátegui, José Carlos. 1928. *Seven Interpretive Essays on Peruvian Reality*. Trans. Marjory Urquidi. University of Texas Press, 1971. Marx, Karl, and Friedrich Engels. 1848. *The Communist Manifesto*. Penguin Books, 2002. Marx, Karl. 1844. *Economic and Philosophic Manuscripts of 1844*. Trans. Martin Milligan. International Publishers, 1964. Marx, Karl. 1845. 'Theses on Feuerbach.' In *Selected Works*, Vol. 1. Progress Publishers, 1969. Marx, Karl. 1859. *A Contribution to the Critique of Political Economy*. International Publishers, 1970. Marx, Karl. 1867. *Capital: A Critique of Political Economy*, Vol. 1. Trans. Ben Fowkes. Penguin Books, 1976. Mencius. c. 300 BCE. Mencius. Trans. D.C. Lau. Penguin Books, 2004. Minsky, Hyman P. 1986. *Stabilizing an Unstable Economy*. Yale University Press. Montesquieu, Charles-Louis de Secondat. 1748. *De l'Esprit des Lois* [The Spirit of the Laws]. Trans. Anne Cohler, Basia Miller, Harold Stone. Cambridge University Press, 1989. Nagel, Thomas. 1974. 'What is it like to be a bat?' *Philosophical Review* 83 (4): 435-450. Nagel, Thomas. 2012. *Mind and Cosmos*. Oxford University Press. Nietzsche, Friedrich. 1883–1885. Also sprach Zarathustra: Ein Buch für Alle und Keinen [Thus Spoke Zarathustra]. Chemnitz: Ernst Schmeitzner. [Trans. Walter Kaufmann, New York: Viking Press, 1954] North, Douglass C. 1990. *Institutions, Institutional Change and Economic Performance*. Cambridge University Press. North, Douglass C., John J. Wallis, and Barry R. Weingast. 2009. *Violence and Social Orders*. Cambridge University Press. Nussbaum, Martha C. 2011. *Creating Capabilities: The Human Development Approach*. Harvard University Press. Olson, Mancur. 1965. *The Logic of Collective Action*. Harvard University Press. Ostrom, Elinor. 1990. *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge University Press. Peirce, Charles Sanders. 1878. 'How to make our ideas clear.' *Popular Science Monthly* 12: 286-302. Perez, Carlota. 2002. *Technological Revolutions and Financial Capital*. Edward Elgar. Pierson, Paul. 2004. *Politics in Time*. Princeton University Press. Polanyi, Karl. 1944. *The Great Transformation*. Farrar & Rinehart. Popper, Karl R. 1959. *The Logic of Scientific Discovery*. Hutchinson. Prigogine, Ilya, and Isabelle Stengers. 1984. *Order Out of Chaos*. Bantam B

ooks. Rawls, John. 1971. *A Theory of Justice*. Harvard University Press. Reinert, Erik S. 2007. *How Rich Countries Got Rich*. Carroll & Graf. Robison, Richard, and Vedi Hadiz. 2004. *Reorganising Power in*

Indonesia: The Politics of Oligarchy in an Age of Markets. RoutledgeCurzon. Rockstrom, Johan, et al. 2009. 'A safe operating space for humanity.' *Nature* 461: 472-475. Rodney, Walter. 1972. *How Europe Underdeveloped Africa*. Bogle-L'Ouverture Publications. Rodrik, Dani. 2007. *One Economics, Many Recipes*. Princeton University Press. Rodrik, Dani. 2011. *The Globalization Paradox*. Norton. Ibn Rushd (Averroes). 1180. *Tahafut al-Tahafut (The Incoherence of the Incoherence)*. Trans. Simon Van Den Bergh. Luzac & Co., 1954. Sartre, Jean-Paul. 1943. *Being and Nothingness: An Essay on Phenomenological Ontology*. Trans. Hazel E. Barnes. Philosophical Library, 1956. Schelling, Thomas C. 1978. *Micromotives and Macrobehavior*. Norton. Schmitt, Carl. 1932. *The Concept of the Political*. Trans. George Schwab. University of Chicago Press, 1996. Schumpeter, Joseph A. 1942. *Capitalism, Socialism and Democracy*. Harper & Brothers. Scott, James C. 1990. *Domination and the Arts of Resistance: Hidden Transcripts*. Yale University Press. Searle, John R. 1964. 'How to Derive \u201cOught\u201d from \u201cIs.\u201d' *Philosophical Review* 73 (1): 43\u201358. Sen, Amartya. 1999. *Development as Freedom*. Anchor Books. Singer, Peter. 1975. *Animal Liberation*. New York Review Books. Sobrino, Jon. 1978. *Christology at the Crossroads*. Trans. John Drury. Orbis Books. Spinoza, Baruch. 1677. *Ethics*. Trans. Edwin Curley. Princeton University Press, 1994. Stigler, George J. 1971. 'The theory of economic regulation.' *Bell Journal of Economics and Management Science* 2 (1): 3-21. Taylor, Charles. 2007. *A Secular Age*. Harvard University Press. Thompson, E. P. 1963. *The Making of the English Working Class*. Victor Gollancz. Tomasello, Michael. 2019. *Becoming Human: A Theory of Ontogeny*. Harvard University Press. Turchin, Peter, and Sergey A. Nefedov. 2009. *Secular Cycles*. Princeton University Press. Tversky, Amos, and Daniel Kahneman. 1974. 'Judgment under uncertainty: Heuristics and biases.' *Science* 185: 1124-1131. Wallerstein, Immanuel. 1974. *The Modern World-System I*. Academic Press. Waltz, Kenneth N. 1979. *Theory of International Politics*. McGraw-Hill. Weber, Max. 1905. *Die protestantische Ethik und der Geist des Kapitalismus [The Protestant Ethic and the Spirit of Capitalism]*. Trans. Talcott Parsons. London: Allen & Unwin, 1930. Weber, Max. 1919. 'Politics as a Vocation.' In *From Max Weber*, ed. H. H. Gerth and C. W. Mills. Oxford University Press, 1946. Wilson, Edward O. 1978. *On Human Nature*. Harvard University Press. Wiredu, Kwame, ed. 2004. *A Companion to African Philosophy*. Blackwell. Wiredu, Kwame. 1980. *Philosophy and an African Culture*. Cambridge University Press. Wright, Erik Olin. 2010. *Envisioning Real Utopias*. Verso. Yergin, Daniel. 1991. *The Prize: The Epic Quest for Oil, Money and Power*

r. Simon & Schuster. Zuboff, Shoshana. 2019. *The Age of Surveillance Capitalism: The Fight for a Human Future at the New Frontier of Power*. PublicAffairs. Zhao Tingyang. 2005. *Tianxia Tixi: Shijie Zhidu Zhexue Daolun. [The Tianxia System: A Philosophy for the World Institution]*. Jiangsu jiaoyu chubanshe. [Partial trans. in various collections]